

Hypercap CC NLP Classifier

Table of contents

1 Workbook for MIMIC Hypercapnia Presenting Chief Concern Analysis	1
1.1 1) Runtime contract and imports	1
1.2 2) Config and constants	4
1.2.1 Core API (single-cell contract)	4
1.3 2B) Local helper functions (embedded core notebook logic)	48
1.4 3) Input load and schema checks	69
1.5 4) Text cleaning and segmentation functions	71
1.6 5) Rule system and ontology mappings	73
1.7 6) Prototype loading, deduplication, and embedding preparation	74
1.8 7) Segment cache build and classification	75
1.9 8) Output shaping and compatibility mapping	76
1.10 9) Diagnostics and QA checks	77
1.11 10) Export	83

1 Workbook for MIMIC Hypercapnia Presenting Chief Concern Analysis

This notebook is intentionally self-contained and deterministic. It performs NLP-based mapping from free-text ED chief complaints to NHAMCS top-level RVC categories, with strict compatibility outputs for downstream analysis notebooks.

1.1 1) Runtime contract and imports

- No runtime package installation.
- Fail fast if required dependencies are missing.
- Keep environment deterministic and explicit.

```

# Purpose: verify required packages, import shared dependencies, and
↳ initialize NLP runtime objects.
# Why this matters: later cells assume these imports/models exist, so we fail
↳ early if the environment is incomplete.

import importlib

REQUIRED_IMPORTS = {
    "numpy": "numpy",
    "pandas": "pandas",
    "torch": "torch",
    "ftfy": "ftfy",
    "spacy": "spacy",
    "sympelly": "sympelly",
    "sentence_transformers": "sentence_transformers",
    "openpyxl": "openpyxl",
    "dotenv": "python-dotenv",
}

missing = []
for import_name, package_name in REQUIRED_IMPORTS.items():
    if importlib.util.find_spec(import_name) is None:
        missing.append((import_name, package_name))

if missing:
    message = "Missing required dependencies (install via uv sync): " + ", "
    .join(
        f"{import_name} ({package_name})" for import_name, package_name in
        missing
    )
    raise ImportError(message)

import hashlib
import json
import logging
import math
import os
import re
import sys
import warnings
from collections import Counter, defaultdict
from dataclasses import dataclass
from datetime import datetime

```

```

from pathlib import Path
from typing import Any, Mapping, Sequence

import ftfy
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import spacy
import torch
from dotenv import load_dotenv
from symspellpy.symspellpy import SymSpell, Verbosity

# Silence widget-only tqdm warning in environments without Jupyter progress
# → widgets.
try:
    from tqdm import TqdmWarning
except Exception: # pragma: no cover - defensive for uncommon tqdm import
    # edge cases
    TqdmWarning = None

if TqdmWarning is not None:
    warnings.filterwarnings(
        "ignore",
        message="IPress not found.*",
        category=TqdmWarning,
    )

# Reduce HF Hub warning noise in notebook output. If HF_TOKEN is set,
# → authenticated requests still work.
try:
    from huggingface_hub import logging as hf_logging

    hf_logging.set_verbosity_error()
except Exception: # pragma: no cover - keeps notebook robust if hub
    # internals change
    pass

logging.getLogger("huggingface_hub").setLevel(logging.ERROR)
logging.getLogger("hf_xet").setLevel(logging.ERROR)

from sentence_transformers import SentenceTransformer, util

```

```

# spaCy model loading is config-driven in the setup cell so behavior is
↳ explicit.
NLP: Any | None = None
SPACY_MODEL_LABEL = "UNINITIALIZED (loaded in setup cell)"
print("spaCy pipeline: pending config-driven load")

```

spaCy pipeline: pending config-driven load

1.2 2) Config and constants

1.2.1 Core API (single-cell contract)

This cell defines the notebook's public in-notebook interfaces:

- ClassifierConfig
- get_cc_column
- normalize_cc
- build_prototype_resources
- build_segment_cache
- classify_segments_cached
- assign_rvc_per_segment_and_visit_cached
- validate_output_schema
- save_with_suffix

```

# Purpose: centralize the notebook's core API (config, constants, data
↳ structures, and reusable functions).
# Why this matters: keeping definitions in one place makes behavior
↳ predictable and easier for new contributors to trace.

@dataclass(frozen=True)
class ClassifierConfig:
    """Configuration for notebook-contained NLP classification."""

    seed: int = 1234
    work_dir_env_var: str = "WORK_DIR"
    data dirname: str = "MIMIC tabular data"
    input_filename: str = "MIMICIV all with CC.xlsx"
    input_filename_env_var: str = "CLASSIFIER_INPUT_FILENAME"
    output_suffix: str = "_with_NLP"
    output_sheet_name: str = "with_NLP"
    model_name: str = "NeuML/bioclinical-modernbert-base-embeddings"
    model_trust_remote_code: bool = True
    model_batch_size: int = 256
    segment_batch_size: int = 512
    scoring_method: str = "max" # {"max", "lse"}
    alpha: float = 12.0
    abstain_threshold: float = 0.40
    max_rfvs: int = 5
    max_segments_per_visit: int = 5

```

```

prototype_dedup_sim_thresh: float = 0.92
summary_weight_strategy: str = "uniform" # {"uniform", "by_rows"}
run_optional_plots: bool = False
require_spacy_model: bool = True
allow_uncodable_primary: bool = True
empty_primary_policy: str = "blank" # {"blank", "uncodable"}
appendix_relpah: str = "Annotation/nhamcs_rvc_2022_appendixII_codes.csv"
summary_relpah: str =
    "Annotation/nhamcs_rvc_2022_summary_by_top_level_17.csv"
hf_token_env_var: str = "HF_TOKEN"

# --- Ontology labels and precedence used by final visit-level outputs. ---
RVC_NAME: dict[str, str] = {
    "RVC-INJ": "Injuries & adverse effects",
    "RVC-SYM-RESP": "Symptom - Respiratory",
    "RVC-SYM-CIRC": "Symptom - Circulatory",
    "RVC-SYM-NERV": "Symptom - Nervous",
    "RVC-SYM-DIG": "Symptom - Digestive",
    "RVC-SYM-GU": "Symptom - Genitourinary",
    "RVC-SYM-MSK": "Symptom - Musculoskeletal",
    "RVC-SYM-SKIN": "Symptom - Skin/Hair/Nails",
    "RVC-SYM-EYE": "Symptom - Eye/Ear",
    "RVC-SYM-GEN": "Symptom - General",
    "RVC-SYM-PSY": "Symptom - Psychological",
    "RVC-DIS": "Diseases (patient-stated)",
    "RVC-TEST": "Abnormal test result",
    "RVC-DIAG": "Diagnostic/Screening/Preventive",
    "RVC-TREAT": "Treatment/Medication",
    "RVC-ADMIN": "Administrative",
    "RVC-UNCL": "Uncodable/Unknown",
}
PRECEDENCE_RVC: list[str] = [
    "RVC-INJ",
    "RVC-SYM-RESP",
    "RVC-SYM-CIRC",
    "RVC-SYM-NERV",
    "RVC-SYM-DIG",
    "RVC-SYM-GU",
    "RVC-SYM-MSK",
    "RVC-SYM-SKIN",
    "RVC-SYM-EYE",
]

```

```

    "RVC-SYM-GEN",
    "RVC-SYM-PSY",
    "RVC-DIS",
    "RVC-TEST",
    "RVC-DIAG",
    "RVC-TREAT",
    "RVC-ADMIN",
    "RVC-UNCL",
]
}

RVC_SHORT: dict[str, str] = {
    "RVC-INJ": "injury",
    "RVC-SYM-RESP": "resp",
    "RVC-SYM-CIRC": "circ",
    "RVC-SYM-NERV": "neuro",
    "RVC-SYM-DIG": "gi",
    "RVC-SYM-GU": "gu",
    "RVC-SYM-MSK": "msk",
    "RVC-SYM-SKIN": "skin",
    "RVC-SYM-EYE": "eye_ear",
    "RVC-SYM-GEN": "general",
    "RVC-SYM-PSY": "psych",
    "RVC-DIS": "disease",
    "RVC-TEST": "abn_test",
    "RVC-DIAG": "screen",
    "RVC-TREAT": "treatment",
    "RVC-ADMIN": "admin",
    "RVC-UNCL": "uncodable",
}
}

# --- Text normalization vocabulary and complaint parsing constants. ---
CC_CANDIDATES: tuple[str, ...] = (
    "ed_triage_cc",
    "chief_complaint",
    "ed_chief_complaint",
    "chiefcomplaint",
    "chief_complaint_text",
    "ed_cc",
    "cc",
)
)

PROTECT = ("n/v", "n/v/d", "s/p", "w/", "w/o", "h/o", "c/o", "h/a", "f/c",
    ↵ "c/p", "lab", "labs")

```

```

SPELL_PROTECT_TOKENS = {
    "lab",
    "labs",
    "abnl",
    "abnormal",
    "fall",
    "falls",
    "head",
    "foot",
    "feet",
    "left",
    "right",
    "cold",
    "copd",
    "cord",
    "hyperglycemia",
    "hyperkalemia",
    "hypoglycemia",
    "hypokalemia",
}
STRICT_SPELL_MIN_TOKEN_LEN = 5
STRICT_SPELL_MAX_EDIT_DISTANCE = 1
POLARITY_GUARD_PREFIX_PAIRS: tuple[tuple[str, str], ...] = (
    ("hyper", "hypo"),
    ("hypo", "hyper"),
)
LATERALITY_GUARD_TOKENS = {"left", "right"}
SPELL_CORRECTION_MODES = {"disabled", "strict", "auto"}
MISSING_RE = re.compile(
    r"^(?:n/?a|na|none|unknown|unk|no chief complaint|no complaint|not
     ↳ applicable|not available)$",
    re.I,
)
NEG_RE = re.compile(r"^(?:denies?|no|not|without|w/o)\b", re.I)

# Preserve a small set of clinically meaningful "absence" complaints instead
# ↳ of dropping as generic negation.
ABSENCE_SYMPTOM_RULES: tuple[tuple[re.Pattern[str], str], ...] = (
    (
        re.compile(r"^(?:no|without|w/o)\s+(?:urine|urinary)\s+(?:output|_
         ↳ void(?:ing)?)\b", re.I),
        "urinary retention",
    ),
)

```

```

(
    re.compile(r"^(?:no|without|w/o)\s+(?:bowel movement|bm|stool(?:\n"
    ↵ output)?)\b", re.I),
    "constipation",
),
)

SEG_OVR: dict[str, str] = {
    "n/v": "nausea vomiting",
    "n/v/d": "nausea vomiting diarrhea",
    "nvd": "nausea vomiting diarrhea",
    "w/": "with",
    "w/o": "without",
    "s/p": "status post",
    "h/o": "history of",
    "c/o": "complains of",
    "h/a": "headache",
    "f/c": "fever chills",
    "c/p": "chest pain",
    "mva": "motor vehicle accident",
    "mvc": "motor vehicle collision",
    "ped vs auto": "pedestrian struck by vehicle",
    "gsw": "gunshot wound",
    "sa": "sexual assault",
}

TOK_OVR: dict[str, str] = {
    "sob": "shortness of breath",
    "doe": "dyspnea on exertion",
    "cp": "chest pain",
    "palps": "palpitations",
    "ha": "headache",
    "h/a": "headache",
    "f/c": "fever chills",
    "c/p": "chest pain",
    "loc": "loss of consciousness",
    "ams": "altered mental status",
    "sz": "seizure",
    "szs": "seizure",
    "cva": "stroke",
    "tia": "transient ischemic attack",
    "abd": "abdominal",
    "abdo": "abdominal",
}

```

```

"rlq": "right lower quadrant",
"ruq": "right upper quadrant",
"llq": "left lower quadrant",
"luq": "left upper quadrant",
"lbp": "low back pain",
"brbpr": "rectal bleeding",
"uti": "urinary tract infection",
"dysuria": "painful urination",
"hematuria": "blood in urine",
"nv": "nausea vomiting",
"nvd": "nausea vomiting diarrhea",
"sorethroat": "sore throat",
"uri": "upper respiratory infection",
"pna": "pneumonia",
# Conservative high-frequency shorthand expansions from uncodable output
    ↵ audit.
"ich": "intracranial hemorrhage",
"sah": "subarachnoid hemorrhage",
"stemi": "st elevation myocardial infarction",
"nstemi": "non st elevation myocardial infarction",
"chf": "congestive heart failure",
"pe": "pulmonary embolism",
"sbo": "small bowel obstruction",
"gib": "gastrointestinal bleeding",
"hypoglycemia": "hypoglycemia",
"hypokalemia": "hypokalemia",
"arf": "acute renal failure",
"ili": "influenza like illness",
"covid": "covid",
"flu": "influenza",
"si": "suicidal ideation",
"hi": "homicidal ideation",
"avh": "auditory visual hallucinations",
"etoh": "alcohol",
"od": "overdose",
"fb": "foreign body",
"vag": "vaginal",
"vb": "vaginal bleeding",
"hg": "hyperemesis gravidarum",
"s/p": "status post",
"w/": "with",
"w/o": "without",
"h/o": "history of",

```

```

    "c/o": "complains of",
    "n/v": "nausea vomiting",
    "n/v/d": "nausea vomiting diarrhea",
}

CTX_OVR: dict[str, tuple[Any, str, str | None]] = {
    "od": (
        lambda seg: not
            ↵ re.search(r"\b(eye|vision|ophth|ophthalm|cornea|ocular)\b", seg),
        "overdose",
        "right eye",
    ),
    "os": (
        lambda seg:
            ↵ re.search(r"\b(eye|vision|ophth|ophthalm|cornea|ocular)\b", seg),
        "left eye",
        None,
    ),
    "ou": (
        lambda seg:
            ↵ re.search(r"\b(eye|vision|ophth|ophthalm|cornea|ocular)\b", seg),
        "both eyes",
        None,
    ),
}
}

ABBR_MAP: dict[str, str] = {}

SAFE_SHORT_TOKENS = {"cp", "ha", "nv", "gi", "gu", "ms", "sx", "bp", "hr",
    ↵ "doe", "sob", "loc"}

_PUNCT_EDGE = re.compile(r"^(?![\w/]+)|([\w/]+$)")
SEP_SPLIT_RE = re.compile(r"\s*(?:[;:,]+|[\+\&])\s*")

CLINICAL_KEYS = {
    "pain",
    "fever",
    "chill",
    "cough",
    "wheez",
    "dyspnea",
    "shortness",
    "hemoptysis",
}

```

"nausea",
"vomit",
"diarrhea",
"bleeding",
"rash",
"injury",
"wound",
"laceration",
"fracture",
"sprain",
"assault",
"fall",
"seizure",
"syncope",
"weakness",
"numbness",
"tingling",
"headache",
"dizziness",
"palpitation",
"chest",
"abdominal",
"flank",
"pelvic",
"back",
"neck",
"urinary",
"dysuria",
"hematuria",
"pregnan",
"vaginal",
"overdose",
"intoxication",
"anxiety",
"depression",
"psychosis",
"sore",
"throat",
"ear",
"eye",
"dental",
"asthma",
"copd",

```

    "flu",
    "covid",
    "pneumonia",
}

_DUR_RE = re.compile(
    r"\b([a-z]{2,})\s*x\s*(\d{1,3})\s*(d|day|h|hr|hour|w|wk|week|m|mo|mon|'
    'month|y|yr|year)s?\b",
    re.I,
)
_UNIT_MAP = {
    "d": "days",
    "day": "days",
    "h": "hours",
    "hr": "hours",
    "hour": "hours",
    "w": "weeks",
    "wk": "weeks",
    "week": "weeks",
    "m": "months",
    "mo": "months",
    "mon": "months",
    "month": "months",
    "y": "years",
    "yr": "years",
    "year": "years",
}
# --- Rule-gate regex definitions for deterministic high-confidence
# overrides. ---
RX_FOCAL_WEAK = re.compile(
    r"\b(focal|left|right|arm|leg|hand|face|hemiparesis|hemiplegia).*\b"
    'weakness\b|\bweakness\b.*\b(focal|left|right|arm|leg|hand|face)\b',
    re.I,
)
RULE_RX: dict[str, re.Pattern[str]] = {
    "RVC-INJ": re.compile(
        r"\b(mvc|mva|collision|ped(\s|)vs(\s|)auto|assault|fell|fall|gsw|'
        'gunshot|stab|laceration|fracture|burn|foreign'
        '|body|poison(ing)?|overdose|od\b|adverse (drug|medication)'
        '|reaction)\b",
        re.I,
    )
}

```

```

),
"RVC-SYM-RESP": re.compile(
    r"\b(shortness of breath|dyspnea|doe\b|wheeze|respiratory
     ↳ distress|cough(ing)?|hemoptysis|sputum)\b",
    re.I,
),
"RVC-SYM-CIRC": re.compile(
    r"\b(chest pain|palpitation(s)?|hypotension|low
     ↳ bp|bradycardia|tachycardia|arrhythmia|irregular
     ↳ heartbeat|edema|leg swelling|ankle swelling|peripheral
     ↳ edema|claudication)\b",
    re.I,
),
"RVC-SYM-NERV": re.compile(
    r"\b(alter(ed)? mental status|ams\b|confusion|letharg(y|ic)|_
     ↳ unresponsive|syncope|faint(ing)?|seizure(s)?|slurred
     ↳ speech|dysarthria|aphasia|facial
     ↳ droop|headache|migraine|dizziness|vertigo|numbness|tingling)\b",
    re.I,
),
"RVC-SYM-DIG": re.compile(
    r"\b(abd(ominal)? pain|abdo pain|nausea|vomit(ting)?|diarrhea|rectal
     ↳ bleeding|hematemesis|dysphagia|jaundice|gastrointestinal
     ↳ bleeding|gi bleeding)\b",
    re.I,
),
"RVC-SYM-GU": re.compile(
    r"\b(dysuria|hematuria|urinary (frequency|urgency)|flank pain|pelvic
     ↳ pain|vaginal (bleeding|discharge)|penile discharge)\b",
    re.I,
),
"RVC-SYM-MSK": re.compile(
    r"\b(back pain|neck pain|shoulder pain|knee pain|hip pain|joint
     ↳ pain|gait problem)\b",
    re.I,
),
"RVC-SYM-SKIN":
    ↳ re.compile(r"\b(rash|pruritus|itch(ing)?|cellulitis|wound(?!?
     ↳ check))\b", re.I),
"RVC-SYM-EYE": re.compile(
    r"\b(eye (pain|redness|irritation)|red eye|pink
     ↳ ?eye|conjunctivitis|blurry vision|vision (loss|change)|ear
     ↳ pain|hearing (loss|change)|tinnitus|ear discharge|vertigo)\b",

```

```

        re.I,
),
"RVC-SYM-GEN": re.compile(r"\b(fever|chills|fatigue|malaise|weakness(?!.*focal))\b",
    re.I),
"RVC-SYM-PSY": re.compile(
    r"\b(anxiety|depression|insomnia|agitation|suicidal
     ↳ ideation|homicidal ideation|si\b|hi\b|intoxication)\b",
    re.I,
),
"RVC-DIS": re.compile(
    r"\b(asthma attack|copd
     ↳ flare|pneumonia|diabetes(?!.*test)|hypertensive crisis|stroke
     ↳ diagnosed|stroke|intracranial hemorrhage|subarachnoid
     ↳ hemorrhage|st elevation myocardial infarction|non st elevation
     ↳ myocardial infarction|congestive heart failure|pulmonary
     ↳ embolism|small bowel obstruction|acute renal failure|influenza
     ↳ like illness|hypoglycemia|hypokalemia)\b",
    re.I,
),
"RVC-TEST": re.compile(
    r"\b(abnormal (lab|labs|imaging|ct|mri|x[- ]?ray|ekg|ecg)|told (to
     ↳ come|results?)|positive (test|culture|blood culture)|blood
     ↳ culture (positive|grew|growth))\b",
    re.I,
),
"RVC-DIAG": re.compile(
    r"\b((needs|for) (covid|flu|strep) test|screen(ing)?|bp
     ↳ (check|recheck)|workup|routine testing|prenatal)\b",
    re.I,
),
"RVC-TREAT": re.compile(
    r"\b(refill|prescription refill|rx refill|dressing change|suture
     ↳ removal|injection|detox clearance|wound check)\b",
    re.I,
),
"RVC-ADMIN": re.compile(r"\b(form|paperwork|note|letter|clearance|insurance|work
    ↳ (note|letter)|transfer(red)?|intubated transfer)\b", re.I),
}

GROUP_MAP: list[tuple[re.Pattern[str], str]] = [
    (re.compile(r"\binjur|poison|overdose|adverse effect", re.I), "RVC-INJ"),

```

```

(re.compile(r"\brespir", re.I), "RVC-SYM-RESP"),
(re.compile(r"\bcircul|cardio|lymph", re.I), "RVC-SYM-CIRC"),
(re.compile(r"\bnerv|neuro", re.I), "RVC-SYM-NERV"),
(re.compile(r"\bdigest|gastro|abd", re.I), "RVC-SYM-DIG"),
(re.compile(r"\bgenitour|urinar|renal|pelvic|vagin|penil|breast", re.I),
"RVC-SYM-GU"),
(re.compile(r"\bmusculoskelet|msk|joint|back|neck", re.I),
"RVC-SYM-MSK"),
(re.compile(r"\bskin|hair|nail|derma", re.I), "RVC-SYM-SKIN"),
(re.compile(r"\beye|ear|vision|hearing|vertigo", re.I), "RVC-SYM-EYE"),
(re.compile(r"\bgeneral|fever|malaise|weakness", re.I), "RVC-SYM-GEN"),
(re.compile(r"\bpsych|mental", re.I), "RVC-SYM-PSY"),
(re.compile(r"\bdisease|diagnos|known|asthma|copd|diabetes|hypsentr|_"
pneumonia|stroke", re.I), "RVC-DIS"),
(re.compile(r"\babnormal .*test|abnormal (lab|imaging|ekg|ecg)|positive"
(test|culture|result)", re.I), "RVC-TEST"),
(re.compile(r"\bscreen|diagnostic|prenatal|immuniz|vaccine|bp"
check|test\b", re.I), "RVC-DIAG"),
(re.compile(r"\btreat|medicat|refill|suture|dressing|injection|therapy|_"
wound check", re.I), "RVC-TREAT"),
(re.compile(r"\badmin|paperwork|form|clearance|insurance|work"
(note|letter)", re.I), "RVC-ADMIN"),
(re.compile(r"\buncodable|unknown|illegible|none", re.I), "RVC-UNCL"),
]

RE_DROP = re.compile(r"\b(nec|nos|unspecified|other|other and"
"unspecified)\b", re.I)
RE_SPACE = re.compile(r"\s+")
BODY_PART_ONLY_RE = re.compile(r"(arm|leg|back|knee|hip|wrist|elbow|ankle|_"
hand|foot|toe|finger|eye|ear)")

CANON_REPLACEERS: list[tuple[re.Pattern[str], str]] = [
    (re.compile(r"\bshortness of breath\b", flags=re.I), "dyspnea"),
    (re.compile(r"\bdyspneic?\b", flags=re.I), "dyspnea"),
    (re.compile(r"\bdyspnea on exertion\b", flags=re.I), "dyspnea_exertion"),
    (re.compile(r"\bstatus post\b", flags=re.I), "status_post"),
]
}

RVC_PROTOs_CURATED: dict[str, list[str]] = {
    "RVC-INJ": [

```

```
"injury: laceration cut",
"injury: fracture",
"injury: burn",
"injury: contusion bruise",
"injury: foreign body",
"injury: motor vehicle collision",
"injury: fall",
"adverse effect: medication reaction",
"poisoning overdose",
"injury: assault",
"gunshot wound",
],
"RVC-SYM-RESP": [
    "respiratory symptom: shortness of breath dyspnea",
    "respiratory symptom: dyspnea on exertion",
    "respiratory symptom: wheeze",
    "respiratory symptom: cough",
    "respiratory symptom: hemoptysis",
    "respiratory symptom: sputum production",
    "respiratory symptom: respiratory distress",
],
"RVC-SYM-CIRC": [
    "circulatory symptom: chest pain suspected cardiac",
    "circulatory symptom: palpitations",
    "circulatory symptom: hypotension low blood pressure",
    "circulatory symptom: bradycardia",
    "circulatory symptom: tachycardia",
    "circulatory symptom: peripheral edema leg swelling ankle swelling",
    "circulatory symptom: claudication",
],
"RVC-SYM-NERV": [
    "neurologic symptom: altered mental status confusion",
    "neurologic symptom: lethargy decreased responsiveness",
    "neurologic symptom: slurred speech dysarthria",
    "neurologic symptom: aphasia word finding difficulty",
    "neurologic symptom: facial droop",
    "neurologic symptom: syncope fainting",
    "neurologic symptom: seizure",
    "neurologic symptom: headache",
    "neurologic symptom: dizziness vertigo",
    "neurologic symptom: numbness tingling focal weakness",
],
"RVC-SYM-DIG": [
```

```

    "digestive symptom: abdominal pain",
    "digestive symptom: nausea vomiting",
    "digestive symptom: diarrhea",
    "digestive symptom: rectal bleeding",
    "digestive symptom: dysphagia",
    "digestive symptom: jaundice",
],
"RVC-SYM-GU": [
    "genitourinary symptom: dysuria painful urination",
    "genitourinary symptom: urinary frequency urgency",
    "genitourinary symptom: hematuria blood in urine",
    "genitourinary symptom: flank pain",
    "genitourinary symptom: pelvic pain",
    "genitourinary symptom: vaginal discharge",
],
"RVC-SYM-MSK": [
    "musculoskeletal symptom: back pain",
    "musculoskeletal symptom: neck pain",
    "musculoskeletal symptom: joint pain",
    "musculoskeletal symptom: hip knee shoulder pain",
    "musculoskeletal symptom: gait problem",
],
"RVC-SYM-SKIN": [
    "skin symptom: rash",
    "skin symptom: pruritus itching",
    "skin symptom: cellulitis redness warmth",
    "skin symptom: nontraumatic wound",
],
"RVC-SYM-EYE": [
    "eye symptom: eye pain",
    "eye symptom: red eye eye redness",
    "eye symptom: blurry vision vision loss",
    "ear symptom: ear pain discharge",
    "vestibular symptom: vertigo spinning",
],
"RVC-SYM-GEN": [
    "general symptom: fever chills",
    "general symptom: fatigue malaise",
    "general symptom: weakness generalized",
    "general symptom: weight change",
    "general symptom: edema swelling",
],
"RVC-SYM-PSY": [

```

```

    "psychological symptom: anxiety",
    "psychological symptom: depression",
    "psychological symptom: agitation",
    "psychological symptom: insomnia",
    "psychological symptom: suicidal ideation",
    "psychological symptom: homicidal ideation",
    "psychological symptom: substance intoxication without overdose",
],
"RVC-DIS": [
    "patient stated diagnosis: asthma attack",
    "patient stated diagnosis: copd flare",
    "patient stated diagnosis: pneumonia",
    "patient stated diagnosis: diabetes problem",
    "patient stated diagnosis: hypertensive crisis",
    "patient stated diagnosis: stroke diagnosed",
],
"RVC-TEST": [
    "abnormal test: abnormal laboratory result",
    "abnormal test: abnormal imaging result",
    "abnormal test: positive culture",
    "abnormal test: abnormal ecg ekg",
    "abnormal test: positive blood culture",
],
"RVC-DIAG": [
    "diagnostic screening: needs covid test",
    "diagnostic screening: blood pressure check",
    "diagnostic screening: routine testing screening",
    "diagnostic screening: prenatal check",
],
"RVC-TREAT": [
    "treatment medication: prescription refill",
    "treatment medication: dressing change",
    "treatment medication: injection therapy",
    "treatment medication: suture removal",
    "treatment: wound check",
],
"RVC-ADMIN": [
    "administrative reason: work form school form insurance paperwork",
    "administrative reason: clearance note",
],
"RVC-UNCL": ["uncodable or unknown reason"],
}

```

```

for _group in RVC_NAME:
    RVC_PROTOS_CURATED.setdefault(_group, [])

LABEL2CODE = {
    re.sub(r"\s+", " ", value.strip().lower().replace("-", "--").replace("-",
    "-")): code
    for code, value in RVC_NAME.items()
}

sym = SymSpell(max_dictionary_edit_distance=2, prefix_length=7)

# --- Lightweight data containers passed between pipeline stages. ---
@dataclass
class SegmentEmbedCache:
    """Embeddings and similarity cache for unique complaint segments."""

    uniq_segments: list[str]
    seg2idx: dict[str, int]
    emb: torch.Tensor
    sim_proto: torch.Tensor

@dataclass
class SegPred:
    """Per-segment prediction record."""

    seg_idx: int
    segment: str
    code: str
    name: str
    sim: float
    rule_code: str | None
    rule_used: bool

def _add_lexicon_from_phrases(phrases: list[str]) -> None:
    for phrase in phrases:
        for token in str(phrase).split():
            if token.isalpha():
                sym.create_dictionary_entry(token, 1)

```

```

_add_lexicon_from_phrases(list(SEG_OVR.values()) + list(TOK_OVR.values()) +
    ↵ list(ABBR_MAP.values()))

# --- Core text preprocessing helpers. ---
def get_cc_column(df: pd.DataFrame) -> str:
    """Return the first recognized chief-complaint column."""
    for column in CC_CANDIDATES:
        if column in df.columns:
            return column
    raise KeyError(f"No chief-complaint column found in candidates:
        ↵ {CC_CANDIDATES}")

def _clean_tok(token: str) -> str:
    return _PUNCT_EDGE.sub("", token)

def _violates_spell_guard(original: str, candidate: str) -> bool:
    original_norm = original.strip().lower()
    candidate_norm = candidate.strip().lower()
    if not original_norm or not candidate_norm:
        return False
    if original_norm == candidate_norm:
        return False
    for prefix_from, prefix_to in POLARITY_GUARD_PREFIX_PAIRS:
        if original_norm.startswith(prefix_from) and
            ↵ candidate_norm.startswith(prefix_to):
            return True
    if original_norm in LATERALITY_GUARD_TOKENS or candidate_norm in
        ↵ LATERALITY_GUARD_TOKENS:
        return original_norm != candidate_norm
    return False

def _spell_candidate(token: str, *, max_edit_distance: int) -> tuple[str, int
    ↵ | None]:
    candidates = sym.lookup(token, Verbosity.CLOSEST,
    ↵ max_edit_distance=max_edit_distance)
    if not candidates:
        return token, None
    match = candidates[0]
    return match.term, int(match.distance)

```

```

def _spell_strict(token: str) -> tuple[str, dict[str, Any] | None]:
    normalized = token.strip().lower()
    if (
        not normalized
        or normalized in SPELL_PROTECT_TOKENS
        or not normalized.isalpha()
        or len(normalized) < STRICT_SPELL_MIN_TOKEN_LEN
    ):
        return token, None
    corrected, distance = _spell_candidate(normalized,
    ↵ max_edit_distance=STRICT_SPELL_MAX_EDIT_DISTANCE)
    corrected_norm = corrected.strip().lower()
    if (
        distance is None
        or corrected_norm == normalized
        or _violates_spell_guard(normalized, corrected_norm)
    ):
        return token, None
    return corrected_norm, {
        "orig": normalized,
        "corrected": corrected_norm,
        "edit_distance": int(distance),
    }

def _spell_token(token: str, *, mode: str) -> tuple[str, dict[str, Any] |
    ↵ None]:
    if mode == "disabled":
        return token, None
    if mode == "strict":
        return _spell_strict(token)
    raise ValueError(f"Unsupported spell correction mode: {mode}")

def _expand_tokens(segment: str) -> str:
    """Apply contextual and curated abbreviation expansions in one
    ↵ segment."""
    output: list[str] = []
    segment_context = segment
    for token in segment.split():
        cleaned = _clean_tok(token)

```

```

        if not cleaned:
            continue
        if cleaned in CTX_OVR:
            predicate, yes_value, no_value = CTX_OVR[cleaned]
            if predicate(segment_context):
                output.extend(yes_value.split())
                continue
            if no_value:
                output.extend(no_value.split())
                continue
        if cleaned in TOK_OVR:
            output.extend(TOK_OVR[cleaned].split())
            continue
        if cleaned in ABBR_MAP:
            output.extend(ABBR_MAP[cleaned].split())
            continue
        output.append(cleaned)
    return " ".join(output)

def _looks_clinical(text: str) -> bool:
    lowered = text.lower()
    return any(key in lowered for key in CLINICAL_KEYS) and
        bool(re.search(r"[a-z]", lowered))

def _split_on_and_if_clinical(segment: str) -> list[str]:
    parts = [part.strip() for part in re.split(r"\band\b", segment,
        flags=re.I)]
    if len(parts) == 1:
        return parts
    output: list[str] = []
    buffer = parts[0]
    for next_part in parts[1:]:
        if _looks_clinical(buffer) and _looks_clinical(next_part):
            output.append(buffer.strip())
            buffer = next_part
        else:
            buffer = f"{buffer} and {next_part}"
    output.append(buffer.strip())
    return output

```

```

def _normalize_absence_symptom(segment: str) -> str | None:
    """Convert selected absence-style phrases into symptom phrases we can
       classify."""
    cleaned = segment.strip()
    if not cleaned:
        return None

    for regex, replacement in ABSENCE_SYMPTOM_RULES:
        if regex.search(cleaned):
            return regex.sub(replacement, cleaned).strip()

    return None

def _expand_duration(segment: str) -> str:
    def repl(match: re.Match[str]) -> str:
        base = match.group(1)
        count = match.group(2)
        unit = match.group(3).lower()
        resolved_unit = _UNIT_MAP.get(unit, unit)
        return f"{base} {count} {resolved_unit}".strip()

    return _DUR_RE.sub(repl, segment)

def _lemmatize_segment(segment: str) -> str:
    doc = NLP(segment)
    lemmas: list[str] = []
    for token in doc:
        lemma = token.lemma_.strip() if token.lemma_ else ""
        if not lemma or lemma == "-PRON-":
            lemma = token.text
        lemmas.append(lemma)
    return " ".join(lemmas).strip()

def segment_cc(raw: str, keep_negated: bool = False) -> list[str]:
    """Split a raw chief complaint into normalized text segments."""
    if not isinstance(raw, str) or not raw.strip():
        return []

    text = ftfy.fix_text(raw).lower().strip()

```

```

for protected in PROTECT:
    text = text.replace(protected, protected.replace("/", "§"))

text = re.sub(r"\bn\s*[,/&+]\s*v\s*[,/&+]\s*d\b", "n§v§d", text)
text = re.sub(r"\bn\s*[,/&+]\s*v\b", "n§v", text)
text = re.sub(r"\bn\s*/\s*v\s*/\s*d\b", "n§v§d", text)
text = re.sub(r"\bn\s*/\s*v\b", "n§v", text)
text = re.sub(r"\bc\s*/\s*p\b", "c§p", text)
text = re.sub(r"\bs\s*/\s*p\b", "s§p", text)
text = re.sub(r"\bw\s*/\s*o\b", "w§o", text)
text = re.sub(r"\bh\s*/\s*o\b", "h§o", text)
text = re.sub(r"\bc\s*/\s*o\b", "c§o", text)

text = text.replace("/", "; ")

segments: list[str] = []
for block in SEP_SPLIT_RE.split(text):
    if not block:
        continue
    block = block.replace("§", "/").strip()
    block = " ".join(block.split())
    if MISSING_RE.fullmatch(block):
        continue
    for segment in _split_on_and_if_clinical(block):
        if not segment:
            continue

        normalized_absence = _normalize_absence_symptom(segment)
        if normalized_absence is not None:
            segment = normalized_absence
        elif not keep_negated and NEG_RE.match(segment):
            continue

        if re.fullmatch(r"[a-z]$", segment):
            continue
        segments.append(segment)
return segments

def normalize_cc(
    raw: str,
    keep_negated: bool = False,
    *,

```

```

spell_mode: str = "strict",
) -> tuple[list[str], list[dict[str, Any]]]:
    """Normalize a raw chief complaint to deduplicated, lemmatized segments.

    Returns:
        tuple(list[str], list[dict]): normalized segments and accepted
    ↵ spell-fix log rows.
    """

if spell_mode not in SPELL_CORRECTION_MODES.difference({"auto"}):
    raise ValueError(f"spell_mode must be one of {{'disabled','strict'}},"
        ↵ got {spell_mode!r}")

output: list[str] = []
seen: set[str] = set()
spellfix_log: list[dict[str, Any]] = []

for segment_index, segment in enumerate(segment_cc(raw,
    ↵ keep_negated=keep_negated)):
    if segment in SEG_OVR:
        segment = SEG_OVR[segment]
    else:
        tokens = segment.split()
        if len(tokens) == 1:
            token = _clean_tok(tokens[0])
            if token in TOK_OVR:
                segment = TOK_OVR[token]
            else:
                segment = ABBR_MAP.get(segment, segment)

        segment = _expand_duration(segment)
        segment = _expand_tokens(segment)
        corrected_tokens: list[str] = []
        for token_index, token in enumerate(segment.split()):
            corrected, correction = _spell_token(token, mode=spell_mode)
            corrected_tokens.append(corrected)
            if correction is not None:
                spellfix_log.append(
                    {
                        "segment_idx": int(segment_index),
                        "token_idx": int(token_index),
                        "segment_raw": segment,
                        **correction,
                    }
                )

```

```

        )
segment = _lemmatize_segment(" ".join(corrected_tokens))

if segment and segment not in seen:
    output.append(segment)
    seen.add(segment)

return output, spellfix_log

def _run_normalization_mode(
    cc_series: pd.Series,
    *,
    spell_mode: str,
    keep_negated: bool = False,
) -> dict[str, Any]:
    cleaned_segments: list[list[str]] = []
    spellfix_logs: list[list[dict[str, Any]]] = []
    for raw in cc_series.tolist():
        if pd.isna(raw):
            cleaned_segments.append([])
            spellfix_logs.append([])
            continue
        segments, fix_log = normalize_cc(
            str(raw),
            keep_negated=keep_negated,
            spell_mode=spell_mode,
        )
        cleaned_segments.append(segments)
        spellfix_logs.append(fix_log)

    cleaned_canon = [canonize_cc_segments(segments) for segments in
        ↪ cleaned_segments]
    cleaned_str = ["; ".join(segments) for segments in cleaned_segments]
    cleaned_canon_str = ["; ".join(segments) for segments in cleaned_canon]
    spellfix_n = [len(entries) for entries in spellfix_logs]
    spellfix_any = [count > 0 for count in spellfix_n]
    spellfix_log_json = [json.dumps(entries, ensure_ascii=False) for entries
        ↪ in spellfix_logs]

    normalized_frame = pd.DataFrame(
        {
            "cc_cleaned": cleaned_segments,

```

```

    "cc_cleaned_str": cleaned_str,
    "cc_cleaned_canon": cleaned_canon,
    "cc_cleaned_canon_str": cleaned_canon_str,
    "cc_spellfix_n": pd.Series(spellfix_n, dtype="Int64"),
    "cc_spellfix_any": pd.Series(spellfix_any, dtype="boolean"),
    "cc_spellfix_log": spellfix_log_json,
},
index=cc_series.index,
)
return {
    "mode": spell_mode,
    "frame": normalized_frame,
    "spellfix_count_total": int(sum(spellfix_n)),
}
}

def _score_normalization_integrity(
    raw_series: pd.Series,
    cleaned_str_series: pd.Series,
) -> dict[str, int]:
    raw = raw_series.fillna("").astype("string").str.lower()
    clean = cleaned_str_series.fillna("").astype("string").str.lower()
    rules = {
        "fall_to_small": raw.str.contains(r"\bfall\b", na=False)
        & ~raw.str.contains(r"\bsmall\b", na=False)
        & clean.str.contains(r"\bsmall\b", na=False),
        "head_to_heart": raw.str.contains(r"\bhead\b", na=False)
        & ~raw.str.contains(r"\bheart\b", na=False)
        & clean.str.contains(r"\bheart\b", na=False),
        "foot_to_post": raw.str.contains(r"\bfoot\b", na=False)
        & ~raw.str.contains(r"\bpost\b", na=False)
        & ~raw.str.contains(r"\bs/p\b|\bstatus\s+post\b|\bsp\b", na=False)
        & clean.str.contains(r"\bpost\b", na=False),
        "copd_or_cold_to_covid": raw.str.contains(r"\b(?:copd|cold|cord)\b",
            na=False)
        & ~raw.str.contains(r"\bcovid\b", na=False)
        & clean.str.contains(r"\bcovid\b", na=False),
        "hyperglycemia_to_hypoglycemia":
            raw.str.contains(r"\bhyperglycemia\b", na=False)
        & clean.str.contains(r"\bhypoglycemia\b", na=False),
        "hyperkalemia_to_hypokalemia": raw.str.contains(r"\bhyperkalemia\b",
            na=False)
        & clean.str.contains(r"\bhypokalemia\b", na=False),
    }

```

```

    }

    return {name: int(mask.sum()) for name, mask in rules.items()}

def choose_spell_mode(
    df: pd.DataFrame,
    *,
    cc_source_column: str,
    requested_mode: str = "auto",
) -> tuple[dict[str, Any], pd.DataFrame]:
    """Run disabled/strict normalization and select a mode by integrity
    score."""
    requested = requested_mode.strip().lower()
    if requested not in SPELL_CORRECTION_MODES:
        raise ValueError(
            f"CC_SPELL_CORRECTION_MODE must be one of"
            f"\n{sorted(SPELL_CORRECTION_MODES)}"
        )
    cc_series = df[cc_source_column].astype("string")
    modes_to_run = ["disabled", "strict"] if requested == "auto" else
    [requested]

    mode_runs: dict[str, dict[str, Any]] = {}
    mode_scores: list[dict[str, Any]] = []
    for mode in modes_to_run:
        run_payload = _run_normalization_mode(cc_series, spell_mode=mode)
        mode_runs[mode] = run_payload
        integrity_counts = _score_normalization_integrity(
            cc_series,
            run_payload["frame"]["cc_cleaned_str"],
        )
        integrityViolationTotal = int(sum(integrityCounts.values()))
        nonemptySegmentRows = int(
            run_payload["frame"]["cc_cleaned"].map(lambda values: len(values))
        > 0).sum()
        mode_scores.append(
            {
                "mode": mode,
                "integrityViolationTotal": integrityViolationTotal,
                "nonemptySegmentRows": nonemptySegmentRows,
                "spellfixCountTotal":
                    int(run_payload["spellfix_count_total"]),
            }
        )

```

```

        **integrity_counts,
    }
)

score_df = pd.DataFrame(mode_scores).sort_values(
    ["integrityViolationTotal", "nonemptySegmentRows",
     "spellfixCountTotal", "mode"],
    ascending=[True, False, False, True],
    kind="stable",
)
selected_mode = str(score_df.iloc[0]["mode"])
selected_payload = mode_runs[selected_mode]
selected_payload["selection_scores"] = score_df.reset_index(drop=True)
selected_payload["requestedMode"] = requested
selected_payload["selectedMode"] = selected_mode
return selected_payload, score_df.reset_index(drop=True)

def canonize_cc_segments(segments: list[str]) -> list[str]:
    """Apply light concept-level canonicalization for
    diagnostics/visualization."""
    output: list[str] = []
    for segment in segments:
        transformed = segment
        for regex, replacement in CANON_REPLACEERS:
            transformed = regex.sub(replacement, transformed)
        output.append(transformed)
    return output

def rule_gate(segment: str) -> tuple[str | None, list[str]]:
    """Apply deterministic rule overrides for high-confidence pattern
    matches."""
    text = f" {segment.lower()} "
    if re.search(r"\babnormal\s+labs?\b", text):
        return "RVC-TEST", ["abnormal labs"]
    if "weakness" in text and RX_FOCAL_WEAK.search(text):
        return "RVC-SYM-NERV", ["focal weakness"]

    hits: list[str] = []
    for code, regex in RULE_RX.items():

```

```

        if regex.search(text):
            hits.append(code)

    if not hits:
        return None, []

    for preferred_code in PRECEDENCE_RVC:
        if preferred_code in hits:
            return preferred_code, [preferred_code]

    return hits[0], [hits[0]]


# --- Ontology file parsing helpers (for Appendix and summary resources). ---
def _guess_col(columns: list[str], candidates: list[str]) -> str | None:
    lowered = [column.lower().strip() for column in columns]
    for candidate in candidates:
        if candidate in lowered:
            return columns[lowered.index(candidate)]
    return None

def _norm_label(label: str) -> str:
    if label is None:
        return ""
    cleaned = str(label).strip().lower().replace("-", "-").replace("–", "–")
    return re.sub(r"\s+", " ", cleaned)

def map_group(label: str) -> str | None:
    for regex, code in GROUP_MAP:
        if regex.search(str(label)):
            return code
    return None

def _label_to_code(label: str) -> str | None:
    normalized = _norm_label(label)
    if normalized in LABEL2CODE:
        return LABEL2CODE[normalized]
    return map_group(label)

```

```

def canon_phrase(text: str) -> str:
    cleaned = str(text).strip().replace("'", "").replace("-", "-")
    cleaned = RE_SPACE.sub(" ", cleaned)
    return cleaned.strip(" .;,:")

def keep_phrase(text: str, group_code: str) -> bool:
    lowered = text.lower()
    if len(lowered) < 3:
        return False
    if RE_DROP.search(lowered):
        return False
    if group_code not in {"RVC-INJ", "RVC-SYM-MSK", "RVC-SYM-EYE",
                           "RVC-SYM-SKIN"}:
        if BODY_PART_ONLY_RE.fullmatch(lowered):
            return False
    return True

def load_appendix_phrases(path: Path) -> pd.DataFrame:
    """Load and normalize Appendix II prototype phrases by RVC code."""
    if not path.exists():
        return pd.DataFrame(columns=["group_code", "phrase"])

    appendix_df = pd.read_csv(path, dtype=str)

    if {"top_level_group17", "text"}.issubset(appendix_df.columns):
        group_col = "top_level_group17"
        phrase_col = "text"
    else:
        group_col = _guess_col(
            list(appendix_df.columns),
            ["top_level_group17", "group", "top_level_group", "rvc_group",
             "module"],
            )
        phrase_col = _guess_col(
            list(appendix_df.columns),
            ["text", "phrase", "subentry", "entry", "label"],
            )
    if group_col is None or phrase_col is None:
        raise ValueError(f"Unable to identify expected columns in
                         appendix file: {path}")

```

```

normalized = appendix_df[[group_col, phrase_col]].rename(
    columns={group_col: "group_raw", phrase_col: "phrase_raw"}
)
normalized["group_code"] = normalized["group_raw"].map(_label_to_code)
normalized = normalized[normalized["group_code"].notna()].copy()
normalized["phrase"] = normalized["phrase_raw"].map(canon_phrase)
normalized = normalized[
    normalized.apply(lambda row: keep_phrase(row["phrase"]),
    row["group_code"]), axis=1
]
return normalized[["group_code", "phrase"]].drop_duplicates()

def load_summary_weights(path: Path, strategy: str = "uniform") -> dict[str,
    float]:
    """Load optional group weights and map them to canonical RVC codes."""
    if not path.exists():
        return {}

    summary_df = pd.read_csv(path, dtype=str)
    group_col = _guess_col(
        list(summary_df.columns),
        ["top_level_group17", "group", "top_level_group", "rvc_group",
    "module"],
    )
    if group_col is None:
        group_col = summary_df.columns[0]

    mapped_codes = summary_df[group_col].map(_label_to_code).dropna()
    if mapped_codes.empty:
        return {}

    counts = mapped_codes.value_counts().astype(float)
    if strategy == "uniform":
        counts[:] = 1.0
    elif strategy != "by_rows":
        raise ValueError("strategy must be one of {'uniform', 'by_rows'}")

    scaled = np.log1p(counts) / np.log1p(counts).max()
    return {code: float(max(1e-3, scaled.get(code, 1.0))) for code in
    PRECEDENCE_RVC}

```

```

def dedup_prototypes(
    rvc_prototypes: dict[str, list[str]],
    encoder: SentenceTransformer,
    sim_thresh: float = 0.92,
    batch_size: int = 256,
) -> dict[str, list[str]]:
    """Deduplicate near-identical prototype phrases within each group."""
    output: dict[str, list[str]] = {}

    for group_code, phrases in rvc_prototypes.items():
        normalized = sorted(dict.fromkeys(canon_phrase(phrase) for phrase in
        ↵ phrases))
        if not normalized:
            output[group_code] = []
            continue

        with torch.inference_mode():
            embedding = encoder.encode(
                normalized,
                convert_to_tensor=True,
                normalize_embeddings=True,
                batch_size=batch_size,
            )

        keep: list[str] = []
        taken = torch.zeros(len(normalized), dtype=torch.bool)
        for idx in range(len(normalized)):
            if taken[idx]:
                continue
            keep.append(normalized[idx])
            if len(normalized) == 1:
                break
            similarities = util.cos_sim(embedding[idx : idx + 1],
        ↵ embedding).squeeze(0).cpu().numpy()
            duplicate_indices = np.where(similarities >= sim_thresh)[0]
            taken[duplicate_indices] = True

        output[group_code] = keep

    return output

# --- Embedding/prototype resource construction. ---

```

```

def build_prototype_resources(cfg: ClassifierConfig) -> dict[str, Any]:
    """Build model, prototype embeddings, scoring indices, and metadata."""
    device = "cuda" if torch.cuda.is_available() else "cpu"
    hf_token = os.getenv(cfg.hf_token_env_var, "").strip() or None
    encoder = SentenceTransformer(
        cfg.model_name,
        device=device,
        trust_remote_code=cfg.model_trust_remote_code,
        token=hf_token,
    )

    work_dir = Path(os.getenv(cfg.work_dir_env_var,
    ↵ Path.cwd())).expanduser().resolve()
    appendix_path = work_dir / cfg.appendix_relpah
    summary_path = work_dir / cfg.summary_relpah

    appendix_df = load_appendix_phrases(appendix_path)
    group_weights = load_summary_weights(summary_path,
    ↵ strategy=cfg.summary_weight_strategy)

    prototypes_augmented = {code: list(values) for code, values in
    ↵ RVC_PROTOS_CURATED.items()}
    for group_code, group_df in appendix_df.groupby("group_code"):
        prototypes_augmented[group_code].extend(group_df["phrase"].tolist())

    prototypes_final = dedup_prototypes(
        prototypes_augmented,
        encoder=encoder,
        sim_thresh=cfg.prototype_dedup_sim_thresh,
        batch_size=cfg.model_batch_size,
    )

    all_prototype_texts: list[str] = []
    proto_to_code: list[str] = []
    for code, prototype_list in prototypes_final.items():
        prefix = RVC_NAME[code].split("-")[0].strip().lower()
        for phrase in prototype_list:
            text = phrase if ":" in phrase else f"{prefix}: {phrase}"
            all_prototype_texts.append(text)
            proto_to_code.append(code)

    with torch.inference_mode():
        proto_emb = encoder.encode(

```

```

        all_prototype_texts,
        convert_to_tensor=True,
        normalize_embeddings=True,
        batch_size=cfg.model_batch_size,
    )

code_to_proto_indices: dict[str, list[int]] = defaultdict(list)
for idx, code in enumerate(proto_to_code):
    code_to_proto_indices[code].append(idx)

proto_weight = np.ones(len(proto_to_code), dtype="float32")
if group_weights:
    for idx, code in enumerate(proto_to_code):
        proto_weight[idx] = max(1e-3, float(group_weights.get(code,
        ↵ 1.0)))

metadata = {
    "model_name": cfg.model_name,
    "device": device,
    "prototype_total": len(all_prototype_texts),
    "prototype_count_by_group": {code:
        ↵ len(code_to_proto_indices.get(code, [])) for code in
        ↵ PRECEDENCE_RVC},
    "appendix_rows_used": int(len(appendix_df)),
    "spacy_model": SPACY_MODEL_LABEL,
    "hf_token_present": bool(hf_token),
}
}

return {
    "encoder": encoder,
    "device": device,
    "proto_emb": proto_emb,
    "proto_to_code": proto_to_code,
    "code_to_proto_indices": code_to_proto_indices,
    "proto_weight": proto_weight,
    "metadata": metadata,
}

# --- Segment caching and scoring helpers. ---
def build_segment_cache(
    df: pd.DataFrame,
    seg_col: str,

```

```

resources: dict[str, Any],
batch_size: int = 512,
) -> SegmentEmbedCache:
    """Encode unique complaint segments once and cache similarity to all
    ↵ prototypes."""
unique_segments: list[str] = []
seen: set[str] = set()

for value in df[seg_col]:
    if isinstance(value, list):
        iterable = value
    elif isinstance(value, str):
        iterable = [value]
    else:
        iterable = []

    for segment in iterable:
        cleaned = str(segment).strip()
        if cleaned and cleaned not in seen:
            unique_segments.append(cleaned)
            seen.add(cleaned)

unique_segments.sort()

if not unique_segments:
    return SegmentEmbedCache(
        uniq_segments=[],
        seg2idx={},
        emb=torch.empty(0, 0),
        sim_proto=torch.empty(0, 0),
    )

encoder: SentenceTransformer = resources["encoder"]
proto_emb: torch.Tensor = resources["proto_emb"]

with torch.inference_mode():
    segment_emb = encoder.encode(
        unique_segments,
        convert_to_tensor=True,
        normalize_embeddings=True,
        batch_size=batch_size,
        device=resources["device"],
    )

```

```

    segment_emb = torch.nn.functional.normalize(segment_emb, p=2, dim=1)
    segment_to_proto = segment_emb @ proto_emb.T

    seg2idx = {segment: idx for idx, segment in enumerate(unique_segments)}
    return SegmentEmbedCache(
        uniq_segments=unique_segments,
        seg2idx=seg2idx,
        emb=segment_emb,
        sim_proto=segment_to_proto,
    )

def group_scores_from_proto_row(
    proto_row: torch.Tensor | np.ndarray,
    resources: dict[str, Any],
    method: str = "max",
    alpha: float = 12.0,
) -> dict[str, float]:
    """Aggregate prototype-level similarities to group-level scores."""
    vector = proto_row.detach().cpu().numpy() if isinstance(proto_row,
    ↵ torch.Tensor) else np.asarray(proto_row)

    scores: dict[str, float] = {}
    code_to_proto = resources["code_to_proto_indices"]
    proto_weight = resources["proto_weight"]

    for group_code, indices in code_to_proto.items():
        values = vector[indices]
        if method == "max":
            score = float(values.max(initial=-1.0))
        elif method == "lse":
            weights = proto_weight[indices]
            logits = alpha * values + np.log(weights)
            max_logit = logits.max()
            score = float(max_logit + np.log(np.exp(logits -
    ↵ max_logit).sum()))
        else:
            raise ValueError("method must be one of {'max', 'lse'}")
        scores[group_code] = score

    if method == "lse" and scores:
        max_score = max(scores.values())
        min_score = min(scores.values())

```

```

        denom = (max_score - min_score) if max_score > min_score else 1.0
        scores = {group: (score - min_score) / denom for group, score in
        ↵ scores.items()}

    return scores

def score_one_segment_soft(
    segment: str,
    resources: dict[str, Any],
    method: str = "max",
    alpha: float = 12.0,
) -> dict[str, float]:
    """Score a single segment directly against prototype embeddings."""
    if not isinstance(segment, str) or not segment.strip():
        return {group: -1.0 for group in RVC_NAME}

    encoder: SentenceTransformer = resources["encoder"]
    proto_emb: torch.Tensor = resources["proto_emb"]

    with torch.inference_mode():
        segment_emb = encoder.encode(
            [segment],
            convert_to_tensor=True,
            normalize_embeddings=True,
            device=resources["device"],
        )
        similarities = util.cos_sim(segment_emb, proto_emb).squeeze(0)

    return group_scores_from_proto_row(similarities, resources=resources,
        ↵ method=method, alpha=alpha)

def classify_segments_cached(
    segments: list[str],
    cache: SegmentEmbedCache,
    resources: dict[str, Any],
    abstain: float = 0.40,
    method: str = "max",
    alpha: float = 12.0,
) -> list[SegPred]:
    """Classify all segments for one visit using cache + rule-gate
    ↵ overrides."""

```

```

predictions: list[SegPred] = []

for seg_idx, segment in enumerate(segments):
    cleaned = str(segment).strip()
    if not cleaned:
        continue

    override_code, _ = rule_gate(cleaned)

    if cleaned in cache.seg2idx:
        row_idx = cache.seg2idx[cleaned]
        score_map = group_scores_from_proto_row(
            cache.sim_proto[row_idx], resources=resources, method=method,
        ↪ alpha=alpha
        )
    else:
        score_map = score_one_segment_soft(
            cleaned, resources=resources, method=method, alpha=alpha
        )

    best_code, best_score = max(score_map.items(), key=lambda kv: kv[1])

    if override_code is not None:
        best_code = override_code

    if best_score < abstain and override_code is None:
        best_code = "RVC-UNCL"

    predictions.append(
        SegPred(
            seg_idx=seg_idx,
            segment=cleaned,
            code=best_code,
            name=RVC_NAME.get(best_code, best_code),
            sim=float(best_score),
            rule_code=override_code,
            rule_used=override_code is not None,
        )
    )

return predictions

```

```

def rfv_from_segment_preds(segpreds: list[SegPred], max_rfv: int = 5) ->
    list[dict[str, Any]]:
    """Aggregate segment-level predictions to visit-level RFV slots."""
    first_position: dict[str, int] = {}
    best_similarity: dict[str, float] = {}
    support_segment: dict[str, str] = {}

    for prediction in segpreds:
        group_code = prediction.code
        if group_code == "RVC-UNCL":
            continue

        if group_code not in first_position or prediction.seg_idx <
            first_position[group_code]:
            first_position[group_code] = prediction.seg_idx
            support_segment[group_code] = prediction.segment

        best_similarity[group_code] = max(best_similarity.get(group_code,
        -1.0), prediction.sim)

    ordered = sorted(
        first_position.keys(),
        key=lambda code: (PRECEDENCE_RVC.index(code), first_position[code]),
    )
    ordered = ordered[:max_rfv]

    return [
        {
            "code": code,
            "name": RVC_NAME[code],
            "support": support_segment[code],
            "sim": best_similarity[code],
        }
        for code in ordered
    ]

# --- Visit-level RFV assignment and output shaping. ---
def assign_rvc_per_segment_and_visit_cached(
    df: pd.DataFrame,
    cache: SegmentEmbedCache,
    resources: dict[str, Any],
    seg_col: str = "cc_cleaned",

```

```

    cfg: ClassifierConfig | None = None,
) -> pd.DataFrame:
    """Attach standardized RFV outputs to each visit using cached segment
    embeddings."""
    active_cfg = cfg or ClassifierConfig()
    max_rfv = active_cfg.max_rfv
    if active_cfg.empty_primary_policy not in {"blank", "uncodable"}:
        raise ValueError("empty_primary_policy must be one of {'blank',
        'uncodable'}")

    rows: list[dict[str, Any]] = []
    for value in df[seg_col].tolist():
        if isinstance(value, list):
            segments = [str(segment).strip() for segment in value if
            str(segment).strip()]
        elif isinstance(value, str) and value.strip():
            segments = [value.strip()]
        else:
            segments = []

    segments = segments[: active_cfg.max_segments_per_visit]
    segment_predictions = classify_segments_cached(
        segments,
        cache=cache,
        resources=resources,
        abstain=active_cfg.abstain_threshold,
        method=active_cfg.scoring_method,
        alpha=active_cfg.alpha,
    )
    visit_rfv = rfv_from_segment_preds(segment_predictions,
    ↵ max_rfv=max_rfv)

    # If all non-empty segment predictions are uncodable, explicitly
    ↵ carry that into RFV1.
    if (
        not visit_rfv
        and active_cfg.allow_uncodable_primary
        and segments
        and segment_predictions
    ):
        uncodable_predictions = [pred for pred in segment_predictions if
    ↵ pred.code == "RVC-UNCL"]
        if uncodable_predictions:

```

```

        chosen = max(uncodable_predictions, key=lambda pred:
←  (pred.sim, -pred.seg_idx))
        visit_rfv = [
            {
                "code": "RVC-UNCL",
                "name": RVC_NAME["RVC-UNCL"],
                "support": chosen.segment,
                "sim": float(chosen.sim),
            }
        ]
    ]

# Optional policy for truly empty complaints.
if (
    not visit_rfv
    and not segments
    and active_cfg.allow_uncodable_primary
    and active_cfg.empty_primary_policy == "uncodable"
):
    visit_rfv = [
        {
            "code": "RVC-UNCL",
            "name": RVC_NAME["RVC-UNCL"],
            "support": "",
            "sim": math.nan,
        }
    ]
]

row: dict[str, Any] = {
    "segment_preds": json.dumps([prediction.__dict__ for prediction
        in segment_predictions], ensure_ascii=False)
}

for idx in range(max_rfv):
    if idx < len(visit_rfv):
        code = visit_rfv[idx]["code"]
        name = visit_rfv[idx]["name"]
        support = visit_rfv[idx]["support"]
        sim = float(visit_rfv[idx]["sim"])
    else:
        code = ""
        name = ""
        support = ""
        sim = math.nan

```

```

slot = idx + 1
row[f"RFV{slot}"] = RVC_SHORT.get(code, "")
row[f"RFV{slot}_name"] = name
row[f"RFV{slot}_support"] = support
row[f"RFV{slot}_sim"] = sim

rows.append(row)

drop_columns = [column for column in df.columns if
    str(column).startswith("RFV") or column == "segment_preds"]
output_df = df.drop(columns=drop_columns, errors="ignore").copy()
output_df = pd.concat([output_df, pd.DataFrame(rows,
    index=output_df.index)], axis=1)
return output_df

# --- Output QA and diagnostic helpers used later in the notebook. ---
def validate_output_schema(
    df: pd.DataFrame,
    max_rfv: int = 5,
    require_primary_for_nonempty: bool = True,
) -> None:
    """Validate required output columns, uniqueness, and JSON shape for
    segment predictions."""
    required = ["segment_preds"]
    for slot in range(1, max_rfv + 1):
        required.extend(
            [
                f"RFV{slot}",
                f"RFV{slot}_name",
                f"RFV{slot}_support",
                f"RFV{slot}_sim",
            ]
        )

    duplicates = df.columns[df.columns.duplicated()].tolist()
    if duplicates:
        raise ValueError(f"Duplicate columns detected: {duplicates}")

    missing = [column for column in required if column not in df.columns]
    if missing:
        raise ValueError(f"Missing required output columns: {missing}")

```

```

bad_json_rows = 0
parsed_payloads: list[list[dict[str, Any]] | None] = []
for value in df["segment_preds"].fillna("[]"):
    if not isinstance(value, str):
        bad_json_rows += 1
        parsed_payloads.append(None)
        continue
    try:
        parsed = json.loads(value)
    except json.JSONDecodeError:
        bad_json_rows += 1
        parsed_payloads.append(None)
        continue
    if not isinstance(parsed, list):
        bad_json_rows += 1
        parsed_payloads.append(None)
        continue
    parsed_payloads.append(parsed)

if bad_json_rows:
    raise ValueError(f"Found {bad_json_rows} malformed segment_preds
        rows")

if require_primary_for_nonempty:
    blank_primary_violations = 0
    for primary_name, payload in zip(df["RFV1_name"], parsed_payloads):
        primary_blank = pd.isna(primary_name) or not
        str(primary_name).strip()
        has_segments = isinstance(payload, list) and len(payload) > 0
        if primary_blank and has_segments:
            blank_primary_violations += 1

    if blank_primary_violations:
        raise ValueError(
            "Found blank RFV1_name for rows with non-empty segment
            predictions: "
            f"{blank_primary_violations}"
        )

print(f"Schema validation passed for {len(df)} rows.")

```

```

def save_with_suffix(
    df: pd.DataFrame,
    working_path: str | Path,
    suffix: str = "_with_NLP",
    sheet_name: str = "with_NLP",
) -> Path:
    """Write workbook with suffix and return the new path."""
    path = Path(working_path)
    output_path = path.with_name(path.stem + suffix + path.suffix)
    df.to_excel(output_path, index=False, sheet_name=sheet_name)
    print(f"Wrote: {output_path}")
    return output_path

def summarize_overrides(df: pd.DataFrame, column: str = "segment_preds") ->
    dict[str, Any]:
    """Summarize rule-override usage from serialized segment prediction
    payloads."""
    if column not in df.columns:
        raise KeyError(f"Missing column: {column}")

    segment_total = 0
    segment_overridden = 0
    by_rule: Counter = Counter()
    by_final: Counter = Counter()
    rows_total = 0
    rows_with_override = 0

    for value in df[column].dropna():
        rows_total += 1
        payload: Any
        if isinstance(value, str):
            try:
                payload = json.loads(value)
            except Exception:
                continue
        elif isinstance(value, list):
            payload = value
        else:
            continue

        if not isinstance(payload, list):
            continue

```

```

row_has_override = False
for segment_prediction in payload:
    if not isinstance(segment_prediction, dict):
        continue
    segment_total += 1
    if segment_prediction.get("rule_used"):
        segment_overridden += 1
        row_has_override = True
    by_rule[segment_prediction.get("rule_code", "UNKNOWN")] += 1
    by_final[segment_prediction.get("code", "UNKNOWN")] += 1

if row_has_override:
    rows_with_override += 1

return {
    "segments_total": segment_total,
    "segments_overridden": segment_overridden,
    "segment_override_rate": (segment_overridden / segment_total) if
        ↵ segment_total else float("nan"),
    "rows_total": rows_total,
    "rows_with_override": rows_with_override,
    "row_override_rate": (rows_with_override / rows_total) if rows_total
        ↵ else float("nan"),
    "by_rule": by_rule,
    "by_final": by_final,
}

```

```

def validate_pipeline(
    df: pd.DataFrame,
    cache: SegmentEmbedCache,
    resources: dict[str, Any],
    seg_col: str = "cc_cleaned",
) -> dict[str, Any]:
    """Run smoke checks for prototype health, segment coverage, and
    ↵ rule/classifier sanity."""
    metadata = resources["metadata"]
    prototype_counts = metadata["prototype_count_by_group"]
    missing_groups = [group for group, count in prototype_counts.items() if
        ↵ count == 0]
    if missing_groups:
        raise ValueError(f"Groups missing prototypes: {missing_groups}")

```

```

row_count = len(df)
nonempty_segment_rows = int(df[seg_col].apply(lambda value:
↪ isinstance(value, list) and len(value) > 0).sum())
unique_segments = len(cache.uniq_segments)

probes = [
    "shortness of breath",
    "bradycardia",
    "positive blood culture",
    "alter mental status",
    "abdominal pain / shortness of breath",
    "right eye redness / pain",
]
probe_results: dict[str, list[tuple[str, str, float, bool]]] = {}
for probe in probes:
    segments = segment_cc(probe)
    predictions = classify_segments_cached(
        segments,
        cache=cache,
        resources=resources,
        method="max",
    )
    probe_results[probe] = [
        (prediction.segment, prediction.code, prediction.sim,
↪ prediction.rule_used)
        for prediction in predictions
    ]

print("Prototype integrity:")
print(prototype_counts)
print(f"Total prototypes: {metadata['prototype_total']}") 
print()
print("Segment column health:")
print(f"rows={row_count:,} | with>=1 segment={nonempty_segment_rows:,} |
↪ unique segments={unique_segments:,}")
print()
print("Probe classifications:")
for probe, result in probe_results.items():
    print(f"{probe} -> {result}")

return {
    "row_count": row_count,
}

```

```

    "rows_with_segments": nonempty_segment_rows,
    "unique_segments": unique_segments,
    "prototype_counts": prototype_counts,
    "probe_results": probe_results,
}

def top_cc_plot(
    df: pd.DataFrame,
    candidates: tuple[str, ...] = CC_CANDIDATES,
    top_n: int = 30,
    normalize_case: str | None = "upper",
) -> pd.DataFrame:
    """Optional EDA: plot top raw chief complaints."""
    cc_column = get_cc_column(df)
    series = (
        df[cc_column]
        .astype(str)
        .str.strip()
        .str.replace(r"\s+", " ", regex=True)
    )

    if normalize_case == "upper":
        series = series.str.upper()
    elif normalize_case == "title":
        series = series.str.title()

    counts = series.value_counts().head(top_n).sort_values(ascending=True)
    axis = counts.plot(kind="barh", figsize=(8, 10))
    axis.set_xlabel("Count")
    axis.set_ylabel("Chief complaint")
    axis.set_title(f"Top {top_n} Chief Complaints ({cc_column})")
    plt.tight_layout()
    plt.show()

    return counts.sort_values(ascending=False).rename_axis("chief_complaint"]
    ↵    ).reset_index(name="count")

```

1.3 2B) Local helper functions (embedded core notebook logic)

Core stage helpers are embedded in this notebook by project policy.

```
DATA_DIRNAME = "MIMIC tabular data"
```

```

PRIOR_RUNS_DIRNAME = "prior runs"

CANONICAL_COHORT_FILENAME = "MIMICIV all with CC.xlsx"
CANONICAL_NLP_FILENAME = "MIMICIV all with CC_with_NLP.xlsx"

CLASSIFIER_TRANSITIONAL_ALIASES: dict[str, tuple[str, ...]] = {
    "age": ("age_at_admit",),
    "hr": ("ed_first_hr_model", "ed_first_hr", "ed_triage_hr_model",
           "ed_triage_hr"),
    "rr": ("ed_first_rr_model", "ed_first_rr", "ed_triage_rr_model",
           "ed_triage_rr"),
    "sbp": ("ed_first_sbp_model", "ed_first_sbp", "ed_triage_sbp_model",
             "ed_triage_sbp"),
    "dbp": ("ed_first_dbp_model", "ed_first_dbp", "ed_triage_dbp_model",
             "ed_triage_dbp"),
    "temp": ("ed_first_temp_model", "ed_first_temp", "ed_triage_temp_model",
              "ed_triage_temp"),
    "spo2": (
        "ed_first_o2sat_model",
        "ed_first_o2sat",
        "ed_triage_o2sat_model",
        "ed_triage_o2sat",
    ),
    "race": ("race_ed_raw",),
}
}

def data_dir(work_dir: Path) -> Path:
    """Return canonical data directory rooted at ``work_dir``."""
    return (work_dir / DATA_DIRNAME).expanduser().resolve()

def resolve_classifier_input_path(
    work_dir: Path, input_filename: str | None = None
) -> Path:
    """Resolve classifier input workbook under canonical data directory."""
    filename = input_filename or CANONICAL_COHORT_FILENAME
    input_path = data_dir(work_dir) / filename
    if not input_path.exists():
        raise FileNotFoundError(
            f"Expected classifier input workbook was not found at "
            f"{input_path}. Run the cohort notebook first or set "
            "CLASSIFIER_INPUT_FILENAME."

```

```

        )
    return input_path

def resolve_analysis_input_path(work_dir: Path, input_filename: str | None =
    None) -> Path:
    """Resolve analysis input workbook under canonical data directory."""
    filename = input_filename or CANONICAL_NLP_FILENAME
    input_path = data_dir(work_dir) / filename
    if not input_path.exists():
        raise FileNotFoundError(
            "Expected analysis input workbook was not found at "
            f"{input_path}. Run the classifier notebook first or set "
            "ANALYSIS_INPUT_FILENAME."
        )
    return input_path

def resolve_rater_nlp_input_path(work_dir: Path, input_filename: str | None =
    None) -> Path:
    """Resolve rater notebook NLP input workbook under canonical data
    directory."""
    filename = input_filename or CANONICAL_NLP_FILENAME
    input_path = data_dir(work_dir) / filename
    if not input_path.exists():
        raise FileNotFoundError(
            "Expected rater NLP input workbook was not found at "
            f"{input_path}. Run the classifier notebook first or set "
            "RATER_NLP_INPUT_FILENAME."
        )
    return input_path

def resolve_classifier_output_paths(
    work_dir: Path,
    *,
    run_dt: datetime | None = None,
) -> tuple[Path, Path]:
    """Return canonical NLP output path and dated archive path."""
    base_data_dir = data_dir(work_dir)
    canonical_path = base_data_dir / CANONICAL_NLP_FILENAME

    run_date = (run_dt or datetime.now()).strftime("%Y-%m-%d")

```

```

archive_dir = base_data_dir / PRIOR_RUNS_DIRNAME
archive_path = archive_dir / f"{run_date} {CANONICAL_NLP_FILENAME}"

return canonical_path, archive_path


def normalize_classifier_input_schema(
    df: pd.DataFrame,
    alias_map: Mapping[str, str | Sequence[str]] | None = None,
) -> pd.DataFrame:
    """Add transitional alias columns for classifier compatibility.

    Existing destination columns are preserved. Source columns are never
    removed.
    """
    resolved_alias_map = alias_map or CLASSIFIER_TRANSITIONAL_ALIASES
    normalized = df.copy()
    for destination, source_spec in resolved_alias_map.items():
        if destination in normalized.columns:
            continue
        if isinstance(source_spec, str):
            candidate_sources = (source_spec,)
        else:
            candidate_sources = tuple(source_spec)
        for source_name in candidate_sources:
            if source_name in normalized.columns:
                normalized[destination] = normalized[source_name]
                break
    return normalized


def ensure_required_columns(
    df: pd.DataFrame,
    required: list[str],
    *,
    context: str,
) -> None:
    """Validate required columns with context-aware error message."""
    missing = sorted(set(required).difference(df.columns))
    if missing:
        raise KeyError(f"{context}: missing required columns: {missing}")

PSEUDO_MISSING_EXACT_TOKENS = {

```

```

    "-",
    "--",
    "-",
    "n",
    "na",
    "n/a",
    "none",
    "unknown",
    "refused",
}
PSEUDO_MISSING_REGEX = re.compile(r"^-+$")
PSEUDO_MISSING_PUNCT_REGEX = re.compile(r"^[^a-z0-9]+$")

def _canonicalize_cc_text(value: object) -> str:
    text = "" if pd.isna(value) else str(value)
    text = text.strip().lower()
    text = text.replace("-", "-").replace("–", "–")
    return " ".join(text.split())

def _to_numeric_series(df: pd.DataFrame, column_name: str) -> pd.Series:
    if column_name in df.columns:
        return pd.to_numeric(df[column_name], errors="coerce")
    return pd.Series(math.nan, index=df.index)

def _to_text_series(df: pd.DataFrame, column_name: str) -> pd.Series:
    if column_name in df.columns:
        return df[column_name].astype("string")
    return pd.Series(pd.NA, index=df.index, dtype="string")

def _to_binary_indicator(df: pd.DataFrame, column_name: str) -> pd.Series:
    if column_name not in df.columns:
        return pd.Series(pd.NA, index=df.index, dtype="Float64")
    numeric = pd.to_numeric(df[column_name], errors="coerce")
    return (numeric > 0).astype("Float64").where(numeric.notna(), pd.NA)

def add_hypcapnia_flags(
    df: pd.DataFrame,
    *,

```

```

art_col: str = "first_abg_pco2",
art_uom_col: str = "first_abg_pco2_uom",
vbg_col: str = "first_vbg_pco2",
vbg_uom_col: str = "first_vbg_pco2_uom",
fallback_art_col: str = "poc_abg_paco2",
fallback_art_uom_col: str = "poc_abg_paco2_uom",
fallback_vbg_col: str = "poc_vbg_paco2",
fallback_vbg_uom_col: str = "poc_vbg_paco2_uom",
authoritative_abg_col: str = "flag_abg_hypercapnia",
authoritative_vbg_col: str = "flag_vbg_hypercapnia",
out_abg: str = "hypercap_by_abg",
out_vbg: str = "hypercap_by_vbg",
out_any: str = "hypercap_by_bg",
) -> pd.DataFrame:
    """Add ABG/VBG hypercapnia flags using cohort-schema defaults.

    Precedence is applied per-row:
    1) authoritative route flags from cohort (`flag_*_hypercapnia`) when
    ↵ present,
    2) first-route pCO2 values (`first_abg_pco2` / `first_vbg_pco2`),
    3) fallback POC pCO2 values.
    """
    def to_mmhg(values: pd.Series, uoms: pd.Series) -> pd.Series:
        numeric = pd.to_numeric(values, errors="coerce")
        unit_text = uoms.astype("string").str.strip().str.lower()
        is_kpa = unit_text.str.contains("kpa", na=False)
        converted = numeric.copy()
        converted.loc[is_kpa] = converted.loc[is_kpa] * 7.50062
        return converted

    def derive_from_value(
        *,
        value_col: str,
        uom_col: str,
        threshold: float,
    ) -> pd.Series:
        values = _to_numeric_series(out, value_col)
        units = _to_text_series(out, uom_col)
        mmhg = to_mmhg(values, units)
        derived = (mmhg >= threshold).astype("Float64")
        return derived.where(mmhg.notna(), pd.NA)

```

```

out = df.copy()
abg_flag = _to_binary_indicator(out, authoritative_abg_col)
if abg_flag.isna().any():
    primary_abg = derive_from_value(
        value_col=art_col,
        uom_col=art_uom_col,
        threshold=45.0,
    )
    abg_flag = abg_flag.where(abg_flag.notna(), primary_abg)
if abg_flag.isna().any():
    fallback_abg = derive_from_value(
        value_col=fallback_art_col,
        uom_col=fallback_art_uom_col,
        threshold=45.0,
    )
    abg_flag = abg_flag.where(abg_flag.notna(), fallback_abg)

vbg_flag = _to_binary_indicator(out, authoritative_vbg_col)
if vbg_flag.isna().any():
    primary_vbg = derive_from_value(
        value_col=vbg_col,
        uom_col=vbg_uom_col,
        threshold=50.0,
    )
    vbg_flag = vbg_flag.where(vbg_flag.notna(), primary_vbg)
if vbg_flag.isna().any():
    fallback_vbg = derive_from_value(
        value_col=fallback_vbg_col,
        uom_col=fallback_vbg_uom_col,
        threshold=50.0,
    )
    vbg_flag = vbg_flag.where(vbg_flag.notna(), fallback_vbg)

out[out_abg] = abg_flag.fillna(0).astype("int8")
out[out_vbg] = vbg_flag.fillna(0).astype("int8")
out[out_any] = ((out[out_abg] == 1) | (out[out_vbg] == 1)).astype("int8")
return out

def annotate_cc_missingness(
    df: pd.DataFrame,
    *,
    cc_column: str,

```

```

nlp_input_column: str = "cc_text_for_nlp",
missing_flag_column: str = "cc_missing_flag",
pseudo_flag_column: str = "cc_pseudomissing_flag",
missing_reason_column: str = "cc_missing_reason",
) -> pd.DataFrame:
    """Annotate true-vs-pseudo missing CCs and derive deterministic NLP input
    ↵ text."""
    if cc_column not in df.columns:
        raise KeyError(f"Chief complaint column not found: {cc_column}")

    out = df.copy()
    canonical = out[cc_column].map(_canonicalize_cc_text)
    is_true_missing = canonical.eq("")
    is_pseudo_missing = (
        canonical.isin(PSEUDO_MISSING_EXACT_TOKENS)
        | canonical.str.match(PSEUDO_MISSING_REGEX, na=False)
        | canonical.str.match(PSEUDO_MISSING_PUNCT_REGEX, na=False)
    )

    reason = pd.Series("valid", index=out.index, dtype="string")
    reason.loc[is_true_missing] = "true_missing"
    reason.loc[is_pseudo_missing] = "pseudo_missing_token"

    out[missing_flag_column] = (is_true_missing |
    ↵ is_pseudo_missing).astype("boolean")
    out[pseudo_flag_column] = is_pseudo_missing.astype("boolean")
    out[missing_reason_column] = reason
    out[nlp_input_column] = out[cc_column].astype("string")
    out.loc[is_true_missing | is_pseudo_missing, nlp_input_column] = pd.NA
    return out

def apply_pseudomissing_uncodable_policy(
    df: pd.DataFrame,
    *,
    max_rfv: int = 5,
    pseudo_flag_column: str = "cc_pseudomissing_flag",
) -> pd.DataFrame:
    """Force pseudo-missing CC rows to a deterministic uncodable RFV
    ↵ payload."""
    if pseudo_flag_column not in df.columns:
        return df

```

```

out = df.copy()
pseudo_mask = out[pseudo_flag_column].fillna(False).astype(bool)
if not pseudo_mask.any():
    return out

payload = json.dumps(
    [
        {
            "seg_idx": 0,
            "segment": "",
            "code": "RVC-UNCL",
            "name": "Uncodable/Unknown",
            "sim": None,
            "rule_code": "pseudo_missing_cc",
            "rule_used": True,
        }
    ]
)
_apply_uncodable_payload(
    out,
    mask=pseudo_mask,
    max_rfv=max_rfv,
    payload=payload,
)
return out

def apply_blank_primary_uncodable_policy(
    df: pd.DataFrame,
    *,
    max_rfv: int = 5,
    missing_reason_column: str = "cc_missing_reason",
) -> pd.DataFrame:
    """Force non-true-missing blank RFV primaries to deterministic uncodable
    output."""
    if "RFV1_name" not in df.columns:
        return df

    out = df.copy()
    blank_primary =
    out["RFV1_name"].fillna("").astype(str).str.strip().eq("")
    true_missing = (
        out[missing_reason_column].astype(str).eq("true_missing")

```

```

        if missing_reason_column in out.columns
            else pd.Series(False, index=out.index)
    )
fallback_mask = blank_primary & ~true_missing
if not fallback_mask.any():
    return out

payload = json.dumps(
    [
        {
            "seg_idx": 0,
            "segment": "",
            "code": "RVC-UNCL",
            "name": "Uncodable/Unknown",
            "sim": None,
            "rule_code": "blank_primary_fallback",
            "rule_used": True,
        }
    ]
)
_apply_uncodable_payload(
    out,
    mask=fallback_mask,
    max_rfsv=max_rfsv,
    payload=payload,
)
return out

def _apply_uncodable_payload(
    out: pd.DataFrame,
    *,
    mask: pd.Series,
    max_rfsv: int,
    payload: str,
) -> None:
    """Apply a deterministic uncodable RFV payload in-place."""
    out.loc[mask, "RFV1"] = "uncodable"
    out.loc[mask, "RFV1_name"] = "Uncodable/Unknown"
    out.loc[mask, "RFV1_support"] = ""
    out.loc[mask, "RFV1_sim"] = math.nan

    for slot in range(2, max_rfsv + 1):

```

```

        out.loc[mask, f"RFV{slot}"] = ""
        out.loc[mask, f"RFV{slot}_name"] = ""
        out.loc[mask, f"RFV{slot}_support"] = ""
        out.loc[mask, f"RFV{slot}_sim"] = math.nan

    out.loc[mask, "segment_preds"] = payload

def build_cc_missing_audit(
    df: pd.DataFrame,
    *,
    cc_column: str,
    missing_reason_column: str = "cc_missing_reason",
    max_examples: int = 5,
) -> pd.DataFrame:
    """Summarize CC missingness categories with example raw values."""
    if cc_column not in df.columns:
        raise KeyError(f"Chief complaint column not found: {cc_column}")
    if missing_reason_column not in df.columns:
        raise KeyError(f"Missing reason column not found:
                       {missing_reason_column}")

    frame = df[[cc_column, missing_reason_column]].copy()
    frame[cc_column] = frame[cc_column].astype("string")
    rows: list[dict[str, Any]] = []
    for reason, group in frame.groupby(missing_reason_column, dropna=False):
        examples = (
            group[cc_column]
            .fillna("<NA>")
            .astype(str)
            .value_counts(dropna=False)
            .head(max_examples)
            .index.tolist()
        )
        rows.append(
            {
                "cc_missing_reason": str(reason),
                "row_count": int(len(group)),
                "examples": "; ".join(examples),
            }
        )
    return pd.DataFrame(rows).sort_values("cc_missing_reason")
        .reset_index(drop=True)

```

```

def build_phrase_regression_audit(
    df: pd.DataFrame,
    *,
    cc_column: str,
    segment_preds_column: str = "segment_preds",
    rfv_column: str = "RFV1_name",
) -> tuple[pd.DataFrame, pd.DataFrame]:
    """Audit phrase-level normalization regressions for abnormal-lab
    ↵ complaints."""
    if cc_column not in df.columns:
        raise KeyError(f"Chief complaint column not found: {cc_column}")
    if segment_preds_column not in df.columns:
        raise KeyError(f"segment prediction column not found:
            ↵ {segment_preds_column}")
    if rfv_column not in df.columns:
        raise KeyError(f"RFV column not found: {rfv_column}")

    raw_cc = df[cc_column].fillna("").astype(str)
    abnormal_lab_mask = raw_cc.str.contains(r"\babnormal\s+labs?\b",
        ↵ case=False, regex=True)
    subset = df.loc[abnormal_lab_mask, [cc_column, segment_preds_column,
        ↵ rfv_column]].copy()

    abnormal_loss_rows = 0
    for payload_text in subset[segment_preds_column].fillna("[]"):
        try:
            payload = json.loads(payload_text) if isinstance(payload_text,
        ↵ str) else []
        except json.JSONDecodeError:
            payload = []
        row_segments: list[str] = []
        for record in payload if isinstance(payload, list) else []:
            if not isinstance(record, dict):
                continue
            segment_text = str(record.get("segment", "")).strip().lower()
            if not segment_text:
                continue
            row_segments.append(segment_text)
        if any("abnormal loss" in segment for segment in row_segments):
            abnormal_loss_rows += 1

```

```

rfv_counts = (
    subset[rfv_column]
    .fillna("NA")
    .astype(str)
    .value_counts(dropna=False)
    .rename_axis("rfv1_name")
    .reset_index(name="row_count")
)
if not rfv_counts.empty:
    rfv_counts["row_rate"] = rfv_counts["row_count"] / max(len(subset),
    ↪ 1)

summary = pd.DataFrame(
    [
        {"metric": "abnormal_lab_rows", "value": int(len(subset))},
        {"metric": "abnormal_loss_rows", "value": int(abnormal_loss_rows)},
        {
            "metric": "abnormal_test_result_rows",
            "value": int(
                subset[rfv_column]
                .fillna("")
                .astype(str)
                .str.contains(r"Abnormal test result", case=False,
                ↪ regex=True)
                .sum()
            ),
        },
        {
            "metric": "eye_ear_rows",
            "value": int(
                subset[rfv_column]
                .fillna("")
                .astype(str)
                .str.contains(r"Eye/Ear", case=False, regex=True)
                .sum()
            ),
        },
    ],
)
return summary, rfv_counts

```

```

def classifier_resource_paths(
    work_dir: Path,
    *,
    appendix_relpah: str,
    summary_relpah: str,
) -> dict[str, Path]:
    """Return required classifier resource paths under ``work_dir``."""
    return {
        "appendix_csv": (work_dir / appendix_relpah).expanduser().resolve(),
        "summary_weights_csv": (work_dir /
            summary_relpah).expanduser().resolve(),
    }

def _sha256(path: Path) -> str:
    hasher = hashlib.sha256()
    with path.open("rb") as handle:
        for chunk in iter(lambda: handle.read(8192), b ""):
            hasher.update(chunk)
    return hasher.hexdigest()

def verify_classifier_resources(
    work_dir: Path,
    *,
    appendix_relpah: str,
    summary_relpah: str,
    manifest_path: str | None = "Annotation/resource_manifest.json",
    strict_hash: bool = True,
) -> dict[str, Any]:
    """Validate required classifier resources and optional hash manifest."""
    checks: list[dict[str, Any]] = []
    missing_paths: list[str] = []
    for name, path in classifier_resource_paths(
        work_dir,
        appendix_relpah=appendix_relpah,
        summary_relpah=summary_relpah,
    ).items():
        exists = path.exists()
        checks.append({"resource": name, "path": str(path), "exists": exists})
        if not exists:
            missing_paths.append(str(path))

```

```

hash_findings: list[dict[str, Any]] = []
manifest_resolved: Path | None = None
if manifest_path:
    manifest_resolved = (work_dir / manifest_path).expanduser().resolve()
if manifest_resolved.exists():
    manifest = json.loads(manifest_resolved.read_text())
    resources = manifest.get("resources", [])
    for entry in resources:
        rel_path = entry.get("path")
        expected = str(entry.get("sha256", "")).strip().lower()
        if not rel_path or not expected:
            continue
        path = (work_dir / rel_path).expanduser().resolve()
        if not path.exists():
            hash_findings.append(
                {
                    "path": str(path),
                    "status": "missing",
                    "expected_sha256": expected,
                    "actual_sha256": None,
                }
            )
            continue
        actual = _sha256(path)
        if actual != expected:
            hash_findings.append(
                {
                    "path": str(path),
                    "status": "hash_mismatch",
                    "expected_sha256": expected,
                    "actual_sha256": actual,
                }
            )
errors: list[str] = []
warnings: list[str] = []

if missing_paths:
    errors.append(
        "Missing required classifier resources: " + "",
        ← ".join(sorted(missing_paths))
    )

```

```

for finding in hash_findings:
    message = (
        f"{{finding['status']}} for {{finding['path']}} "
        f"(expected={{finding['expected_sha256']}}, "
        f"actual={{finding['actual_sha256']}})"
    )
    if strict_hash:
        errors.append(message)
    else:
        warnings.append(message)

status = "pass"
if errors:
    status = "fail"
elif warnings:
    status = "warning"

return {
    "status": status,
    "strict_hash": bool(strict_hash),
    "manifest_path": str(manifest_resolved) if manifest_resolved else
        None,
    "checks": checks,
    "hash_findings": hash_findings,
    "errors": errors,
    "warnings": warnings,
}

def validate_classifier_contract(
    df: pd.DataFrame,
    *,
    max_rfv: int = 5,
    pseudo_flag_column: str = "cc_pseudomissing_flag",
    missing_reason_column: str = "cc_missing_reason",
    missing_flag_column: str = "cc_missing_flag",
    authoritative_disagreement_tolerance: float = 0.0,
) -> dict[str, Any]:
    """Run deterministic post-export classifier contract checks."""
    findings: list[dict[str, str]] = []

    if "hadm_id" in df.columns:

```

```

duplicates = int(df.duplicated(subset=["hadm_id"]).sum())
if duplicates:
    findings.append(
        {
            "severity": "error",
            "code": "hadm_id_not_unique",
            "message": f"Found {duplicates} duplicate hadm_id rows.",
        }
    )

if "segment_preds" not in df.columns:
    findings.append(
        {
            "severity": "error",
            "code": "missing_segment_preds",
            "message": "segment_preds column missing from classifier
                        ↵ output.",
        }
    )
else:
    bad_json_rows = 0
    for value in df["segment_preds"].fillna("[]"):
        try:
            parsed = json.loads(value) if isinstance(value, str) else
            ↵ value
        except json.JSONDecodeError:
            bad_json_rows += 1
            continue
        if not isinstance(parsed, list):
            bad_json_rows += 1
    if bad_json_rows:
        findings.append(
            {
                "severity": "error",
                "code": "segment_preds_malformed",
                "message": f"Found {bad_json_rows} malformed
                            ↵ segment_preds rows.",
            }
        )

if {"hypercap_by_abg", "hypercap_by_vbg",
    ↵ "hypercap_by_bg"}.issubset(df.columns):

```

```

abg = pd.to_numeric(df["hypercap_by_abg"], 
↪ errors="coerce").fillna(0).astype(int)
vbg = pd.to_numeric(df["hypercap_by_vbg"], 
↪ errors="coerce").fillna(0).astype(int)
bg = pd.to_numeric(df["hypercap_by_bg"], 
↪ errors="coerce").fillna(0).astype(int)
mismatch = int(((abg | vbg) != bg).sum())
if mismatch:
    findings.append(
    {
        "severity": "error",
        "code": "bg_union_mismatch",
        "message": f"hypercap_by_bg mismatch in {mismatch} 
↪ rows.",
    }
)
abg_sum = int(abg.sum())
abg_authoritative_positive = 0
if "flag_abg_hypercapnia" in df.columns:
    abg_authoritative_numeric = pd.to_numeric(
        df["flag_abg_hypercapnia"], errors="coerce"
    )
    abg_authoritative_positive = int(
        (abg_authoritative_numeric.fillna(0).astype(int) > 0).sum()
    )
if abg_sum == 0:
    if abg_authoritative_positive > 0:
        findings.append(
        {
            "severity": "error",
            "code": "abg_all_zero_with_authoritative_positive",
            "message": (
                "hypercap_by_abg has zero positives despite
                ↪ non-zero "
                f"flag_abg_hypercapnia positives
                ↪ ({abg_authoritative_positive})."
            ),
        }
    )
else:
    findings.append(
    {
        "severity": "warning",

```

```

        "code": "abg_all_zero",
        "message": "hypercap_by_abg has zero positives.",
    }
)
if "flag_abg_hypcapnia" in df.columns:
    abg_authoritative = pd.to_numeric(
        df["flag_abg_hypcapnia"], errors="coerce"
    )
    valid_abg = abg_authoritative.notna()
    if bool(valid_abg.any()):
        authoritative_binary = (
            abg_authoritative.loc[valid_abg].fillna(0).astype(int) >
        ← 0
            ).astype(int)
        disagreement = int((abg.loc[valid_abg] !=
    ← authoritative_binary).sum())
        rate = float(disagreement / int(valid_abg.sum()))
        if rate > authoritative_disagreement_tolerance:
            findings.append(
                {
                    "severity": "error",
                    "code": "abg_authoritative_disagreement",
                    "message": (
                        "hypercap_by_abg disagrees with "
                        ← flag_abg_hypcapnia in "
                        f"{disagreement} rows ({rate:.4%}), "
                        ← tolerance="
                        f"{authoritative_disagreement_tolerance:.4%}]"
                        ← .
                    ),
                }
            )
        if "flag_vbg_hypcapnia" in df.columns:
            vbg_authoritative = pd.to_numeric(
                df["flag_vbg_hypcapnia"], errors="coerce"
            )
            valid_vbg = vbg_authoritative.notna()
            if bool(valid_vbg.any()):
                authoritative_binary = (
                    vbg_authoritative.loc[valid_vbg].fillna(0).astype(int) >
    ← 0
                    ).astype(int)

```

```

        disagreement = int((vbg.loc[valid_vbg] !=
↪ authoritative_binary).sum())
        rate = float(disagreement / int(valid_vbg.sum()))
        if rate > authoritative_disagreement_tolerance:
            findings.append(
                {
                    "severity": "error",
                    "code": "vbg_authoritative_disagreement",
                    "message": (
                        "hypercap_by_vbg disagrees with"
                        "flag_vbg_hypercapnia in "
                        f"{disagreement} rows ({rate:.4%}),"
                        "tolerance="
                        f"{authoritative_disagreement_tolerance:.4%}"
                        "."
                    ),
                }
            )
        )

required_cc_columns = {missing_reason_column, pseudo_flag_column,
↪ missing_flag_column}
missing_cc_columns = sorted(required_cc_columns.difference(df.columns))
if missing_cc_columns:
    findings.append(
        {
            "severity": "error",
            "code": "missing_cc_missingness_columns",
            "message": f"Missing CC missingness columns:
↪ {missing_cc_columns}",
        }
    )

if "RFV1_name" in df.columns:
    blank_rfv1 =
↪ df["RFV1_name"].fillna("").astype(str).str.strip().eq("")
    if missing_reason_column in df.columns:
        allowed_blank =
↪ df[missing_reason_column].astype(str).eq("true_missing")
        unexpected_blank = int((blank_rfv1 & ~allowed_blank).sum())
    else:
        unexpected_blank = int(blank_rfv1.sum())
    if unexpected_blank:
        findings.append(

```

```

    {
        "severity": "error",
        "code": "unexpected_blank_rf1",
        "message": (
            "Found blank RFV1_name rows outside allowed
             ↳ true-missing policy: "
            f"{{unexpected_blank}}"
        ),
    }
)

pseudo_count = 0
if pseudo_flag_column in df.columns:
    pseudo_count = int(df[pseudo_flag_column].fillna(False).sum())

error_count = sum(1 for finding in findings if finding["severity"] ==
← "error")
warning_count = sum(1 for finding in findings if finding["severity"] ==
← "warning")
status = "pass"
if error_count:
    status = "fail"
elif warning_count:
    status = "warning"

return {
    "status": status,
    "error_count": error_count,
    "warning_count": warning_count,
    "pseudo_missing_rows": pseudo_count,
    "findings": findings,
    "row_count": int(len(df)),
    "column_count": int(df.shape[1]),
    "max_rf1": int(max_rf1),
}
}

def render_latex_longtable(
    table_df: pd.DataFrame,
    *,
    caption: str,
    label: str,
    landscape: bool = False,
    index: bool = True,

```

```

) -> str:
    latex_text = table_df.to_latex(
        index=index,
        escape=False,
        longtable=True,
        caption=caption,
        label=label,
    )
    if landscape:
        latex_text = "\\begin{landscape}\\n" + latex_text +
    ↵ "\\n\\end{landscape}\\n"
    return latex_text

```

1.4 3) Input load and schema checks

```

# Purpose: instantiate config, set deterministic seeds, load the input
    ↵ workbook, and validate required identifiers.
# Why this matters: stable setup reduces run-to-run variability and catches
    ↵ input issues before expensive model steps.

cfg = ClassifierConfig()

try:
    NLP = spacy.load("en_core_web_sm", disable=["ner", "parser", "textcat",
    ↵ "senter"])
    SPACY_MODEL_LABEL = "en_core_web_sm"
except OSError as exc:
    if cfg.require_spacy_model:
        raise ImportError(
            "spaCy model 'en_core_web_sm' is required but not installed in
            ↵ this environment."
            "Install it once with: python -m spacy download en_core_web_sm"
        ) from exc
    NLP = spacy.blank("en")
    SPACY_MODEL_LABEL = "spacy.blank('en')"

print(f"spaCy pipeline: {SPACY_MODEL_LABEL}")

np.random.seed(cfg.seed)
torch.manual_seed(cfg.seed)
if torch.cuda.is_available():
    torch.cuda.manual_seed_all(cfg.seed)

```

```

load_dotenv()
WORK_DIR = Path(os.getenv(cfg.work_dir_env_var,
    ↪ Path.cwd()).expanduser().resolve())
SRC_DIR = WORK_DIR / "src"
if str(SRC_DIR) not in sys.path:
    sys.path.insert(0, str(SRC_DIR))

from hypercap_cc_nlp.pipeline_contracts import write_contract_report
from hypercap_cc_nlp.pipeline_audit import collect_run_manifest
ensure_required_workflow_columns = ensure_required_columns

classifier_input_filename = os.getenv(cfg.input_filename_env_var,
    ↪ cfg.input_filename)
file_path = resolve_classifier_input_path(WORK_DIR,
    ↪ classifier_input_filename)

strict_resource_hash = os.getenv("CLASSIFIER_STRICT_RESOURCE_HASH",
    ↪ "1").strip() == "1"
resource_report = verify_classifier_resources(
    WORK_DIR,
    appendix_relpah=cfg.appendix_relpah,
    summary_relpah=cfg.summary_relpah,
    strict_hash=strict_resource_hash,
)
if resource_report["status"] == "fail":
    raise FileNotFoundError(
        "Classifier resource verification failed: "
        + "; ".join(resource_report["errors"]))
if resource_report["warnings"]:
    print("Resource verification warnings:")
    for warning_message in resource_report["warnings"]:
        print("-", warning_message)

df = pd.read_excel(file_path, sheet_name=0, engine="openpyxl")
df.columns = df.columns.str.strip()
pre_alias_columns = set(df.columns)
df = normalize_classifier_input_schema(df)
added_aliases = sorted(set(df.columns).difference(pre_alias_columns))

ensure_required_workflow_columns(
    df,

```

```

        ["hadm_id", "subject_id"],
        context="classifier input",
    )

print(f"Loaded input rows={len(df):,}, cols={df.shape[1]:,}")
print(f"Input path: {file_path}")
print(f"Added transitional aliases: {added_aliases if added_aliases else
    'None'}")
cc_column = get_cc_column(df)
print(f"Detected chief complaint column: {cc_column}")

spaCy pipeline: en_core_web_sm
Loaded input rows=12,020, cols=301
Input path: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Pro
CC-NLP/MIMIC tabular data/MIMICIV all with CC.xlsx
Added transitional aliases: ['age', 'dbp', 'hr', 'race', 'rr', 'sbp', 'spo2', 'temp']
Detected chief complaint column: ed_triage_cc

```

1.5 4) Text cleaning and segmentation functions

```

# Purpose: define data-enrichment helpers and generate normalized
    ↵ chief-complaint text used by the classifier.
# Why this matters: high-quality text normalization is foundational for both
    ↵ rule matching and embedding similarity.

df = add_hypcapnia_flags(df)
if {"flag_abg_hypcapnia", "hypcap_by_abg"}.issubset(df.columns):
    authoritative_abg_positive = int(
        pd.to_numeric(df["flag_abg_hypcapnia"]),
    ↵ errors="coerce").fillna(0).astype(int) > 0).sum()
    )
    derived_abg_positive = int(
        pd.to_numeric(df["hypcap_by_abg"]),
    ↵ errors="coerce").fillna(0).astype(int).sum()
    )
    if authoritative_abg_positive > 0 and derived_abg_positive == 0:
        raise ValueError(
            "ABG hypcapnia derivation failure: flag_abg_hypcapnia has
                ↵ positives "
            f"({authoritative_abg_positive}) but hypcap_by_abg is all
                ↵ zero."
        )

```

```

df = annotate_cc_missingness(df, cc_column=cc_column)
requested_spell_mode = os.getenv("CC_SPELL_CORRECTION_MODE",
    "auto").strip().lower()
normalization_payload, spell_mode_scores = choose_spell_mode(
    df,
    cc_source_column="cc_text_for_nlp",
    requested_mode=requested_spell_mode,
)
df = pd.concat([df, normalization_payload["frame"]], axis=1)
df["cc_spell_mode_selected"] = normalization_payload["selected_mode"]

cc_missing_audit = build_cc_missing_audit(df, cc_column=cc_column)
print("CC missingness audit:")
print(cc_missing_audit.to_string(index=False))
print("Pseudo-missing rows:",
    int(df["cc_pseudomissing_flag"].fillna(False).sum()))
print("Spell correction mode requested:",
    normalization_payload["requested_mode"])
print("Spell correction mode selected:",
    normalization_payload["selected_mode"])
print("Spell mode integrity comparison:")
print(spell_mode_scores.to_string(index=False))
print("Total accepted spell fixes:",
    int(df["cc_spellfix_n"].fillna(0).sum()))

print("Text normalization complete.")
print(
    df[
        [
            cc_column,
            "cc_missing_flag",
            "cc_missing_reason",
            "cc_pseudomissing_flag",
            "cc_spell_mode_selected",
            "cc_spellfix_n",
            "cc_cleaned",
            "cc_cleaned_str",
        ]
    ]
    .head(5)
    .to_string()
)

```

```

CC missingness audit:
  cc_missing_reason  row_count
  pseudo_missing_token        4
; N; ---
  valid      12016 Dyspnea; Altered mental status; Chest pain; Respiratory disti
Pseudo-missing rows: 4
Spell correction mode requested: auto
Spell correction mode selected: strict
Spell mode integrity comparison:
  mode  integrityViolation_total  nonempty_segment_rows  spellfix_count_total  fall_to_sma
  strict                      0                      12016                  109
disabled                     0                      12016                  0
Total accepted spell fixes: 109
Text normalization complete.
      ed_triage_cc  cc_missing_flag cc_missing_reason  cc_pseudomissing_flag cc_spell
0          DYSPNEA       False        valid            False
1          DYSPNEA       False        valid            False
2             SOB       False        valid            False
3  POSITIVE BLOOD CULTURES       False        valid            False
4        ABDO PAIN       False        valid            False

```

1.6 5) Rule system and ontology mappings

```

# Purpose: run a quick rule-coverage spot check on representative phrases.
# Why this matters: confirms deterministic override logic is wired correctly
→ before full inference.

print("Rule coverage preview:")
probe_texts = [
    "shortness of breath",
    "bradycardia",
    "positive blood culture",
    "alter mental status",
    "abdominal pain",
    "right eye redness / pain",
    "paperwork clearance",
]
for text in probe_texts:
    override, _ = rule_gate(text)
    print(f"{text}!r -> {override}")

```

Rule coverage preview:

```
'shortness of breath' -> RVC-SYM-RESP
'bradycardia' -> RVC-SYM-CIRC
'positive blood culture' -> RVC-TEST
'alter mental status' -> RVC-SYM-NERV
'abdominal pain' -> RVC-SYM-DIG
'right eye redness / pain' -> RVC-SYM-EYE
'paperwork clearance' -> RVC-ADMIN
```

1.7 6) Prototype loading, deduplication, and embedding preparation

```
# Purpose: build prototype embeddings/resources once and print metadata for
#           ↳ reproducibility auditing.
# Why this matters: model/device/prototype counts help explain classification
#           ↳ behavior during review.

resources = build_prototype_resources(cfg)
metadata = resources["metadata"]

print("Model metadata:")
print(f"- model_name: {metadata['model_name']}")"
print(f"- device: {metadata['device']}")"
print(f"- spacy_model: {metadata['spacy_model']}")"
print(f"- hf_token_present: {metadata['hf_token_present']}")"
print(f"- appendix_rows_used: {metadata['appendix_rows_used']}")"
print(f"- total_prototypes: {metadata['prototype_total']}")"
print("- prototype_count_by_group:")
print(metadata["prototype_count_by_group"])

Model metadata:
- model_name: NeuML/bioclinical-modernbert-base-embeddings
- device: cpu
- spacy_model: en_core_web_sm
- hf_token_present: True
- appendix_rows_used: 606
- total_prototypes: 693
- prototype_count_by_group:
{'RVC-INJ': 75, 'RVC-SYM-RESP': 43, 'RVC-SYM-CIRC': 18, 'RVC-SYM-NERV': 25, 'RVC-SYM-DIG': 52, 'RVC-SYM-GU': 76, 'RVC-SYM-MSK': 22, 'RVC-SYM-SKIN': 29, 'RVC-SYM-EYE': 50, 'RVC-SYM-GEN': 38, 'RVC-SYM-PSY': 40, 'RVC-DIS': 92, 'RVC-TEST': 12, 'RVC-DIAG': 47, 'RVC-TREAT': 54, 'RVC-ADMIN': 14, 'RVC-UNCL': 6}
```

1.8 7) Segment cache build and classification

```
# Purpose: build the segment cache and run visit-level RFV assignment across
#           the full dataset.
# Why this matters: caching unique segments improves runtime while preserving
#           deterministic outputs.

cache = build_segment_cache(
    df=df,
    seg_col="cc_cleaned",
    resources=resources,
    batch_size=cfg.segment_batch_size,
)

df_out = assign_rvc_per_segment_and_visit_cached(
    df=df,
    cache=cache,
    resources=resources,
    seg_col="cc_cleaned",
    cfg=cfg,
)
df_out = apply_pseudomissing_uncodable_policy(
    df_out,
    max_rfvs=cfg.max_rfvs,
    pseudo_flag_column="cc_pseudomissing_flag",
)
df_out = apply_blank_primary_uncodable_policy(
    df_out,
    max_rfvs=cfg.max_rfvs,
    missing_reason_column="cc_missing_reason",
)

print(f"Unique segments cached: {len(cache.uniq_segments)}")
print(f"Cache emb shape: {tuple(cache.emb.shape)}")
print(f"Cache sim shape: {tuple(cache.sim_proto.shape)}")

Unique segments cached: 1,334
Cache emb shape: (1334, 768)
Cache sim shape: (1334, 693)
```

1.9 8) Output shaping and compatibility mapping

```

Standardized output columns: - RFV1..RFV5 (legacy short-code compatibility) -  

RFV1_name..RFV5_name - RFV1_support..RFV5_support - RFV1_sim..RFV5_sim -  

segment_preds

# Purpose: validate the output schema contract and preview standardized RFV  

# columns.  

# Why this matters: downstream notebooks rely on these exact column names and  

# structures.

validate_output_schema(  

    df_out,  

    max_rfv=cfg.max_rfv,  

    require_primary_for_nonempty=cfg.allow_uncodable_primary,  

)  
  

display_cols = ["cc_cleaned"]  

for slot in range(1, cfg.max_rfv + 1):  

    display_cols.extend(  

        [  

            f"RFV{slot}",  

            f"RFV{slot}_name",  

            f"RFV{slot}_support",  

            f"RFV{slot}_sim",  

        ]  

    )  
  

print(df_out[display_cols].head(10).to_string())

```

Schema validation passed for 12,020 rows.

	cc_cleaned	RFV1	RFV1_name	RFV1
0	[dyspnea]	resp	Symptom - Respiratory	
1	[dyspnea]	resp	Symptom - Respiratory	
2	[shortness of breath]	resp	Symptom - Respiratory	shortness o
3	[positive blood culture]	abn_test	Abnormal test result	positive bloo
4	[abdominal pain]	gi	Symptom - Digestive	abdom
5	[status post fall]	injury	Injuries & adverse effects	status p
6	[alter mental status, vomit]	neuro	Symptom - Nervous	alter menta
7	[anemia, dyspnea]	resp	Symptom - Respiratory	
stated)	anemia 0.647051		NAN	
8	[dyspnea]	resp	Symptom - Respiratory	
9	[subarachnoid hemorrhage, transfer]	disease	Diseases (patient-stated)	subarachnoid h

1.10 9) Diagnostics and QA checks

```
# Purpose: run QA smoke checks and summarize regression metrics after
    ↵ classification.
# Why this matters: immediate diagnostics make it easier to detect unintended
    ↵ behavioral drift.

qa_summary = validate_pipeline(df_out, cache=cache, resources=resources,
    ↵ seg_col="cc_cleaned")

override_summary = summarize_overrides(df_out, column="segment_preds")

print("\nRegression summary")
print("-----")
print(f"Input rows: {len(df)}:{,}")
print(f"Output rows: {len(df_out)}:{,}")
print(f"Row count preserved: {len(df) == len(df_out)}")

print("\nTop RFV1_name counts:")
print(df_out["RFV1_name"].value_counts(dropna=True).head(15).to_string())

print("\nOverride usage:")
print(f"Segments total: {override_summary['segments_total']}:{,}")
print(f"Segments overridden: {override_summary['segments_overridden']}:{,}")
print(f"Segment override rate:
    ↵ {override_summary['segment_override_rate']:.4f}")
print(f"Rows total: {override_summary['rows_total']}:{,}")
print(f"Rows with >=1 override: {override_summary['rows_with_override']}:{,}")
print(f"Row override rate: {override_summary['row_override_rate']:.4f}")

phrase_audit_summary, phrase_audit_rfv = build_phrase_regression_audit(
    df_out,
    cc_column=cc_column,
)
print("\nAbnormal-labs phrase audit:")
print(phrase_audit_summary.to_string(index=False))
if not phrase_audit_rfv.empty:
    print(phrase_audit_rfv.head(10).to_string(index=False))

abnormal_loss_rows = int(
    phrase_audit_summary.loc[
        phrase_audit_summary["metric"].eq("abnormal_loss_rows"), "value"
    ].iloc[0]
```

```

)
if abnormal_loss_rows > 0:
    raise AssertionError(
        "Detected abnormal-labs normalization regression: "
        f"{abnormal_loss_rows} rows still contain 'abnormal loss' segments."
)

normalization_integrity = _score_normalization_integrity(
    df_out["cc_text_for_nlp"].astype("string"),
    df_out["cc_cleaned_str"].astype("string"),
)
integrityViolation_total = int(sum(normalization_integrity.values()))
print("\nNormalization integrity checks:")
for metric_name, metric_value in normalization_integrity.items():
    print(f"{metric_name}: {metric_value},")
print(f"integrityViolation_total: {integrityViolation_total},")
if integrityViolation_total > 0:
    raise AssertionError(
        "Detected clinically invalid normalization substitutions: "
        f"{normalization_integrity}"
    )

if {"ed_intime", "cc_cleaned_str"}.issubset(df_out.columns):
    years = pd.to_datetime(df_out["ed_intime"], errors="coerce").dt.year
    covidArtifact_rows = int(
        (
            years.between(2011, 2019, inclusive="both")
            & df_out["cc_cleaned_str"].astype("string").str.contains(r"\b"
                ↵ covid\b", case=False, na=False)
        ).sum()
    )
    print(f"covid_token_rows_2011_2019: {covidArtifact_rows},")
    if covidArtifact_rows > 0:
        raise AssertionError(
            "Found covid tokens in pre-2020 visits after
            ↵ normalization/classification."
        )

parsedSegmentPreds = df_out["segment_preds"].fillna("[]").map(
    lambda value: json.loads(value) if isinstance(value, str) else []
)
blank_rfv1 = df_out["RFV1_name"].fillna("").astype(str).str.strip().eq("")

```

```

nonempty_segment_rows = parsed_segment_preds.map(lambda payload:
    ↪ isinstance(payload, list) and len(payload) > 0)
all_uncodable_rows = parsed_segment_preds.map(
    lambda payload: isinstance(payload, list)
    and len(payload) > 0
    and all(isinstance(item, dict) and item.get("code") == "RVC-UNCL" for
        ↪ item in payload)
)

blank_rfv1_total = int(blank_rfv1.sum())
blank_rfv1_with_nonempty_segments = int((blank_rfv1 &
    ↪ nonempty_segment_rows).sum())
all_uncodable_total = int(all_uncodable_rows.sum())

print("\nPrimary-label completeness:")
print(f"blank_rfv1_total: {blank_rfv1_total:,}")
print(f"blank_rfv1_with_nonempty_segments:
    ↪ {blank_rfv1_with_nonempty_segments:,}")
print(f"all_uncodable_rows: {all_uncodable_total:,}")

uncodable_segment_counter: Counter = Counter()
for payload in parsed_segment_preds[all_uncodable_rows]:
    for item in payload:
        if isinstance(item, dict):
            uncodable_segment_counter[item.get("segment", "")] += 1

print("\nTop uncodable segments (post-fix):")
for segment, count in uncodable_segment_counter.most_common(15):
    print(f"{count:4d} | {segment}")

if cfg.allow_uncodable_primary and blank_rfv1_with_nonempty_segments != 0:
    raise AssertionError(
        "Expected zero blank RFV1_name values for rows with non-empty segment
        ↪ predictions, "
        f"found {blank_rfv1_with_nonempty_segments}."
    )

if cfg.run_optional_plots:
    top_table = top_cc_plot(df, top_n=30, normalize_case="upper")
    print(top_table.head(10).to_string(index=False))
else:
    print("\nOptional plots skipped (cfg.run_optional_plots=False).")

```

```
Prototype integrity:  
{'RVC-INJ': 75, 'RVC-SYM-RESP': 43, 'RVC-SYM-CIRC': 18, 'RVC-SYM-NERV': 25, 'RVC-SYM-DIG': 52, 'RVC-SYM-GU': 76, 'RVC-SYM-MSK': 22, 'RVC-SYM-SKIN': 29, 'RVC-SYM-EYE': 50, 'RVC-SYM-GEN': 38, 'RVC-SYM-PSY': 40, 'RVC-DIS': 92, 'RVC-TEST': 12, 'RVC-DIAG': 47, 'RVC-TREAT': 54, 'RVC-ADMIN': 14, 'RVC-UNCL': 6}  
Total prototypes: 693
```

Segment column health:

```
rows=12,020 | with>=1 segment=12,016 | unique segments=1,334
```

Probe classifications:

```
shortness of breath -> [('shortness of breath', 'RVC-SYM-RESP', 0.8401217460632324, True)]  
bradycardia -> [('bradycardia', 'RVC-SYM-CIRC', 0.6315359473228455, True)]  
positive blood culture -> [('positive blood culture', 'RVC-TEST', 0.8695230484008789, True)]  
alter mental status -> [('alter mental status', 'RVC-SYM-NERV', 0.7240737080574036, True)]  
abdominal pain / shortness of breath -> [('abdominal pain', 'RVC-SYM-DIG', 0.830388069152832  
SYM-RESP', 0.8401217460632324, True)]  
right eye redness / pain -> [('right eye redness', 'RVC-SYM-EYE', 0.9090079069137573, True),  
SYM-DIG', 0.7539336085319519, False)]
```

Regression summary

```
-----  
Input rows: 12,020  
Output rows: 12,020  
Row count preserved: True
```

Top RFV1_name counts:

RFV1_name	
Symptom - Respiratory	3616
Injuries & adverse effects	1683
Symptom - Digestive	1412
Symptom - Nervous	1330
Symptom - Circulatory	1135
Diseases (patient-stated)	662
Symptom - General	551
Symptom - Skin/Hair/Nails	308
Symptom - Musculoskeletal	273
Uncodable/Unknown	246
Abnormal test result	224
Administrative	183
Symptom - Genitourinary	160
Diagnostic/Screening/Preventive	79
Symptom - Psychological	56

```
Override usage:  
Segments total: 18,184  
Segments overridden: 13,748  
Segment override rate: 0.7560  
Rows total: 12,020  
Rows with >=1 override: 9,793  
Row override rate: 0.8147
```

```
Abnormal-labs phrase audit:
```

metric	value
abnormal_lab_rows	288
abnormal_loss_rows	0
abnormal_test_result_rows	130
eye_ear_rows	0
rfv1_name	row_count row_rate
Abnormal test result	130 0.451389
Symptom - Respiratory	29 0.100694
Symptom - Digestive	29 0.100694
Symptom - Nervous	27 0.093750
Symptom - Circulatory	24 0.083333
Symptom - General	20 0.069444
Diseases (patient-stated)	8 0.027778
Symptom - Genitourinary	6 0.020833
Symptom - Skin/Hair/Nails	5 0.017361
Symptom - Musculoskeletal	4 0.013889

```
Normalization integrity checks:
```

```
fall_to_small: 0  
head_to_heart: 0  
foot_to_post: 0  
copd_or_cold_to_covid: 0  
hyperglycemia_to_hypoglycemia: 0  
hyperkalemia_to_hypokalemia: 0  
integrityViolation_total: 0  
covid_token_rows_2011_2019: 0
```

```
Primary-label completeness:
```

```
blank_rfv1_total: 0  
blank_rfv1_with_nonempty_segments: 0  
all_uncodable_rows: 246
```

```
Top uncodable segments (post-fix):
```

```

18 | ms change
17 | unknown - cc
15 | mcc
11 | gtube eval
11 | alter ms
7 | aortic dissection
5 | ftt
4 |
4 | unable to ambulate
4 | clot fistula
4 | aortic disection
3 | hang
3 | delta ms
3 | mca
3 | change ms

Optional plots skipped (cfg.run_optional_plots=False).

cc_missing_render = cc_missing_audit.copy()
if "examples" in cc_missing_render.columns:
    cc_missing_render["examples"] =
    ↵ cc_missing_render["examples"].astype(str).str.slice(0, 180)
print(
    render_latex_longtable(
        cc_missing_render,
        caption="Chief complaint missingness audit (pseudo-missing and
    ↵ true-missing policy).",
        label="tab:classifier_cc_missing_audit",
        index=False,
    )
)

\begin{longtable}{lrl}
\caption{Chief complaint missingness audit (pseudo-missing and true-missing policy).} \label{tab:classifier_cc_missing_audit}
\toprule
cc_missing_reason & row_count & examples \\
\midrule
\endfirsthead
\caption[]{}{Chief complaint missingness audit (pseudo-missing and true-missing policy).} \\
\toprule
cc_missing_reason & row_count & examples \\
\midrule

```

```

\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
pseudo_missing_token & 4 & -; N; ___ \\
valid & 12016 & Dyspnea; Altered mental status; Chest pain; Respiratory distress; Abd pain \\
\end{longtable}

```

1.11 10) Export

```

# Purpose: export the NLP-augmented workbook using canonical handoff names
    ↵ and a dated archive copy.
# Why this matters: downstream analysis reads a stable canonical path while
    ↵ preserving run-history snapshots.

canonical_output_path, archive_output_path =
    ↵ resolve_classifier_output_paths(WORK_DIR)
archive_output_path.parent.mkdir(parents=True, exist_ok=True)
write_archive_xlsx = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0").strip() ==
    ↵ "1"

def _sha256_file(path: Path) -> str:
    hasher = hashlib.sha256()
    with path.open("rb") as file_handle:
        for chunk in iter(lambda: file_handle.read(1024 * 1024), b ""):
            hasher.update(chunk)
    return hasher.hexdigest()

df_out.to_excel(canonical_output_path, index=False,
    ↵ sheet_name=cfg.output_sheet_name)
if write_archive_xlsx:
    df_out.to_excel(archive_output_path, index=False,
    ↵ sheet_name=cfg.output_sheet_name)
canonical_output_sha256 = _sha256_file(canonical_output_path)
archive_output_sha256 = _sha256_file(archive_output_path) if
    ↵ write_archive_xlsx else None
print(
    "EXPORT_SHAPE classifier",
    f"rows={len(df_out)}:,}",

```

```

f"cols={df_out.shape[1]:,}",
f"canonical_sha256={canonical_output_sha256}" ,
)

run_date = pd.Timestamp.now().strftime("%Y-%m-%d")
cc_missing_audit_path = archive_output_path.parent / f"{run_date}"
↪ classifier_cc_missing_audit.csv"
cc_missing_audit.to_csv(cc_missing_audit_path, index=False)
phrase_audit_path = archive_output_path.parent / f"{run_date}"
↪ classifier_phrase_audit.csv"
phrase_audit = phrase_audit_summary.copy()
phrase_audit["audit_type"] = "summary"
if not phrase_audit_rfv.empty:
    phrase_detail = phrase_audit_rfv.copy()
    phrase_detail["metric"] = "rfv_distribution"
    phrase_detail["value"] = phrase_detail["row_count"]
    phrase_detail["audit_type"] = "rfv_distribution"
    phrase_audit = pd.concat([phrase_audit, phrase_detail],
↪ ignore_index=True, sort=False)
phrase_audit.to_csv(phrase_audit_path, index=False)

spell_mode_comparison_path = archive_output_path.parent / f"{run_date}"
↪ classifier_spell_mode_comparison.csv"
spell_mode_scores.to_csv(spell_mode_comparison_path, index=False)
spellfix_log_rows = (
    df_out.loc[
        df_out["cc_spellfix_any"].fillna(False).astype(bool),
        ["hadm_id", "subject_id", "cc_text_for_nlp", "cc_cleaned_str",
↪ "cc_spell_mode_selected", "cc_spellfix_log"],
    ]
    .copy()
)
spellfix_log_path = archive_output_path.parent / f"{run_date}"
↪ classifier_spellfix_log.csv"
spellfix_log_rows.to_csv(spellfix_log_path, index=False)

hypercap_flags_audit = pd.DataFrame(
    [
        {
            "row_count": int(len(df_out)),
            "flag_abg_hypercapnia_positive": int(
                (pd.to_numeric(df_out.get("flag_abg_hypercapnia", 0),
↪ errors="coerce").fillna(0).astype(int) > 0).sum()

```

```

),
"flag_vbg_hypercapnia_positive": int(
    pd.to_numeric(df_out.get("flag_vbg_hypercapnia", 0),
    errors="coerce").fillna(0).astype(int) > 0).sum()
),
"hypercap_by_abg_positive": int(
    pd.to_numeric(df_out.get("hypercap_by_abg", 0),
    errors="coerce").fillna(0).astype(int).sum()
),
"hypercap_by_vbg_positive": int(
    pd.to_numeric(df_out.get("hypercap_by_vbg", 0),
    errors="coerce").fillna(0).astype(int).sum()
),
"hypercap_by_bg_positive": int(
    pd.to_numeric(df_out.get("hypercap_by_bg", 0),
    errors="coerce").fillna(0).astype(int).sum()
),
"bg_union_mismatch_n": int(
(
(
    pd.to_numeric(df_out.get("hypercap_by_abg", 0),
    errors="coerce").fillna(0).astype(int)
    | pd.to_numeric(df_out.get("hypercap_by_vbg", 0),
    errors="coerce").fillna(0).astype(int)
)
!= pd.to_numeric(df_out.get("hypercap_by_bg", 0),
    errors="coerce").fillna(0).astype(int)
).sum()
),
}
],
)
hypercap_flags_audit_path = archive_output_path.parent / f"{run_date}"
    classifier_hypercap_flags_audit.csv"
hypercap_flags_audit.to_csv(hypercap_flags_audit_path, index=False)

classifier_contract = validate_classifier_contract(df_out,
    max_rfv=cfg.max_rfv)
classifier_contract["resource_report"] = resource_report
contract_report_path = archive_output_path.parent / f"{run_date}"
    classifier_contract_report.json"
contract_report_path.write_text(json.dumps(classifier_contract, indent=2))
contract_artifact_payload = {

```

```

"generated_utc": pd.Timestamp.utcnow().isoformat(),
"status": classifier_contract["status"],
"contracts": {"classifier": classifier_contract},
"findings": classifier_contract["findings"],
}
contract_artifact_paths = write_contract_report(contract_artifact_payload,
↪ work_dir=WORK_DIR)

stage_manifest = collect_run_manifest(
    WORK_DIR,
    run_id=f"classifier_{pd.Timestamp.now().strftime('%Y%m%d_%H%M%S')}",
)
stage_manifest["stage"] = "classifier"
stage_manifest["model_metadata"] = {
    "model_name": metadata.get("model_name"),
    "spacy_model": metadata.get("spacy_model"),
    "device": metadata.get("device"),
    "appendix_rows_used": metadata.get("appendix_rows_used"),
    "prototype_total": metadata.get("prototype_total"),
}
stage_manifest["resource_verification"] = resource_report
stage_manifest["outputs"] = {
    "canonical_output_path": str(canonical_output_path),
    "canonical_output_sha256": canonical_output_sha256,
    "archive_output_path": str(archive_output_path) if write_archive_xlsx
↪ else None,
    "archive_output_sha256": archive_output_sha256,
    "cc_missing_audit_path": str(cc_missing_audit_path),
    "phrase_audit_path": str(phrase_audit_path),
    "spell_mode_comparison_path": str(spell_mode_comparison_path),
    "spellfix_log_path": str(spellfix_log_path),
    "hypercap_flags_audit_path": str(hypercap_flags_audit_path),
    "classifier_contract_report_path": str(contract_report_path),
}
stage_manifest_path = archive_output_path.parent / f"{run_date}"
↪ classifier_run_manifest.json"
stage_manifest_path.write_text(json.dumps(stage_manifest, indent=2))

contract_mode = os.getenv("PIPELINE_CONTRACT_MODE", "fail").strip().lower()
if contract_mode not in {"fail", "warn"}:
    raise ValueError("PIPELINE_CONTRACT_MODE must be one of {'fail',
↪ 'warn'}")
if classifier_contract["status"] == "fail" and contract_mode == "fail":

```

```

        raise ValueError(
            "Classifier contract failed. See report at "
            f"{contract_report_path} for details."
    )

    print(f"Wrote canonical output: {canonical_output_path}")
    if write_archive_xlsx:
        print(f"Wrote archive output: {archive_output_path}")
    else:
        print("Skipped archive workbook export (WRITE_ARCHIVE_XLSX_EXPORTS!=1).")
    print(f"Wrote CC missingness audit: {cc_missing_audit_path}")
    print(f"Wrote phrase regression audit: {phrase_audit_path}")
    print(f"Wrote spell-mode comparison audit: {spell_mode_comparison_path}")
    print(f"Wrote spellfix log audit: {spellfix_log_path}")
    print(f"Wrote hypercapnia flags audit: {hypercap_flags_audit_path}")
    print(f"Wrote classifier contract report: {contract_report_path}")
    print(f"Wrote classifier run manifest: {stage_manifest_path}")
    print(f"Wrote pipeline contract artifact:
        ↳ {contract_artifact_paths['contract_report_path']}"))
    if contract_artifact_paths["failed_contract_path"].exists():
        print(f"Wrote failed contract artifact:
            ↳ {contract_artifact_paths['failed_contract_path']}")
    print(f"Classifier contract status: {classifier_contract['status']}"))
    print(
        "Schema summary: "
        f"rows={len(df_out):,}, cols={df_out.shape[1]:,}, "
        f"rfv_slots={cfg.max_rfv}")
)
print(f"Export complete: {canonical_output_path}")

```

```

EXPORT_SHAPE classifier rows=12,020 cols=337 canonical_sha256=f0fb85abd68b82b9046428dccc58b5
Wrote canonical output: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke
CC-NLP/MIMIC tabular data/MIMICIV all with CC_with_NLP.xlsx
Skipped archive workbook export (WRITE_ARCHIVE_XLSX_EXPORTS!=1).
Wrote CC missingness audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/L
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 classifier_cc_missing_audit.csv
Wrote phrase regression audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Wor
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 classifier_phrase_audit.csv
Wrote spell-mode comparison audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 classifier_spell_mode_comparison.csv
Wrote spellfix log audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Lo
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 classifier_spellfix_log.csv

```

```
Wrote hypercapnia flags audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 classifier_hypercap_flags_audit.csv
Wrote classifier contract report: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 classifier_contract_report.json
Wrote classifier run manifest: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 classifier_run_manifest.json
Wrote pipeline contract artifact: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/CC-NLP/debug/contracts/20260220_193338/contract_report.json
Classifier contract status: pass
Schema summary: rows=12,020, cols=337, rfv_slots=5
Export complete: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research/CC-NLP/MIMIC tabular data/MIMICIV all with CC_with_NLP.xlsx
```