Hypercap CC NLP Classifier

Table of contents

1 Workbook for MIMIC Hypercapnia Presenting Chief Concern Analysis

This notebook is intentionally self-contained and deterministic. It performs NLP-based mapping from free-text ED chief complaints to NHAMCS top-level RVC categories, with strict compatibility outputs for downstream analysis notebooks.

1.1 1) Runtime contract and imports

- No runtime package installation.
- Fail fast if required dependencies are missing.
- Keep environment deterministic and explicit.

```python
# Purpose: verify required packages, import shared dependencies, and
↳  initialize NLP runtime objects.
# Why this matters: later cells assume these imports/models exist, so we fail
↳  early if the environment is incomplete.

import importlib

REQUIRED_IMPORTS = {
    "numpy": "numpy",
    "pandas": "pandas",
    "torch": "torch",
    "ftfy": "ftfy",
    "spacy": "spacy",
    "symspellpy": "symspellpy",
    "sentence_transformers": "sentence_transformers",
    "openpyxl": "openpyxl",
    "dotenv": "python-dotenv",
}

missing = []
for import_name, package_name in REQUIRED_IMPORTS.items():
    if importlib.util.find_spec(import_name) is None:
        missing.append((import_name, package_name))

if missing:
    message = "Missing required dependencies (install via uv sync): " + ", 
↳  ".join(
        f"{import_name} ({package_name})" for import_name, package_name in
            ↳  missing
    )
    raise ImportError(message)

import hashlib
import json
import logging
import math
import os
import re
import sys
import warnings
from collections import Counter, defaultdict
from dataclasses import dataclass
from datetime import datetime
```

```python
from pathlib import Path
from typing import Any, Mapping, Sequence

import ftfy
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import spacy
import torch
from dotenv import load_dotenv
from symspellpy.symspellpy import SymSpell, Verbosity

# Silence widget-only tqdm warning in environments without Jupyter progress
↪  widgets.
try:
    from tqdm import TqdmWarning
except Exception:  # pragma: no cover - defensive for uncommon tqdm import
↪  edge cases
    TqdmWarning = None

if TqdmWarning is not None:
    warnings.filterwarnings(
        "ignore",
        message="IProgress not found.*",
        category=TqdmWarning,
    )

# Reduce HF Hub warning noise in notebook output. If HF_TOKEN is set,
↪  authenticated requests still work.
try:
    from huggingface_hub import logging as hf_logging

    hf_logging.set_verbosity_error()
except Exception:  # pragma: no cover - keeps notebook robust if hub
↪  internals change
    pass

logging.getLogger("huggingface_hub").setLevel(logging.ERROR)
logging.getLogger("hf_xet").setLevel(logging.ERROR)

from sentence_transformers import SentenceTransformer, util
```

```
# spaCy model loading is config-driven in the setup cell so behavior is
↪  explicit.
NLP: Any | None = None
SPACY_MODEL_LABEL = "UNINITIALIZED (loaded in setup cell)"
print("spaCy pipeline: pending config-driven load")


spaCy pipeline: pending config-driven load
```

1.2 2) Config and constants

1.2.1 Core API (single-cell contract)

This cell defines the notebook's public in-notebook interfaces: - ClassifierConfig - get_cc_column - normalize_cc - build_prototype_resources - build_segment_cache - classify_segments_cached - assign_rvc_per_segment_and_visit_cached - validate_output_schema - save_with_suffix

```
# Purpose: centralize the notebook's core API (config, constants, data
↪  structures, and reusable functions).
# Why this matters: keeping definitions in one place makes behavior
↪  predictable and easier for new contributors to trace.

@dataclass(frozen=True)
class ClassifierConfig:
    """Configuration for notebook-contained NLP classification."""

    seed: int = 1234
    work_dir_env_var: str = "WORK_DIR"
    data_dirname: str = "MIMIC tabular data"
    input_filename: str = "MIMICIV all with CC.xlsx"
    input_filename_env_var: str = "CLASSIFIER_INPUT_FILENAME"
    output_suffix: str = "_with_NLP"
    output_sheet_name: str = "with_NLP"
    model_name: str = "NeuML/bioclinical-modernbert-base-embeddings"
    model_trust_remote_code: bool = True
    model_batch_size: int = 256
    segment_batch_size: int = 512
    scoring_method: str = "max"  # {"max", "lse"}
    alpha: float = 12.0
    abstain_threshold: float = 0.40
    max_rfv: int = 5
    max_segments_per_visit: int = 5
```

```python
    prototype_dedup_sim_thresh: float = 0.92
    summary_weight_strategy: str = "uniform"  # {"uniform", "by_rows"}
    run_optional_plots: bool = False
    require_spacy_model: bool = True
    allow_uncodable_primary: bool = True
    empty_primary_policy: str = "blank"  # {"blank", "uncodable"}
    appendix_relpath: str = "Annotation/nhamcs_rvc_2022_appendixII_codes.csv"
    summary_relpath: str =
 ↪  "Annotation/nhamcs_rvc_2022_summary_by_top_level_17.csv"
    hf_token_env_var: str = "HF_TOKEN"


# --- Ontology labels and precedence used by final visit-level outputs. ---
RVC_NAME: dict[str, str] = {
    "RVC-INJ": "Injuries & adverse effects",
    "RVC-SYM-RESP": "Symptom - Respiratory",
    "RVC-SYM-CIRC": "Symptom - Circulatory",
    "RVC-SYM-NERV": "Symptom - Nervous",
    "RVC-SYM-DIG": "Symptom - Digestive",
    "RVC-SYM-GU": "Symptom - Genitourinary",
    "RVC-SYM-MSK": "Symptom - Musculoskeletal",
    "RVC-SYM-SKIN": "Symptom - Skin/Hair/Nails",
    "RVC-SYM-EYE": "Symptom - Eye/Ear",
    "RVC-SYM-GEN": "Symptom - General",
    "RVC-SYM-PSY": "Symptom - Psychological",
    "RVC-DIS": "Diseases (patient-stated)",
    "RVC-TEST": "Abnormal test result",
    "RVC-DIAG": "Diagnostic/Screening/Preventive",
    "RVC-TREAT": "Treatment/Medication",
    "RVC-ADMIN": "Administrative",
    "RVC-UNCL": "Uncodable/Unknown",
}

PRECEDENCE_RVC: list[str] = [
    "RVC-INJ",
    "RVC-SYM-RESP",
    "RVC-SYM-CIRC",
    "RVC-SYM-NERV",
    "RVC-SYM-DIG",
    "RVC-SYM-GU",
    "RVC-SYM-MSK",
    "RVC-SYM-SKIN",
    "RVC-SYM-EYE",
```

```python
    "RVC-SYM-GEN",
    "RVC-SYM-PSY",
    "RVC-DIS",
    "RVC-TEST",
    "RVC-DIAG",
    "RVC-TREAT",
    "RVC-ADMIN",
    "RVC-UNCL",
]

RVC_SHORT: dict[str, str] = {
    "RVC-INJ": "injury",
    "RVC-SYM-RESP": "resp",
    "RVC-SYM-CIRC": "circ",
    "RVC-SYM-NERV": "neuro",
    "RVC-SYM-DIG": "gi",
    "RVC-SYM-GU": "gu",
    "RVC-SYM-MSK": "msk",
    "RVC-SYM-SKIN": "skin",
    "RVC-SYM-EYE": "eye_ear",
    "RVC-SYM-GEN": "general",
    "RVC-SYM-PSY": "psych",
    "RVC-DIS": "disease",
    "RVC-TEST": "abn_test",
    "RVC-DIAG": "screen",
    "RVC-TREAT": "treatment",
    "RVC-ADMIN": "admin",
    "RVC-UNCL": "uncodable",
}

# --- Text normalization vocabulary and complaint parsing constants. ---
CC_CANDIDATES: tuple[str, ...] = (
    "ed_triage_cc",
    "chief_complaint",
    "ed_chief_complaint",
    "chiefcomplaint",
    "chief_complaint_text",
    "ed_cc",
    "cc",
)

PROTECT = ("n/v", "n/v/d", "s/p", "w/", "w/o", "h/o", "c/o", "h/a", "f/c",
 ↪  "c/p", "lab", "labs")
```

```python
SPELL_PROTECT_PHRASES = {
    "found down",
}
SPELL_PROTECT_TOKENS = {
    "lab",
    "labs",
    "abnl",
    "abnormal",
    "fall",
    "falls",
    "head",
    "foot",
    "feet",
    "left",
    "right",
    "cold",
    "copd",
    "cord",
    "hyperglycemia",
    "hyperkalemia",
    "hypoglycemia",
    "hypokalemia",
    "found",
    "trach",
    "dental",
    "black",
}
SPELL_DENYLIST_SUBSTITUTIONS = {
    ("found", "wound"),
    ("trach", "tract"),
    ("dental", "mental"),
}
STRICT_SPELL_MIN_TOKEN_LEN = 5
STRICT_SPELL_MAX_EDIT_DISTANCE = 1
POLARITY_GUARD_PREFIX_PAIRS: tuple[tuple[str, str], ...] = (
    ("hyper", "hypo"),
    ("hypo", "hyper"),
)
LATERALITY_GUARD_TOKENS = {"left", "right"}
SPELL_CORRECTION_MODES = {"disabled", "strict", "auto"}
MISSING_RE = re.compile(
    r"^(?:n/?a|na|none|unknown|unk|no chief complaint|no complaint|not
    ↪  applicable|not available)$",
```

```python
    re.I,
)
NEG_RE = re.compile(r"^(?:denies?|no|not|without|w/o)\b", re.I)

# Preserve a small set of clinically meaningful "absence" complaints instead
↪  of dropping as generic negation.
ABSENCE_SYMPTOM_RULES: tuple[tuple[re.Pattern[str], str], ...] = (
    (
        re.compile(r"^(?:no|without|w/o)\s+(?:urine|urinary)\s+(?:output|⌋
↪  void(?:ing)?)\b", re.I),
        "urinary retention",
    ),
    (
        re.compile(r"^(?:no|without|w/o)\s+(?:bowel movement|bm|stool(?:
↪  output)?)\b", re.I),
        "constipation",
    ),
)

SEG_OVR: dict[str, str] = {
    "n/v": "nausea vomiting",
    "n/v/d": "nausea vomiting diarrhea",
    "nvd": "nausea vomiting diarrhea",
    "w/": "with",
    "w/o": "without",
    "s/p": "status post",
    "h/o": "history of",
    "c/o": "complains of",
    "h/a": "headache",
    "f/c": "fever chills",
    "c/p": "chest pain",
    "mva": "motor vehicle accident",
    "mvc": "motor vehicle collision",
    "ped vs auto": "pedestrian struck by vehicle",
    "gsw": "gunshot wound",
    "sa": "sexual assault",
}

TOK_OVR: dict[str, str] = {
    "sob": "shortness of breath",
    "doe": "dyspnea on exertion",
    "cp": "chest pain",
    "palps": "palpitations",
```

```python
    "ha": "headache",
    "h/a": "headache",
    "f/c": "fever chills",
    "c/p": "chest pain",
    "loc": "loss of consciousness",
    "ams": "altered mental status",
    "sz": "seizure",
    "szs": "seizure",
    "cva": "stroke",
    "tia": "transient ischemic attack",
    "abd": "abdominal",
    "abdo": "abdominal",
    "rlq": "right lower quadrant",
    "ruq": "right upper quadrant",
    "llq": "left lower quadrant",
    "luq": "left upper quadrant",
    "lbp": "low back pain",
    "brbpr": "rectal bleeding",
    "uti": "urinary tract infection",
    "dysuria": "painful urination",
    "hematuria": "blood in urine",
    "nv": "nausea vomiting",
    "nvd": "nausea vomiting diarrhea",
    "sorethroat": "sore throat",
    "uri": "upper respiratory infection",
    "pna": "pneumonia",
    # Conservative high-frequency shorthand expansions from uncodable output
    ↪  audit.
    "ich": "intracranial hemorrhage",
    "sah": "subarachnoid hemorrhage",
    "stemi": "st elevation myocardial infarction",
    "nstemi": "non st elevation myocardial infarction",
    "chf": "congestive heart failure",
    "pe": "pulmonary embolism",
    "sbo": "small bowel obstruction",
    "gib": "gastrointestinal bleeding",
    "hypoglycemia": "hypoglycemia",
    "hypokalemia": "hypokalemia",
    "arf": "acute renal failure",
    "ili": "influenza like illness",
    "covid": "covid",
    "flu": "influenza",
    "si": "suicidal ideation",
```

```python
    "hi": "homicidal ideation",
    "avh": "auditory visual hallucinations",
    "etoh": "alcohol",
    "od": "overdose",
    "fb": "foreign body",
    "vag": "vaginal",
    "vb": "vaginal bleeding",
    "hg": "hyperemesis gravidarum",
    "s/p": "status post",
    "w/": "with",
    "w/o": "without",
    "h/o": "history of",
    "c/o": "complains of",
    "n/v": "nausea vomiting",
    "n/v/d": "nausea vomiting diarrhea",
}

CTX_OVR: dict[str, tuple[Any, str, str | None]] = {
    "od": (
        lambda seg: not
        ↪  re.search(r"\b(eye|vision|ophth|ophthalm|cornea|ocular)\b", seg),
        "overdose",
        "right eye",
    ),
    "os": (
        lambda seg:
        ↪  re.search(r"\b(eye|vision|ophth|ophthalm|cornea|ocular)\b", seg),
        "left eye",
        None,
    ),
    "ou": (
        lambda seg:
        ↪  re.search(r"\b(eye|vision|ophth|ophthalm|cornea|ocular)\b", seg),
        "both eyes",
        None,
    ),
}

ABBR_MAP: dict[str, str] = {}

SAFE_SHORT_TOKENS = {"cp", "ha", "nv", "gi", "gu", "ms", "sx", "bp", "hr",
↪  "doe", "sob", "loc"}
```

```python
_PUNCT_EDGE = re.compile(r"(^[^\w/]+)|([^\w/]+$)")
SEP_SPLIT_RE = re.compile(r"\s*(?:[;:,]|[+&])\s*")

CLINICAL_KEYS = {
    "pain",
    "fever",
    "chill",
    "cough",
    "wheez",
    "dyspnea",
    "shortness",
    "hemopty",
    "nausea",
    "vomit",
    "diarrhea",
    "bleeding",
    "rash",
    "injury",
    "wound",
    "laceration",
    "fracture",
    "sprain",
    "assault",
    "fall",
    "seizure",
    "syncope",
    "weakness",
    "numbness",
    "tingling",
    "headache",
    "dizziness",
    "palpitation",
    "chest",
    "abdominal",
    "flank",
    "pelvic",
    "back",
    "neck",
    "urinary",
    "dysuria",
    "hematuria",
    "pregnan",
    "vaginal",
```

```python
    "overdose",
    "intoxication",
    "anxiety",
    "depression",
    "psychosis",
    "sore",
    "throat",
    "ear",
    "eye",
    "dental",
    "asthma",
    "copd",
    "flu",
    "covid",
    "pneumonia",
}

_DUR_RE = re.compile(
    r"\b([a-z]{2,})\s*x\s*(\d{1,3})\s*(d|day|h|hr|hour|w|wk|week|m|mo|mon| ⌋
    ↪  month|y|yr|year)s?\b",
    re.I,
)
_UNIT_MAP = {
    "d": "days",
    "day": "days",
    "h": "hours",
    "hr": "hours",
    "hour": "hours",
    "w": "weeks",
    "wk": "weeks",
    "week": "weeks",
    "m": "months",
    "mo": "months",
    "mon": "months",
    "month": "months",
    "y": "years",
    "yr": "years",
    "year": "years",
}

# --- Rule-gate regex definitions for deterministic high-confidence
↪  overrides. ---
RX_FOCAL_WEAK = re.compile(
```

```python
        r"\b(focal|left|right|arm|leg|hand|face|hemiparesis|hemiplegia).*\b⌋
        ↪ weakness\b|\bweakness\b.*\b(focal|left|right|arm|leg|hand|face)\b",
        re.I,
)

RULE_RX: dict[str, re.Pattern[str]] = {
    "RVC-INJ": re.compile(
        r"\b(mvc|mva|collision|ped(\s|)vs(\s|)auto|assault|fell|fall|gsw|⌋
        ↪ gunshot|stab|laceration|fracture|burn|foreign
        ↪ body|poison(ing)?|overdose|od\b|adverse (drug|medication)
        ↪ reaction)\b",
        re.I,
    ),
    "RVC-SYM-RESP": re.compile(
        r"\b(shortness of breath|dyspnea|doe\b|wheeze|respiratory
        ↪ distress|cough(ing)?|hemoptysis|sputum)\b",
        re.I,
    ),
    "RVC-SYM-CIRC": re.compile(
        r"\b(chest pain|palpitation(s)?|hypotension|low
        ↪ bp|bradycardia|tachycardia|arrhythmia|irregular
        ↪ heartbeat|edema|leg swelling|ankle swelling|peripheral
        ↪ edema|claudication)\b",
        re.I,
    ),
    "RVC-SYM-NERV": re.compile(
        r"\b(alter(ed)? mental status|ams\b|confusion|letharg(y|ic)|⌋
        ↪ unresponsive|syncope|faint(ing)?|seizure(s)?|slurred
        ↪ speech|dysarthria|aphasia|facial
        ↪ droop|headache|migraine|dizziness|vertigo|numbness|tingling)\b",
        re.I,
    ),
    "RVC-SYM-DIG": re.compile(
        r"\b(abd(ominal)? pain|abdo pain|nausea|vomit(ting)?|diarrhea|rectal
        ↪ bleeding|hematemesis|dysphagia|jaundice|gastrointestinal
        ↪ bleeding|gi bleeding)\b",
        re.I,
    ),
    "RVC-SYM-GU": re.compile(
        r"\b(dysuria|hematuria|urinary (frequency|urgency)|flank pain|pelvic
        ↪ pain|vaginal (bleeding|discharge)|penile discharge)\b",
        re.I,
    ),
```

```
"RVC-SYM-MSK": re.compile(
    r"\b(back pain|neck pain|shoulder pain|knee pain|hip pain|joint
    ↪  pain|gait problem)\b",
    re.I,
),
"RVC-SYM-SKIN":
↪  re.compile(r"\b(rash|pruritus|itch(ing)?|cellulitis|wound(?!
↪  check))\b", re.I),
"RVC-SYM-EYE": re.compile(
    r"\b(eye (pain|redness|irritation)|red eye|pink
    ↪  ?eye|conjunctivitis|blurry vision|vision (loss|change)|ear
    ↪  pain|hearing (loss|change)|tinnitus|ear discharge|vertigo)\b",
    re.I,
),
"RVC-SYM-GEN":
↪  re.compile(r"\b(fever|chills|fatigue|malaise|weakness(?!.*focal))\b",
↪  re.I),
"RVC-SYM-PSY": re.compile(
    r"\b(anxiety|depression|insomnia|agitation|suicidal
    ↪  ideation|homicidal ideation|si\b|hi\b|intoxication)\b",
    re.I,
),
"RVC-DIS": re.compile(
    r"\b(asthma attack|copd
    ↪  flare|pneumonia|diabetes(?!.*test)|hypertensive crisis|stroke
    ↪  diagnosed|stroke|intracranial hemorrhage|subarachnoid
    ↪  hemorrhage|st elevation myocardial infarction|non st elevation
    ↪  myocardial infarction|congestive heart failure|pulmonary
    ↪  embolism|small bowel obstruction|acute renal failure|influenza
    ↪  like illness|hypoglycemia|hypokalemia)\b",
    re.I,
),
"RVC-TEST": re.compile(
    r"\b(abnormal (lab|labs|imaging|ct|mri|x[- ]?ray|ekg|ecg)|told (to
    ↪  come|results?)|positive (test|culture|blood culture)|blood
    ↪  culture (positive|grew|growth))\b",
    re.I,
),
"RVC-DIAG": re.compile(
    r"\b((needs|for) (covid|flu|strep) test|screen(ing)?|bp
    ↪  (check|recheck)|workup|routine testing|prenatal)\b",
    re.I,
),
```

```python
    "RVC-TREAT": re.compile(
        r"\b(refill|prescription refill|rx refill|dressing change|suture
        ↪  removal|injection|detox clearance|wound check)\b",
        re.I,
    ),
    "RVC-ADMIN":
    ↪  re.compile(r"\b(form|paperwork|note|letter|clearance|insurance|work
    ↪  (note|letter)|transfer(red)?|intubated transfer)\b", re.I),
}

GROUP_MAP: list[tuple[re.Pattern[str], str]] = [
    (re.compile(r"\binjur|poison|overdose|adverse effect", re.I), "RVC-INJ"),
    (re.compile(r"\brespir", re.I), "RVC-SYM-RESP"),
    (re.compile(r"\bcircul|cardio|lymph", re.I), "RVC-SYM-CIRC"),
    (re.compile(r"\bnerv|neuro", re.I), "RVC-SYM-NERV"),
    (re.compile(r"\bdigest|gastro|abd", re.I), "RVC-SYM-DIG"),
    (re.compile(r"\bgenitour|urinar|renal|pelvic|vagin|penil|breast", re.I),
    ↪  "RVC-SYM-GU"),
    (re.compile(r"\bmusculoskelet|msk|joint|back|neck", re.I),
    ↪  "RVC-SYM-MSK"),
    (re.compile(r"\bskin|hair|nail|derma", re.I), "RVC-SYM-SKIN"),
    (re.compile(r"\beye|ear|vision|hearing|vertigo", re.I), "RVC-SYM-EYE"),
    (re.compile(r"\bgeneral|fever|malaise|weakness", re.I), "RVC-SYM-GEN"),
    (re.compile(r"\bpsych|mental", re.I), "RVC-SYM-PSY"),
    (re.compile(r"\bdisease|diagnos|known|asthma|copd|diabetes|hypertens|⌋
    ↪  pneumonia|stroke", re.I), "RVC-DIS"),
    (re.compile(r"\babnormal .*test|abnormal (lab|imaging|ekg|ecg)|positive
    ↪  (test|culture|result)", re.I), "RVC-TEST"),
    (re.compile(r"\bscreen|diagnostic|prenatal|immuniz|vaccine|bp
    ↪  check|test\b", re.I), "RVC-DIAG"),
    (re.compile(r"\btreat|medicat|refill|suture|dressing|injection|therapy|⌋
    ↪  wound check", re.I), "RVC-TREAT"),
    (re.compile(r"\badmin|paperwork|form|clearance|insurance|work
    ↪  (note|letter)", re.I), "RVC-ADMIN"),
    (re.compile(r"\buncodable|unknown|illegible|none", re.I), "RVC-UNCL"),
]

RE_DROP = re.compile(r"\b(nec|nos|unspecified|other|other and
↪  unspecified)\b", re.I)
RE_SPACE = re.compile(r"\s+")
BODY_PART_ONLY_RE = re.compile(r"(arm|leg|back|knee|hip|wrist|elbow|ankle|⌋
↪  hand|foot|toe|finger|eye|ear)")
```

```python
CANON_REPLACERS: list[tuple[re.Pattern[str], str]] = [
    (re.compile(r"\bshortness of breath\b", flags=re.I), "dyspnea"),
    (re.compile(r"\bdyspneic?\b", flags=re.I), "dyspnea"),
    (re.compile(r"\bdyspnea on exertion\b", flags=re.I), "dyspnea_exertion"),
    (re.compile(r"\bstatus post\b", flags=re.I), "status_post"),
]




RVC_PROTOS_CURATED: dict[str, list[str]] = {
    "RVC-INJ": [
        "injury: laceration cut",
        "injury: fracture",
        "injury: burn",
        "injury: contusion bruise",
        "injury: foreign body",
        "injury: motor vehicle collision",
        "injury: fall",
        "adverse effect: medication reaction",
        "poisoning overdose",
        "injury: assault",
        "gunshot wound",
    ],
    "RVC-SYM-RESP": [
        "respiratory symptom: shortness of breath dyspnea",
        "respiratory symptom: dyspnea on exertion",
        "respiratory symptom: wheeze",
        "respiratory symptom: cough",
        "respiratory symptom: hemoptysis",
        "respiratory symptom: sputum production",
        "respiratory symptom: respiratory distress",
    ],
    "RVC-SYM-CIRC": [
        "circulatory symptom: chest pain suspected cardiac",
        "circulatory symptom: palpitations",
        "circulatory symptom: hypotension low blood pressure",
        "circulatory symptom: bradycardia",
        "circulatory symptom: tachycardia",
        "circulatory symptom: peripheral edema leg swelling ankle swelling",
        "circulatory symptom: claudication",
    ],
    "RVC-SYM-NERV": [
```

```
        "neurologic symptom: altered mental status confusion",
        "neurologic symptom: lethargy decreased responsiveness",
        "neurologic symptom: slurred speech dysarthria",
        "neurologic symptom: aphasia word finding difficulty",
        "neurologic symptom: facial droop",
        "neurologic symptom: syncope fainting",
        "neurologic symptom: seizure",
        "neurologic symptom: headache",
        "neurologic symptom: dizziness vertigo",
        "neurologic symptom: numbness tingling focal weakness",
    ],
    "RVC-SYM-DIG": [
        "digestive symptom: abdominal pain",
        "digestive symptom: nausea vomiting",
        "digestive symptom: diarrhea",
        "digestive symptom: rectal bleeding",
        "digestive symptom: dysphagia",
        "digestive symptom: jaundice",
    ],
    "RVC-SYM-GU": [
        "genitourinary symptom: dysuria painful urination",
        "genitourinary symptom: urinary frequency urgency",
        "genitourinary symptom: hematuria blood in urine",
        "genitourinary symptom: flank pain",
        "genitourinary symptom: pelvic pain",
        "genitourinary symptom: vaginal discharge",
    ],
    "RVC-SYM-MSK": [
        "musculoskeletal symptom: back pain",
        "musculoskeletal symptom: neck pain",
        "musculoskeletal symptom: joint pain",
        "musculoskeletal symptom: hip knee shoulder pain",
        "musculoskeletal symptom: gait problem",
    ],
    "RVC-SYM-SKIN": [
        "skin symptom: rash",
        "skin symptom: pruritus itching",
        "skin symptom: cellulitis redness warmth",
        "skin symptom: nontraumatic wound",
    ],
    "RVC-SYM-EYE": [
        "eye symptom: eye pain",
        "eye symptom: red eye eye redness",
```

```
        "eye symptom: blurry vision vision loss",
        "ear symptom: ear pain discharge",
        "vestibular symptom: vertigo spinning",
    ],
    "RVC-SYM-GEN": [
        "general symptom: fever chills",
        "general symptom: fatigue malaise",
        "general symptom: weakness generalized",
        "general symptom: weight change",
        "general symptom: edema swelling",
    ],
    "RVC-SYM-PSY": [
        "psychological symptom: anxiety",
        "psychological symptom: depression",
        "psychological symptom: agitation",
        "psychological symptom: insomnia",
        "psychological symptom: suicidal ideation",
        "psychological symptom: homicidal ideation",
        "psychological symptom: substance intoxication without overdose",
    ],
    "RVC-DIS": [
        "patient stated diagnosis: asthma attack",
        "patient stated diagnosis: copd flare",
        "patient stated diagnosis: pneumonia",
        "patient stated diagnosis: diabetes problem",
        "patient stated diagnosis: hypertensive crisis",
        "patient stated diagnosis: stroke diagnosed",
    ],
    "RVC-TEST": [
        "abnormal test: abnormal laboratory result",
        "abnormal test: abnormal imaging result",
        "abnormal test: positive culture",
        "abnormal test: abnormal ecg ekg",
        "abnormal test: positive blood culture",
    ],
    "RVC-DIAG": [
        "diagnostic screening: needs covid test",
        "diagnostic screening: blood pressure check",
        "diagnostic screening: routine testing screening",
        "diagnostic screening: prenatal check",
    ],
    "RVC-TREAT": [
        "treatment medication: prescription refill",
```

```python
        "treatment medication: dressing change",
        "treatment medication: injection therapy",
        "treatment medication: suture removal",
        "treatment: wound check",
    ],
    "RVC-ADMIN": [
        "administrative reason: work form school form insurance paperwork",
        "administrative reason: clearance note",
    ],
    "RVC-UNCL": ["uncodable or unknown reason"],
}

for _group in RVC_NAME:
    RVC_PROTOS_CURATED.setdefault(_group, [])

LABEL2CODE = {
    re.sub(r"\s+", " ", value.strip().lower().replace("-", "-").replace("-",
 ↪   "-")): code
    for code, value in RVC_NAME.items()
}

sym = SymSpell(max_dictionary_edit_distance=2, prefix_length=7)


# --- Lightweight data containers passed between pipeline stages. ---
@dataclass
class SegmentEmbedCache:
    """Embeddings and similarity cache for unique complaint segments."""

    uniq_segments: list[str]
    seg2idx: dict[str, int]
    emb: torch.Tensor
    sim_proto: torch.Tensor


@dataclass
class SegPred:
    """Per-segment prediction record."""

    seg_idx: int
    segment: str
    code: str
    name: str
```

```python
        sim: float
        rule_code: str | None
        rule_used: bool


def _add_lexicon_from_phrases(phrases: list[str]) -> None:
    for phrase in phrases:
        for token in str(phrase).split():
            if token.isalpha():
                sym.create_dictionary_entry(token, 1)


_add_lexicon_from_phrases(list(SEG_OVR.values()) + list(TOK_OVR.values()) +
 ↪  list(ABBR_MAP.values()))


# --- Core text preprocessing helpers. ---
def get_cc_column(df: pd.DataFrame) -> str:
    """Return the first recognized chief-complaint column."""
    for column in CC_CANDIDATES:
        if column in df.columns:
            return column
    raise KeyError(f"No chief-complaint column found in candidates:
     ↪  {CC_CANDIDATES}")


def _clean_tok(token: str) -> str:
    return _PUNCT_EDGE.sub("", token)


def _violates_spell_guard(original: str, candidate: str) -> bool:
    original_norm = original.strip().lower()
    candidate_norm = candidate.strip().lower()
    if not original_norm or not candidate_norm:
        return False
    if original_norm == candidate_norm:
        return False
    for prefix_from, prefix_to in POLARITY_GUARD_PREFIX_PAIRS:
        if original_norm.startswith(prefix_from) and
         ↪  candidate_norm.startswith(prefix_to):
            return True
    if original_norm in LATERALITY_GUARD_TOKENS or candidate_norm in
     ↪  LATERALITY_GUARD_TOKENS:
```

```python
        return original_norm != candidate_norm
    return False


def _spell_candidate(token: str, *, max_edit_distance: int) -> tuple[str, int
↪ | None]:
    candidates = sym.lookup(token, Verbosity.CLOSEST,
↪ max_edit_distance=max_edit_distance)
    if not candidates:
        return token, None
    match = candidates[0]
    return match.term, int(match.distance)


def _spell_strict(token: str) -> tuple[str, dict[str, Any] | None]:
    normalized = token.strip().lower()
    if (
        not normalized
        or normalized in SPELL_PROTECT_TOKENS
        or not normalized.isalpha()
        or len(normalized) < STRICT_SPELL_MIN_TOKEN_LEN
    ):
        return token, None
    corrected, distance = _spell_candidate(normalized,
↪ max_edit_distance=STRICT_SPELL_MAX_EDIT_DISTANCE)
    corrected_norm = corrected.strip().lower()
    if (
        distance is None
        or corrected_norm == normalized
        or corrected_norm in SPELL_PROTECT_TOKENS
        or (normalized, corrected_norm) in SPELL_DENYLIST_SUBSTITUTIONS
        or _violates_spell_guard(normalized, corrected_norm)
    ):
        return token, None
    return corrected_norm, {
        "orig": normalized,
        "corrected": corrected_norm,
        "edit_distance": int(distance),
    }


def _spell_token(token: str, *, mode: str) -> tuple[str, dict[str, Any] |
↪ None]:
```

```python
    if mode == "disabled":
        return token, None
    if mode == "strict":
        return _spell_strict(token)
    raise ValueError(f"Unsupported spell correction mode: {mode}")


def _expand_tokens(segment: str) -> str:
    """Apply contextual and curated abbreviation expansions in one
     ↪ segment."""
    output: list[str] = []
    segment_context = segment
    for token in segment.split():
        cleaned = _clean_tok(token)
        if not cleaned:
            continue
        if cleaned in CTX_OVR:
            predicate, yes_value, no_value = CTX_OVR[cleaned]
            if predicate(segment_context):
                output.extend(yes_value.split())
                continue
            if no_value:
                output.extend(no_value.split())
                continue
        if cleaned in TOK_OVR:
            output.extend(TOK_OVR[cleaned].split())
            continue
        if cleaned in ABBR_MAP:
            output.extend(ABBR_MAP[cleaned].split())
            continue
        output.append(cleaned)
    return " ".join(output)


def _looks_clinical(text: str) -> bool:
    lowered = text.lower()
    return any(key in lowered for key in CLINICAL_KEYS) and
     ↪ bool(re.search(r"[a-z]", lowered))


def _split_on_and_if_clinical(segment: str) -> list[str]:
    parts = [part.strip() for part in re.split(r"\band\b", segment,
 ↪ flags=re.I)]
```

```python
    if len(parts) == 1:
        return parts
    output: list[str] = []
    buffer = parts[0]
    for next_part in parts[1:]:
        if _looks_clinical(buffer) and _looks_clinical(next_part):
            output.append(buffer.strip())
            buffer = next_part
        else:
            buffer = f"{buffer} and {next_part}"
    output.append(buffer.strip())
    return output


def _normalize_absence_symptom(segment: str) -> str | None:
    """Convert selected absence-style phrases into symptom phrases we can
    ↪  classify."""
    cleaned = segment.strip()
    if not cleaned:
        return None

    for regex, replacement in ABSENCE_SYMPTOM_RULES:
        if regex.search(cleaned):
            return regex.sub(replacement, cleaned).strip()

    return None


def _expand_duration(segment: str) -> str:
    def repl(match: re.Match[str]) -> str:
        base = match.group(1)
        count = match.group(2)
        unit = match.group(3).lower()
        resolved_unit = _UNIT_MAP.get(unit, unit)
        return f"{base} for {count} {resolved_unit}".strip()

    return _DUR_RE.sub(repl, segment)


def _lemmatize_segment(segment: str) -> str:
    doc = NLP(segment)
    lemmas: list[str] = []
    for token in doc:
```

```python
            lemma = token.lemma_.strip() if token.lemma_ else ""
            if not lemma or lemma == "-PRON-":
                lemma = token.text
            lemmas.append(lemma)
    return " ".join(lemmas).strip()


def segment_cc(raw: str, keep_negated: bool = False) -> list[str]:
    """Split a raw chief complaint into normalized text segments."""
    if not isinstance(raw, str) or not raw.strip():
        return []

    text = ftfy.fix_text(raw).lower().strip()

    for protected in PROTECT:
        text = text.replace(protected, protected.replace("/", "§"))

    text = re.sub(r"\bn\s*[,/&+]\s*v\s*[,/&+]\s*d\b", "n§v§d", text)
    text = re.sub(r"\bn\s*[,/&+]\s*v\b", "n§v", text)
    text = re.sub(r"\bn\s*/\s*v\s*/\s*d\b", "n§v§d", text)
    text = re.sub(r"\bn\s*/\s*v\b", "n§v", text)
    text = re.sub(r"\bc\s*/\s*p\b", "c§p", text)
    text = re.sub(r"\bs\s*/\s*p\b", "s§p", text)
    text = re.sub(r"\bw\s*/\s*o\b", "w§o", text)
    text = re.sub(r"\bh\s*/\s*o\b", "h§o", text)
    text = re.sub(r"\bc\s*/\s*o\b", "c§o", text)

    text = text.replace("/", "; ")

    segments: list[str] = []
    for block in SEP_SPLIT_RE.split(text):
        if not block:
            continue
        block = block.replace("§", "/").strip()
        block = " ".join(block.split())
        if MISSING_RE.fullmatch(block):
            continue
        for segment in _split_on_and_if_clinical(block):
            if not segment:
                continue

            normalized_absence = _normalize_absence_symptom(segment)
            if normalized_absence is not None:
```

```python
            segment = normalized_absence
        elif not keep_negated and NEG_RE.match(segment):
            continue

        if re.fullmatch(r"[a-z]$", segment):
            continue
        segments.append(segment)
    return segments


def normalize_cc(
    raw: str,
    keep_negated: bool = False,
    *,
    spell_mode: str = "strict",
) -> tuple[list[str], list[dict[str, Any]]]:
    """Normalize a raw chief complaint to deduplicated, lemmatized segments.

    Returns:
        tuple(list[str], list[dict]): normalized segments and accepted
↪   spell-fix log rows.
    """
    if spell_mode not in SPELL_CORRECTION_MODES.difference({"auto"}):
        raise ValueError(f"spell_mode must be one of {{'disabled','strict'}},
        ↪   got {spell_mode!r}")

    output: list[str] = []
    seen: set[str] = set()
    spellfix_log: list[dict[str, Any]] = []

    for segment_index, segment in enumerate(segment_cc(raw,
    ↪   keep_negated=keep_negated)):
        if segment in SEG_OVR:
            segment = SEG_OVR[segment]
        else:
            tokens = segment.split()
            if len(tokens) == 1:
                token = _clean_tok(tokens[0])
                if token in TOK_OVR:
                    segment = TOK_OVR[token]
                else:
                    segment = ABBR_MAP.get(segment, segment)
```

25

```python
            segment = _expand_duration(segment)
            segment = _expand_tokens(segment)
            segment_lower = segment.strip().lower()
            phrase_locked = any(phrase in segment_lower for phrase in
 ↪  SPELL_PROTECT_PHRASES)
            corrected_tokens: list[str] = []
            for token_index, token in enumerate(segment.split()):
                if phrase_locked:
                    corrected, correction = _clean_tok(token), None
                else:
                    corrected, correction = _spell_token(token, mode=spell_mode)
                corrected_tokens.append(corrected)
                if correction is not None:
                    spellfix_log.append(
                        {
                            "segment_idx": int(segment_index),
                            "token_idx": int(token_index),
                            "segment_raw": segment,
                            **correction,
                        }
                    )
            corrected_segment = " ".join(token for token in corrected_tokens if
 ↪  token)
            segment = corrected_segment if phrase_locked else
 ↪  _lemmatize_segment(corrected_segment)

            if segment and segment not in seen:
                output.append(segment)
                seen.add(segment)

    return output, spellfix_log


def _run_normalization_mode(
    cc_series: pd.Series,
    *,
    spell_mode: str,
    keep_negated: bool = False,
) -> dict[str, Any]:
    cleaned_segments: list[list[str]] = []
    spellfix_logs: list[list[dict[str, Any]]] = []
    for raw in cc_series.tolist():
        if pd.isna(raw):
```

```python
            cleaned_segments.append([])
            spellfix_logs.append([])
            continue
        segments, fix_log = normalize_cc(
            str(raw),
            keep_negated=keep_negated,
            spell_mode=spell_mode,
        )
        cleaned_segments.append(segments)
        spellfix_logs.append(fix_log)

    cleaned_canon = [canonize_cc_segments(segments) for segments in
↪   cleaned_segments]
    cleaned_str = ["; ".join(segments) for segments in cleaned_segments]
    cleaned_canon_str = ["; ".join(segments) for segments in cleaned_canon]
    spellfix_n = [len(entries) for entries in spellfix_logs]
    spellfix_any = [count > 0 for count in spellfix_n]
    spellfix_log_json = [json.dumps(entries, ensure_ascii=False) for entries
↪   in spellfix_logs]

    normalized_frame = pd.DataFrame(
        {
            "cc_cleaned": cleaned_segments,
            "cc_cleaned_str": cleaned_str,
            "cc_cleaned_canon": cleaned_canon,
            "cc_cleaned_canon_str": cleaned_canon_str,
            "cc_spellfix_n": pd.Series(spellfix_n, dtype="Int64"),
            "cc_spellfix_any": pd.Series(spellfix_any, dtype="boolean"),
            "cc_spellfix_log": spellfix_log_json,
        },
        index=cc_series.index,
    )
    return {
        "mode": spell_mode,
        "frame": normalized_frame,
        "spellfix_count_total": int(sum(spellfix_n)),
    }


def _score_normalization_integrity(
    raw_series: pd.Series,
    cleaned_str_series: pd.Series,
) -> dict[str, int]:
```

```python
    raw = raw_series.fillna("").astype("string").str.lower()
    clean = cleaned_str_series.fillna("").astype("string").str.lower()
    rules = {
        "fall_to_small": raw.str.contains(r"\bfall\b", na=False)
        & ~raw.str.contains(r"\bsmall\b", na=False)
        & clean.str.contains(r"\bsmall\b", na=False),
        "head_to_heart": raw.str.contains(r"\bhead\b", na=False)
        & ~raw.str.contains(r"\bheart\b", na=False)
        & clean.str.contains(r"\bheart\b", na=False),
        "foot_to_post": raw.str.contains(r"\bfoot\b", na=False)
        & ~raw.str.contains(r"\bpost\b", na=False)
        & ~raw.str.contains(r"\bs/p\b|\bstatus\s+post\b|\bsp\b", na=False)
        & clean.str.contains(r"\bpost\b", na=False),
        "copd_or_cold_to_covid": raw.str.contains(r"\b(?:copd|cold|cord)\b",
        ↪  na=False)
        & ~raw.str.contains(r"\bcovid\b", na=False)
        & clean.str.contains(r"\bcovid\b", na=False),
        "hyperglycemia_to_hypoglycemia":
        ↪  raw.str.contains(r"\bhyperglycemia\b", na=False)
        & clean.str.contains(r"\bhypoglycemia\b", na=False),
        "hyperkalemia_to_hypokalemia": raw.str.contains(r"\bhyperkalemia\b",
        ↪  na=False)
        & clean.str.contains(r"\bhypokalemia\b", na=False),
        "found_down_changed": raw.str.contains(r"\bfound\s+down\b", na=False)
        & ~clean.str.contains(r"\bfound\s+down\b", na=False),
        "found_to_wound": raw.str.contains(r"\bfound\b", na=False)
        & ~raw.str.contains(r"\bwound\b", na=False)
        & clean.str.contains(r"\bwound\b", na=False),
        "trach_to_tract": raw.str.contains(r"\btrach\b", na=False)
        & ~raw.str.contains(r"\btract\b", na=False)
        & clean.str.contains(r"\btract\b", na=False),
        "dental_to_mental": raw.str.contains(r"\bdental\b", na=False)
        & ~raw.str.contains(r"\bmental\b", na=False)
        & clean.str.contains(r"\bmental\b", na=False),
        "black_to_block": raw.str.contains(r"\bblack\b", na=False)
        & ~raw.str.contains(r"\bblock\b", na=False)
        & clean.str.contains(r"\bblock\b", na=False),
    }
    return {name: int(mask.sum()) for name, mask in rules.items()}


def choose_spell_mode(
    df: pd.DataFrame,
```

```python
    *,
    cc_source_column: str,
    requested_mode: str = "auto",
) -> tuple[dict[str, Any], pd.DataFrame]:
    """Run disabled/strict normalization and select a mode by integrity
    ↪  score."""
    requested = requested_mode.strip().lower()
    if requested not in SPELL_CORRECTION_MODES:
        raise ValueError(
            f"CC_SPELL_CORRECTION_MODE must be one of
                ↪  {sorted(SPELL_CORRECTION_MODES)}"
        )
    cc_series = df[cc_source_column].astype("string")
    modes_to_run = ["disabled", "strict"] if requested == "auto" else
↪  [requested]

    mode_runs: dict[str, dict[str, Any]] = {}
    mode_scores: list[dict[str, Any]] = []
    for mode in modes_to_run:
        run_payload = _run_normalization_mode(cc_series, spell_mode=mode)
        mode_runs[mode] = run_payload
        integrity_counts = _score_normalization_integrity(
            cc_series,
            run_payload["frame"]["cc_cleaned_str"],
        )
        integrity_violation_total = int(sum(integrity_counts.values()))
        nonempty_segment_rows = int(
            run_payload["frame"]["cc_cleaned"].map(lambda values: len(values)
↪  > 0).sum()
        )
        mode_scores.append(
            {
                "mode": mode,
                "integrity_violation_total": integrity_violation_total,
                "nonempty_segment_rows": nonempty_segment_rows,
                "spellfix_count_total":
                    ↪  int(run_payload["spellfix_count_total"]),
                **integrity_counts,
            }
        )

    score_df = pd.DataFrame(mode_scores).sort_values(
```

```python
            ["integrity_violation_total", "nonempty_segment_rows",
↪   "spellfix_count_total", "mode"],
            ascending=[True, False, False, True],
            kind="stable",
        )
    selected_mode = str(score_df.iloc[0]["mode"])
    selected_payload = mode_runs[selected_mode]
    selected_payload["selection_scores"] = score_df.reset_index(drop=True)
    selected_payload["requested_mode"] = requested
    selected_payload["selected_mode"] = selected_mode
    return selected_payload, score_df.reset_index(drop=True)


def canonize_cc_segments(segments: list[str]) -> list[str]:
    """Apply light concept-level canonicalization for
     ↪   diagnostics/visualization."""
    output: list[str] = []
    for segment in segments:
        transformed = segment
        for regex, replacement in CANON_REPLACERS:
            transformed = regex.sub(replacement, transformed)
        output.append(transformed)
    return output


def rule_gate(segment: str) -> tuple[str | None, list[str]]:
    """Apply deterministic rule overrides for high-confidence pattern
     ↪   matches."""
    text = f" {segment.lower()} "
    if re.search(r"\babnormal\s+labs?\b", text):
        return "RVC-TEST", ["abnormal labs"]
    if "weakness" in text and RX_FOCAL_WEAK.search(text):
        return "RVC-SYM-NERV", ["focal weakness"]

    hits: list[str] = []
    for code, regex in RULE_RX.items():
        if regex.search(text):
            hits.append(code)

    if not hits:
        return None, []
```

```python
        for preferred_code in PRECEDENCE_RVC:
            if preferred_code in hits:
                return preferred_code, [preferred_code]

        return hits[0], [hits[0]]



# --- Ontology file parsing helpers (for Appendix and summary resources). ---
def _guess_col(columns: list[str], candidates: list[str]) -> str | None:
    lowered = [column.lower().strip() for column in columns]
    for candidate in candidates:
        if candidate in lowered:
            return columns[lowered.index(candidate)]
    return None


def _norm_label(label: str) -> str:
    if label is None:
        return ""
    cleaned = str(label).strip().lower().replace("-", "-").replace("-", "-")
    return re.sub(r"\s+", " ", cleaned)


def map_group(label: str) -> str | None:
    for regex, code in GROUP_MAP:
        if regex.search(str(label)):
            return code
    return None


def _label_to_code(label: str) -> str | None:
    normalized = _norm_label(label)
    if normalized in LABEL2CODE:
        return LABEL2CODE[normalized]
    return map_group(label)


def canon_phrase(text: str) -> str:
    cleaned = str(text).strip().replace("'", "'").replace("-", "-")
    cleaned = RE_SPACE.sub(" ", cleaned)
    return cleaned.strip(" .;:,")
```

```python
def keep_phrase(text: str, group_code: str) -> bool:
    lowered = text.lower()
    if len(lowered) < 3:
        return False
    if RE_DROP.search(lowered):
        return False
    if group_code not in {"RVC-INJ", "RVC-SYM-MSK", "RVC-SYM-EYE",
      ↪  "RVC-SYM-SKIN"}:
        if BODY_PART_ONLY_RE.fullmatch(lowered):
            return False
    return True


def load_appendix_phrases(path: Path) -> pd.DataFrame:
    """Load and normalize Appendix II prototype phrases by RVC code."""
    if not path.exists():
        return pd.DataFrame(columns=["group_code", "phrase"])

    appendix_df = pd.read_csv(path, dtype=str)

    if {"top_level_group17", "text"}.issubset(appendix_df.columns):
        group_col = "top_level_group17"
        phrase_col = "text"
    else:
        group_col = _guess_col(
            list(appendix_df.columns),
            ["top_level_group17", "group", "top_level_group", "rvc_group",
      ↪  "module"],
        )
        phrase_col = _guess_col(
            list(appendix_df.columns),
            ["text", "phrase", "subentry", "entry", "label"],
        )
        if group_col is None or phrase_col is None:
            raise ValueError(f"Unable to identify expected columns in
              ↪  appendix file: {path}")

    normalized = appendix_df[[group_col, phrase_col]].rename(
        columns={group_col: "group_raw", phrase_col: "phrase_raw"}
    )
    normalized["group_code"] = normalized["group_raw"].map(_label_to_code)
    normalized = normalized[normalized["group_code"].notna()].copy()
```

```python
    normalized["phrase"] = normalized["phrase_raw"].map(canon_phrase)
    normalized = normalized[
        normalized.apply(lambda row: keep_phrase(row["phrase"],
↪   row["group_code"]), axis=1)
    ]
    return normalized[["group_code", "phrase"]].drop_duplicates()


def load_summary_weights(path: Path, strategy: str = "uniform") -> dict[str,
↪   float]:
    """Load optional group weights and map them to canonical RVC codes."""
    if not path.exists():
        return {}

    summary_df = pd.read_csv(path, dtype=str)
    group_col = _guess_col(
        list(summary_df.columns),
        ["top_level_group17", "group", "top_level_group", "rvc_group",
↪   "module"],
    )
    if group_col is None:
        group_col = summary_df.columns[0]

    mapped_codes = summary_df[group_col].map(_label_to_code).dropna()
    if mapped_codes.empty:
        return {}

    counts = mapped_codes.value_counts().astype(float)
    if strategy == "uniform":
        counts[:] = 1.0
    elif strategy != "by_rows":
        raise ValueError("strategy must be one of {'uniform', 'by_rows'}")

    scaled = np.log1p(counts) / np.log1p(counts).max()
    return {code: float(max(1e-3, scaled.get(code, 1.0))) for code in
↪       PRECEDENCE_RVC}


def dedup_prototypes(
    rvc_prototypes: dict[str, list[str]],
    encoder: SentenceTransformer,
    sim_thresh: float = 0.92,
    batch_size: int = 256,
```

```python
) -> dict[str, list[str]]:
    """Deduplicate near-identical prototype phrases within each group."""
    output: dict[str, list[str]] = {}

    for group_code, phrases in rvc_prototypes.items():
        normalized = sorted(dict.fromkeys(canon_phrase(phrase) for phrase in
↪   phrases))
        if not normalized:
            output[group_code] = []
            continue

        with torch.inference_mode():
            embedding = encoder.encode(
                normalized,
                convert_to_tensor=True,
                normalize_embeddings=True,
                batch_size=batch_size,
            )

        keep: list[str] = []
        taken = torch.zeros(len(normalized), dtype=torch.bool)
        for idx in range(len(normalized)):
            if taken[idx]:
                continue
            keep.append(normalized[idx])
            if len(normalized) == 1:
                break
            similarities = util.cos_sim(embedding[idx : idx + 1],
↪   embedding).squeeze(0).cpu().numpy()
            duplicate_indices = np.where(similarities >= sim_thresh)[0]
            taken[duplicate_indices] = True

        output[group_code] = keep

    return output


# --- Embedding/prototype resource construction. ---
def build_prototype_resources(cfg: ClassifierConfig) -> dict[str, Any]:
    """Build model, prototype embeddings, scoring indices, and metadata."""
    device = "cuda" if torch.cuda.is_available() else "cpu"
    hf_token = os.getenv(cfg.hf_token_env_var, "").strip() or None
    encoder = SentenceTransformer(
```

```python
        cfg.model_name,
        device=device,
        trust_remote_code=cfg.model_trust_remote_code,
        token=hf_token,
    )

    work_dir = Path(os.getenv(cfg.work_dir_env_var,
↪ Path.cwd())).expanduser().resolve()
    appendix_path = work_dir / cfg.appendix_relpath
    summary_path = work_dir / cfg.summary_relpath

    appendix_df = load_appendix_phrases(appendix_path)
    group_weights = load_summary_weights(summary_path,
↪ strategy=cfg.summary_weight_strategy)

    prototypes_augmented = {code: list(values) for code, values in
↪ RVC_PROTOS_CURATED.items()}
    for group_code, group_df in appendix_df.groupby("group_code"):
        prototypes_augmented[group_code].extend(group_df["phrase"].tolist())

    prototypes_final = dedup_prototypes(
        prototypes_augmented,
        encoder=encoder,
        sim_thresh=cfg.prototype_dedup_sim_thresh,
        batch_size=cfg.model_batch_size,
    )

    all_prototype_texts: list[str] = []
    proto_to_code: list[str] = []
    for code, prototype_list in prototypes_final.items():
        prefix = RVC_NAME[code].split("-")[0].strip().lower()
        for phrase in prototype_list:
            text = phrase if ":" in phrase else f"{prefix}: {phrase}"
            all_prototype_texts.append(text)
            proto_to_code.append(code)

    with torch.inference_mode():
        proto_emb = encoder.encode(
            all_prototype_texts,
            convert_to_tensor=True,
            normalize_embeddings=True,
            batch_size=cfg.model_batch_size,
        )
```

```python
    code_to_proto_indices: dict[str, list[int]] = defaultdict(list)
    for idx, code in enumerate(proto_to_code):
        code_to_proto_indices[code].append(idx)

    proto_weight = np.ones(len(proto_to_code), dtype="float32")
    if group_weights:
        for idx, code in enumerate(proto_to_code):
            proto_weight[idx] = max(1e-3, float(group_weights.get(code,
↪   1.0)))

    metadata = {
        "model_name": cfg.model_name,
        "device": device,
        "prototype_total": len(all_prototype_texts),
        "prototype_count_by_group": {code:
        ↪   len(code_to_proto_indices.get(code, [])) for code in
        ↪   PRECEDENCE_RVC},
        "appendix_rows_used": int(len(appendix_df)),
        "spacy_model": SPACY_MODEL_LABEL,
        "hf_token_present": bool(hf_token),
    }

    return {
        "encoder": encoder,
        "device": device,
        "proto_emb": proto_emb,
        "proto_to_code": proto_to_code,
        "code_to_proto_indices": code_to_proto_indices,
        "proto_weight": proto_weight,
        "metadata": metadata,
    }


# --- Segment caching and scoring helpers. ---
def build_segment_cache(
    df: pd.DataFrame,
    seg_col: str,
    resources: dict[str, Any],
    batch_size: int = 512,
) -> SegmentEmbedCache:
    """Encode unique complaint segments once and cache similarity to all
    ↪   prototypes."""
```

```python
unique_segments: list[str] = []
seen: set[str] = set()

for value in df[seg_col]:
    if isinstance(value, list):
        iterable = value
    elif isinstance(value, str):
        iterable = [value]
    else:
        iterable = []

    for segment in iterable:
        cleaned = str(segment).strip()
        if cleaned and cleaned not in seen:
            unique_segments.append(cleaned)
            seen.add(cleaned)

unique_segments.sort()

if not unique_segments:
    return SegmentEmbedCache(
        uniq_segments=[],
        seg2idx={},
        emb=torch.empty(0, 0),
        sim_proto=torch.empty(0, 0),
    )

encoder: SentenceTransformer = resources["encoder"]
proto_emb: torch.Tensor = resources["proto_emb"]

with torch.inference_mode():
    segment_emb = encoder.encode(
        unique_segments,
        convert_to_tensor=True,
        normalize_embeddings=True,
        batch_size=batch_size,
        device=resources["device"],
    )
    segment_emb = torch.nn.functional.normalize(segment_emb, p=2, dim=1)
    segment_to_proto = segment_emb @ proto_emb.T

seg2idx = {segment: idx for idx, segment in enumerate(unique_segments)}
return SegmentEmbedCache(
```

```python
        uniq_segments=unique_segments,
        seg2idx=seg2idx,
        emb=segment_emb,
        sim_proto=segment_to_proto,
    )


def group_scores_from_proto_row(
    proto_row: torch.Tensor | np.ndarray,
    resources: dict[str, Any],
    method: str = "max",
    alpha: float = 12.0,
) -> dict[str, float]:
    """Aggregate prototype-level similarities to group-level scores."""
    vector = proto_row.detach().cpu().numpy() if isinstance(proto_row,
↪   torch.Tensor) else np.asarray(proto_row)

    scores: dict[str, float] = {}
    code_to_proto = resources["code_to_proto_indices"]
    proto_weight = resources["proto_weight"]

    for group_code, indices in code_to_proto.items():
        values = vector[indices]
        if method == "max":
            score = float(values.max(initial=-1.0))
        elif method == "lse":
            weights = proto_weight[indices]
            logits = alpha * values + np.log(weights)
            max_logit = logits.max()
            score = float(max_logit + np.log(np.exp(logits -
↪   max_logit).sum()))
        else:
            raise ValueError("method must be one of {'max', 'lse'}")
        scores[group_code] = score

    if method == "lse" and scores:
        max_score = max(scores.values())
        min_score = min(scores.values())
        denom = (max_score - min_score) if max_score > min_score else 1.0
        scores = {group: (score - min_score) / denom for group, score in
↪   scores.items()}

    return scores
```

```python
def score_one_segment_soft(
    segment: str,
    resources: dict[str, Any],
    method: str = "max",
    alpha: float = 12.0,
) -> dict[str, float]:
    """Score a single segment directly against prototype embeddings."""
    if not isinstance(segment, str) or not segment.strip():
        return {group: -1.0 for group in RVC_NAME}

    encoder: SentenceTransformer = resources["encoder"]
    proto_emb: torch.Tensor = resources["proto_emb"]

    with torch.inference_mode():
        segment_emb = encoder.encode(
            [segment],
            convert_to_tensor=True,
            normalize_embeddings=True,
            device=resources["device"],
        )
        similarities = util.cos_sim(segment_emb, proto_emb).squeeze(0)

    return group_scores_from_proto_row(similarities, resources=resources,
    ↪   method=method, alpha=alpha)


def classify_segments_cached(
    segments: list[str],
    cache: SegmentEmbedCache,
    resources: dict[str, Any],
    abstain: float = 0.40,
    method: str = "max",
    alpha: float = 12.0,
) -> list[SegPred]:
    """Classify all segments for one visit using cache + rule-gate
    ↪   overrides."""
    predictions: list[SegPred] = []

    for seg_idx, segment in enumerate(segments):
        cleaned = str(segment).strip()
        if not cleaned:
```

```python
            continue

        override_code, _ = rule_gate(cleaned)

        if cleaned in cache.seg2idx:
            row_idx = cache.seg2idx[cleaned]
            score_map = group_scores_from_proto_row(
                cache.sim_proto[row_idx], resources=resources, method=method,
↪  alpha=alpha
            )
        else:
            score_map = score_one_segment_soft(
                cleaned, resources=resources, method=method, alpha=alpha
            )

        best_code, best_score = max(score_map.items(), key=lambda kv: kv[1])

        if override_code is not None:
            best_code = override_code

        if best_score < abstain and override_code is None:
            best_code = "RVC-UNCL"

        predictions.append(
            SegPred(
                seg_idx=seg_idx,
                segment=cleaned,
                code=best_code,
                name=RVC_NAME.get(best_code, best_code),
                sim=float(best_score),
                rule_code=override_code,
                rule_used=override_code is not None,
            )
        )

    return predictions


def rfv_from_segment_preds(segpreds: list[SegPred], max_rfv: int = 5) ->
↪  list[dict[str, Any]]:
    """Aggregate segment-level predictions to visit-level RFV slots."""
    first_position: dict[str, int] = {}
    best_similarity: dict[str, float] = {}
```

```python
    support_segment: dict[str, str] = {}

    for prediction in segpreds:
        group_code = prediction.code
        if group_code == "RVC-UNCL":
            continue

        if group_code not in first_position or prediction.seg_idx <
        ↪  first_position[group_code]:
            first_position[group_code] = prediction.seg_idx
            support_segment[group_code] = prediction.segment

        best_similarity[group_code] = max(best_similarity.get(group_code,
↪  -1.0), prediction.sim)

    ordered = sorted(
        first_position.keys(),
        key=lambda code: (PRECEDENCE_RVC.index(code), first_position[code]),
    )
    ordered = ordered[:max_rfv]

    return [
        {
            "code": code,
            "name": RVC_NAME[code],
            "support": support_segment[code],
            "sim": best_similarity[code],
        }
        for code in ordered
    ]


# --- Visit-level RFV assignment and output shaping. ---
def assign_rvc_per_segment_and_visit_cached(
    df: pd.DataFrame,
    cache: SegmentEmbedCache,
    resources: dict[str, Any],
    seg_col: str = "cc_cleaned",
    cfg: ClassifierConfig | None = None,
) -> pd.DataFrame:
    """Attach standardized RFV outputs to each visit using cached segment
    ↪  embeddings."""
    active_cfg = cfg or ClassifierConfig()
```

```python
    max_rfv = active_cfg.max_rfv
    if active_cfg.empty_primary_policy not in {"blank", "uncodable"}:
        raise ValueError("empty_primary_policy must be one of {'blank',
        ↪  'uncodable'}")

    rows: list[dict[str, Any]] = []
    for value in df[seg_col].tolist():
        if isinstance(value, list):
            segments = [str(segment).strip() for segment in value if
↪  str(segment).strip()]
        elif isinstance(value, str) and value.strip():
            segments = [value.strip()]
        else:
            segments = []

        segments = segments[: active_cfg.max_segments_per_visit]
        segment_predictions = classify_segments_cached(
            segments,
            cache=cache,
            resources=resources,
            abstain=active_cfg.abstain_threshold,
            method=active_cfg.scoring_method,
            alpha=active_cfg.alpha,
        )
        visit_rfv = rfv_from_segment_preds(segment_predictions,
↪  max_rfv=max_rfv)

        # If all non-empty segment predictions are uncodable, explicitly
        ↪  carry that into RFV1.
        if (
            not visit_rfv
            and active_cfg.allow_uncodable_primary
            and segments
            and segment_predictions
        ):
            uncodable_predictions = [pred for pred in segment_predictions if
↪  pred.code == "RVC-UNCL"]
            if uncodable_predictions:
                chosen = max(uncodable_predictions, key=lambda pred:
↪  (pred.sim, -pred.seg_idx))
                visit_rfv = [
                    {
                        "code": "RVC-UNCL",
```

```python
                    "name": RVC_NAME["RVC-UNCL"],
                    "support": chosen.segment,
                    "sim": float(chosen.sim),
                }
            ]

    # Optional policy for truly empty complaints.
    if (
        not visit_rfv
        and not segments
        and active_cfg.allow_uncodable_primary
        and active_cfg.empty_primary_policy == "uncodable"
    ):
        visit_rfv = [
            {
                "code": "RVC-UNCL",
                "name": RVC_NAME["RVC-UNCL"],
                "support": "",
                "sim": math.nan,
            }
        ]

    row: dict[str, Any] = {
        "segment_preds": json.dumps([prediction.__dict__ for prediction
        ↪   in segment_predictions], ensure_ascii=False)
    }

    for idx in range(max_rfv):
        if idx < len(visit_rfv):
            code = visit_rfv[idx]["code"]
            name = visit_rfv[idx]["name"]
            support = visit_rfv[idx]["support"]
            sim = float(visit_rfv[idx]["sim"])
        else:
            code = ""
            name = ""
            support = ""
            sim = math.nan

        slot = idx + 1
        row[f"RFV{slot}"] = RVC_SHORT.get(code, "")
        row[f"RFV{slot}_name"] = name
        row[f"RFV{slot}_support"] = support
```

```python
            row[f"RFV{slot}_sim"] = sim

        rows.append(row)

    drop_columns = [column for column in df.columns if
↪   str(column).startswith("RFV") or column == "segment_preds"]
    output_df = df.drop(columns=drop_columns, errors="ignore").copy()
    output_df = pd.concat([output_df, pd.DataFrame(rows,
↪   index=output_df.index)], axis=1)
    return output_df


# --- Output QA and diagnostic helpers used later in the notebook. ---
def validate_output_schema(
    df: pd.DataFrame,
    max_rfv: int = 5,
    require_primary_for_nonempty: bool = True,
) -> None:
    """Validate required output columns, uniqueness, and JSON shape for
    ↪   segment predictions."""
    required = ["segment_preds"]
    for slot in range(1, max_rfv + 1):
        required.extend(
            [
                f"RFV{slot}",
                f"RFV{slot}_name",
                f"RFV{slot}_support",
                f"RFV{slot}_sim",
            ]
        )

    duplicates = df.columns[df.columns.duplicated()].tolist()
    if duplicates:
        raise ValueError(f"Duplicate columns detected: {duplicates}")

    missing = [column for column in required if column not in df.columns]
    if missing:
        raise ValueError(f"Missing required output columns: {missing}")

    bad_json_rows = 0
    parsed_payloads: list[list[dict[str, Any]] | None] = []
    for value in df["segment_preds"].fillna("[]"):
        if not isinstance(value, str):
```

```python
            bad_json_rows += 1
            parsed_payloads.append(None)
            continue
        try:
            parsed = json.loads(value)
        except json.JSONDecodeError:
            bad_json_rows += 1
            parsed_payloads.append(None)
            continue
        if not isinstance(parsed, list):
            bad_json_rows += 1
            parsed_payloads.append(None)
            continue
        parsed_payloads.append(parsed)

    if bad_json_rows:
        raise ValueError(f"Found {bad_json_rows} malformed segment_preds
        ↪  rows")

    if require_primary_for_nonempty:
        blank_primary_violations = 0
        for primary_name, payload in zip(df["RFV1_name"], parsed_payloads):
            primary_blank = pd.isna(primary_name) or not
↪  str(primary_name).strip()
            has_segments = isinstance(payload, list) and len(payload) > 0
            if primary_blank and has_segments:
                blank_primary_violations += 1

        if blank_primary_violations:
            raise ValueError(
                "Found blank RFV1_name for rows with non-empty segment
                ↪  predictions: "
                f"{blank_primary_violations}"
            )

    print(f"Schema validation passed for {len(df):,} rows.")


def save_with_suffix(
    df: pd.DataFrame,
    working_path: str | Path,
    suffix: str = "_with_NLP",
    sheet_name: str = "with_NLP",
```

45

```python
) -> Path:
    """Write workbook with suffix and return the new path."""
    path = Path(working_path)
    output_path = path.with_name(path.stem + suffix + path.suffix)
    df.to_excel(output_path, index=False, sheet_name=sheet_name)
    print(f"Wrote: {output_path}")
    return output_path


def summarize_overrides(df: pd.DataFrame, column: str = "segment_preds") ->
↪  dict[str, Any]:
    """Summarize rule-override usage from serialized segment prediction
    ↪  payloads."""
    if column not in df.columns:
        raise KeyError(f"Missing column: {column}")

    segment_total = 0
    segment_overridden = 0
    by_rule: Counter = Counter()
    by_final: Counter = Counter()
    rows_total = 0
    rows_with_override = 0

    for value in df[column].dropna():
        rows_total += 1
        payload: Any
        if isinstance(value, str):
            try:
                payload = json.loads(value)
            except Exception:
                continue
        elif isinstance(value, list):
            payload = value
        else:
            continue

        if not isinstance(payload, list):
            continue

        row_has_override = False
        for segment_prediction in payload:
            if not isinstance(segment_prediction, dict):
                continue
```

```python
                segment_total += 1
                if segment_prediction.get("rule_used"):
                    segment_overridden += 1
                    row_has_override = True
                    by_rule[segment_prediction.get("rule_code", "UNKNOWN")] += 1
                    by_final[segment_prediction.get("code", "UNKNOWN")] += 1

        if row_has_override:
            rows_with_override += 1

    return {
        "segments_total": segment_total,
        "segments_overridden": segment_overridden,
        "segment_override_rate": (segment_overridden / segment_total) if
        ↪   segment_total else float("nan"),
        "rows_total": rows_total,
        "rows_with_override": rows_with_override,
        "row_override_rate": (rows_with_override / rows_total) if rows_total
        ↪   else float("nan"),
        "by_rule": by_rule,
        "by_final": by_final,
    }


def validate_pipeline(
    df: pd.DataFrame,
    cache: SegmentEmbedCache,
    resources: dict[str, Any],
    seg_col: str = "cc_cleaned",
) -> dict[str, Any]:
    """Run smoke checks for prototype health, segment coverage, and
    ↪   rule/classifier sanity."""
    metadata = resources["metadata"]
    prototype_counts = metadata["prototype_count_by_group"]
    missing_groups = [group for group, count in prototype_counts.items() if
↪   count == 0]
    if missing_groups:
        raise ValueError(f"Groups missing prototypes: {missing_groups}")

    row_count = len(df)
    nonempty_segment_rows = int(df[seg_col].apply(lambda value:
↪   isinstance(value, list) and len(value) > 0).sum())
    unique_segments = len(cache.uniq_segments)
```

```python
    probes = [
        "shortness of breath",
        "bradycardia",
        "positive blood culture",
        "alter mental status",
        "abdominal pain / shortness of breath",
        "right eye redness / pain",
    ]
    probe_results: dict[str, list[tuple[str, str, float, bool]]] = {}
    for probe in probes:
        segments = segment_cc(probe)
        predictions = classify_segments_cached(
            segments,
            cache=cache,
            resources=resources,
            method="max",
        )
        probe_results[probe] = [
            (prediction.segment, prediction.code, prediction.sim,
  ↪ prediction.rule_used)
            for prediction in predictions
        ]

    print("Prototype integrity:")
    print(prototype_counts)
    print(f"Total prototypes: {metadata['prototype_total']}")
    print()
    print("Segment column health:")
    print(f"rows={row_count:,} | with>=1 segment={nonempty_segment_rows:,} |
  ↪  unique segments={unique_segments:,}")
    print()
    print("Probe classifications:")
    for probe, result in probe_results.items():
        print(f"{probe} -> {result}")

    return {
        "row_count": row_count,
        "rows_with_segments": nonempty_segment_rows,
        "unique_segments": unique_segments,
        "prototype_counts": prototype_counts,
        "probe_results": probe_results,
    }
```

```python
def top_cc_plot(
    df: pd.DataFrame,
    candidates: tuple[str, ...] = CC_CANDIDATES,
    top_n: int = 30,
    normalize_case: str | None = "upper",
) -> pd.DataFrame:
    """Optional EDA: plot top raw chief complaints."""
    cc_column = get_cc_column(df)
    series = (
        df[cc_column]
        .astype(str)
        .str.strip()
        .str.replace(r"\s+", " ", regex=True)
    )

    if normalize_case == "upper":
        series = series.str.upper()
    elif normalize_case == "title":
        series = series.str.title()

    counts = series.value_counts().head(top_n).sort_values(ascending=True)
    axis = counts.plot(kind="barh", figsize=(8, 10))
    axis.set_xlabel("Count")
    axis.set_ylabel("Chief complaint")
    axis.set_title(f"Top {top_n} Chief Complaints ({cc_column})")
    plt.tight_layout()
    plt.show()

    return counts.sort_values(ascending=False).rename_axis("chief_complaint"⌋
      ↪  ).reset_index(name="count")
```

1.3 2B) Local helper functions (embedded core notebook logic)

```python
# Core stage helpers are embedded in this notebook by project policy.

DATA_DIRNAME = "MIMIC tabular data"
PRIOR_RUNS_DIRNAME = "prior runs"

CANONICAL_COHORT_FILENAME = "MIMICIV all with CC.xlsx"
CANONICAL_NLP_FILENAME = "MIMICIV all with CC_with_NLP.xlsx"
```

```python
CLASSIFIER_TRANSITIONAL_ALIASES: dict[str, tuple[str, ...]] = {
    "age": ("age_at_admit",),
    "hr": ("ed_first_hr_model", "ed_first_hr", "ed_triage_hr_model",
     ↪  "ed_triage_hr"),
    "rr": ("ed_first_rr_model", "ed_first_rr", "ed_triage_rr_model",
     ↪  "ed_triage_rr"),
    "sbp": ("ed_first_sbp_model", "ed_first_sbp", "ed_triage_sbp_model",
     ↪  "ed_triage_sbp"),
    "dbp": ("ed_first_dbp_model", "ed_first_dbp", "ed_triage_dbp_model",
     ↪  "ed_triage_dbp"),
    "temp": ("ed_first_temp_model", "ed_first_temp", "ed_triage_temp_model",
     ↪  "ed_triage_temp"),
    "spo2": (
        "ed_first_o2sat_model",
        "ed_first_o2sat",
        "ed_triage_o2sat_model",
        "ed_triage_o2sat",
    ),
    "race": ("ed_race", "race_ed_raw", "hosp_race", "race_hosp_raw"),
}


def data_dir(work_dir: Path) -> Path:
    """Return canonical data directory rooted at ``work_dir``."""
    return (work_dir / DATA_DIRNAME).expanduser().resolve()


def resolve_classifier_input_path(
    work_dir: Path, input_filename: str | None = None
) -> Path:
    """Resolve classifier input workbook under canonical data directory."""
    filename = input_filename or CANONICAL_COHORT_FILENAME
    input_path = data_dir(work_dir) / filename
    if not input_path.exists():
        raise FileNotFoundError(
            "Expected classifier input workbook was not found at "
            f"{input_path}. Run the cohort notebook first or set "
            "CLASSIFIER_INPUT_FILENAME."
        )
    return input_path
```

```python
def resolve_analysis_input_path(work_dir: Path, input_filename: str | None =
↪   None) -> Path:
    """Resolve analysis input workbook under canonical data directory."""
    filename = input_filename or CANONICAL_NLP_FILENAME
    input_path = data_dir(work_dir) / filename
    if not input_path.exists():
        raise FileNotFoundError(
            "Expected analysis input workbook was not found at "
            f"{input_path}. Run the classifier notebook first or set "
            "ANALYSIS_INPUT_FILENAME."
        )
    return input_path


def resolve_rater_nlp_input_path(work_dir: Path, input_filename: str | None =
↪   None) -> Path:
    """Resolve rater notebook NLP input workbook under canonical data
    ↪   directory."""
    filename = input_filename or CANONICAL_NLP_FILENAME
    input_path = data_dir(work_dir) / filename
    if not input_path.exists():
        raise FileNotFoundError(
            "Expected rater NLP input workbook was not found at "
            f"{input_path}. Run the classifier notebook first or set "
            "RATER_NLP_INPUT_FILENAME."
        )
    return input_path


def resolve_classifier_output_paths(
    work_dir: Path,
    *,
    run_dt: datetime | None = None,
) -> tuple[Path, Path]:
    """Return canonical NLP output path and dated archive path."""
    base_data_dir = data_dir(work_dir)
    canonical_path = base_data_dir / CANONICAL_NLP_FILENAME

    run_date = (run_dt or datetime.now()).strftime("%Y-%m-%d")
    archive_dir = base_data_dir / PRIOR_RUNS_DIRNAME
    archive_path = archive_dir / f"{run_date} {CANONICAL_NLP_FILENAME}"

    return canonical_path, archive_path
```

```python
def normalize_classifier_input_schema(
    df: pd.DataFrame,
    alias_map: Mapping[str, str | Sequence[str]] | None = None,
) -> pd.DataFrame:
    """Add transitional alias columns for classifier compatibility.

    Existing destination columns are preserved. Source columns are never
 ↪  removed.
    """
    resolved_alias_map = alias_map or CLASSIFIER_TRANSITIONAL_ALIASES
    normalized = df.copy()
    for destination, source_spec in resolved_alias_map.items():
        if destination in normalized.columns:
            continue
        if isinstance(source_spec, str):
            candidate_sources = (source_spec,)
        else:
            candidate_sources = tuple(source_spec)
        for source_name in candidate_sources:
            if source_name in normalized.columns:
                normalized[destination] = normalized[source_name]
                break
    return normalized


def ensure_required_columns(
    df: pd.DataFrame,
    required: list[str],
    *,
    context: str,
) -> None:
    """Validate required columns with context-aware error message."""
    missing = sorted(set(required).difference(df.columns))
    if missing:
        raise KeyError(f"{context}: missing required columns: {missing}")

PSEUDO_MISSING_EXACT_TOKENS = {
    "-",
    "--",
    "_",
    "n",
```

```python
        "na",
        "n/a",
        "none",
        "unknown",
        "refused",
}
PSEUDO_MISSING_REGEX = re.compile(r"^_+$")
PSEUDO_MISSING_PUNCT_REGEX = re.compile(r"^[^a-z0-9]+$")


def _canonicalize_cc_text(value: object) -> str:
    text = "" if pd.isna(value) else str(value)
    text = text.strip().lower()
    text = text.replace("-", "-").replace("-", "-")
    return " ".join(text.split())


def annotate_cc_missingness(
    df: pd.DataFrame,
    *,
    cc_column: str,
    nlp_input_column: str = "cc_text_for_nlp",
    missing_flag_column: str = "cc_missing_flag",
    pseudo_flag_column: str = "cc_pseudomissing_flag",
    missing_reason_column: str = "cc_missing_reason",
) -> pd.DataFrame:
    """Annotate true-vs-pseudo missing CCs and derive deterministic NLP input
    ↪  text."""
    if cc_column not in df.columns:
        raise KeyError(f"Chief complaint column not found: {cc_column}")

    out = df.copy()
    canonical = out[cc_column].map(_canonicalize_cc_text)
    is_true_missing = canonical.eq("")
    is_pseudo_missing = (
        canonical.isin(PSEUDO_MISSING_EXACT_TOKENS)
        | canonical.str.match(PSEUDO_MISSING_REGEX, na=False)
        | canonical.str.match(PSEUDO_MISSING_PUNCT_REGEX, na=False)
    )

    reason = pd.Series("valid", index=out.index, dtype="string")
    reason.loc[is_true_missing] = "true_missing"
    reason.loc[is_pseudo_missing] = "pseudo_missing_token"
```

```python
    out[missing_flag_column] = (is_true_missing |
↪    is_pseudo_missing).astype("boolean")
    out[pseudo_flag_column] = is_pseudo_missing.astype("boolean")
    out[missing_reason_column] = reason
    out[nlp_input_column] = out[cc_column].astype("string")
    out.loc[is_true_missing | is_pseudo_missing, nlp_input_column] = pd.NA
    return out


def apply_pseudomissing_uncodable_policy(
    df: pd.DataFrame,
    *,
    max_rfv: int = 5,
    pseudo_flag_column: str = "cc_pseudomissing_flag",
) -> pd.DataFrame:
    """Force pseudo-missing CC rows to a deterministic uncodable RFV
↪    payload."""
    if pseudo_flag_column not in df.columns:
        return df

    out = df.copy()
    pseudo_mask = out[pseudo_flag_column].fillna(False).astype(bool)
    if not pseudo_mask.any():
        return out

    payload = json.dumps(
        [
            {
                "seg_idx": 0,
                "segment": "",
                "code": "RVC-UNCL",
                "name": "Uncodable/Unknown",
                "sim": None,
                "rule_code": "pseudo_missing_cc",
                "rule_used": True,
            }
        ]
    )
    _apply_uncodable_payload(
        out,
        mask=pseudo_mask,
        max_rfv=max_rfv,
```

```python
        payload=payload,
    )
    return out


def apply_blank_primary_uncodable_policy(
    df: pd.DataFrame,
    *,
    max_rfv: int = 5,
    missing_reason_column: str = "cc_missing_reason",
) -> pd.DataFrame:
    """Force non-true-missing blank RFV primaries to deterministic uncodable
     ↪   output."""
    if "RFV1_name" not in df.columns:
        return df

    out = df.copy()
    blank_primary =
 ↪   out["RFV1_name"].fillna("").astype(str).str.strip().eq("")
    true_missing = (
        out[missing_reason_column].astype(str).eq("true_missing")
        if missing_reason_column in out.columns
        else pd.Series(False, index=out.index)
    )
    fallback_mask = blank_primary & ~true_missing
    if not fallback_mask.any():
        return out

    payload = json.dumps(
        [
            {
                "seg_idx": 0,
                "segment": "",
                "code": "RVC-UNCL",
                "name": "Uncodable/Unknown",
                "sim": None,
                "rule_code": "blank_primary_fallback",
                "rule_used": True,
            }
        ]
    )
    _apply_uncodable_payload(
        out,
```

```python
        mask=fallback_mask,
        max_rfv=max_rfv,
        payload=payload,
    )
    return out


def _apply_uncodable_payload(
    out: pd.DataFrame,
    *,
    mask: pd.Series,
    max_rfv: int,
    payload: str,
) -> None:
    """Apply a deterministic uncodable RFV payload in-place."""
    out.loc[mask, "RFV1"] = "uncodable"
    out.loc[mask, "RFV1_name"] = "Uncodable/Unknown"
    out.loc[mask, "RFV1_support"] = ""
    out.loc[mask, "RFV1_sim"] = math.nan

    for slot in range(2, max_rfv + 1):
        out.loc[mask, f"RFV{slot}"] = ""
        out.loc[mask, f"RFV{slot}_name"] = ""
        out.loc[mask, f"RFV{slot}_support"] = ""
        out.loc[mask, f"RFV{slot}_sim"] = math.nan

    out.loc[mask, "segment_preds"] = payload


def build_cc_missing_audit(
    df: pd.DataFrame,
    *,
    cc_column: str,
    missing_reason_column: str = "cc_missing_reason",
    max_examples: int = 5,
) -> pd.DataFrame:
    """Summarize CC missingness categories with example raw values."""
    if cc_column not in df.columns:
        raise KeyError(f"Chief complaint column not found: {cc_column}")
    if missing_reason_column not in df.columns:
        raise KeyError(f"Missing reason column not found:
        ↪    {missing_reason_column}")
```

```python
    frame = df[[cc_column, missing_reason_column]].copy()
    frame[cc_column] = frame[cc_column].astype("string")
    rows: list[dict[str, Any]] = []
    for reason, group in frame.groupby(missing_reason_column, dropna=False):
        examples = (
            group[cc_column]
            .fillna("<NA>")
            .astype(str)
            .value_counts(dropna=False)
            .head(max_examples)
            .index.tolist()
        )
        rows.append(
            {
                "cc_missing_reason": str(reason),
                "row_count": int(len(group)),
                "examples": "; ".join(examples),
            }
        )
    return pd.DataFrame(rows).sort_values("cc_missing_reason"
    ↪    ).reset_index(drop=True)


def build_phrase_regression_audit(
    df: pd.DataFrame,
    *,
    cc_column: str,
    segment_preds_column: str = "segment_preds",
    rfv_column: str = "RFV1_name",
) -> tuple[pd.DataFrame, pd.DataFrame]:
    """Audit phrase-level normalization regressions for abnormal-lab
    ↪    complaints."""
    if cc_column not in df.columns:
        raise KeyError(f"Chief complaint column not found: {cc_column}")
    if segment_preds_column not in df.columns:
        raise KeyError(f"segment prediction column not found:
        ↪    {segment_preds_column}")
    if rfv_column not in df.columns:
        raise KeyError(f"RFV column not found: {rfv_column}")

    raw_cc = df[cc_column].fillna("").astype(str)
    abnormal_lab_mask = raw_cc.str.contains(r"\babnormal\s+labs?\b",
↪    case=False, regex=True)
```

```python
    subset = df.loc[abnormal_lab_mask, [cc_column, segment_preds_column,
↪  rfv_column]].copy()

    abnormal_loss_rows = 0
    for payload_text in subset[segment_preds_column].fillna("[]"):
        try:
            payload = json.loads(payload_text) if isinstance(payload_text,
↪  str) else []
        except json.JSONDecodeError:
            payload = []
        row_segments: list[str] = []
        for record in payload if isinstance(payload, list) else []:
            if not isinstance(record, dict):
                continue
            segment_text = str(record.get("segment", "")).strip().lower()
            if not segment_text:
                continue
            row_segments.append(segment_text)
        if any("abnormal loss" in segment for segment in row_segments):
            abnormal_loss_rows += 1

    rfv_counts = (
        subset[rfv_column]
        .fillna("NA")
        .astype(str)
        .value_counts(dropna=False)
        .rename_axis("rfv1_name")
        .reset_index(name="row_count")
    )
    if not rfv_counts.empty:
        rfv_counts["row_rate"] = rfv_counts["row_count"] / max(len(subset),
↪  1)

    summary = pd.DataFrame(
        [
            {"metric": "abnormal_lab_rows", "value": int(len(subset))},
            {"metric": "abnormal_loss_rows", "value":
↪  int(abnormal_loss_rows)},
            {
                "metric": "abnormal_test_result_rows",
                "value": int(
                    subset[rfv_column]
                    .fillna("")
```

```python
                    .astype(str)
                    .str.contains(r"Abnormal test result", case=False,
↪   regex=True)
                    .sum()
                ),
            },
            {
                "metric": "eye_ear_rows",
                "value": int(
                    subset[rfv_column]
                    .fillna("")
                    .astype(str)
                    .str.contains(r"Eye/Ear", case=False, regex=True)
                    .sum()
                ),
            },
        ]
    )
    return summary, rfv_counts


def classifier_resource_paths(
    work_dir: Path,
    *,
    appendix_relpath: str,
    summary_relpath: str,
) -> dict[str, Path]:
    """Return required classifier resource paths under ``work_dir``."""
    return {
        "appendix_csv": (work_dir / appendix_relpath).expanduser().resolve(),
        "summary_weights_csv": (work_dir /
        ↪   summary_relpath).expanduser().resolve(),
    }


def _sha256(path: Path) -> str:
    hasher = hashlib.sha256()
    with path.open("rb") as handle:
        for chunk in iter(lambda: handle.read(8192), b""):
            hasher.update(chunk)
    return hasher.hexdigest()
```

```python
def verify_classifier_resources(
    work_dir: Path,
    *,
    appendix_relpath: str,
    summary_relpath: str,
    manifest_path: str | None = "Annotation/resource_manifest.json",
    strict_hash: bool = True,
) -> dict[str, Any]:
    """Validate required classifier resources and optional hash manifest."""
    checks: list[dict[str, Any]] = []
    missing_paths: list[str] = []
    for name, path in classifier_resource_paths(
        work_dir,
        appendix_relpath=appendix_relpath,
        summary_relpath=summary_relpath,
    ).items():
        exists = path.exists()
        checks.append({"resource": name, "path": str(path), "exists":
↪ exists})
        if not exists:
            missing_paths.append(str(path))

    hash_findings: list[dict[str, Any]] = []
    manifest_resolved: Path | None = None
    if manifest_path:
        manifest_resolved = (work_dir / manifest_path).expanduser().resolve()
        if manifest_resolved.exists():
            manifest = json.loads(manifest_resolved.read_text())
            resources = manifest.get("resources", [])
            for entry in resources:
                rel_path = entry.get("path")
                expected = str(entry.get("sha256", "")).strip().lower()
                if not rel_path or not expected:
                    continue
                path = (work_dir / rel_path).expanduser().resolve()
                if not path.exists():
                    hash_findings.append(
                        {
                            "path": str(path),
                            "status": "missing",
                            "expected_sha256": expected,
                            "actual_sha256": None,
                        }
```

```python
                )
                continue
            actual = _sha256(path)
            if actual != expected:
                hash_findings.append(
                    {
                        "path": str(path),
                        "status": "hash_mismatch",
                        "expected_sha256": expected,
                        "actual_sha256": actual,
                    }
                )

    errors: list[str] = []
    warnings: list[str] = []

    if missing_paths:
        errors.append(
            "Missing required classifier resources: " + ",
            ↪  ".join(sorted(missing_paths))
        )

    for finding in hash_findings:
        message = (
            f"{finding['status']} for {finding['path']} "
            f"(expected={finding['expected_sha256']},
            ↪  actual={finding['actual_sha256']})"
        )
        if strict_hash:
            errors.append(message)
        else:
            warnings.append(message)

    status = "pass"
    if errors:
        status = "fail"
    elif warnings:
        status = "warning"

    return {
        "status": status,
        "strict_hash": bool(strict_hash),
```

```python
        "manifest_path": str(manifest_resolved) if manifest_resolved else
        ↪   None,
        "checks": checks,
        "hash_findings": hash_findings,
        "errors": errors,
        "warnings": warnings,
    }


def validate_classifier_contract(
    df: pd.DataFrame,
    *,
    max_rfv: int = 5,
    pseudo_flag_column: str = "cc_pseudomissing_flag",
    missing_reason_column: str = "cc_missing_reason",
    missing_flag_column: str = "cc_missing_flag",
    authoritative_disagreement_tolerance: float = 0.0,
) -> dict[str, Any]:
    """Run deterministic post-export classifier contract checks."""
    findings: list[dict[str, str]] = []

    if "hadm_id" in df.columns:
        duplicates = int(df.duplicated(subset=["hadm_id"]).sum())
        if duplicates:
            findings.append(
                {
                    "severity": "error",
                    "code": "hadm_id_not_unique",
                    "message": f"Found {duplicates} duplicate hadm_id rows.",
                }
            )

    if "segment_preds" not in df.columns:
        findings.append(
            {
                "severity": "error",
                "code": "missing_segment_preds",
                "message": "segment_preds column missing from classifier
                ↪   output.",
            }
        )
    else:
        bad_json_rows = 0
```

```python
        for value in df["segment_preds"].fillna("[]"):
            try:
                parsed = json.loads(value) if isinstance(value, str) else
↪    value
            except json.JSONDecodeError:
                bad_json_rows += 1
                continue
            if not isinstance(parsed, list):
                bad_json_rows += 1
        if bad_json_rows:
            findings.append(
                {
                    "severity": "error",
                    "code": "segment_preds_malformed",
                    "message": f"Found {bad_json_rows} malformed
                      ↪    segment_preds rows.",
                }
            )

    required_cc_columns = {missing_reason_column}
    missing_cc_columns = sorted(required_cc_columns.difference(df.columns))
    if missing_cc_columns:
        findings.append(
            {
                "severity": "error",
                "code": "missing_cc_missingness_columns",
                "message": f"Missing CC missingness columns:
                  ↪    {missing_cc_columns}",
            }
        )

    if "RFV1_name" in df.columns:
        blank_rfv1 =
↪    df["RFV1_name"].fillna("").astype(str).str.strip().eq("")
        if missing_reason_column in df.columns:
            allowed_blank =
↪    df[missing_reason_column].astype(str).eq("true_missing")
            unexpected_blank = int((blank_rfv1 & ~allowed_blank).sum())
        else:
            unexpected_blank = int(blank_rfv1.sum())
        if unexpected_blank:
            findings.append(
                {
```

```python
                    "severity": "error",
                    "code": "unexpected_blank_rfv1",
                    "message": (
                        "Found blank RFV1_name rows outside allowed
                         ↪  true-missing policy: "
                        f"{unexpected_blank}"
                    ),
                }
            )

    pseudo_count = 0
    if missing_reason_column in df.columns:
        pseudo_count = int(df[missing_reason_column].astype(str).eq(
↪   "pseudo_missing_token").sum())

    if pseudo_flag_column in df.columns and missing_reason_column in
     ↪  df.columns:
        pseudo_flag_series =
↪   df[pseudo_flag_column].fillna(False).astype(bool)
        pseudo_expected_series =
↪   df[missing_reason_column].astype(str).eq("pseudo_missing_token")
        pseudo_flag_mismatch = int((pseudo_flag_series !=
↪   pseudo_expected_series).sum())
        if pseudo_flag_mismatch:
            findings.append(
                {
                    "severity": "warning",
                    "code": "cc_pseudomissing_flag_mismatch",
                    "message": (
                        "cc_pseudomissing_flag mismatches cc_missing_reason
                         ↪  pseudo_missing_token rows: "
                        f"{pseudo_flag_mismatch}"
                    ),
                }
            )

    if missing_flag_column in df.columns and missing_reason_column in
     ↪  df.columns:
        missing_flag_series =
↪   df[missing_flag_column].fillna(False).astype(bool)
        missing_expected_series =
↪   df[missing_reason_column].astype(str).ne("valid")
```

64

```python
        missing_flag_mismatch = int((missing_flag_series !=
↪  missing_expected_series).sum())
        if missing_flag_mismatch:
            findings.append(
                {
                    "severity": "warning",
                    "code": "cc_missing_flag_mismatch",
                    "message": (
                        "cc_missing_flag mismatches cc_missing_reason
                          ↪  non-valid rows: "
                        f"{missing_flag_mismatch}"
                    ),
                }
            )

    error_count = sum(1 for finding in findings if finding["severity"] ==
↪  "error")
    warning_count = sum(1 for finding in findings if finding["severity"] ==
↪  "warning")
    status = "pass"
    if error_count:
        status = "fail"
    elif warning_count:
        status = "warning"

    return {
        "status": status,
        "error_count": error_count,
        "warning_count": warning_count,
        "pseudo_missing_rows": pseudo_count,
        "findings": findings,
        "row_count": int(len(df)),
        "column_count": int(df.shape[1]),
        "max_rfv": int(max_rfv),
    }


def render_latex_longtable(
    table_df: pd.DataFrame,
    *,
    caption: str,
    label: str,
    landscape: bool = False,
    index: bool = True,
```

```python
) -> str:
    latex_text = table_df.to_latex(
        index=index,
        escape=False,
        longtable=True,
        caption=caption,
        label=label,
    )
    if landscape:
        latex_text = "\\begin{landscape}\\n" + latex_text +
↪    "\\n\\end{landscape}\\n"
    return latex_text
```

1.4 3) Input load and schema checks

```python
# Purpose: instantiate config, set deterministic seeds, load the input
↪    workbook, and validate required identifiers.
# Why this matters: stable setup reduces run-to-run variability and catches
↪    input issues before expensive model steps.

cfg = ClassifierConfig()

try:
    NLP = spacy.load("en_core_web_sm", disable=["ner", "parser", "textcat",
↪    "senter"])
    SPACY_MODEL_LABEL = "en_core_web_sm"
except OSError as exc:
    if cfg.require_spacy_model:
        raise ImportError(
            "spaCy model 'en_core_web_sm' is required but not installed in
                ↪    this environment. "
            "Install it once with: python -m spacy download en_core_web_sm"
        ) from exc
    NLP = spacy.blank("en")
    SPACY_MODEL_LABEL = "spacy.blank('en')"

print(f"spaCy pipeline: {SPACY_MODEL_LABEL}")

np.random.seed(cfg.seed)
torch.manual_seed(cfg.seed)
if torch.cuda.is_available():
    torch.cuda.manual_seed_all(cfg.seed)
```

```python
load_dotenv()
WORK_DIR = Path(os.getenv(cfg.work_dir_env_var,
 ↪  Path.cwd())).expanduser().resolve()
SRC_DIR = WORK_DIR / "src"
if str(SRC_DIR) not in sys.path:
    sys.path.insert(0, str(SRC_DIR))

from hypercap_cc_nlp.pipeline_contracts import write_contract_report
from hypercap_cc_nlp.pipeline_audit import collect_run_manifest
ensure_required_workflow_columns = ensure_required_columns

classifier_input_filename = os.getenv(cfg.input_filename_env_var,
 ↪  cfg.input_filename)
file_path = resolve_classifier_input_path(WORK_DIR,
 ↪  classifier_input_filename)

strict_resource_hash = os.getenv("CLASSIFIER_STRICT_RESOURCE_HASH",
 ↪  "1").strip() == "1"
resource_report = verify_classifier_resources(
    WORK_DIR,
    appendix_relpath=cfg.appendix_relpath,
    summary_relpath=cfg.summary_relpath,
    strict_hash=strict_resource_hash,
)
if resource_report["status"] == "fail":
    raise FileNotFoundError(
        "Classifier resource verification failed: "
        + "; ".join(resource_report["errors"])
    )
if resource_report["warnings"]:
    print("Resource verification warnings:")
    for warning_message in resource_report["warnings"]:
        print("-", warning_message)

df = pd.read_excel(file_path, sheet_name=0, engine="openpyxl")
df.columns = df.columns.str.strip()
pre_alias_columns = set(df.columns)
df = normalize_classifier_input_schema(df)
added_aliases = sorted(set(df.columns).difference(pre_alias_columns))

ensure_required_workflow_columns(
    df,
```

```
        ["hadm_id", "subject_id"],
        context="classifier input",
    )

    print(f"Loaded input rows={len(df):,}, cols={df.shape[1]:,}")
    print(f"Input path: {file_path}")
    print(f"Added transitional aliases: {added_aliases if added_aliases else
     ↪  'None'}")
    cc_column = get_cc_column(df)
    print(f"Detected chief complaint column: {cc_column}")
```

```
spaCy pipeline: en_core_web_sm
Loaded input rows=7,152, cols=305
Input path: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Pro
CC-NLP/MIMIC tabular data/MIMICIV all with CC.xlsx
Added transitional aliases: ['age', 'dbp', 'hr', 'race', 'rr', 'sbp', 'spo2', 'temp']
Detected chief complaint column: ed_triage_cc
```

1.5 4) Text cleaning and segmentation functions

```
# Purpose: define data-enrichment helpers and generate normalized
 ↪  chief-complaint text used by the classifier.
# Why this matters: high-quality text normalization is foundational for both
 ↪  rule matching and embedding similarity.

df = annotate_cc_missingness(df, cc_column=cc_column)
requested_spell_mode = os.getenv("CC_SPELL_CORRECTION_MODE",
 ↪  "auto").strip().lower()
normalization_payload, spell_mode_scores = choose_spell_mode(
    df,
    cc_source_column="cc_text_for_nlp",
    requested_mode=requested_spell_mode,
)
df = pd.concat([df, normalization_payload["frame"]], axis=1)
df["cc_spell_mode_selected"] = normalization_payload["selected_mode"]

cc_missing_audit = build_cc_missing_audit(df, cc_column=cc_column)
print("CC missingness audit:")
print(cc_missing_audit.to_string(index=False))
print("Pseudo-missing rows:",
 ↪  int(df["cc_pseudomissing_flag"].fillna(False).sum()))
```

```
print("Spell correction mode requested:",
↪  normalization_payload["requested_mode"])
print("Spell correction mode selected:",
↪  normalization_payload["selected_mode"])
print("Spell mode integrity comparison:")
print(spell_mode_scores.to_string(index=False))
print("Total accepted spell fixes:",
↪  int(df["cc_spellfix_n"].fillna(0).sum()))

print("Text normalization complete.")
print(
    df[
        [
            cc_column,
            "cc_missing_flag",
            "cc_missing_reason",
            "cc_pseudomissing_flag",
            "cc_spell_mode_selected",
            "cc_spellfix_n",
            "cc_cleaned",
            "cc_cleaned_str",
        ]
    ]
    .head(5)
    .to_string()
)
```

```
CC missingness audit:
   cc_missing_reason  row_count                                                  example
pseudo_missing_token          2                                                        -
; ___
              valid       7150 Dyspnea; Altered mental status; Respiratory distress; MVC; S(
Pseudo-missing rows: 2
Spell correction mode requested: auto
Spell correction mode selected: strict
Spell mode integrity comparison:
    mode  integrity_violation_total  nonempty_segment_rows  spellfix_count_total  fall_to_sma
  strict                          0                   7150                    16
disabled                          0                   7150                     0
Total accepted spell fixes: 16
Text normalization complete.
                   ed_triage_cc  cc_missing_flag cc_missing_reason  cc_pseudomissing_flag
0                       DYSPNEA            False             valid                  False
```

```
1                              SOB         False         valid            False
2  Altered mental status, Vomiting         False         valid            False
3              Anemia, Dyspnea            False         valid            False
4                          Dyspnea         False         valid            False
```

1.6 5) Rule system and ontology mappings

```python
# Purpose: run a quick rule-coverage spot check on representative phrases.
# Why this matters: confirms deterministic override logic is wired correctly
↪   before full inference.

print("Rule coverage preview:")
probe_texts = [
    "shortness of breath",
    "bradycardia",
    "positive blood culture",
    "alter mental status",
    "abdominal pain",
    "right eye redness / pain",
    "paperwork clearance",
]
for text in probe_texts:
    override, _ = rule_gate(text)
    print(f"{text!r} -> {override}")


Rule coverage preview:
'shortness of breath' -> RVC-SYM-RESP
'bradycardia' -> RVC-SYM-CIRC
'positive blood culture' -> RVC-TEST
'alter mental status' -> RVC-SYM-NERV
'abdominal pain' -> RVC-SYM-DIG
'right eye redness / pain' -> RVC-SYM-EYE
'paperwork clearance' -> RVC-ADMIN
```

1.7 6) Prototype loading, deduplication, and embedding preparation

```python
# Purpose: build prototype embeddings/resources once and print metadata for
↪   reproducibility auditing.
# Why this matters: model/device/prototype counts help explain classification
↪   behavior during review.
```

```python
resources = build_prototype_resources(cfg)
metadata = resources["metadata"]

print("Model metadata:")
print(f"- model_name: {metadata['model_name']}")
print(f"- device: {metadata['device']}")
print(f"- spacy_model: {metadata['spacy_model']}")
print(f"- hf_token_present: {metadata['hf_token_present']}")
print(f"- appendix_rows_used: {metadata['appendix_rows_used']}")
print(f"- total_prototypes: {metadata['prototype_total']}")
print("- prototype_count_by_group:")
print(metadata["prototype_count_by_group"])
```

```
Model metadata:
- model_name: NeuML/bioclinical-modernbert-base-embeddings
- device: cpu
- spacy_model: en_core_web_sm
- hf_token_present: True
- appendix_rows_used: 606
- total_prototypes: 693
- prototype_count_by_group:
{'RVC-INJ': 75, 'RVC-SYM-RESP': 43, 'RVC-SYM-CIRC': 18, 'RVC-SYM-NERV': 25, 'RVC-
SYM-DIG': 52, 'RVC-SYM-GU': 76, 'RVC-SYM-MSK': 22, 'RVC-SYM-SKIN': 29, 'RVC-
SYM-EYE': 50, 'RVC-SYM-GEN': 38, 'RVC-SYM-PSY': 40, 'RVC-DIS': 92, 'RVC-
TEST': 12, 'RVC-DIAG': 47, 'RVC-TREAT': 54, 'RVC-ADMIN': 14, 'RVC-UNCL': 6}
```

1.8 7) Segment cache build and classification

```python
# Purpose: build the segment cache and run visit-level RFV assignment across
↪   the full dataset.
# Why this matters: caching unique segments improves runtime while preserving
↪   deterministic outputs.

cache = build_segment_cache(
    df=df,
    seg_col="cc_cleaned",
    resources=resources,
    batch_size=cfg.segment_batch_size,
)

df_out = assign_rvc_per_segment_and_visit_cached(
    df=df,
```

```
        cache=cache,
        resources=resources,
        seg_col="cc_cleaned",
        cfg=cfg,
    )
df_out = apply_pseudomissing_uncodable_policy(
        df_out,
        max_rfv=cfg.max_rfv,
        pseudo_flag_column="cc_pseudomissing_flag",
    )
df_out = apply_blank_primary_uncodable_policy(
        df_out,
        max_rfv=cfg.max_rfv,
        missing_reason_column="cc_missing_reason",
    )

print(f"Unique segments cached: {len(cache.uniq_segments):,}")
print(f"Cache emb shape: {tuple(cache.emb.shape)}")
print(f"Cache sim shape: {tuple(cache.sim_proto.shape)}")


Unique segments cached: 905
Cache emb shape: (905, 768)
Cache sim shape: (905, 693)
```

1.9 8) Output shaping and compatibility mapping

Standardized output columns: - RFV1..RFV5 (legacy short-code compatibility) -
RFV1_name..RFV5_name - RFV1_support..RFV5_support - RFV1_sim..RFV5_sim -
segment_preds

```
# Purpose: validate the output schema contract and preview standardized RFV
↪  columns.
# Why this matters: downstream notebooks rely on these exact column names and
↪  structures.

validate_output_schema(
        df_out,
        max_rfv=cfg.max_rfv,
        require_primary_for_nonempty=cfg.allow_uncodable_primary,
    )

display_cols = ["cc_cleaned"]
```

```python
for slot in range(1, cfg.max_rfv + 1):
    display_cols.extend(
        [
            f"RFV{slot}",
            f"RFV{slot}_name",
            f"RFV{slot}_support",
            f"RFV{slot}_sim",
        ]
    )


print(df_out[display_cols].head(10).to_string())
```

```
Schema validation passed for 7,152 rows.
                             cc_cleaned      RFV1                RFV1_name             RFV
0                             [dyspnea]      resp     Symptom - Respiratory
1                  [shortness of breath]      resp     Symptom - Respiratory        shortness
2           [alter mental status, vomit]     neuro        Symptom - Nervous        alter ment
3                     [anemia, dyspnea]      resp     Symptom - Respiratory
stated)        anemia   0.647051                                       NaN
4                             [dyspnea]      resp     Symptom - Respiratory
5    [subarachnoid hemorrhage, transfer]   disease  Diseases (patient-stated)  subarachnoid he
6         [abdominal pain, swallow knife]        gi       Symptom - Digestive           abdomi
7                          [found down]  uncodable          Uncodable/Unknown               fc
8                 [abnormal lab, transfer]   abn_test        Abnormal test result           abnc
9                 [transfer, expose bowel]        gi       Symptom - Digestive            expc
```

1.10 9) Diagnostics and QA checks

```python
# Purpose: run QA smoke checks and summarize regression metrics after
↪  classification.
# Why this matters: immediate diagnostics make it easier to detect unintended
↪  behavioral drift.

qa_summary = validate_pipeline(df_out, cache=cache, resources=resources,
↪  seg_col="cc_cleaned")

override_summary = summarize_overrides(df_out, column="segment_preds")

print("\nRegression summary")
print("------------------")
print(f"Input rows: {len(df):,}")
print(f"Output rows: {len(df_out):,}")
```

```python
    print(f"Row count preserved: {len(df) == len(df_out)}")

    print("\nTop RFV1_name counts:")
    print(df_out["RFV1_name"].value_counts(dropna=True).head(15).to_string())

    print("\nOverride usage:")
    print(f"Segments total: {override_summary['segments_total']:,}")
    print(f"Segments overridden: {override_summary['segments_overridden']:,}")
    print(f"Segment override rate:
    ↪ {override_summary['segment_override_rate']:.4f}")
    print(f"Rows total: {override_summary['rows_total']:,}")
    print(f"Rows with >=1 override: {override_summary['rows_with_override']:,}")
    print(f"Row override rate: {override_summary['row_override_rate']:.4f}")

    phrase_audit_summary, phrase_audit_rfv = build_phrase_regression_audit(
        df_out,
        cc_column=cc_column,
    )
    print("\nAbnormal-labs phrase audit:")
    print(phrase_audit_summary.to_string(index=False))
    if not phrase_audit_rfv.empty:
        print(phrase_audit_rfv.head(10).to_string(index=False))

    abnormal_loss_rows = int(
        phrase_audit_summary.loc[
            phrase_audit_summary["metric"].eq("abnormal_loss_rows"), "value"
        ].iloc[0]
    )
    if abnormal_loss_rows > 0:
        raise AssertionError(
            "Detected abnormal-labs normalization regression: "
            f"{abnormal_loss_rows} rows still contain 'abnormal loss' segments."
        )

    normalization_integrity = _score_normalization_integrity(
        df_out["cc_text_for_nlp"].astype("string"),
        df_out["cc_cleaned_str"].astype("string"),
    )
    normalization_guardrail_audit = (
        pd.DataFrame(
            [
                {"metric": metric_name, "count": int(metric_value)}
                for metric_name, metric_value in normalization_integrity.items()
```

```python
        ]
    )
    .sort_values("metric")
    .reset_index(drop=True)
)
integrity_violation_total = int(sum(normalization_integrity.values()))
print("\nNormalization integrity checks:")
for metric_name, metric_value in normalization_integrity.items():
    print(f"{metric_name}: {metric_value:,}")
print(f"integrity_violation_total: {integrity_violation_total:,}")
if integrity_violation_total > 0:
    raise AssertionError(
        "Detected clinically invalid normalization substitutions: "
        f"{normalization_integrity}"
    )


if {"ed_intime", "cc_cleaned_str"}.issubset(df_out.columns):
    years = pd.to_datetime(df_out["ed_intime"], errors="coerce").dt.year
    covid_artifact_rows = int(
        (
            years.between(2011, 2019, inclusive="both")
            & df_out["cc_cleaned_str"].astype("string").str.contains(r"\b↵
            ↪  covid\b", case=False, na=False)
        ).sum()
    )
    print(f"covid_token_rows_2011_2019: {covid_artifact_rows:,}")
    if covid_artifact_rows > 0:
        raise AssertionError(
            "Found covid tokens in pre-2020 visits after
            ↪  normalization/classification."
        )

parsed_segment_preds = df_out["segment_preds"].fillna("[]").map(
    lambda value: json.loads(value) if isinstance(value, str) else []
)
blank_rfv1 = df_out["RFV1_name"].fillna("").astype(str).str.strip().eq("")
nonempty_segment_rows = parsed_segment_preds.map(lambda payload:
 ↪  isinstance(payload, list) and len(payload) > 0)
all_uncodable_rows = parsed_segment_preds.map(
    lambda payload: isinstance(payload, list)
    and len(payload) > 0
    and all(isinstance(item, dict) and item.get("code") == "RVC-UNCL" for
     ↪  item in payload)
```

```python
)

blank_rfv1_total = int(blank_rfv1.sum())
blank_rfv1_with_nonempty_segments = int((blank_rfv1 &
 ↪  nonempty_segment_rows).sum())
all_uncodable_total = int(all_uncodable_rows.sum())

print("\nPrimary-label completeness:")
print(f"blank_rfv1_total: {blank_rfv1_total:,}")
print(f"blank_rfv1_with_nonempty_segments:
 ↪  {blank_rfv1_with_nonempty_segments:,}")
print(f"all_uncodable_rows: {all_uncodable_total:,}")

uncodable_segment_counter: Counter = Counter()
for payload in parsed_segment_preds[all_uncodable_rows]:
    for item in payload:
        if isinstance(item, dict):
            uncodable_segment_counter[item.get("segment", "")] += 1

print("\nTop uncodable segments (post-fix):")
for segment, count in uncodable_segment_counter.most_common(15):
    print(f"{count:4d} | {segment}")

if cfg.allow_uncodable_primary and blank_rfv1_with_nonempty_segments != 0:
    raise AssertionError(
        "Expected zero blank RFV1_name values for rows with non-empty segment
         ↪  predictions, "
        f"found {blank_rfv1_with_nonempty_segments}."
    )

if cfg.run_optional_plots:
    top_table = top_cc_plot(df, top_n=30, normalize_case="upper")
    print(top_table.head(10).to_string(index=False))
else:
    print("\nOptional plots skipped (cfg.run_optional_plots=False).")


Prototype integrity:
{'RVC-INJ': 75, 'RVC-SYM-RESP': 43, 'RVC-SYM-CIRC': 18, 'RVC-SYM-NERV': 25, 'RVC-
SYM-DIG': 52, 'RVC-SYM-GU': 76, 'RVC-SYM-MSK': 22, 'RVC-SYM-SKIN': 29, 'RVC-
SYM-EYE': 50, 'RVC-SYM-GEN': 38, 'RVC-SYM-PSY': 40, 'RVC-DIS': 92, 'RVC-
TEST': 12, 'RVC-DIAG': 47, 'RVC-TREAT': 54, 'RVC-ADMIN': 14, 'RVC-UNCL': 6}
Total prototypes: 693
```

```
Segment column health:
rows=7,152 | with>=1 segment=7,150 | unique segments=905

Probe classifications:
shortness of breath -> [('shortness of breath', 'RVC-SYM-RESP', 0.8401217460632324, True)]
bradycardia -> [('bradycardia', 'RVC-SYM-CIRC', 0.6315358281135559, True)]
positive blood culture -> [('positive blood culture', 'RVC-TEST', 0.8695230484008789, True)]
alter mental status -> [('alter mental status', 'RVC-SYM-NERV', 0.7240737080574036, True)]
abdominal pain / shortness of breath -> [('abdominal pain', 'RVC-SYM-DIG', 0.830388069152832
SYM-RESP', 0.8401217460632324, True)]
right eye redness / pain -> [('right eye redness', 'RVC-SYM-EYE', 0.9090079069137573, True),
SYM-DIG', 0.7539336085319519, False)]

Regression summary
------------------
Input rows: 7,152
Output rows: 7,152
Row count preserved: True

Top RFV1_name counts:
RFV1_name
Symptom - Respiratory             2404
Injuries & adverse effects        1107
Symptom - Nervous                  830
Symptom - Digestive                708
Symptom - Circulatory              547
Diseases (patient-stated)          391
Symptom - General                  268
Uncodable/Unknown                  187
Symptom - Skin/Hair/Nails          134
Administrative                     131
Symptom - Musculoskeletal          117
Abnormal test result                97
Symptom - Genitourinary             70
Diagnostic/Screening/Preventive     58
Symptom - Psychological             37

Override usage:
Segments total: 10,651
Segments overridden: 7,829
Segment override rate: 0.7350
Rows total: 7,152
Rows with >=1 override: 5,675
```

```
Row override rate: 0.7935

Abnormal-labs phrase audit:
                  metric  value
       abnormal_lab_rows    140
      abnormal_loss_rows      0
abnormal_test_result_rows     60
            eye_ear_rows      0
               rfv1_name  row_count  row_rate
     Abnormal test result         60  0.428571
         Symptom - Nervous        19  0.135714
     Symptom - Respiratory        15  0.107143
       Symptom - Digestive        13  0.092857
         Symptom - General        11  0.078571
     Symptom - Circulatory         9  0.064286
Diseases (patient-stated)         6  0.042857
   Symptom - Genitourinary         3  0.021429
Symptom - Musculoskeletal         2  0.014286
Symptom - Skin/Hair/Nails         1  0.007143

Normalization integrity checks:
fall_to_small: 0
head_to_heart: 0
foot_to_post: 0
copd_or_cold_to_covid: 0
hyperglycemia_to_hypoglycemia: 0
hyperkalemia_to_hypokalemia: 0
found_down_changed: 0
found_to_wound: 0
trach_to_tract: 0
dental_to_mental: 0
black_to_block: 0
integrity_violation_total: 0
covid_token_rows_2011_2019: 0

Primary-label completeness:
blank_rfv1_total: 0
blank_rfv1_with_nonempty_segments: 0
all_uncodable_rows: 187

Top uncodable segments (post-fix):
  25 | found down
  15 | mcc
```

```
  14 | ms change
  11 | unknown - cc
   8 | alter ms
   7 | gtube eval
   6 | aortic dissection
   3 | mca
   3 | ftt
   2 | chb
   2 | delta ms
   2 | unable to ambulate
   2 | hang
   2 | o aaa
   2 |
```

Optional plots skipped (cfg.run_optional_plots=False).

```python
cc_missing_render = cc_missing_audit.copy()
if "examples" in cc_missing_render.columns:
    cc_missing_render["examples"] =
↪   cc_missing_render["examples"].astype(str).str.slice(0, 180)
print(
    render_latex_longtable(
        cc_missing_render,
        caption="Chief complaint missingness audit (pseudo-missing and
↪   true-missing policy).",
        label="tab:classifier_cc_missing_audit",
        index=False,
    )
)
```

```
\begin{longtable}{lrl}
\caption{Chief complaint missingness audit (pseudo-missing and true-missing policy).} \label-
\toprule
cc_missing_reason & row_count & examples \\
\midrule
\endfirsthead
\caption[]{Chief complaint missingness audit (pseudo-missing and true-missing policy).} \\
\toprule
cc_missing_reason & row_count & examples \\
\midrule
\endhead
\midrule
```

```latex
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
pseudo_missing_token & 2 & -; ___ \\
valid & 7150 & Dyspnea; Altered mental status; Respiratory distress; MVC; SOB \\
\end{longtable}
```

1.11 10) Export

```python
# Purpose: export the NLP-augmented workbook using canonical handoff names
↪   and a dated archive copy.
# Why this matters: downstream analysis reads a stable canonical path while
↪   preserving run-history snapshots.

canonical_output_path, archive_output_path =
↪   resolve_classifier_output_paths(WORK_DIR)
archive_output_path.parent.mkdir(parents=True, exist_ok=True)
write_archive_xlsx = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0").strip() ==
↪   "1"
classifier_export_drop_columns = [
    "cc_missing_flag",
    "cc_pseudomissing_flag",
]
df_export = df_out.drop(
    columns=[column for column in classifier_export_drop_columns if column
↪   in df_out.columns]
).copy()

def _sha256_file(path: Path) -> str:
    hasher = hashlib.sha256()
    with path.open("rb") as file_handle:
        for chunk in iter(lambda: file_handle.read(1024 * 1024), b""):
            hasher.update(chunk)
    return hasher.hexdigest()

df_export.to_excel(canonical_output_path, index=False,
↪   sheet_name=cfg.output_sheet_name)
if write_archive_xlsx:
    df_export.to_excel(archive_output_path, index=False,
↪   sheet_name=cfg.output_sheet_name)
```

```python
canonical_output_sha256 = _sha256_file(canonical_output_path)
archive_output_sha256 = _sha256_file(archive_output_path) if
↪   write_archive_xlsx else None
print(
    "EXPORT_SHAPE classifier",
    f"rows={len(df_export):,}",
    f"cols={df_export.shape[1]:,}",
    f"canonical_sha256={canonical_output_sha256}",
)


run_date = pd.Timestamp.now().strftime("%Y-%m-%d")
cc_missing_audit_path = archive_output_path.parent / f"{run_date}
↪   classifier_cc_missing_audit.csv"
cc_missing_audit.to_csv(cc_missing_audit_path, index=False)
phrase_audit_path = archive_output_path.parent / f"{run_date}
↪   classifier_phrase_audit.csv"
phrase_audit = phrase_audit_summary.copy()
phrase_audit["audit_type"] = "summary"
if not phrase_audit_rfv.empty:
    phrase_detail = phrase_audit_rfv.copy()
    phrase_detail["metric"] = "rfv_distribution"
    phrase_detail["value"] = phrase_detail["row_count"]
    phrase_detail["audit_type"] = "rfv_distribution"
    phrase_audit = pd.concat([phrase_audit, phrase_detail],
↪   ignore_index=True, sort=False)
phrase_audit.to_csv(phrase_audit_path, index=False)
spellfix_guardrail_audit_path = archive_output_path.parent / f"{run_date}
↪   classifier_spellfix_guardrail_audit.csv"
normalization_guardrail_audit.to_csv(spellfix_guardrail_audit_path,
↪   index=False)


spell_mode_comparison_path = archive_output_path.parent / f"{run_date}
↪   classifier_spell_mode_comparison.csv"
spell_mode_scores.to_csv(spell_mode_comparison_path, index=False)
spellfix_log_rows = (
    df_out.loc[
        df_out["cc_spellfix_any"].fillna(False).astype(bool),
        ["hadm_id", "subject_id", "cc_text_for_nlp", "cc_cleaned_str",
↪   "cc_spell_mode_selected", "cc_spellfix_log"],
    ]
    .copy()
)
```

```python
spellfix_log_path = archive_output_path.parent / f"{run_date}
↪  classifier_spellfix_log.csv"
spellfix_log_rows.to_csv(spellfix_log_path, index=False)

hypercap_flags_audit = pd.DataFrame(
    [
        {
            "row_count": int(len(df_out)),
            "abg_hypercap_threshold_positive": int(
                pd.to_numeric(df_out.get("abg_hypercap_threshold", 0),
↪  errors="coerce").fillna(0).astype(int).sum()
            ),
            "vbg_hypercap_threshold_positive": int(
                pd.to_numeric(df_out.get("vbg_hypercap_threshold", 0),
↪  errors="coerce").fillna(0).astype(int).sum()
            ),
            "unknown_hypercap_threshold_positive": int(
                pd.to_numeric(df_out.get("unknown_hypercap_threshold", 0),
↪  errors="coerce").fillna(0).astype(int).sum()
            ),
            "pco2_threshold_any_positive": int(
                pd.to_numeric(df_out.get("pco2_threshold_any", 0),
↪  errors="coerce").fillna(0).astype(int).sum()
            ),
            "threshold_union_mismatch_n": int(
                (
                    (
                        pd.to_numeric(df_out.get("abg_hypercap_threshold",
↪  0), errors="coerce").fillna(0).astype(int)
                        | pd.to_numeric(df_out.get("vbg_hypercap_threshold",
                          ↪  0), errors="coerce").fillna(0).astype(int)
                        | pd.to_numeric(df_out.get(⌋
                          ↪  "unknown_hypercap_threshold", 0),
                          ↪  errors="coerce").fillna(0).astype(int)
                    )
                    != pd.to_numeric(df_out.get("pco2_threshold_any", 0),
                       ↪  errors="coerce").fillna(0).astype(int)
                ).sum()
            ),
        }
    ]
)
```

```python
hypercap_flags_audit_path = archive_output_path.parent / f"{run_date}
↪   classifier_hypercap_flags_audit.csv"
hypercap_flags_audit.to_csv(hypercap_flags_audit_path, index=False)

classifier_contract = validate_classifier_contract(df_export,
↪   max_rfv=cfg.max_rfv)
classifier_contract["resource_report"] = resource_report
contract_report_path = archive_output_path.parent / f"{run_date}
↪   classifier_contract_report.json"
contract_report_path.write_text(json.dumps(classifier_contract, indent=2))
contract_artifact_payload = {
    "generated_utc": pd.Timestamp.utcnow().isoformat(),
    "status": classifier_contract["status"],
    "contracts": {"classifier": classifier_contract},
    "findings": classifier_contract["findings"],
}
contract_artifact_paths = write_contract_report(contract_artifact_payload,
↪   work_dir=WORK_DIR)

stage_manifest = collect_run_manifest(
    WORK_DIR,
    run_id=f"classifier_{pd.Timestamp.now().strftime('%Y%m%d_%H%M%S')}",
)
stage_manifest["stage"] = "classifier"
stage_manifest["model_metadata"] = {
    "model_name": metadata.get("model_name"),
    "spacy_model": metadata.get("spacy_model"),
    "device": metadata.get("device"),
    "appendix_rows_used": metadata.get("appendix_rows_used"),
    "prototype_total": metadata.get("prototype_total"),
}
stage_manifest["resource_verification"] = resource_report
stage_manifest["outputs"] = {
    "canonical_output_path": str(canonical_output_path),
    "canonical_output_sha256": canonical_output_sha256,
    "archive_output_path": str(archive_output_path) if write_archive_xlsx
        ↪   else None,
    "archive_output_sha256": archive_output_sha256,
    "cc_missing_audit_path": str(cc_missing_audit_path),
    "phrase_audit_path": str(phrase_audit_path),
    "spellfix_guardrail_audit_path": str(spellfix_guardrail_audit_path),
    "spell_mode_comparison_path": str(spell_mode_comparison_path),
    "spellfix_log_path": str(spellfix_log_path),
```

```python
        "hypercap_flags_audit_path": str(hypercap_flags_audit_path),
        "classifier_contract_report_path": str(contract_report_path),
    }
    stage_manifest_path = archive_output_path.parent / f"{run_date}
    ↪  classifier_run_manifest.json"
    stage_manifest_path.write_text(json.dumps(stage_manifest, indent=2))

    contract_mode = os.getenv("PIPELINE_CONTRACT_MODE", "fail").strip().lower()
    if contract_mode not in {"fail", "warn"}:
        raise ValueError("PIPELINE_CONTRACT_MODE must be one of {'fail',
        ↪  'warn'}")
    if classifier_contract["status"] == "fail" and contract_mode == "fail":
        raise ValueError(
            "Classifier contract failed. See report at "
            f"{contract_report_path} for details."
        )

    print(f"Wrote canonical output: {canonical_output_path}")
    if write_archive_xlsx:
        print(f"Wrote archive output: {archive_output_path}")
    else:
        print("Skipped archive workbook export (WRITE_ARCHIVE_XLSX_EXPORTS!=1).")
    print(f"Wrote CC missingness audit: {cc_missing_audit_path}")
    print(f"Wrote phrase regression audit: {phrase_audit_path}")
    print(f"Wrote spellfix guardrail audit: {spellfix_guardrail_audit_path}")
    print(f"Wrote spell-mode comparison audit: {spell_mode_comparison_path}")
    print(f"Wrote spellfix log audit: {spellfix_log_path}")
    print(f"Wrote hypercapnia flags audit: {hypercap_flags_audit_path}")
    print(f"Wrote classifier contract report: {contract_report_path}")
    print(f"Wrote classifier run manifest: {stage_manifest_path}")
    print(f"Wrote pipeline contract artifact:
    ↪  {contract_artifact_paths['contract_report_path']}")
    if contract_artifact_paths["failed_contract_path"].exists():
        print(f"Wrote failed contract artifact:
        ↪  {contract_artifact_paths['failed_contract_path']}")
    print(f"Classifier contract status: {classifier_contract['status']}")
    print(
        "Schema summary: "
        f"rows={len(df_export):,}, cols={df_export.shape[1]:,}, "
        f"rfv_slots={cfg.max_rfv}"
    )
    print(f"Export complete: {canonical_output_path}")
```

```
EXPORT_SHAPE classifier rows=7,152 cols=336 canonical_sha256=392fd7b77aef9397990051f334caf484
Wrote canonical output: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke
CC-NLP/MIMIC tabular data/MIMICIV all with CC_with_NLP.xlsx
Skipped archive workbook export (WRITE_ARCHIVE_XLSX_EXPORTS!=1).
Wrote CC missingness audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Lo
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_cc_missing_audit.csv
Wrote phrase regression audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_phrase_audit.csv
Wrote spellfix guardrail audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Wor
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_spellfix_guardrail_audit.csv
Wrote spell-mode comparison audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_spell_mode_comparison.csv
Wrote spellfix log audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Lock
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_spellfix_log.csv
Wrote hypercapnia flags audit: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_hypercap_flags_audit.csv
Wrote classifier contract report: /Users/blocke/Box Sync/Residency Personal Files/Scholarly W
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_contract_report.json
Wrote classifier run manifest: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 classifier_run_manifest.json
Wrote pipeline contract artifact: /Users/blocke/Box Sync/Residency Personal Files/Scholarly W
CC-NLP/debug/contracts/20260223_173743/contract_report.json
Classifier contract status: pass
Schema summary: rows=7,152, cols=336, rfv_slots=5
Export complete: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Researc
CC-NLP/MIMIC tabular data/MIMICIV all with CC_with_NLP.xlsx
```