MIMICIV_hypercap_EXT_cohort

Table of contents

1 MIMIC-IV Hypercapnia Cohort — ICD Physiologic Thresholds (BigQuery)

TODO:

[ ] want to add ED-rendered diagnoses - a flag to split off whether the hypercapnic respiratory failure ICD was rendered in the ED or during the hospital stay. [ ] are the installation instructions at the beginning of this notebook consistent with the way we intend the notebook to be used? is it also consistent with the README.md instructions?

Goal: Build an admissions-level tabular dataset that **enrolls** any hospital admission (`hadm_id`) meeting *__any__* of:

1. **ICD** codes for hypercapnic respiratory failure (legacy cohort).
2. **Any arterial blood gas** (LAB or POC) with **PaCO   45.0 mmHg** anywhere during the episode.
3. **Any venous blood gas** (LAB or POC) with **PaCO   50.0 mmHg** anywhere during the episode.

Then, keep all downstream columns/logic from the current workflow: - Per-code ICD indicators and an `any_hypercap_icd` flag. - Robust extraction of **first ABG** and **first VBG** (across LAB + POC) with standardized units (mmHg) and pairing logic. - Demographics/outcomes, NIH/OMB race & ethnicity, ED triage + first ED vitals, ICU meta (first stay + LOS), ventilation flags. - Sanity checks.

> **Assumptions** - You already configured BigQuery auth (`gcloud auth application-default login`) and `.env` variables as in the previous notebook. - The PhysioNet hosting project is `physionet-data`. - Datasets exist (e.g., `mimiciv_3_1_hosp`, `mimiciv_3_1_icu`, and an ED dataset such as `mimiciv_ed`). This notebook auto-detects the ED dataset.

Execution timing note: To capture per-cell runtimes, enable cell execution timing in VS Code (Notebook: Show Cell Execution Time) or enable ExecuteTime in Jupyter. Then re-run and save to allow runtime summaries.

1.1 Pipeline stages (readable map)

1. **Stage A — Config & helpers**: dataset IDs, thresholds, shared helper functions.
2. **Stage B — Cohort spine**: ICD + gas thresholds → `hadm_list` / `ed_stay_list`.
3. **Stage C — Bulk pulls**: one query per table where possible.
4. **Stage D — Transforms**: panels, flags, timing, comorbidities, vitals.
5. **Stage E — QA checks**: deterministic assertions and range checks.
6. **Stage F — Outputs**: parquet + Excel exports.

1.2 SQL registry (centralized query templates)

All BigQuery SQL is defined in one place below, then referenced by name in subsequent cells.

```
# Purpose: Build ABG/VBG hypercapnia threshold flags from lab and ICU POC
↪  pCO2 measurements.

import sys
print(sys.executable)

# Central SQL registry (define all query templates here)
SQL = {}

def sql(name: str) -> str:
    if name not in SQL:
        raise KeyError(f"SQL template not found: {name}")
    return SQL[name]

# SQL templates (populated below in-place to keep notebook linear)
# Names: admit_sql, co2_thresholds_sql, cohort_icd_sql, counts_sql, demo_sql,
↪  ditems_sql, ed_counts_sql, ed_first_vitals_sql, ed_spine_sql,
↪  ed_to_icu_sql, ed_triage_sql, ed_vitals_sql, icd_sql, icu_meta_sql,
↪  icu_sql, labitems_sql, labs_sql, vent_chart_sql, vent_sql
```

/Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/Hyperc
CC-NLP/.venv/bin/python

Rationale: Define a reproducible, admission-level cohort that captures hypercapnia using complementary diagnostic (ICD) and physiologic (blood gas) criteria.

2 MIMIC-IV on BigQuery

2.1 Environment Bootstrap & Smoke Test

Purpose: make a clean, reproducible start on a new machine.

Outcome: verify auth, project config, and dataset access; provide a reusable BigQuery runner for the build notebook.

Rationale: Establish the BigQuery environment and dataset configuration so queries are consistent and reproducible across runs.

2.2 0. Prerequisites (one-time)

Accounts & access - PhysioNet access to MIMIC-IV on BigQuery; in BigQuery Console star project `physionet-data`. - A Google Cloud **Project ID** with BigQuery API enabled (this is your **billing** project).

CLI & environment - Google Cloud SDK (gcloud) installed and on PATH. - Python environment created with `uv` (see README) and Jupyter kernel selected. - A project-local `.env` with the variables below.

.env variables

```
MIMIC_BACKEND=bigquery
WORK_PROJECT=<your-billing-project-id>
BQ_PHYSIONET_PROJECT=physionet-data
BQ_DATASET_HOSP=mimiciv_3_1_hosp
BQ_DATASET_ICU=mimiciv_3_1_icu
BQ_DATASET_ED=mimiciv_ed
WORK_DIR=/path/to/Hypercap-CC-NLP
# GOOGLE_APPLICATION_CREDENTIALS=/path/to/service-account.json
```

Command line quickstart

```
brew install --cask google-cloud-sdk
gcloud init
gcloud auth application-default login
gcloud services enable bigquery.googleapis.com --project
 ↪  <your-billing-project-id>
ls -l ~/.config/gcloud/application_default_credentials.json
```

Rationale: Verify access and credentials up front to prevent silent failures later in the pipeline.

```python
# Purpose: Set up project paths, environment variables, and BigQuery client
↪   connections for reproducible execution.

# --- Imports & environment
import os, re, json, math, textwrap, sys
from pathlib import Path
from typing import Any, Mapping, Sequence

import numpy as np
import pandas as pd

from google.cloud import bigquery
from google.oauth2 import service_account
from dotenv import load_dotenv

load_dotenv()

WORK_DIR = Path(os.getenv("WORK_DIR", Path.cwd())).expanduser().resolve()
DATA_DIR = WORK_DIR / "MIMIC tabular data"
DATA_DIR.mkdir(parents=True, exist_ok=True)
SRC_DIR = WORK_DIR / "src"
if SRC_DIR.exists() and str(SRC_DIR) not in sys.path:
    sys.path.insert(0, str(SRC_DIR))

# Local helper functions (core cohort logic must remain embedded in this
↪   notebook)
OMR_RESULT_NAMES = ("bmi", "height", "weight")
OMR_OUTPUT_COLUMNS = (
    "bmi_closest_pre_ed",
    "height_closest_pre_ed",
    "weight_closest_pre_ed",
)

OMR_PROVENANCE_COLUMNS = (
    "anthro_timing_tier",
    "anthro_days_offset",
    "anthro_chartdate",
    "anthro_timing_uncertain",
    "anthro_source",
    "anthro_obstime",
    "anthro_hours_offset",
    "anthro_timing_basis",
)
```

```python
OMR_TIMING_TIERS = ("pre_ed_365", "post_ed_365", "missing")

EXPECTED_STRUCTURAL_NULL_FIELDS = (
    "poc_abg_ph_uom",
    "poc_vbg_ph_uom",
    "poc_other_ph_uom",
)

PACO2_VALUE_UOM_PAIRS = (
    ("poc_abg_paco2", "poc_abg_paco2_uom"),
    ("poc_vbg_paco2", "poc_vbg_paco2_uom"),
    ("poc_other_paco2", "poc_other_paco2_uom"),
)

DEFAULT_VITALS_MODEL_RANGES: dict[str, tuple[float, float]] = {
    "ed_triage_hr": (20.0, 250.0),
    "ed_first_hr": (20.0, 250.0),
    "ed_triage_rr": (4.0, 80.0),
    "ed_first_rr": (4.0, 80.0),
}

DEFAULT_GAS_MODEL_RANGES: dict[str, tuple[float, float]] = {
    "first_ph": (6.8, 7.8),
    "first_pco2": (10.0, 200.0),
    "first_other_pco2": (10.0, 200.0),
    "first_lactate": (0.0, 30.0),
}

_NUMERIC_TOKEN_PATTERN = re.compile(r"(-?\d+(?:\.\d+)?)")
_ARTERIAL_HINT_PATTERN = re.compile(
    r"(arterial|abg|a[- ]?line|art line|\bart\b|\bartery\b)",
    re.I,
)
_VENOUS_HINT_PATTERN = re.compile(
    r"(venous|vbg|cvbg|mixed venous|central venous|\bven\b)",
    re.I,
)


def _to_int64(series: pd.Series) -> pd.Series:
    """Return pandas nullable integer series for key columns."""
    return pd.to_numeric(series, errors="coerce").astype("Int64")
```

```python
def prepare_omr_records(omr_df: pd.DataFrame) -> pd.DataFrame:
    """Normalize OMR records into a deterministic schema.

    Required source columns: ``subject_id``, ``chartdate``, ``result_name``,
    ``result_value``.

    Returns:
        DataFrame with columns:
        ``subject_id`` (Int64), ``chartdate_dt`` (datetime64),
        ``result_name`` (lowercase), and ``result_value_num`` (float).
    """
    required = {"subject_id", "chartdate", "result_name", "result_value"}
    missing = sorted(required.difference(omr_df.columns))
    if missing:
        raise KeyError(f"prepare_omr_records missing required columns:
        ↪  {missing}")

    prepared = omr_df.copy()
    prepared["subject_id"] = _to_int64(prepared["subject_id"])
    prepared["chartdate_dt"] = pd.to_datetime(prepared["chartdate"],
↪  errors="coerce")
    prepared["result_name"] = (
        prepared["result_name"].astype(str).str.strip().str.lower()
    )
    prepared["result_value_num"] = pd.to_numeric(
        prepared["result_value"]
        .astype(str)
        .str.extract(_NUMERIC_TOKEN_PATTERN, expand=False),
        errors="coerce",
    )

    prepared =
↪  prepared.loc[prepared["result_name"].isin(OMR_RESULT_NAMES)].copy()
    prepared = prepared.loc[
        prepared["subject_id"].notna() & prepared["chartdate_dt"].notna()
    ].copy()

    return prepared[["subject_id", "chartdate_dt", "result_name",
    ↪  "result_value_num"]]
```

```python
def attach_closest_pre_ed_omr(
    ed_df: pd.DataFrame,
    omr_df: pd.DataFrame,
    window_days: int = 365,
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Attach closest anthropometrics with tiered timing fallback.

    Args:
        ed_df: ED-stay grain dataframe with ``ed_stay_id``, ``subject_id``,
            and ``ed_intime``.
        omr_df: Prepared OMR records from ``prepare_omr_records``.
        window_days: Inclusion window in days before/after ED arrival.
    """

    def _with_default_anthro_columns(frame: pd.DataFrame) -> pd.DataFrame:
        updated = frame.copy()
        for column_name in OMR_OUTPUT_COLUMNS:
            if column_name not in updated.columns:
                updated[column_name] = pd.NA
        if "anthro_timing_tier" not in updated.columns:
            updated["anthro_timing_tier"] = "missing"
        else:
            updated["anthro_timing_tier"] =
  ↪ updated["anthro_timing_tier"].fillna("missing")
        if "anthro_days_offset" not in updated.columns:
            updated["anthro_days_offset"] = pd.Series([pd.NA] * len(updated),
  ↪ dtype="Int64")
        else:
            updated["anthro_days_offset"] = pd.to_numeric(
                updated["anthro_days_offset"], errors="coerce"
            ).astype("Int64")
        if "anthro_chartdate" not in updated.columns:
            updated["anthro_chartdate"] = pd.NaT
        else:
            updated["anthro_chartdate"] = pd.to_datetime(
                updated["anthro_chartdate"], errors="coerce"
            )
        if "anthro_timing_uncertain" not in updated.columns:
            updated["anthro_timing_uncertain"] = pd.Series(
                [pd.NA] * len(updated), dtype="boolean"
            )
        else:
            updated["anthro_timing_uncertain"] = updated[
```

```python
            "anthro_timing_uncertain"
        ].astype("boolean")
    if "anthro_source" not in updated.columns:
        updated["anthro_source"] = "missing"
    else:
        updated["anthro_source"] = (
            updated["anthro_source"].astype(str).replace({"":
↪  "missing"}).fillna("missing")
        )
    if "anthro_obstime" not in updated.columns:
        updated["anthro_obstime"] = pd.NaT
    else:
        updated["anthro_obstime"] =
↪  pd.to_datetime(updated["anthro_obstime"], errors="coerce")
    if "anthro_hours_offset" not in updated.columns:
        updated["anthro_hours_offset"] = pd.NA
    updated["anthro_hours_offset"] = pd.to_numeric(
        updated["anthro_hours_offset"], errors="coerce"
    )
    if "anthro_timing_basis" not in updated.columns:
        updated["anthro_timing_basis"] = "missing"
    else:
        updated["anthro_timing_basis"] = (
            updated["anthro_timing_basis"].astype(str).replace({"":
↪  "missing"}).fillna("missing")
        )

    missing_mask = updated["anthro_timing_tier"].eq("missing")
    pre_mask = updated["anthro_timing_tier"].eq("pre_ed_365")
    post_mask = updated["anthro_timing_tier"].eq("post_ed_365")
    updated.loc[pre_mask, "anthro_timing_uncertain"] = False
    updated.loc[post_mask, "anthro_timing_uncertain"] = True
    updated.loc[missing_mask, "anthro_timing_uncertain"] = pd.NA
    updated.loc[pre_mask | post_mask, "anthro_source"] = "omr"
    updated.loc[missing_mask, "anthro_source"] = "missing"
    updated.loc[pre_mask, "anthro_timing_basis"] = "pre"
    updated.loc[post_mask, "anthro_timing_basis"] = "post"
    updated.loc[missing_mask, "anthro_timing_basis"] = "missing"
    updated.loc[missing_mask, "anthro_obstime"] = pd.NaT
    updated.loc[missing_mask, "anthro_hours_offset"] = pd.NA

    return updated
```

```python
def _tier_counts(frame: pd.DataFrame) -> dict[str, int]:
    tier_series =
↪  frame["anthro_timing_tier"].astype(str).fillna("missing")
    return {
        "pre_ed_365": int((tier_series == "pre_ed_365").sum()),
        "post_ed_365": int((tier_series == "post_ed_365").sum()),
        "missing": int((tier_series == "missing").sum()),
    }

required_ed = {"ed_stay_id", "subject_id", "ed_intime"}
missing_ed = sorted(required_ed.difference(ed_df.columns))
if missing_ed:
    raise KeyError(f"attach_closest_pre_ed_omr missing ED columns:
    ↪  {missing_ed}")

required_omr = {"subject_id", "chartdate_dt", "result_name",
↪  "result_value_num"}
missing_omr = sorted(required_omr.difference(omr_df.columns))
if missing_omr:
    raise KeyError(f"attach_closest_pre_ed_omr missing OMR columns:
    ↪  {missing_omr}")

diagnostics: dict[str, Any] = {
    "window_days": int(window_days),
    "source_rows_prepared": int(len(omr_df)),
    "parsed_value_rows": int(omr_df["result_value_num"].notna().sum()),
    "ed_rows": int(len(ed_df)),
}

empty_window_diagnostics = {
    "nonnegative_candidate_rows": 0,
    "pre_window_candidate_rows": 0,
    "post_window_candidate_rows": 0,
    "closest_absolute_candidate_rows": 0,
    "within_window_candidate_rows": 0,
}

ed_norm = ed_df[["ed_stay_id", "subject_id", "ed_intime"]].copy()
ed_norm["subject_id"] = _to_int64(ed_norm["subject_id"])
ed_norm["ed_intime_dt"] = pd.to_datetime(ed_norm["ed_intime"],
↪  errors="coerce")
ed_norm["ed_date_dt"] = ed_norm["ed_intime_dt"].dt.floor("D")
```

```python
        ed_norm = ed_norm.loc[ed_norm["subject_id"].notna() &
↪   ed_norm["ed_date_dt"].notna()]

    if omr_df.empty or ed_norm.empty:
        updated = _with_default_anthro_columns(ed_df)
        tier_counts = _tier_counts(updated)
        total_rows = max(int(len(updated)), 1)
        diagnostics.update(
            {
                "ed_rows_eligible_for_join": int(len(ed_norm)),
                "subject_overlap_count": 0,
                "candidate_rows_after_subject_join": 0,
                **empty_window_diagnostics,
                "within_window_candidate_rows": 0,
                "eligible_ed_stays_with_candidates": 0,
                "attached_non_null_counts": {
                    column_name: int(updated[column_name].notna().sum())
                    for column_name in OMR_OUTPUT_COLUMNS
                },
                "attached_any_non_null_rows": int(

↪   updated[list(OMR_OUTPUT_COLUMNS)].notna().any(axis=1).sum()
                ),
                "selected_tier_counts": tier_counts,
                "selected_tier_rates": {
                    key: float(value / total_rows) for key, value in
↪   tier_counts.items()
                },
                "timing_uncertain_count": int(
                    updated["anthro_timing_uncertain"].fillna(False).sum()
                ),
            }
        )
        return updated, diagnostics

    omr_pivot = (
        omr_df.pivot_table(
            index=["subject_id", "chartdate_dt"],
            columns="result_name",
            values="result_value_num",
            aggfunc="first",
        )
        .reset_index()
```

```python
        .copy()
)

shared_subjects =
↪  set(ed_norm["subject_id"].dropna().astype(int)).intersection(
        set(omr_pivot["subject_id"].dropna().astype(int))
)
diagnostics["ed_rows_eligible_for_join"] = int(len(ed_norm))
diagnostics["subject_overlap_count"] = int(len(shared_subjects))

merged = ed_norm.merge(omr_pivot, on="subject_id", how="left")
diagnostics["candidate_rows_after_subject_join"] = int(
        merged["chartdate_dt"].notna().sum()
)

merged["days_before"] = (
        merged["ed_date_dt"] - pd.to_datetime(merged["chartdate_dt"],
↪  errors="coerce")
).dt.days
valid_days = merged["days_before"].dropna()
diagnostics["days_before_min"] = (
        int(valid_days.min()) if not valid_days.empty else None
)
diagnostics["days_before_max"] = (
        int(valid_days.max()) if not valid_days.empty else None
)
pre_window_mask = (
        merged["days_before"].notna()
        & merged["days_before"].ge(0)
        & merged["days_before"].le(window_days)
)
post_window_mask = (
        merged["days_before"].notna()
        & merged["days_before"].lt(0)
        & merged["days_before"].ge(-window_days)
)
abs_window_mask = (
        merged["days_before"].notna()
        & merged["days_before"].abs().le(window_days)
)

diagnostics["nonnegative_candidate_rows"] = int((merged["days_before"] >=
↪  0).sum())
```

```python
    diagnostics["pre_window_candidate_rows"] = int(pre_window_mask.sum())
    diagnostics["post_window_candidate_rows"] = int(post_window_mask.sum())
    diagnostics["closest_absolute_candidate_rows"] =
↪ int(abs_window_mask.sum())

    pre_candidates = merged.loc[pre_window_mask].copy()
    post_candidates = merged.loc[post_window_mask].copy()
    diagnostics["within_window_candidate_rows"] = int(len(pre_candidates))

    selected_parts: list[pd.DataFrame] = []
    selected_stays: set[Any] = set()

    if not pre_candidates.empty:
        pre_selected = (
            pre_candidates.sort_values(
                ["ed_stay_id", "days_before", "chartdate_dt"],
                ascending=[True, True, False],
            )
            .groupby("ed_stay_id", as_index=False)
            .first()
        )
        pre_selected["anthro_timing_tier"] = "pre_ed_365"
        selected_parts.append(pre_selected)
        selected_stays.update(pre_selected["ed_stay_id"].tolist())

    if not post_candidates.empty:
        post_candidates["days_after_abs"] =
↪ post_candidates["days_before"].abs()
        post_selected = (
            post_candidates.sort_values(
                ["ed_stay_id", "days_after_abs", "chartdate_dt"],
                ascending=[True, True, True],
            )
            .groupby("ed_stay_id", as_index=False)
            .first()
        )
        post_selected = post_selected.loc[
            ~post_selected["ed_stay_id"].isin(selected_stays)
        ].copy()
        if not post_selected.empty:
            post_selected["anthro_timing_tier"] = "post_ed_365"
            selected_parts.append(post_selected)
```

13

```python
if selected_parts:
    selected = pd.concat(selected_parts, ignore_index=True)
    selected = selected.rename(
        columns={
            "bmi": "bmi_closest_pre_ed",
            "height": "height_closest_pre_ed",
            "weight": "weight_closest_pre_ed",
            "chartdate_dt": "anthro_chartdate",
            "days_before": "anthro_days_offset",
        }
    )
    selected["anthro_chartdate"] = pd.to_datetime(
        selected["anthro_chartdate"], errors="coerce"
    )
    selected["anthro_days_offset"] = pd.to_numeric(
        selected["anthro_days_offset"], errors="coerce"
    ).astype("Int64")
    selected["anthro_timing_uncertain"] =
selected["anthro_timing_tier"].eq(
        "post_ed_365"
    )
    selected["anthro_source"] = "omr"
    selected["anthro_obstime"] = pd.to_datetime(
        selected["anthro_chartdate"], errors="coerce"
    )
    selected["anthro_hours_offset"] = pd.to_numeric(
        selected["anthro_days_offset"], errors="coerce"
    ) * 24.0
    selected["anthro_timing_basis"] = selected["anthro_timing_tier"].map(
        {"pre_ed_365": "pre", "post_ed_365": "post"}
    ).fillna("missing")

    updates = selected[
        [
            "ed_stay_id",
            *OMR_OUTPUT_COLUMNS,
            *OMR_PROVENANCE_COLUMNS,
        ]
    ].copy()
    updated = ed_df.merge(updates, on="ed_stay_id", how="left")
else:
    updated = ed_df.copy()
```

```python
    updated = _with_default_anthro_columns(updated)

    diagnostics["eligible_ed_stays_with_candidates"] = int(
        updated["anthro_timing_tier"].isin({"pre_ed_365",
↪   "post_ed_365"}).sum()
    )
    diagnostics["attached_non_null_counts"] = {
        column_name: int(updated[column_name].notna().sum())
        for column_name in OMR_OUTPUT_COLUMNS
    }
    diagnostics["attached_any_non_null_rows"] = int(
        updated[list(OMR_OUTPUT_COLUMNS)].notna().any(axis=1).sum()
    )
    tier_counts = _tier_counts(updated)
    diagnostics["selected_tier_counts"] = tier_counts
    total_rows = max(int(len(updated)), 1)
    diagnostics["selected_tier_rates"] = {
        key: float(value / total_rows) for key, value in tier_counts.items()
    }
    diagnostics["timing_uncertain_count"] = int(
        updated["anthro_timing_uncertain"].fillna(False).sum()
    )
    diagnostics["anthro_source_counts"] = {
        str(key): int(value)
        for key, value in updated["anthro_source"].fillna("missing")
        .astype(str)
        .value_counts(dropna=False)
        .items()
    }

    return updated, diagnostics


def _infer_route_hint_text(value: object) -> str | None:
    text = "" if pd.isna(value) else str(value).strip().lower()
    if not text or text in {"nan", "none"}:
        return None
    if _ARTERIAL_HINT_PATTERN.search(text):
        return "arterial"
    if _VENOUS_HINT_PATTERN.search(text):
        return "venous"
    return None
```

```python
def _resolve_route_hints(values: Sequence[str]) -> tuple[str | None, bool,
↪   int]:
    hints = [str(value).strip().lower() for value in values if
↪   str(value).strip()]
    arterial_n = sum(1 for value in hints if value == "arterial")
    venous_n = sum(1 for value in hints if value == "venous")
    conflict = arterial_n > 0 and venous_n > 0
    if conflict:
        return None, True, int(len(hints))
    if arterial_n > 0:
        return "arterial", False, int(len(hints))
    if venous_n > 0:
        return "venous", False, int(len(hints))
    return None, False, int(len(hints))


def infer_panel_gas_source_metadata(
    panel_df: pd.DataFrame,
    labs_df: pd.DataFrame,
    labitems_df: pd.DataFrame,
    *,
    specimen_source_itemids: Sequence[int] | None = None,
    pco2_itemids: Sequence[int] | None = None,
    mode: str = "metadata_only",
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Infer gas panel source from metadata/text hints with tier diagnostics.

    Tier precedence:
      1. Specimen/source row text (``value_text`` from specimen/source
↪   itemids),
      2. ``d_labitems`` label/fluid hints,
      3. Panel co-occurrence hints (pCO2-labeled rows + free-text),
      4. Fallback ``other``.
    """
    if mode != "metadata_only":
        raise ValueError(
            "infer_panel_gas_source_metadata supports only
            ↪   mode='metadata_only'."
        )

    required_panel = {"ed_stay_id", "specimen_id"}
    missing_panel = sorted(required_panel.difference(panel_df.columns))
```

16

```python
    if missing_panel:
        raise KeyError(
            "infer_panel_gas_source_metadata missing panel columns: "
            f"{missing_panel}"
        )
    required_labs = {"ed_stay_id", "specimen_id", "itemid"}
    missing_labs = sorted(required_labs.difference(labs_df.columns))
    if missing_labs:
        raise KeyError(
            "infer_panel_gas_source_metadata missing labs columns: "
            f"{missing_labs}"
        )
    required_labitems = {"itemid"}
    missing_labitems =
↪   sorted(required_labitems.difference(labitems_df.columns))
    if missing_labitems:
        raise KeyError(
            "infer_panel_gas_source_metadata missing labitems columns: "
            f"{missing_labitems}"
        )

    key_cols = ["ed_stay_id", "specimen_id"]
    panel = panel_df.copy()
    panel["source"] = pd.Series(pd.NA, index=panel.index, dtype="string")
    panel["source_inference_tier"] = "fallback_other"
    panel["source_hint_conflict"] = pd.Series(False, index=panel.index,
↪   dtype="boolean")
    panel["source_hint_count"] = pd.Series(pd.NA, index=panel.index,
↪   dtype="Int64")

    labs = labs_df.copy()
    labs["itemid"] = pd.to_numeric(labs["itemid"],
↪   errors="coerce").astype("Int64")
    labs["specimen_id"] = pd.to_numeric(labs["specimen_id"],
↪   errors="coerce").astype(
        "Int64"
    )
    labs = labs.loc[labs["itemid"].notna() &
↪   labs["specimen_id"].notna()].copy()
    if labs.empty:
        panel["source"] = "other"
        diagnostics = summarize_gas_source(panel)
        diagnostics.update(
```

```python
            {
                "mode": mode,
                "resolved_rows": 0,
                "resolved_rate": 0.0,
                "unresolved_specimen_id_count": 0,
                "unresolved_specimen_id_examples": [],
                "unresolved_value_text_top": {},
                "unresolved_label_top": {},
            }
        )
        return panel, diagnostics

    if "value_text" in labs.columns:
        labs["value_text_norm"] = (
            labs["value_text"]
            .astype("string")
            .fillna("")
            .str.strip()
            .str.lower()
        )
    else:
        labs["value_text_norm"] = ""

    labitems = labitems_df.copy()
    labitems["itemid"] = pd.to_numeric(labitems["itemid"],
↪   errors="coerce").astype(
        "Int64"
    )
    for text_col in ("label", "fluid"):
        if text_col not in labitems.columns:
            labitems[text_col] = ""
    labitems["label_norm"] =
↪   labitems["label"].astype("string").fillna("").str.strip()
    labitems["fluid_norm"] =
↪   labitems["fluid"].astype("string").fillna("").str.strip()
    labitems["item_route_hint"] = (
        (labitems["label_norm"] + " " + labitems["fluid_norm"])
        .str.strip()
        .map(_infer_route_hint_text)
    )
    item_meta = (
        labitems.dropna(subset=["itemid"])
        .drop_duplicates(subset=["itemid"])
```

```python
        .set_index("itemid")[["label_norm", "fluid_norm", "item_route_hint"]]
    )
    labs = labs.merge(item_meta, left_on="itemid", right_index=True,
↪  how="left")
    labs["item_route_hint"] = labs["item_route_hint"].astype("string")
    labs["text_route_hint"] =
↪  labs["value_text_norm"].map(_infer_route_hint_text).astype(
        "string"
    )
    labs["label_fluid_text"] = (
        labs["label_norm"].fillna("").astype(str).str.strip()
        + " | "
        + labs["fluid_norm"].fillna("").astype(str).str.strip()
    ).str.strip(" |")

    def _build_tier_mapping(
        frame: pd.DataFrame,
        *,
        hint_column: str,
        tier_name: str,
    ) -> pd.DataFrame:
        subset = frame[key_cols + [hint_column]].copy()
        subset = subset.loc[subset[hint_column].notna()].copy()
        if subset.empty:
            return pd.DataFrame(
                columns=[
                    *key_cols,
                    "tier_source",
                    "tier_conflict",
                    "tier_hint_count",
                    "tier_name",
                ]
            )

        grouped = (
            subset.groupby(key_cols, dropna=False)[hint_column]
            .agg(list)
            .reset_index(name="hints")
        )
        rows: list[dict[str, Any]] = []
        for _, grouped_row in grouped.iterrows():
            source, conflict, hint_count =
↪  _resolve_route_hints(grouped_row["hints"])
```

```python
        rows.append(
            {
                "ed_stay_id": grouped_row["ed_stay_id"],
                "specimen_id": grouped_row["specimen_id"],
                "tier_source": source,
                "tier_conflict": bool(conflict),
                "tier_hint_count": int(hint_count),
                "tier_name": tier_name,
            }
        )
    resolved = pd.DataFrame(rows)
    return resolved.loc[resolved["tier_source"
     ↪  ].notna()].reset_index(drop=True)

specimen_source_itemid_set = {
    int(item) for item in (specimen_source_itemids or []) if
     ↪  pd.notna(item)
}
pco2_itemid_set = {int(item) for item in (pco2_itemids or []) if
 ↪  pd.notna(item)}

tier1 = _build_tier_mapping(
    labs.loc[labs["itemid"].astype("Int64"
 ↪  ).isin(specimen_source_itemid_set)],
    hint_column="text_route_hint",
    tier_name="specimen_text",
)
tier2 = _build_tier_mapping(
    labs,
    hint_column="item_route_hint",
    tier_name="label_fluid",
)
tier_cluster = pd.DataFrame(
    columns=[
        *key_cols,
        "tier_source",
        "tier_conflict",
        "tier_hint_count",
        "tier_name",
    ]
)
if "panel_time" in panel_df.columns:
    panel_times = panel_df[key_cols + ["panel_time"]].copy()
```

```
        panel_times["panel_time"] = pd.to_datetime(panel_times["panel_time"],
↪  errors="coerce")

        tier_seed = pd.concat([tier1, tier2], ignore_index=True)
        if not tier_seed.empty:
            tier_seed = (
                tier_seed[key_cols + ["tier_source"]]
                .dropna(subset=["tier_source"])
                .drop_duplicates(subset=key_cols)
            )
            panel_seed = panel_times.merge(tier_seed, on=key_cols,
↪  how="left")
            known = panel_seed.loc[panel_seed["tier_source"].notna()].copy()
            unknown = panel_seed.loc[panel_seed["tier_source"].isna()].copy()
            if not known.empty and not unknown.empty:
                pairs = unknown.merge(
                    known[["ed_stay_id", "specimen_id", "panel_time",
↪  "tier_source"]],
                    on="ed_stay_id",
                    how="left",
                    suffixes=("_unknown", "_known"),
                )
                pairs["delta_minutes"] = (
                    (pairs["panel_time_known"] -
↪  pairs["panel_time_unknown"]).abs().dt.total_seconds()
                    / 60.0
                )
                pairs = pairs.loc[pairs["delta_minutes"].notna() &
↪  pairs["delta_minutes"].le(60.0)].copy()
                if not pairs.empty:
                    keys = ["ed_stay_id", "specimen_id_unknown"]
                    min_delta =
↪  pairs.groupby(keys)["delta_minutes"].transform("min")
                    nearest =
↪  pairs.loc[pairs["delta_minutes"].eq(min_delta)].copy()
                    conflict_stats = (
                        nearest.groupby(keys)["tier_source_known"]
                        .nunique(dropna=True)
                        .reset_index(name="source_nunique")
                    )
                    nearest = nearest.merge(conflict_stats, on=keys,
↪  how="left")
```

```python
                    nearest =
↪   nearest.loc[nearest["source_nunique"].eq(1)].copy()
                    if not nearest.empty:
                        nearest = nearest.sort_values(keys +
↪   ["delta_minutes"])
                        nearest = nearest.groupby(keys,
↪   as_index=False).first()
                        tier_cluster = nearest.rename(
                            columns={
                                "specimen_id_unknown": "specimen_id",
                                "tier_source_known": "tier_source",
                            }
                        )[
                            ["ed_stay_id", "specimen_id", "tier_source"]
                        ].copy()
                        tier_cluster["tier_conflict"] = False
                        tier_cluster["tier_hint_count"] = 1
                        tier_cluster["tier_name"] = "cluster_inheritance"

    tier3_candidates = labs.copy()
    if pco2_itemid_set:
        tier3_candidates = tier3_candidates.loc[
            tier3_candidates["itemid"].astype("Int64").isin(pco2_itemid_set)
            | tier3_candidates["text_route_hint"].notna()
        ].copy()
    tier3_candidates["panel_route_hint"] =
↪   tier3_candidates["item_route_hint"].fillna(
        tier3_candidates["text_route_hint"]
    )
    tier3 = _build_tier_mapping(
        tier3_candidates,
        hint_column="panel_route_hint",
        tier_name="panel_cooccurrence",
    )

    for tier_frame in (tier1, tier2, tier_cluster, tier3):
        if tier_frame.empty:
            continue
        panel = panel.merge(
            tier_frame[
                key_cols
                + ["tier_source", "tier_conflict", "tier_hint_count",
                ↪   "tier_name"]
```

22

```python
    ],
    on=key_cols,
    how="left",
)
assign_mask = panel["source"].isna() & panel["tier_source"].notna()
panel.loc[assign_mask, "source"] = panel.loc[assign_mask,
↳ "tier_source"]
panel.loc[assign_mask, "source_inference_tier"] = panel.loc[
    assign_mask, "tier_name"
]
selected_conflict = panel.loc[assign_mask,
↳ "tier_conflict"].astype("boolean")
panel.loc[assign_mask, "source_hint_conflict"] =
↳ selected_conflict.fillna(False)
panel.loc[assign_mask, "source_hint_count"] = pd.to_numeric(
    panel.loc[assign_mask, "tier_hint_count"], errors="coerce"
).astype("Int64")
panel = panel.drop(
    columns=["tier_source", "tier_conflict", "tier_hint_count",
↳ "tier_name"]
)

panel["source"] = panel["source"].fillna("other").astype("string")
panel["source"] = panel["source"].replace({"unknown": "other"})
panel.loc[panel["source"].eq("other"), "source_inference_tier"] =
↳ panel.loc[
    panel["source"].eq("other"), "source_inference_tier"
].replace({"": "fallback_other"}).fillna("fallback_other")
panel["source_inference_tier"] =
↳ panel["source_inference_tier"].astype("string")
panel["source_hint_conflict"] = (
    panel["source_hint_conflict"].fillna(False).astype("boolean")
)

diagnostics = summarize_gas_source(panel)
total = max(int(len(panel)), 1)
diagnostics["mode"] = mode
diagnostics["resolved_rows"] = int(
    panel["source"].isin({"arterial", "venous"}).sum()
)
diagnostics["resolved_rate"] = float(diagnostics["resolved_rows"] /
↳ total)
```

```python
    unresolved = panel.loc[panel["source"].eq("other"),
↪   key_cols].drop_duplicates()
    diagnostics["unresolved_specimen_id_count"] =
↪   int(unresolved["specimen_id"].nunique())
    diagnostics["unresolved_specimen_id_examples"] = (
        unresolved["specimen_id"].dropna().astype(int).head(20).tolist()
    )
    if unresolved.empty:
        diagnostics["unresolved_value_text_top"] = {}
        diagnostics["unresolved_label_top"] = {}
        return panel, diagnostics

    unresolved_labs = labs.merge(unresolved, on=key_cols, how="inner")
    unresolved_value_text = (
        unresolved_labs["value_text_norm"]
        .replace({"": pd.NA})
        .dropna()
        .astype(str)
        .value_counts()
        .head(15)
    )
    unresolved_label_text = (
        unresolved_labs["label_fluid_text"]
        .replace({"": pd.NA})
        .dropna()
        .astype(str)
        .value_counts()
        .head(15)
    )
    diagnostics["unresolved_value_text_top"] = {
        str(key): int(value) for key, value in unresolved_value_text.items()
    }
    diagnostics["unresolved_label_top"] = {
        str(key): int(value) for key, value in unresolved_label_text.items()
    }
    return panel, diagnostics


def summarize_gas_source(panel_df: pd.DataFrame) -> dict[str, Any]:
    """Summarize gas source composition from panel-level records."""
    total_rows = int(len(panel_df))
    if total_rows == 0:
        return {
```

```python
        "panel_rows": 0,
        "source_present": bool("source" in panel_df.columns),
        "source_counts": {},
        "source_rates": {},
        "all_other_or_unknown": False,
        "tier_counts": {},
        "tier_rates": {},
    }

if "source" not in panel_df.columns:
    return {
        "panel_rows": total_rows,
        "source_present": False,
        "source_counts": {},
        "source_rates": {},
        "all_other_or_unknown": True,
        "tier_counts": {},
        "tier_rates": {},
    }

source = (
    panel_df["source"]
    .astype(str)
    .str.strip()
    .str.lower()
    .replace({"": "other", "nan": "other", "none": "other"})
    .fillna("other")
)
source_counts = {str(key): int(value) for key, value in
↪  source.value_counts(dropna=False).items()}
source_rates = {key: float(value / total_rows) for key, value in
↪  source_counts.items()}
all_other_or_unknown = bool(
    total_rows > 0 and set(source_counts.keys()).issubset({"other",
↪  "unknown"})
)
tier_counts: dict[str, int] = {}
tier_rates: dict[str, float] = {}
if "source_inference_tier" in panel_df.columns:
    tier = (
        panel_df["source_inference_tier"]
        .astype(str)
        .str.strip()
```

```python
            .str.lower()
            .replace({"": "unknown", "nan": "unknown", "none": "unknown"})
            .fillna("unknown")
        )
        tier_counts = {
            str(key): int(value) for key, value in
            ↪   tier.value_counts(dropna=False).items()
        }
        tier_rates = {key: float(value / total_rows) for key, value in
    ↪   tier_counts.items()}

    return {
        "panel_rows": total_rows,
        "source_present": True,
        "source_counts": source_counts,
        "source_rates": source_rates,
        "all_other_or_unknown": all_other_or_unknown,
        "tier_counts": tier_counts,
        "tier_rates": tier_rates,
    }


def assert_gas_source_coverage(
    gas_source_audit: Mapping[str, Any],
    *,
    fail_on_all_other_source: bool = True,
) -> None:
    """Fail fast when source attribution collapses to all other/unknown."""
    if not fail_on_all_other_source:
        return
    if int(gas_source_audit.get("panel_rows", 0)) <= 0:
        return
    if bool(gas_source_audit.get("all_other_or_unknown", False)):
        raise ValueError(
            "Gas source attribution classified all panel rows as
            ↪   other/unknown. "
            "Set COHORT_FAIL_ON_ALL_OTHER_SOURCE=0 to bypass this guard."
        )


def build_gas_source_overlap_summary(ed_df: pd.DataFrame) -> pd.DataFrame:
    """Build ABG/VBG/OTHER overlap counts and percentages."""
    frame = ed_df.copy()
```

```python
    abg = pd.to_numeric(frame.get("flag_abg_hypercapnia", 0),
↪   errors="coerce").fillna(0).astype(int)
    vbg = pd.to_numeric(frame.get("flag_vbg_hypercapnia", 0),
↪   errors="coerce").fillna(0).astype(int)
    other = pd.to_numeric(frame.get("flag_other_hypercapnia", 0),
↪   errors="coerce").fillna(0).astype(int)

    labels = []
    for a, v, o in zip(abg.tolist(), vbg.tolist(), other.tolist(),
↪   strict=False):
        parts: list[str] = []
        if a == 1:
            parts.append("ABG")
        if v == 1:
            parts.append("VBG")
        if o == 1:
            parts.append("OTHER")
        labels.append("+".join(parts) if parts else "NO_GAS")

    counts = pd.Series(labels, dtype="string").value_counts(dropna=False
↪   ).rename_axis("gas_overlap").reset_index(name="count")
    total = max(int(counts["count"].sum()), 1)
    counts["percent"] = counts["count"].astype(float) / total * 100.0
    counts["percent"] = counts["percent"].round(2)
    return counts.sort_values(["count", "gas_overlap"], ascending=[False,
↪   True]).reset_index(drop=True)


def add_gas_model_fields(
    ed_df: pd.DataFrame,
    *,
    ranges: Mapping[str, tuple[float, float]] | None = None,
) -> tuple[pd.DataFrame, pd.DataFrame]:
    """Create cleaned gas model fields and a field-level outlier audit."""
    resolved_ranges = dict(ranges or DEFAULT_GAS_MODEL_RANGES)
    updated = ed_df.copy()
    audit_rows: list[dict[str, Any]] = []

    for raw_column, (lower_bound, upper_bound) in resolved_ranges.items():
        if raw_column not in updated.columns:
            continue
        numeric = pd.to_numeric(updated[raw_column], errors="coerce")
```

```python
        out_of_range = numeric.notna() & ((numeric < lower_bound) | (numeric
↪   > upper_bound))
        model_column = f"{raw_column}_model"
        outlier_flag_column = f"{raw_column}_outlier_flag"
        updated[model_column] = numeric.where(~out_of_range)
        updated[outlier_flag_column] = out_of_range.astype("boolean")
        nonnull_n = int(numeric.notna().sum())
        outlier_n = int(out_of_range.sum())
        examples = (
            numeric.loc[out_of_range]
            .round(4)
            .value_counts()
            .head(5)
            .index.astype(str)
            .tolist()
        )
        audit_rows.append(
            {
                "domain": "gas",
                "raw_column": raw_column,
                "model_column": model_column,
                "outlier_flag_column": outlier_flag_column,
                "lower_bound": float(lower_bound),
                "upper_bound": float(upper_bound),
                "nonnull_n": nonnull_n,
                "out_of_range_n": outlier_n,
                "out_of_range_pct": float(outlier_n / nonnull_n) if nonnull_n
                ↪   else 0.0,
                "example_outlier_values": "; ".join(examples),
            }
        )

    audit =
↪   pd.DataFrame(audit_rows).sort_values("raw_column").reset_index(drop=True)
    return updated, audit


def attach_charted_anthro_fallback(
    ed_df: pd.DataFrame,
    charted_df: pd.DataFrame,
    *,
    nearest_anytime: bool = True,
) -> tuple[pd.DataFrame, dict[str, Any]]:
```

```python
    """Fill missing anthropometrics from charted records with nearest-time
    ↪  selection."""
    required_ed = {"ed_stay_id", "subject_id", "ed_intime"}
    missing_ed = sorted(required_ed.difference(ed_df.columns))
    if missing_ed:
        raise KeyError(f"attach_charted_anthro_fallback missing ED columns:
        ↪  {missing_ed}")

    required_charted = {"subject_id", "obs_time", "result_name",
↪ "result_value_num"}
    missing_charted = sorted(required_charted.difference(charted_df.columns))
    if missing_charted:
        raise KeyError(
            "attach_charted_anthro_fallback missing charted columns: "
            f"{missing_charted}"
        )

    updated = ed_df.copy()
    for output_column in OMR_OUTPUT_COLUMNS:
        if output_column not in updated.columns:
            updated[output_column] = pd.NA
    if "anthro_source" not in updated.columns:
        updated["anthro_source"] = "missing"
    else:
        updated["anthro_source"] = (
            updated["anthro_source"].astype(str).replace({"":
↪ "missing"}).fillna("missing")
        )
    if "anthro_obstime" not in updated.columns:
        updated["anthro_obstime"] = pd.NaT
    else:
        updated["anthro_obstime"] = pd.to_datetime(updated["anthro_obstime"],
↪ errors="coerce")
    if "anthro_hours_offset" not in updated.columns:
        updated["anthro_hours_offset"] = pd.NA
    updated["anthro_hours_offset"] =
↪ pd.to_numeric(updated["anthro_hours_offset"], errors="coerce")
    if "anthro_timing_basis" not in updated.columns:
        updated["anthro_timing_basis"] = "missing"
    if "anthro_timing_uncertain" not in updated.columns:
        updated["anthro_timing_uncertain"] = pd.Series([pd.NA] *
↪ len(updated), dtype="boolean")
    else:
```

```python
        updated["anthro_timing_uncertain"] =
↪   updated["anthro_timing_uncertain"].astype("boolean")

    charted = charted_df.copy()
    charted["subject_id"] = _to_int64(charted["subject_id"])
    charted["obs_time"] = pd.to_datetime(charted["obs_time"],
↪   errors="coerce")
    charted["result_name"] =
↪   charted["result_name"].astype(str).str.strip().str.lower()
    charted["result_value_num"] = pd.to_numeric(charted["result_value_num"],
↪   errors="coerce")
    if "source" not in charted.columns:
        charted["source"] = "icu_charted"
    charted["source"] = charted["source"].astype(str).replace({"":
↪   "icu_charted"}).fillna("icu_charted")
    charted = charted.loc[
        charted["subject_id"].notna()
        & charted["obs_time"].notna()
        & charted["result_name"].isin(OMR_RESULT_NAMES)
        & charted["result_value_num"].notna()
    ].copy()

    diagnostics: dict[str, Any] = {
        "charted_rows_input": int(len(charted_df)),
        "charted_rows_usable": int(len(charted)),
        "nearest_anytime": bool(nearest_anytime),
        "filled_counts": {column: 0 for column in OMR_OUTPUT_COLUMNS},
        "rows_with_any_charted_fill": 0,
    }
    if charted.empty:
        return updated, diagnostics

    ed_norm = updated[["ed_stay_id", "subject_id", "ed_intime"]].copy()
    ed_norm["subject_id"] = _to_int64(ed_norm["subject_id"])
    ed_norm["ed_intime_dt"] = pd.to_datetime(ed_norm["ed_intime"],
↪   errors="coerce")
    ed_norm = ed_norm.loc[ed_norm["subject_id"].notna() &
↪   ed_norm["ed_intime_dt"].notna()].copy()
    if ed_norm.empty:
        return updated, diagnostics

    merged = ed_norm.merge(charted, on="subject_id", how="inner")
    if merged.empty:
```

```python
        diagnostics["subject_overlap_count"] = 0
        return updated, diagnostics
    diagnostics["subject_overlap_count"] =
↪   int(merged["subject_id"].nunique())
    merged["hours_offset"] = (
        (merged["obs_time"] - merged["ed_intime_dt"]).dt.total_seconds() /
↪   3600.0
    )
    merged["abs_hours_offset"] = merged["hours_offset"].abs()
    if not nearest_anytime:
        merged = merged.loc[merged["abs_hours_offset"] <= 24.0].copy()

    output_to_name = {
        "bmi_closest_pre_ed": "bmi",
        "height_closest_pre_ed": "height",
        "weight_closest_pre_ed": "weight",
    }
    any_fill_mask = pd.Series(False, index=updated.index)
    for output_column, result_name in output_to_name.items():
        subset = merged.loc[merged["result_name"] == result_name].copy()
        if subset.empty:
            continue
        selected = (
            subset.sort_values(
                ["ed_stay_id", "abs_hours_offset", "obs_time"],
                ascending=[True, True, True],
            )
            .groupby("ed_stay_id", as_index=False)
            .first()
        )
        selected = selected.rename(
            columns={
                "result_value_num": f"{output_column}__candidate",
                "obs_time": f"{output_column}__obs_time",
                "hours_offset": f"{output_column}__hours_offset",
                "source": f"{output_column}__source",
            }
        )
        keep_columns = [
            "ed_stay_id",
            f"{output_column}__candidate",
            f"{output_column}__obs_time",
            f"{output_column}__hours_offset",
```

```
            f"{output_column}__source",
        ]
        updated = updated.merge(selected[keep_columns], on="ed_stay_id",
↪  how="left")

        fill_mask = updated[output_column].isna() &
↪  updated[f"{output_column}__candidate"].notna()
        updated.loc[fill_mask, output_column] = updated.loc[fill_mask,
↪  f"{output_column}__candidate"]

        provenance_mask = fill_mask &
↪  updated["anthro_source"].isin({"missing", "nan"})
        updated.loc[provenance_mask, "anthro_source"] = updated.loc[
            provenance_mask, f"{output_column}__source"
        ]
        updated.loc[provenance_mask, "anthro_obstime"] = pd.to_datetime(
            updated.loc[provenance_mask, f"{output_column}__obs_time"],
↪  errors="coerce"
        )
        updated.loc[provenance_mask, "anthro_hours_offset"] = pd.to_numeric(
            updated.loc[provenance_mask, f"{output_column}__hours_offset"],
↪  errors="coerce"
        )
        updated.loc[provenance_mask, "anthro_timing_basis"] =
↪  "nearest_anytime"
        updated.loc[provenance_mask, "anthro_timing_uncertain"] = True

        any_fill_mask = any_fill_mask | fill_mask
        diagnostics["filled_counts"][output_column] = int(fill_mask.sum())

        drop_columns = [
            f"{output_column}__candidate",
            f"{output_column}__obs_time",
            f"{output_column}__hours_offset",
            f"{output_column}__source",
        ]
        updated = updated.drop(columns=drop_columns)

    diagnostics["rows_with_any_charted_fill"] = int(any_fill_mask.sum())
    diagnostics["fallback_source_counts"] = {
        str(key): int(value)
        for key, value in updated.loc[any_fill_mask, "anthro_source"]
        .fillna("missing")
```

```python
                .astype(str)
                .value_counts(dropna=False)
                .items()
        }
        return updated, diagnostics


def build_anthro_coverage_audit(ed_df: pd.DataFrame) -> dict[str, Any]:
    """Summarize anthropometric coverage and provenance rates."""
    total_rows = max(int(len(ed_df)), 1)
    field_counts = {
        column: int(pd.to_numeric(ed_df[column],
    ↪ errors="coerce").notna().sum())
        for column in OMR_OUTPUT_COLUMNS
        if column in ed_df.columns
    }
    field_rates = {column: float(count / total_rows) for column, count in
    ↪ field_counts.items()}

    source_counts: dict[str, int] = {}
    source_rates: dict[str, float] = {}
    if "anthro_source" in ed_df.columns:
        source_counts = {
            str(key): int(value)
            for key, value in ed_df["anthro_source"]
            .fillna("missing")
            .astype(str)
            .value_counts(dropna=False)
            .items()
        }
        source_rates = {key: float(value / total_rows) for key, value in
    ↪ source_counts.items()}

    timing_basis_counts: dict[str, int] = {}
    timing_basis_rates: dict[str, float] = {}
    if "anthro_timing_basis" in ed_df.columns:
        timing_basis_counts = {
            str(key): int(value)
            for key, value in ed_df["anthro_timing_basis"]
            .fillna("missing")
            .astype(str)
            .value_counts(dropna=False)
            .items()
```

```python
        }
        timing_basis_rates = {
            key: float(value / total_rows) for key, value in
↪   timing_basis_counts.items()
        }

    return {
        "row_count": int(len(ed_df)),
        "field_nonnull_counts": field_counts,
        "field_nonnull_rates": field_rates,
        "source_counts": source_counts,
        "source_rates": source_rates,
        "timing_basis_counts": timing_basis_counts,
        "timing_basis_rates": timing_basis_rates,
    }



def build_first_other_pco2_audit(ed_df: pd.DataFrame) -> pd.DataFrame:
    """Build route-stratified audit summary for first_other_pco2 values.

    This helper is intentionally tolerant for notebook QA execution: if the
    required fields are not present, it returns a sentinel audit row instead
↪   of
    raising.
    """
    columns = [
        "source",
        "count_nonnull",
        "mean",
        "median",
        "q25",
        "q75",
        "p95",
        "max",
        "pct_ge_80",
        "pct_ge_100",
        "pct_ge_150",
        "pct_eq_160",
        "top_values",
        "status",
        "missing_columns",
    ]
    required = {"first_other_pco2", "first_other_src"}
```

```python
    missing = sorted(required.difference(ed_df.columns))
    if missing:
        return pd.DataFrame(
            [
                {
                    "source": "UNAVAILABLE",
                    "count_nonnull": 0,
                    "mean": None,
                    "median": None,
                    "q25": None,
                    "q75": None,
                    "p95": None,
                    "max": None,
                    "pct_ge_80": 0.0,
                    "pct_ge_100": 0.0,
                    "pct_ge_150": 0.0,
                    "pct_eq_160": 0.0,
                    "top_values": {},
                    "status": "missing_columns",
                    "missing_columns": ",".join(missing),
                }
            ],
            columns=columns,
        )

    frame = ed_df[["first_other_src", "first_other_pco2"]].copy()
    frame["first_other_src"] = (

↪   frame["first_other_src"].astype(str).str.strip().str.upper().replace({"":
↪   "UNKNOWN", "NAN": "UNKNOWN"})
    )
    frame["first_other_pco2"] = pd.to_numeric(frame["first_other_pco2"],
↪   errors="coerce")
    frame = frame.loc[frame["first_other_pco2"].notna()].copy()

    if frame.empty:
        return pd.DataFrame(
            [
                {
                    "source": "UNAVAILABLE",
                    "count_nonnull": 0,
                    "mean": None,
                    "median": None,
```

```python
                    "q25": None,
                    "q75": None,
                    "p95": None,
                    "max": None,
                    "pct_ge_80": 0.0,
                    "pct_ge_100": 0.0,
                    "pct_ge_150": 0.0,
                    "pct_eq_160": 0.0,
                    "top_values": {},
                    "status": "no_nonnull_values",
                    "missing_columns": "",
                }
            ],
            columns=columns,
        )

    rows: list[dict[str, Any]] = []
    for source_name, group in frame.groupby("first_other_src"):
        values = group["first_other_pco2"]
        rows.append(
            {
                "source": source_name,
                "count_nonnull": int(values.shape[0]),
                "mean": float(values.mean()),
                "median": float(values.median()),
                "q25": float(values.quantile(0.25)),
                "q75": float(values.quantile(0.75)),
                "p95": float(values.quantile(0.95)),
                "max": float(values.max()),
                "pct_ge_80": float((values >= 80).mean()),
                "pct_ge_100": float((values >= 100).mean()),
                "pct_ge_150": float((values >= 150).mean()),
                "pct_eq_160": float((values == 160).mean()),
                "top_values": {
                    str(key): int(value)
                    for key, value in values.value_counts().head(10).items()
                },
                "status": "ok",
                "missing_columns": "",
            }
        )
    return pd.DataFrame(rows,
    ↪  columns=columns).sort_values("source").reset_index(drop=True)
```

```python
def normalize_temperature_to_f(temp: pd.Series) -> pd.DataFrame:
    """Normalize mixed-unit temperature values to cleaned Fahrenheit output.

    Rules:
    - values in ``[20, 50]`` are treated as Celsius-like and converted to
↪   Fahrenheit,
    - all other numeric values are treated as Fahrenheit already,
    - cleaned output is valid only for Fahrenheit ``[50, 120]``.
    """
    numeric = pd.to_numeric(temp, errors="coerce")
    celsius_like = numeric.between(20.0, 50.0, inclusive="both")
    temp_f = numeric.where(~celsius_like, numeric * 9.0 / 5.0 + 32.0)
    out_of_range = ~temp_f.between(50.0, 120.0, inclusive="both")
    temp_f_clean = temp_f.where(~out_of_range)
    return pd.DataFrame(
        {
            "temp_f_clean": temp_f_clean,
            "temp_was_celsius_like": celsius_like.astype("boolean"),
            "temp_out_of_range": out_of_range.astype("boolean"),
        }
    )


def clean_pain_score(pain: pd.Series) -> pd.DataFrame:
    """Clean pain score with explicit sentinel handling."""
    numeric = pd.to_numeric(pain, errors="coerce")
    sentinel = numeric.eq(13)
    out_of_range = numeric.notna() & (~sentinel) & ((numeric < 0) | (numeric
↪   > 10))
    clean = numeric.where(~(sentinel | out_of_range))
    return pd.DataFrame(
        {
            "pain_clean": clean,
            "pain_is_sentinel_13": sentinel.astype("boolean"),
            "pain_out_of_range": out_of_range.astype("boolean"),
        }
    )


def clean_bp(sbp: pd.Series, dbp: pd.Series) -> pd.DataFrame:
    """Range-clean SBP/DBP values with deterministic bounds."""
```

```python
        sbp_num = pd.to_numeric(sbp, errors="coerce")
        dbp_num = pd.to_numeric(dbp, errors="coerce")
        sbp_out_of_range = sbp_num.notna() & ((sbp_num < 20) | (sbp_num > 300))
        dbp_out_of_range = dbp_num.notna() & ((dbp_num < 10) | (dbp_num > 200))
        return pd.DataFrame(
            {
                "sbp_clean": sbp_num.where(~sbp_out_of_range),
                "dbp_clean": dbp_num.where(~dbp_out_of_range),
                "sbp_out_of_range": sbp_out_of_range.astype("boolean"),
                "dbp_out_of_range": dbp_out_of_range.astype("boolean"),
            }
        )


    def clean_o2sat(o2sat: pd.Series) -> pd.DataFrame:
        """Clean oxygen saturation and keep explicit >100 and zero flags."""
        numeric = pd.to_numeric(o2sat, errors="coerce")
        gt_100 = numeric > 100
        out_of_range = numeric.notna() & (numeric < 0)
        clean = numeric.where(~(gt_100 | out_of_range))
        return pd.DataFrame(
            {
                "o2sat_clean": clean,
                "o2sat_gt_100": gt_100.astype("boolean"),
                "o2sat_out_of_range": out_of_range.astype("boolean"),
                "o2sat_zero": numeric.eq(0).astype("boolean"),
            }
        )


    def _series_distribution_stats(values: pd.Series) -> dict[str, Any]:
        """Return compact distribution metrics for audit CSV artifacts."""
        numeric = pd.to_numeric(values, errors="coerce")
        total_rows = int(len(numeric))
        non_missing = int(numeric.notna().sum())
        if non_missing == 0:
            return {
                "n_non_missing": 0,
                "missing_pct": 100.0 if total_rows else 0.0,
                "min": np.nan,
                "p1": np.nan,
                "median": np.nan,
                "mean": np.nan,
```

```python
            "p99": np.nan,
            "max": np.nan,
        }
    return {
        "n_non_missing": non_missing,
        "missing_pct": round(float((1.0 - non_missing / max(total_rows, 1)) *
        ↪   100.0), 2),
        "min": float(numeric.min()),
        "p1": float(numeric.quantile(0.01)),
        "median": float(numeric.median()),
        "mean": float(numeric.mean()),
        "p99": float(numeric.quantile(0.99)),
        "max": float(numeric.max()),
    }


def add_vitals_model_fields(
    ed_df: pd.DataFrame,
    *,
    ranges: Mapping[str, tuple[float, float]] | None = None,
) -> tuple[pd.DataFrame, pd.DataFrame]:
    """Create cleaned ED vitals fields and backward-compatible ``*_model``
    ↪   aliases."""
    updated = ed_df.copy()
    resolved_ranges = dict(ranges or DEFAULT_VITALS_MODEL_RANGES)
    audit_rows: list[dict[str, Any]] = []

    def _append_audit_row(
        *,
        raw_column: str,
        model_column: str,
        outlier_flag_column: str,
        lower_bound: float | None,
        upper_bound: float | None,
        numeric: pd.Series,
        outlier_flag: pd.Series,
    ) -> None:
        outlier_flag_numeric = outlier_flag.fillna(False).astype(bool)
        nonnull_n = int(numeric.notna().sum())
        outlier_n = int(outlier_flag_numeric.sum())
        examples = (
            numeric.loc[outlier_flag_numeric]
            .round(4)
```

```
                    .value_counts()
                    .head(5)
                    .index.astype(str)
                    .tolist()
            )
        audit_rows.append(
            {
                "domain": "vitals",
                "raw_column": raw_column,
                "model_column": model_column,
                "outlier_flag_column": outlier_flag_column,
                "lower_bound": lower_bound,
                "upper_bound": upper_bound,
                "nonnull_n": nonnull_n,
                "out_of_range_n": outlier_n,
                "out_of_range_pct": float(outlier_n / nonnull_n) if nonnull_n
                ↪    else 0.0,
                "example_outlier_values": "; ".join(examples),
            }
        )

    for raw_column, (lower_bound, upper_bound) in resolved_ranges.items():
        if raw_column not in updated.columns:
            continue
        numeric = pd.to_numeric(updated[raw_column], errors="coerce")
        out_of_range = numeric.notna() & ((numeric < lower_bound) | (numeric
↪  > upper_bound))
        model_column = f"{raw_column}_model"
        outlier_flag_column = f"{raw_column}_outlier_flag"
        updated[model_column] = numeric.where(~out_of_range)
        updated[outlier_flag_column] = out_of_range.astype("boolean")
        _append_audit_row(
            raw_column=raw_column,
            model_column=model_column,
            outlier_flag_column=outlier_flag_column,
            lower_bound=float(lower_bound),
            upper_bound=float(upper_bound),
            numeric=numeric,
            outlier_flag=out_of_range,
        )

    for prefix in ("ed_triage", "ed_first"):
        temp_raw_col = f"{prefix}_temp"
```

```python
        pain_raw_col = f"{prefix}_pain"
        sbp_raw_col = f"{prefix}_sbp"
        dbp_raw_col = f"{prefix}_dbp"
        o2sat_raw_col = f"{prefix}_o2sat"

        if temp_raw_col in updated.columns:
            temp_clean = normalize_temperature_to_f(updated[temp_raw_col])
            temp_clean_col = f"{prefix}_temp_f_clean"
            temp_celsius_flag_col = f"{prefix}_temp_was_celsius_like"
            temp_range_flag_col = f"{prefix}_temp_out_of_range"
            updated[temp_clean_col] = temp_clean["temp_f_clean"]
            updated[temp_celsius_flag_col] =
↪  temp_clean["temp_was_celsius_like"]
            updated[temp_range_flag_col] = temp_clean["temp_out_of_range"]
            updated[f"{prefix}_temp_f_model"] = updated[temp_clean_col]
            updated[f"{prefix}_temp_c_model"] =
↪  (updated[f"{prefix}_temp_f_model"] - 32.0) * (5.0 / 9.0)
            updated[f"{prefix}_temp_model"] =
↪  updated[f"{prefix}_temp_f_model"]
            updated[f"{prefix}_temp_outlier_flag"] =
↪  updated[temp_range_flag_col].astype("boolean")
            _append_audit_row(
                raw_column=temp_raw_col,
                model_column=f"{prefix}_temp_model",
                outlier_flag_column=f"{prefix}_temp_outlier_flag",
                lower_bound=50.0,
                upper_bound=120.0,
                numeric=pd.to_numeric(updated[temp_raw_col],
↪  errors="coerce"),
                outlier_flag=updated[temp_range_flag_col],
            )

        if pain_raw_col in updated.columns:
            pain_clean = clean_pain_score(updated[pain_raw_col])
            pain_clean_col = f"{prefix}_pain_clean"
            pain_sentinel_col = f"{prefix}_pain_is_sentinel_13"
            pain_range_col = f"{prefix}_pain_out_of_range"
            updated[pain_clean_col] = pain_clean["pain_clean"]
            updated[pain_sentinel_col] = pain_clean["pain_is_sentinel_13"]
            updated[pain_range_col] = pain_clean["pain_out_of_range"]
            updated[f"{prefix}_pain_model"] = updated[pain_clean_col]
            pain_outlier = updated[pain_sentinel_col].fillna(False) |
↪  updated[pain_range_col].fillna(False)
```

```python
        updated[f"{prefix}_pain_outlier_flag"] =
↪  pain_outlier.astype("boolean")
        _append_audit_row(
            raw_column=pain_raw_col,
            model_column=f"{prefix}_pain_model",
            outlier_flag_column=f"{prefix}_pain_outlier_flag",
            lower_bound=0.0,
            upper_bound=10.0,
            numeric=pd.to_numeric(updated[pain_raw_col],
↪  errors="coerce"),
            outlier_flag=pain_outlier,
        )

    if sbp_raw_col in updated.columns and dbp_raw_col in updated.columns:
        bp_clean = clean_bp(updated[sbp_raw_col], updated[dbp_raw_col])
        sbp_clean_col = f"{prefix}_sbp_clean"
        dbp_clean_col = f"{prefix}_dbp_clean"
        sbp_range_col = f"{prefix}_sbp_out_of_range"
        dbp_range_col = f"{prefix}_dbp_out_of_range"
        updated[sbp_clean_col] = bp_clean["sbp_clean"]
        updated[dbp_clean_col] = bp_clean["dbp_clean"]
        updated[sbp_range_col] = bp_clean["sbp_out_of_range"]
        updated[dbp_range_col] = bp_clean["dbp_out_of_range"]
        updated[f"{prefix}_sbp_model"] = updated[sbp_clean_col]
        updated[f"{prefix}_dbp_model"] = updated[dbp_clean_col]
        updated[f"{prefix}_sbp_outlier_flag"] =
↪  updated[sbp_range_col].astype("boolean")
        updated[f"{prefix}_dbp_outlier_flag"] =
↪  updated[dbp_range_col].astype("boolean")
        _append_audit_row(
            raw_column=sbp_raw_col,
            model_column=f"{prefix}_sbp_model",
            outlier_flag_column=f"{prefix}_sbp_outlier_flag",
            lower_bound=20.0,
            upper_bound=300.0,
            numeric=pd.to_numeric(updated[sbp_raw_col], errors="coerce"),
            outlier_flag=updated[sbp_range_col],
        )
        _append_audit_row(
            raw_column=dbp_raw_col,
            model_column=f"{prefix}_dbp_model",
            outlier_flag_column=f"{prefix}_dbp_outlier_flag",
            lower_bound=10.0,
```

```python
                upper_bound=200.0,
                numeric=pd.to_numeric(updated[dbp_raw_col], errors="coerce"),
                outlier_flag=updated[dbp_range_col],
            )

        if o2sat_raw_col in updated.columns:
            o2sat_clean = clean_o2sat(updated[o2sat_raw_col])
            o2sat_clean_col = f"{prefix}_o2sat_clean"
            o2sat_gt_100_col = f"{prefix}_o2sat_gt_100"
            o2sat_range_col = f"{prefix}_o2sat_out_of_range"
            o2sat_zero_col = f"{prefix}_o2sat_zero"
            updated[o2sat_clean_col] = o2sat_clean["o2sat_clean"]
            updated[o2sat_gt_100_col] = o2sat_clean["o2sat_gt_100"]
            updated[o2sat_range_col] = o2sat_clean["o2sat_out_of_range"]
            updated[o2sat_zero_col] = o2sat_clean["o2sat_zero"]
            updated[f"{prefix}_o2sat_model"] = updated[o2sat_clean_col]
            o2sat_outlier = updated[o2sat_gt_100_col].fillna(False) |
↪   updated[o2sat_range_col].fillna(False)
            updated[f"{prefix}_o2sat_outlier_flag"] =
↪   o2sat_outlier.astype("boolean")
            _append_audit_row(
                raw_column=o2sat_raw_col,
                model_column=f"{prefix}_o2sat_model",
                outlier_flag_column=f"{prefix}_o2sat_outlier_flag",
                lower_bound=0.0,
                upper_bound=100.0,
                numeric=pd.to_numeric(updated[o2sat_raw_col],
↪   errors="coerce"),
                outlier_flag=o2sat_outlier,
            )

    audit =
↪   pd.DataFrame(audit_rows).sort_values("raw_column").reset_index(drop=True)
    return updated, audit


def build_ed_vitals_audit_artifacts(
    ed_df: pd.DataFrame,
) -> tuple[pd.DataFrame, pd.DataFrame, pd.DataFrame]:
    """Build ED-vitals distribution, extreme-example, and raw-to-clean delta
    ↪   artifacts."""
    specs = [
        {
```

```
        "vital_name": "triage_temp",
        "raw_col": "ed_triage_temp",
        "clean_col": "ed_triage_temp_f_clean",
        "flags": ["ed_triage_temp_was_celsius_like",
        ↪  "ed_triage_temp_out_of_range"],
        "extra_cols": [],
    },
    {
        "vital_name": "first_temp",
        "raw_col": "ed_first_temp",
        "clean_col": "ed_first_temp_f_clean",
        "flags": ["ed_first_temp_was_celsius_like",
        ↪  "ed_first_temp_out_of_range"],
        "extra_cols": ["ed_first_vitals_time"],
    },
    {
        "vital_name": "triage_pain",
        "raw_col": "ed_triage_pain",
        "clean_col": "ed_triage_pain_clean",
        "flags": ["ed_triage_pain_is_sentinel_13",
        ↪  "ed_triage_pain_out_of_range"],
        "extra_cols": [],
    },
    {
        "vital_name": "first_pain",
        "raw_col": "ed_first_pain",
        "clean_col": "ed_first_pain_clean",
        "flags": ["ed_first_pain_is_sentinel_13",
        ↪  "ed_first_pain_out_of_range"],
        "extra_cols": ["ed_first_vitals_time"],
    },
    {
        "vital_name": "triage_sbp",
        "raw_col": "ed_triage_sbp",
        "clean_col": "ed_triage_sbp_clean",
        "flags": ["ed_triage_sbp_out_of_range"],
        "extra_cols": [],
    },
    {
        "vital_name": "triage_dbp",
        "raw_col": "ed_triage_dbp",
        "clean_col": "ed_triage_dbp_clean",
        "flags": ["ed_triage_dbp_out_of_range"],
```

```python
            "extra_cols": [],
        },
        {
            "vital_name": "first_sbp",
            "raw_col": "ed_first_sbp",
            "clean_col": "ed_first_sbp_clean",
            "flags": ["ed_first_sbp_out_of_range"],
            "extra_cols": ["ed_first_vitals_time"],
        },
        {
            "vital_name": "first_dbp",
            "raw_col": "ed_first_dbp",
            "clean_col": "ed_first_dbp_clean",
            "flags": ["ed_first_dbp_out_of_range"],
            "extra_cols": ["ed_first_vitals_time"],
        },
        {
            "vital_name": "triage_o2sat",
            "raw_col": "ed_triage_o2sat",
            "clean_col": "ed_triage_o2sat_clean",
            "flags": ["ed_triage_o2sat_gt_100",
            ↪  "ed_triage_o2sat_out_of_range", "ed_triage_o2sat_zero"],
            "extra_cols": [],
        },
        {
            "vital_name": "first_o2sat",
            "raw_col": "ed_first_o2sat",
            "clean_col": "ed_first_o2sat_clean",
            "flags": ["ed_first_o2sat_gt_100", "ed_first_o2sat_out_of_range",
            ↪  "ed_first_o2sat_zero"],
            "extra_cols": ["ed_first_vitals_time"],
        },
    ]

    summary_rows: list[dict[str, Any]] = []
    delta_rows: list[dict[str, Any]] = []
    extreme_frames: list[pd.DataFrame] = []

    for spec in specs:
        raw_col = spec["raw_col"]
        clean_col = spec["clean_col"]
        if raw_col not in ed_df.columns or clean_col not in ed_df.columns:
            continue
```

```python
raw = pd.to_numeric(ed_df[raw_col], errors="coerce")
clean = pd.to_numeric(ed_df[clean_col], errors="coerce")
raw_stats = _series_distribution_stats(raw)
clean_stats = _series_distribution_stats(clean)
row: dict[str, Any] = {
    "vital_name": spec["vital_name"],
    "raw_column": raw_col,
    "clean_column": clean_col,
    "raw_n_non_missing": raw_stats["n_non_missing"],
    "raw_missing_pct": raw_stats["missing_pct"],
    "raw_min": raw_stats["min"],
    "raw_p1": raw_stats["p1"],
    "raw_median": raw_stats["median"],
    "raw_mean": raw_stats["mean"],
    "raw_p99": raw_stats["p99"],
    "raw_max": raw_stats["max"],
    "clean_n_non_missing": clean_stats["n_non_missing"],
    "clean_missing_pct": clean_stats["missing_pct"],
    "clean_min": clean_stats["min"],
    "clean_p1": clean_stats["p1"],
    "clean_median": clean_stats["median"],
    "clean_mean": clean_stats["mean"],
    "clean_p99": clean_stats["p99"],
    "clean_max": clean_stats["max"],
}

delta_rows.append(
    {
        "vital_name": spec["vital_name"],
        "raw_column": raw_col,
        "clean_column": clean_col,
        "raw_n": int(raw.notna().sum()),
        "clean_n": int(clean.notna().sum()),
        "raw_removed_n": int(raw.notna().sum() -
        ↪  clean.notna().sum()),
        "raw_removed_pct": (
            float((raw.notna().sum() - clean.notna().sum()) /
            ↪  max(int(raw.notna().sum()), 1))
            if int(raw.notna().sum()) > 0
            else 0.0
        ),
        "converted_celsius_like_n": int(
```

```
                    ed_df.get(spec["flags"][0], pd.Series(False,
↪  index=ed_df.index))
                        .fillna(False)
                        .astype(bool)
                        .sum()
                )
                if "temp_was_celsius_like" in spec["flags"][0]
                else 0,
            }
        )

        available_flags = [flag for flag in spec["flags"] if flag in
↪  ed_df.columns]
        for flag_col in available_flags:
            flag_series = ed_df[flag_col].fillna(False).astype(bool)
            row[f"{flag_col}_count"] = int(flag_series.sum())
            row[f"{flag_col}_pct"] = float(flag_series.mean()) if
↪  len(flag_series) else 0.0
            if not flag_series.any():
                continue
            include_cols = ["hadm_id", "ed_stay_id", raw_col, clean_col,
↪  flag_col]
            include_cols += [col for col in spec["extra_cols"] if col in
↪  ed_df.columns]
            include_cols = [col for col in include_cols if col in
↪  ed_df.columns]
            flagged = ed_df.loc[flag_series, include_cols].copy()
            flagged = flagged.rename(columns={raw_col: "raw_value",
↪  clean_col: "clean_value", flag_col: "flag_value"})
            low = flagged.sort_values("raw_value", ascending=True,
↪  na_position="last").head(5).copy()
            low["tail"] = "low"
            high = flagged.sort_values("raw_value", ascending=False,
↪  na_position="last").head(5).copy()
            high["tail"] = "high"
            extremes = pd.concat([low, high],
↪  ignore_index=True).drop_duplicates()
            extremes.insert(0, "vital_name", spec["vital_name"])
            extremes.insert(1, "raw_column", raw_col)
            extremes.insert(2, "clean_column", clean_col)
            extremes.insert(3, "flag_type", flag_col)
            extreme_frames.append(extremes)
```

```python
        summary_rows.append(row)

    distribution = pd.DataFrame(summary_rows)
    delta = pd.DataFrame(delta_rows)
    if extreme_frames:
        extremes = pd.concat(extreme_frames, ignore_index=True, sort=False)
    else:
        extremes = pd.DataFrame(
            columns=[
                "vital_name",
                "raw_column",
                "clean_column",
                "flag_type",
                "hadm_id",
                "ed_stay_id",
                "raw_value",
                "clean_value",
                "flag_value",
                "tail",
            ]
        )
    return distribution, extremes, delta


def evaluate_uom_expectations(ed_df: pd.DataFrame) -> dict[str, Any]:
    """Evaluate expected-null and value/uom consistency rules."""
    structural_nulls: dict[str, dict[str, Any]] = {}
    for field_name in EXPECTED_STRUCTURAL_NULL_FIELDS:
        if field_name not in ed_df.columns:
            continue
        structural_nulls[field_name] = {
            "present": True,
            "missing_n": int(ed_df[field_name].isna().sum()),
            "missing_pct": float(ed_df[field_name].isna().mean()),
            "all_null": bool(ed_df[field_name].isna().all()),
        }

    paco2_checks: dict[str, dict[str, Any]] = {}
    for value_column, uom_column in PACO2_VALUE_UOM_PAIRS:
        if value_column not in ed_df.columns or uom_column not in
        ↪   ed_df.columns:
            paco2_checks[uom_column] = {
                "present": False,
```

```python
                "reason": "missing_value_or_uom_column",
            }
            continue

        value_present = ed_df[value_column].notna()
        uom_lower = ed_df[uom_column].astype(str).str.strip().str.lower()
        missing_uom_with_value = int((value_present &
  ↪  ed_df[uom_column].isna()).sum())
        non_mmhg_uom_with_value = int(
            (value_present & ed_df[uom_column].notna() &
  ↪  uom_lower.ne("mmhg")).sum()
        )

        paco2_checks[uom_column] = {
            "present": True,
            "paired_value_column": value_column,
            "value_rows": int(value_present.sum()),
            "missing_uom_when_value_present": missing_uom_with_value,
            "non_mmhg_uom_when_value_present": non_mmhg_uom_with_value,
            "passes": bool(
                missing_uom_with_value == 0 and non_mmhg_uom_with_value == 0
            ),
        }

    return {
        "expected_structural_null_fields":
          ↪  list(EXPECTED_STRUCTURAL_NULL_FIELDS),
        "structural_null_checks": structural_nulls,
        "paco2_uom_checks": paco2_checks,
    }


def classify_missingness_expectations(
    ed_df: pd.DataFrame,
    target_fields: list[str],
    *,
    expected_sparse_fields: set[str] | None = None,
) -> pd.DataFrame:
    """Classify field-level missingness into expected and unexpected
      ↪  categories."""
    rows: list[dict[str, Any]] = []
    total_rows = max(int(len(ed_df)), 1)
    expected_structural = set(EXPECTED_STRUCTURAL_NULL_FIELDS)
```

```python
    expected_sparse = set(expected_sparse_fields or set())

    for field_name in target_fields:
        if field_name not in ed_df.columns:
            rows.append(
                {
                    "field": field_name,
                    "missing_n": int(len(ed_df)),
                    "missing_pct": 1.0,
                    "expectation": "missing_column",
                }
            )
            continue

        missing_n = int(ed_df[field_name].isna().sum())
        missing_pct = float(missing_n / total_rows)

        if field_name in expected_structural:
            expectation = "expected_structural_null"
        elif field_name in expected_sparse and missing_pct >= 1.0:
            expectation = "expected_sparse"
        elif missing_pct >= 1.0:
            expectation = "unexpected_full_null"
        elif missing_pct > 0.0:
            expectation = "conditional_sparse"
        else:
            expectation = "complete"

        rows.append(
            {
                "field": field_name,
                "missing_n": missing_n,
                "missing_pct": missing_pct,
                "expectation": expectation,
            }
        )

    return pd.DataFrame(rows)


def render_latex_longtable(
    table_df: pd.DataFrame,
    *,
```

```python
    caption: str,
    label: str,
    landscape: bool = False,
    index: bool = True,
) -> str:
    latex_text = table_df.to_latex(
        index=index,
        escape=False,
        longtable=True,
        caption=caption,
        label=label,
    )
    if landscape:
        latex_text = "\\begin{landscape}\\n" + latex_text +
↪  "\\n\\end{landscape}\\n"
    return latex_text

from hypercap_cc_nlp.pipeline_audit import collect_run_manifest
from hypercap_cc_nlp.pipeline_contracts import validate_cohort_contract,
↪  write_contract_report

# ---- Backend selection (we use BigQuery)
BACKEND = os.getenv("MIMIC_BACKEND", "bigquery").strip().lower()
assert BACKEND == "bigquery", "This notebook is BigQuery-specific."

WORK_PROJECT = os.getenv("WORK_PROJECT", "").strip()  # your billing project
PHYS = os.getenv("BQ_PHYSIONET_PROJECT", "physionet-data").strip()  # hosting
↪  project (read-only)

# Dataset preferences: resolved to accessible datasets in the next setup
↪  cell.
HOSP = os.getenv("BQ_DATASET_HOSP", "mimiciv_3_1_hosp").strip()
ICU  = os.getenv("BQ_DATASET_ICU",  "mimiciv_3_1_icu").strip()
ED   = os.getenv("BQ_DATASET_ED",    "").strip()

# BigQuery client
client = bigquery.Client(project=WORK_PROJECT)

print("Project:", WORK_PROJECT)
print("PhysioNet host:", PHYS)
print("HOSP (pref):", HOSP, "ICU (pref):", ICU, "ED (pref):", ED)
print("WORK_DIR:", WORK_DIR)
```

```
Project: mimic-hypercapnia
PhysioNet host: physionet-data
HOSP (pref): mimiciv_3_1_hosp ICU (pref): mimiciv_3_1_icu ED (pref): mimiciv_ed
WORK_DIR: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Proje
CC-NLP
```

```python
# Purpose: Define reusable BigQuery helpers and resolve dataset names across
↪  known naming variants.

from datetime import datetime, timezone

# Fail fast on long-running queries instead of hanging indefinitely in
↪  nbconvert.
BQ_QUERY_TIMEOUT_SECS = int(os.getenv("BQ_QUERY_TIMEOUT_SECS", "1800"))

# --- Helper: run SQL with optional named parameters
def run_sql_bq(sql: str, params: dict | None = None) -> pd.DataFrame:
    job_config = bigquery.QueryJobConfig()
    job_config.labels = {"pipeline": "hypercapcc", "notebook": "cohort"}
    if params:
        bq_params = []
        for k, v in params.items():
            if isinstance(v, (list, tuple, np.ndarray, pd.Series)):
                # BigQuery ARRAY<INT64> if all ints; else ARRAY<STRING>
                v_list = list(v)
                if all(isinstance(x, (int, np.integer)) for x in v_list):
                    bq_params.append(bigquery.ArrayQueryParameter(k, "INT64",
↪  list(map(int, v_list))))
                else:
                    bq_params.append(bigquery.ArrayQueryParameter(k,
↪  "STRING", list(map(str, v_list))))
            else:
                # scalar
                if isinstance(v, (int, np.integer)):
                    bq_params.append(bigquery.ScalarQueryParameter(k,
↪  "INT64", int(v)))
                elif isinstance(v, float):
                    bq_params.append(bigquery.ScalarQueryParameter(k,
↪  "FLOAT64", float(v)))
                else:
                    bq_params.append(bigquery.ScalarQueryParameter(k,
↪  "STRING", str(v)))
        job_config.query_parameters = bq_params
```

```python
    job = client.query(sql, job_config=job_config)
    try:
        result = job.result(timeout=BQ_QUERY_TIMEOUT_SECS)
    except Exception as exc:
        try:
            job.cancel()
        except Exception:
            pass
        raise RuntimeError(
            f"BigQuery query failed or timed out after
            ↪   {BQ_QUERY_TIMEOUT_SECS}s (job_id={job.job_id})."
        ) from exc

    try:
        return result.to_dataframe(create_bqstorage_client=True)
    except TypeError:
        return result.to_dataframe()

# --- Helper: test if a fully-qualified table exists and is accessible
def table_exists(fqtn: str) -> bool:
    try:
        _ = run_sql_bq(f"SELECT 1 FROM `{fqtn}` LIMIT 1")
        return True
    except Exception:
        return False

def resolve_dataset(preferred: str, candidates: list[str], probe_table: str,
↪   label: str) -> tuple[str, str]:
    # Keep preferred first, then try known aliases.
    ordered = []
    if preferred:
        ordered.append(preferred)
    for cand in candidates:
        if cand not in ordered:
            ordered.append(cand)

    for dataset in ordered:
        fqtn = f"{PHYS}.{dataset}.{probe_table}"
        if table_exists(fqtn):
            return dataset, fqtn

    raise RuntimeError(
```

```python
        f"No accessible {label} dataset found for probe table
        ↪  '{probe_table}'. Tried: {ordered}"
    )


# Resolve HOSP/ICU/ED names so notebook runs across v3.1 naming variants.
HOSP, HOSP_PROBE = resolve_dataset(
    preferred=HOSP,
    candidates=["mimiciv_3_1_hosp", "mimiciv_v3_1_hosp", "mimiciv_hosp"],
    probe_table="admissions",
    label="HOSP",
)
ICU, ICU_PROBE = resolve_dataset(
    preferred=ICU,
    candidates=["mimiciv_3_1_icu", "mimiciv_v3_1_icu", "mimiciv_icu"],
    probe_table="icustays",
    label="ICU",
)
if not ED:
    ED = "mimiciv_ed"
ED, ED_PROBE = resolve_dataset(
    preferred=ED,
    candidates=["mimiciv_ed", "mimiciv_3_1_ed", "mimiciv_v3_1_ed"],
    probe_table="edstays",
    label="ED",
)


RUN_METADATA = {
    "run_utc": datetime.now(timezone.utc).isoformat(),
    "work_project": WORK_PROJECT,
    "physionet_project": PHYS,
    "datasets": {"hosp": HOSP, "icu": ICU, "ed": ED},
    "dataset_probes": {"hosp": HOSP_PROBE, "icu": ICU_PROBE, "ed": ED_PROBE},
    "notebook": "MIMICIV_hypercap_EXT_cohort.ipynb",
}

print("Resolved datasets:", RUN_METADATA["datasets"])
print("Dataset probes:", RUN_METADATA["dataset_probes"])


Resolved datasets: {'hosp': 'mimiciv_3_1_hosp', 'icu': 'mimiciv_3_1_icu', 'ed': 'mimiciv_ed']
Dataset probes: {'hosp': 'physionet-data.mimiciv_3_1_hosp.admissions', 'icu': 'physionet-
data.mimiciv_3_1_icu.icustays', 'ed': 'physionet-data.mimiciv_ed.edstays'}
```

```python
# Purpose: Create reusable data-quality and merge guardrail helpers to
↪    prevent silent join errors.

# --- Helper utilities for reproducibility and safe joins
def require_cols(df: pd.DataFrame, cols: list[str], name: str) -> None:
    missing = [c for c in cols if c not in df.columns]
    if missing:
        raise KeyError(f"{name} missing columns: {missing}")

def assert_unique(df: pd.DataFrame, key: str, name: str) -> None:
    if df[key].duplicated().any():
        n = int(df[key].duplicated().sum())
        raise ValueError(f"{name} has {n} duplicate {key} values")

def safe_merge(left: pd.DataFrame, right: pd.DataFrame, on: list[str] | str,
↪    how: str, name: str) -> pd.DataFrame:
    # guard against accidental duplicate columns
    overlap = set(left.columns) & set(right.columns)
    if isinstance(on, str):
        on_cols = {on}
    else:
        on_cols = set(on)
    overlap = overlap - on_cols
    if overlap:
        raise ValueError(f"{name} merge would duplicate columns:
↪    {sorted(overlap)}")
    return left.merge(right, on=on, how=how)

def check_ranges(df: pd.DataFrame, ranges: dict[str, tuple[float, float]])
↪    -> pd.DataFrame:
    rows = []
    for col, (lo, hi) in ranges.items():
        if col not in df.columns:
            continue
        bad = df[col].notna() & ((df[col] < lo) | (df[col] > hi))
        rows.append({"col": col, "n_bad": int(bad.sum())})
    return pd.DataFrame(rows)
```

2.3 Data Generation

2.3.1 1) ICD cohort flags (hypercapnic respiratory failure)

Rationale: Capture diagnosis-based hypercapnia from ED and hospital discharge codes to define a broad, clinically recognized cohort.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.

# Target ICD codes (dotless, uppercase)
ICD10_CODES = ['J9602','J9612','J9622','J9692','E662']
ICD9_CODES  = ['27803']

SQL["cohort_icd_sql"] = f"""
-- ICD-based cohort flags per admission
WITH target_codes AS (
  SELECT 'J9602' AS code, 10 AS ver UNION ALL
  SELECT 'J9612', 10 UNION ALL
  SELECT 'J9622', 10 UNION ALL
  SELECT 'J9692', 10 UNION ALL
  SELECT 'E662',  10 UNION ALL
  SELECT '27803', 9
),

-- Hospital ICDs restricted to target codes
hosp_dx AS (
  SELECT
    d.subject_id,
    d.hadm_id,
    UPPER(REPLACE(d.icd_code, '.', '')) AS code_norm,
    d.icd_version
  FROM `{PHYS}.{HOSP}.diagnoses_icd` d
  JOIN target_codes t
    ON t.ver = d.icd_version AND t.code = UPPER(REPLACE(d.icd_code, '.', ''))
  WHERE d.hadm_id IS NOT NULL
),

-- Hospital flags per admission
hosp_flags AS (
  SELECT
    subject_id, hadm_id,
    MAX(IF(icd_version=10 AND code_norm='J9602',1,0)) AS ICD10_J9602,
```

```
        MAX(IF(icd_version=10 AND code_norm='J9612',1,0)) AS ICD10_J9612,
        MAX(IF(icd_version=10 AND code_norm='J9622',1,0)) AS ICD10_J9622,
        MAX(IF(icd_version=10 AND code_norm='J9692',1,0)) AS ICD10_J9692,
        MAX(IF(icd_version=10 AND code_norm='E662', 1,0)) AS ICD10_E662,
        MAX(IF(icd_version=9  AND code_norm='27803',1,0)) AS ICD9_27803
    FROM hosp_dx
    GROUP BY subject_id, hadm_id
),

-- ED ICDs restricted to target codes (map to hadm via edstays)
ed_dx AS (
    SELECT
        s.subject_id,
        s.hadm_id,
        s.stay_id,
        s.intime AS ed_intime,
        UPPER(REPLACE(d.icd_code, '.', '')) AS code_norm,
        d.icd_version
    FROM `{PHYS}.{ED}.diagnosis` d
    JOIN `{PHYS}.{ED}.edstays` s
        ON s.subject_id = d.subject_id AND s.stay_id = d.stay_id
    JOIN target_codes t
        ON t.ver = d.icd_version AND t.code = UPPER(REPLACE(d.icd_code, '.', ''))
    WHERE s.hadm_id IS NOT NULL
),

-- ED flags per ED stay (so we can both: OR flags across stays and also pick
↪  earliest stay_id)
ed_flags_by_stay AS (
    SELECT
        subject_id, hadm_id, stay_id, MIN(ed_intime) AS ed_intime,
        MAX(IF(icd_version=10 AND code_norm='J9602',1,0)) AS ICD10_J9602,
        MAX(IF(icd_version=10 AND code_norm='J9612',1,0)) AS ICD10_J9612,
        MAX(IF(icd_version=10 AND code_norm='J9622',1,0)) AS ICD10_J9622,
        MAX(IF(icd_version=10 AND code_norm='J9692',1,0)) AS ICD10_J9692,
        MAX(IF(icd_version=10 AND code_norm='E662', 1,0)) AS ICD10_E662,
        MAX(IF(icd_version=9  AND code_norm='27803',1,0)) AS ICD9_27803
    FROM ed_dx
    GROUP BY subject_id, hadm_id, stay_id
),

-- OR the ED flags across all ED stays mapped to the same hadm
ed_flags_or AS (
```

```sql
  SELECT
    subject_id, hadm_id,
    MAX(ICD10_J9602) AS ICD10_J9602,
    MAX(ICD10_J9612) AS ICD10_J9612,
    MAX(ICD10_J9622) AS ICD10_J9622,
    MAX(ICD10_J9692) AS ICD10_J9692,
    MAX(ICD10_E662 ) AS ICD10_E662,
    MAX(ICD9_27803) AS ICD9_27803
  FROM ed_flags_by_stay
  GROUP BY subject_id, hadm_id
),

-- Earliest ED stay_id per hadm (NO UNNEST of aggregates; use [OFFSET(0)])
ed_earliest AS (
  SELECT
    subject_id,
    hadm_id,
    (ARRAY_AGG(STRUCT(stay_id, ed_intime) ORDER BY ed_intime LIMIT
 ↪  1))[OFFSET(0)].stay_id AS stay_id
  FROM ed_flags_by_stay
  GROUP BY subject_id, hadm_id
),

-- Bring flags and earliest stay_id together
ed_by_hadm AS (
  SELECT
    f.subject_id,
    f.hadm_id,
    e.stay_id,
    f.ICD10_J9602,
    f.ICD10_J9612,
    f.ICD10_J9622,
    f.ICD10_J9692,
    f.ICD10_E662,
    f.ICD9_27803
  FROM ed_flags_or f
  LEFT JOIN ed_earliest e
    USING (subject_id, hadm_id)
),

-- Combine ED and hospital flags at the admission level
combined AS (
  SELECT
```

58

```
    COALESCE(h.subject_id, e.subject_id) AS subject_id,
    COALESCE(h.hadm_id,    e.hadm_id)    AS hadm_id,
    GREATEST(IFNULL(h.ICD10_J9602,0), IFNULL(e.ICD10_J9602,0)) AS
↪   ICD10_J9602,
    GREATEST(IFNULL(h.ICD10_J9612,0), IFNULL(e.ICD10_J9612,0)) AS
↪   ICD10_J9612,
    GREATEST(IFNULL(h.ICD10_J9622,0), IFNULL(e.ICD10_J9622,0)) AS
↪   ICD10_J9622,
    GREATEST(IFNULL(h.ICD10_J9692,0), IFNULL(e.ICD10_J9692,0)) AS
↪   ICD10_J9692,
    GREATEST(IFNULL(h.ICD10_E662 ,0), IFNULL(e.ICD10_E662 ,0)) AS ICD10_E662,
    GREATEST(IFNULL(h.ICD9_27803,0), IFNULL(e.ICD9_27803,0)) AS ICD9_27803,

↪   IF((IFNULL(h.ICD10_J9602,0)+IFNULL(h.ICD10_J9612,0)+IFNULL(h.ICD10_J9622,0)+IFNULL(h.ICD
↪   > 0, 1, 0) AS any_hypercap_icd_hosp,

↪   IF((IFNULL(e.ICD10_J9602,0)+IFNULL(e.ICD10_J9612,0)+IFNULL(e.ICD10_J9622,0)+IFNULL(e.ICD
↪   > 0, 1, 0) AS any_hypercap_icd_ed
  FROM hosp_flags h
  FULL OUTER JOIN ed_by_hadm e
    ON h.hadm_id = e.hadm_id
)

SELECT
  subject_id, hadm_id,
  ICD10_J9602, ICD10_J9612, ICD10_J9622, ICD10_J9692, ICD10_E662, ICD9_27803,
  IF((ICD10_J9602+ICD10_J9612+ICD10_J9622+ICD10_J9692+ICD10_E662+ICD9_27803)
↪   > 0, 1, 0) AS any_hypercap_icd,
  any_hypercap_icd_hosp,
  any_hypercap_icd_ed,
  CASE
    WHEN any_hypercap_icd_hosp=1 AND any_hypercap_icd_ed=1 THEN 'ED+HOSP'
    WHEN any_hypercap_icd_ed=1 THEN 'ED'
    WHEN any_hypercap_icd_hosp=1 THEN 'HOSP'
    ELSE 'NONE'
  END AS icd_source
FROM combined
"""

cohort_icd = run_sql_bq(sql("cohort_icd_sql"))
print("ICD cohort admissions:", len(cohort_icd))
cohort_icd.head(3)
```

```
ICD cohort admissions: 4237
```

	subject_id	hadm_id	ICD10_J9602	ICD10_J9612	ICD10_J9622	ICD10_J9692	ICD10_E662	
0	10652693	28272409	0	0	1	0	1	
1	16809525	25990450	0	0	1	0	1	
2	18662357	27604952	0	0	1	0	1	

2.3.2 2) Blood gas inclusion thresholds — ANY PaCO meeting cutoffs (LAB + POC)

Rationale: Identify physiologic hypercapnia using ABG/VBG thresholds, independent of diagnostic coding.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.

SQL["co2_thresholds_sql"] = f"""
/* ---- LAB (HOSP) pCO2 across entire dataset ---- */
WITH hosp_cand AS (
  SELECT
    le.hadm_id, le.charttime, le.specimen_id,
    COALESCE(
      CAST(le.valuenum AS FLOAT64),
      SAFE_CAST(REGEXP_EXTRACT(LOWER(le.value), r'(-?\d+(?:\.\d+)?)') AS
↪  FLOAT64)
    ) AS val,
    LOWER(REPLACE(COALESCE(le.valueuom,''),' ','')) AS uom_nospace,
    LOWER(di.label) AS lbl,
    LOWER(COALESCE(di.fluid,'')) AS fl
  FROM `{PHYS}.{HOSP}.labevents` le
  JOIN `{PHYS}.{HOSP}.d_labitems` di ON di.itemid = le.itemid
  WHERE (le.valuenum IS NOT NULL OR le.value IS NOT NULL)
    AND (
      LOWER(COALESCE(di.category,'')) LIKE '%blood gas%' OR
      LOWER(di.label) LIKE '%pco2%' OR
      REGEXP_CONTAINS(LOWER(di.label), r'\bpa?\s*co(?:2| )\b')
    )
    AND NOT REGEXP_CONTAINS(LOWER(di.label),
        r'(et\s*co2|end[- ]?tidal|t\s*co2|tco2|total\s*co2|hco3|bicar)')
),
hosp_spec AS (
  SELECT le.specimen_id, LOWER(COALESCE(le.value,'')) AS spec_val
```

```sql
  FROM `{PHYS}.{HOSP}.labevents` le
  JOIN `{PHYS}.{HOSP}.d_labitems` di ON di.itemid = le.itemid
  WHERE le.specimen_id IS NOT NULL
    AND REGEXP_CONTAINS(LOWER(di.label), r'(specimen|sample)')
),
hosp_pco2 AS (
  SELECT
    c.hadm_id, c.charttime,
    CASE
      WHEN REGEXP_CONTAINS(s.spec_val, r'arter') OR
↪ REGEXP_CONTAINS(s.spec_val, r'\bart\b') THEN 'arterial'
      WHEN REGEXP_CONTAINS(s.spec_val, r'ven|mixed|central') THEN 'venous'
      WHEN REGEXP_CONTAINS(c.fl, r'arter') THEN 'arterial'
      WHEN REGEXP_CONTAINS(c.fl, r'ven') THEN 'venous'
      WHEN c.fl LIKE '%arterial%' OR REGEXP_CONTAINS(c.lbl,
↪ r'\b(abg|art|arterial|a[- ]?line)\b') THEN 'arterial'
      WHEN c.fl LIKE '%ven%'      OR REGEXP_CONTAINS(c.lbl,
↪ r'\b(vbg|ven|venous|mixed|central)\b') THEN 'venous'
      ELSE 'other'
    END AS site,
    CASE WHEN c.uom_nospace='kpa' THEN c.val*7.50062 ELSE c.val END AS
↪ pco2_mmHg
  FROM hosp_cand c
  LEFT JOIN hosp_spec s USING (specimen_id)
  WHERE c.val IS NOT NULL
),
hosp_pco2_std AS (
  SELECT hadm_id, site, charttime, pco2_mmHg
  FROM hosp_pco2
  WHERE site IN ('arterial','venous','other') AND pco2_mmHg BETWEEN 5 AND 200
),

/* ---- ICU (POC) pCO2 across entire dataset ---- */
icu_pco2_items AS (
  SELECT
    itemid,
    LOWER(label) AS lbl
  FROM `{PHYS}.{ICU}.d_items`
  WHERE REGEXP_CONTAINS(LOWER(label),
↪ r'(^|[^a-z])p(?:a)?\s*co\s*(?:2| )([^a-z]|$)')
    AND NOT REGEXP_CONTAINS(
      LOWER(label),
```

```
        r'((^|[^a-z])p(?:a)?\s*o\s*2([^a-z]|$)|\bpo2\b|\bpao2\b|oxygen|o2\
↪   s*(sat|saturation)|sao2|spo2|fio2|et\s*co2|end[-
↪   ]?tidal|t\s*co2|tco2|total\s*co2|hco3|bicar|v\s*co2|vco2|co2\
↪   s*(prod|elimin|production|elimination))'
      )
),
icu_raw AS (
  SELECT
    ie.hadm_id,
    ce.stay_id,
    ce.charttime,
    di.lbl,
    LOWER(REPLACE(COALESCE(ce.valueuom,''),' ','')) AS uom_nospace,
    COALESCE(
      CAST(ce.valuenum AS FLOAT64),
      SAFE_CAST(REGEXP_EXTRACT(LOWER(ce.value), r'(-?\d+(?:\.\d+)?)') AS
↪   FLOAT64)
    ) AS val
  FROM `{PHYS}.{ICU}.chartevents` ce
  JOIN icu_pco2_items di ON di.itemid = ce.itemid
  JOIN `{PHYS}.{ICU}.icustays` ie ON ie.stay_id = ce.stay_id
),
icu_cand AS (
  SELECT
    hadm_id, stay_id, charttime, lbl, uom_nospace, val,
    'pco2' AS analyte,
    CASE
      WHEN REGEXP_CONTAINS(lbl, r'\b(abg|art|arterial|a[- ]?line)\b') THEN
↪   'arterial'
      WHEN REGEXP_CONTAINS(lbl, r'\b(vbg|ven|venous|mixed|central)\b') THEN
↪   'venous'
      ELSE 'other'
    END AS site
  FROM icu_raw
  WHERE val IS NOT NULL
),
icu_co2_std AS (
  SELECT
    hadm_id,
    site,
    charttime,
    CASE WHEN uom_nospace='kpa' THEN val*7.50062 ELSE val END AS pco2_mmHg
  FROM icu_cand
```

```
    WHERE analyte='pco2'
      AND site IN ('arterial','venous','other')
      AND (CASE WHEN uom_nospace='kpa' THEN val*7.50062 ELSE val END) BETWEEN 5
  ↪   AND 200
),

/* ---- Combine and threshold per admission ---- */
all_pco2 AS (
  SELECT * FROM hosp_pco2_std
  UNION ALL
  SELECT * FROM icu_co2_std
),
thresh AS (
  SELECT
    hadm_id,
    MAX(IF(site='arterial' AND pco2_mmHg >= 45.0, 1, 0)) AS
  ↪ abg_hypercap_threshold,
    MAX(IF(site='venous'   AND pco2_mmHg >= 50.0, 1, 0)) AS
  ↪ vbg_hypercap_threshold,
    MAX(IF(site='other'    AND pco2_mmHg >= 50.0, 1, 0)) AS
  ↪ other_hypercap_threshold
  FROM all_pco2
  GROUP BY hadm_id
)
SELECT * FROM thresh
"""

co2_thresh = run_sql_bq(sql("co2_thresholds_sql"))
print("Admissions meeting thresholds:", len(co2_thresh))
co2_thresh.head(3)


Admissions meeting thresholds: 98693
```

	hadm_id	abg_hypercap_threshold	vbg_hypercap_threshold	other_hypercap_threshold
0	26483561	0	0	1
1	20423793	0	0	1
2	28437846	0	0	1

2.3.3 3) Cohort union (ICD thresholds) and `hadm_list` for downstream queries

Rationale: Combine ICD and physiologic routes to maximize sensitivity, then define a stable hadm_id list for downstream joins.

```
# Purpose: Combine ICD and gas-threshold ascertainment routes to produce the
↪   final hadm inclusion list.

# Outer-join because thresholds can identify hadm_id with no ICD codes and
↪   vice versa
cohort_any = cohort_icd.merge(co2_thresh, how="outer", on="hadm_id")

# Fill missing flags with 0 where appropriate
icd_cols = ["ICD10_J9602","ICD10_J9612","ICD10_J9622","ICD10_J9692",⌋
↪   "ICD10_E662","ICD9_27803","any_hypercap_icd","any_hypercap_icd_hosp",⌋
↪   "any_hypercap_icd_ed"]
for c in icd_cols:
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].fillna(0).astype(int)

for c in ["abg_hypercap_threshold","vbg_hypercap_threshold",⌋
↪   "other_hypercap_threshold"]:
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].fillna(0).astype(int)

# Final enrollment flag
cohort_any["pco2_threshold_any"] = ((cohort_any["abg_hypercap_threshold"]==1)
↪   | (cohort_any["vbg_hypercap_threshold"]==1) |
↪   (cohort_any["other_hypercap_threshold"]==1)).astype(int)
cohort_any["enrolled_any"] = ((cohort_any["any_hypercap_icd"]==1) |
↪   (cohort_any["pco2_threshold_any"]==1)).astype(int)

print("ICD-only admissions        :",
↪   int((cohort_any["any_hypercap_icd"]==1).sum()))
print("Threshold-only admissions  :",
↪   int(((cohort_any["pco2_threshold_any"]==1) &
↪   (cohort_any["any_hypercap_icd"]==0)).sum()))
print("Both ICD and threshold      :",
↪   int(((cohort_any["pco2_threshold_any"]==1) &
↪   (cohort_any["any_hypercap_icd"]==1)).sum()))
print("Total enrolled (union)      :",
↪   int((cohort_any["enrolled_any"]==1).sum()))
```

```
# New hadm list used for the rest of the notebook
hadm_list = cohort_any.loc[cohort_any["enrolled_any"]==1,
↪  "hadm_id"].dropna().astype("int64").tolist()
len(hadm_list)
```

```
ICD-only admissions        : 4237
Threshold-only admissions  : 80358
Both ICD and threshold     : 3549
Total enrolled (union)     : 84595
```

84594

2.3.4 4) First ABG and First VBG (LAB + POC, standardized to mmHg)

Rationale: Extract earliest ABG/VBG measurements to characterize baseline physiology with standardized units.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.

params = {"hadms": hadm_list}

bg_pairs_sql = rf"""
WITH hadms AS (SELECT hadm_id FROM UNNEST(@hadms) AS hadm_id),

/* --------------- LAB (HOSP) --------------- */
hosp_cand AS (
  SELECT
    le.subject_id, le.hadm_id, le.charttime, le.specimen_id,
    CAST(le.valuenum AS FLOAT64) AS val,
    LOWER(COALESCE(le.valueuom,'')) AS uom,
    LOWER(di.label) AS lbl,
    LOWER(COALESCE(di.fluid,'')) AS fl,
    LOWER(COALESCE(di.category,'')) AS cat
  FROM `{PHYS}.{HOSP}.labevents`  le
  JOIN `{PHYS}.{HOSP}.d_labitems` di ON di.itemid = le.itemid
  JOIN hadms h ON h.hadm_id = le.hadm_id
  WHERE le.valuenum IS NOT NULL
    AND (
        LOWER(COALESCE(di.category,'')) LIKE '%blood gas%' OR
        LOWER(di.label) LIKE '%pco2%' OR
```

```
          REGEXP_CONTAINS(LOWER(di.label), r'\bph\b') OR
          REGEXP_CONTAINS(LOWER(di.label), r'\bpa?\s*co(?:2| )\b')
        )
    AND NOT REGEXP_CONTAINS(LOWER(di.label), r'(et\s*co2|end[-
↪  ]?tidal|t\s*co2|tco2|total\s*co2|hco3|bicar)')
),
hosp_spec AS (
  SELECT le.specimen_id, LOWER(COALESCE(le.value,'')) AS spec_val
  FROM `{PHYS}.{HOSP}.labevents` le
  JOIN `{PHYS}.{HOSP}.d_labitems` di ON di.itemid = le.itemid
  WHERE le.specimen_id IS NOT NULL
    AND REGEXP_CONTAINS(LOWER(di.label), r'(specimen|sample)')
),
hosp_class AS (
  SELECT
    c.hadm_id, c.charttime, c.specimen_id, c.val, c.uom, c.lbl, c.fl,
    CASE
      WHEN REGEXP_CONTAINS(c.lbl, r'\b(?:blood\s*)?ph\b') THEN 'ph'
      WHEN (c.lbl LIKE '%pco2%' OR REGEXP_CONTAINS(c.lbl,
↪  r'\bpa?\s*co(?:2| )\b')) THEN 'pco2'
      ELSE NULL
    END AS analyte,
    CASE
      WHEN REGEXP_CONTAINS(s.spec_val, r'arter') OR
↪  REGEXP_CONTAINS(s.spec_val, r'\bart\b') THEN 'arterial'
      WHEN REGEXP_CONTAINS(s.spec_val, r'ven|mixed|central') THEN 'venous'
      WHEN c.fl LIKE '%arterial%' OR REGEXP_CONTAINS(c.lbl,
↪  r'\b(abg|art|arterial|a[- ]?line)\b') THEN 'arterial'
      WHEN c.fl LIKE '%ven%'      OR REGEXP_CONTAINS(c.lbl,
↪  r'\b(vbg|ven|venous|mixed|central)\b') THEN 'venous'
      ELSE 'other'
    END AS site
  FROM hosp_cand c
  LEFT JOIN hosp_spec s USING (specimen_id)
),
hosp_pairs AS (
  SELECT
    hadm_id, specimen_id,
    MIN(charttime) AS sample_time,
    MAX(IF(analyte='ph',   val, NULL)) AS ph,
    MAX(IF(analyte='pco2', val, NULL)) AS pco2_raw,
    (ARRAY_AGG(IF(analyte='pco2', uom, NULL) IGNORE NULLS LIMIT
↪  1))[OFFSET(0)] AS pco2_uom,
```

```
      (ARRAY_AGG(IF(analyte='ph',   uom, NULL) IGNORE NULLS LIMIT
↪  1))[OFFSET(0)] AS ph_uom,
      (ARRAY_AGG(site IGNORE NULLS LIMIT 1))[OFFSET(0)] AS site
  FROM hosp_class
  GROUP BY hadm_id, specimen_id
  HAVING (ph IS NOT NULL OR pco2_raw IS NOT NULL) AND site IN
↪  ('arterial','venous','other')
),
hosp_pairs_std AS (
  SELECT
    hadm_id, specimen_id, sample_time, site,
    ph, ph_uom,
    CASE WHEN pco2_uom = 'kpa' THEN pco2_raw * 7.50062 ELSE pco2_raw END AS
↪  pco2_mmHg,
    'mmhg' AS pco2_uom_norm
  FROM hosp_pairs
  WHERE (ph IS NULL OR (ph BETWEEN 6.3 AND 7.8))
    AND (pco2_raw IS NULL OR (CASE WHEN pco2_uom='kpa' THEN pco2_raw*7.50062
↪  ELSE pco2_raw END) BETWEEN 5 AND 200)
),
lab_abg AS (
  SELECT hadm_id,
         ph             AS lab_abg_ph,
         ph_uom         AS lab_abg_ph_uom,
         pco2_mmHg      AS lab_abg_paco2,
         'mmhg'         AS lab_abg_paco2_uom,
         sample_time    AS lab_abg_time
  FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
↪  sample_time) rn
       FROM hosp_pairs_std WHERE site='arterial') WHERE rn=1
),
lab_vbg AS (
  SELECT hadm_id,
         ph             AS lab_vbg_ph,
         ph_uom         AS lab_vbg_ph_uom,
         pco2_mmHg      AS lab_vbg_paco2,
         'mmhg'         AS lab_vbg_paco2_uom,
         sample_time    AS lab_vbg_time
  FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
↪  sample_time) rn
       FROM hosp_pairs_std WHERE site='venous') WHERE rn=1
),
lab_other AS (
```

```sql
    SELECT hadm_id,
           ph              AS lab_other_ph,
           ph_uom          AS lab_other_ph_uom,
           pco2_mmHg       AS lab_other_paco2,
           'mmhg'          AS lab_other_paco2_uom,
           sample_time     AS lab_other_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
↪   sample_time) rn
         FROM hosp_pairs_std WHERE site='other') WHERE rn=1
),

/* --------------- POC (ICU) --------------- */
icu_item_map AS (
  SELECT
    itemid,
    LOWER(label) AS lbl,
    CASE
      WHEN REGEXP_CONTAINS(LOWER(label), r'(^|[^a-z])ph([^a-z]|$)') THEN 'ph'
      WHEN REGEXP_CONTAINS(LOWER(label),
↪   r'(^|[^a-z])p(?:a)?\s*co\s*(?:2| )([^a-z]|$)')
           AND NOT REGEXP_CONTAINS(
             LOWER(label),
             r'((^|[^a-z])p(?:a)?\s*o\s*2([^a-z]|$)|\bpo2\b|\bpao2\b|oxygen|⌋
↪   o2\s*(sat|saturation)|sao2|spo2|fio2|et\s*co2|end[-
↪   ]?tidal|t\s*co2|tco2|total\s*co2|hco3|bicar|v\s*co2|vco2|co2\s*(prod|⌋
↪   elimin|production|elimination))'
           )
      THEN 'pco2'
      ELSE NULL
    END AS analyte
  FROM `{PHYS}.{ICU}.d_items`
  WHERE
    REGEXP_CONTAINS(LOWER(label), r'(^|[^a-z])ph([^a-z]|$)')
    OR (
      REGEXP_CONTAINS(LOWER(label),
↪   r'(^|[^a-z])p(?:a)?\s*co\s*(?:2| )([^a-z]|$)')
      AND NOT REGEXP_CONTAINS(
        LOWER(label),
        r'((^|[^a-z])p(?:a)?\s*o\s*2([^a-z]|$)|\bpo2\b|\bpao2\b|oxygen|o2\s*⌋
↪   (sat|saturation)|sao2|spo2|fio2|et\s*co2|end[-
↪   ]?tidal|t\s*co2|tco2|total\s*co2|hco3|bicar|v\s*co2|vco2|co2\s*(prod|⌋
↪   elimin|production|elimination))'
      )
    )
```

```
    )
),
icu_raw AS (
  SELECT
    ie.hadm_id, ce.stay_id, ce.itemid, ce.charttime,
    di.lbl,
    di.analyte,
    LOWER(REPLACE(COALESCE(ce.valueuom,''),' ','')) AS uom_nospace,
    LOWER(COALESCE(ce.value,'')) AS valstr,
    COALESCE(
      CAST(ce.valuenum AS FLOAT64),
      SAFE_CAST(REGEXP_EXTRACT(LOWER(ce.value), r'(-?\d+(?:\.\d+)?)') AS
↪  FLOAT64)
    ) AS val
  FROM `{PHYS}.{ICU}.chartevents` ce
  JOIN icu_item_map di ON di.itemid = ce.itemid
  JOIN `{PHYS}.{ICU}.icustays` ie ON ie.stay_id = ce.stay_id
  JOIN hadms h ON h.hadm_id = ie.hadm_id
),
icu_cand AS (
  SELECT
    hadm_id, stay_id, itemid, charttime, lbl, analyte, uom_nospace, valstr,
↪  val,
    CASE
      WHEN REGEXP_CONTAINS(lbl, r'\b(abg|art|arterial|a[- ]?line)\b') THEN
↪  'arterial'
      WHEN REGEXP_CONTAINS(lbl, r'\b(vbg|ven|venous|mixed|central)\b') THEN
↪  'venous'
      ELSE 'other'
    END AS site
  FROM icu_raw
  WHERE val IS NOT NULL
    AND analyte IN ('ph', 'pco2')
),
icu_ph AS (
  SELECT hadm_id, stay_id, charttime, val AS ph, site AS site_ph
  FROM icu_cand WHERE analyte='ph'
),
icu_co2 AS (
  SELECT hadm_id, stay_id, charttime, val AS pco2_raw, uom_nospace, valstr,
↪  site AS site_co2
  FROM icu_cand WHERE analyte='pco2'
),
```

```
icu_pair_win AS (
  SELECT
    p.hadm_id, p.stay_id,
    COALESCE(p.site_ph, c.site_co2) AS site,
    p.charttime AS ph_time, c.charttime AS co2_time,
    p.ph,
    CASE
      WHEN c.uom_nospace='kpa' THEN 'kpa'
      WHEN c.uom_nospace='mmhg' THEN 'mmhg'
      ELSE c.uom_nospace
    END AS pco2_uom_norm_raw,
    c.pco2_raw,
    ABS(TIMESTAMP_DIFF(c.charttime, p.charttime, SECOND)) AS dt_sec
  FROM icu_ph p
  JOIN icu_co2 c
    ON c.hadm_id = p.hadm_id
   AND c.stay_id = p.stay_id
   AND (COALESCE(p.site_ph, c.site_co2) IN ('arterial','venous','other'))
   AND ABS(TIMESTAMP_DIFF(c.charttime, p.charttime, MINUTE)) <= 10
  QUALIFY ROW_NUMBER() OVER (
    PARTITION BY p.hadm_id, p.stay_id, p.charttime
    ORDER BY dt_sec
  ) = 1
),
icu_pairs_std AS (
  SELECT
    hadm_id, stay_id, site,
    LEAST(ph_time, co2_time) AS sample_time,
    ph,
    CAST(NULL AS STRING) AS ph_uom,              -- POC pH is unitless/null
    CASE WHEN pco2_uom_norm_raw='kpa' THEN pco2_raw*7.50062 ELSE pco2_raw END
↪  AS pco2_mmHg,
    'mmhg' AS pco2_uom_norm
  FROM icu_pair_win
  WHERE (ph BETWEEN 6.3 AND 7.8 OR ph IS NULL)
    AND (CASE WHEN pco2_uom_norm_raw='kpa' THEN pco2_raw*7.50062 ELSE
↪  pco2_raw END) BETWEEN 5 AND 200
),
icu_solo_pco2_std AS (
  SELECT
    hadm_id, stay_id, site_co2 AS site,
    charttime AS sample_time,
    CAST(NULL AS FLOAT64) AS ph,
```

```
        CAST(NULL AS STRING)  AS ph_uom,            -- no pH here
        CASE WHEN uom_nospace='kpa' THEN pco2_raw*7.50062 ELSE pco2_raw END AS
  ↪  pco2_mmHg,
        'mmhg' AS pco2_uom_norm
  FROM icu_co2
  WHERE site_co2 IN ('arterial','venous','other')
    AND pco2_raw BETWEEN 5 AND 200
),
icu_all AS (
  SELECT * FROM icu_pairs_std
  UNION ALL
  SELECT * FROM icu_solo_pco2_std
),

poc_abg AS (
  SELECT hadm_id,
         ph              AS poc_abg_ph,
         ph_uom          AS poc_abg_ph_uom,
         pco2_mmHg       AS poc_abg_paco2,
         'mmhg'          AS poc_abg_paco2_uom,
         sample_time    AS poc_abg_time
  FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
  ↪  sample_time) rn
        FROM icu_all WHERE site='arterial') WHERE rn=1
),
poc_vbg AS (
  SELECT hadm_id,
         ph              AS poc_vbg_ph,
         ph_uom          AS poc_vbg_ph_uom,
         pco2_mmHg       AS poc_vbg_paco2,
         'mmhg'          AS poc_vbg_paco2_uom,
         sample_time    AS poc_vbg_time
  FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
  ↪  sample_time) rn
        FROM icu_all WHERE site='venous') WHERE rn=1
),
poc_other AS (
  SELECT hadm_id,
         ph              AS poc_other_ph,
         ph_uom          AS poc_other_ph_uom,
         pco2_mmHg       AS poc_other_paco2,
         'mmhg'          AS poc_other_paco2_uom,
         sample_time    AS poc_other_time
```

```
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
↪   sample_time) rn
        FROM icu_all WHERE site='other') WHERE rn=1
)

/* --------------- Final one row per hadm --------------- */
SELECT
  h.hadm_id,
  -- LAB-ABG / LAB-VBG / LAB-OTHER
  la.lab_abg_ph, la.lab_abg_ph_uom, la.lab_abg_paco2, la.lab_abg_paco2_uom,
↪   la.lab_abg_time,
  lv.lab_vbg_ph, lv.lab_vbg_ph_uom, lv.lab_vbg_paco2, lv.lab_vbg_paco2_uom,
↪   lv.lab_vbg_time,
  lo.lab_other_ph, lo.lab_other_ph_uom, lo.lab_other_paco2,
↪   lo.lab_other_paco2_uom, lo.lab_other_time,
  -- POC-ABG / POC-VBG / POC-OTHER
  pa.poc_abg_ph, pa.poc_abg_ph_uom, pa.poc_abg_paco2, pa.poc_abg_paco2_uom,
↪   pa.poc_abg_time,
  pv.poc_vbg_ph, pv.poc_vbg_ph_uom, pv.poc_vbg_paco2, pv.poc_vbg_paco2_uom,
↪   pv.poc_vbg_time,
  po.poc_other_ph, po.poc_other_ph_uom, po.poc_other_paco2,
↪   po.poc_other_paco2_uom, po.poc_other_time,
  -- First ABG across LAB+POC
  (SELECT AS STRUCT src, t, ph, pco2
   FROM (SELECT 'LAB' AS src, la.lab_abg_time AS t, la.lab_abg_ph AS ph,
↪   la.lab_abg_paco2 AS pco2
        UNION ALL
        SELECT 'POC', pa.poc_abg_time, pa.poc_abg_ph, pa.poc_abg_paco2)
   WHERE t IS NOT NULL
   ORDER BY t LIMIT 1) AS first_abg,
  -- First VBG across LAB+POC
  (SELECT AS STRUCT src, t, ph, pco2
   FROM (SELECT 'LAB' AS src, lv.lab_vbg_time AS t, lv.lab_vbg_ph AS ph,
↪   lv.lab_vbg_paco2 AS pco2
        UNION ALL
        SELECT 'POC', pv.poc_vbg_time, pv.poc_vbg_ph, pv.poc_vbg_paco2)
   WHERE t IS NOT NULL
   ORDER BY t LIMIT 1) AS first_vbg,
  -- First OTHER-source pCO2 across LAB+POC
  (SELECT AS STRUCT src, t, ph, pco2
   FROM (SELECT 'LAB' AS src, lo.lab_other_time AS t, lo.lab_other_ph AS ph,
↪   lo.lab_other_paco2 AS pco2
        UNION ALL
```

```
        SELECT 'POC', po.poc_other_time, po.poc_other_ph,
↪   po.poc_other_paco2)
    WHERE t IS NOT NULL
    ORDER BY t LIMIT 1) AS first_other
FROM hadms h
LEFT JOIN lab_abg la USING (hadm_id)
LEFT JOIN lab_vbg lv USING (hadm_id)
LEFT JOIN lab_other lo USING (hadm_id)
LEFT JOIN poc_abg pa USING (hadm_id)
LEFT JOIN poc_vbg pv USING (hadm_id)
LEFT JOIN poc_other po USING (hadm_id)
"""

bg_pairs = run_sql_bq(bg_pairs_sql, params)

# Flatten STRUCTs for first_abg, first_vbg, and first_other
for col in ["first_abg","first_vbg","first_other"]:
    if col in bg_pairs.columns:
        bg_pairs[f"{col}_src"]  = bg_pairs[col].apply(lambda x: x.get("src")
↪   if isinstance(x, dict) else None)
        bg_pairs[f"{col}_time"] = bg_pairs[col].apply(lambda x: x.get("t")
↪   if isinstance(x, dict) else None)
        bg_pairs[f"{col}_ph"]   = bg_pairs[col].apply(lambda x: x.get("ph")
↪   if isinstance(x, dict) else None)
        bg_pairs[f"{col}_pco2"] = bg_pairs[col].apply(lambda x:
↪   x.get("pco2")if isinstance(x, dict) else None)
        bg_pairs = bg_pairs.drop(columns=[col])

bg_pairs.head(3)
```

	hadm_id	lab_abg_ph	lab_abg_ph_uom	lab_abg_paco2	lab_abg_paco2_uom	lab_abg_time
0	20000094	NaN	None	NaN	None	NaT
1	20000147	7.41	units	35.0	mmhg	2121-08-30 17:38
2	20000235	7.45	units	43.0	mmhg	2139-11-27 11:46

2.3.5 5) Demographics & outcomes

Rationale: Build baseline covariates used for descriptive statistics and potential confounding adjustment.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪   consistent across notebook stages.
```

```python
SQL["demo_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
  a.hadm_id,
  a.subject_id,
  a.admittime,
  a.dischtime,
  a.deathtime,
  a.admission_type,
  a.admission_location,
  a.discharge_location,
  a.insurance,
  -- LOS (days)
  TIMESTAMP_DIFF(a.dischtime, a.admittime, HOUR) / 24.0 AS hosp_los_days,
  -- in-hospital death
  IF(a.deathtime IS NOT NULL, 1, 0) AS death_in_hosp,
  -- demographics
  p.gender,
  SAFE_CAST(ROUND(p.anchor_age + (EXTRACT(YEAR FROM a.admittime) -
↪  p.anchor_year), 1) AS FLOAT64) AS age_at_admit,
  -- 30-day all-cause mortality from admission
  IF(p.dod IS NOT NULL AND DATE_DIFF(DATE(p.dod), DATE(a.admittime), DAY)
↪  BETWEEN 0 AND 30, 1, 0) AS death_30d
FROM `{PHYS}.{HOSP}.admissions` a
JOIN hadms h USING (hadm_id)
JOIN `{PHYS}.{HOSP}.patients` p USING (subject_id)
"""
demo = run_sql_bq(sql("demo_sql"), {"hadms": hadm_list})
print("Demo rows:", len(demo))
demo.head(3)
```

```
Demo rows: 84594
```

	hadm_id	subject_id	admittime	dischtime	deathtime	admission_type
0	23485217	10584718	2165-02-12 15:41:00	2165-03-06 08:20:00	2165-03-06 08:20:00	EW EMER.
1	27004173	10297948	2135-08-23 18:46:00	2135-09-08 15:50:00	NaT	URGENT
2	22661627	10032409	2130-01-12 18:42:00	2130-01-21 14:32:00	NaT	EW EMER.

```python
# Purpose: Create reusable data-quality and merge guardrail helpers to
↪  prevent silent join errors.
```

```python
# ==== Drop-in: safe merge utilities (one cell, run once) ====
import pandas as pd
from typing import Iterable, Optional, Literal

def _ensure_Int64(s: pd.Series) -> pd.Series:
    """Coerce to pandas nullable Int64 (preserves NA)."""
    return pd.to_numeric(s, errors="coerce").astype("Int64")

def strip_subject_cols(fr: pd.DataFrame) -> pd.DataFrame:
    """Remove any subject_id-like columns from a frame (e.g., 'subject_id',
    ↪  'Subject_ID')."""
    return fr.drop(columns=[c for c in fr.columns if
    ↪  c.lower().startswith("subject_id")],
                  errors="ignore")

def safe_merge_on_hadm(
    left: pd.DataFrame,
    right: pd.DataFrame,
    *,
    right_name: str,
    take: Optional[Iterable[str]] = None,
    order_by: Optional[Iterable[str]] = None,
    check_subject: Literal[False, "warn", "raise"] = False,
) -> pd.DataFrame:
    """
    Left-merge 'right' into 'left' on hadm_id, returning a copy of left with
    ↪  right's columns.
    - Dedupes right on hadm_id (optionally using order_by to pick the first
    ↪  row).
    - Optionally restricts right columns via `take`.
    - Optionally audits subject_id agreement before dropping subject_id from
    ↪  right.
    - Always strips subject_id-like columns from the right to prevent *_x/_y
    ↪  suffixes.
    - Raises if any *_x/_y suffixes still appear (indicates overlapping names
    ↪  besides hadm_id).
    """
    if "hadm_id" not in left.columns:
        raise KeyError(f"left frame lacks hadm_id before merging
        ↪  {right_name}")
    if "hadm_id" not in right.columns:
        raise KeyError(f"{right_name} lacks hadm_id")
```

```python
    L = left.copy()
    R = right.copy()

    # Standardize dtypes of keys
    L["hadm_id"] = _ensure_Int64(L["hadm_id"])
    R["hadm_id"] = _ensure_Int64(R["hadm_id"])
    if "subject_id" in L.columns:
        L["subject_id"] = _ensure_Int64(L["subject_id"])
    if "subject_id" in R.columns:
        R["subject_id"] = _ensure_Int64(R["subject_id"])

    # Dedupe RIGHT by hadm_id (optionally order_by first)
    if order_by:
        R = (R.sort_values(list(order_by))
                .drop_duplicates(subset=["hadm_id"], keep="first"))
    else:
        R = R.drop_duplicates(subset=["hadm_id"], keep="first")

    # Optional subject_id consistency audit (before stripping)
    if check_subject and ("subject_id" in L.columns) and ("subject_id" in
    ↪  R.columns):
        # Join only on hadm_id where both sides have subject_id
        tmp = (L[["hadm_id", "subject_id"]]
                .merge(R[["hadm_id", "subject_id"]],
                        on="hadm_id", how="inner", suffixes=("_L","_R")))
        mism = (tmp["subject_id_L"].notna() & tmp["subject_id_R"].notna() &
                (tmp["subject_id_L"] != tmp["subject_id_R"]))
        n_mism = int(mism.sum())
        if n_mism > 0:
            sample_ids = tmp.loc[mism, "hadm_id"].head(10).tolist()
            msg = (f"[{right_name}] subject_id mismatch on {n_mism}
    ↪ hadm_id(s). "
                    f"Examples: {sample_ids}")
            if check_subject == "raise":
                raise ValueError(msg)
            else:
                print("WARNING:", msg)

    # Limit right columns (avoid accidental overlaps)
    if take is not None:
        keep = ["hadm_id"] + [c for c in take if c != "hadm_id"]
        R = R[[c for c in keep if c in R.columns]]
```

```python
    # Always strip subject_id-like columns from right to prevent *_x/_y
    R = strip_subject_cols(R)

    # Final merge
    out = L.merge(R, on="hadm_id", how="left", suffixes=("", ""))

    # Guard: no suffixes should be present
    bad = [c for c in out.columns if c.endswith("_x") or c.endswith("_y")]
    if bad:
        raise RuntimeError(
            f"Merge with {right_name} produced suffixed columns {bad}. "
            "You likely have overlapping column names other than hadm_id."
        )
    return out

print("Safe merge helpers loaded.")
```

```
Safe merge helpers loaded.
```

2.3.6 6) NIH/OMB race & ethnicity (ED + Hospital)

Rationale: Harmonize race/ethnicity across sources using NIH/OMB categories for consistent reporting.

```python
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.

race_eth_sql = rf"""
WITH hadms AS (
  SELECT x AS hadm_id
  FROM UNNEST(@hadms) AS x
),

-- Hospital admission "race" text
hosp AS (
  SELECT a.hadm_id, LOWER(TRIM(a.race)) AS race_hosp_raw
  FROM `{PHYS}.{HOSP}.admissions` a
  JOIN hadms hm USING (hadm_id)
),

-- Earliest ED stay leading to the admission; take its "race" text if present
ed_first AS (
```

77

```
    SELECT
      e.hadm_id,
      (ARRAY_AGG(STRUCT(e.intime AS intime, LOWER(TRIM(e.race)) AS race_ed_raw)
                 ORDER BY e.intime ASC LIMIT 1))[OFFSET(0)] AS pick
    FROM `{PHYS}.{ED}.edstays` e
    JOIN hadms hm USING (hadm_id)
    GROUP BY e.hadm_id
),
ed AS (
  SELECT hadm_id, pick.race_ed_raw
  FROM ed_first
),

-- Combine ED + Hospital for maximum coverage
comb AS (
  SELECT
    hm.hadm_id,
    ho.race_hosp_raw,
    ed.race_ed_raw,
    TRIM(REGEXP_REPLACE(CONCAT(COALESCE(ho.race_hosp_raw,''), ' ',
  ↪ COALESCE(ed.race_ed_raw,'')), r'\s+', ' ')) AS race_text_any
  FROM hadms hm
  LEFT JOIN hosp ho USING (hadm_id)
  LEFT JOIN ed   ed USING (hadm_id)
),

-- Tokenization to OMB families + Hispanic ethnicity
tok AS (
  SELECT
    hadm_id, race_hosp_raw, race_ed_raw, race_text_any,

    -- Ethnicity (Hispanic)
    REGEXP_CONTAINS(race_text_any, r'\b(hispanic|latinx|latino|latina)\b') AS
  ↪ is_hisp,

    -- Race families (use boundaries to reduce false positives)
    REGEXP_CONTAINS(race_text_any, r'american\s+indian|\balaska\b') AS
  ↪ is_aian,
    REGEXP_CONTAINS(race_text_any, r'\basian\b') AS is_asian,
    REGEXP_CONTAINS(race_text_any, r'\b(black|african\s+american)\b') AS
  ↪ is_black,
    REGEXP_CONTAINS(race_text_any, r'hawaiian|pacific\s+islander') AS
  ↪ is_nhopi,
```

```
      REGEXP_CONTAINS(race_text_any, r'\bwhite\b|caucasian') AS is_white,

      -- Unknown/other indicators
      REGEXP_CONTAINS(race_text_any,
↪  r'unknown|other|declined|unable|not\s+reported|missing|null') AS
↪  is_unknown_any,

      -- Multi-race hints
      REGEXP_CONTAINS(race_text_any,
↪  r'(two|2)\s+or\s+more|multi|biracial|multiracial') AS is_multi_hint
    FROM comb
),

-- Decide ethnicity per NIH
ethn AS (
  SELECT
    hadm_id, race_hosp_raw, race_ed_raw, race_text_any,
    CASE
      WHEN is_hisp THEN 'Hispanic or Latino'
      WHEN (race_text_any IS NULL OR race_text_any = '' OR is_unknown_any)
↪  THEN 'Unknown or Not Reported'
      ELSE 'Not Hispanic or Latino'
    END AS nih_ethnicity,
    (CAST(is_aian AS INT64) + CAST(is_asian AS INT64) + CAST(is_black AS
↪  INT64)
     + CAST(is_nhopi AS INT64) + CAST(is_white AS INT64)) AS race_hits,
    is_aian, is_asian, is_black, is_nhopi, is_white, is_multi_hint,
↪  is_unknown_any
  FROM tok
),

-- Decide race per NIH/OMB (1997)
race_assign AS (
  SELECT
    hadm_id, race_hosp_raw, race_ed_raw, race_text_any, nih_ethnicity,
    CASE
      WHEN race_hits >= 2 OR is_multi_hint THEN 'More than one race'
      WHEN is_aian THEN 'American Indian or Alaska Native'
      WHEN is_asian THEN 'Asian'
      WHEN is_black THEN 'Black or African American'
      WHEN is_nhopi THEN 'Native Hawaiian or Other Pacific Islander'
      WHEN is_white THEN 'White'
```

```
        WHEN is_unknown_any OR race_text_any IS NULL OR race_text_any = '' THEN
↪  'Unknown or Not Reported'
        ELSE 'Unknown or Not Reported'
    END AS nih_race
  FROM ethn
)

SELECT hadm_id, race_hosp_raw, race_ed_raw, nih_race, nih_ethnicity
FROM race_assign
"""
race_eth = run_sql_bq(race_eth_sql, {"hadms": hadm_list})
print("Race/Eth rows:", len(race_eth))
race_eth.head(3)
```

```
Race/Eth rows: 84594
```

	hadm_id	race_hosp_raw	race_ed_raw	nih_race	nih_ethnicity
0	20000094	white	None	White	Not Hispanic or Latino
1	20000147	white - other european	None	White	Unknown or Not Reported
2	20000235	white	None	White	Not Hispanic or Latino

2.3.7 7) ED triage (linked to hadm) and first ED vitals

Rationale: Capture ED presentation features (vitals and chief complaint) for symptom and severity analyses.

```
# Purpose: Pull ED triage/first-vitals by loading ED tables once and
↪  filtering to cohort hadm_ids in pandas.

hadms_for_ed = set(
    pd.Series(hadm_list)
    .dropna()
    .astype("int64")
    .tolist()
)

# 1) ED stay map (full table), then restrict to cohort hadm_ids locally.
SQL["edmap_all_sql"] = f"""
SELECT stay_id, hadm_id, intime
FROM `{PHYS}.{ED}.edstays`
WHERE hadm_id IS NOT NULL
```

```
"""
edmap_all = run_sql_bq(sql("edmap_all_sql"))
edmap_all["hadm_id"] = pd.to_numeric(edmap_all["hadm_id"],
↪   errors="coerce").astype("Int64")
edmap_all["stay_id"] = pd.to_numeric(edmap_all["stay_id"],
↪   errors="coerce").astype("Int64")
edmap_all["intime"] = pd.to_datetime(edmap_all["intime"], errors="coerce")

edmap = edmap_all[edmap_all["hadm_id"].isin(hadms_for_ed)].copy()
edmap = edmap.drop_duplicates(subset=["stay_id", "hadm_id"])
print("ED map rows (cohort):", len(edmap))

if edmap.empty:
    ed_triage = pd.DataFrame(columns=[
        "hadm_id", "ed_triage_temp", "ed_triage_hr", "ed_triage_rr",
        "ed_triage_o2sat", "ed_triage_sbp", "ed_triage_dbp",
        ↪   "ed_triage_pain",
        "ed_triage_acuity", "ed_triage_cc",
    ])
    ed_first = pd.DataFrame(columns=[
        "hadm_id", "ed_first_vitals_time", "ed_first_temp", "ed_first_hr",
        "ed_first_rr", "ed_first_o2sat", "ed_first_sbp", "ed_first_dbp",
        "ed_first_rhythm", "ed_first_pain",
    ])
else:
    # 2) ED triage (full table), then reduce to earliest ED stay per hadm.
    SQL["ed_triage_all_sql"] = f"""
    SELECT
      stay_id,
      temperature    AS ed_triage_temp,
      heartrate      AS ed_triage_hr,
      resprate       AS ed_triage_rr,
      o2sat          AS ed_triage_o2sat,
      sbp            AS ed_triage_sbp,
      dbp            AS ed_triage_dbp,
      pain           AS ed_triage_pain,
      acuity         AS ed_triage_acuity,
      chiefcomplaint AS ed_triage_cc
    FROM `{PHYS}.{ED}.triage`
    """
    tri_all = run_sql_bq(sql("ed_triage_all_sql"))
    tri_all["stay_id"] = pd.to_numeric(tri_all["stay_id"],
↪   errors="coerce").astype("Int64")
```

81

```python
tri_merged = (
    edmap[["stay_id", "hadm_id", "intime"]]
    .merge(tri_all, on="stay_id", how="left")
)
ed_triage = (
    tri_merged
    .sort_values(["hadm_id", "intime"], na_position="last")
    .drop_duplicates(subset=["hadm_id"], keep="first")
    .drop(columns=["stay_id", "intime"])
    .reset_index(drop=True)
)
print("ED triage rows:", len(ed_triage))

# 3) First vitals per stay across full ED vitals table, then reduce to
#    earliest vitals-time per hadm.
SQL["ed_first_vitals_all_sql"] = f"""
WITH vs_ranked AS (
  SELECT
    v.stay_id,
    v.charttime,
    v.temperature,
    v.heartrate,
    v.resprate,
    v.o2sat,
    v.sbp,
    v.dbp,
    v.rhythm,
    v.pain,
    ROW_NUMBER() OVER (PARTITION BY v.stay_id ORDER BY v.charttime) AS rn
  FROM `{PHYS}.{ED}.vitalsign` v
)
SELECT
  stay_id,
  charttime    AS ed_first_vitals_time,
  temperature  AS ed_first_temp,
  heartrate    AS ed_first_hr,
  resprate     AS ed_first_rr,
  o2sat        AS ed_first_o2sat,
  sbp          AS ed_first_sbp,
  dbp          AS ed_first_dbp,
  rhythm       AS ed_first_rhythm,
  pain         AS ed_first_pain
```

```
    FROM vs_ranked
    WHERE rn = 1
    """
    first_stay_all = run_sql_bq(sql("ed_first_vitals_all_sql"))
    first_stay_all["stay_id"] = pd.to_numeric(first_stay_all["stay_id"],
↪  errors="coerce").astype("Int64")
    first_stay_all["ed_first_vitals_time"] =
↪  pd.to_datetime(first_stay_all["ed_first_vitals_time"], errors="coerce")

    first_merged = (
        edmap[["stay_id", "hadm_id"]]
        .merge(first_stay_all, on="stay_id", how="left")
    )
    ed_first = (
        first_merged
        .sort_values(["hadm_id", "ed_first_vitals_time"], na_position="last")
        .drop_duplicates(subset=["hadm_id"], keep="first")
        .drop(columns=["stay_id"])
        .reset_index(drop=True)
    )
    print("ED first vitals rows:", len(ed_first))
```

```
ED map rows (cohort): 28026
ED triage rows: 27975
ED first vitals rows: 27975
```

2.3.8 8) ICU meta (first ICU stay, LOS days)

Rationale: Summarize ICU exposure and length of stay to contextualize disease severity.

```
# Purpose: Build first-ICU metadata via full-table pull + local hadm filter
↪  to avoid large ARRAY parameter stalls.

SQL["icu_all_sql"] = f"""
SELECT
  hadm_id,
  stay_id AS first_icu_stay_id,
  intime  AS icu_intime,
  outtime AS icu_outtime
FROM `{PHYS}.{ICU}.icustays`
WHERE hadm_id IS NOT NULL
"""
```

```
icu_all = run_sql_bq(sql("icu_all_sql"))
icu_all["hadm_id"] = pd.to_numeric(icu_all["hadm_id"],
 ↪  errors="coerce").astype("Int64")
icu_all["first_icu_stay_id"] = pd.to_numeric(icu_all["first_icu_stay_id"],
 ↪  errors="coerce").astype("Int64")
icu_all["icu_intime"] = pd.to_datetime(icu_all["icu_intime"],
 ↪  errors="coerce")
icu_all["icu_outtime"] = pd.to_datetime(icu_all["icu_outtime"],
 ↪  errors="coerce")

hadms_for_icu = set(pd.Series(hadm_list).dropna().astype("int64").tolist())
icu_all = icu_all[icu_all["hadm_id"].isin(hadms_for_icu)].copy()

icu_meta = (
    icu_all
    .sort_values(["hadm_id", "icu_intime", "first_icu_stay_id"],
 ↪  na_position="last")
    .drop_duplicates(subset=["hadm_id"], keep="first")
    .reset_index(drop=True)
)
icu_meta["icu_los_days"] = (
    (icu_meta["icu_outtime"] - icu_meta["icu_intime"]).dt.total_seconds() /
 ↪  86400.0
)

print("ICU meta rows:", len(icu_meta))


ICU meta rows: 54630
```

2.3.9 9) Ventilation flags (ICD procedures)

Rationale: Identify IMV/NIV exposure as clinically relevant respiratory support indicators.

```
# Purpose: Build IMV/NIV ICD flags from targeted procedure codes, then filter
 ↪  to cohort hadm_ids locally.

SQL["vent_proc_sql"] = f"""
SELECT
  hadm_id,
  icd_version,
  REPLACE(icd_code, '.', '') AS code_norm
FROM `{PHYS}.{HOSP}.procedures_icd`
```

```
WHERE
  (icd_version = 10 AND icd_code IN
↪  ('5A1935Z','5A1945Z','5A1955Z','0BH17EZ','0BH18EZ','5A09357','5A09457','5A09557'))
  OR
  (icd_version = 9 AND REPLACE(icd_code, '.', '') IN
↪  ('9670','9671','9672','9604','9390','9391','9399'))
"""

vent_proc = run_sql_bq(sql("vent_proc_sql"))
vent_proc["hadm_id"] = pd.to_numeric(vent_proc["hadm_id"],
↪  errors="coerce").astype("Int64")
vent_proc["icd_version"] = pd.to_numeric(vent_proc["icd_version"],
↪  errors="coerce").astype("Int64")
vent_proc["code_norm"] = vent_proc["code_norm"].astype(str)

hadms_for_vent = set(pd.Series(hadm_list).dropna().astype("int64").tolist())
vent_proc = vent_proc[vent_proc["hadm_id"].isin(hadms_for_vent)].copy()

imv_codes_10 = {"5A1935Z", "5A1945Z", "5A1955Z", "0BH17EZ", "0BH18EZ"}
imv_codes_9 = {"9670", "9671", "9672", "9604"}
niv_codes_10 = {"5A09357", "5A09457", "5A09557"}
niv_codes_9 = {"9390", "9391", "9399"}

vent_proc["imv_hit"] = (
    ((vent_proc["icd_version"] == 10) &
↪  vent_proc["code_norm"].isin(imv_codes_10)) |
    ((vent_proc["icd_version"] == 9) &
↪  vent_proc["code_norm"].isin(imv_codes_9))
)
vent_proc["niv_hit"] = (
    ((vent_proc["icd_version"] == 10) &
↪  vent_proc["code_norm"].isin(niv_codes_10)) |
    ((vent_proc["icd_version"] == 9) &
↪  vent_proc["code_norm"].isin(niv_codes_9))
)

vent = (
    vent_proc
    .groupby("hadm_id", as_index=False)
    .agg(
        imv_flag=("imv_hit", "max"),
        niv_flag=("niv_hit", "max"),
    )
```

```
)
vent["imv_flag"] = vent["imv_flag"].astype(int)
vent["niv_flag"] = vent["niv_flag"].astype(int)
vent["any_vent_flag"] = ((vent["imv_flag"] == 1) | (vent["niv_flag"] ==
↪  1)).astype(int)

# Preserve one row per hadm in hadm_list even if no vent procedure codes were
↪  found.
vent = pd.DataFrame({"hadm_id": pd.Series(hadm_list,
↪  dtype="Int64")}).drop_duplicates().merge(vent, on="hadm_id", how="left")
vent[["imv_flag", "niv_flag", "any_vent_flag"]] = vent[["imv_flag",
↪  "niv_flag", "any_vent_flag"]].fillna(0).astype(int)

print("Vent rows:", len(vent))


Vent rows: 84594
```

and from charts

```
# Purpose: Extract earliest NIV/IMV charttimes using prefiltered ventilation
↪  itemids for better performance.

SQL["vent_chart_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),

stays AS (
  SELECT DISTINCT hadm_id, stay_id
  FROM `{PHYS}.{ICU}.icustays`
  WHERE hadm_id IN (SELECT hadm_id FROM hadms)
),

vent_itemids AS (
  SELECT itemid, LOWER(label) AS lbl
  FROM `{PHYS}.{ICU}.d_items`
  WHERE REGEXP_CONTAINS(LOWER(label), r'(vent|ventilator|mode|bipap|bi[-
↪  ]?pap|cpap|nippv|niv|mask|ett|endotracheal)')
),

cand AS (
  SELECT
    s.hadm_id,
    ce.charttime,
    vi.lbl,
```

```
      LOWER(COALESCE(ce.value, '')) AS valstr
    FROM `{PHYS}.{ICU}.chartevents` ce
    JOIN stays s ON s.stay_id = ce.stay_id
    JOIN vent_itemids vi ON vi.itemid = ce.itemid
),

flags AS (
  SELECT
    hadm_id,
    MIN(IF(
         REGEXP_CONTAINS(lbl, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
↪ ]?pap|cpap)')
         OR REGEXP_CONTAINS(valstr, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
↪ ]?pap|cpap)'),
         charttime, NULL)) AS first_niv_time,
    MIN(IF(
         REGEXP_CONTAINS(lbl, r'(invasive
↪ ventilation|endotracheal|ett|mech(|anical)? vent|ventilator
↪ mode|ac/|simv|prvc|aprv|pcv|vcv|assist\s*control)')
         OR REGEXP_CONTAINS(valstr, r'(invasive
↪ ventilation|endotracheal|ett|ac/|simv|prvc|aprv|pcv|vcv|assist\
↪ s*control)'),
         charttime, NULL)) AS first_imv_time,
    MAX(CASE
         WHEN REGEXP_CONTAINS(lbl, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
↪ ]?pap|cpap)')
            OR REGEXP_CONTAINS(valstr, r'(non[-
↪ ]?invasive|niv|nippv|bipap|bi[- ]?pap|cpap)')
         THEN 1 ELSE 0 END) AS niv_chart_flag,
    MAX(CASE
         WHEN REGEXP_CONTAINS(lbl, r'(invasive
↪ ventilation|endotracheal|ett|mech(|anical)? vent|ventilator
↪ mode|ac/|simv|prvc|aprv|pcv|vcv|assist\s*control)')
            OR REGEXP_CONTAINS(valstr, r'(invasive
↪ ventilation|endotracheal|ett|ac/|simv|prvc|aprv|pcv|vcv|assist\
↪ s*control)')
         THEN 1 ELSE 0 END) AS imv_chart_flag
  FROM cand
  GROUP BY hadm_id
)

SELECT hadm_id, niv_chart_flag, imv_chart_flag, first_niv_time,
↪ first_imv_time
```

```
FROM flags
"""

try:
    vent_chart = run_sql_bq(sql("vent_chart_sql"), {"hadms": hadm_list})
except Exception as e:
    print("WARNING: vent chart extraction failed; falling back to ICD-only
    ↪  vent timing.", e)
    vent_chart = pd.DataFrame(columns=["hadm_id", "niv_chart_flag",
↪  "imv_chart_flag", "first_niv_time", "first_imv_time"])

# Purpose: Derive respiratory support flags and timing fields for IMV/NIV
↪  exposure.

# If your existing ICD-only result is called `vent`, rename for clarity:
vent_proc = vent.copy()

# Outer merge so we keep hadm_ids that appear in only one source
vent_combined = vent_proc.merge(vent_chart, on="hadm_id", how="outer")

# Fill missing with 0 before taking maxima
for c in ["imv_flag","niv_flag","any_vent_flag","imv_chart_flag",⌋
↪  "niv_chart_flag"]:
    if c in vent_combined.columns:
        vent_combined[c] = vent_combined[c].fillna(0).astype("Int64")

# Final "any-source" flags
vent_combined["imv_flag"]        =
↪  vent_combined[["imv_flag","imv_chart_flag"]].max(axis=1).astype("Int64")
vent_combined["niv_flag"]        =
↪  vent_combined[["niv_flag","niv_chart_flag"]].max(axis=1).astype("Int64")
vent_combined["any_vent_flag"]   =
↪  vent_combined[["imv_flag","niv_flag"]].max(axis=1).astype("Int64")

vent_combined = vent_combined[["hadm_id","imv_flag","niv_flag",⌋
↪  "any_vent_flag","first_imv_time","first_niv_time"]]
print("After combining ICD + chart signals:",
      "\nIMV=1:", int((vent_combined["imv_flag"]==1).sum()),
      "\nNIV=1:", int((vent_combined["niv_flag"]==1).sum()))


After combining ICD + chart signals:
IMV=1: 36606
NIV=1: 23677
```

2.3.10 10) Assemble final DataFrame

Rationale: Merge all derived features into a single analytic table keyed by hadm_id.

```
# Purpose: Extract and standardize first ABG/VBG physiology fields for
↪  baseline characterization.

# Canonical base (carries authoritative subject_id)
df = demo.copy()

# Cohort flags / thresholds / labs / etc.
df = safe_merge_on_hadm(df, cohort_any, right_name="cohort_any",
↪  check_subject="warn")
df = safe_merge_on_hadm(df, bg_pairs,    right_name="bg_pairs")
df = safe_merge_on_hadm(df, race_eth,    right_name="race_eth")
df = safe_merge_on_hadm(df, ed_triage,   right_name="ed_triage")
df = safe_merge_on_hadm(df, ed_first,    right_name="ed_first")
df = safe_merge_on_hadm(df, icu_meta,    right_name="icu_meta")
df = safe_merge_on_hadm(df, vent_combined,      right_name="vent_combined")

# Anchor to first ED presentation (per admission)
if "ed_intime_first" in globals():
    df = safe_merge_on_hadm(df, ed_intime_first,
↪  right_name="ed_intime_first")

# Derived timing: first NIV/IMV relative to ED presentation
if "ed_intime_first" in df.columns and "first_imv_time" in df.columns:
    df["dt_first_imv_hours"] = (df["first_imv_time"] -
↪  df["ed_intime_first"]).dt.total_seconds() / 3600.0
if "ed_intime_first" in df.columns and "first_niv_time" in df.columns:
    df["dt_first_niv_hours"] = (df["first_niv_time"] -
↪  df["ed_intime_first"]).dt.total_seconds() / 3600.0

# ABG/VBG before IMV (hadm-level)
if {"first_abg_time", "first_imv_time"}.issubset(df.columns):
    df["abg_before_imv"] = (
        df["first_abg_time"].notna() & df["first_imv_time"].notna() &
        (df["first_abg_time"] < df["first_imv_time"])
    ).astype("Int64")
if {"first_vbg_time", "first_imv_time"}.issubset(df.columns):
    df["vbg_before_imv"] = (
        df["first_vbg_time"].notna() & df["first_imv_time"].notna() &
        (df["first_vbg_time"] < df["first_imv_time"])
```

```python
    ).astype("Int64")

# Canonical hadm-level threshold normalization.
# This consolidates explicit threshold columns with parsed ABG/VBG/other pCO2
  ↪  fields
# so later merges cannot silently zero-out source-specific thresholds.
def _num_series(frame: pd.DataFrame, col: str) -> pd.Series:
    if col in frame.columns:
        return pd.to_numeric(frame[col], errors="coerce")
    return pd.Series(np.nan, index=frame.index)


def _any_ge(frame: pd.DataFrame, cols: list[str], threshold: float) ->
  ↪  pd.Series:
    avail = [c for c in cols if c in frame.columns]
    if not avail:
        return pd.Series(0, index=frame.index, dtype="Int64")
    mat = pd.concat([_num_series(frame, c) for c in avail], axis=1)
    return mat.ge(threshold).any(axis=1).astype("Int64")


existing_abg = _num_series(df,
  ↪  "abg_hypercap_threshold").fillna(0).astype(int)
existing_vbg = _num_series(df,
  ↪  "vbg_hypercap_threshold").fillna(0).astype(int)
existing_other = _num_series(df,
  ↪  "other_hypercap_threshold").fillna(0).astype(int)

abg_from_values = _any_ge(df, ["lab_abg_paco2", "poc_abg_paco2",
  ↪  "first_abg_pco2"], 45.0)
vbg_from_values = _any_ge(df, ["lab_vbg_paco2", "poc_vbg_paco2",
  ↪  "first_vbg_pco2"], 50.0)
other_from_values = _any_ge(df, ["first_other_pco2"], 50.0)

df["abg_hypercap_threshold"] = ((existing_abg == 1) | (abg_from_values ==
  ↪  1)).astype("Int64")
df["vbg_hypercap_threshold"] = ((existing_vbg == 1) | (vbg_from_values ==
  ↪  1)).astype("Int64")
df["other_hypercap_threshold"] = ((existing_other == 1) | (other_from_values
  ↪  == 1)).astype("Int64")

df["pco2_threshold_any"] = (
    (df["abg_hypercap_threshold"] == 1)
```

```python
        | (df["vbg_hypercap_threshold"] == 1)
        | (df["other_hypercap_threshold"] == 1)
).astype("Int64")


if "any_hypercap_icd" in df.columns:
    df["enrolled_any"] = ((pd.to_numeric(df["any_hypercap_icd"],
 ↳  errors="coerce").fillna(0).astype(int) == 1) | (df["pco2_threshold_any"]
 ↳  == 1)).astype("Int64")


print(
    "Threshold sanity (hadm-level):",
    "ABG=", int(df["abg_hypercap_threshold"].fillna(0).sum()),
    "VBG=", int(df["vbg_hypercap_threshold"].fillna(0).sum()),
    "OTHER=", int(df["other_hypercap_threshold"].fillna(0).sum()),
    "ANY=", int(df["pco2_threshold_any"].fillna(0).sum()),
)


# Harmonize NIH race/ethnicity into a collapsed reporting label at
 ↳  cohort-build time.
RACE_NIH_MAP = {
    "WHITE": "WHITE",
    "BLACK OR AFRICAN AMERICAN": "BLACK",
    "ASIAN": "ASIAN",
    "AMERICAN INDIAN OR ALASKA NATIVE": "AI/AN",
    "NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER": "NH/PI",
    "MORE THAN ONE RACE": "MULTIRACIAL",
    "UNKNOWN OR NOT REPORTED": "UNKNOWN",
    "OTHER": "OTHER",
    "UNKNOWN": "UNKNOWN",
}


ETH_NIH_MAP = {
    "HISPANIC OR LATINO": "HISPANIC",
    "NOT HISPANIC OR LATINO": "NOT HISPANIC",
    "UNKNOWN OR NOT REPORTED": "UNKNOWN",
    "UNKNOWN": "UNKNOWN",
}


LABEL_MAP = {
    "WHITE": "Non-Hispanic White",
    "BLACK": "Non-Hispanic Black",
    "ASIAN": "Non-Hispanic Asian",
    "AI/AN": "Non-Hispanic American Indian/Alaska Native",
```

```python
        "NH/PI": "Non-Hispanic Native Hawaiian/Pacific Islander",
        "MULTIRACIAL": "Non-Hispanic Multiracial/Other",
        "OTHER": "Unknown/Other",
        "UNKNOWN": "Unknown/Other",
    }


    def add_collapsed_race_eth(
        frame: pd.DataFrame,
        race_col: str = "nih_race",
        eth_col: str = "nih_ethnicity",
        out_col: str = "race_eth_collapsed",
    ) -> pd.DataFrame:
        """Collapse NIH race/ethnicity fields into publication-friendly
        ↪  categories."""
        if race_col not in frame.columns or eth_col not in frame.columns:
            return frame

        def _std(value):
            if pd.isna(value):
                return None
            return re.sub(r"\s+", " ", str(value).strip()).upper()

        race_std = frame[race_col].map(lambda value:
    ↪  RACE_NIH_MAP.get(_std(value), "OTHER"))
        eth_std = frame[eth_col].map(lambda value: ETH_NIH_MAP.get(_std(value),
    ↪  "UNKNOWN"))

        collapsed = []
        for race_value, eth_value in zip(race_std, eth_std):
            if eth_value == "HISPANIC":
                collapsed.append("Hispanic/Latino")
            else:
                collapsed.append(LABEL_MAP.get(race_value, "Unknown/Other"))

        category_order = [
            "Hispanic/Latino",
            "Non-Hispanic White",
            "Non-Hispanic Black",
            "Non-Hispanic Asian",
            "Non-Hispanic American Indian/Alaska Native",
            "Non-Hispanic Native Hawaiian/Pacific Islander",
            "Non-Hispanic Multiracial/Other",
```

```python
        "Unknown/Other",
    ]

    frame[out_col] = pd.Categorical(collapsed, categories=category_order,
↪    ordered=True)
    frame["nih_race_std"] = race_std
    frame["nih_ethnicity_std"] = eth_std
    return frame


df = add_collapsed_race_eth(df)

print("Final df rows:", len(df), "cols:", len(df.columns))

# Safety checks
assert "subject_id" in df.columns, "subject_id missing from final df"
assert not any(c.endswith("_x") or c.endswith("_y") for c in df.columns),
↪    "Found suffixed columns"
print("Final df rows:", len(df), "cols:", len(df.columns))
df.head(3)
```

```
Threshold sanity (hadm-level): ABG= 17090 VBG= 39016 OTHER= 64911 ANY= 83907
Final df rows: 84594 cols: 107
Final df rows: 84594 cols: 107
```

	hadm_id	subject_id	admittime	dischtime	deathtime	admission_type
0	23485217	10584718	2165-02-12 15:41:00	2165-03-06 08:20:00	2165-03-06 08:20:00	EW EMER.
1	27004173	10297948	2135-08-23 18:46:00	2135-09-08 15:50:00	NaT	URGENT
2	22661627	10032409	2130-01-12 18:42:00	2130-01-21 14:32:00	NaT	EW EMER.

```python
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪    consistent across notebook stages.

# --- Cohort flow counts (ED / ICU / blood gas / hypercapnia / CC) ---

# 1) Dataset-level ED counts
SQL["ed_counts_sql"] = f"""
SELECT
  COUNT(*) AS total_ed_encounters,
  COUNTIF(hadm_id IS NOT NULL) AS ed_encounters_with_hadm
FROM `{PHYS}.{ED}.edstays`
```

```python
"""

SQL["ed_to_icu_sql"] = f"""
SELECT
  COUNT(DISTINCT e.hadm_id) AS ed_to_icu_hadm,
  COUNT(DISTINCT e.stay_id) AS ed_to_icu_edstays
FROM `{PHYS}.{ED}.edstays` e
JOIN `{PHYS}.{ICU}.icustays` i USING (hadm_id)
WHERE e.hadm_id IS NOT NULL
"""

try:
    ed_counts = run_sql_bq(sql("ed_counts_sql"))
    ed_to_icu = run_sql_bq(sql("ed_to_icu_sql"))
except Exception as e:
    print("Warning: ED/ICU counts query failed:", e)
    ed_counts = None
    ed_to_icu = None

# 2) Cohort-level counts (admission-level)
cohort_union = int((cohort_any["enrolled_any"] == 1).sum()) if "cohort_any"
 ↪  in globals() and "enrolled_any" in cohort_any.columns else len(hadm_list)
cohort_df_n = len(df)

# Any blood gas present (ABG/VBG, LAB/POC)
co2_cols = [c for c in [
    "lab_abg_paco2", "lab_vbg_paco2", "poc_abg_paco2", "poc_vbg_paco2"
] if c in df.columns]
any_bg = int(df[co2_cols].notna().any(axis=1).sum()) if co2_cols else None

# Hypercapnia thresholds and ICD
icd_count = int((df["any_hypercap_icd"] == 1).sum()) if "any_hypercap_icd"
 ↪  in df.columns else None
threshold_count = int((df["pco2_threshold_any"] == 1).sum()) if
 ↪  "pco2_threshold_any" in df.columns else None

# ED chief complaint missing / present (within cohort)
if "ed_triage_cc" in df.columns:
    mask_cc_present = df["ed_triage_cc"].notna() &
 ↪  (df["ed_triage_cc"].astype(str).str.strip() != "")
    cc_present = int(mask_cc_present.sum())
    cc_missing = int((~mask_cc_present).sum())
else:
```

```python
    cc_present = None
    cc_missing = None

# ED→ICU within cohort (admissions with ED triage data and ICU stay)
if "first_icu_stay_id" in df.columns and "ed_triage_cc" in df.columns:
    cohort_ed_to_icu = int((df["first_icu_stay_id"].notna() &
↪   mask_cc_present).sum())
else:
    cohort_ed_to_icu = None


rows = []
if ed_counts is not None:
    rows.append({"step": "Total ED encounters (edstays)", "count":
↪   int(ed_counts.loc[0, "total_ed_encounters"]), "scope": "All ED dataset"})
    rows.append({"step": "ED encounters with hadm_id", "count":
↪   int(ed_counts.loc[0, "ed_encounters_with_hadm"]), "scope": "All ED
↪   dataset"})
if ed_to_icu is not None:
    rows.append({"step": "ED→ICU admissions (distinct hadm_id)", "count":
↪   int(ed_to_icu.loc[0, "ed_to_icu_hadm"]), "scope": "All ED+ICU"})
    rows.append({"step": "ED→ICU ED-stays (distinct stay_id)", "count":
↪   int(ed_to_icu.loc[0, "ed_to_icu_edstays"]), "scope": "All ED+ICU"})

rows.append({"step": "Cohort admissions (union ICD  thresholds)", "count":
↪   cohort_union, "scope": "Cohort"})
rows.append({"step": "Cohort admissions after merges (df rows)", "count":
↪   cohort_df_n, "scope": "Cohort"})
if any_bg is not None:
    rows.append({"step": "Cohort with any ABG/VBG (LAB or POC)", "count":
↪   any_bg, "scope": "Cohort"})
if threshold_count is not None:
    rows.append({"step": "Cohort meeting hypercapnia thresholds", "count":
↪   threshold_count, "scope": "Cohort"})
if icd_count is not None:
    rows.append({"step": "Cohort meeting ICD code criteria", "count":
↪   icd_count, "scope": "Cohort"})
if cc_present is not None:
    rows.append({"step": "Cohort with ED chief complaint present", "count":
↪   cc_present, "scope": "Cohort"})
if cc_missing is not None:
    rows.append({"step": "Cohort excluded for missing ED chief complaint",
↪   "count": cc_missing, "scope": "Cohort"})
if cohort_ed_to_icu is not None:
```

```
    rows.append({"step": "Cohort ED→ICU (ED CC present + ICU stay)", "count":
 ↪  cohort_ed_to_icu, "scope": "Cohort"})

flow_counts = pd.DataFrame(rows)
flow_counts
```

	step	count	scope
0	Total ED encounters (edstays)	425087	All ED dataset
1	ED encounters with hadm_id	203016	All ED dataset
2	ED→ICU admissions (distinct hadm_id)	31862	All ED+ICU
3	ED→ICU ED-stays (distinct stay_id)	31916	All ED+ICU
4	Cohort admissions (union ICD thresholds)	84595	Cohort
5	Cohort admissions after merges (df rows)	84594	Cohort
6	Cohort with any ABG/VBG (LAB or POC)	75870	Cohort
7	Cohort meeting hypercapnia thresholds	83907	Cohort
8	Cohort meeting ICD code criteria	4237	Cohort
9	Cohort with ED chief complaint present	27975	Cohort
10	Cohort excluded for missing ED chief complaint	56619	Cohort
11	Cohort ED→ICU (ED CC present + ICU stay)	18507	Cohort

```
# Purpose: Quantify overlap between ascertainment routes so non-exclusive
 ↪  cohorts are explicit.

# --- Ascertainment overlap counts (ABG/VBG/ICD) ---

required = ["abg_hypercap_threshold", "vbg_hypercap_threshold",
 ↪  "other_hypercap_threshold", "any_hypercap_icd"]
missing = [c for c in required if c not in df.columns]
if missing:
    raise KeyError(f"Missing required columns for overlap counts: {missing}")

abg = pd.to_numeric(df["abg_hypercap_threshold"],
 ↪  errors="coerce").fillna(0).astype(int)
vbg = pd.to_numeric(df["vbg_hypercap_threshold"],
 ↪  errors="coerce").fillna(0).astype(int)
other = pd.to_numeric(df["other_hypercap_threshold"],
 ↪  errors="coerce").fillna(0).astype(int)

gas_any = (
    pd.to_numeric(df.get("pco2_threshold_any", None), errors="coerce")
    if "pco2_threshold_any" in df.columns else (abg | vbg | other)
```

```python
)
if hasattr(gas_any, "fillna"):
    gas_any = gas_any.fillna(0).astype(int)
else:
    gas_any = gas_any.astype(int)

icd = pd.to_numeric(df["any_hypercap_icd"],
↪    errors="coerce").fillna(0).astype(int)

total_n = len(df)
ngas = int((gas_any == 1).sum())

abg_vbg_overlap = pd.DataFrame([
    {"group": "ABG-only", "count": int(((abg==1) & (vbg==0) &
↪    (other==0)).sum())},
    {"group": "VBG-only", "count": int(((vbg==1) & (abg==0) &
↪    (other==0)).sum())},
    {"group": "Other-only", "count": int(((other==1) & (abg==0) &
↪    (vbg==0)).sum())},
    {"group": "Mixed (>=2 routes)", "count": int((((abg + vbg + other) >=
↪    2).sum()))},
])
if ngas > 0:
    abg_vbg_overlap["pct_of_gas"] = (abg_vbg_overlap["count"] / ngas *
↪    100).round(1)
else:
    abg_vbg_overlap["pct_of_gas"] = 0.0
abg_vbg_overlap["pct_of_cohort"] = (abg_vbg_overlap["count"] / max(total_n,1)
↪    * 100).round(1)

icd_gas_overlap = pd.DataFrame([
    {"group": "ICD+Gas", "count": int(((icd==1) & (gas_any==1)).sum())},
    {"group": "ICD-only", "count": int(((icd==1) & (gas_any==0)).sum())},
    {"group": "Gas-only", "count": int(((icd==0) & (gas_any==1)).sum())},
    {"group": "Neither", "count": int(((icd==0) & (gas_any==0)).sum())},
])
icd_gas_overlap["pct_of_cohort"] = (icd_gas_overlap["count"] / max(total_n,1)
↪    * 100).round(1)

print("ABG/VBG/Other overlap (among gas-positive):")
print(abg_vbg_overlap.to_string(index=False))
print("ICD vs Gas overlap (cohort-level):")
print(icd_gas_overlap.to_string(index=False))
```

```
ABG/VBG/Other overlap (among gas-positive):
            group   count  pct_of_gas  pct_of_cohort
         ABG-only       1         0.0            0.0
         VBG-only   18955        22.6           22.4
       Other-only   33779        40.3           39.9
Mixed (>=2 routes)   31172        37.2           36.8
ICD vs Gas overlap (cohort-level):
   group   count  pct_of_cohort
 ICD+Gas    3550            4.2
ICD-only     687            0.8
Gas-only   80357           95.0
 Neither       0            0.0
```

```python
# Purpose: Capture ED presentation context (chief complaint, acuity, and
↪   early vitals) at the ED-stay level.

# --- Missingness summary (chief complaint, race/ethnicity, ED triage/vitals)
↪   ---

import numpy as np

summary_rows = []

# Chief complaint missingness (ED triage CC)
if "ed_triage_cc" in df.columns:
    cc_present = df["ed_triage_cc"].notna() &
↪   (df["ed_triage_cc"].astype(str).str.strip() != "")
    summary_rows.append({
        "variable": "ed_triage_cc_present",
        "missing_n": int((~cc_present).sum()),
        "missing_pct": float((~cc_present).mean())
    })

# Race/Ethnicity missingness (NIH categories + raw sources)
unknown_tokens = {
    "unknown or not reported",
    "unknown",
    "not reported",
    "missing",
    "declined",
    "unable"
}
```

```python
def _missing_rate(series):
    if series is None:
        return None, None
    s = series.astype(str).str.strip()
    is_missing = series.isna() | (s == "") |
↪    s.str.lower().isin(unknown_tokens)
    return int(is_missing.sum()), float(is_missing.mean())

for col in ["nih_race", "nih_ethnicity", "race_hosp_raw", "race_ed_raw"]:
    if col in df.columns:
        m_n, m_p = _missing_rate(df[col])
        summary_rows.append({
            "variable": col,
            "missing_n": m_n,
            "missing_pct": m_p
        })

missing_summary = pd.DataFrame(summary_rows)
print("Missingness summary (key variables):")
missing_summary

# ED triage + first ED vitals missingness
triage_cols = [c for c in df.columns if c.startswith("ed_triage_")]
first_cols  = [c for c in df.columns if c.startswith("ed_first_")]

vital_cols = triage_cols + first_cols
if vital_cols:
    miss_tbl = (
        pd.DataFrame({"variable": vital_cols})
        .assign(
            missing_n=lambda d: [int(df[c].isna().sum()) for c in
↪    d["variable"]],
            missing_pct=lambda d: [float(df[c].isna().mean()) for c in
↪    d["variable"]]
        )
        .sort_values("missing_pct", ascending=False)
    )
    print("Missingness summary (ED triage + first ED vitals):")
    miss_tbl
else:
    miss_tbl = pd.DataFrame(columns=["variable", "missing_n", "missing_pct"])


Missingness summary (key variables):
```

Missingness summary (ED triage + first ED vitals):

2.3.11 11) Sanity checks

Rationale: Run QC checks to validate units, flags, and basic data integrity before export.

```python
# Purpose: Derive respiratory support flags and timing fields for IMV/NIV
↪   exposure.


# ICU LOS negative?
if {"icu_los_days","first_icu_stay_id"}.issubset(df.columns):
    neg_los = int((df["icu_los_days"] < 0).fillna(False).sum())
    print("Negative ICU LOS rows:", neg_los)

# Vent flags consistency
vent_cols = {"imv_flag","niv_flag","any_vent_flag"}
if vent_cols.issubset(df.columns):
    any_calc = ((df["imv_flag"]==1) |
↪   (df["niv_flag"]==1)).fillna(False).astype(int)
    any_flag = pd.to_numeric(df["any_vent_flag"],
↪   errors="coerce").fillna(0).astype(int)
    mism = int((any_calc != any_flag).sum())
    print("any_vent_flag mismatches vs (imv|niv):", mism)

# UOMs: expect mmhg only
uom_cols = [c for c in df.columns if c.endswith("_paco2_uom")]
for c in uom_cols:
    vals = sorted(pd.Series(df[c]).dropna().astype(str).str.lower().str⌋
↪   .strip().unique().tolist())
    print(c, vals)

# ABG/VBG coverage QC
def qc_pair(df, ph_col, co2_col, label, ph_lo=6.3, ph_hi=7.8, co2_lo=5,
↪   co2_hi=200):
    ph  = pd.to_numeric(df.get(ph_col), errors="coerce")
    co2 = pd.to_numeric(df.get(co2_col), errors="coerce")
    return {
        "pair": label,
        "present_any":  int(((ph.notna()) | (co2.notna())).sum()),
        "present_both": int(((ph.notna()) & (co2.notna())).sum()),
        "only_ph":      int(((ph.notna()) & (~co2.notna())).sum()),
        "only_pco2":    int(((co2.notna()) & (~ph.notna())).sum()),
```

```
        "ph_oob":       int(((((ph  < ph_lo)  | (ph  > ph_hi))  &
        ↪  ph.notna())).sum()),
        "pco2_oob":     int(((((co2 < co2_lo) | (co2 > co2_hi)) &
        ↪  co2.notna())).sum()),
    }

qc = pd.DataFrame([
    qc_pair(df, "lab_abg_ph","lab_abg_paco2","LAB ABG"),
    qc_pair(df, "lab_vbg_ph","lab_vbg_paco2","LAB VBG"),
    qc_pair(df, "poc_abg_ph","poc_abg_paco2","POC ABG"),
    qc_pair(df, "poc_vbg_ph","poc_vbg_paco2","POC VBG"),
])
qc
```

```
Negative ICU LOS rows: 0
any_vent_flag mismatches vs (imv|niv): 0
lab_abg_paco2_uom ['mmhg']
lab_vbg_paco2_uom ['mmhg']
lab_other_paco2_uom ['mmhg']
poc_abg_paco2_uom []
poc_vbg_paco2_uom []
poc_other_paco2_uom ['mmhg']
```

	pair	present_any	present_both	only_ph	only_pco2	ph_oob	pco2_oob
0	LAB ABG	55913	55452	449	12	0	0
1	LAB VBG	43197	35309	7880	8	0	0
2	POC ABG	0	0	0	0	0	0
3	POC VBG	0	0	0	0	0	0

2.3.12 PDF-ready long tables

```
print(
    render_latex_longtable(
        flow_counts,
        caption=f"Cohort flow and key denominator checkpoints
↪  (N={len(df):,}).",
        label="tab:cohort_flow_counts",
        index=False,
    )
)
```

```python
print(
    render_latex_longtable(
        abg_vbg_overlap,
        caption="ABG/VBG/OTHER overlap among gas-positive encounters.",
        label="tab:gas_overlap_abg_vbg_other",
        index=False,
    )
)
print(
    render_latex_longtable(
        icd_gas_overlap,
        caption="ICD versus gas ascertainment overlap in the full cohort.",
        label="tab:icd_gas_overlap",
        index=False,
    )
)
print(
    render_latex_longtable(
        missing_summary,
        caption="Missingness summary for key cohort variables.",
        label="tab:key_missingness",
        index=False,
    )
)
if "miss_tbl" in globals() and not miss_tbl.empty:
    print(
        render_latex_longtable(
            miss_tbl,
            caption="ED triage and first ED vital-sign missingness summary.",
            label="tab:ed_vital_missingness",
            landscape=True,
            index=False,
        )
    )
print(
    render_latex_longtable(
        qc,
        caption="ABG/VBG pH and pCO2 plausibility quality checks.",
        label="tab:abg_vbg_qc",
        landscape=True,
        index=False,
    )
)
```

```latex
\begin{longtable}{lrl}
\caption{Cohort flow and key denominator checkpoints (N=84,594).} \label{tab:cohort_flow_cou
\toprule
step & count & scope \\
\midrule
\endfirsthead
\caption[]{Cohort flow and key denominator checkpoints (N=84,594).} \\
\toprule
step & count & scope \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
Total ED encounters (edstays) & 425087 & All ED dataset \\
ED encounters with hadm_id & 203016 & All ED dataset \\
ED→ICU admissions (distinct hadm_id) & 31862 & All ED+ICU \\
ED→ICU ED-stays (distinct stay_id) & 31916 & All ED+ICU \\
Cohort admissions (union ICD  thresholds) & 84595 & Cohort \\
Cohort admissions after merges (df rows) & 84594 & Cohort \\
Cohort with any ABG/VBG (LAB or POC) & 75870 & Cohort \\
Cohort meeting hypercapnia thresholds & 83907 & Cohort \\
Cohort meeting ICD code criteria & 4237 & Cohort \\
Cohort with ED chief complaint present & 27975 & Cohort \\
Cohort excluded for missing ED chief complaint & 56619 & Cohort \\
Cohort ED→ICU (ED CC present + ICU stay) & 18507 & Cohort \\
\end{longtable}

\begin{longtable}{lrrr}
\caption{ABG/VBG/OTHER overlap among gas-positive encounters.} \label{tab:gas_overlap_abg_vb
\toprule
group & count & pct_of_gas & pct_of_cohort \\
\midrule
\endfirsthead
\caption[]{ABG/VBG/OTHER overlap among gas-positive encounters.} \\
\toprule
group & count & pct_of_gas & pct_of_cohort \\
\midrule
\endhead
\midrule
```

```latex
\multicolumn{4}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ABG-only & 1 & 0.000000 & 0.000000 \\
VBG-only & 18955 & 22.600000 & 22.400000 \\
Other-only & 33779 & 40.300000 & 39.900000 \\
Mixed (>=2 routes) & 31172 & 37.200000 & 36.800000 \\
\end{longtable}

\begin{longtable}{lrr}
\caption{ICD versus gas ascertainment overlap in the full cohort.} \label{tab:icd_gas_overlap}
\toprule
group & count & pct_of_cohort \\
\midrule
\endfirsthead
\caption[]{ICD versus gas ascertainment overlap in the full cohort.} \\
\toprule
group & count & pct_of_cohort \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ICD+Gas & 3550 & 4.200000 \\
ICD-only & 687 & 0.800000 \\
Gas-only & 80357 & 95.000000 \\
Neither & 0 & 0.000000 \\
\end{longtable}

\begin{longtable}{lrr}
\caption{Missingness summary for key cohort variables.} \label{tab:key_missingness} \\
\toprule
variable & missing_n & missing_pct \\
\midrule
\endfirsthead
\caption[]{Missingness summary for key cohort variables.} \\
\toprule
variable & missing_n & missing_pct \\
```

```
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ed_triage_cc_present & 56619 & 0.669303 \\
nih_race & 14782 & 0.174741 \\
nih_ethnicity & 13420 & 0.158640 \\
race_hosp_raw & 5999 & 0.070915 \\
race_ed_raw & 58868 & 0.695889 \\
\end{longtable}

\begin{landscape}\n\begin{longtable}{lrr}
\caption{ED triage and first ED vital-sign missingness summary.} \label{tab:ed_vital_missingr
\toprule
variable & missing_n & missing_pct \\
\midrule
\endfirsthead
\caption[]{ED triage and first ED vital-sign missingness summary.} \\
\toprule
variable & missing_n & missing_pct \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ed_first_rhythm & 83071 & 0.981996 \\
ed_first_temp & 66955 & 0.791486 \\
ed_triage_temp & 63525 & 0.750940 \\
ed_triage_o2sat & 62923 & 0.743823 \\
ed_triage_rr & 62910 & 0.743670 \\
ed_triage_dbp & 62675 & 0.740892 \\
ed_triage_sbp & 62589 & 0.739875 \\
ed_triage_hr & 62535 & 0.739237 \\
ed_first_o2sat & 62481 & 0.738598 \\
ed_first_dbp & 61780 & 0.730312 \\
ed_first_sbp & 61780 & 0.730312 \\
```

```
ed_first_rr & 61718 & 0.729579 \\
ed_first_pain & 61511 & 0.727132 \\
ed_first_hr & 61432 & 0.726198 \\
ed_triage_pain & 60163 & 0.711197 \\
ed_triage_acuity & 59152 & 0.699246 \\
ed_first_vitals_time & 58502 & 0.691562 \\
ed_triage_cc & 56619 & 0.669303 \\
\end{longtable}
\n\end{landscape}\n
\begin{landscape}\n\begin{longtable}{lrrrrrr}
\caption{ABG/VBG pH and pCO2 plausibility quality checks.} \label{tab:abg_vbg_qc} \\
\toprule
pair & present_any & present_both & only_ph & only_pco2 & ph_oob & pco2_oob \\
\midrule
\endfirsthead
\caption[]{ABG/VBG pH and pCO2 plausibility quality checks.} \\
\toprule
pair & present_any & present_both & only_ph & only_pco2 & ph_oob & pco2_oob \\
\midrule
\endhead
\midrule
\multicolumn{7}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
LAB ABG & 55913 & 55452 & 449 & 12 & 0 & 0 \\
LAB VBG & 43197 & 35309 & 7880 & 8 & 0 & 0 \\
POC ABG & 0 & 0 & 0 & 0 & 0 & 0 \\
POC VBG & 0 & 0 & 0 & 0 & 0 & 0 \\
\end{longtable}
\n\end{landscape}\n
```

2.3.13 12) Save to Excel

Rationale: Persist cohort outputs for downstream annotation and NLP analyses.

```
# Purpose: Optional archive export of the full hadm-level cohort workbook.

from datetime import datetime
```

```
WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
↪  "1"

if WRITE_ARCHIVE_XLSX_EXPORTS:
    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    out_path = prior_runs_dir / f"mimic_hypercap_EXT_bq_abg_vbg_{datetime.
↪  now().strftime('%Y%m%d_%H%M%S')}.xlsx"
    with pd.ExcelWriter(out_path, engine="openpyxl") as xw:
        df.to_excel(xw, sheet_name="cohort", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    out_path
else:
    print("Skipping archive hadm-level workbook export (set
↪  WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")

Skipping archive hadm-level workbook export (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).
```

2.3.14 Create Annotation Dataset

Rationale: Create ED chief-complaint subsets and a reproducible sample for manual annotation.

```
# Purpose: Capture ED presentation context (chief complaint, acuity, and
↪  early vitals) at the ED-stay level.

# ---- Extra exports: (1) ED chief-complaint only; (2) random sample of 160
↪  patients ----
from datetime import datetime

# Dated artifacts go to archive folder (optional)
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')
WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
↪  "1"

# 1) Filter to rows with a non-empty ED chief complaint
if "ed_triage_cc" not in df.columns:
```

```python
    raise KeyError(
        "Column 'ed_triage_cc' not found in df. "
        "Ensure the ED triage merge cell ran earlier."
    )

mask_cc = df["ed_triage_cc"].notna() &
↪  (df["ed_triage_cc"].astype(str).str.strip() != "")
df_cc = df.loc[mask_cc].copy()

print(f"ED-CC present rows: {len(df_cc)} of {len(df)} "
      f"({(len(df_cc) / max(len(df),1)):.1%} of cohort).")

# Save ED-CC-only cohort only when archive exports are enabled
if WRITE_ARCHIVE_XLSX_EXPORTS:
    out_path_cc = prior_runs_dir /
↪  f"mimic_hypercap_EXT_EDcc_only_bq_abg_vbg_{timestamp}.xlsx"
    with pd.ExcelWriter(out_path_cc, engine="openpyxl") as xw:
        df_cc.to_excel(xw, sheet_name="cohort_cc_only", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    print("Saved:", out_path_cc)

# 2) Random sample of n = 160 patients (distinct subject_id), one row per
↪  patient
if "subject_id" not in df_cc.columns:
    raise KeyError("Column 'subject_id' missing; cannot sample by patient.")

# Make a one-row-per-patient frame by earliest admission
if "admittime" in df_cc.columns:
    df_cc_one = (
        df_cc.sort_values(["subject_id", "admittime"])
            .groupby("subject_id", as_index=False)
            .head(1)
    )
else:
    # Fallback if admittime not present: choose the smallest hadm_id per
    ↪  patient
    df_cc_one = (
        df_cc.sort_values(["subject_id", "hadm_id"])
            .groupby("subject_id", as_index=False)
            .head(1)
```

```
    )

N = 160
n_avail = len(df_cc_one)
n_take = min(N, n_avail)
if n_avail < N:
    print(f"Warning: only {n_avail} unique patients with ED chief complaint;
    ↪  sampling all of them.")

RANDOM_SEED = 42
df_cc_sample = df_cc_one.sample(n=n_take, random_state=RANDOM_SEED)

# Save the sample only when archive exports are enabled
if WRITE_ARCHIVE_XLSX_EXPORTS:
    out_path_cc_sample = prior_runs_dir /
    ↪  f"mimic_hypercap_EXT_EDcc_sample{n_take}_bq_abg_vbg_{timestamp}.xlsx"
    with pd.ExcelWriter(out_path_cc_sample, engine="openpyxl") as xw:
        df_cc_sample.to_excel(xw, sheet_name="cohort_cc_sample", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    print("Saved:", out_path_cc_sample)
else:
    print("Skipping ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1
    ↪  to enable).")


ED-CC present rows: 27975 of 84594 (33.1% of cohort).
Skipping ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).
```

3 ED-stay cohort expansion (timing, severity, comorbidity, outcomes)

Rationale: Build a one-row-per-ED-stay analytic extract with time-anchored gas phenotypes, key comorbidities, and outcomes.

3.0.1 Phase 0 — Inventory & missing-field registry

Rationale: Detect which fields already exist and only add missing fields to avoid redundant extraction or join explosions.

```python
# Purpose: Inventory available fields and identify missing target variables
↪  before enrichment steps.

from pathlib import Path
import json
import pandas as pd

# Identify current cohort dataframe (admission-level)
if 'df' not in globals():
    raise NameError("Expected admission-level df to exist before inventory
        ↪  step.")

ED_KEY = "ed_stay_id"  # target key for ED-level cohort

print("Current admission-level df:", df.shape)
print("Current columns count:", len(df.columns))
if ED_KEY in df.columns:
    print("ED stay unique count:", int(df[ED_KEY].nunique()))
else:
    print("ED stay unique count: ED_KEY not in columns")

# Persist columns snapshot
cols_out = WORK_DIR / "current_columns.json"
cols_out.write_text(json.dumps(sorted(df.columns), indent=2))
print("Wrote:", cols_out)




# Persist ED-stay columns snapshot (if ed_df exists)
if "ed_df" in globals():
    ed_cols_out = WORK_DIR / "ed_columns.json"
    ed_cols_out.write_text(json.dumps(sorted(ed_df.columns), indent=2))
    print("Wrote:", ed_cols_out)
```

```
Current admission-level df: (84594, 107)
Current columns count: 107
ED stay unique count: ED_KEY not in columns
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/current_columns.json
```

```python
# Purpose: Capture ED presentation context (chief complaint, acuity, and
↪  early vitals) at the ED-stay level.
```

```python
# Target field registry
TARGET_RAW_FIELDS = {
    "ed_edstays": [
        "ed_stay_id","subject_id","hadm_id","ed_intime","ed_outtime",⌋
        ↪  "ed_intime_first",
        "arrival_transport","disposition","ed_gender","ed_race",
    ],
    "ed_triage": [
        "ed_triage_temp","ed_triage_hr","ed_triage_rr","ed_triage_o2sat",
        "ed_triage_sbp","ed_triage_dbp","ed_triage_pain","ed_triage_acuity",⌋
        ↪  "ed_triage_cc",
    ],
    "ed_vitals_first": [
        "ed_first_vitals_time","ed_first_temp","ed_first_hr","ed_first_rr",
        "ed_first_o2sat","ed_first_sbp","ed_first_dbp","ed_first_rhythm",⌋
        ↪  "ed_first_pain",
    ],
    "admissions": [
        "admittime","dischtime","deathtime","hospital_expire_flag",
        "admission_type","admission_location","discharge_location",
        "insurance","language","marital_status","hosp_race",
    ],
    "icu": [
        "icu_stay_id","icu_intime_first","icu_outtime_last","icu_los_total",⌋
        ↪  "n_icu_stays",
        "first_careunit","last_careunit",
    ],
    "labs_gas": [
        "first_gas_time","first_pco2","first_ph","first_hco3",⌋
        ↪  "first_lactate",
        "max_pco2_0_6h","min_ph_0_6h","max_pco2_0_24h","min_ph_0_24h",
        "flag_abg_hypercapnia","flag_vbg_hypercapnia",⌋
        ↪  "flag_other_hypercapnia","flag_any_gas_hypercapnia",
        "gas_source_other_rate",
        "dt_first_imv_hours","dt_first_niv_hours","first_other_time",
    ],
    "omr": [
        "bmi_closest_pre_ed","height_closest_pre_ed","weight_closest_pre_ed",
        "anthro_timing_tier","anthro_days_offset","anthro_chartdate",⌋
        ↪  "anthro_timing_uncertain",
        "anthro_source","anthro_obstime","anthro_hours_offset",⌋
        ↪  "anthro_timing_basis",
    ],
```

```python
    "dx_flags": [
        "flag_copd","flag_osa_ohs","flag_chf","flag_neuromuscular",
        "flag_opioid_substance","flag_pneumonia",
    ],
    "timing": [
        "dt_first_qualifying_gas_hours","presenting_hypercapnia",
        ↪  "late_hypercapnia",
        "dt_first_imv_hours","dt_first_niv_hours","abg_before_imv",
        ↪  "vbg_before_imv",
        "ph_band","hco3_band","lactate_band",
    ],
}


TARGET_DERIVED_FIELDS = [
    "hospital_los_hours","in_hospital_death",
]


TARGET_FIELDS = sorted({c for v in TARGET_RAW_FIELDS.values() for c in v} |
 ↪  set(TARGET_DERIVED_FIELDS))

# Grouped missing report
missing_by_group = {}
for group, cols in TARGET_RAW_FIELDS.items():
    missing_by_group[group] = [c for c in cols if c not in df.columns]

missing_derived = [c for c in TARGET_DERIVED_FIELDS if c not in df.columns]
missing_by_group["derived"] = missing_derived

missing_flat = [c for cols in missing_by_group.values() for c in cols]
print("Missing fields total:", len(missing_flat))
for group, cols in missing_by_group.items():
    if cols:
        print(f"- {group}: {cols}")


Missing fields total: 62
- ed_edstays: ['ed_stay_id', 'ed_intime', 'ed_outtime', 'ed_intime_first', 'arrival_transport
- admissions: ['hospital_expire_flag', 'language', 'marital_status', 'hosp_race']
- icu: ['icu_stay_id', 'icu_intime_first', 'icu_outtime_last', 'icu_los_total', 'n_icu_stays
- labs_gas: ['first_gas_time', 'first_pco2', 'first_ph', 'first_hco3', 'first_lactate', 'max_
- omr: ['bmi_closest_pre_ed', 'height_closest_pre_ed', 'weight_closest_pre_ed', 'anthro_timi
- dx_flags: ['flag_copd', 'flag_osa_ohs', 'flag_chf', 'flag_neuromuscular', 'flag_opioid_subs
- timing: ['dt_first_qualifying_gas_hours', 'presenting_hypercapnia', 'late_hypercapnia', 'dt
- derived: ['hospital_los_hours', 'in_hospital_death']
```

3.0.2 Phase 1 — ED encounter spine and ED enrichment (one row per ED stay)

Rationale: Create a dedicated ED-stay-level cohort with ED-specific attributes to avoid mixing admission- and ED-level grains.

```python
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪    consistent across notebook stages.

# ED stay spine (rename stay_id to ed_stay_id)

SQL["ed_spine_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
  s.stay_id AS ed_stay_id,
  s.subject_id,
  s.hadm_id,
  s.intime AS ed_intime,
  s.outtime AS ed_outtime,
  s.arrival_transport,
  s.disposition,
  s.gender AS ed_gender,
  s.race    AS ed_race
FROM `{PHYS}.{ED}.edstays` s
JOIN hadms h ON h.hadm_id = s.hadm_id
"""

ed_spine = run_sql_bq(sql("ed_spine_sql"), {"hadms": hadm_list})
print("ED spine rows:", len(ed_spine), "unique ed_stay_id:",
↪    ed_spine["ed_stay_id"].nunique())

# ensure uniqueness
if ed_spine["ed_stay_id"].nunique() != len(ed_spine):
    raise ValueError("ed_stay_id not unique in ED spine")

# First ED presentation time per admission
ed_intime_first = (
    ed_spine.groupby("hadm_id", as_index=False)["ed_intime"]
    .min()
    .rename(columns={"ed_intime": "ed_intime_first"})
)

# Start ED-level df
ed_df = ed_spine.copy()
ed_df = ed_df.merge(ed_intime_first, on="hadm_id", how="left")
```

```
ED spine rows: 28026 unique ed_stay_id: 28026


# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪   consistent across notebook stages.

# ED triage and first ED vitals (reuse existing logic if present; otherwise
↪   join)

def _needs_cols(df, cols):
    return (df is None) or any(c not in df.columns for c in cols)

# Use existing ed_triage / ed_first if already in memory from earlier cells
try:
    _ = ed_triage
except NameError:
    ed_triage = None

try:
    _ = ed_first
except NameError:
    ed_first = None

# Force re-query if required keys are missing
if _needs_cols(locals().get('ed_triage', None), ['ed_stay_id', 'hadm_id']):
    ed_triage = None
if _needs_cols(locals().get('ed_first', None), ['ed_stay_id']):
    ed_first = None

# If missing, prefer in-memory tables loaded earlier in the notebook.
# This avoids re-running expensive ED queries during nbconvert.
if ed_triage is None:
    if 'edmap' in globals() and 'tri_all' in globals():
        ed_triage = (
            edmap[['stay_id', 'hadm_id']]
            .merge(tri_all, on='stay_id', how='left')
            .rename(columns={'stay_id': 'ed_stay_id'})
        )
        print('ED triage rows (from in-memory tables):', len(ed_triage))
    else:
        ed_triage_sql = f"""
        WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
        SELECT
          s.stay_id AS ed_stay_id,
```

```
            s.hadm_id,
            t.temperature    AS ed_triage_temp,
            t.heartrate      AS ed_triage_hr,
            t.resprate       AS ed_triage_rr,
            t.o2sat          AS ed_triage_o2sat,
            t.sbp            AS ed_triage_sbp,
            t.dbp            AS ed_triage_dbp,
            t.pain           AS ed_triage_pain,
            t.acuity         AS ed_triage_acuity,
            t.chiefcomplaint AS ed_triage_cc
        FROM `{PHYS}.{ED}.edstays` s
        JOIN hadms h ON h.hadm_id = s.hadm_id
        LEFT JOIN `{PHYS}.{ED}.triage` t ON t.stay_id = s.stay_id
        """
        ed_triage = run_sql_bq(ed_triage_sql, {'hadms': hadm_list})
        print('ED triage rows (fallback query):', len(ed_triage))

if ed_first is None:
    if 'edmap' in globals() and 'first_stay_all' in globals():
        ed_first = (
            edmap[['stay_id', 'hadm_id']]
            .merge(first_stay_all, on='stay_id', how='left')
            .rename(columns={'stay_id': 'ed_stay_id'})
        )
        print('ED first vitals rows (from in-memory tables):', len(ed_first))
    else:
        ed_first_vitals_sql = f"""
        WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
        edmap AS (
          SELECT stay_id, hadm_id
          FROM `{PHYS}.{ED}.edstays`
          WHERE hadm_id IN (SELECT hadm_id FROM hadms)
        ),
        vs AS (
          SELECT stay_id, charttime, temperature, heartrate, resprate, o2sat,
↪  sbp, dbp, rhythm, pain
          FROM `{PHYS}.{ED}.vitalsign`
        ),
        first_vs AS (
          SELECT
            v.stay_id,
            (ARRAY_AGG(STRUCT(v.charttime, v.temperature, v.heartrate,
↪  v.resprate, v.o2sat, v.sbp, v.dbp, v.rhythm, v.pain)
```

```
                        ORDER BY v.charttime LIMIT 1))[OFFSET(0)] AS pick
          FROM vs v
          JOIN edmap m USING (stay_id)
          GROUP BY v.stay_id
        )
        SELECT
          f.stay_id AS ed_stay_id,
          m.hadm_id,
          pick.charttime AS ed_first_vitals_time,
          pick.temperature AS ed_first_temp,
          pick.heartrate AS ed_first_hr,
          pick.resprate AS ed_first_rr,
          pick.o2sat AS ed_first_o2sat,
          pick.sbp AS ed_first_sbp,
          pick.dbp AS ed_first_dbp,
          pick.rhythm AS ed_first_rhythm,
          pick.pain AS ed_first_pain
        FROM first_vs f
        JOIN edmap m ON m.stay_id = f.stay_id
        """

        ed_first = run_sql_bq(ed_first_vitals_sql, {'hadms': hadm_list})
        print('ED first vitals rows (fallback query):', len(ed_first))

# Debug columns before merge
print('ed_triage cols:', list(ed_triage.columns))
print('ed_first cols:', list(ed_first.columns))
print('ed_df cols:', list(ed_df.columns))

if 'ed_stay_id' not in ed_df.columns:
    raise KeyError('ed_df missing ed_stay_id; ensure ED spine cell ran.')

# Merge ED triage + vitals onto ed_df with available keys
merge_keys_triage = [k for k in ["ed_stay_id", "hadm_id"] if k in
 ↪  ed_df.columns and k in ed_triage.columns]
if not merge_keys_triage:
    raise KeyError("No common keys between ed_df and ed_triage")
ed_df = ed_df.merge(ed_triage, on=merge_keys_triage, how="left")
merge_keys_first = [k for k in ["ed_stay_id", "hadm_id"] if k in
 ↪  ed_df.columns and k in ed_first.columns]
if not merge_keys_first:
    raise KeyError("No common keys between ed_df and ed_first")
ed_df = ed_df.merge(ed_first, on=merge_keys_first, how="left")
```

```
ED triage rows (from in-memory tables): 28026
ED first vitals rows (from in-memory tables): 28026
ed_triage cols: ['ed_stay_id', 'hadm_id', 'ed_triage_temp', 'ed_triage_hr', 'ed_triage_rr',
ed_first cols: ['ed_stay_id', 'hadm_id', 'ed_first_vitals_time', 'ed_first_temp', 'ed_first_
ed_df cols: ['ed_stay_id', 'subject_id', 'hadm_id', 'ed_intime', 'ed_outtime', 'arrival_trans
```

3.0.3 Phase 2 — Hospital admission context and outcomes

Rationale: Add admission-level outcomes and demographics for downstream stratification
and analysis.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.

# Admissions fields
SQL["admit_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
  a.hadm_id,
  a.admittime,
  a.dischtime,
  a.deathtime,
  a.hospital_expire_flag,
  a.admission_type,
  a.admission_location,
  a.discharge_location,
  a.insurance,
  a.language,
  a.marital_status,
  a.race AS hosp_race
FROM `{PHYS}.{HOSP}.admissions` a
JOIN hadms h USING (hadm_id)
"""

admit = run_sql_bq(sql("admit_sql"), {"hadms": hadm_list})
print("Admissions rows:", len(admit))

ed_df = ed_df.merge(admit, on="hadm_id", how="left")

# Merge ventilation flags/times (hadm-level)
if "vent_combined" in globals():
    ed_df = ed_df.merge(vent_combined, on="hadm_id", how="left")
```

```
# Derived outcomes
ed_df["hospital_los_hours"] = (ed_df["dischtime"] -
↪  ed_df["admittime"]).dt.total_seconds() / 3600.0
ed_df["in_hospital_death"] = ((ed_df["hospital_expire_flag"] == 1) |
↪  ed_df["deathtime"].notna()).astype("int64")

# Concordance check
discord = (
    ((ed_df["hospital_expire_flag"] == 1) & ed_df["deathtime"].isna()) |
    ((ed_df["hospital_expire_flag"] == 0) & ed_df["deathtime"].notna())
)
print("Admissions discordance (expire_flag vs deathtime):",
↪  int(discord.sum()))


Admissions rows: 84594
Admissions discordance (expire_flag vs deathtime): 4
```

3.0.4 Phase 3 — ICU timing and LOS

Rationale: Capture ICU exposure, timing, and total LOS for severity stratification.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.

# ICU stays (aggregate per hadm)
SQL["icu_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
  i.hadm_id,
  i.stay_id AS icu_stay_id,
  i.intime,
  i.outtime,
  i.los,
  i.first_careunit,
  i.last_careunit
FROM `{PHYS}.{ICU}.icustays` i
JOIN hadms h USING (hadm_id)
"""

icu = run_sql_bq(sql("icu_sql"), {"hadms": hadm_list})
print("ICU stay rows:", len(icu))

if len(icu) > 0:
```

```python
    icu_agg = (
        icu.sort_values(["hadm_id", "intime"]).groupby("hadm_id",
↪   as_index=False)
        .agg(
            icu_intime_first=("intime", "min"),
            icu_outtime_last=("outtime", "max"),
            icu_los_total=("los", "sum"),
            n_icu_stays=("icu_stay_id", "nunique"),
            first_careunit=("first_careunit", "first"),
            last_careunit=("last_careunit", "last"),
        )
    )
    ed_df = ed_df.merge(icu_agg, on="hadm_id", how="left")
else:
    print("No ICU stays found for cohort.")


ICU stay rows: 62908
```

3.0.5 Phase 4 — ED longitudinal vitals (0–6h)

Rationale: Summarize early ED vitals for severity phenotyping.

```python
# Purpose: Build ED vitals features robustly without a single large BigQuery
↪   join that can stall.

# ED vitals long + aggregates (0-6h)

# Build the ED stay map from in-memory cohort tables to keep query scope
↪   tight and deterministic.
if "ed_df" not in globals():
    raise NameError("ed_df is required before ED vitals extraction")

required_cols = ["ed_stay_id", "hadm_id", "ed_intime"]
missing_required = [c for c in required_cols if c not in ed_df.columns]
if missing_required:
    raise KeyError(f"ed_df missing required columns for ED vitals extraction:
    ↪   {missing_required}")

edmap_local = (
    ed_df[required_cols]
    .dropna(subset=["ed_stay_id"])
    .drop_duplicates(subset=["ed_stay_id"])
```

```python
        .copy()
    )
edmap_local["ed_stay_id"] = pd.to_numeric(edmap_local["ed_stay_id"],
 ↪  errors="coerce")
edmap_local = edmap_local.dropna(subset=["ed_stay_id"])
edmap_local["ed_stay_id"] = edmap_local["ed_stay_id"].astype(int)

stay_ids = sorted(edmap_local["ed_stay_id"].unique().tolist())
print("ED stays for vitals pull:", len(stay_ids))

SQL["ed_vitals_sql"] = f"""
SELECT
  stay_id AS ed_stay_id,
  charttime,
  temperature,
  heartrate,
  resprate,
  o2sat,
  sbp,
  dbp,
  rhythm,
  pain
FROM `{PHYS}.{ED}.vitalsign`
WHERE stay_id IN UNNEST(@stay_ids)
"""

# Query in chunks so we avoid one very large parameter payload/job.
chunk_size = 5000
vitals_chunks = []
for i in range(0, len(stay_ids), chunk_size):
    chunk = stay_ids[i:i + chunk_size]
    part = run_sql_bq(sql("ed_vitals_sql"), {"stay_ids": chunk})
    if len(part) > 0:
        part["ed_stay_id"] = pd.to_numeric(part["ed_stay_id"],
 ↪  errors="coerce")
        part = part.dropna(subset=["ed_stay_id"])
        part["ed_stay_id"] = part["ed_stay_id"].astype(int)
        vitals_chunks.append(part)
    print(f"ED vitals chunk {i // chunk_size + 1}: rows={len(part)}")

if vitals_chunks:
    ed_vitals_long = pd.concat(vitals_chunks, ignore_index=True)
else:
```

```python
    ed_vitals_long = pd.DataFrame(
        columns=[
            "ed_stay_id", "charttime", "temperature", "heartrate",
             ↪  "resprate",
            "o2sat", "sbp", "dbp", "rhythm", "pain"
        ]
    )

ed_vitals_long = ed_vitals_long.merge(edmap_local, on="ed_stay_id",
 ↪  how="inner")
print("ED vitals long rows:", len(ed_vitals_long))

# Window filter: 0-6h from ED intime
ed_vitals_long["dt_hours"] = (ed_vitals_long["charttime"] -
 ↪  ed_vitals_long["ed_intime"]).dt.total_seconds() / 3600.0
in_6h = ed_vitals_long["dt_hours"].between(0, 6, inclusive="both")

agg = (
    ed_vitals_long.loc[in_6h]
    .groupby("ed_stay_id", as_index=False)
    .agg(
        max_heartrate_0_6h=("heartrate", "max"),
        max_resprate_0_6h=("resprate", "max"),
        min_o2sat_0_6h=("o2sat", "min"),
        min_sbp_0_6h=("sbp", "min"),
        n_vitals_0_6h=("charttime", "count"),
    )
)

ed_df = ed_df.merge(agg, on="ed_stay_id", how="left")

# Range warnings (report only, do not drop)
range_checks = {
    "heartrate": (0, 300),
    "resprate": (0, 80),
    "o2sat": (0, 100),
    "sbp": (0, 300),
}
for col, (lo, hi) in range_checks.items():
    bad = ed_vitals_long[col].notna() & (~ed_vitals_long[col].between(lo,
 ↪  hi))
    if bad.any():
        print(f"Warning: {col} out of range count:", int(bad.sum()))
```

```
ED stays for vitals pull: 28026
ED vitals chunk 1: rows=31314
ED vitals chunk 2: rows=30451
ED vitals chunk 3: rows=31393
ED vitals chunk 4: rows=31420
ED vitals chunk 5: rows=30885
ED vitals chunk 6: rows=18799
ED vitals long rows: 174262
Warning: heartrate out of range count: 1
Warning: resprate out of range count: 4
Warning: o2sat out of range count: 3
```

3.0.6 Phase 5 — Robust lab discovery + gas panels

Rationale: Capture blood gas and key chemistry labs with label-robust item discovery and unit normalization.

```python
# Purpose: Discover lab item IDs robustly and record cohort-aware item
↪   frequency for auditability.

import re
import json

# Discover itemids from d_labitems
SQL["labitems_sql"] = f"""
SELECT itemid, label, fluid, category
FROM `{PHYS}.{HOSP}.d_labitems`
"""

labitems = run_sql_bq(sql("labitems_sql"))

patterns = {
    "gas_pco2": re.compile(r"\bp\s*co2\b|pco2|pco ", re.I),
    "gas_ph": re.compile(r"\bph\b", re.I),
    "gas_hco3": re.compile(r"hco3|bicarbonate", re.I),
    "gas_lactate": re.compile(r"lactate", re.I),
    "gas_specimen": re.compile(r"specimen|source|type", re.I),
    "chem_creatinine": re.compile(r"creatinine", re.I),
    "chem_sodium": re.compile(r"\bsodium\b", re.I),
    "chem_chloride": re.compile(r"\bchloride\b", re.I),
    "chem_total_co2": re.compile(r"carbon dioxide|total co2|\bco2\b", re.I),
    "cbc_hemoglobin": re.compile(r"hemoglobin", re.I),
}
```

```python
# Category filters limit false matches (for example chemistry CO2 vs
↪  blood-gas pCO2)
cat_gas = re.compile(r"blood\s*gas|blood gas|arterial|venous", re.I)
cat_chem = re.compile(r"chemistry|chem|blood", re.I)
cat_cbc = re.compile(r"hematology|cbc", re.I)

matches = {}
for name, pat in patterns.items():
    dfm = labitems.copy()
    dfm = dfm[dfm["label"].str.contains(pat, na=False)]
    if name.startswith("gas_"):
        dfm = dfm[dfm["category"].str.contains(cat_gas, na=False)]
    elif name.startswith("chem_"):
        dfm = dfm[dfm["category"].str.contains(cat_chem, na=False)]
    elif name.startswith("cbc_"):
        dfm = dfm[dfm["category"].str.contains(cat_cbc, na=False)]
    matches[name] = dfm[["itemid", "label", "category"]]

# Build lab_item_map with cohort counts. Run hadm counts in chunks to avoid
↪  one oversized query.
itemids_all = sorted({int(i) for dfm in matches.values() for i in
↪  dfm["itemid"].tolist()})

SQL["counts_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT itemid, COUNT(*) AS n
FROM `{PHYS}.{HOSP}.labevents`
WHERE hadm_id IN (SELECT hadm_id FROM hadms)
  AND itemid IN UNNEST(@itemids)
GROUP BY itemid
"""

counts_cols = ["itemid", "n"]
if itemids_all and len(hadm_list) > 0:
    hadm_chunk_size = 10000
    count_parts = []
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
        part = run_sql_bq(sql("counts_sql"), {"hadms": hadm_chunk, "itemids":
↪  itemids_all})
        if len(part) > 0:
            count_parts.append(part)
```

```python
        print(f"Lab item count chunk {i // hadm_chunk_size + 1}:
         ↪  rows={len(part)}")

    if count_parts:
        counts = (
            pd.concat(count_parts, ignore_index=True)
            .groupby("itemid", as_index=False)["n"]
            .sum()
        )
    else:
        counts = pd.DataFrame(columns=counts_cols)
else:
    counts = pd.DataFrame(columns=counts_cols)

lab_item_map = {}
for name, dfm in matches.items():
    tmp = dfm.merge(counts, on="itemid", how="left").fillna({"n": 0})
    lab_item_map[name] = {
        "pattern": patterns[name].pattern,
        "items": tmp.sort_values("n",
         ↪  ascending=False).to_dict(orient="records"),
    }

lab_item_map_path = WORK_DIR / "lab_item_map.json"
lab_item_map_path.write_text(json.dumps(lab_item_map, indent=2))
print("Wrote:", lab_item_map_path)


Lab item count chunk 1: rows=41
Lab item count chunk 2: rows=39
Lab item count chunk 3: rows=39
Lab item count chunk 4: rows=41
Lab item count chunk 5: rows=39
Lab item count chunk 6: rows=39
Lab item count chunk 7: rows=39
Lab item count chunk 8: rows=40
Lab item count chunk 9: rows=37
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/lab_item_map.json


# Purpose: Extract cohort labs and apply ED-time windows in pandas for
 ↪  predictable runtime.
```

```python
# Extract labevents and apply ED 0-24h windows locally

# Assemble itemid lists
itemid_sets = {k: [int(x["itemid"]) for x in v["items"]] for k, v in
↪   lab_item_map.items()}
all_itemids = sorted({i for v in itemid_sets.values() for i in v})

if "ed_df" not in globals():
    raise NameError("ed_df is required before lab window extraction")

ed_windows = (
    ed_df[["ed_stay_id", "hadm_id", "ed_intime"]]
    .dropna(subset=["ed_stay_id", "hadm_id", "ed_intime"])
    .drop_duplicates()
    .copy()
)
ed_windows["hadm_id"] = pd.to_numeric(ed_windows["hadm_id"],
↪   errors="coerce").astype("Int64")
ed_windows = ed_windows.dropna(subset=["hadm_id"])
ed_windows["hadm_id"] = ed_windows["hadm_id"].astype(int)

SQL["labs_sql"] = f"""
SELECT
  subject_id,
  hadm_id,
  itemid,
  charttime,
  specimen_id,
  valuenum,
  valueuom,
  value AS value_text
FROM `{PHYS}.{HOSP}.labevents`
WHERE hadm_id IN UNNEST(@hadms)
  AND itemid IN UNNEST(@itemids)
"""

# Run in hadm chunks; merge to ED windows in pandas to keep each query
↪   lightweight.
labs_parts = []
if all_itemids and len(hadm_list) > 0:
    hadm_chunk_size = 5000
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
```

```python
        part = run_sql_bq(sql("labs_sql"), {"hadms": hadm_chunk, "itemids":
↪  all_itemids})
        raw_n = len(part)

        if raw_n > 0:
            part["hadm_id"] = pd.to_numeric(part["hadm_id"], errors="coerce")
            part = part.dropna(subset=["hadm_id"])
            part["hadm_id"] = part["hadm_id"].astype(int)

            part = part.merge(ed_windows, on="hadm_id", how="inner")
            in_window = (
                (part["charttime"] >= part["ed_intime"]) &
                (part["charttime"] <= (part["ed_intime"] +
↪  pd.Timedelta(hours=24)))
            )
            part = part.loc[in_window, [
                "ed_stay_id", "subject_id", "hadm_id", "itemid", "charttime",
                  ↪  "specimen_id", "valuenum", "valueuom", "value_text"
            ]]
            if len(part) > 0:
                labs_parts.append(part)

        print(f"Labs window chunk {i // hadm_chunk_size + 1}: raw={raw_n},
          ↪  kept={len(part) if raw_n > 0 else 0}")

if labs_parts:
    labs_long = pd.concat(labs_parts, ignore_index=True)
else:
    labs_long = pd.DataFrame(
        columns=["ed_stay_id", "subject_id", "hadm_id", "itemid",
  ↪  "charttime", "specimen_id", "valuenum", "valueuom", "value_text"]
    )

print("Labs long rows:", len(labs_long))

# Unit audit for pCO2
pco2_ids = itemid_sets.get("gas_pco2", [])
unit_audit = (
    labs_long.loc[labs_long["itemid"].isin(pco2_ids)]
    .groupby("valueuom", dropna=False)
    .size().reset_index(name="n")
)
unit_audit_path = WORK_DIR / "lab_unit_audit.csv"
```

```python
unit_audit.to_csv(unit_audit_path, index=False)
print("Wrote:", unit_audit_path)
```

Labs window chunk 1: raw=467203, kept=24247
Labs window chunk 2: raw=479072, kept=24272
Labs window chunk 3: raw=471202, kept=25225
Labs window chunk 4: raw=490350, kept=24893
Labs window chunk 5: raw=469123, kept=25182
Labs window chunk 6: raw=494959, kept=26200
Labs window chunk 7: raw=468077, kept=25135
Labs window chunk 8: raw=480764, kept=24858
Labs window chunk 9: raw=487740, kept=24174
Labs window chunk 10: raw=486042, kept=24087
Labs window chunk 11: raw=481435, kept=25533
Labs window chunk 12: raw=496301, kept=25123
Labs window chunk 13: raw=484743, kept=24528
Labs window chunk 14: raw=479159, kept=25312
Labs window chunk 15: raw=481069, kept=24031
Labs window chunk 16: raw=453474, kept=23140
Labs window chunk 17: raw=450146, kept=22988
Labs long rows: 418928
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/lab_unit_audit.csv

```python
# Purpose: Reconstruct gas panels and normalize units before ED-stay level
↪  aggregation.

# Reconstruct gas panels by specimen_id within ED stay

pco2_ids = set(itemid_sets.get("gas_pco2", []))
ph_ids = set(itemid_sets.get("gas_ph", []))
hco3_ids = set(itemid_sets.get("gas_hco3", []))
lact_ids = set(itemid_sets.get("gas_lactate", []))

labs = labs_long.copy()

# Convert pCO2 kPa to mmHg when needed
is_kpa = labs["valueuom"].astype(str).str.lower().str.contains("kpa",
↪  na=False)
mask_pco2 = labs["itemid"].isin(pco2_ids)
if mask_pco2.any() and is_kpa.any():
    labs.loc[mask_pco2 & is_kpa, "valuenum"] = labs.loc[mask_pco2 & is_kpa,
↪  "valuenum"] * 7.50062
```

```python
        labs.loc[mask_pco2 & is_kpa, "valueuom"] = "mmHg"

# panel by specimen_id
panel = (
    labs.groupby(["ed_stay_id","specimen_id"], as_index=False)
    .agg(panel_time=("charttime","min"))
)

# attach analytes

def pick_analyte(df, ids, name):
    tmp = df.loc[df["itemid"].isin(ids),
    ↪   ["ed_stay_id","specimen_id","valuenum"]]
    tmp = tmp.rename(columns={"valuenum": name})
    return tmp.groupby(["ed_stay_id","specimen_id"], as_index=False).first()

if pco2_ids:
    panel = panel.merge(pick_analyte(labs, pco2_ids, "pco2"),
    ↪   on=["ed_stay_id","specimen_id"], how="left")
if ph_ids:
    panel = panel.merge(pick_analyte(labs, ph_ids, "ph"),
    ↪   on=["ed_stay_id","specimen_id"], how="left")
if hco3_ids:
    panel = panel.merge(pick_analyte(labs, hco3_ids, "hco3"),
    ↪   on=["ed_stay_id","specimen_id"], how="left")
if lact_ids:
    panel = panel.merge(pick_analyte(labs, lact_ids, "lactate"),
    ↪   on=["ed_stay_id","specimen_id"], how="left")

# First panel per ED stay
first_panel = (
    panel.sort_values(["ed_stay_id","panel_time"]).groupby("ed_stay_id",
    ↪   as_index=False).first()
)

# 0-6h and 0-24h extrema
panel = panel.merge(ed_df[["ed_stay_id","ed_intime"]], on="ed_stay_id",
↪   how="left")
panel["dt_hours"] = (panel["panel_time"] -
↪   panel["ed_intime"]).dt.total_seconds() / 3600.0

# Ensure expected panel columns exist even if analyte is absent
for col in ["pco2", "ph", "hco3", "lactate"]:
```

```python
        if col not in panel.columns:
            panel[col] = pd.NA

p06 = panel.loc[panel["dt_hours"].between(0,6, inclusive="both")]
p24 = panel.loc[panel["dt_hours"].between(0,24, inclusive="both")]

agg06 = p06.groupby("ed_stay_id",
↪   as_index=False).agg(max_pco2_0_6h=("pco2","max"),
↪   min_ph_0_6h=("ph","min"))
agg24 = p24.groupby("ed_stay_id",
↪   as_index=False).agg(max_pco2_0_24h=("pco2","max"),
↪   min_ph_0_24h=("ph","min"))

ed_df = ed_df.merge(first_panel[["ed_stay_id","panel_time","pco2","ph",⌋
↪   "hco3","lactate"]].rename(
        columns={"panel_time":"first_gas_time","pco2":"first_pco2","ph":⌋
↪   "first_ph","hco3":"first_hco3","lactate":"first_lactate"}
), on="ed_stay_id", how="left")

ed_df = ed_df.merge(agg06, on="ed_stay_id", how="left")
ed_df = ed_df.merge(agg24, on="ed_stay_id", how="left")

# Purpose: Reconstruct gas panels and normalize units before ED-stay level
↪   aggregation.

# ABG vs VBG classification (specimen- and label-based inference)
gas_source_mode = os.getenv("COHORT_GAS_SOURCE_INFERENCE_MODE",
↪   "metadata_only").strip().lower()
panel, gas_source_audit = infer_panel_gas_source_metadata(
    panel,
    labs,
    labitems,
    specimen_source_itemids=itemid_sets.get("gas_specimen", []),
    pco2_itemids=itemid_sets.get("gas_pco2", []),
    mode=gas_source_mode,
)

fail_on_all_other_source = os.getenv("COHORT_FAIL_ON_ALL_OTHER_SOURCE",
↪   "1").strip() == "1"
if gas_source_audit["all_other_or_unknown"]:
    if fail_on_all_other_source:
        assert_gas_source_coverage(
            gas_source_audit,
```

```python
                fail_on_all_other_source=fail_on_all_other_source,
            )
        else:
            print(
                "WARNING: Gas source attribution collapsed to other/unknown; "
                "continuing because COHORT_FAIL_ON_ALL_OTHER_SOURCE=0."
            )

    warn_other_rate_threshold = float(os.getenv("COHORT_WARN_OTHER_RATE",
    ↪   "0.50"))
    fail_other_rate_raw = os.getenv("COHORT_FAIL_OTHER_RATE", "").strip()
    fail_other_rate_threshold = float(fail_other_rate_raw) if fail_other_rate_raw
    ↪   else None
    other_source_rate = float(gas_source_audit.get("source_rates",
    ↪   {}).get("other", 0.0))

    if fail_other_rate_threshold is not None and other_source_rate >=
    ↪   fail_other_rate_threshold:
        raise ValueError(
            "gas_source_other_rate exceeded configured fail threshold "
            f"{fail_other_rate_threshold:.3f}: {other_source_rate:.4f}."
        )
    if other_source_rate >= warn_other_rate_threshold:
        print(
            "WARNING: gas_source_other_rate exceeded warning threshold "
            f"{warn_other_rate_threshold:.3f}: {other_source_rate:.4f}."
        )

    print("Gas source rates:", gas_source_audit.get("source_rates", {}))
    print("Gas source inference tiers:", gas_source_audit.get("tier_rates", {}))
    print("Unresolved specimen IDs:",
    ↪   gas_source_audit.get("unresolved_specimen_id_count", 0))

    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    out_date = datetime.now().strftime("%Y-%m-%d")
    gas_source_audit_path = prior_runs_dir / f"{out_date} gas_source_audit.json"

    reduction_target = float(os.getenv("COHORT_OTHER_RELATIVE_REDUCTION_MIN",
    ↪   "0.10"))
    baseline_other_rate = None
    baseline_audit_path = None
```

```python
for candidate in sorted(prior_runs_dir.glob("* gas_source_audit.json"),
↪   reverse=True):
    if candidate.resolve() == gas_source_audit_path.resolve():
        continue
    try:
        baseline_payload = json.loads(candidate.read_text())
        baseline_other_rate = float(
            baseline_payload.get("source_rates", {}).get("other",
↪   float("nan"))
        )
        if math.isnan(baseline_other_rate):
            baseline_other_rate = None
            continue
        baseline_audit_path = candidate
        break
    except Exception:
        continue

if baseline_other_rate is None or baseline_other_rate <= 0:
    gas_source_audit["baseline_other_rate"] = None
    gas_source_audit["baseline_path"] = str(baseline_audit_path) if
↪   baseline_audit_path else None
    gas_source_audit["relative_reduction_vs_baseline"] = None
    print(
        "WARNING: Unable to evaluate gas_source_other_rate relative reduction
            ↪   "
        "because no valid baseline gas_source_audit was found."
    )
else:
    relative_reduction = (baseline_other_rate - other_source_rate) /
↪   baseline_other_rate
    gas_source_audit["baseline_other_rate"] = baseline_other_rate
    gas_source_audit["baseline_path"] = str(baseline_audit_path)
    gas_source_audit["relative_reduction_vs_baseline"] = relative_reduction
    gas_source_audit["relative_reduction_target"] = reduction_target
    print(
        "Gas source OTHER relative reduction vs baseline: "
        f"{relative_reduction:.4f} (target={reduction_target:.4f})"
    )
    if relative_reduction < reduction_target:
        message = (
            "gas_source_other_rate relative reduction target not met: "
```

```python
        f"baseline={baseline_other_rate:.4f},
          ↪  current={other_source_rate:.4f}, "
        f"reduction={relative_reduction:.4f},
          ↪  target={reduction_target:.4f}."
    )
    contract_mode = os.getenv("PIPELINE_CONTRACT_MODE",
 ↪  "fail").strip().lower()
    if contract_mode == "fail":
        raise ValueError(message)
    print(f"WARNING: {message}")

gas_source_audit_path.write_text(json.dumps(gas_source_audit, indent=2))
print("Wrote:", gas_source_audit_path)

# flags
panel["flag_abg_hypercapnia"] = ((panel["source"]=="arterial") &
 ↪  (panel["pco2"]>=45)).astype(int)
panel["flag_vbg_hypercapnia"] = ((panel["source"]=="venous") &
 ↪  (panel["pco2"]>=50)).astype(int)
panel["flag_other_hypercapnia"] = ((panel["source"]=="other") &
 ↪  (panel["pco2"]>=50)).astype(int)
panel["flag_any_gas_hypercapnia"] = ((panel["flag_abg_hypercapnia"]==1) |
 ↪  (panel["flag_vbg_hypercapnia"]==1) |
 ↪  (panel["flag_other_hypercapnia"]==1)).astype(int)

# per-stay unknown source rate
if len(panel) > 0:
    unk_rate = (
        panel.assign(_unk=(panel["source"].fillna("other") == "other"))
            .groupby("ed_stay_id", as_index=False)["_unk"].mean()
            .rename(columns={"_unk": "gas_source_other_rate"})
    )
else:
    unk_rate = panel[["ed_stay_id"]].drop_duplicates()
    unk_rate["gas_source_other_rate"] = 1.0

if len(panel) > 0:
    panel_for_diag = panel.copy()
    panel_for_diag["source_inference_tier"] = (
        panel_for_diag["source_inference_tier"]
        .astype("string")
        .fillna("fallback_other")
    )
```

```python
    panel_for_diag["source_hint_conflict"] = (
        panel_for_diag["source_hint_conflict"]
        .fillna(False)
        .astype(bool)
    )
    panel_for_diag["source"] = (
        panel_for_diag["source"].astype("string").fillna("other").str.lower()
    )
    panel_for_diag["source_resolved"] = panel_for_diag["source"].isin(
        {"arterial", "venous"}
    )
    panel_for_diag["tier_specimen_text"] =
↪   panel_for_diag["source_inference_tier"].eq(
        "specimen_text"
    )
    panel_for_diag["tier_label_fluid"] =
↪   panel_for_diag["source_inference_tier"].eq(
        "label_fluid"
    )
    panel_for_diag["tier_cluster_inheritance"] = panel_for_diag[
        "source_inference_tier"
    ].eq("cluster_inheritance")
    panel_for_diag["tier_panel_cooccurrence"] = panel_for_diag[
        "source_inference_tier"
    ].eq("panel_cooccurrence")
    panel_for_diag["tier_fallback_other"] =
↪   panel_for_diag["source_inference_tier"].eq(
        "fallback_other"
    )

    def _tier_mode(series: pd.Series) -> str:
        mode_values = series.mode(dropna=True)
        if mode_values.empty:
            return "fallback_other"
        return str(mode_values.iloc[0])

    source_diag = (
        panel_for_diag.groupby("ed_stay_id", as_index=False)
        .agg(
            gas_source_inference_primary_tier=("source_inference_tier",
↪   _tier_mode),
            gas_source_hint_conflict_rate=("source_hint_conflict", "mean"),
            gas_source_resolved_rate=("source_resolved", "mean"),
```

133

```python
        gas_source_tier_specimen_text_rate=("tier_specimen_text",
↪    "mean"),
            gas_source_tier_label_fluid_rate=("tier_label_fluid", "mean"),

↪    gas_source_tier_cluster_inheritance_rate=("tier_cluster_inheritance",
↪    "mean"),

↪    gas_source_tier_panel_cooccurrence_rate=("tier_panel_cooccurrence",
↪    "mean"),
            gas_source_tier_fallback_other_rate=("tier_fallback_other",
↪    "mean"),
        )
        .copy()
    )
else:
    source_diag = ed_df[["ed_stay_id"]].drop_duplicates().copy()
    source_diag["gas_source_inference_primary_tier"] = "fallback_other"
    source_diag["gas_source_hint_conflict_rate"] = 0.0
    source_diag["gas_source_resolved_rate"] = 0.0
    source_diag["gas_source_tier_specimen_text_rate"] = 0.0
    source_diag["gas_source_tier_label_fluid_rate"] = 0.0
    source_diag["gas_source_tier_cluster_inheritance_rate"] = 0.0
    source_diag["gas_source_tier_panel_cooccurrence_rate"] = 0.0
    source_diag["gas_source_tier_fallback_other_rate"] = 1.0

# collapse to ED stay flags
flags = panel.groupby("ed_stay_id", as_index=False).agg(
    flag_abg_hypercapnia=("flag_abg_hypercapnia","max"),
    flag_vbg_hypercapnia=("flag_vbg_hypercapnia","max"),
    flag_other_hypercapnia=("flag_other_hypercapnia","max"),
    flag_any_gas_hypercapnia=("flag_any_gas_hypercapnia","max"),
)

ed_df = ed_df.merge(flags, on="ed_stay_id", how="left")
ed_df = ed_df.merge(unk_rate, on="ed_stay_id", how="left")
ed_df = ed_df.merge(source_diag, on="ed_stay_id", how="left")
ed_df["gas_source_other_rate"] = ed_df["gas_source_other_rate"].fillna(1.0)


Gas source rates: {'other': 0.4554416743409502, 'venous': 0.2937318218859524, 'arterial': 0.2
Gas source inference tiers: {'fallback_other': 0.4554416743409502, 'specimen_text': 0.278891
Unresolved specimen IDs: 68700
Gas source OTHER relative reduction vs baseline: 0.1690 (target=0.0000)
```

```
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 gas_source_audit.json
```

3.0.7 Phase 5C — ICU POC blood gases (chartevents, optional)

Rationale: Capture ICU point-of-care gases if central lab labevents miss early measurements.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.

# Discover ICU POC itemids from d_items

SQL["ditems_sql"] = f"""
SELECT itemid, label, category
FROM `{PHYS}.{ICU}.d_items`
"""

ditems = run_sql_bq(sql("ditems_sql"))

icu_patterns = {
    "pco2": re.compile(r"(?:^|[^a-z])p(?:a)?\s*co(?:2|)(?:[^a-z]|$)", re.I),
    "ph": re.compile(r"\bph\b", re.I),
    "hco3": re.compile(r"hco3|bicarbonate", re.I),
    "lactate": re.compile(r"lactate", re.I),
    "specimen": re.compile(r"specimen|source|type", re.I),
}

icu_cat = re.compile(r"blood\s*gas|blood gas|resp|arterial|venous", re.I)
icu_pco2_exclusion = re.compile(
    r"(?:^|[^a-z])p(?:a)?\s*o\s*2(?:[^a-z]|$)|\bpo2\b|\bpao2\b|oxygen|o2\s*⌋
    ↪  (?:sat|saturation)|sao2|spo2|fio2|co2\s*(?:prod|production|elimin|⌋
    ↪  elimination)|\bvco2\b",
    re.I,
)

icu_matches = {}
for name, pat in icu_patterns.items():
    dfm = ditems[ditems["label"].str.contains(pat, na=False)]
    dfm = dfm[dfm["category"].str.contains(icu_cat, na=False)]
    icu_matches[name] = dfm[["itemid","label","category"]]
# Exclude non-blood pCO2 measurements from ICU candidates
if "pco2" in icu_matches and len(icu_matches["pco2"]) > 0:
    icu_matches["pco2"] = icu_matches["pco2"][
```

```python
            ~icu_matches["pco2"]["label"].str.contains(icu_pco2_exclusion,
            ↪  na=False)
    ]

icu_itemid_map_frames = [
    frame.assign(analyte=name)
    for name, frame in icu_matches.items()
    if not frame.empty
]
if not icu_itemid_map_frames:
    icu_itemid_map = pd.DataFrame(columns=["analyte", "itemid", "label",
↪  "category"])
else:
    icu_itemid_map = pd.concat(icu_itemid_map_frames, ignore_index=True)
    icu_itemid_map = (
        icu_itemid_map[["analyte", "itemid", "label", "category"]]
        .drop_duplicates()
        .sort_values(["analyte", "itemid"], kind="stable")
        .reset_index(drop=True)
    )

icu_itemids = sorted({int(i) for dfm in icu_matches.values() for i in
↪  dfm["itemid"].tolist()})
print("ICU POC candidate itemids:", len(icu_itemids))
if not icu_itemid_map.empty:
    print("ICU POC itemid map preview:")
    print(icu_itemid_map.head(20).to_string(index=False))

audit_date = pd.Timestamp.now().strftime("%Y-%m-%d")
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
icu_itemid_map_path = prior_runs_dir / f"{audit_date} icu_poc_itemid_map.csv"
icu_itemid_map.to_csv(icu_itemid_map_path, index=False)
print("Wrote:", icu_itemid_map_path)


ICU POC candidate itemids: 6
ICU POC itemid map preview:
 analyte  itemid              label    category
specimen  223836       Airway Type Respiratory
specimen  223848 Ventilator Type Respiratory
specimen  224372    Sputum Source Respiratory
specimen  224829 Trach Tube Type Respiratory
specimen  224832          ETT Type Respiratory
```

```
specimen    230146    Blocker Type Respiratory
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 icu_poc_itemid_map.csv


# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪    consistent across notebook stages.

# Extract ICU chartevents within ED 0-24h window for cohort ICU stays

if len(icu_itemids) == 0:
    icu_poc_long = pd.DataFrame()
    print("No ICU POC itemids found.")
else:
    # ensure icu stay ids available
    if 'icu' not in globals():
        raise NameError("ICU stays table 'icu' not found; run ICU phase
        ↪    first.")

    icu_stays = icu[["icu_stay_id","hadm_id","intime"]].copy()
    icu_stays = icu_stays.dropna(subset=["icu_stay_id"])

    icu_poc_sql = f"""
    WITH icu_stays AS (
      SELECT stay_id AS icu_stay_id, hadm_id
      FROM `{PHYS}.{ICU}.icustays`
      WHERE stay_id IN UNNEST(@icu_stay_ids)
    ),
    eds AS (
      SELECT stay_id AS ed_stay_id, hadm_id, intime AS ed_intime
      FROM `{PHYS}.{ED}.edstays`
      WHERE hadm_id IN (SELECT hadm_id FROM icu_stays)
    )
    SELECT
      s.icu_stay_id,
      e.ed_stay_id,
      e.hadm_id,
      ce.charttime,
      ce.itemid,
      ce.valuenum,
      ce.valueuom
    FROM `{PHYS}.{ICU}.chartevents` ce
    JOIN icu_stays s ON s.icu_stay_id = ce.stay_id
    JOIN eds e ON e.hadm_id = s.hadm_id
```

137

```
    WHERE ce.itemid IN UNNEST(@itemids)
      AND ce.charttime BETWEEN e.ed_intime AND TIMESTAMP_ADD(e.ed_intime,
↪  INTERVAL 24 HOUR)
    """

    icu_poc_long = run_sql_bq(icu_poc_sql, {"icu_stay_ids":
↪  icu_stays["icu_stay_id"].astype(int).tolist(), "itemids": icu_itemids})
    print("ICU POC long rows:", len(icu_poc_long))

# Build panels by 5-minute bins per ICU stay
if len(icu_poc_long) > 0:
    icu_poc_long["time_bin"] = icu_poc_long["charttime"].dt.floor("5min")

    def pick_from_ids(df, ids, name):
        tmp = df.loc[df["itemid"].isin(ids),
↪  ["icu_stay_id","time_bin","valuenum"]]
        tmp = tmp.rename(columns={"valuenum": name})
        return tmp.groupby(["icu_stay_id","time_bin"],
           ↪  as_index=False).first()

    pco2_ids = set(icu_matches.get("pco2", pd.DataFrame()).get("itemid",
↪  []).tolist())
    ph_ids = set(icu_matches.get("ph", pd.DataFrame()).get("itemid",
↪  []).tolist())
    hco3_ids = set(icu_matches.get("hco3", pd.DataFrame()).get("itemid",
↪  []).tolist())
    lact_ids = set(icu_matches.get("lactate", pd.DataFrame()).get("itemid",
↪  []).tolist())

    panel_poc = (
        icu_poc_long.groupby(["icu_stay_id","time_bin"], as_index=False)
        .agg(panel_time=("charttime","min"))
    )
    if pco2_ids:
        panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, pco2_ids,
↪  "pco2"), on=["icu_stay_id","time_bin"], how="left")
    if ph_ids:
        panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, ph_ids,
↪  "ph"), on=["icu_stay_id","time_bin"], how="left")
    if hco3_ids:
        panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, hco3_ids,
↪  "hco3"), on=["icu_stay_id","time_bin"], how="left")
    if lact_ids:
```

```python
        panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, lact_ids,
↪   "lactate"), on=["icu_stay_id","time_bin"], how="left")

        # map to ED stay via hadm_id
        panel_poc = panel_poc.merge(icu[["icu_stay_id","hadm_id"]],
↪   on="icu_stay_id", how="left")
        panel_poc = panel_poc.merge(ed_df[["ed_stay_id","hadm_id","ed_intime"]],
↪   on="hadm_id", how="left")
        panel_poc["dt_hours"] = (panel_poc["panel_time"] -
↪   panel_poc["ed_intime"]).dt.total_seconds() / 3600.0

        icu_itemid_usage = (
            icu_poc_long.groupby("itemid", as_index=False)
            .size()
            .rename(columns={"size": "row_count"})
        )
        icu_itemid_usage["itemid"] = pd.to_numeric(
            icu_itemid_usage["itemid"], errors="coerce"
        ).astype("Int64")
        if not icu_itemid_map.empty:
            icu_itemid_usage = icu_itemid_usage.merge(
                icu_itemid_map[["analyte", "itemid", "label",
↪   "category"]].drop_duplicates(),
                on="itemid",
                how="left",
            )
        icu_itemid_usage = icu_itemid_usage.sort_values(
            ["row_count", "itemid"], ascending=[False, True], kind="stable"
        ).reset_index(drop=True)
        icu_itemid_usage_path = prior_runs_dir / f"{audit_date}
↪   icu_poc_itemid_usage.csv"
        icu_itemid_usage.to_csv(icu_itemid_usage_path, index=False)
        print("Wrote:", icu_itemid_usage_path)
else:
    panel_poc = pd.DataFrame()

# Ensure expected panel_poc columns exist even if analyte is absent
for col in ["pco2", "ph", "hco3", "lactate"]:
    if col not in panel_poc.columns:
        panel_poc[col] = pd.NA

poc_flags = pd.DataFrame(columns=["ed_stay_id",
↪   "flag_any_gas_hypercapnia_poc"])
```

```python
if len(panel_poc) > 0:
    p24_poc = panel_poc.loc[panel_poc["dt_hours"].between(0,24,
    ↪  inclusive="both")].copy()
    p24_poc.loc[:, "flag_any_gas_hypercapnia_poc"] = (p24_poc["pco2"] >=
    ↪  50).astype("Int64")

    poc_flags = p24_poc.groupby("ed_stay_id", as_index=False).agg(
        flag_any_gas_hypercapnia_poc=("flag_any_gas_hypercapnia_poc","max")
    )

# Avoid duplicate column on re-run
if "flag_any_gas_hypercapnia_poc" in ed_df.columns:
    ed_df = ed_df.drop(columns=["flag_any_gas_hypercapnia_poc"])

if len(poc_flags) > 0:
    ed_df = ed_df.merge(poc_flags, on="ed_stay_id", how="left")

# incremental yield
base = ed_df.get("flag_any_gas_hypercapnia", pd.Series(0,
↪  index=ed_df.index)).fillna(0).astype(int)
poc = ed_df.get("flag_any_gas_hypercapnia_poc", pd.Series(0,
↪  index=ed_df.index)).fillna(0).astype(int)
inc = ((base == 0) & (poc == 1)).sum()
print("ICU POC incremental hypercapnia cases (ED stays):", int(inc))

# optional export
if len(panel_poc) > 0:
    from datetime import datetime

    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    poc_path = prior_runs_dir /
↪  f"gas_panels_poc_{datetime.now().strftime('%Y%m%d_%H%M%S')}.parquet"
    panel_poc.to_parquet(poc_path, index=False)
    print("Wrote:", poc_path)
```

```
ICU POC long rows: 78452
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 icu_poc_itemid_usage.csv
ICU POC incremental hypercapnia cases (ED stays): 0
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/gas_panels_poc_20260217_094347.parquet
```

3.0.8 Phase 6 — BMI/anthropometrics (OMR)

Rationale: Add BMI/height/weight closest to ED presentation when available.

```python
# Purpose: Attach closest anthropometrics with pre-ED preference and bounded
↪   post-ED fallback.

from datetime import datetime

omr_sql = f"""
SELECT subject_id, chartdate, result_name, result_value
FROM `{PHYS}.{HOSP}.omr`
WHERE LOWER(result_name) IN ('bmi','height','weight')
"""

allow_omr_query_failure = os.getenv("COHORT_ALLOW_OMR_QUERY_FAILURE",
↪   "0").strip() == "1"
try:
    omr_raw = run_sql_bq(omr_sql)
except Exception as exc:
    if allow_omr_query_failure:
        print("OMR query failed; continuing with empty OMR frame because
        ↪   COHORT_ALLOW_OMR_QUERY_FAILURE=1:", exc)
        omr_raw = pd.DataFrame(columns=["subject_id", "chartdate",
↪   "result_name", "result_value"])
    else:
        raise

print("OMR rows:", len(omr_raw))
omr_prepared = prepare_omr_records(omr_raw)
ed_df, omr_diagnostics = attach_closest_pre_ed_omr(ed_df, omr_prepared,
↪   window_days=365)
omr_diagnostics["source_rows_raw"] = int(len(omr_raw))
omr_diagnostics["prepared_rows"] = int(len(omr_prepared))

diag_table = (
    pd.DataFrame({"metric": list(omr_diagnostics.keys()), "value":
↪   list(omr_diagnostics.values())})
    .assign(value=lambda d: d["value"].astype(str))
)
print(diag_table)

prior_runs_dir = DATA_DIR / "prior runs"
```

```python
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")
omr_diag_path = prior_runs_dir / f"{out_date} omr_diagnostics.json"
omr_diag_path.write_text(json.dumps(omr_diagnostics, indent=2))
print("Wrote:", omr_diag_path)

# Optional charted anthropometric fallback (nearest-anytime selection).
use_charted_fallback = os.getenv("COHORT_ANTHRO_CHARTED_FALLBACK",
 ↪  "1").strip() == "1"
use_nearest_anytime = os.getenv("COHORT_ANTHRO_NEAREST_ANYTIME", "1").strip()
 ↪  == "1"
charted_anthro_diagnostics = {
    "enabled": bool(use_charted_fallback),
    "nearest_anytime": bool(use_nearest_anytime),
    "charted_rows_input": 0,
    "charted_rows_usable": 0,
    "filled_counts": {},
    "rows_with_any_charted_fill": 0,
}

if use_charted_fallback and len(hadm_list) > 0:
    charted_anthro_sql = f"""
    WITH hadms AS (
      SELECT x AS hadm_id
      FROM UNNEST(@hadms) AS x
    ),
    candidate_items AS (
      SELECT itemid, label, category
      FROM `{PHYS}.{ICU}.d_items`
      WHERE REGEXP_CONTAINS(LOWER(label),
 ↪  r'(height|weight|admission\\s+weight)')
    ),
    ce AS (
      SELECT
        ce.subject_id,
        ce.hadm_id,
        ce.charttime AS obs_time,
        ci.label,
        ce.valuenum,
        LOWER(REPLACE(COALESCE(ce.valueuom, ''), ' ', '')) AS valueuom
      FROM `{PHYS}.{ICU}.chartevents` ce
      JOIN hadms h
        ON ce.hadm_id = h.hadm_id
```

```
      JOIN candidate_items ci
        ON ce.itemid = ci.itemid
      WHERE ce.valuenum IS NOT NULL
    )
    SELECT
      subject_id,
      obs_time,
      CASE
        WHEN REGEXP_CONTAINS(LOWER(label), r'height') THEN 'height'
        WHEN REGEXP_CONTAINS(LOWER(label), r'weight') THEN 'weight'
        ELSE NULL
      END AS result_name,
      CASE
        WHEN REGEXP_CONTAINS(LOWER(label), r'height') THEN
          CASE
            WHEN valueuom IN ('cm', 'centimeter', 'centimeters') THEN
↪ valuenum
            WHEN valueuom IN ('m', 'meter', 'meters') THEN valuenum * 100.0
            WHEN valueuom IN ('in', 'inch', 'inches') THEN valuenum * 2.54
            ELSE valuenum
          END
        WHEN REGEXP_CONTAINS(LOWER(label), r'weight') THEN
          CASE
            WHEN valueuom IN ('kg', 'kilogram', 'kilograms') THEN valuenum
            WHEN valueuom IN ('lb', 'lbs', 'pound', 'pounds') THEN valuenum *
↪ 0.45359237
            WHEN valueuom IN ('oz', 'ounce', 'ounces') THEN valuenum *
↪ 0.0283495231
            ELSE valuenum
          END
        ELSE NULL
      END AS result_value_num,
      'icu_charted' AS source
    FROM ce
    WHERE REGEXP_CONTAINS(LOWER(label), r'(height|weight)')
    """
    charted_raw = run_sql_bq(charted_anthro_sql, {"hadms": hadm_list})
    ed_df, charted_anthro_diagnostics = attach_charted_anthro_fallback(
        ed_df,
        charted_raw,
        nearest_anytime=use_nearest_anytime,
    )
```

```python
    height_cm = pd.to_numeric(ed_df.get("height_closest_pre_ed"),
↪   errors="coerce")
    weight_kg = pd.to_numeric(ed_df.get("weight_closest_pre_ed"),
↪   errors="coerce")
    bmi_existing = pd.to_numeric(ed_df.get("bmi_closest_pre_ed"),
↪   errors="coerce")
    bmi_calc = weight_kg / (height_cm / 100.0) ** 2
    bmi_fill_mask = bmi_existing.isna() & bmi_calc.notna() & height_cm.gt(0)
    ed_df.loc[bmi_fill_mask, "bmi_closest_pre_ed"] = bmi_calc[bmi_fill_mask]
    charted_anthro_diagnostics["bmi_filled_from_height_weight"] =
↪   int(bmi_fill_mask.sum())
    charted_anthro_diagnostics["enabled"] = True
    charted_anthro_diagnostics["nearest_anytime"] = bool(use_nearest_anytime)

anthro_coverage_audit = build_anthro_coverage_audit(ed_df)
anthro_coverage_audit["omr"] = omr_diagnostics
anthro_coverage_audit["charted_fallback"] = charted_anthro_diagnostics
anthro_coverage_audit_path = prior_runs_dir / f"{out_date}
↪   anthropometrics_coverage_audit.json"
anthro_coverage_audit_path.write_text(json.dumps(anthro_coverage_audit,
↪   indent=2))
print("Wrote:", anthro_coverage_audit_path)

attached_non_null = omr_diagnostics.get("attached_non_null_counts", {})
attached_total = int(sum(int(v) for v in attached_non_null.values())) if
↪   attached_non_null else 0
subject_overlap = int(omr_diagnostics.get("subject_overlap_count", 0))
pre_window_rows = int(omr_diagnostics.get("pre_window_candidate_rows", 0))
post_window_rows = int(omr_diagnostics.get("post_window_candidate_rows", 0))
within_window_rows = int(omr_diagnostics.get("within_window_candidate_rows",
↪   0))
selected_tier_counts = omr_diagnostics.get("selected_tier_counts", {})
selected_tier_rates = omr_diagnostics.get("selected_tier_rates", {})
allow_empty_omr = os.getenv("COHORT_ALLOW_EMPTY_OMR", "0").strip() == "1"
fail_on_omr_attach_inconsistency = (
    os.getenv("COHORT_FAIL_ON_OMR_ATTACH_INCONSISTENCY", "1").strip() == "1"
)

omr_window_diagnostics = {
    "subject_overlap_count": subject_overlap,
    "pre_window_candidate_rows": pre_window_rows,
    "post_window_candidate_rows": post_window_rows,
```

```python
            "closest_absolute_candidate_rows":
            ↪  int(omr_diagnostics.get("closest_absolute_candidate_rows", 0)),
            "within_window_candidate_rows": within_window_rows,
            "attached_total_non_null_values": attached_total,
            "selected_tier_counts": selected_tier_counts,
            "selected_tier_rates": selected_tier_rates,
}

if len(omr_raw) > 0 and subject_overlap > 0:
    if pre_window_rows == 0 and post_window_rows > 0:
        print(
            "INFO: OMR subject overlap exists with no pre-ED candidates but "
            ↪  post-ED candidates within the 365-day window. "
            "Anthropometric fallback will use post-ED observations and mark "
            ↪  them as timing-uncertain."
        )
    elif pre_window_rows == 0 and post_window_rows == 0:
        print(
            "INFO: OMR subject overlap exists but no candidates are within "
            ↪  +/-365 days of ED arrival. "
            "Anthropometrics will remain missing for this run."
        )
    if attached_total == 0 and (pre_window_rows > 0 or post_window_rows > 0):
        msg = (
            "OMR attachment produced all-null outputs despite available "
            ↪  +/-365-day candidates. "
            "Set COHORT_ALLOW_EMPTY_OMR=1 or "
            ↪  COHORT_FAIL_ON_OMR_ATTACH_INCONSISTENCY=0 to bypass this "
            ↪  guard."
        )
        if allow_empty_omr or not fail_on_omr_attach_inconsistency:
            print("WARNING:", msg)
        else:
            raise ValueError(msg)


OMR rows: 46494
                              metric  \
0                        window_days
1                source_rows_prepared
2                  parsed_value_rows
3                            ed_rows
4            ed_rows_eligible_for_join
5                subject_overlap_count
```

```
6          candidate_rows_after_subject_join
7                          days_before_min
8                          days_before_max
9               nonnegative_candidate_rows
10               pre_window_candidate_rows
11              post_window_candidate_rows
12          closest_absolute_candidate_rows
13             within_window_candidate_rows
14          eligible_ed_stays_with_candidates
15                 attached_non_null_counts
16               attached_any_non_null_rows
17                     selected_tier_counts
18                      selected_tier_rates
19                   timing_uncertain_count
20                      anthro_source_counts
21                          source_rows_raw
22                            prepared_rows
```

```
                                                value
0                                                 365
1                                               46494
2                                               46489
3                                               28026
4                                               28026
5                                                1588
6                                                5670
7                                               -4043
8                                                -262
9                                                   0
10                                                  0
11                                                 81
12                                                 81
13                                                  0
14                                                 64
15  {'bmi_closest_pre_ed': 32, 'height_closest_pre...
16                                                 64
17  {'pre_ed_365': 0, 'post_ed_365': 64, 'missing'...
18  {'pre_ed_365': 0.0, 'post_ed_365': 0.002283593...
19                                                 64
20                       {'missing': 27962, 'omr': 64}
21                                               46494
22                                               46494
```
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects

3.0.9 Phase 7 — ICD comorbidity flags

Rationale: Derive comorbidity indicators from index admission ICD codes for stratified
analyses.

```
# Purpose: Compute comorbidity flags from hospital + ED ICD codes with
↪   chunked hadm queries.

# ICD code pulls for comorbidity flags (hospital + ED; combined OR)
# NOTE: Use prefix filters to reduce CPU and avoid regex-heavy scans.

FLAGS = [
    "flag_copd", "flag_osa_ohs", "flag_chf", "flag_neuromuscular",
    "flag_opioid_substance", "flag_pneumonia",
]

SQL["icd_flags_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
icd AS (
  SELECT
    hadm_id,
    UPPER(REPLACE(icd_code, ".", "")) AS code_norm
  FROM `{PHYS}.{HOSP}.diagnoses_icd`
  WHERE hadm_id IN (SELECT hadm_id FROM hadms)
    AND (
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J43") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J44") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G473") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "E662") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "I50") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G12") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G70") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G71") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "F11") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "T40") OR
      STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "F13") OR
```

```
          STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J12") OR
          STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J13") OR
          STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J14") OR
          STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J15") OR
          STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J16") OR
          STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J17") OR
          STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J18")
      )
)
SELECT
  hadm_id,
  MAX(IF(STARTS_WITH(code_norm, "J43") OR STARTS_WITH(code_norm, "J44"), 1,
↪  0)) AS flag_copd,
  MAX(IF(STARTS_WITH(code_norm, "G473") OR STARTS_WITH(code_norm, "E662"), 1,
↪  0)) AS flag_osa_ohs,
  MAX(IF(STARTS_WITH(code_norm, "I50"), 1, 0)) AS flag_chf,
  MAX(IF(STARTS_WITH(code_norm, "G12") OR STARTS_WITH(code_norm, "G70") OR
↪  STARTS_WITH(code_norm, "G71"), 1, 0)) AS flag_neuromuscular,
  MAX(IF(STARTS_WITH(code_norm, "F11") OR STARTS_WITH(code_norm, "T40") OR
↪  STARTS_WITH(code_norm, "F13"), 1, 0)) AS flag_opioid_substance,
  MAX(IF(STARTS_WITH(code_norm, "J12") OR STARTS_WITH(code_norm, "J13") OR
↪  STARTS_WITH(code_norm, "J14") OR STARTS_WITH(code_norm, "J15") OR
↪  STARTS_WITH(code_norm, "J16") OR STARTS_WITH(code_norm, "J17") OR
↪  STARTS_WITH(code_norm, "J18"), 1, 0)) AS flag_pneumonia
FROM icd
GROUP BY hadm_id
"""

SQL["ed_icd_flags_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
ed_dx AS (
  SELECT
    s.hadm_id,
    UPPER(REPLACE(d.icd_code, ".", "")) AS code_norm
  FROM `{PHYS}.{ED}.diagnosis` d
  JOIN `{PHYS}.{ED}.edstays` s
    ON s.stay_id = d.stay_id
  WHERE s.hadm_id IN (SELECT hadm_id FROM hadms)
    AND (
      STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J43") OR
      STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J44") OR
      STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G473") OR
      STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "E662") OR
```

```
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "I50") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G12") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G70") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G71") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "F11") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "T40") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "F13") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J12") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J13") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J14") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J15") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J16") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J17") OR
                STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J18")
        )
)
SELECT
  hadm_id,
  MAX(IF(STARTS_WITH(code_norm, "J43") OR STARTS_WITH(code_norm, "J44"), 1,
↪  0)) AS flag_copd,
  MAX(IF(STARTS_WITH(code_norm, "G473") OR STARTS_WITH(code_norm, "E662"), 1,
↪  0)) AS flag_osa_ohs,
  MAX(IF(STARTS_WITH(code_norm, "I50"), 1, 0)) AS flag_chf,
  MAX(IF(STARTS_WITH(code_norm, "G12") OR STARTS_WITH(code_norm, "G70") OR
↪  STARTS_WITH(code_norm, "G71"), 1, 0)) AS flag_neuromuscular,
  MAX(IF(STARTS_WITH(code_norm, "F11") OR STARTS_WITH(code_norm, "T40") OR
↪  STARTS_WITH(code_norm, "F13"), 1, 0)) AS flag_opioid_substance,
  MAX(IF(STARTS_WITH(code_norm, "J12") OR STARTS_WITH(code_norm, "J13") OR
↪  STARTS_WITH(code_norm, "J14") OR STARTS_WITH(code_norm, "J15") OR
↪  STARTS_WITH(code_norm, "J16") OR STARTS_WITH(code_norm, "J17") OR
↪  STARTS_WITH(code_norm, "J18"), 1, 0)) AS flag_pneumonia
FROM ed_dx
GROUP BY hadm_id
"""


def _run_flag_query_chunked(sql_name: str, label: str) -> pd.DataFrame:
    parts = []
    hadm_chunk_size = 10000
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
        part = run_sql_bq(sql(sql_name), {"hadms": hadm_chunk})
        if len(part) > 0:
            parts.append(part)
```

```python
        print(f"{label} chunk {i // hadm_chunk_size + 1}: rows={len(part)}")

    if parts:
        out = pd.concat(parts, ignore_index=True)
        out = out.groupby("hadm_id", as_index=False)[FLAGS].max()
        return out

    return pd.DataFrame(columns=["hadm_id"] + FLAGS)


flag_hosp = _run_flag_query_chunked("icd_flags_sql", "Hosp ICD flags")
flag_ed = _run_flag_query_chunked("ed_icd_flags_sql", "ED ICD flags")

flag_hosp = flag_hosp.rename(columns={k: f"{k}_hosp" for k in FLAGS})
flag_ed = flag_ed.rename(columns={k: f"{k}_ed" for k in FLAGS})

flag_df = flag_hosp.merge(flag_ed, on="hadm_id", how="outer")
for k in FLAGS:
    hosp_col = f"{k}_hosp"
    ed_col = f"{k}_ed"
    if hosp_col not in flag_df.columns:
        flag_df[hosp_col] = 0
    if ed_col not in flag_df.columns:
        flag_df[ed_col] = 0
    flag_df[k] = ((flag_df[hosp_col].fillna(0).astype(int) == 1) |
 ↪  (flag_df[ed_col].fillna(0).astype(int) == 1)).astype(int)

# Merge into admission-level and ED-stay-level frames (override if present)
for _df_name in ["df", "ed_df"]:
    if _df_name in globals():
        _df = globals()[_df_name]
        drop_cols = [c for c in flag_df.columns if c != "hadm_id" and c in
 ↪  _df.columns]
        if drop_cols:
            _df = _df.drop(columns=drop_cols)
        _df = _df.merge(flag_df, on="hadm_id", how="left")
        globals()[_df_name] = _df

# Prevalence (combined flags)
for k in FLAGS:
    if k in ed_df.columns:
        print(k, int(ed_df[k].fillna(0).sum()))


Hosp ICD flags chunk 1: rows=3142
```

```
Hosp ICD flags chunk 2: rows=3040
Hosp ICD flags chunk 3: rows=3103
Hosp ICD flags chunk 4: rows=3115
Hosp ICD flags chunk 5: rows=3133
Hosp ICD flags chunk 6: rows=3101
Hosp ICD flags chunk 7: rows=3066
Hosp ICD flags chunk 8: rows=2965
Hosp ICD flags chunk 9: rows=1427
ED ICD flags chunk 1: rows=289
ED ICD flags chunk 2: rows=282
ED ICD flags chunk 3: rows=313
ED ICD flags chunk 4: rows=324
ED ICD flags chunk 5: rows=324
ED ICD flags chunk 6: rows=295
ED ICD flags chunk 7: rows=336
ED ICD flags chunk 8: rows=255
ED ICD flags chunk 9: rows=151
flag_copd 3373
flag_osa_ohs 2334
flag_chf 5810
flag_neuromuscular 132
flag_opioid_substance 1053
flag_pneumonia 3490
```

3.0.10 Phase 8 — Timing phenotypes and derived bands

Rationale: Compute time-anchored hypercapnia/acidemia phenotypes for ED presentation vs later course.

```
# Purpose: Create time-anchored phenotypes and severity bands relative to
↪  first ED presentation.

# Timing phenotypes

anchor = "ed_intime_first" if "ed_intime_first" in ed_df.columns else
↪  "ed_intime"

ed_df["dt_first_qualifying_gas_hours"] = (ed_df["first_gas_time"] -
↪  ed_df[anchor]).dt.total_seconds() / 3600.0
ed_df["presenting_hypercapnia"] = (ed_df["dt_first_qualifying_gas_hours"] <=
↪  6).astype("Int64")
ed_df["late_hypercapnia"] = (ed_df["dt_first_qualifying_gas_hours"] >
↪  6).astype("Int64")
```

```python
# NIV/IMV timing relative to first ED presentation
if "first_imv_time" in ed_df.columns:
    ed_df["dt_first_imv_hours"] = (ed_df["first_imv_time"] -
↪   ed_df[anchor]).dt.total_seconds() / 3600.0
if "first_niv_time" in ed_df.columns:
    ed_df["dt_first_niv_hours"] = (ed_df["first_niv_time"] -
↪   ed_df[anchor]).dt.total_seconds() / 3600.0

# Bands
bins_ph = [-1, 7.20, 7.30, 7.35, 99]
labels_ph = ["<7.20","7.20-7.29","7.30-7.34"," 7.35"]
ed_df["ph_band"] = pd.cut(ed_df["first_ph"], bins=bins_ph, labels=labels_ph)

bins_hco3 = [-1, 24, 30, 999]
labels_hco3 = ["<24","24-29"," 30"]
ed_df["hco3_band"] = pd.cut(ed_df["first_hco3"], bins=bins_hco3,
↪   labels=labels_hco3)

bins_lac = [-1, 2, 4, 999]
labels_lac = ["<2","2-4",">4"]
ed_df["lactate_band"] = pd.cut(ed_df["first_lactate"], bins=bins_lac,
↪   labels=labels_lac)




# Align ED-stay gas source flags to hadm-level thresholds and ICU POC signal.
# This keeps source-specific flags internally consistent in the final ED-stay
↪   export.
if "df" in globals() and "hadm_id" in ed_df.columns and "hadm_id" in
↪   df.columns:
    hadm_thr_cols = [c for c in ["abg_hypercap_threshold",
↪   "vbg_hypercap_threshold", "other_hypercap_threshold"] if c in df.columns]
    if hadm_thr_cols:
        hadm_thr = df[["hadm_id"] + hadm_thr_cols].drop_duplicates("hadm_id")
        hadm_thr_map = hadm_thr.set_index("hadm_id")

        for col in hadm_thr_cols:
            mapped = pd.to_numeric(ed_df["hadm_id"].map(hadm_thr_map[col]),
↪   errors="coerce")
            existing = pd.to_numeric(ed_df[col], errors="coerce") if col in
↪   ed_df.columns else pd.Series(np.nan, index=ed_df.index)
            ed_df[col] = existing.fillna(mapped).fillna(0).astype("Int64")
```

```python
# Source-specific flags must include any matching hadm-level threshold.
flag_specs = [
    ("flag_abg_hypercapnia", "abg_hypercap_threshold"),
    ("flag_vbg_hypercapnia", "vbg_hypercap_threshold"),
    ("flag_other_hypercapnia", "other_hypercap_threshold"),
]
for flag_col, thr_col in flag_specs:
    base = pd.to_numeric(ed_df.get(flag_col, pd.Series(0,
 ↪  index=ed_df.index)), errors="coerce").fillna(0).astype(int)
    thr = pd.to_numeric(ed_df.get(thr_col, pd.Series(0, index=ed_df.index)),
 ↪  errors="coerce").fillna(0).astype(int)
    ed_df[flag_col] = ((base == 1) | (thr == 1)).astype("Int64")

any_from_sources = (
    (pd.to_numeric(ed_df.get("flag_abg_hypercapnia", 0),
 ↪  errors="coerce").fillna(0).astype(int) == 1)
    | (pd.to_numeric(ed_df.get("flag_vbg_hypercapnia", 0),
     ↪  errors="coerce").fillna(0).astype(int) == 1)
    | (pd.to_numeric(ed_df.get("flag_other_hypercapnia", 0),
     ↪  errors="coerce").fillna(0).astype(int) == 1)
)

poc_any = pd.to_numeric(ed_df.get("flag_any_gas_hypercapnia_poc",
 ↪  pd.Series(0, index=ed_df.index)), errors="coerce").fillna(0).astype(int)
 ↪  == 1
ed_df["flag_any_gas_hypercapnia"] = (any_from_sources |
 ↪  poc_any).astype("Int64")
```

3.1 QA & Data Fidelity

Rationale: Validate joins, missingness, lab completeness, and produce reproducibility arti-facts.

```python
# Purpose: Create reusable data-quality and merge guardrail helpers to
 ↪  prevent silent join errors.

# --- QA checks (lightweight, deterministic)
STRICT_QA = False

# 1) Key uniqueness
if 'ed_df' in globals() and 'ed_stay_id' in ed_df.columns:
    try:
```

```python
        assert_unique(ed_df, 'ed_stay_id', 'ed_df')
        print('QA: ed_stay_id unique OK')
    except Exception as e:
        print('QA FAIL:', e)
        if STRICT_QA:
            raise


# 2) Inclusion sanity: any_hypercap_icd==0 implies gas criteria met
if 'ed_df' in globals() and 'any_hypercap_icd' in ed_df.columns:
    gas_flag = ed_df.get('flag_any_gas_hypercapnia', pd.Series(0,
 ↪  index=ed_df.index)).fillna(0).astype(int)
    icd_flag = ed_df['any_hypercap_icd'].fillna(0).astype(int)
    viol = (icd_flag == 0) & (gas_flag == 0)
    n_viol = int(viol.sum())
    print('QA: ICD==0 & Gas==0 count:', n_viol)
    if STRICT_QA and n_viol > 0:
        raise AssertionError('Found rows without ICD and without gas
         ↪  criteria.')


# 3) Temporal sanity: first_gas_time >= ed_intime (allow small negative
 ↪  drift)
if 'ed_df' in globals() and 'first_gas_time' in ed_df.columns and 'ed_intime'
 ↪  in ed_df.columns:
    dt = (pd.to_datetime(ed_df['first_gas_time']) -
 ↪  pd.to_datetime(ed_df['ed_intime'])).dt.total_seconds() / 3600
    n_neg = int((dt < -1).sum())  # allow 1h clock drift
    print('QA: first_gas_time < ed_intime by >1h:', n_neg)
    if STRICT_QA and n_neg > 0:
        raise AssertionError('first_gas_time before ed_intime by >1h')


# 4) Range checks (report only)
if 'ed_df' in globals():
    ranges = {
        'first_ph': (6.8, 7.8),
        'first_pco2': (10, 200),
        'first_lactate': (0, 30),
        'creatinine': (0, 20),
    }
    rc = check_ranges(ed_df, ranges)
    if not rc.empty:
        print('QA range check (n out-of-range):')
        print(rc.to_string(index=False))
```

```
QA: ed_stay_id unique OK
QA: first_gas_time < ed_intime by >1h: 0
QA range check (n out-of-range):
          col   n_bad
     first_ph      12
   first_pco2       9
first_lactate       0
```

```python
# Purpose: Report missingness for key variables to support transparent
↪   data-quality reporting.

import json
from datetime import datetime

# T1 - uniqueness and join explosion guard
if ed_df["ed_stay_id"].nunique() != len(ed_df):
    raise ValueError("ed_stay_id not unique after merges")

# T1b - create cleaned vitals model fields while preserving raw values.
ed_df, vitals_outlier_audit = add_vitals_model_fields(ed_df)
ed_df, gas_outlier_audit = add_gas_model_fields(ed_df)
outliers_audit = pd.concat([vitals_outlier_audit, gas_outlier_audit],
↪   ignore_index=True, sort=False)

# T1c - fail-fast cleaned-vitals bounds checks.
clean_bounds = {
    "ed_triage_o2sat_clean": (0.0, 100.0),
    "ed_first_o2sat_clean": (0.0, 100.0),
    "ed_triage_sbp_clean": (20.0, 300.0),
    "ed_first_sbp_clean": (20.0, 300.0),
    "ed_triage_dbp_clean": (10.0, 200.0),
    "ed_first_dbp_clean": (10.0, 200.0),
}
for column_name, (lower_bound, upper_bound) in clean_bounds.items():
    if column_name not in ed_df.columns:
        continue
    numeric = pd.to_numeric(ed_df[column_name], errors="coerce")
    invalid_n = int((numeric.notna() & ((numeric < lower_bound) | (numeric >
↪   upper_bound))).sum())
    if invalid_n:
        raise ValueError(
            f"Cleaned vital bounds failure for {column_name}: {invalid_n}
↪   values outside [{lower_bound}, {upper_bound}]"
        )
```

```python
for temp_clean_column in ("ed_triage_temp_f_clean", "ed_first_temp_f_clean"):
    if temp_clean_column not in ed_df.columns:
        continue
    numeric = pd.to_numeric(ed_df[temp_clean_column], errors="coerce")
    residual_celsius_like_n = int(numeric.between(20.0, 50.0,
 ↪  inclusive="both").sum())
    if residual_celsius_like_n:
        raise ValueError(
            f"Found {residual_celsius_like_n} cleaned values still in 20-50°F
             ↪  band in {temp_clean_column}."
        )

ed_vitals_distribution_summary, ed_vitals_extreme_examples,
 ↪  ed_vitals_model_delta = build_ed_vitals_audit_artifacts(ed_df)

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")

vitals_outlier_audit_path = prior_runs_dir / f"{out_date}
 ↪  vitals_outlier_audit.csv"
vitals_outlier_audit.to_csv(vitals_outlier_audit_path, index=False)
print("Wrote:", vitals_outlier_audit_path)
outliers_audit_path = prior_runs_dir / f"{out_date} outliers_audit.csv"
outliers_audit.to_csv(outliers_audit_path, index=False)
print("Wrote:", outliers_audit_path)
ed_vitals_distribution_summary_path = prior_runs_dir / f"{out_date}
 ↪  ed_vitals_distribution_summary.csv"
ed_vitals_distribution_summary.to_csv(ed_vitals_distribution_summary_path,
 ↪  index=False)
print("Wrote:", ed_vitals_distribution_summary_path)
ed_vitals_extreme_examples_path = prior_runs_dir / f"{out_date}
 ↪  ed_vitals_extreme_examples.csv"
ed_vitals_extreme_examples.to_csv(ed_vitals_extreme_examples_path,
 ↪  index=False)
print("Wrote:", ed_vitals_extreme_examples_path)
ed_vitals_model_delta_path = prior_runs_dir / f"{out_date}
 ↪  ed_vitals_model_delta.csv"
ed_vitals_model_delta.to_csv(ed_vitals_model_delta_path, index=False)
print("Wrote:", ed_vitals_model_delta_path)

print("ED vitals cleaning summary (raw->clean removals):")
```

```python
display(
    ed_vitals_model_delta[
        [
            "vital_name",
            "raw_n",
            "clean_n",
            "raw_removed_n",
            "raw_removed_pct",
            "converted_celsius_like_n",
        ]
    ].sort_values("vital_name")
)

# T2 - missingness summary with expected-vs-unexpected classifications.
expected_sparse_fields = set()
omr_pre_window_rows = int(omr_diagnostics.get("pre_window_candidate_rows",
 ↪  0)) if "omr_diagnostics" in globals() else 0
omr_post_window_rows = int(omr_diagnostics.get("post_window_candidate_rows",
 ↪  0)) if "omr_diagnostics" in globals() else 0
if omr_pre_window_rows == 0 and omr_post_window_rows == 0:
    expected_sparse_fields.update({"bmi_closest_pre_ed",
 ↪  "height_closest_pre_ed", "weight_closest_pre_ed"})

miss = classify_missingness_expectations(
    ed_df,
    TARGET_FIELDS,
    expected_sparse_fields=expected_sparse_fields,
)
print(miss.sort_values(["missing_pct", "field"], ascending=[False,
 ↪  True]).head(30))

unexpected_full_null_fields = miss.loc[
    miss["expectation"].eq("unexpected_full_null"), "field"
].tolist()
expected_structural_null_fields = sorted(
    set(EXPECTED_STRUCTURAL_NULL_FIELDS).intersection(ed_df.columns)
)

# Compute UOM checks against canonical export frame contract (or
 ↪  deterministic preview).
if "final_cc" in globals():
    uom_scope_df = final_cc
    uom_scope_frame = "final_cc"
```

```python
elif "df" in globals() and "ed_triage_cc" in df.columns:
    _mask_df_cc = df["ed_triage_cc"].notna() &
↪   (df["ed_triage_cc"].astype(str).str.strip() != "")
    _df_cc = df.loc[_mask_df_cc].copy()
    _mask_cc = ed_df["ed_triage_cc"].notna() &
↪   (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
    _ed_cc_only = ed_df.loc[_mask_cc].copy()
    _ed_tmp = _ed_cc_only.copy()
    if "ed_intime" in _ed_tmp.columns:
        _ed_tmp = _ed_tmp.sort_values(["hadm_id", "ed_intime", "ed_stay_id"],
↪   na_position="last")
    else:
        _ed_tmp = _ed_tmp.sort_values(["hadm_id", "ed_stay_id"],
↪   na_position="last")
    _ed_first = _ed_tmp.drop_duplicates("hadm_id", keep="first")
    _add_ed_cols = [c for c in _ed_first.columns if c not in _df_cc.columns]
    uom_scope_df = _df_cc.merge(_ed_first[["hadm_id"] + _add_ed_cols],
↪   on="hadm_id", how="left")
    uom_scope_frame = "final_cc_preview"
else:
    uom_scope_df = ed_df
    uom_scope_frame = "ed_df"


uom_expectation_checks = evaluate_uom_expectations(uom_scope_df)
uom_expectation_scope = {"frame": uom_scope_frame, "row_count":
↪   int(len(uom_scope_df))}

# T2b - route-stratified first_other_pco2 audit using canonical QA scope.
first_other_scope_df = (
    uom_scope_df
    if {"first_other_src", "first_other_pco2"}.issubset(uom_scope_df.columns)
    else ed_df
)
first_other_scope_frame = "uom_scope_df" if first_other_scope_df is
↪   uom_scope_df else "ed_df"
first_other_pco2_audit = build_first_other_pco2_audit(first_other_scope_df)
first_other_pco2_audit_path = prior_runs_dir / f"{out_date}
↪   first_other_pco2_audit.csv"
first_other_pco2_audit.to_csv(first_other_pco2_audit_path, index=False)
print("Wrote:", first_other_pco2_audit_path)
print("first_other_pco2 audit scope:", first_other_scope_frame)


print("Expected structural null fields:", expected_structural_null_fields)
```

```python
print("Unexpected full-null fields (top 20):",
 ↪  unexpected_full_null_fields[:20])
print("UOM expectation scope:", uom_expectation_scope)
print("UOM expectation checks:", json.dumps(uom_expectation_checks,
 ↪  indent=2))


# T3 - lab capture completeness.
pct_any_6h = float(p06["ed_stay_id"].nunique() /
 ↪  max(ed_df["ed_stay_id"].nunique(),1))
pct_any_24h = float(p24["ed_stay_id"].nunique() /
 ↪  max(ed_df["ed_stay_id"].nunique(),1))
print("% any gas panel 0-6h:", round(pct_any_6h*100,1))
print("% any gas panel 0-24h:", round(pct_any_24h*100,1))


if "gas_source_audit" not in globals():
    gas_source_audit = summarize_gas_source(panel if "panel" in globals()
 ↪  else pd.DataFrame())


gas_overlap_summary = build_gas_source_overlap_summary(ed_df)
gas_overlap_summary_path = prior_runs_dir / f"{out_date}
 ↪  gas_source_overlap_summary.csv"
gas_overlap_summary.to_csv(gas_overlap_summary_path, index=False)
print("Wrote:", gas_overlap_summary_path)


gas_source_other_rate = float(gas_source_audit.get("source_rates",
 ↪  {}).get("other", 0.0))
source_unknown_rate = float(gas_source_audit.get("source_rates",
 ↪  {}).get("unknown", 0.0))
print("% source other (panel-level):", round(gas_source_other_rate * 100, 1))
print("% source unknown (panel-level):", round(source_unknown_rate * 100, 1))


hadm_rows = int(len(df)) if "df" in globals() else None
hadm_cc_rows = None
if "df" in globals() and "ed_triage_cc" in df.columns:
    _hadm_cc_mask = df["ed_triage_cc"].notna() &
 ↪  (df["ed_triage_cc"].astype(str).str.strip() != "")
    hadm_cc_rows = int(_hadm_cc_mask.sum())


uom_failed_checks = []
for uom_column, details in uom_expectation_checks.get("paco2_uom_checks",
 ↪  {}).items():
    if details.get("present") and not details.get("passes", True):
        uom_failed_checks.append(uom_column)
```

```python
qa_warnings = []
qa_infos = []
if unexpected_full_null_fields:
    qa_warnings.append("unexpected_full_null_fields_detected")
if gas_source_audit.get("all_other_or_unknown", False):
    qa_warnings.append("gas_source_all_other_or_unknown")
if omr_pre_window_rows == 0 and omr_post_window_rows > 0:
    qa_infos.append("omr_pre_window_empty_post_window_fallback_used")
if omr_pre_window_rows == 0 and omr_post_window_rows == 0:
    qa_infos.append("omr_pre_post_window_empty_expected_sparse")
if uom_failed_checks:
    qa_warnings.append("uom_checks_failed")
if "status" in first_other_pco2_audit.columns and
 ↪  (first_other_pco2_audit["status"] == "missing_columns").any():
    qa_warnings.append("first_other_pco2_audit_missing_columns")
qa_other_warn_threshold = float(os.getenv("COHORT_WARN_OTHER_RATE", "0.50"))
qa_other_fail_raw = os.getenv("COHORT_FAIL_OTHER_RATE", "").strip()
qa_other_fail_threshold = float(qa_other_fail_raw) if qa_other_fail_raw else
 ↪  None
if qa_other_fail_threshold is not None and gas_source_other_rate >=
 ↪  qa_other_fail_threshold:
    qa_warnings.append("gas_source_other_rate_fail_threshold_exceeded")
elif gas_source_other_rate >= qa_other_warn_threshold:
    qa_infos.append("gas_source_other_rate_high")

first_other_pco2_audit_records =
 ↪  json.loads(first_other_pco2_audit.to_json(orient="records"))
vitals_outlier_audit_records =
 ↪  json.loads(vitals_outlier_audit.to_json(orient="records"))
outliers_audit_records = json.loads(outliers_audit.to_json(orient="records"))
ed_vitals_distribution_summary_records =
 ↪  json.loads(ed_vitals_distribution_summary.to_json(orient="records"))
ed_vitals_model_delta_records =
 ↪  json.loads(ed_vitals_model_delta.to_json(orient="records"))
gas_overlap_summary_records =
 ↪  json.loads(gas_overlap_summary.to_json(orient="records")) if
 ↪  "gas_overlap_summary" in globals() else []

qa_status = "warning" if qa_warnings else "pass"

# T4 - QA summary artifact.
qa_summary = {
```

```python
    "ed_spine_rows": int(len(ed_df)),
    "ed_rows": int(len(ed_df)),
    "ed_unique": int(ed_df["ed_stay_id"].nunique()),
    "hadm_rows": hadm_rows,
    "hadm_cc_rows": hadm_cc_rows,
    "icu_link_rate": float(ed_df["icu_intime_first"].notna().mean()) if
    ↪   "icu_intime_first" in ed_df.columns else None,
    "pct_any_gas_0_6h": pct_any_6h,
    "pct_any_gas_0_24h": pct_any_24h,
    "gas_source_other_rate": gas_source_other_rate,
    "gas_source_other_rate_warn_threshold": qa_other_warn_threshold,
    "gas_source_other_rate_fail_threshold": qa_other_fail_threshold,
    "source_unknown_rate": source_unknown_rate,
    "gas_source_audit": gas_source_audit,
    "gas_source_overlap_summary": gas_overlap_summary_records,
    "omr_diagnostics": omr_diagnostics if "omr_diagnostics" in globals() else
    ↪   None,
    "omr_window_diagnostics": omr_window_diagnostics if
    ↪   "omr_window_diagnostics" in globals() else None,
    "anthro_timing_tier_counts": omr_diagnostics.get("selected_tier_counts",
    ↪   {}) if "omr_diagnostics" in globals() else {},
    "anthro_timing_tier_rates": omr_diagnostics.get("selected_tier_rates",
    ↪   {}) if "omr_diagnostics" in globals() else {},
    "anthro_coverage_audit": anthro_coverage_audit if "anthro_coverage_audit"
    ↪   in globals() else None,
    "charted_anthro_diagnostics": charted_anthro_diagnostics if
    ↪   "charted_anthro_diagnostics" in globals() else None,
    "unexpected_full_null_fields": unexpected_full_null_fields,
    "expected_structural_null_fields": expected_structural_null_fields,
    "uom_expectation_scope": uom_expectation_scope,
    "uom_expectation_checks": uom_expectation_checks,
    "first_other_pco2_audit_scope": first_other_scope_frame,
    "first_other_pco2_audit": first_other_pco2_audit_records,
    "vitals_outlier_audit": vitals_outlier_audit_records,
    "ed_vitals_distribution_summary": ed_vitals_distribution_summary_records,
    "ed_vitals_model_delta": ed_vitals_model_delta_records,
    "outliers_audit": outliers_audit_records,
    "ed_vitals_audit_paths": {
        "distribution_summary": str(ed_vitals_distribution_summary_path),
        "extreme_examples": str(ed_vitals_extreme_examples_path),
        "model_delta": str(ed_vitals_model_delta_path),
    },
    "qa_status": qa_status,
```

```python
    "qa_warnings": qa_warnings,
    "qa_infos": qa_infos,
    "run_metadata": RUN_METADATA if "RUN_METADATA" in globals() else None,
}
qa_path = WORK_DIR / "qa_summary.json"
qa_path.write_text(json.dumps(qa_summary, indent=2))
print("Wrote:", qa_path)
```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 vitals_outlier_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 outliers_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 ed_vitals_distribution_summary.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 ed_vitals_extreme_examples.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 ed_vitals_model_delta.csv
ED vitals cleaning summary (raw->clean removals):

	vital_name	raw_n	clean_n	raw_removed_n	raw_removed_pct	converted_celsius_like_n
7	first_dbp	22847	22827	20	0.000875	0
9	first_o2sat	22143	22141	2	0.000090	0
3	first_pain	17932	17932	0	0.000000	0
6	first_sbp	22847	22828	19	0.000832	0
1	first_temp	17666	17665	1	0.000057	104
5	triage_dbp	21958	21933	25	0.001139	0
8	triage_o2sat	21708	21704	4	0.000184	0
2	triage_pain	20566	19294	1272	0.061850	0
4	triage_sbp	22044	22032	12	0.000544	0
0	triage_temp	21105	21101	4	0.000190	48

	field	missing_n	missing_pct	expectation
0	abg_before_imv	28026	1.000000	missing_column
49	first_other_time	28026	1.000000	missing_column
72	icu_stay_id	28026	1.000000	missing_column
88	vbg_before_imv	28026	1.000000	missing_column
4	anthro_chartdate	27962	0.997716	conditional_sparse
5	anthro_days_offset	27962	0.997716	conditional_sparse
25	ed_first_rhythm	26503	0.945658	conditional_sparse

162

```
47            first_hco3      26484    0.944980    conditional_sparse
64             hco3_band      26484    0.944980    conditional_sparse
14             deathtime      24486    0.873689    conditional_sparse
81           max_pco2_0_6h    22038    0.786341    conditional_sparse
83            min_ph_0_6h     21796    0.777706    conditional_sparse
19       dt_first_niv_hours   20217    0.721366    conditional_sparse
54               flag_chf     17954    0.640619    conditional_sparse
55              flag_copd     17954    0.640619    conditional_sparse
56        flag_neuromuscular  17954    0.640619    conditional_sparse
57     flag_opioid_substance  17954    0.640619    conditional_sparse
58            flag_osa_ohs    17954    0.640619    conditional_sparse
60           flag_pneumonia   17954    0.640619    conditional_sparse
18       dt_first_imv_hours   17223    0.614537    conditional_sparse
13          bmi_closest_pre_ed 13605   0.485442    conditional_sparse
65       height_closest_pre_ed 13598  0.485192    conditional_sparse
50             first_pco2     11419    0.407443    conditional_sparse
80           max_pco2_0_24h   11419    0.407443    conditional_sparse
48           first_lactate    11124    0.396917    conditional_sparse
75            lactate_band    11124    0.396917    conditional_sparse
51              first_ph      10931    0.390031    conditional_sparse
82           min_ph_0_24h     10931    0.390031    conditional_sparse
85               ph_band      10931    0.390031    conditional_sparse
28            ed_first_temp   10360    0.369657    conditional_sparse
```
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 first_other_pco2_audit.csv
first_other_pco2 audit scope: uom_scope_df
Expected structural null fields: []
Unexpected full-null fields (top 20): []
UOM expectation scope: {'frame': 'final_cc_preview', 'row_count': 27975}
UOM expectation checks: {
  "expected_structural_null_fields": [
    "poc_abg_ph_uom",
    "poc_vbg_ph_uom",
    "poc_other_ph_uom"
  ],
  "structural_null_checks": {
    "poc_abg_ph_uom": {
      "present": true,
      "missing_n": 27975,
      "missing_pct": 1.0,
      "all_null": true
    },
    "poc_vbg_ph_uom": {
```

```
      "present": true,
      "missing_n": 27975,
      "missing_pct": 1.0,
      "all_null": true
    },
    "poc_other_ph_uom": {
      "present": true,
      "missing_n": 27975,
      "missing_pct": 1.0,
      "all_null": true
    }
  },
  "paco2_uom_checks": {
    "poc_abg_paco2_uom": {
      "present": true,
      "paired_value_column": "poc_abg_paco2",
      "value_rows": 0,
      "missing_uom_when_value_present": 0,
      "non_mmhg_uom_when_value_present": 0,
      "passes": true
    },
    "poc_vbg_paco2_uom": {
      "present": true,
      "paired_value_column": "poc_vbg_paco2",
      "value_rows": 0,
      "missing_uom_when_value_present": 0,
      "non_mmhg_uom_when_value_present": 0,
      "passes": true
    },
    "poc_other_paco2_uom": {
      "present": true,
      "paired_value_column": "poc_other_paco2",
      "value_rows": 0,
      "missing_uom_when_value_present": 0,
      "non_mmhg_uom_when_value_present": 0,
      "passes": true
    }
  }
}
```
% any gas panel 0-6h: 42.2
% any gas panel 0-24h: 96.0
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 gas_source_overlap_summary.csv

```
% source other (panel-level): 45.5
% source unknown (panel-level): 0.0
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/qa_summary.json
```

3.2 Outputs (ED-stay cohort + long tables)

Rationale: Persist ED-stay analytic datasets and supporting long tables.

```python
# Purpose: Write finalized cohort outputs and derivative tables for
↪   downstream analysis workflows.

# Save outputs (dated parquet artifacts archived under prior runs)
from datetime import datetime

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")

cohort_path = prior_runs_dir / f"cohort_ed_stay_{timestamp}.parquet"
ed_vitals_path = prior_runs_dir / f"ed_vitals_long_{timestamp}.parquet"
labs_path = prior_runs_dir / f"labs_long_{timestamp}.parquet"
gas_panels_path = prior_runs_dir / f"gas_panels_{timestamp}.parquet"

ed_df.to_parquet(cohort_path, index=False)
ed_vitals_long.to_parquet(ed_vitals_path, index=False)
labs_long.to_parquet(labs_path, index=False)
panel.to_parquet(gas_panels_path, index=False)

print("Wrote:", cohort_path)
print("Wrote:", ed_vitals_path)
print("Wrote:", labs_path)
print("Wrote:", gas_panels_path)
```

```
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/cohort_ed_stay_20260217_094455.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/ed_vitals_long_20260217_094455.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/labs_long_20260217_094455.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/gas_panels_20260217_094455.parquet
```

```python
# Purpose: Report missingness for key variables to support transparent
↪  data-quality reporting.

# Also export to Excel (tabular dataset) - optional archive output
from datetime import datetime

WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
↪  "1"
if WRITE_ARCHIVE_XLSX_EXPORTS:
    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')

    xlsx_path = prior_runs_dir /
↪  f"mimic_hypercap_EXT_EDstay_bq_gas_{timestamp}.xlsx"

    with pd.ExcelWriter(xlsx_path, engine="openpyxl") as xw:
        ed_df.to_excel(xw, sheet_name="cohort_ed_stay", index=False)
        # Optional: include QA tables if present
        try:
            miss.to_excel(xw, sheet_name="missingness", index=False)
        except Exception:
            pass

    print("Saved:", xlsx_path)
else:
    print("Skipping ED-stay archive workbook export (set
↪  WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")
```

Skipping ED-stay archive workbook export (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

```python
# Purpose: Capture ED presentation context (chief complaint, acuity, and
↪  early vitals) at the ED-stay level.

# Final Excel outputs requested: all encounters + ED chief-complaint only
from datetime import datetime

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')

# 1) All encounters meeting inclusion criteria (admission-level), with ED
↪  linkage flags
```

```python
ed_link = (
    ed_df.groupby("hadm_id", as_index=False)
        .agg(ed_stay_id_first=("ed_stay_id", "first"),
↪  n_ed_stays=("ed_stay_id", "nunique"))
)
all_encounters = df.merge(ed_link, on="hadm_id", how="left")
all_encounters["has_ed_encounter"] =
↪  all_encounters["n_ed_stays"].fillna(0).astype(int).gt(0).astype(int)
all_encounters["encounter_source"] =
↪  all_encounters["has_ed_encounter"].map({1: "ED-linked", 0:
↪  "Inpatient-only"})

if WRITE_ARCHIVE_XLSX_EXPORTS:
    all_path = prior_runs_dir /
↪  f"mimic_hypercap_EXT_all_encounters_bq_{timestamp}.xlsx"
    with pd.ExcelWriter(all_path, engine="openpyxl") as xw:
        all_encounters.to_excel(xw, sheet_name="all_encounters", index=False)
    print("Saved:", all_path)

# 2) ED chief-complaint-only (ED-stay level)
if "ed_triage_cc" not in ed_df.columns:
    raise KeyError("Column 'ed_triage_cc' not found in ed_df. Ensure ED
↪  triage merge ran.")

mask_cc = ed_df["ed_triage_cc"].notna() &
↪  (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
ed_cc_only = ed_df.loc[mask_cc].copy()

if WRITE_ARCHIVE_XLSX_EXPORTS:
    cc_path = prior_runs_dir /
↪  f"mimic_hypercap_EXT_EDcc_only_edstay_bq_{timestamp}.xlsx"
    with pd.ExcelWriter(cc_path, engine="openpyxl") as xw:
        ed_cc_only.to_excel(xw, sheet_name="ed_cc_only", index=False)
    print("Saved:", cc_path)
else:
    print("Skipping all-encounters/ED-CC archive exports (set
↪  WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")

Skipping all-encounters/ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

# Purpose: Capture ED presentation context (chief complaint, acuity, and
↪  early vitals) at the ED-stay level.
```

```python
# Final CC output (cohort-only; excludes NLP-derived columns)
from datetime import datetime

# Ensure cc-only frames exist (build from df if upstream optional export
↪  cells were skipped)
if "df_cc" not in globals():
    if "df" not in globals():
        raise NameError("df not found. Run cohort assembly cells first.")
    if "ed_triage_cc" not in df.columns:
        raise KeyError("ed_triage_cc missing in df; cannot build df_cc.")
    _mask_df_cc = df["ed_triage_cc"].notna() &
↪  (df["ed_triage_cc"].astype(str).str.strip() != "")
    df_cc = df.loc[_mask_df_cc].copy()


if "ed_cc_only" not in globals():
    if "ed_df" not in globals():
        raise NameError("ed_df not found. Run the ED-stay build cells
            ↪  first.")
    if "ed_triage_cc" not in ed_df.columns:
        raise KeyError("ed_triage_cc missing in ed_df; cannot build
            ↪  ed_cc_only.")
    _mask_cc = ed_df["ed_triage_cc"].notna() &
↪  (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
    ed_cc_only = ed_df.loc[_mask_cc].copy()

# Add ED-stay columns (dedup to earliest ED stay per hadm_id)
ed_tmp = ed_cc_only.copy()
if "ed_intime" in ed_tmp.columns:
    ed_tmp = ed_tmp.sort_values(["hadm_id", "ed_intime", "ed_stay_id"],
↪  na_position="last")
else:
    ed_tmp = ed_tmp.sort_values(["hadm_id", "ed_stay_id"],
↪  na_position="last")
ed_first = ed_tmp.drop_duplicates("hadm_id", keep="first")

add_ed_cols = [c for c in ed_first.columns if c not in df_cc.columns]

final_cc = df_cc.merge(ed_first[["hadm_id"] + add_ed_cols], on="hadm_id",
↪  how="left")

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")
```

```python
cohort_contract = validate_cohort_contract(final_cc)
cohort_contract_report = {
    "generated_utc": datetime.utcnow().isoformat() + "Z",
    "status": cohort_contract["status"],
    "contracts": {"cohort": cohort_contract},
    "findings": cohort_contract["findings"],
}
cohort_contract_paths = write_contract_report(cohort_contract_report,
  work_dir=WORK_DIR)
print("Wrote:", cohort_contract_paths["contract_report_path"])
if cohort_contract_paths["failed_contract_path"].exists():
    print("Wrote:", cohort_contract_paths["failed_contract_path"])


contract_mode = os.getenv("PIPELINE_CONTRACT_MODE", "fail").strip().lower()
if contract_mode not in {"fail", "warn"}:
    raise ValueError("PIPELINE_CONTRACT_MODE must be one of {'fail',
      'warn'}")
if cohort_contract["status"] == "fail" and contract_mode == "fail":
    raise ValueError(
        "Cohort contract checks failed. "
        f"See {cohort_contract_paths['failed_contract_path']}"
    )


dated_out_path = prior_runs_dir / f"{out_date} MIMICIV all with CC.xlsx"
canonical_out_path = DATA_DIR / "MIMICIV all with CC.xlsx"

# Write via temporary files then atomically replace targets to avoid
  zero-byte outputs on interruption.
def _atomic_write_excel(df_obj, target_path: Path) -> None:
    tmp_path = target_path.with_suffix(target_path.suffix + ".tmp")
    df_obj.to_excel(tmp_path, index=False)
    os.replace(tmp_path, target_path)

_atomic_write_excel(final_cc, dated_out_path)
_atomic_write_excel(final_cc, canonical_out_path)

cohort_manifest = collect_run_manifest(
    WORK_DIR,
    run_id=f"cohort_{datetime.now().strftime('%Y%m%d_%H%M%S')}",
)
cohort_manifest["stage"] = "cohort"
cohort_manifest["outputs"] = {
```

```python
        "canonical_output_path": str(canonical_out_path),
        "dated_output_path": str(dated_out_path),
        "gas_source_audit_path": str(gas_source_audit_path) if
        ↪   "gas_source_audit_path" in globals() else None,
        "gas_source_overlap_summary_path": str(gas_overlap_summary_path) if
        ↪   "gas_overlap_summary_path" in globals() else None,
        "anthropometrics_coverage_audit_path": str(anthro_coverage_audit_path) if
        ↪   "anthro_coverage_audit_path" in globals() else None,
        "ed_vitals_distribution_summary_path":
        ↪   str(ed_vitals_distribution_summary_path) if
        ↪   "ed_vitals_distribution_summary_path" in globals() else None,
        "ed_vitals_extreme_examples_path": str(ed_vitals_extreme_examples_path)
        ↪   if "ed_vitals_extreme_examples_path" in globals() else None,
        "ed_vitals_model_delta_path": str(ed_vitals_model_delta_path) if
        ↪   "ed_vitals_model_delta_path" in globals() else None,
        "outliers_audit_path": str(outliers_audit_path) if "outliers_audit_path"
        ↪   in globals() else None,
    }
    cohort_manifest["gas_source"] = {
        "mode": gas_source_mode if "gas_source_mode" in globals() else None,
        "source_rates": gas_source_audit.get("source_rates", {}) if
        ↪   "gas_source_audit" in globals() else {},
        "tier_rates": gas_source_audit.get("tier_rates", {}) if
        ↪   "gas_source_audit" in globals() else {},
    }
    cohort_manifest["anthropometrics"] = (
        anthro_coverage_audit if "anthro_coverage_audit" in globals() else None
    )
    cohort_manifest_path = prior_runs_dir / f"{out_date}
    ↪   cohort_run_manifest.json"
    cohort_manifest_path.write_text(json.dumps(cohort_manifest, indent=2))

    print("Saved:", dated_out_path)
    print("Saved:", canonical_out_path)
    print("Wrote:", cohort_manifest_path)
    print("Base rows:", len(df_cc), "| Added from ED-stay:", len(add_ed_cols))
    print("Total columns:", len(final_cc.columns))
```

```
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/debug/contracts/20260217_164457/contract_report.json
Saved: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 MIMICIV all with CC.xlsx
```

```
Saved: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/MIMICIV all with CC.xlsx
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-17 cohort_run_manifest.json
Base rows: 27975 | Added from ED-stay: 151
Total columns: 258
```

```python
# Purpose: Extract and standardize first ABG/VBG physiology fields for
↪   baseline characterization.

# Data dictionary (OSF-style) for final cohort output
from datetime import datetime
import numpy as np

# Choose target dataset for dictionary
if "final_cc" in globals():
    dd_target = final_cc
elif "df_cc" in globals():
    dd_target = df_cc
elif "df" in globals():
    dd_target = df
else:
    raise NameError("No cohort dataframe found for data dictionary (expected
    ↪   final_cc/df_cc/df).")

# Helper functions

def _readable(name: str) -> str:
    return name.replace("_", " ").strip().title()


def _infer_units(col: str) -> str:
    if col.endswith("_hours"):
        return "hours"
    if col.endswith("_time") or col.endswith("_intime") or
    ↪   col.endswith("_outtime") or col.endswith("_date"):
        return "datetime"
    if col.startswith("flag_") or col.endswith("_flag") or
    ↪   col.endswith("_before_imv"):
        return "binary (0/1)"
    if col.endswith("_paco2") or col.endswith("_pco2"):
        return "mmHg"
    if col.endswith("_ph"):
```

```python
        return "pH"
    if col.endswith("_lactate"):
        return "mmol/L"
    return ""


def _allowed_values(series):
    s = series.dropna()
    if s.empty:
        return ""
    if s.dtype.kind in "biu":
        vals = sorted(map(int, s.unique()))
        if len(vals) <= 15:
            return ", ".join(map(str, vals))
        return f"{min(vals)}-{max(vals)}"
    if s.dtype.kind in "f":
        return f"{np.nanmin(s):.3g}-{np.nanmax(s):.3g}"
    # object/categorical: use bounded sampling to keep dictionary generation
    ↪  fast.
    sample_n = 5000
    s_str = s.astype(str)
    if len(s_str) > sample_n:
        s_sample = s_str.head(sample_n)
        sample_note = f"sampled {sample_n}/{len(s_str)} rows"
    else:
        s_sample = s_str
        sample_note = None
    unique_n = int(s_sample.nunique(dropna=True))
    top_vals = s_sample.value_counts(dropna=False).head(15).index.tolist()
    if unique_n > len(top_vals):
        tail = f" … ({unique_n} unique"
        if sample_note:
            tail += f"; {sample_note}"
        tail += ")"
        return ", ".join(top_vals) + tail
    if sample_note:
        return ", ".join(top_vals) + f" ({sample_note})"
    return ", ".join(top_vals)


META = {
    # IDs
    "subject_id": {
```

```
        "label": "Subject identifier",
        "definition": "Identifier for a unique patient in MIMIC-IV.",
        "source": "mimiciv_hosp.patients",
        "derivation": "as recorded",
    },
    "hadm_id": {
        "label": "Hospital admission identifier",
        "definition": "Identifier for a hospital admission.",
        "source": "mimiciv_hosp.admissions",
        "derivation": "as recorded",
    },
    "ed_stay_id": {
        "label": "ED stay identifier",
        "definition": "Identifier for an ED stay.",
        "source": "mimiciv_ed.edstays",
        "derivation": "as recorded",
    },
    "ed_intime": {
        "label": "ED arrival time",
        "definition": "ED arrival time for this ED stay.",
        "source": "mimiciv_ed.edstays",
        "derivation": "as recorded",
        "units": "datetime",
    },
    "ed_outtime": {
        "label": "ED departure time",
        "definition": "ED departure time for this ED stay.",
        "source": "mimiciv_ed.edstays",
        "derivation": "as recorded",
        "units": "datetime",
    },
    "ed_intime_first": {
        "label": "First ED arrival time (admission)",
        "definition": "Earliest ED arrival time among ED stays linked to the
         ↪  admission.",
        "source": "mimiciv_ed.edstays",
        "derivation": "min(edstays.intime) per hadm_id",
        "units": "datetime",
    },
    # Demographics / triage
    "ed_gender": {
        "label": "ED gender",
        "definition": "Gender as recorded in ED stay table.",
```

```
        "source": "mimiciv_ed.edstays",
        "derivation": "as recorded",
    },
    "ed_race": {
        "label": "ED race",
        "definition": "Race as recorded in ED stay table.",
        "source": "mimiciv_ed.edstays",
        "derivation": "as recorded",
    },
    "hosp_race": {
        "label": "Hospital race",
        "definition": "Race as recorded in admissions table.",
        "source": "mimiciv_hosp.admissions",
        "derivation": "as recorded",
    },
    "ed_triage_cc": {
        "label": "ED chief complaint",
        "definition": "Chief complaint recorded at ED triage.",
        "source": "mimiciv_ed.triage",
        "derivation": "as recorded",
    },
    "ed_triage_acuity": {
        "label": "ED triage acuity",
        "definition": "Triage acuity level recorded at ED presentation.",
        "source": "mimiciv_ed.triage",
        "derivation": "as recorded",
    },
    # Outcomes
    "hospital_expire_flag": {
        "label": "Hospital expire flag",
        "definition": "Hospital mortality indicator from admissions table.",
        "source": "mimiciv_hosp.admissions",
        "derivation": "as recorded",
        "units": "binary (0/1)",
    },
    "in_hospital_death": {
        "label": "In-hospital death",
        "definition": "Indicator for in-hospital death during admission.",
        "source": "mimiciv_hosp.admissions",
        "derivation": "(hospital_expire_flag == 1) OR (deathtime not null)",
        "units": "binary (0/1)",
    },
    # Ventilation
```

```
"imv_flag": {
    "label": "IMV flag",
    "definition": "Indicator of invasive mechanical ventilation during
    ↪  admission.",
    "source": "ICD procedures + ICU chartevents",
    "derivation": "max(IMV ICD flag, IMV chart flag)",
    "units": "binary (0/1)",
},
"niv_flag": {
    "label": "NIV flag",
    "definition": "Indicator of noninvasive ventilation during
    ↪  admission.",
    "source": "ICD procedures + ICU chartevents",
    "derivation": "max(NIV ICD flag, NIV chart flag)",
    "units": "binary (0/1)",
},
"first_imv_time": {
    "label": "First IMV time",
    "definition": "Earliest charted time of IMV in ICU chartevents (if
    ↪  present).",
    "source": "mimiciv_icu.chartevents",
    "derivation": "min charttime among IMV charted events",
    "units": "datetime",
},
"first_niv_time": {
    "label": "First NIV time",
    "definition": "Earliest charted time of NIV in ICU chartevents (if
    ↪  present).",
    "source": "mimiciv_icu.chartevents",
    "derivation": "min charttime among NIV charted events",
    "units": "datetime",
},
"dt_first_imv_hours": {
    "label": "Hours to first IMV",
    "definition": "Hours from first ED arrival to first IMV time.",
    "source": "Derived",
    "derivation": "(first_imv_time - ed_intime_first) in hours",
    "units": "hours",
},
"dt_first_niv_hours": {
    "label": "Hours to first NIV",
    "definition": "Hours from first ED arrival to first NIV time.",
    "source": "Derived",
```

```
        "derivation": "(first_niv_time - ed_intime_first) in hours",
        "units": "hours",
    },
    "abg_before_imv": {
        "label": "ABG before IMV",
        "definition": "Indicator that first ABG time precedes first IMV
        ↪  time.",
        "source": "Derived",
        "derivation": "first_abg_time < first_imv_time",
        "units": "binary (0/1)",
    },
    "vbg_before_imv": {
        "label": "VBG before IMV",
        "definition": "Indicator that first VBG time precedes first IMV
        ↪  time.",
        "source": "Derived",
        "derivation": "first_vbg_time < first_imv_time",
        "units": "binary (0/1)",
    },
    # Gas/ABG/VBG
    "abg_hypercap_threshold": {
        "label": "ABG hypercapnia threshold met",
        "definition": "Indicator that any arterial pCO2   45 mmHg in hosp/ICU
        ↪  pCO2 extraction.",
        "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
        "derivation": "max(arterial pCO2   45) per hadm_id",
        "units": "binary (0/1)",
    },
    "vbg_hypercap_threshold": {
        "label": "VBG hypercapnia threshold met",
        "definition": "Indicator that any venous pCO2   50 mmHg in hosp/ICU
        ↪  pCO2 extraction.",
        "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
        "derivation": "max(venous pCO2   50) per hadm_id",
        "units": "binary (0/1)",
    },
    "other_hypercap_threshold": {
        "label": "Other-source hypercapnia threshold met",
        "definition": "Indicator that any other/unspecified-source pCO2   50
        ↪  mmHg in hosp/ICU pCO2 extraction.",
        "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
        "derivation": "max(other-source pCO2   50) per hadm_id",
        "units": "binary (0/1)",
```

```
    },
    "first_gas_time": {
        "label": "First gas panel time",
        "definition": "Earliest gas panel time in the ED 0-24h window.",
        "source": "mimiciv_hosp.labevents (gas panels)",
        "derivation": "min panel_time within ED 0-24h window",
        "units": "datetime",
    },
    "dt_first_qualifying_gas_hours": {
        "label": "Hours to first gas panel",
        "definition": "Hours from first ED arrival to first gas panel time.",
        "source": "Derived",
        "derivation": "(first_gas_time - ed_intime_first) in hours",
        "units": "hours",
    },
    "first_other_time": {
        "label": "First other-source gas time",
        "definition": "Earliest time of a pCO2 measurement with
         ↪  other/unspecified source classification.",
        "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
        "derivation": "min(other-source pCO2 charttime) across LAB and POC",
        "units": "datetime",
    },
    "presenting_hypercapnia": {
        "label": "Presenting ( 6h) gas timing",
        "definition": "Indicator that first gas panel occurred within 6h of
         ↪  first ED arrival.",
        "source": "Derived",
        "derivation": "dt_first_qualifying_gas_hours   6",
        "units": "binary (0/1)",
    },
    "late_hypercapnia": {
        "label": "Late (>6h) gas timing",
        "definition": "Indicator that first gas panel occurred >6h after
         ↪  first ED arrival.",
        "source": "Derived",
        "derivation": "dt_first_qualifying_gas_hours > 6",
        "units": "binary (0/1)",
    },
    "gas_source_other_rate": {
        "label": "Gas source other/unknown rate",
        "definition": "Proportion of gas panels classified as
         ↪  other/unspecified source per ED stay.",
```

```
        "source": "Derived from gas panel source inference",
        "derivation": "mean(source == 'other') per ed_stay_id",
        "units": "proportion",
    },
    # Comorbidities
    "flag_copd": {
        "label": "COPD/emphysema flag",
        "definition": "Indicator for COPD/emphysema ICD codes in ED or
        ↪  hospital diagnoses.",
        "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
        "derivation": "ICD prefixes J43/J44 (ED OR hospital)",
        "units": "binary (0/1)",
    },
    "flag_osa_ohs": {
        "label": "OSA/OHS flag",
        "definition": "Indicator for OSA/OHS ICD codes in ED or hospital
        ↪  diagnoses.",
        "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
        "derivation": "ICD prefixes G473/E662 (ED OR hospital)",
        "units": "binary (0/1)",
    },
    "flag_chf": {
        "label": "CHF flag",
        "definition": "Indicator for CHF ICD codes in ED or hospital
        ↪  diagnoses.",
        "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
        "derivation": "ICD prefix I50 (ED OR hospital)",
        "units": "binary (0/1)",
    },
    "flag_neuromuscular": {
        "label": "Neuromuscular flag",
        "definition": "Indicator for neuromuscular ICD codes in ED or
        ↪  hospital diagnoses.",
        "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
        "derivation": "ICD prefixes G12/G70/G71 (ED OR hospital)",
        "units": "binary (0/1)",
    },
    "flag_opioid_substance": {
        "label": "Opioid/substance flag",
        "definition": "Indicator for opioid/substance ICD codes in ED or
        ↪  hospital diagnoses.",
        "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
        "derivation": "ICD prefixes F11/T40/F13 (ED OR hospital)",
```

```python
            "units": "binary (0/1)",
        },
        "flag_pneumonia": {
            "label": "Pneumonia flag",
            "definition": "Indicator for pneumonia ICD codes in ED or hospital
            ↪   diagnoses.",
            "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
            "derivation": "ICD prefixes J12-J18 (ED OR hospital)",
            "units": "binary (0/1)",
        },
}

rows = []
for col in dd_target.columns:
    info = META.get(col, {})
    series = dd_target[col]
    rows.append({
        "variable": col,
        "label": info.get("label", _readable(col)),
        "units": info.get("units", _infer_units(col)),
        "allowed_values": info.get("allowed_values",
        ↪   _allowed_values(series)),
        "definition": info.get("definition", "UNCONFIRMED"),
        "synonyms": info.get("synonyms", ""),
        "description": info.get("description", ""),
        "source": info.get("source", "UNCONFIRMED"),
        "derivation": info.get("derivation", "UNCONFIRMED"),
        "dtype": str(series.dtype),
        "example_value": series.dropna().iloc[0] if series.dropna().shape[0]
        ↪   else None,
    })

data_dict = pd.DataFrame(rows)

out_date = datetime.now().strftime("%Y-%m-%d")
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_xlsx = prior_runs_dir / f"{out_date} MIMICIV all with
↪   CC_data_dictionary.xlsx"
out_csv = prior_runs_dir / f"{out_date} MIMICIV all with
↪   CC_data_dictionary.csv"

data_dict.to_excel(out_xlsx, index=False)
```

```
data_dict.to_csv(out_csv, index=False)

print("Saved data dictionary:", out_xlsx)
print("Saved data dictionary:", out_csv)
print("Rows:", len(data_dict))
```