

MIMICIV_hypercap_EXT_cohort

Table of contents

1 MIMIC-IV Hypercapnia Cohort — ICD Physiologic Thresholds (BigQuery)	2
1.1 Pipeline stages (readable map)	3
1.2 SQL registry (centralized query templates)	3
2 MIMIC-IV on BigQuery	4
2.1 Environment Bootstrap & Smoke Test	4
2.2 0. Prerequisites (one-time)	4
2.3 Data Generation	77
2.3.1 1) ICD cohort flags (hypercapnic respiratory failure)	77
2.3.2 2) Blood gas inclusion thresholds — qualifying hypercapnic gases in 0–24h (LAB + POC)	81
2.3.3 3) Cohort enrollment and <code>hadm_list</code> for downstream queries	91
2.3.4 4) First ABG and First VBG (LAB + POC, standardized to mmHg)	94
2.3.5 5) Demographics & outcomes	105
2.3.6 6) NIH/OMB race & ethnicity (ED + Hospital)	109
2.3.7 7) ED triage (linked to hadm) and first ED vitals	112
2.3.8 8) ICU meta (first ICU stay, LOS days)	115
2.3.9 9) Ventilation flags (ICD procedures)	116
2.3.10 10) Assemble final DataFrame	120
2.3.11 11) Sanity checks	134
2.3.12 PDF-ready long tables	135
2.3.13 12) Save to Excel	140
2.3.14 Create Annotation Dataset	141
3 ED-stay cohort expansion (timing, severity, comorbidity, outcomes)	143
3.0.1 Phase 0 — Inventory & missing-field registry	143
3.0.2 Phase 1 — ED encounter spine and ED enrichment (one row per ED stay)	147
3.0.3 Phase 2 — Hospital admission context and outcomes	151
3.0.4 Phase 3 — ICU timing and LOS	153
3.0.5 Phase 4 — ED longitudinal vitals (0–6h)	154

3.0.6	Phase 5 — Robust lab discovery + gas panels	157
3.0.7	Phase 5C — ICU POC blood gases (chartevents, optional)	179
3.0.8	Phase 6 — BMI/anthropometrics (OMR)	191
3.0.9	Phase 7 — ICD comorbidity flags	208
3.0.10	Phase 8 — Timing phenotypes and derived bands	212
3.1	QA & Data Fidelity	222
3.2	Outputs (ED-stay cohort + long tables)	249

1 MIMIC-IV Hypercapnia Cohort — ICD Physiologic Thresholds (BigQuery)

TODO:

[] want to add ED-rendered diagnoses - a flag to split off whether the hypercapnic respiratory failure ICD was rendered in the ED or during the hospital stay. [] are the installation instructions at the beginning of this notebook consistent with the way we intend the notebook to be used? is it also consistent with the README.md instructions?

Goal: Build an admissions-level tabular dataset that **enrolls** any hospital admission (`hadm_id`) with a **first qualifying hypercapnic blood gas** (LAB or POC) within **24 hours** of first ED arrival:

1. Arterial blood gas (ABG): $\text{PaCO} > 45.0 \text{ mmHg}$
2. Venous blood gas (VBG): $\text{PaCO} > 50.0 \text{ mmHg}$
3. Unknown/indeterminate blood gas source: $\text{PaCO} > 50.0 \text{ mmHg}$

Then, keep all downstream columns/logic from the current workflow: - Per-code ICD indicators and an `any_hypercap_icd` flag. - Deterministic extraction of the **first qualifying hypercapnic blood gas** with source/site provenance. - Demographics/outcomes, NIH/OMB race & ethnicity, ED triage + first ED vitals, ICU meta (first stay + LOS), ventilation flags. - Sanity checks.

Assumptions - You already configured BigQuery auth (`gcloud auth application-default login`) and `.env` variables as in the previous notebook. - The PhysioNet hosting project is `physionet-data`. - Datasets exist (e.g., `mimiciv_3_1_hosp`, `mimiciv_3_1_icu`, and an ED dataset such as `mimiciv_ed`). This notebook auto-detects the ED dataset.

Execution timing note: To capture per-cell runtimes, enable cell execution timing in VS Code (Notebook: Show Cell Execution Time) or enable ExecuteTime in Jupyter. Then re-run and save to allow runtime summaries.

1.1 Pipeline stages (readable map)

1. Stage A — Config & helpers: dataset IDs, thresholds, shared helper functions.
2. Stage B — Cohort spine: ICD + gas thresholds → hadm_list / ed_stay_list.
3. Stage C — Bulk pulls: one query per table where possible.
4. Stage D — Transforms: panels, flags, timing, comorbidities, vitals.
5. Stage E — QA checks: deterministic assertions and range checks.
6. Stage F — Outputs: parquet + Excel exports.

1.2 SQL registry (centralized query templates)

All BigQuery SQL is defined in one place below, then referenced by name in subsequent cells.

```
# Purpose: Build ABG/VBG hypercapnia threshold flags from lab and ICU POC
#           ↳ pCO2 measurements.

import sys
print(sys.executable)

# Central SQL registry (define all query templates here)
SQL = {}

def sql(name: str) -> str:
    if name not in SQL:
        raise KeyError(f"SQL template not found: {name}")
    return SQL[name]

# SQL templates (populated below in-place to keep notebook linear)
# Names: admit_sql, co2_thresholds_sql, cohort_icd_sql, counts_sql, demo_sql,
#         ↳ ditems_sql, ed_counts_sql, ed_first_vitals_sql, ed_spine_sql,
#         ↳ ed_to_icu_sql, ed_triage_sql, ed_vitals_sql, icd_sql, icu_meta_sql,
#         ↳ icu_sql, labitems_sql, labs_sql, vent_chart_sql, vent_sql

/Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/HyperCC-NLP/.venv/bin/python
```

Rationale: Define a reproducible, admission-level cohort that captures hypercapnia using complementary diagnostic (ICD) and physiologic (blood gas) criteria.

2 MIMIC-IV on BigQuery

2.1 Environment Bootstrap & Smoke Test

Purpose: make a clean, reproducible start on a new machine.

Outcome: verify auth, project config, and dataset access; provide a reusable BigQuery runner for the build notebook.

Rationale: Establish the BigQuery environment and dataset configuration so queries are consistent and reproducible across runs.

2.2 0. Prerequisites (one-time)

Accounts & access - PhysioNet access to MIMIC-IV on BigQuery; in BigQuery Console star project `physionet-data`. - A Google Cloud **Project ID** with BigQuery API enabled (this is your **billing** project).

CLI & environment - Google Cloud SDK (`gcloud`) installed and on PATH. - Python environment created with `uv` (see README) and Jupyter kernel selected. - A project-local `.env` with the variables below.

.env variables

```
MIMIC_BACKEND=bigquery
WORK_PROJECT=<your-billing-project-id>
BQ_PHYSIONET_PROJECT=physionet-data
BQ_DATASET_HOSP=mimiciv_3_1_hosp
BQ_DATASET_ICU=mimiciv_3_1_icu
BQ_DATASET_ED=mimiciv_ed
WORK_DIR=/path/to/Hypercap-CC-NLP
# GOOGLE_APPLICATION_CREDENTIALS=/path/to/service-account.json
```

Command line quickstart

```
brew install --cask google-cloud-sdk
gcloud init
gcloud auth application-default login
gcloud services enable bq.googleapis.com --project
  ↳ <your-billing-project-id>
ls -l ~/.config/gcloud/application_default_credentials.json
```

Rationale: Verify access and credentials up front to prevent silent failures later in the pipeline.

```

# Purpose: Set up project paths, environment variables, and BigQuery client
    ↵ connections for reproducible execution.

# --- Imports & environment
import os, re, json, math, textwrap, sys, hashlib
from pathlib import Path
from typing import Any, Mapping, Sequence

import numpy as np
import pandas as pd

from google.cloud import bigquery
from google.oauth2 import service_account
from dotenv import load_dotenv

load_dotenv()

WORK_DIR = Path(os.getenv("WORK_DIR", Path.cwd())).expanduser().resolve()
DATA_DIR = WORK_DIR / "MIMIC tabular data"
DATA_DIR.mkdir(parents=True, exist_ok=True)
SRC_DIR = WORK_DIR / "src"
if SRC_DIR.exists() and str(SRC_DIR) not in sys.path:
    sys.path.insert(0, str(SRC_DIR))

# Local helper functions (core cohort logic must remain embedded in this
    ↵ notebook)
OMR_RESULT_NAMES = ("bmi", "height", "weight")
OMR_OUTPUT_COLUMNS = (
    "bmi_closest_pre_ed",
    "height_closest_pre_ed",
    "weight_closest_pre_ed",
)
ANTHRO_UNIT_COLUMNS: dict[str, str] = {
    "bmi": "bmi_closest_pre_ed_uom",
    "height": "height_closest_pre_ed_uom",
    "weight": "weight_closest_pre_ed_uom",
}
ANTHRO_TIME_COLUMNS: dict[str, str] = {
    "bmi": "bmi_closest_pre_ed_time",
    "height": "height_closest_pre_ed_time",
    "weight": "weight_closest_pre_ed_time",
}

```

```

}

ANTHRO_VALUE_COLUMNS: dict[str, str] = {
    "bmi": "bmi_closest_pre_ed",
    "height": "height_closest_pre_ed",
    "weight": "weight_closest_pre_ed",
}
}

ANTHRO_CANONICAL_UNITS: dict[str, str] = {
    "bmi": "kg/m2",
    "height": "cm",
    "weight": "kg",
}
}

ANTHRO_SOURCE_PRIORITY: dict[str, int] = {
    "ED": 0,
    "ICU": 1,
    "HOSPITAL": 2,
    "missing": 99,
}
}

ANTHRO_BMI_PAIR_WINDOW_HOURS = 24.0 * 7.0

OMR_PROVENANCE_COLUMNS = (
    "anthro_timing_tier",
    "anthro_days_offset",
    "anthro_chartdate",
    "anthro_timing_uncertain",
    "anthro_source",
    "anthro_obstime",
    "anthro_hours_offset",
    "anthro_timing_basis",
)
)

OMR_TIMING_TIERS = ("pre_ed_365", "post_ed_365", "missing")

EXPECTED_STRUCTURAL_NULL_FIELDS: tuple[str, ...] = ()

PACO2_VALUE_UOM_PAIRS = (
    ("poc_abg_paco2", "poc_abg_paco2_uom"),
    ("poc_vbg_paco2", "poc_vbg_paco2_uom"),
    ("poc_other_paco2", "poc_other_paco2_uom"),
)
)
```

```

DEFAULT_VITALS_MODEL_RANGES: dict[str, tuple[float, float]] = {
    "ed_triage_hr": (20.0, 250.0),
    "ed_first_hr": (20.0, 250.0),
    "ed_triage_rr": (4.0, 80.0),
    "ed_first_rr": (4.0, 80.0),
}

DEFAULT_GAS_MODEL_RANGES: dict[str, tuple[float, float]] = {
    "first_ph": (6.8, 7.8),
    "first_pco2": (10.0, 200.0),
    "first_other_pco2": (10.0, 200.0),
    "first_lactate": (0.0, 30.0),
}

_NUMERIC_TOKEN_PATTERN = re.compile(r"(-?\d+(?:\.\d+)?)")
_ARTERIAL_HINT_PATTERN = re.compile(
    r"(arterial|abg|a[- ]?line|art line|\bart\b|\bartery\b",
    re.I,
)
_VENOUS_HINT_PATTERN = re.compile(
    r"(venous|vbg|cvbg|mixed venous|central venous|\bven\b",
    re.I,
)
_P02_LABEL_DENY_PATTERN_DEFAULT = re.compile(
    r"(:|^|[a-z])p(:a)?\s*\o\s*2(:|[a-z]|$)|\bpo2\b|\bpao2\b|\oxygen|o2\s*"
    "\sat|saturation|sao2|spo2|fio2|et\s*co2|end[-"
    "]?tidal|t\s*co2|tco2|total\s*co2|hco3|bicar|v\s*co2|vco2|co2\s*(?:"
    "prod|production|elimin|elimination)",
    re.I,
)

BLOOD_GAS_SPEC_PATH = WORK_DIR / "specs" / "blood_gas_itemids.json"

def _coerce_itemid_list(values: Sequence[Any], *, key_name: str) ->
    list[int]:
    cleaned: list[int] = []
    for raw in values:
        try:
            cleaned.append(int(raw))
        except (TypeError, ValueError) as exc:

```

```

        raise ValueError(f"Invalid itemid in {key_name}: {raw!r}") from
        ↵     exc
    return sorted(set(cleaned))

def _sql_int_array(values: Sequence[int], *, empty_sentinel: int =
    ↵ -999999999) -> str:
    """Return a BigQuery INT64 array literal from Python integers."""
    resolved = [int(v) for v in values]
    if not resolved:
        resolved = [int(empty_sentinel)]
    return "[" + ", ".join(str(v) for v in resolved) + "]"

def load_blood_gas_itemid_manifest(path: Path) -> dict[str, Any]:
    """Load and validate versioned blood-gas itemid manifest."""
    if not path.exists():
        raise FileNotFoundError(
            f"Blood-gas manifest missing at {path}. Add"
            ↵     specs/blood_gas_itemids.json."
        )
    payload = json.loads(path.read_text())
    if not isinstance(payload, dict):
        raise ValueError("Blood-gas manifest must be a JSON object.")

    for section in ("lab", "icu", "denylist"):
        if section not in payload or not isinstance(payload[section], dict):
            raise ValueError(f"Blood-gas manifest missing required section"
                ↵     '{section}'.")

    lab = payload["lab"]
    icu = payload["icu"]
    denylist = payload["denylist"]

    lab_pco2 = _coerce_itemid_list(lab.get("pco2_itemids", []),
    ↵     key_name="lab.pco2_itemids")
    lab_pco2_excluded = _coerce_itemid_list(
        lab.get("excluded_pco2_itemids", []),
        key_name="lab.excluded_pco2_itemids",
    )
    lab_po2 = _coerce_itemid_list(lab.get("po2_itemids", []),
    ↵     key_name="lab.po2_itemids")
    lab_po2_excluded = _coerce_itemid_list(

```

```

        lab.get("excluded_po2_itemids", []),
        key_name="lab.excluded_po2_itemids",
    )
    lab_ph = _coerce_itemid_list(lab.get("ph_itemids", []),
        key_name="lab.ph_itemids")
    lab_hco3 = _coerce_itemid_list(lab.get("hco3_itemids", []),
        key_name="lab.hco3_itemids")
    lab_spec = _coerce_itemid_list(
        lab.get("specimen_type_itemids", []),
        key_name="lab.specimen_type_itemids",
    )
    if not lab_pco2:
        raise ValueError("Blood-gas manifest requires at least one
            ↵ lab.pco2_itemids value.")
    if not lab_po2:
        raise ValueError("Blood-gas manifest requires at least one
            ↵ lab.po2_itemids value.")
    if not lab_ph:
        raise ValueError("Blood-gas manifest requires at least one
            ↵ lab.ph_itemids value.")
    if not lab_hco3:
        raise ValueError("Blood-gas manifest requires at least one
            ↵ lab.hco3_itemids value.")
    if not lab_spec:
        raise ValueError(
            "Blood-gas manifest requires at least one
            ↵ lab.specimen_type_itemids value."
        )

    blood_fluid_terms = [str(term).strip().lower() for term in
        ↵ lab.get("blood_fluid_terms", ["blood"])]
    blood_fluid_terms = sorted(set(term for term in blood_fluid_terms if
        ↵ term))
    if not blood_fluid_terms:
        blood_fluid_terms = ["blood"]

    icu_pco2 = _coerce_itemid_list(icu.get("pco2_itemids", []),
        ↵ key_name="icu.pco2_itemids")
    icu_pco2_abg = _coerce_itemid_list(
        icu.get("pco2_abg_itemids", []),
        key_name="icu.pco2_abg_itemids",
    )
    icu_pco2_vbg = _coerce_itemid_list(

```

```

        icu.get("pco2_vbg_itemids", []),
        key_name="icu.pco2_vbg_itemids",
    )
    icu_pco2_excluded = _coerce_itemid_list(
        icu.get("excluded_pco2_itemids", []),
        key_name="icu.excluded_pco2_itemids",
    )
    icu_po2 = _coerce_itemid_list(icu.get("po2_itemids", []),
        key_name="icu.po2_itemids")
    icu_po2_abg = _coerce_itemid_list(
        icu.get("po2_abg_itemids", []),
        key_name="icu.po2_abg_itemids",
    )
    icu_po2_vbg = _coerce_itemid_list(
        icu.get("po2_vbg_itemids", []),
        key_name="icu.po2_vbg_itemids",
    )
    icu_po2_excluded = _coerce_itemid_list(
        icu.get("excluded_po2_itemids", []),
        key_name="icu.excluded_po2_itemids",
    )
    icu_ph = _coerce_itemid_list(icu.get("ph_itemids", []),
        key_name="icu.ph_itemids")
    icu_hco3 = _coerce_itemid_list(icu.get("hco3_itemids", []),
        key_name="icu.hco3_itemids")
    icu_spec = _coerce_itemid_list(
        icu.get("specimen_type_itemids", []),
        key_name="icu.specimen_type_itemids",
    )
if icu_pco2_excluded:
    excluded_set = set(icu_pco2_excluded)
    icu_pco2 = sorted(set(icu_pco2).difference(excluded_set))
    icu_pco2_abg = sorted(set(icu_pco2_abg).difference(excluded_set))
    icu_pco2_vbg = sorted(set(icu_pco2_vbg).difference(excluded_set))
if icu_po2_excluded:
    excluded_set_po2 = set(icu_po2_excluded)
    icu_po2 = sorted(set(icu_po2).difference(excluded_set_po2))
    icu_po2_abg = sorted(set(icu_po2_abg).difference(excluded_set_po2))
    icu_po2_vbg = sorted(set(icu_po2_vbg).difference(excluded_set_po2))
if icu_pco2_abg and not set(icu_pco2_abg).issubset(set(icu_pco2)):
    raise ValueError(
        "Blood-gas manifest icu.pco2_abg_itemids must be a subset of
        ↵ icu.pco2_itemids."

```

```

        )
if icu_pco2_vbg and not set(icu_pco2_vbg).issubset(set(icu_pco2)):
    raise ValueError(
        "Blood-gas manifest icu.pco2_vbg_itemids must be a subset of
         ↳ icu.pco2_itemids."
    )
if icu_po2_abg and not set(icu_po2_abg).issubset(set(icu_po2)):
    raise ValueError(
        "Blood-gas manifest icu.po2_abg_itemids must be a subset of
         ↳ icu.po2_itemids."
    )
if icu_po2_vbg and not set(icu_po2_vbg).issubset(set(icu_po2)):
    raise ValueError(
        "Blood-gas manifest icu.po2_vbg_itemids must be a subset of
         ↳ icu.po2_itemids."
    )

deny_tokens = [
    str(token).strip().lower()
    for token in denylist.get("label_tokens", [])
    if str(token).strip()
]
allow_patternFallback = bool(icu.get("allow_patternFallback", True))
inclusion_mode = str(icu.get("inclusion_mode",
    "validated_only").strip().lower())
if inclusion_mode not in {"validated_only", "disabled"}:
    raise ValueError(
        "Blood-gas manifest icu.inclusion_mode must be one of "
        "{'validated_only', 'disabled'}."
    )

pco2_qc_payload = icu.get("pco2_qc", {})
if pco2_qc_payload is None:
    pco2_qc_payload = {}
if not isinstance(pco2_qc_payload, dict):
    raise ValueError("Blood-gas manifest icu.pco2_qc must be a JSON
        ↳ object.")
qc_allowed_uoms = [
    str(value).strip().lower()
    for value in pco2_qc_payload.get("allowed_uoms", ["", "mmhg", "kpa"])
]
qc_allowed_uoms = sorted(set(qc_allowed_uoms))
if not qc_allowed_uoms:

```

```

qc_allowed_uoms = ["", "mmhg", "kpa"]
qc_label_include = [
    str(value).strip().lower()
    for value in pco2_qc_payload.get("label_must_include_tokens",
        ["pco2"])
    if str(value).strip()
]
qc_label_include = sorted(set(qc_label_include))
qc_label_exclude = [
    str(value).strip().lower()
    for value in pco2_qc_payload.get(
        "label_must_exclude_tokens",
        ["po2", "pao2", "oxygen", "sao2", "spo2", "fio2"],
    )
    if str(value).strip()
]
qc_label_exclude = sorted(set(qc_label_exclude))
qc_median_min = float(pco2_qc_payload.get("median_min", 45.0))
qc_median_max = float(pco2_qc_payload.get("median_max", 80.0))
if qc_median_min > qc_median_max:
    raise ValueError(
        "Blood-gas manifest icu.pco2_qc median_min must be <=
        median_max."
    )

manifest = {
    "version": str(payload.get("version", "UNSPECIFIED")),
    "lab": {
        "pco2_itemids": lab_pco2,
        "excluded_pco2_itemids": lab_pco2_excluded,
        "po2_itemids": lab_po2,
        "excluded_po2_itemids": lab_po2_excluded,
        "ph_itemids": lab_ph,
        "hco3_itemids": lab_hco3,
        "specimen_type_itemids": lab_spec,
        "blood_fluid_terms": blood_fluid_terms,
    },
    "icu": {
        "pco2_itemids": icu_pco2,
        "pco2_abg_itemids": icu_pco2_abg,
        "pco2_vbg_itemids": icu_pco2_vbg,
        "excluded_pco2_itemids": icu_pco2_excluded,
        "po2_itemids": icu_po2,
    }
}

```

```

    "po2_abg_itemids": icu_po2_abg,
    "po2_vbg_itemids": icu_po2_vbg,
    "excluded_po2_itemids": icu_po2_excluded,
    "ph_itemids": icu_ph,
    "hco3_itemids": icu_hco3,
    "specimen_type_itemids": icu_spec,
    "allow_patternFallback": allow_patternFallback,
    "inclusion_mode": inclusion_mode,
    "pco2_qc": {
        "allowed_uoms": qc_allowed_uoms,
        "median_min": qc_median_min,
        "median_max": qc_median_max,
        "label_must_include_tokens": qc_label_include,
        "label_must_exclude_tokens": qc_label_exclude,
    },
},
"denylist": {"label_tokens": deny_tokens},
"path": str(path),
}
return manifest

BLOOD_GAS_MANIFEST = load_blood_gas_itemid_manifest(BLOOD_GAS_SPEC_PATH)
LAB_PCO2_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["pco2_itemids"]
LAB_EXCLUDED_PCO2_ITEMIDS =
    ↪ BLOOD_GAS_MANIFEST["lab"]["excluded_pco2_itemids"]
LAB_P02_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["po2_itemids"]
LAB_EXCLUDED_P02_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["excluded_po2_itemids"]
if LAB_EXCLUDED_PCO2_ITEMIDS:
    LAB_PCO2_ITEMIDS = sorted(
        set(LAB_PCO2_ITEMIDS).difference(set(LAB_EXCLUDED_PCO2_ITEMIDS)))
)
if not LAB_PCO2_ITEMIDS:
    raise ValueError(
        "All lab.pco2_itemids were excluded by lab.excluded_pco2_itemids. "
        "Update specs/blood_gas_itemids.json."
    )
if LAB_EXCLUDED_P02_ITEMIDS:
    LAB_P02_ITEMIDS = sorted(
        set(LAB_P02_ITEMIDS).difference(set(LAB_EXCLUDED_P02_ITEMIDS)))
)
if not LAB_P02_ITEMIDS:
    raise ValueError(

```

```

    "All lab.po2_itemids were excluded by lab.excluded_po2_itemids. "
    "Update specs/blood_gas_itemids.json."
)
LAB_PH_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["ph_itemids"]
LAB_HCO3_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["hco3_itemids"]
LAB_SPECIMEN_TYPE_ITEMIDS =
    ↳ BLOOD_GAS_MANIFEST["lab"]["specimen_type_itemids"]
LAB_BLOOD_FLUID_TERMS = BLOOD_GAS_MANIFEST["lab"]["blood_fluid_terms"]
ICU_PCO2_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["pc02_itemids"]
ICU_PCO2_ABG_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["pc02_abg_itemids"]
ICU_PCO2_VBG_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["pc02_vbg_itemids"]
ICU_EXCLUDED_PCO2_ITEMIDS =
    ↳ BLOOD_GAS_MANIFEST["icu"]["excluded_pc02_itemids"]
ICU_PO2_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["po2_itemids"]
ICU_PO2_ABG_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["po2_abg_itemids"]
ICU_PO2_VBG_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["po2_vbg_itemids"]
ICU_EXCLUDED_PO2_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["excluded_po2_itemids"]
ICU_PH_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["ph_itemids"]
ICU_HCO3_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["hco3_itemids"]
ICU_SPECIMEN_TYPE_ITEMIDS =
    ↳ BLOOD_GAS_MANIFEST["icu"]["specimen_type_itemids"]
ICU_ALLOW_PATTERN_FALLBACK =
    ↳ bool(BLOOD_GAS_MANIFEST["icu"]["allow_patternFallback"])
ICU_INCLUSION_MODE =
    ↳ str(BLOOD_GAS_MANIFEST["icu"]["inclusionMode"]).strip().lower()
ICU_PCO2_QC = BLOOD_GAS_MANIFEST["icu"]["pc02_qc"]
ICU_LABEL_DENY_TOKENS = BLOOD_GAS_MANIFEST["denylist"]["label_tokens"]
ICU_PCO2_QC_ALLOWED_UOMS = {
    str(value).strip().lower() for value in ICU_PCO2_QC["allowed_uoms"]
}
ICU_PCO2_QC_MEDIAN_MIN = float(ICU_PCO2_QC["median_min"])
ICU_PCO2_QC_MEDIAN_MAX = float(ICU_PCO2_QC["median_max"])
ICU_PCO2_QC_INCLUDE_TOKENS = [
    str(value).strip().lower()
    for value in ICU_PCO2_QC["label_must_include_tokens"]
]
ICU_PCO2_QC_EXCLUDE_TOKENS = [
    str(value).strip().lower()
    for value in ICU_PCO2_QC["label_must_exclude_tokens"]
]

LAB_PCO2_SQL_ARRAY = _sql_int_array(LAB_PCO2_ITEMIDS)
LAB_PO2_SQL_ARRAY = _sql_int_array(LAB_PO2_ITEMIDS)

```

```

LAB_PH_SQL_ARRAY = _sql_int_array(LAB_PH_ITEMIDS)
LAB_HCO3_SQL_ARRAY = _sql_int_array(LAB_HCO3_ITEMIDS)
LAB_SPECIMEN_SQL_ARRAY = _sql_int_array(LAB_SPECIMEN_TYPE_ITEMIDS)
ICU_PCO2_SQL_ARRAY = _sql_int_array(ICU_PCO2_ITEMIDS)
ICU_PCO2_ABG_SQL_ARRAY = _sql_int_array(ICU_PCO2_ABG_ITEMIDS)
ICU_PCO2_VBG_SQL_ARRAY = _sql_int_array(ICU_PCO2_VBG_ITEMIDS)
ICU_P02_SQL_ARRAY = _sql_int_array(ICU_P02_ITEMIDS)
ICU_P02_ABG_SQL_ARRAY = _sql_int_array(ICU_P02_ABG_ITEMIDS)
ICU_P02_VBG_SQL_ARRAY = _sql_int_array(ICU_P02_VBG_ITEMIDS)
ICU_PH_SQL_ARRAY = _sql_int_array(ICU_PH_ITEMIDS)
ICU_HCO3_SQL_ARRAY = _sql_int_array(ICU_HCO3_ITEMIDS)
ICU_SPECIMEN_SQL_ARRAY = _sql_int_array(ICU_SPECIMEN_TYPE_ITEMIDS)
LAB_BLOOD_FLUID_SQL_LIST = ", ".join(f'{term}' for term in
    LAB_BLOOD_FLUID_TERMS)
ICU_LABEL_DENY_REGEX = (
    "|".join(re.escape(token) for token in ICU_LABEL_DENY_TOKENS)
    if ICU_LABEL_DENY_TOKENS
    else _P02_LABEL_DENY_PATTERN_DEFAULT.pattern
)

def _to_int64(series: pd.Series) -> pd.Series:
    """Return pandas nullable integer series for key columns."""
    return pd.to_numeric(series, errors="coerce").astype("Int64")

def normalize_anthro_source(value: object) -> str:
    """Map anthropometric source labels to canonical values."""
    text = "" if pd.isna(value) else str(value).strip().lower()
    if not text or text in {"nan", "none", "missing"}:
        return "missing"
    if text in {"ed", "ed_charted", "ed_triage", "ed_vitalsign"} or
        text.startswith(
            "ed_charted"
        ):
        return "ED"
    if text in {"icu", "icu_charted"} or text.startswith("icu"):
        return "ICU"
    if text in {"hospital", "omr", "hospital.omr"} or text.startswith("omr"):
        return "HOSPITAL"
    return "missing"

```

```

def normalize_anthro_metric_name(value: object) -> str | None:
    """Normalize raw anthropometric metric labels to bmi/height/weight."""
    text = "" if pd.isna(value) else str(value).strip().lower()
    if not text:
        return None
    if "bmi" in text or "body_mass_index" in text:
        return "bmi"
    if "height" in text or text in {"ht"}:
        return "height"
    if "weight" in text or text in {"wt", "admission_weight"}:
        return "weight"
    return None

def normalize_anthro_unit_text(value: object) -> str:
    """Normalize unit text for deterministic matching."""
    text = "" if pd.isna(value) else str(value).strip().lower()
    text = text.replace("²", "2")
    text = re.sub(r"\s+", "", text)
    return text

def infer_omr_unit(metric: str, result_name_raw: object, result_value_raw:
    object) -> str | None:
    """Infer OMR units from result labels/value text when possible."""
    label_text = "" if pd.isna(result_name_raw) else
    str(result_name_raw).strip().lower()
    value_text = "" if pd.isna(result_value_raw) else
    str(result_value_raw).strip().lower()
    text = f"{label_text} {value_text}".strip()
    if metric == "bmi":
        return "kg/m2"
    if metric == "height":
        if "inch" in text or re.search(r"\bin\b", text):
            return "in"
        if "centimeter" in text or re.search(r"\bcm\b", text):
            return "cm"
        if "meter" in text or re.search(r"\bm\b", text):
            return "m"
    if metric == "weight":
        if "kilogram" in text or re.search(r"\bkg\b", text):
            return "kg"

```

```

if "pound" in text or re.search(r"\blb\b", text) or
    re.search(r"\blbs\b", text):
    return "lb"
if "ounce" in text or re.search(r"\boz\b", text):
    return "oz"
if "gram" in text or re.search(r"\bg\b", text):
    return "g"
return None

def prepare_omr_records(omr_df: pd.DataFrame) -> pd.DataFrame:
    """Normalize OMR records into a deterministic schema.

    Required source columns: ``subject_id``, ``chartdate``, ``result_name``,
    ``result_value``.

    Returns:
        DataFrame with columns:
        ``subject_id`` (Int64), ``chartdate_dt``/``obs_time`` (datetime64),
        ``result_name`` (bmi/height/weight), ``result_value_num`` (float),
        ``result_unit_raw`` (string), and ``source`` (HOSPITAL).
    """
    required = {"subject_id", "chartdate", "result_name", "result_value"}
    missing = sorted(required.difference(omr_df.columns))
    if missing:
        raise KeyError(f"prepare_omr_records missing required columns:
        {missing}")

    prepared = omr_df.copy()
    prepared["subject_id"] = _to_int64(prepared["subject_id"])
    prepared["chartdate_dt"] = pd.to_datetime(prepared["chartdate"],
    errors="coerce")
    prepared["result_name_raw"] = prepared["result_name"]
    prepared["result_name"] =
    prepared["result_name_raw"].map(normalize_anthro_metric_name)
    prepared["result_unit_raw"] = [
        infer_omr_unit(metric, name_raw, value_raw)
        for metric, name_raw, value_raw in zip(
            prepared["result_name"],
            prepared["result_name_raw"],
            prepared["result_value"],
            strict=False,
        )
    ]

```

```

        ]
prepared["result_value_num"] = pd.to_numeric(
    prepared["result_value"]
    .astype(str)
    .str.extract(_NUMERIC_TOKEN_PATTERN, expand=False),
    errors="coerce",
)
prepared["source"] = "HOSPITAL"
prepared["obs_time"] = prepared["chartdate_dt"]

prepared =
↪ prepared.loc[prepared["result_name"].isin(OMR_RESULT_NAMES)].copy()
prepared = prepared.loc[
    prepared["subject_id"].notna() & prepared["chartdate_dt"].notna()
].copy()

return prepared[
    [
        "subject_id",
        "chartdate_dt",
        "obs_time",
        "result_name",
        "result_value_num",
        "result_unit_raw",
        "source",
    ]
]
]

def attach_closest_pre_ed_omr(
    ed_df: pd.DataFrame,
    omr_df: pd.DataFrame,
    window_days: int = 365,
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Attach closest anthropometrics with tiered timing fallback.

Args:
    ed_df: ED-stay grain dataframe with ``ed_stay_id``, ``subject_id```,
           and ``ed_intime``.
    omr_df: Prepared OMR records from ``prepare_omr_records``.
    window_days: Inclusion window in days before/after ED arrival.
    """

```

```

def _with_default_anthro_columns(frame: pd.DataFrame) -> pd.DataFrame:
    updated = frame.copy()
    for column_name in OMR_OUTPUT_COLUMNS:
        if column_name not in updated.columns:
            updated[column_name] = pd.NA
    for metric_name, unit_column in ANTHRO_UNIT_COLUMNS.items():
        if unit_column not in updated.columns:
            updated[unit_column] = pd.NA
        else:
            updated[unit_column] = (
                updated[unit_column].astype("string").replace({"":
        ↵  pd.NA})
            )
    for metric_name, time_column in ANTHRO_TIME_COLUMNS.items():
        if time_column not in updated.columns:
            updated[time_column] = pd.NaT
        else:
            updated[time_column] = pd.to_datetime(updated[time_column],
        ↵  errors="coerce")
        if "anthro_timing_tier" not in updated.columns:
            updated["anthro_timing_tier"] = "missing"
        else:
            updated["anthro_timing_tier"] =
        ↵  updated["anthro_timing_tier"].fillna("missing")
        if "anthro_days_offset" not in updated.columns:
            updated["anthro_days_offset"] = pd.Series([pd.NA] * len(updated),
        ↵  dtype="Int64")
        else:
            updated["anthro_days_offset"] = pd.to_numeric(
                updated["anthro_days_offset"], errors="coerce"
            ).astype("Int64")
    if "anthro_chartdate" not in updated.columns:
        updated["anthro_chartdate"] = pd.NaT
    else:
        updated["anthro_chartdate"] = pd.to_datetime(
            updated["anthro_chartdate"], errors="coerce"
        )
    if "anthro_timing_uncertain" not in updated.columns:
        updated["anthro_timing_uncertain"] = pd.Series(
            [pd.NA] * len(updated), dtype="boolean"
        )
    else:
        updated["anthro_timing_uncertain"] = updated[

```

```

        "anthro_timing_uncertain"
    ].astype("boolean")
if "anthro_source" not in updated.columns:
    updated["anthro_source"] = "missing"
else:
    updated["anthro_source"] =
↪ updated["anthro_source"].map(normalize_anthro_source)
    if "anthro_obstime" not in updated.columns:
        updated["anthro_obstime"] = pd.NaT
    else:
        updated["anthro_obstime"] =
↪ pd.to_datetime(updated["anthro_obstime"], errors="coerce")
        if "anthro_hours_offset" not in updated.columns:
            updated["anthro_hours_offset"] = pd.NA
        updated["anthro_hours_offset"] = pd.to_numeric(
            updated["anthro_hours_offset"], errors="coerce"
        )
        if "anthro_timing_basis" not in updated.columns:
            updated["anthro_timing_basis"] = "missing"
        else:
            updated["anthro_timing_basis"] = (
                updated["anthro_timing_basis"].astype(str).replace({"":
↪ "missing"}).fillna("missing")
            )

missing_mask = updated["anthro_timing_tier"].eq("missing")
pre_mask = updated["anthro_timing_tier"].eq("pre_ed_365")
post_mask = updated["anthro_timing_tier"].eq("post_ed_365")
updated.loc[pre_mask, "anthro_timing_uncertain"] = False
updated.loc[post_mask, "anthro_timing_uncertain"] = True
updated.loc[missing_mask, "anthro_timing_uncertain"] = pd.NA
updated.loc[pre_mask | post_mask, "anthro_source"] = "HOSPITAL"
updated.loc[missing_mask, "anthro_source"] = "missing"
updated.loc[pre_mask, "anthro_timing_basis"] = "pre"
updated.loc[post_mask, "anthro_timing_basis"] = "post"
updated.loc[missing_mask, "anthro_timing_basis"] = "missing"
updated.loc[missing_mask, "anthro_obstime"] = pd.NaT
updated.loc[missing_mask, "anthro_hours_offset"] = pd.NA
for metric_name, output_column in ANTHRO_VALUE_COLUMNS.items():
    metric_mask = pre_mask | post_mask
    unit_column = ANTHRO_UNIT_COLUMNS[metric_name]
    time_column = ANTHRO_TIME_COLUMNS[metric_name]

```

```

        updated.loc[metric_mask, unit_column] =
    ↵ ANTHRO_CANONICAL_UNITS[metric_name]
        updated.loc[missing_mask, unit_column] = pd.NA
        updated.loc[metric_mask, time_column] = pd.to_datetime(
            updated.loc[metric_mask, "anthro_obstime"], errors="coerce"
        )
        updated.loc[missing_mask, time_column] = pd.NaT

    return updated

def _tier_counts(frame: pd.DataFrame) -> dict[str, int]:
    tier_series =
    ↵ frame["anthro_timing_tier"].astype(str).fillna("missing")
    return {
        "pre_ed_365": int((tier_series == "pre_ed_365").sum()),
        "post_ed_365": int((tier_series == "post_ed_365").sum()),
        "missing": int((tier_series == "missing").sum()),
    }

required_ed = {"ed_stay_id", "subject_id", "ed_intime"}
missing_ed = sorted(required_ed.difference(ed_df.columns))
if missing_ed:
    raise KeyError(f"attach_closest_pre_ed.omr missing ED columns:
    ↵ {missing_ed}")

required_omr = {"subject_id", "chartdate_dt", "result_name",
    ↵ "result_value_num"}
missing_omr = sorted(required_omr.difference(omr_df.columns))
if missing_omr:
    raise KeyError(f"attach_closest_pre_ed.omr missing OMR columns:
    ↵ {missing_omr}")

diagnostics: dict[str, Any] = {
    "window_days": int(window_days),
    "source_rows_prepared": int(len(omr_df)),
    "parsed_value_rows": int(omr_df["result_value_num"].notna().sum()),
    "ed_rows": int(len(ed_df)),
}
empty_window_diagnostics = {
    "nonnegative_candidate_rows": 0,
    "pre_window_candidate_rows": 0,
    "post_window_candidate_rows": 0,
}

```

```

        "closest_absolute_candidate_rows": 0,
        "within_window_candidate_rows": 0,
    }

    ed_norm = ed_df[["ed_stay_id", "subject_id", "ed_intime"]].copy()
    ed_norm["subject_id"] = _to_int64(ed_norm["subject_id"])
    ed_norm["ed_intime_dt"] = pd.to_datetime(ed_norm["ed_intime"],
    ↪ errors="coerce")
    ed_norm["ed_date_dt"] = ed_norm["ed_intime_dt"].dt.floor("D")
    ed_norm = ed_norm.loc[ed_norm["subject_id"].notna() &
    ↪ ed_norm["ed_date_dt"].notna()]

    if omr_df.empty or ed_norm.empty:
        updated = _with_default_anthro_columns(ed_df)
        tier_counts = _tier_counts(updated)
        total_rows = max(int(len(updated)), 1)
        diagnostics.update(
            {
                "ed_rows_eligible_for_join": int(len(ed_norm)),
                "subject_overlap_count": 0,
                "candidate_rows_after_subject_join": 0,
                **empty_window_diagnostics,
                "within_window_candidate_rows": 0,
                "eligible_ed_stays_with_candidates": 0,
                "attached_non_null_counts": {
                    column_name: int(updated[column_name].notna().sum())
                    for column_name in OMR_OUTPUT_COLUMNS
                },
                "attached_any_non_null_rows": int(
                    updated[list(OMR_OUTPUT_COLUMNS)].notna().any(axis=1).sum()
                ),
                "selected_tier_counts": tier_counts,
                "selected_tier_rates": {
                    key: float(value / total_rows) for key, value in
    ↪ tier_counts.items()
                },
                "timing_uncertain_count": int(
                    updated["anthro_timing_uncertain"].fillna(False).sum()
                ),
            }
        )
    return updated, diagnostics

```

```

omr_pivot = (
    omr_df.pivot_table(
        index=["subject_id", "chartdate_dt"],
        columns="result_name",
        values="result_value_num",
        aggfunc="first",
    )
    .reset_index()
    .copy()
)

shared_subjects =
↪ set(ed_norm["subject_id"].dropna().astype(int)).intersection(
    set(omr_pivot["subject_id"].dropna().astype(int))
)
diagnostics["ed_rows_eligible_for_join"] = int(len(ed_norm))
diagnostics["subject_overlap_count"] = int(len(shared_subjects))

merged = ed_norm.merge(omr_pivot, on="subject_id", how="left")
diagnostics["candidate_rows_after_subject_join"] = int(
    merged["chartdate_dt"].notna().sum()
)

merged["days_before"] = (
    merged["ed_date_dt"] - pd.to_datetime(merged["chartdate_dt"]),
↪ errors="coerce"
).dt.days
valid_days = merged["days_before"].dropna()
diagnostics["days_before_min"] = (
    int(valid_days.min()) if not valid_days.empty else None
)
diagnostics["days_before_max"] = (
    int(valid_days.max()) if not valid_days.empty else None
)
pre_window_mask = (
    merged["days_before"].notna()
    & merged["days_before"].ge(0)
    & merged["days_before"].le(window_days)
)
post_window_mask = (
    merged["days_before"].notna()
    & merged["days_before"].lt(0)
)

```

```

    & merged["days_before"].ge(-window_days)
)
abs_window_mask = (
    merged["days_before"].notna()
    & merged["days_before"].abs().le(window_days)
)

diagnostics["nonnegative_candidate_rows"] = int((merged["days_before"] >=
0).sum())
diagnostics["pre_window_candidate_rows"] = int(pre_window_mask.sum())
diagnostics["post_window_candidate_rows"] = int(post_window_mask.sum())
diagnostics["closest_absolute_candidate_rows"] =
int(abs_window_mask.sum())

pre_candidates = merged.loc[pre_window_mask].copy()
post_candidates = merged.loc[post_window_mask].copy()
diagnostics["within_window_candidate_rows"] = int(len(pre_candidates))

selected_parts: list[pd.DataFrame] = []
selected_stays: set[Any] = set()

if not pre_candidates.empty:
    pre_selected = (
        pre_candidates.sort_values(
            ["ed_stay_id", "days_before", "chartdate_dt"],
            ascending=[True, True, False],
        )
        .groupby("ed_stay_id", as_index=False)
        .first()
    )
    pre_selected["anthro_timing_tier"] = "pre_ed_365"
    selected_parts.append(pre_selected)
    selected_stays.update(pre_selected["ed_stay_id"].tolist())

if not post_candidates.empty:
    post_candidates["days_after_abs"] =
    post_candidates["days_before"].abs()
    post_selected = (
        post_candidates.sort_values(
            ["ed_stay_id", "days_after_abs", "chartdate_dt"],
            ascending=[True, True, True],
        )
        .groupby("ed_stay_id", as_index=False)

```

```

        .first()
    )
post_selected = post_selected.loc[
    ~post_selected["ed_stay_id"].isin(selected_stays)
].copy()
if not post_selected.empty:
    post_selected["anthro_timing_tier"] = "post_ed_365"
    selected_parts.append(post_selected)

if selected_parts:
    selected = pd.concat(selected_parts, ignore_index=True)
    selected = selected.rename(
        columns={
            "bmi": "bmi_closest_pre_ed",
            "height": "height_closest_pre_ed",
            "weight": "weight_closest_pre_ed",
            "chartdate_dt": "anthro_chartdate",
            "days_before": "anthro_days_offset",
        }
    )
    selected["anthro_chartdate"] = pd.to_datetime(
        selected["anthro_chartdate"], errors="coerce"
    )
    selected["anthro_days_offset"] = pd.to_numeric(
        selected["anthro_days_offset"], errors="coerce"
    ).astype("Int64")
    selected["anthro_timing_uncertain"] =
    ↪ selected["anthro_timing_tier"].eq(
        "post_ed_365"
    )
    selected["anthro_source"] = "HOSPITAL"
    selected["anthro_obstime"] = pd.to_datetime(
        selected["anthro_chartdate"], errors="coerce"
    )
    selected["anthro_hours_offset"] = pd.to_numeric(
        selected["anthro_days_offset"], errors="coerce"
    ) * 24.0
    selected["anthro_timing_basis"] = selected["anthro_timing_tier"].map(
        {"pre_ed_365": "pre", "post_ed_365": "post"}
    ).fillna("missing")
    selected["bmi_closest_pre_ed_uom"] = ANTHRO_CANONICAL_UNITS["bmi"]
    selected["height_closest_pre_ed_uom"] =
    ↪ ANTHRO_CANONICAL_UNITS["height"]

```

```

        selected["weight_closest_pre_ed_uom"] =
    ↵ ANTHRO_CANONICAL_UNITS["weight"]
        selected["bmi_closest_pre_ed_time"] = pd.to_datetime(
            selected["anthro_obstime"], errors="coerce"
        )
        selected["height_closest_pre_ed_time"] = pd.to_datetime(
            selected["anthro_obstime"], errors="coerce"
        )
        selected["weight_closest_pre_ed_time"] = pd.to_datetime(
            selected["anthro_obstime"], errors="coerce"
        )

    updates = selected[
        [
            "ed_stay_id",
            *OMR_OUTPUT_COLUMNS,
            *ANTHRO_UNIT_COLUMNS.values(),
            *ANTHRO_TIME_COLUMNS.values(),
            *OMR_PROVENANCE_COLUMNS,
        ]
    ].copy()
    updated = ed_df.merge(updates, on="ed_stay_id", how="left")
else:
    updated = ed_df.copy()

updated = _with_default_anthro_columns(updated)

diagnostics["eligible_ed_stays_with_candidates"] = int(
    updated["anthro_timing_tier"].isin({"pre_ed_365",
    ↵ "post_ed_365"}).sum()
)
diagnostics["attached_non_null_counts"] = {
    column_name: int(updated[column_name].notna().sum())
    for column_name in OMR_OUTPUT_COLUMNS
}
diagnostics["attached_any_non_null_rows"] = int(
    updated[list(OMR_OUTPUT_COLUMNS)].notna().any(axis=1).sum()
)
tier_counts = _tier_counts(updated)
diagnostics["selected_tier_counts"] = tier_counts
total_rows = max(int(len(updated)), 1)
diagnostics["selected_tier_rates"] = {
    key: float(value / total_rows) for key, value in tier_counts.items()
}

```

```

    }

    diagnostics["timing_uncertain_count"] = int(
        updated["anthro_timing_uncertain"].fillna(False).sum()
    )

    diagnostics["anthro_source_counts"] = {
        str(key): int(value)
        for key, value in updated["anthro_source"].fillna("missing")
            .astype(str)
            .value_counts(dropna=False)
            .items()
    }

}

return updated, diagnostics


def _infer_route_hint_text(value: object) -> str | None:
    text = "" if pd.isna(value) else str(value).strip().lower()
    if not text or text in {"nan", "none"}:
        return None
    if _ARTERIAL_HINT_PATTERN.search(text):
        return "arterial"
    if _VENOUS_HINT_PATTERN.search(text):
        return "venous"
    return None


def _resolve_route_hints(values: Sequence[str]) -> tuple[str | None, bool,
    int]:
    hints = [str(value).strip().lower() for value in values if
    str(value).strip()]
    arterial_n = sum(1 for value in hints if value == "arterial")
    venous_n = sum(1 for value in hints if value == "venous")
    conflict = arterial_n > 0 and venous_n > 0
    if conflict:
        return None, True, int(len(hints))
    if arterial_n > 0:
        return "arterial", False, int(len(hints))
    if venous_n > 0:
        return "venous", False, int(len(hints))
    return None, False, int(len(hints))

def _classify_specimen_type_text(value: object) -> str | None:

```

```

text = "" if pd.isna(value) else str(value).strip().lower()
if not text:
    return None
if re.search(r"arter|(?:^|\b)art(?:\b|$)", text):
    return "arterial"
if re.search(r"ven|mixed|central", text):
    return "venous"
return None

def assign_panel_source_from_specimen(
    panel_df: pd.DataFrame,
    labs_df: pd.DataFrame,
    *,
    specimen_source_itemids: Sequence[int] | None = None,
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Assign gas source using specimen-type rows only (deterministic core
    ↵ policy)."""
    required_panel = {"ed_stay_id", "specimen_id"}
    missing_panel = sorted(required_panel.difference(panel_df.columns))
    if missing_panel:
        raise KeyError(
            f"assign_panel_source_from_specimen missing panel columns: "
            f"{missing_panel}"
        )
    required_labs = {"ed_stay_id", "specimen_id", "itemid"}
    missing_labs = sorted(required_labs.difference(labs_df.columns))
    if missing_labs:
        raise KeyError(
            f"assign_panel_source_from_specimen missing labs columns: "
            f"{missing_labs}"
        )

    key_cols = ["ed_stay_id", "specimen_id"]
    panel = panel_df.copy()
    panel["source"] = "unknown"
    panel["source_inference_tier"] = "fallback_unknown"
    panel["source_hint_conflict"] = False
    panel["source_hint_count"] = pd.Series(0, index=panel.index,
    ↵ dtype="Int64")
    panel["source_hint_specimen_text"] = pd.Series(pd.NA, index=panel.index,
    ↵ dtype="string")

```

```

specimen_ids = [
    int(value) for value in (specimen_source_itemids or []) if
        pd.notna(value)
]
if not specimen_ids or "value_text" not in labs_df.columns:
    diagnostics = summarize_gas_source(panel)
    diagnostics.update(
        {
            "mode": "specimen_only",
            "resolved_rows": int(panel["source"].isin({"arterial",
                "venous"}).sum()),
            "resolved_rate": float(panel["source"].isin({"arterial",
                "venous"}).mean())
            if len(panel)
            else 0.0,
            "unresolved_specimen_id_count": int(
                pd.to_numeric(panel["specimen_id"]),
                errors="coerce").nunique(dropna=True)
            ),
            "unresolved_specimen_id_examples": [],
        }
    )
    return panel, diagnostics

labs = labs_df.copy()
labs["itemid"] = pd.to_numeric(labs["itemid"],
    errors="coerce").astype("Int64")
labs["specimen_id"] = pd.to_numeric(labs["specimen_id"],
    errors="coerce").astype("Int64")
labs = labs.loc[labs["itemid"].isin(specimen_ids) &
    labs["specimen_id"].notna()].copy()
if labs.empty:
    diagnostics = summarize_gas_source(panel)
    diagnostics.update(
        {
            "mode": "specimen_only",
            "resolved_rows": int(panel["source"].isin({"arterial",
                "venous"}).sum()),
            "resolved_rate": float(panel["source"].isin({"arterial",
                "venous"}).mean())
            if len(panel)
            else 0.0,
            "unresolved_specimen_id_count": int(

```

```

        pd.to_numeric(panel["specimen_id"],
←   errors="coerce").nunique(dropna=True)
    ),
    "unresolved_specimen_id_examples": [] ,
}
)
return panel, diagnostics

labs["specimen_text_raw"] = (
    labs["value_text"].astype("string").fillna("").str.strip()
)
labs["specimen_text_norm"] = labs["specimen_text_raw"].str.lower()
labs["specimen_route_hint"] =
← labs["specimen_text_norm"].map(_classify_specimen_type_text)

grouped = (
    labs.groupby(key_cols, dropna=False)
    .agg(
        route_hints=("specimen_route_hint", list),
        specimen_texts=("specimen_text_raw", list),
    )
    .reset_index()
)

mapping_rows: list[dict[str, Any]] = []
for _, grouped_row in grouped.iterrows():
    source, conflict, hint_count =
← _resolve_route_hints(grouped_row["route_hints"])
    specimen_texts = [
        str(value).strip()
        for value in grouped_row["specimen_texts"]
        if str(value).strip()
    ]
    specimen_text = specimen_texts[0] if specimen_texts else None
    mapping_rows.append(
    {
        "ed_stay_id": grouped_row["ed_stay_id"],
        "specimen_id": grouped_row["specimen_id"],
        "source": source if source is not None else "unknown",
        "source_inference_tier": "specimen_text" if source is not
        ← None else "fallback_unknown",
        "source_hint_conflict": bool(conflict),
        "source_hint_count": int(hint_count),
    }
)

```

```

        "source_hint_specimen_text": specimen_text,
    }
)

mapping = pd.DataFrame(mapping_rows)
if not mapping.empty:
    panel = panel.merge(mapping, on=key_cols, how="left", suffixes=("",
"specimen"))
    panel["source"] = (
        panel.get("source_specimen")
        .fillna(panel["source"])
        .astype("string")
        .fillna("unknown")
    )
    panel["source_inference_tier"] = (
        panel.get("source_inference_tier_specimen")
        .fillna(panel["source_inference_tier"])
        .astype("string")
        .fillna("fallback_unknown")
    )
    panel["source_hint_conflict"] = (
        panel.get("source_hint_conflict_specimen")
        .fillna(panel["source_hint_conflict"])
        .astype(bool)
    )
    panel["source_hint_count"] = pd.to_numeric(
        panel.get("source_hint_count_specimen").fillna(panel[
"source_hint_count"]),
        errors="coerce",
    ).fillna(0).astype("Int64")
    panel["source_hint_specimen_text"] = (
        panel.get("source_hint_specimen_text_specimen")
        .fillna(panel["source_hint_specimen_text"])
        .astype("string")
    )
    panel = panel.drop(
        columns=[
            column_name
            for column_name in (
                "source_specimen",
                "source_inference_tier_specimen",
                "source_hint_conflict_specimen",
                "source_hint_count_specimen",
            )
        ]
    )

```

```

        "source_hint_specimen_text_specimen",
    )
    if column_name in panel.columns
]
)

diagnostics = summarize_gas_source(panel)
unresolved = panel.loc[panel["source"].eq("unknown"),
↪ key_cols].drop_duplicates()
diagnostics.update(
{
    "mode": "specimen_only",
    "resolved_rows": int(panel["source"].isin({"arterial",
↪ "venous"}).sum()),
    "resolved_rate": float(panel["source"].isin({"arterial",
↪ "venous"}).mean())
    if len(panel)
    else 0.0,
    "unresolved_specimen_id_count": int(
        pd.to_numeric(unresolved["specimen_id"],
↪ errors="coerce").nunique(dropna=True)
    ),
    "unresolved_specimen_id_examples": unresolved["specimen_id"]
        .dropna()
        .astype(int)
        .head(20)
        .tolist(),
}
)
return panel, diagnostics

def infer_panel_gas_source_metadata(
    panel_df: pd.DataFrame,
    labs_df: pd.DataFrame,
    labitems_df: pd.DataFrame,
    *,
    specimen_source_itemids: Sequence[int] | None = None,
    pco2_itemids: Sequence[int] | None = None,
    mode: str = "metadata_only",
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Infer gas panel source from metadata/text hints with tier diagnostics.

```

```

Tier precedence:
    1. Specimen/source row text (``value_text`` from specimen/source
   ↵ itemids),
    2. ``d_labitems`` label/fluid hints,
    3. Panel co-occurrence hints (pCO2-labeled rows + free-text),
    4. Fallback ``unknown``.

"""
if mode != "metadata_only":
    raise ValueError(
        "infer_panel_gas_source_metadata supports only
         ↵ mode='metadata_only'."
    )

required_panel = {"ed_stay_id", "specimen_id"}
missing_panel = sorted(required_panel.difference(panel_df.columns))
if missing_panel:
    raise KeyError(
        "infer_panel_gas_source_metadata missing panel columns: "
        f"{missing_panel}"
    )
required_labs = {"ed_stay_id", "specimen_id", "itemid"}
missing_labs = sorted(required_labs.difference(labs_df.columns))
if missing_labs:
    raise KeyError(
        "infer_panel_gas_source_metadata missing labs columns: "
        f"{missing_labs}"
    )
required_labitems = {"itemid"}
missing_labitems =
   ↵ sorted(required_labitems.difference(labitems_df.columns))
if missing_labitems:
    raise KeyError(
        "infer_panel_gas_source_metadata missing labitems columns: "
        f"{missing_labitems}"
    )

key_cols = ["ed_stay_id", "specimen_id"]
panel = panel_df.copy()
panel["source"] = pd.Series(pd.NA, index=panel.index, dtype="string")
panel["source_inference_tier"] = "fallback_unknown"
panel["source_hint_conflict"] = pd.Series(False, index=panel.index,
   ↵ dtype="boolean")

```

```

panel["source_hint_count"] = pd.Series(pd.NA, index=panel.index,
                                         dtype="Int64")

    ↵ labs = labs_df.copy()
    ↵ labs["itemid"] = pd.to_numeric(labs["itemid"],
                                     errors="coerce").astype("Int64")
    ↵ labs["specimen_id"] = pd.to_numeric(labs["specimen_id"],
                                         errors="coerce").astype(
                                             "Int64"
                                         )
    ↵ labs = labs.loc[labs["itemid"].notna() &
    ↵ labs["specimen_id"].notna()].copy()
    if labs.empty:
        panel["source"] = "unknown"
        diagnostics = summarize_gas_source(panel)
        diagnostics.update(
            {
                "mode": mode,
                "resolved_rows": 0,
                "resolved_rate": 0.0,
                "unresolved_specimen_id_count": 0,
                "unresolved_specimen_id_examples": [],
                "unresolved_value_text_top": {},
                "unresolved_label_top": {}
            }
        )
    return panel, diagnostics

if "value_text" in labs.columns:
    labs["value_text_norm"] = (
        labs["value_text"]
        .astype("string")
        .fillna("")
        .str.strip()
        .str.lower()
    )
else:
    labs["value_text_norm"] = ""

labitems = labitems_df.copy()
labitems["itemid"] = pd.to_numeric(labitems["itemid"],
                                     errors="coerce").astype(
                                         "Int64"

```

```

)
for text_col in ("label", "fluid"):
    if text_col not in labitems.columns:
        labitems[text_col] = ""
labitems["label_norm"] =
    labitems["label"].astype("string").fillna("").str.strip()
labitems["fluid_norm"] =
    labitems["fluid"].astype("string").fillna("").str.strip()
labitems["item_route_hint"] = (
    (labitems["label_norm"] + " " + labitems["fluid_norm"])
    .str.strip()
    .map(_infer_route_hint_text)
)
item_meta = (
    labitems.dropna(subset=["itemid"])
    .drop_duplicates(subset=["itemid"])
    .set_index("itemid")[["label_norm", "fluid_norm", "item_route_hint"]]
)
labs = labs.merge(item_meta, left_on="itemid", right_index=True,
    how="left")
labs["item_route_hint"] = labs["item_route_hint"].astype("string")
labs["text_route_hint"] =
    labs["value_text_norm"].map(_infer_route_hint_text).astype(
        "string"
)
labs["label_fluid_text"] = (
    labs["label_norm"].fillna("").astype(str).str.strip()
    + " | "
    + labs["fluid_norm"].fillna("").astype(str).str.strip()
).str.strip(" |")

def _build_tier_mapping(
    frame: pd.DataFrame,
    *,
    hint_column: str,
    tier_name: str,
) -> pd.DataFrame:
    subset = frame[key_cols + [hint_column]].copy()
    subset = subset.loc[subset[hint_column].notna()].copy()
    if subset.empty:
        return pd.DataFrame(
            columns=[
                *key_cols,

```

```

        "tier_source",
        "tier_conflict",
        "tier_hint_count",
        "tier_name",
    ]
)
grouped = (
    subset.groupby(key_cols, dropna=False)[hint_column]
    .agg(list)
    .reset_index(name="hints")
)
rows: list[dict[str, Any]] = []
for _, grouped_row in grouped.iterrows():
    source, conflict, hint_count =
    ↪ _resolve_route_hints(grouped_row["hints"])
    rows.append(
        {
            "ed_stay_id": grouped_row["ed_stay_id"],
            "specimen_id": grouped_row["specimen_id"],
            "tier_source": source,
            "tier_conflict": bool(conflict),
            "tier_hint_count": int(hint_count),
            "tier_name": tier_name,
        }
    )
resolved = pd.DataFrame(rows)
return resolved.loc[resolved["tier_source"] |
    ↪ ].notna()].reset_index(drop=True)

specimen_source_itemid_set = {
    int(item) for item in (specimen_source_itemids or [])
    ↪ pd.notna(item)
}
pco2_itemid_set = {int(item) for item in (pco2_itemids or [])
    ↪ pd.notna(item)}

tier1 = _build_tier_mapping(
    labs.loc[labs["itemid"].astype("Int64") |
    ↪ ).isin(specimen_source_itemid_set)],
    hint_column="text_route_hint",
    tier_name="specimen_text",
)

```

```

tier2 = _build_tier_mapping(
    labs,
    hint_column="item_route_hint",
    tier_name="label_fluid",
)
tier_cluster = pd.DataFrame(
    columns=[
        *key_cols,
        "tier_source",
        "tier_conflict",
        "tier_hint_count",
        "tier_name",
    ]
)
if "panel_time" in panel_df.columns:
    panel_times = panel_df[key_cols + ["panel_time"]].copy()
    panel_times["panel_time"] = pd.to_datetime(panel_times["panel_time"],
    ↵ errors="coerce")

    tier_seed = pd.concat([tier1, tier2], ignore_index=True)
    if not tier_seed.empty:
        tier_seed = (
            tier_seed[key_cols + ["tier_source"]]
            .dropna(subset=["tier_source"])
            .drop_duplicates(subset=key_cols)
        )
        panel_seed = panel_times.merge(tier_seed, on=key_cols,
    ↵ how="left")
        known = panel_seed.loc[panel_seed["tier_source"].notna()].copy()
        unknown = panel_seed.loc[panel_seed["tier_source"].isna()].copy()
        if not known.empty and not unknown.empty:
            pairs = unknown.merge(
                known[["ed_stay_id", "specimen_id", "panel_time",
    ↵ "tier_source"]],
                on="ed_stay_id",
                how="left",
                suffixes=("_unknown", "_known"),
            )
            pairs["delta_minutes"] = (
                pairs["panel_time_known"] -
    ↵ pairs["panel_time_unknown"]).abs().dt.total_seconds()
                / 60.0
            )

```

```

        pairs = pairs.loc[pairs["delta_minutes"].notna() &
↪  pairs["delta_minutes"].le(60.0)].copy()
        if not pairs.empty:
            keys = ["ed_stay_id", "specimen_id_unknown"]
            min_delta =
↪  pairs.groupby(keys)["delta_minutes"].transform("min")
            nearest =
↪  pairs.loc[pairs["delta_minutes"].eq(min_delta)].copy()
            conflict_stats = (
                nearest.groupby(keys)["tier_source_known"]
                .nunique(dropna=True)
                .reset_index(name="source_nunique")
            )
            nearest = nearest.merge(conflict_stats, on=keys,
↪  how="left")
            nearest =
↪  nearest.loc[nearest["source_nunique"].eq(1)].copy()
            if not nearest.empty:
                nearest = nearest.sort_values(keys +
↪  ["delta_minutes"])
                nearest = nearest.groupby(keys,
↪  as_index=False).first()
                tier_cluster = nearest.rename(
                    columns={
                        "specimen_id_unknown": "specimen_id",
                        "tier_source_known": "tier_source",
                    }
                )[["ed_stay_id", "specimen_id", "tier_source"]]
                .copy()
                tier_cluster["tier_conflict"] = False
                tier_cluster["tier_hint_count"] = 1
                tier_cluster["tier_name"] = "cluster_inheritance"

tier3_candidates = labs.copy()
if pco2_itemid_set:
    tier3_candidates = tier3_candidates.loc[
        tier3_candidates["itemid"].astype("Int64").isin(pco2_itemid_set)
        | tier3_candidates["text_route_hint"].notna()
    ].copy()
    tier3_candidates["panel_route_hint"] =
↪  tier3_candidates["item_route_hint"].fillna(
    tier3_candidates["text_route_hint"]

```

```

)
tier3 = _build_tier_mapping(
    tier3_candidates,
    hint_column="panel_route_hint",
    tier_name="panel_cooccurrence",
)
for tier_frame in (tier1, tier2, tier_cluster, tier3):
    if tier_frame.empty:
        continue
    panel = panel.merge(
        tier_frame[
            key_cols
            + ["tier_source", "tier_conflict", "tier_hint_count",
               "tier_name"]
        ],
        on=key_cols,
        how="left",
    )
    assign_mask = panel["source"].isna() & panel["tier_source"].notna()
    panel.loc[assign_mask, "source"] = panel.loc[assign_mask,
        "tier_source"]
    panel.loc[assign_mask, "source_inference_tier"] = panel.loc[
        assign_mask, "tier_name"
    ]
    selected_conflict = panel.loc[assign_mask,
        "tier_conflict"].astype("boolean")
    panel.loc[assign_mask, "source_hint_conflict"] =
        selected_conflict.fillna(False)
    panel.loc[assign_mask, "source_hint_count"] = pd.to_numeric(
        panel.loc[assign_mask, "tier_hint_count"], errors="coerce"
    ).astype("Int64")
    panel = panel.drop(
        columns=["tier_source", "tier_conflict", "tier_hint_count",
        "tier_name"]
    )
    panel["source"] = panel["source"].fillna("unknown").astype("string")
    panel.loc[panel["source"].eq("unknown"), "source_inference_tier"] =
        panel.loc[
            panel["source"].eq("unknown"), "source_inference_tier"
        ].replace({"": "fallback_unknown"}).fillna("fallback_unknown")

```

```

panel["source_inference_tier"] =
↪ panel["source_inference_tier"].astype("string")
panel["source_hint_conflict"] = (
    panel["source_hint_conflict"].fillna(False).astype("boolean")
)

diagnostics = summarize_gas_source(panel)
total = max(int(len(panel)), 1)
diagnostics["mode"] = mode
diagnostics["resolved_rows"] = int(
    panel["source"].isin({"arterial", "venous"}).sum()
)
diagnostics["resolved_rate"] = float(diagnostics["resolved_rows"] /
total)

unresolved = panel.loc[panel["source"].eq("unknown"),
↪ key_cols].drop_duplicates()
diagnostics["unresolved_specimen_id_count"] =
↪ int(unresolved["specimen_id"].nunique())
diagnostics["unresolved_specimen_id_examples"] = (
    unresolved["specimen_id"].dropna().astype(int).head(20).tolist()
)
if unresolved.empty:
    diagnostics["unresolved_value_text_top"] = {}
    diagnostics["unresolved_label_top"] = {}
return panel, diagnostics

unresolved_labs = labs.merge(unresolved, on=key_cols, how="inner")
unresolved_value_text = (
    unresolved_labs["value_text_norm"]
    .replace({"": pd.NA})
    .dropna()
    .astype(str)
    .value_counts()
    .head(15)
)
unresolved_label_text = (
    unresolved_labs["label_fluid_text"]
    .replace({"": pd.NA})
    .dropna()
    .astype(str)
    .value_counts()
    .head(15)
)

```

```

        )
    diagnostics["unresolved_value_text_top"] = {
        str(key): int(value) for key, value in unresolved_value_text.items()
    }
    diagnostics["unresolved_label_top"] = {
        str(key): int(value) for key, value in unresolved_label_text.items()
    }
    return panel, diagnostics

def summarize_gas_source(panel_df: pd.DataFrame) -> dict[str, Any]:
    """Summarize gas source composition from panel-level records."""
    total_rows = int(len(panel_df))
    if total_rows == 0:
        return {
            "panel_rows": 0,
            "source_present": bool("source" in panel_df.columns),
            "source_counts": {},
            "source_rates": {},
            "all_other_or_unknown": False,
            "tier_counts": {},
            "tier_rates": {},
        }
    if "source" not in panel_df.columns:
        return {
            "panel_rows": total_rows,
            "source_present": False,
            "source_counts": {},
            "source_rates": {},
            "all_other_or_unknown": True,
            "tier_counts": {},
            "tier_rates": {},
        }
    source = (
        panel_df["source"]
        .astype(str)
        .str.strip()
        .str.lower()
        .replace({"": "unknown", "nan": "unknown", "none": "unknown",
        "other": "unknown"})
        .fillna("unknown")
    )

```

```

)
source_counts = {str(key): int(value) for key, value in
    source.value_counts(dropna=False).items()}
source_rates = {key: float(value / total_rows) for key, value in
    source_counts.items()}
all_other_or_unknown = bool(total_rows > 0 and
    set(source_counts.keys()).issubset({"unknown"}))
tier_counts: dict[str, int] = {}
tier_rates: dict[str, float] = {}
if "source_inference_tier" in panel_df.columns:
    tier = (
        panel_df["source_inference_tier"]
        .astype(str)
        .str.strip()
        .str.lower()
        .replace({"": "unknown", "nan": "unknown", "none": "unknown"})
        .fillna("unknown")
    )
    tier_counts = {
        str(key): int(value) for key, value in
            tier.value_counts(dropna=False).items()
    }
    tier_rates = {key: float(value / total_rows) for key, value in
        tier_counts.items()}

return {
    "panel_rows": total_rows,
    "source_present": True,
    "source_counts": source_counts,
    "source_rates": source_rates,
    "all_other_or_unknown": all_other_or_unknown,
    "tier_counts": tier_counts,
    "tier_rates": tier_rates,
}
}

def assert_gas_source_coverage(
    gas_source_audit: Mapping[str, Any],
    *,
    fail_on_all_other_source: bool = True,
) -> None:
    """Fail fast when source attribution collapses to all other/unknown."""
    if not fail_on_all_other_source:

```

```

    return
if int(gas_source_audit.get("panel_rows", 0)) <= 0:
    return
if bool(gas_source_audit.get("all_other_or_unknown", False)):
    raise ValueError(
        "Gas source attribution classified all panel rows as
        ↪ other/unknown.
        "Set COHORT_FAIL_ON_ALL_OTHER_SOURCE=0 to bypass this guard."
    )

def build_gas_source_overlap_summary(ed_df: pd.DataFrame) -> pd.DataFrame:
    """Build ABG/VBG/OTHER overlap counts and percentages."""
    frame = ed_df.copy()
    abg = pd.to_numeric(frame.get("abg_hypcap_threshold", 0),
    ↪ errors="coerce").fillna(0).astype(int)
    vbg = pd.to_numeric(frame.get("vbg_hypcap_threshold", 0),
    ↪ errors="coerce").fillna(0).astype(int)
    other = pd.to_numeric(frame.get("unknown_hypcap_threshold", 0),
    ↪ errors="coerce").fillna(0).astype(int)

    labels = []
    for a, v, o in zip(abg.tolist(), vbg.tolist(), other.tolist(),
    ↪ strict=False):
        parts: list[str] = []
        if a == 1:
            parts.append("ABG")
        if v == 1:
            parts.append("VBG")
        if o == 1:
            parts.append("OTHER")
        labels.append("+".join(parts) if parts else "NO_GAS")

    counts = pd.Series(labels, dtype="string").value_counts(dropna=False,
    ↪ ).rename_axis("gas_overlap").reset_index(name="count")
    total = max(int(counts["count"].sum()), 1)
    counts["percent"] = counts["count"].astype(float) / total * 100.0
    counts["percent"] = counts["percent"].round(2)
    return counts.sort_values(["count", "gas_overlap"], ascending=[False,
    ↪ True]).reset_index(drop=True)

def add_gas_model_fields(

```

```

ed_df: pd.DataFrame,
*,
ranges: Mapping[str, tuple[float, float]] | None = None,
) -> tuple[pd.DataFrame, pd.DataFrame]:
    """Create cleaned gas model fields and a field-level outlier audit."""
    resolved_ranges = dict(ranges or DEFAULT_GAS_MODEL_RANGES)
    updated = ed_df.copy()
    audit_rows: list[dict[str, Any]] = []

    for raw_column, (lower_bound, upper_bound) in resolved_ranges.items():
        if raw_column not in updated.columns:
            continue
        numeric = pd.to_numeric(updated[raw_column], errors="coerce")
        out_of_range = numeric.notna() & ((numeric < lower_bound) | (numeric
        < upper_bound))
        model_column = f"{raw_column}_model"
        outlier_flag_column = f"{raw_column}_outlier_flag"
        updated[model_column] = numeric.where(~out_of_range)
        updated[outlier_flag_column] = out_of_range.astype("boolean")
        nonnull_n = int(numeric.notna().sum())
        outlier_n = int(out_of_range.sum())
        examples = (
            numeric.loc[out_of_range]
            .round(4)
            .value_counts()
            .head(5)
            .index.astype(str)
            .tolist()
        )
        audit_rows.append(
            {
                "domain": "gas",
                "raw_column": raw_column,
                "model_column": model_column,
                "outlier_flag_column": outlier_flag_column,
                "lower_bound": float(lower_bound),
                "upper_bound": float(upper_bound),
                "nonnull_n": nonnull_n,
                "out_of_range_n": outlier_n,
                "out_of_range_pct": float(outlier_n / nonnull_n) if nonnull_n
                < else 0.0,
                "example_outlier_values": "; ".join(examples),
            }
        )

```

```

        )

audit =
↳ pd.DataFrame(audit_rows).sort_values("raw_column").reset_index(drop=True)
return updated, audit

def attach_chARTed_anthro_fallback(
    ed_df: pd.DataFrame,
    charted_df: pd.DataFrame,
    *,
    nearest_anytime: bool = True,
    source_preference: Sequence[str] = ("ED", "ICU", "HOSPITAL"),
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Resolve anthropometrics from candidate rows with strict unit
    handling."""
    required_ed = {"ed_stay_id", "subject_id", "ed_intime"}
    missing_ed = sorted(required_ed.difference(ed_df.columns))
    if missing_ed:
        raise KeyError(f"attach_chARTed_anthro_fallback missing ED columns:
        {missing_ed}")

    required_chARTed = {
        "subject_id",
        "obs_time",
        "result_name",
        "result_value_num",
        "result_unit_raw",
        "source",
    }
    missing_chARTed = sorted(required_chARTed.difference(chARTed_df.columns))
    if missing_chARTed:
        raise KeyError(
            f"attach_chARTed_anthro_fallback missing chARTed columns:
            {missing_chARTed}"
        )

    updated = ed_df.copy()
    for metric_name, output_column in ANTHRO_VALUE_COLUMNS.items():
        if output_column not in updated.columns:
            updated[output_column] = pd.NA
        unit_column = ANTHRO_UNIT_COLUMNS[metric_name]
        time_column = ANTHRO_TIME_COLUMNS[metric_name]

```

```

        if unit_column not in updated.columns:
            updated[unit_column] = pd.NA
        else:
            updated[unit_column] = (
                updated[unit_column].astype("string").replace({"": pd.NA})
            )
        if time_column not in updated.columns:
            updated[time_column] = pd.NaT
        else:
            updated[time_column] = pd.to_datetime(updated[time_column],
        ↵ errors="coerce")

    metric_source_cols = {
        metric_name: f"__anthro_{metric_name}_source"
        for metric_name in ANTHRO_VALUE_COLUMNS
    }
    metric_hours_cols = {
        metric_name: f"__anthro_{metric_name}_hours_offset"
        for metric_name in ANTHRO_VALUE_COLUMNS
    }
    for metric_name in ANTHRO_VALUE_COLUMNS:
        source_col = metric_source_cols[metric_name]
        hours_col = metric_hours_cols[metric_name]
        if source_col not in updated.columns:
            updated[source_col] = "missing"
        else:
            updated[source_col] =
        ↵ updated[source_col].map(normalize_anthro_source)
        if hours_col not in updated.columns:
            updated[hours_col] = pd.NA
        else:
            updated[hours_col] = pd.to_numeric(updated[hours_col],
        ↵ errors="coerce")

    if "anthro_source" not in updated.columns:
        updated["anthro_source"] = "missing"
    updated["anthro_source"] =
    ↵ updated["anthro_source"].map(normalize_anthro_source)
    if "anthro_obstime" not in updated.columns:
        updated["anthro_obstime"] = pd.NaT
    else:
        updated["anthro_obstime"] = pd.to_datetime(updated["anthro_obstime"],
    ↵ errors="coerce")

```

```

if "anthro_hours_offset" not in updated.columns:
    updated["anthro_hours_offset"] = pd.NA
updated["anthro_hours_offset"] =
↪ pd.to_numeric(updated["anthro_hours_offset"], errors="coerce")
if "anthro_timing_basis" not in updated.columns:
    updated["anthro_timing_basis"] = "missing"
if "anthro_timing_uncertain" not in updated.columns:
    updated["anthro_timing_uncertain"] = pd.Series([pd.NA] *
↪ len(updated), dtype="boolean")
else:
    updated["anthro_timing_uncertain"] =
↪ updated["anthro_timing_uncertain"].astype("boolean")

charted = charted_df.copy()
charted["subject_id"] = _to_int64(charted["subject_id"])
charted["obs_time"] = pd.to_datetime(charted["obs_time"],
↪ errors="coerce")
charted["result_name"] =
↪ charted["result_name"].map(normalize_anthro_metric_name)
charted["result_value_num"] = pd.to_numeric(charted["result_value_num"],
↪ errors="coerce")
charted["source"] = charted["source"].map(normalize_anthro_source)
charted["result_unit_raw"] = charted["result_unit_raw"].astype("string")
charted["unit_norm"] =
↪ charted["result_unit_raw"].map(normalize_anthro_unit_text)
charted["canonical_value"] = np.nan
charted["canonical_unit"] = pd.Series(pd.NA, index=charted.index,
↪ dtype="string")
charted["conversion_applied"] = False
charted["drop_reason"] = pd.Series(pd.NA, index=charted.index,
↪ dtype="string")
charted = charted.loc[
    charted["subject_id"].notna()
    & charted["obs_time"].notna()
    & charted["result_name"].isin(OMR_RESULT_NAMES)
    & charted["result_value_num"].notna()
].copy()

source_rank = {source: rank for rank, source in
↪ enumerate(source_preference)}
source_rank.setdefault("missing", 999)

diagnostics: dict[str, Any] = {

```

```

    "charted_rows_input": int(lencharted_df)),
    "charted_rows_usable": 0,
    "nearest_anytime": bool(nearest_anytime),
    "source_preference": list(source_preference),
    "selected_counts": {column: 0 for column in OMR_OUTPUT_COLUMNS},
    "rows_with_any_selected_metric": 0,
    "source_metric_input_counts": {},
    "source_metric_usable_counts": {},
    "unit_conversion_counts": {},
    "unsupported_unit_drop_counts": {},
    "bmi_missing_before_pair_fill": 0,
    "bmi_pair_within_7d_n": 0,
    "bmi_pair_outside_7d_n": 0,
    "bmi_pair_missing_time_n": 0,
    "bmi_pair_gap_hours_quantiles": {}
}

if not charted.empty:
    diagnostics["source_metric_input_counts"] = {
        f"{source}|{metric}": int(count)
        for (source, metric), count in (
            charted.groupby(["source", "result_name"], dropna=False)
            .size()
            .items()
        )
    }

if charted.empty:
    return updated, diagnostics

for metric_name in OMR_RESULT_NAMES:
    metric_mask = charted["result_name"].eq(metric_name)
    if not bool(metric_mask.any()):
        continue
    values = pd.to_numeric(charted.loc[metric_mask, "result_value_num"],
    ↪ errors="coerce")
    unit_norm = charted.loc[metric_mask, "unit_norm"].fillna("")

    if metric_name == "weight":
        kg_mask = unit_norm.isin({"kg", "kgs", "kilogram", "kilograms"})
        lb_mask = unit_norm.isin({"lb", "lbs", "pound", "pounds"})
        oz_mask = unit_norm.isin({"oz", "ounce", "ounces"})
        g_mask = unit_norm.isin({"g", "gram", "grams"})

```

```

factor = pd.Series(np.nan, index=values.index, dtype="float64")
factor.loc[kg_mask] = 1.0
factor.loc[lb_mask] = 0.45359237
factor.loc[oz_mask] = 0.028349523125
factor.loc[g_mask] = 0.001
supported = factor.notna()
canonical = values * factor

elif metric_name == "height":
    cm_mask = unit_norm.isin({"cm", "centimeter", "centimeters"})
    m_mask = unit_norm.isin({"m", "meter", "meters"})
    in_mask = unit_norm.isin({"in", "inch", "inches"})
    ft_mask = unit_norm.isin({"ft", "foot", "feet"})
    factor = pd.Series(np.nan, index=values.index, dtype="float64")
    factor.loc[cm_mask] = 1.0
    factor.loc[m_mask] = 100.0
    factor.loc[in_mask] = 2.54
    factor.loc[ft_mask] = 30.48
    supported = factor.notna()
    canonical = values * factor
else:
    supported = unit_norm.isin(
        {
            "",
            "kg/m2",
            "kg/m^2",
            "kg/m²",
            "bmi",
            "kgperm2",
            "kgperm^2",
        }
    )
    canonical = values
    factor = pd.Series(1.0, index=values.index, dtype="float64")

positive_mask = values.gt(0)
usable_metric = supported & positive_mask & canonical.notna()
conversion_applied = usable_metric & factor.ne(1.0)

charted.loc[metric_mask, "canonical_value"] = canonical
charted.loc[metric_mask, "canonical_unit"] =
↪ ANTHRO_CANONICAL_UNITS[metric_name]
↪ charted.loc[metric_mask & conversion_applied, "conversion_applied"] =
↪ True

```

```

charted.loc[metric_mask & ~supported, "drop_reason"] = "unknown_unit"
charted.loc[metric_mask & supported & ~positive_mask, "drop_reason"]
↪ = "nonpositive_value"
charted.loc[metric_mask & ~usable_metric, "canonical_value"] = np.nan
charted.loc[metric_mask & ~usable_metric, "canonical_unit"] = pd.NA

usable = charted.loc[
    charted["canonical_value"].notna()
    & charted["canonical_unit"].notna()
    & charted["source"].ne("missing")
].copy()
diagnostics["charted_rows_usable"] = int(len(usable))
if not usable.empty:
    diagnostics["source_metric_usable_counts"] = {
        f"{source}|{metric}": int(count)
        for (source, metric), count in (
            usable.groupby(["source", "result_name"], dropna=False)
            .size()
            .items()
        )
    }
diagnostics["unit_conversion_counts"] = {
    f"{source}|{metric}": int(count)
    for (source, metric), count in (
        usable.loc[usable["conversion_applied"]]
        .groupby(["source", "result_name"], dropna=False)
        .size()
        .items()
    )
}
diagnostics["unsupported_unit_drop_counts"] = {
    f"{source}|{metric}|{reason}": int(count)
    for (source, metric, reason), count in (
        charted.loc[charted["drop_reason"].notna()]
        .groupby(["source", "result_name", "drop_reason"], dropna=False)
        .size()
        .items()
    )
}
if usable.empty:
    return updated, diagnostics

```

```

ed_norm = updated[["ed_stay_id", "subject_id", "ed_intime"]].copy()
ed_norm["subject_id"] = _to_int64(ed_norm["subject_id"])
ed_norm["ed_intime_dt"] = pd.to_datetime(ed_norm["ed_intime"],
                                         errors="coerce")
ed_norm = ed_norm.loc[ed_norm["subject_id"].notna() &
                      ed_norm["ed_intime_dt"].notna()].copy()
if ed_norm.empty:
    return updated, diagnostics

merged = ed_norm.merge(usable, on="subject_id", how="inner")
if merged.empty:
    diagnostics["subject_overlap_count"] = 0
    return updated, diagnostics
diagnostics["subject_overlap_count"] =
int(merged["subject_id"].nunique())
merged["hours_offset"] = (
    (merged["obs_time"] - merged["ed_intime_dt"]).dt.total_seconds() /
    3600.0
)
merged["abs_hours_offset"] = merged["hours_offset"].abs()
merged["source_priority"] =
merged["source"].map(source_rank).fillna(999).astype(int)
if not nearest_anytime:
    merged = merged.loc[merged["abs_hours_offset"] <= 24.0].copy()
if merged.empty:
    return updated, diagnostics

any_fill_mask = pd.Series(False, index=updated.index)
for metric_name, output_column in ANTHRO_VALUE_COLUMNS.items():
    subset = merged.loc[merged["result_name"] == metric_name].copy()
    if subset.empty:
        continue
    selected = (
        subset.sort_values(
            ["ed_stay_id", "abs_hours_offset", "obs_time",
             "source_priority", "canonical_value"],
            ascending=[True, True, True, True, False],
        )
        .groupby("ed_stay_id", as_index=False)
        .first()
    )
    selected = selected.rename(
        columns={

```

```

        "canonical_value": f"{output_column}__candidate",
        "canonical_unit": f"{output_column}__uom",
        "obs_time": f"{output_column}__obs_time",
        "hours_offset": f"{output_column}__hours_offset",
        "source": f"{output_column}__source",
    }
)
unit_column = ANTHRO_UNIT_COLUMNS[metric_name]
time_column = ANTHRO_TIME_COLUMNS[metric_name]
source_col = metric_source_cols[metric_name]
hours_col = metric_hours_cols[metric_name]
keep_columns = [
    "ed_stay_id",
    f"{output_column}__candidate",
    f"{output_column}__uom",
    f"{output_column}__obs_time",
    f"{output_column}__hours_offset",
    f"{output_column}__source",
]
updated = updated.merge(selected[keep_columns], on="ed_stay_id",
← how="left")
candidate_mask = updated[f"{output_column}__candidate"].notna()
updated.loc[candidate_mask, output_column] = updated.loc[
    candidate_mask, f"{output_column}__candidate"
]
updated.loc[candidate_mask, unit_column] = (
    updated.loc[candidate_mask, f"{output_column}__uom"]
    .astype("string")
    .replace({"": pd.NA})
)
updated.loc[candidate_mask, time_column] = pd.to_datetime(
    updated.loc[candidate_mask, f"{output_column}__obs_time"],
← errors="coerce"
)
updated.loc[candidate_mask, source_col] = (
    updated.loc[candidate_mask, f"{output_column}__source"]
    .map(normalize_anthro_source)
    .fillna("missing")
)
updated.loc[candidate_mask, hours_col] = pd.to_numeric(
    updated.loc[candidate_mask, f"{output_column}__hours_offset"],
← errors="coerce"
)

```

```

any_fill_mask = any_fill_mask | candidate_mask
diagnostics["selected_counts"][output_column] =
↪ int(candidate_mask.sum())

drop_columns = [
    f"{output_column}__candidate",
    f"{output_column}__uom",
    f"{output_column}__obs_time",
    f"{output_column}__hours_offset",
    f"{output_column}__source",
]
updated = updated.drop(columns=drop_columns)

diagnostics["rows_with_any_selected_metric"] = int(any_fill_mask.sum())
diagnostics["selected_source_counts"] = {
    str(key): int(value)
    for key, value in pd.Series(
        np.where(
            updated.get(metric_source_cols["bmi"], "missing").astype(str)
        )
        != "missing",
        updated.get(metric_source_cols["bmi"], "missing"),
        np.where(
            updated.get(metric_source_cols["weight"], "missing"),
            "missing").astype(str) != "missing",
            updated.get(metric_source_cols["weight"], "missing"),
            updated.get(metric_source_cols["height"], "missing"),
        ),
    ),
    dtype="string",
)
.fillna("missing")
.astype(str)
.value_counts(dropna=False)
.items()
}
return updated, diagnostics

```

```

def build_anthro_coverage_audit(ed_df: pd.DataFrame) -> dict[str, Any]:
    """Summarize anthropometric coverage and provenance rates."""
    total_rows = max(int(len(ed_df)), 1)
    field_counts = {

```

```

        column: int(pd.to_numeric(ed_df[column],
        ↵ errors="coerce").notna().sum())
            for column in OMR_OUTPUT_COLUMNS
                if column in ed_df.columns
        }
        field_rates = {column: float(count / total_rows) for column, count in
        ↵ field_counts.items()}
        unit_nonnull_counts = {
            unit_column: int(ed_df[unit_column].astype("string").notna().sum())
                for unit_column in ANTHRO_UNIT_COLUMNS.values()
                    if unit_column in ed_df.columns
        }
        time_nonnull_counts = {
            time_column: int(pd.to_datetime(ed_df[time_column],
        ↵ errors="coerce").notna().sum())
                for time_column in ANTHRO_TIME_COLUMNS.values()
                    if time_column in ed_df.columns
        }

        source_counts: dict[str, int] = {}
        source_rates: dict[str, float] = {}
        if "anthro_source" in ed_df.columns:
            source_counts = {
                str(key): int(value)
                    for key, value in ed_df["anthro_source"]
                        .map(normalize_anthro_source)
                            .fillna("missing")
                                .astype("string")
                                    .value_counts(dropna=False)
                                        .items()
            }
            source_rates = {key: float(value / total_rows) for key, value in
        ↵ source_counts.items()}

        timing_basis_counts: dict[str, int] = {}
        timing_basis_rates: dict[str, float] = {}
        if "anthro_timing_basis" in ed_df.columns:
            timing_basis_counts = {
                str(key): int(value)
                    for key, value in ed_df["anthro_timing_basis"]
                        .fillna("missing")
                            .astype(str)
                                .value_counts(dropna=False)

```

```

        .items()
    }
    timing_basis_rates = {
        key: float(value / total_rows) for key, value in
    ↪ timing_basis_counts.items()
    }

    return {
        "row_count": int(len(ed_df)),
        "field_nonnull_counts": field_counts,
        "field_nonnull_rates": field_rates,
        "unit_nonnull_counts": unit_nonnull_counts,
        "time_nonnull_counts": time_nonnull_counts,
        "source_counts": source_counts,
        "source_rates": source_rates,
        "timing_basis_counts": timing_basis_counts,
        "timing_basis_rates": timing_basis_rates,
    }
}

def build_first_other_pco2_audit(ed_df: pd.DataFrame) -> pd.DataFrame:
    """Build route-stratified audit summary for first_other_pco2 values.

    This helper is intentionally tolerant for notebook QA execution: if the
    required fields are not present, it returns a sentinel audit row instead
    ↪ of
    raising.
    """
    columns = [
        "source",
        "count_nonnull",
        "mean",
        "median",
        "q25",
        "q75",
        "p95",
        "max",
        "pct_ge_80",
        "pct_ge_100",
        "pct_ge_150",
        "pct_eq_160",
        "top_values",
        "status",
    ]

```

```

        "missing_columns",
    ]
required = {"first_other_pco2", "first_other_src"}
missing = sorted(required.difference(ed_df.columns))
if missing:
    return pd.DataFrame(
        [
            {
                "source": "UNAVAILABLE",
                "count_nonnull": 0,
                "mean": None,
                "median": None,
                "q25": None,
                "q75": None,
                "p95": None,
                "max": None,
                "pct_ge_80": 0.0,
                "pct_ge_100": 0.0,
                "pct_ge_150": 0.0,
                "pct_eq_160": 0.0,
                "top_values": {},
                "status": "missing_columns",
                "missing_columns": ",".join(missing),
            }
        ],
        columns=columns,
    )

frame = ed_df[["first_other_src", "first_other_pco2"]].copy()
frame["first_other_src"] = (
    frame["first_other_src"].astype(str).str.strip().str.upper().replace({
        "UNKNOWN": "NAN"
    })
)
frame["first_other_pco2"] = pd.to_numeric(frame["first_other_pco2"],
                                         errors="coerce")
frame = frame.loc[frame["first_other_pco2"].notna()].copy()

if frame.empty:
    return pd.DataFrame(
        [
            {
                "source": "UNAVAILABLE",

```

```

        "count_nonnull": 0,
        "mean": None,
        "median": None,
        "q25": None,
        "q75": None,
        "p95": None,
        "max": None,
        "pct_ge_80": 0.0,
        "pct_ge_100": 0.0,
        "pct_ge_150": 0.0,
        "pct_eq_160": 0.0,
        "top_values": {},
        "status": "no_nonnull_values",
        "missing_columns": ""
    }
],
columns=columns,
)
)

rows: list[dict[str, Any]] = []
for source_name, group in frame.groupby("first_other_src"):
    values = group["first_other_pco2"]
    rows.append(
    {
        "source": source_name,
        "count_nonnull": int(values.shape[0]),
        "mean": float(values.mean()),
        "median": float(values.median()),
        "q25": float(values.quantile(0.25)),
        "q75": float(values.quantile(0.75)),
        "p95": float(values.quantile(0.95)),
        "max": float(values.max()),
        "pct_ge_80": float((values >= 80).mean()),
        "pct_ge_100": float((values >= 100).mean()),
        "pct_ge_150": float((values >= 150).mean()),
        "pct_eq_160": float((values == 160).mean()),
        "top_values": [
            str(key): int(value)
            for key, value in values.value_counts().head(10).items()
        ],
        "status": "ok",
        "missing_columns": ""
    }
)

```

```

        )
    return pd.DataFrame(rows,
        columns=columns).sort_values("source").reset_index(drop=True)

def normalize_temperature_to_f(temp: pd.Series) -> pd.DataFrame:
    """Normalize mixed-unit temperature values to cleaned Fahrenheit output.

    Rules:
    - values in ``[20, 50]`` are treated as Celsius-like and converted to
    Farenheit,
    - all other numeric values are treated as Farenheit already,
    - cleaned output is valid only for Farenheit ``[50, 120]``.
    """
    numeric = pd.to_numeric(temp, errors="coerce")
    celsius_like = numeric.between(20.0, 50.0, inclusive="both")
    temp_f = numeric.where(~celsius_like, numeric * 9.0 / 5.0 + 32.0)
    out_of_range = temp_f.notna() & ~temp_f.between(50.0, 120.0,
        inclusive="both")
    temp_f_clean = temp_f.where(~out_of_range)
    return pd.DataFrame(
        {
            "temp_f_clean": temp_f_clean,
            "temp_was_celsius_like": celsius_like.astype("boolean"),
            "temp_out_of_range": out_of_range.astype("boolean"),
        }
    )

def clean_pain_score(pain: pd.Series) -> pd.DataFrame:
    """Clean pain score with explicit sentinel handling."""
    numeric = pd.to_numeric(pain, errors="coerce")
    sentinel = numeric.eq(13)
    out_of_range = numeric.notna() & (~sentinel) & ((numeric < 0) | (numeric
        > 10))
    clean = numeric.where(~(sentinel | out_of_range))
    return pd.DataFrame(
        {
            "pain_clean": clean,
            "pain_is_sentinel_13": sentinel.astype("boolean"),
            "pain_out_of_range": out_of_range.astype("boolean"),
        }
    )

```

```

def clean_bp(sbp: pd.Series, dbp: pd.Series) -> pd.DataFrame:
    """Range-clean SBP/DBP values with deterministic bounds."""
    sbp_num = pd.to_numeric(sbp, errors="coerce")
    dbp_num = pd.to_numeric(dbp, errors="coerce")
    sbp_out_of_range = sbp_num.notna() & ((sbp_num < 20) | (sbp_num > 300))
    dbp_out_of_range = dbp_num.notna() & ((dbp_num < 10) | (dbp_num > 200))
    return pd.DataFrame(
        {
            "sbp_clean": sbp_num.where(~sbp_out_of_range),
            "dbp_clean": dbp_num.where(~dbp_out_of_range),
            "sbp_out_of_range": sbp_out_of_range.astype("boolean"),
            "dbp_out_of_range": dbp_out_of_range.astype("boolean"),
        }
    )

def clean_o2sat(o2sat: pd.Series) -> pd.DataFrame:
    """Clean oxygen saturation and keep explicit >100 and zero flags."""
    numeric = pd.to_numeric(o2sat, errors="coerce")
    gt_100 = numeric > 100
    out_of_range = numeric.notna() & (numeric < 0)
    clean = numeric.where(~(gt_100 | out_of_range))
    return pd.DataFrame(
        {
            "o2sat_clean": clean,
            "o2sat_gt_100": gt_100.astype("boolean"),
            "o2sat_out_of_range": out_of_range.astype("boolean"),
            "o2sat_zero": numeric.eq(0).astype("boolean"),
        }
    )

def _series_distribution_stats(values: pd.Series) -> dict[str, Any]:
    """Return compact distribution metrics for audit CSV artifacts."""
    numeric = pd.to_numeric(values, errors="coerce")
    total_rows = int(len(numeric))
    non_missing = int(numeric.notna().sum())
    if non_missing == 0:
        return {
            "n_non_missing": 0,
            "missing_pct": 100.0 if total_rows else 0.0,
        }

```



```

outlier_n = int(outlier_flag_numeric.sum())
examples = (
    numeric.loc[outlier_flag_numeric]
    .round(4)
    .value_counts()
    .head(5)
    .index.astype(str)
    .tolist()
)
audit_rows.append(
{
    "domain": "vitals",
    "raw_column": raw_column,
    "model_column": model_column,
    "outlier_flag_column": outlier_flag_column,
    "lower_bound": lower_bound,
    "upper_bound": upper_bound,
    "nonnull_n": nonnull_n,
    "out_of_range_n": outlier_n,
    "out_of_range_pct": float(outlier_n / nonnull_n) if nonnull_n
    ↵ else 0.0,
    "example_outlier_values": "; ".join(examples),
}
)
for raw_column, (lower_bound, upper_bound) in resolved_ranges.items():
    if raw_column not in updated.columns:
        continue
    numeric = pd.to_numeric(updated[raw_column], errors="coerce")
    out_of_range = numeric.notna() & ((numeric < lower_bound) | (numeric
    ↵ > upper_bound))
    model_column = f"{raw_column}_model"
    outlier_flag_column = f"{raw_column}_outlier_flag"
    updated[model_column] = numeric.where(~out_of_range)
    updated[outlier_flag_column] = out_of_range.astype("boolean")
    _append_audit_row(
        raw_column=raw_column,
        model_column=model_column,
        outlier_flag_column=outlier_flag_column,
        lower_bound=float(lower_bound),
        upper_bound=float(upper_bound),
        numeric=numeric,
        outlier_flag=out_of_range,
    )

```

```

        )

for prefix in ("ed_triage", "ed_first"):
    temp_raw_col = f"{prefix}_temp"
    pain_raw_col = f"{prefix}_pain"
    sbp_raw_col = f"{prefix}_sbp"
    dbp_raw_col = f"{prefix}_dbp"
    o2sat_raw_col = f"{prefix}_o2sat"

    if temp_raw_col in updated.columns:
        temp_clean = normalize_temperature_to_f(updated[temp_raw_col])
        temp_clean_col = f"{prefix}_temp_f_clean"
        temp_celsius_flag_col = f"{prefix}_temp_was_celsius_like"
        temp_range_flag_col = f"{prefix}_temp_out_of_range"
        updated[temp_clean_col] = temp_clean["temp_f_clean"]
        updated[temp_celsius_flag_col] =
        ↵ temp_clean["temp_was_celsius_like"]
        updated[temp_range_flag_col] = temp_clean["temp_out_of_range"]
        updated[f"{prefix}_temp_f_model"] = updated[temp_clean_col]
        updated[f"{prefix}_temp_c_model"] =
        ↵ (updated[f"{prefix}_temp_f_model"] - 32.0) * (5.0 / 9.0)
        updated[f"{prefix}_temp_model"] =
        ↵ updated[f"{prefix}_temp_f_model"]
        updated[f"{prefix}_temp_outlier_flag"] =
        ↵ updated[temp_range_flag_col].astype("boolean")
        _append_audit_row(
            raw_column=temp_raw_col,
            model_column=f"{prefix}_temp_model",
            outlier_flag_column=f"{prefix}_temp_outlier_flag",
            lower_bound=50.0,
            upper_bound=120.0,
            numeric=pd.to_numeric(updated[temp_raw_col],
        ↵ errors="coerce"),
            outlier_flag=updated[temp_range_flag_col],
        )

    if pain_raw_col in updated.columns:
        pain_clean = clean_pain_score(updated[pain_raw_col])
        pain_clean_col = f"{prefix}_pain_clean"
        pain_sentinel_col = f"{prefix}_pain_is_sentinel_13"
        pain_range_col = f"{prefix}_pain_out_of_range"
        updated[pain_clean_col] = pain_clean["pain_clean"]
        updated[pain_sentinel_col] = pain_clean["pain_is_sentinel_13"]

```

```

        updated[pain_range_col] = pain_clean["pain_out_of_range"]
        updated[f"{prefix}_pain_model"] = updated[pain_clean_col]
        pain_outlier = updated[pain_sentinel_col].fillna(False) |
        ↵ updated[pain_range_col].fillna(False)
            updated[f"{prefix}_pain_outlier_flag"] =
        ↵ pain_outlier.astype("boolean")
            _append_audit_row(
                raw_column=pain_raw_col,
                model_column=f"{prefix}_pain_model",
                outlier_flag_column=f"{prefix}_pain_outlier_flag",
                lower_bound=0.0,
                upper_bound=10.0,
                numeric=pd.to_numeric(updated[pain_raw_col],
        ↵ errors="coerce"),
                outlier_flag=pain_outlier,
            )
    if sbp_raw_col in updated.columns and dbp_raw_col in updated.columns:
        bp_clean = clean_bp(updated[sbp_raw_col], updated[dbp_raw_col])
        sbp_clean_col = f"{prefix}_sbp_clean"
        dbp_clean_col = f"{prefix}_dbp_clean"
        sbp_range_col = f"{prefix}_sbp_out_of_range"
        dbp_range_col = f"{prefix}_dbp_out_of_range"
        updated[sbp_clean_col] = bp_clean["sbp_clean"]
        updated[dbp_clean_col] = bp_clean["dbp_clean"]
        updated[sbp_range_col] = bp_clean["sbp_out_of_range"]
        updated[dbp_range_col] = bp_clean["dbp_out_of_range"]
        updated[f"{prefix}_sbp_model"] = updated[sbp_clean_col]
        updated[f"{prefix}_dbp_model"] = updated[dbp_clean_col]
        updated[f"{prefix}_sbp_outlier_flag"] =
        ↵ updated[sbp_range_col].astype("boolean")
            updated[f"{prefix}_dbp_outlier_flag"] =
        ↵ updated[dbp_range_col].astype("boolean")
            _append_audit_row(
                raw_column=sbp_raw_col,
                model_column=f"{prefix}_sbp_model",
                outlier_flag_column=f"{prefix}_sbp_outlier_flag",
                lower_bound=20.0,
                upper_bound=300.0,
                numeric=pd.to_numeric(updated[sbp_raw_col], errors="coerce"),
                outlier_flag=updated[sbp_range_col],
            )
            _append_audit_row(

```

```

        raw_column=dbp_raw_col,
        model_column=f"{prefix}_dbp_model",
        outlier_flag_column=f"{prefix}_dbp_outlier_flag",
        lower_bound=10.0,
        upper_bound=200.0,
        numeric=pd.to_numeric(updated[dbp_raw_col], errors="coerce"),
        outlier_flag=updated[dbp_range_col],
    )

    if o2sat_raw_col in updated.columns:
        o2sat_clean = clean_o2sat(updated[o2sat_raw_col])
        o2sat_clean_col = f"{prefix}_o2sat_clean"
        o2sat_gt_100_col = f"{prefix}_o2sat_gt_100"
        o2sat_range_col = f"{prefix}_o2sat_out_of_range"
        o2sat_zero_col = f"{prefix}_o2sat_zero"
        updated[o2sat_clean_col] = o2sat_clean["o2sat_clean"]
        updated[o2sat_gt_100_col] = o2sat_clean["o2sat_gt_100"]
        updated[o2sat_range_col] = o2sat_clean["o2sat_out_of_range"]
        updated[o2sat_zero_col] = o2sat_clean["o2sat_zero"]
        updated[f"{prefix}_o2sat_model"] = updated[o2sat_clean_col]
        o2sat_outlier = updated[o2sat_gt_100_col].fillna(False) |
        ↵ updated[o2sat_range_col].fillna(False)
        updated[f"{prefix}_o2sat_outlier_flag"] =
        ↵ o2sat_outlier.astype("boolean")
        _append_audit_row(
            raw_column=o2sat_raw_col,
            model_column=f"{prefix}_o2sat_model",
            outlier_flag_column=f"{prefix}_o2sat_outlier_flag",
            lower_bound=0.0,
            upper_bound=100.0,
            numeric=pd.to_numeric(updated[o2sat_raw_col],
        ↵ errors="coerce"),
            outlier_flag=o2sat_outlier,
        )
    else:
        audit =
        ↵ pd.DataFrame(audit_rows).sort_values("raw_column").reset_index(drop=True)
    return updated, audit

def build_ed_vitals_audit_artifacts(
    ed_df: pd.DataFrame,
) -> tuple[pd.DataFrame, pd.DataFrame, pd.DataFrame]:

```

```

"""Build ED-vitals distribution, extreme-example, and raw-to-clean delta
↳ artifacts."""

specs = [
    {
        "vital_name": "triage_temp",
        "raw_col": "ed_triage_temp",
        "clean_col": "ed_triage_temp_f_clean",
        "flags": ["ed_triage_temp_was_celsius_like",
                  ↳ "ed_triage_temp_out_of_range"],
        "extra_cols": [],
    },
    {
        "vital_name": "first_temp",
        "raw_col": "ed_first_temp",
        "clean_col": "ed_first_temp_f_clean",
        "flags": ["ed_first_temp_was_celsius_like",
                  ↳ "ed_first_temp_out_of_range"],
        "extra_cols": ["ed_first_vitals_time"],
    },
    {
        "vital_name": "triage_pain",
        "raw_col": "ed_triage_pain",
        "clean_col": "ed_triage_pain_clean",
        "flags": ["ed_triage_pain_is_sentinel_13",
                  ↳ "ed_triage_pain_out_of_range"],
        "extra_cols": [],
    },
    {
        "vital_name": "first_pain",
        "raw_col": "ed_first_pain",
        "clean_col": "ed_first_pain_clean",
        "flags": ["ed_first_pain_is_sentinel_13",
                  ↳ "ed_first_pain_out_of_range"],
        "extra_cols": ["ed_first_vitals_time"],
    },
    {
        "vital_name": "triage_sbp",
        "raw_col": "ed_triage_sbp",
        "clean_col": "ed_triage_sbp_clean",
        "flags": ["ed_triage_sbp_out_of_range"],
        "extra_cols": [],
    },
]

```

```

        "vital_name": "triage_dbp",
        "raw_col": "ed_triage_dbp",
        "clean_col": "ed_triage_dbp_clean",
        "flags": ["ed_triage_dbp_out_of_range"],
        "extra_cols": [],
    },
    {
        "vital_name": "first_sbp",
        "raw_col": "ed_first_sbp",
        "clean_col": "ed_first_sbp_clean",
        "flags": ["ed_first_sbp_out_of_range"],
        "extra_cols": ["ed_first_vitals_time"],
    },
    {
        "vital_name": "first_dbp",
        "raw_col": "ed_first_dbp",
        "clean_col": "ed_first_dbp_clean",
        "flags": ["ed_first_dbp_out_of_range"],
        "extra_cols": ["ed_first_vitals_time"],
    },
    {
        "vital_name": "triage_o2sat",
        "raw_col": "ed_triage_o2sat",
        "clean_col": "ed_triage_o2sat_clean",
        "flags": ["ed_triage_o2sat_gt_100",
                  "ed_triage_o2sat_out_of_range", "ed_triage_o2sat_zero"],
        "extra_cols": [],
    },
    {
        "vital_name": "first_o2sat",
        "raw_col": "ed_first_o2sat",
        "clean_col": "ed_first_o2sat_clean",
        "flags": ["ed_first_o2sat_gt_100", "ed_first_o2sat_out_of_range",
                  "ed_first_o2sat_zero"],
        "extra_cols": ["ed_first_vitals_time"],
    },
],
summary_rows: list[dict[str, Any]] = []
delta_rows: list[dict[str, Any]] = []
extreme_frames: list[pd.DataFrame] = []

for spec in specs:

```

```

raw_col = spec["raw_col"]
clean_col = spec["clean_col"]
if raw_col not in ed_df.columns or clean_col not in ed_df.columns:
    continue

raw = pd.to_numeric(ed_df[raw_col], errors="coerce")
clean = pd.to_numeric(ed_df[clean_col], errors="coerce")
raw_stats = _series_distribution_stats(raw)
clean_stats = _series_distribution_stats(clean)
row: dict[str, Any] = {
    "vital_name": spec["vital_name"],
    "raw_column": raw_col,
    "clean_column": clean_col,
    "raw_n_non_missing": raw_stats["n_non_missing"],
    "raw_missing_pct": raw_stats["missing_pct"],
    "raw_min": raw_stats["min"],
    "raw_p1": raw_stats["p1"],
    "raw_median": raw_stats["median"],
    "raw_mean": raw_stats["mean"],
    "raw_p99": raw_stats["p99"],
    "raw_max": raw_stats["max"],
    "clean_n_non_missing": clean_stats["n_non_missing"],
    "clean_missing_pct": clean_stats["missing_pct"],
    "clean_min": clean_stats["min"],
    "clean_p1": clean_stats["p1"],
    "clean_median": clean_stats["median"],
    "clean_mean": clean_stats["mean"],
    "clean_p99": clean_stats["p99"],
    "clean_max": clean_stats["max"],
}
delta_rows.append(
{
    "vital_name": spec["vital_name"],
    "raw_column": raw_col,
    "clean_column": clean_col,
    "raw_n": int(raw.notna().sum()),
    "clean_n": int(clean.notna().sum()),
    "raw_removed_n": int(raw.notna().sum() -
        ↴ clean.notna().sum()),
    "raw_removed_pct": (
        float((raw.notna().sum() - clean.notna().sum()) /
        ↴ max(int(raw.notna().sum()), 1)))
}

```

```

        if int(raw.notna().sum()) > 0
        else 0.0
    ),
    "converted_celsius_like_n": int(
        ed_df.get(spec["flags"])[0], pd.Series(False,
    index=ed_df.index)
        .fillna(False)
        .astype(bool)
        .sum()
    )
    if "temp_was_celsius_like" in spec["flags"][0]
    else 0,
}
)

available_flags = [flag for flag in spec["flags"] if flag in
ed_df.columns]
for flag_col in available_flags:
    flag_series = ed_df[flag_col].fillna(False).astype(bool)
    row[f"{flag_col}_count"] = int(flag_series.sum())
    row[f"{flag_col}_pct"] = float(flag_series.mean()) if
len(flag_series) else 0.0
    if not flag_series.any():
        continue
    include_cols = ["hadm_id", "ed_stay_id", raw_col, clean_col,
flag_col]
    include_cols += [col for col in spec["extra_cols"] if col in
ed_df.columns]
    include_cols = [col for col in include_cols if col in
ed_df.columns]
    flagged = ed_df.loc[flag_series, include_cols].copy()
    flagged = flagged.rename(columns={raw_col: "raw_value",
clean_col: "clean_value", flag_col: "flag_value"})
    low = flagged.sort_values("raw_value", ascending=True,
na_position="last").head(5).copy()
    low["tail"] = "low"
    high = flagged.sort_values("raw_value", ascending=False,
na_position="last").head(5).copy()
    high["tail"] = "high"
    extremes = pd.concat([low, high],
ignore_index=True).drop_duplicates()
    extremes.insert(0, "vital_name", spec["vital_name"])
    extremes.insert(1, "raw_column", raw_col)

```

```

extremes.insert(2, "clean_column", clean_col)
extremes.insert(3, "flag_type", flag_col)
extreme_frames.append(extremes)

summary_rows.append(row)

distribution = pd.DataFrame(summary_rows)
delta = pd.DataFrame(delta_rows)
if extreme_frames:
    extremes = pd.concat(extreme_frames, ignore_index=True, sort=False)
else:
    extremes = pd.DataFrame(
        columns=[
            "vital_name",
            "raw_column",
            "clean_column",
            "flag_type",
            "hadm_id",
            "ed_stay_id",
            "raw_value",
            "clean_value",
            "flag_value",
            "tail",
        ]
    )
return distribution, extremes, delta

def evaluate_uom_expectations(ed_df: pd.DataFrame) -> dict[str, Any]:
    """Evaluate expected-null and value/uom consistency rules."""
    structural_nulls: dict[str, dict[str, Any]] = {}
    for field_name in EXPECTED_STRUCTURAL_NULL_FIELDS:
        if field_name not in ed_df.columns:
            continue
        structural_nulls[field_name] = {
            "present": True,
            "missing_n": int(ed_df[field_name].isna().sum()),
            "missing_pct": float(ed_df[field_name].isna().mean()),
            "all_null": bool(ed_df[field_name].isna().all()),
        }
    paco2_checks: dict[str, dict[str, Any]] = {}
    for value_column, uom_column in PACO2_VALUE_UOM_PAIRS:

```

```

        if value_column not in ed_df.columns or uom_column not in
        ↵ ed_df.columns:
            paco2_checks[uom_column] = {
                "present": False,
                "reason": "missing_value_or_uom_column",
            }
            continue

            value_present = ed_df[value_column].notna()
            uom_lower = ed_df[uom_column].astype(str).str.strip().str.lower()
            missing_uom_with_value = int((value_present &
        ↵ ed_df[uom_column].isna()).sum())
            non_mmhg_uom_with_value = int(
                (value_present & ed_df[uom_column].notna()) &
        ↵ uom_lower.ne("mmhg")).sum()
            )

            paco2_checks[uom_column] = {
                "present": True,
                "paired_value_column": value_column,
                "value_rows": int(value_present.sum()),
                "missing_uom_when_value_present": missing_uom_with_value,
                "non_mmhg_uom_when_value_present": non_mmhg_uom_with_value,
                "passes": bool(
                    missing_uom_with_value == 0 and non_mmhg_uom_with_value == 0
                ),
            }

        return {
            "expected_structural_null_fields":
            ↵ list(EXPECTED_STRUCTURAL_NULL_FIELDS),
            "structural_null_checks": structural_nulls,
            "paco2_uom_checks": paco2_checks,
        }

def classify_missingness_expectations(
    ed_df: pd.DataFrame,
    target_fields: list[str],
    *,
    expected_sparse_fields: set[str] | None = None,
) -> pd.DataFrame:

```

```

"""Classify field-level missingness into expected and unexpected
   categories."""
rows: list[dict[str, Any]] = []
total_rows = max(int(len(ed_df)), 1)
expected_structural = set(EXPECTED_STRUCTURAL_NULL_FIELDS)
expected_sparse = set(expected_sparse_fields or set())

for field_name in target_fields:
    if field_name not in ed_df.columns:
        rows.append(
            {
                "field": field_name,
                "missing_n": int(len(ed_df)),
                "missing_pct": 1.0,
                "expectation": "missing_column",
            }
        )
    continue

missing_n = int(ed_df[field_name].isna().sum())
missing_pct = float(missing_n / total_rows)

if field_name in expected_structural:
    expectation = "expected_structural_null"
elif field_name in expected_sparse and missing_pct >= 1.0:
    expectation = "expected_sparse"
elif missing_pct >= 1.0:
    expectation = "unexpected_full_null"
elif missing_pct > 0.0:
    expectation = "conditional_sparse"
else:
    expectation = "complete"

rows.append(
    {
        "field": field_name,
        "missing_n": missing_n,
        "missing_pct": missing_pct,
        "expectation": expectation,
    }
)

return pd.DataFrame(rows)

```

```

def render_latex_longtable(
    table_df: pd.DataFrame,
    *,
    caption: str,
    label: str,
    landscape: bool = False,
    index: bool = True,
) -> str:
    latex_text = table_df.to_latex(
        index=index,
        escape=False,
        longtable=True,
        caption=caption,
        label=label,
    )
    if landscape:
        latex_text = "\\begin{landscape}\n" + latex_text +
    ↵ "\\end{landscape}\n"
    return latex_text

from hypercap_cc_nlp.pipeline_audit import collect_run_manifest
from hypercap_cc_nlp.pipeline_contracts import validate_cohort_contract,
    write_contract_report

# ---- Backend selection (we use BigQuery)
BACKEND = os.getenv("MIMIC_BACKEND", "bigquery").strip().lower()
assert BACKEND == "bigquery", "This notebook is BigQuery-specific."

WORK_PROJECT = os.getenv("WORK_PROJECT", "").strip() # your billing project
PHYS = os.getenv("BQ_PHYSIONET_PROJECT", "physionet-data").strip() # hosting
    ↵ project (read-only)

# Dataset preferences: resolved to accessible datasets in the next setup
    ↵ cell.
HOSP = os.getenv("BQ_DATASET_HOSP", "mimiciv_3_1_hosp").strip()
ICU = os.getenv("BQ_DATASET_ICU", "mimiciv_3_1_icu").strip()
ED = os.getenv("BQ_DATASET_ED", "").strip()

# BigQuery client
client = bigquery.Client(project=WORK_PROJECT)

```

```

print("Project:", WORK_PROJECT)
print("PhysioNet host:", PHYS)
print("HOSP (pref):", HOSP, "ICU (pref):", ICU, "ED (pref):", ED)
print("WORK_DIR:", WORK_DIR)

Project: mimic-hypercapnia
PhysioNet host: physionet-data
HOSP (pref): mimiciv_3_1_hosp ICU (pref): mimiciv_3_1_icu ED (pref): mimiciv_ed
WORK_DIR: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Proj
CC-NLP

# Purpose: Define reusable BigQuery helpers and resolve dataset names across
# known naming variants.

from datetime import datetime, timezone

# Fail fast on long-running queries instead of hanging indefinitely in
# nbconvert.
BQ_QUERY_TIMEOUT_SECS = int(os.getenv("BQ_QUERY_TIMEOUT_SECS", "1800"))

# --- Helper: run SQL with optional named parameters
def run_sql_bq(sql: str, params: dict | None = None) -> pd.DataFrame:
    job_config = bigquery.QueryJobConfig()
    job_config.labels = {"pipeline": "hypercapcc", "notebook": "cohort"}
    if params:
        bq_params = []
        for k, v in params.items():
            if isinstance(v, (list, tuple, np.ndarray, pd.Series)):
                # BigQuery ARRAY<INT64> if all ints; else ARRAY<STRING>
                v_list = list(v)
                if all(isinstance(x, np.integer) for x in v_list):
                    bq_params.append(bigquery.ArrayQueryParameter(k, "INT64",
# list(map(int, v_list))))
                else:
                    bq_params.append(bigquery.ArrayQueryParameter(k,
# "STRING", list(map(str, v_list))))
            else:
                # scalar
                if isinstance(v, np.integer):
                    bq_params.append(bigquery.ScalarQueryParameter(k,
# "INT64", int(v)))
                elif isinstance(v, float):

```

```

                bq_params.append(bigquery.ScalarQueryParameter(k,
        ↵ "FLOAT64", float(v)))
            else:
                bq_params.append(bigquery.ScalarQueryParameter(k,
        ↵ "STRING", str(v)))
        job_config.query_parameters = bq_params

job = client.query(sql, job_config=job_config)
try:
    result = job.result(timeout=BQ_QUERY_TIMEOUT_SECS)
except Exception as exc:
    try:
        job.cancel()
    except Exception:
        pass
    raise RuntimeError(
        f"BigQuery query failed or timed out after
        ↵ {BQ_QUERY_TIMEOUT_SECS}s (job_id={job.job_id})."
    ) from exc

try:
    return result.to_dataframe(create_bqstorage_client=True)
except TypeError:
    return result.to_dataframe()

# --- Helper: test if a fully-qualified table exists and is accessible
def table_exists(fqtn: str) -> bool:
    try:
        _ = run_sql_bq(f"SELECT 1 FROM `fqtn` LIMIT 1")
        return True
    except Exception:
        return False

def resolve_dataset(preferred: str, candidates: list[str], probe_table: str,
    ↵ label: str) -> tuple[str, str]:
    # Keep preferred first, then try known aliases.
    ordered = []
    if preferred:
        ordered.append(preferred)
    for cand in candidates:
        if cand not in ordered:
            ordered.append(cand)

```

```

for dataset in ordered:
    fqtn = f"{PHYS}.{dataset}.{probe_table}"
    if table_exists(fqtn):
        return dataset, fqtn

raise RuntimeError(
    f"No accessible {label} dataset found for probe table
    ↵ '{probe_table}'. Tried: {ordered}"
)

# Resolve HOSP/ICU/ED names so notebook runs across v3.1 naming variants.
HOSP, HOSP_PROBE = resolve_dataset(
    preferred=HOSP,
    candidates=["mimiciv_3_1_hosp", "mimiciv_v3_1_hosp", "mimiciv_hosp"],
    probe_table="admissions",
    label="HOSP",
)
ICU, ICU_PROBE = resolve_dataset(
    preferred=ICU,
    candidates=["mimiciv_3_1_icu", "mimiciv_v3_1_icu", "mimiciv_icu"],
    probe_table="icustays",
    label="ICU",
)
if not ED:
    ED = "mimiciv_ed"
ED, ED_PROBE = resolve_dataset(
    preferred=ED,
    candidates=["mimiciv_ed", "mimiciv_3_1_ed", "mimiciv_v3_1_ed"],
    probe_table="edstays",
    label="ED",
)

RUN_METADATA = {
    "run_utc": datetime.now(timezone.utc).isoformat(),
    "work_project": WORK_PROJECT,
    "physionet_project": PHYS,
    "datasets": {"hosp": HOSP, "icu": ICU, "ed": ED},
    "dataset_probes": {"hosp": HOSP_PROBE, "icu": ICU_PROBE, "ed": ED_PROBE},
    "notebook": "MIMICIV_hypercap_EXT_cohort.ipynb",
    "blood_gas_manifest_path": BLOOD_GAS_MANIFEST["path"],
    "blood_gas_manifest_version": BLOOD_GAS_MANIFEST["version"],
}

```

```

print("Resolved datasets:", RUN_METADATA["datasets"])
print("Dataset probes:", RUN_METADATA["dataset_probes"])

Resolved datasets: {'hosp': 'mimiciv_3_1_hosp', 'icu': 'mimiciv_3_1_icu', 'ed': 'mimiciv_ed'}
Dataset probes: {'hosp': 'physionet-data.mimiciv_3_1_hosp.admissions', 'icu': 'physionet-
data.mimiciv_3_1_icu.icustays', 'ed': 'physionet-data.mimiciv_ed.edstays'}
```

Purpose: Create reusable data-quality and merge guardrail helpers to
 ↳ prevent silent join errors.

--- Helper utilities for reproducibility and safe joins

```

def require_cols(df: pd.DataFrame, cols: list[str], name: str) -> None:
    missing = [c for c in cols if c not in df.columns]
    if missing:
        raise KeyError(f"{name} missing columns: {missing}")

def assert_unique(df: pd.DataFrame, key: str, name: str) -> None:
    if df[key].duplicated().any():
        n = int(df[key].duplicated().sum())
        raise ValueError(f"{name} has {n} duplicate {key} values")

def safe_merge(left: pd.DataFrame, right: pd.DataFrame, on: list[str] | str,
              how: str, name: str) -> pd.DataFrame:
    # guard against accidental duplicate columns
    overlap = set(left.columns) & set(right.columns)
    if isinstance(on, str):
        on_cols = {on}
    else:
        on_cols = set(on)
    overlap = overlap - on_cols
    if overlap:
        raise ValueError(f"{name} merge would duplicate columns:
                         ↳ {sorted(overlap)}")
    return left.merge(right, on=on, how=how)

def check_ranges(df: pd.DataFrame, ranges: dict[str, tuple[float, float]]) -> pd.DataFrame:
    rows = []
    for col, (lo, hi) in ranges.items():
        if col not in df.columns:
            continue
        bad = df[col].notna() & ((df[col] < lo) | (df[col] > hi))

```

```

        rows.append({"col": col, "n_bad": int(bad.sum())})
    return pd.DataFrame(rows)

```

2.3 Data Generation

2.3.1 1) ICD cohort flags (hypercapnic respiratory failure)

Rationale: Capture diagnosis-based hypercapnia from ED and hospital discharge codes to define a broad, clinically recognized cohort.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
#           ↵ consistent across notebook stages.

# Target ICD codes (dotless, uppercase)
ICD10_CODES = ['J9602','J9612','J9622','J9692','E662']
ICD9_CODES   = ['27803']

SQL["cohort_icd_sql"] = f"""
-- ICD-based cohort flags per admission
WITH target_codes AS (
    SELECT 'J9602' AS code, 10 AS ver UNION ALL
    SELECT 'J9612', 10 UNION ALL
    SELECT 'J9622', 10 UNION ALL
    SELECT 'J9692', 10 UNION ALL
    SELECT 'E662', 10 UNION ALL
    SELECT '27803', 9
),
-- Hospital ICDs restricted to target codes
hosp_dx AS (
    SELECT
        d.subject_id,
        d.hadm_id,
        UPPER(REPLACE(d.icd_code, '.', '')) AS code_norm,
        d.icd_version
    FROM `{{PHYS}}.{{HOSP}}.diagnoses_icd` d
    JOIN target_codes t
        ON t.ver = d.icd_version AND t.code = UPPER(REPLACE(d.icd_code, '.', ''))
    WHERE d.hadm_id IS NOT NULL
),
-- Hospital flags per admission
hosp_flags AS (

```

```

SELECT
    subject_id, hadm_id,
    MAX(IF(icd_version=10 AND code_norm='J9602',1,0)) AS ICD10_J9602,
    MAX(IF(icd_version=10 AND code_norm='J9612',1,0)) AS ICD10_J9612,
    MAX(IF(icd_version=10 AND code_norm='J9622',1,0)) AS ICD10_J9622,
    MAX(IF(icd_version=10 AND code_norm='J9692',1,0)) AS ICD10_J9692,
    MAX(IF(icd_version=10 AND code_norm='E662', 1,0)) AS ICD10_E662,
    MAX(IF(icd_version=9  AND code_norm='27803',1,0)) AS ICD9_27803
FROM hosp_dx
GROUP BY subject_id, hadm_id
),

-- ED ICDs restricted to target codes (map to hadm via edstays)
ed_dx AS (
SELECT
    s.subject_id,
    s.hadm_id,
    s.stay_id,
    s.intime AS ed_intime,
    UPPER(REPLACE(d.icd_code, '.', '')) AS code_norm,
    d.icd_version
FROM `'{PHYS}`.{ED}.diagnosis` d
JOIN `'{PHYS}`.{ED}.edstays` s
    ON s.subject_id = d.subject_id AND s.stay_id = d.stay_id
JOIN target_codes t
    ON t.ver = d.icd_version AND t.code = UPPER(REPLACE(d.icd_code, '.', ''))
WHERE s.hadm_id IS NOT NULL
),

-- ED flags per ED stay (so we can both: OR flags across stays and also pick
-- earliest stay_id)
ed_flags_by_stay AS (
SELECT
    subject_id, hadm_id, stay_id, MIN(ed_intime) AS ed_intime,
    MAX(IF(icd_version=10 AND code_norm='J9602',1,0)) AS ICD10_J9602,
    MAX(IF(icd_version=10 AND code_norm='J9612',1,0)) AS ICD10_J9612,
    MAX(IF(icd_version=10 AND code_norm='J9622',1,0)) AS ICD10_J9622,
    MAX(IF(icd_version=10 AND code_norm='J9692',1,0)) AS ICD10_J9692,
    MAX(IF(icd_version=10 AND code_norm='E662', 1,0)) AS ICD10_E662,
    MAX(IF(icd_version=9  AND code_norm='27803',1,0)) AS ICD9_27803
FROM ed_dx
GROUP BY subject_id, hadm_id, stay_id
),

```

```

-- OR the ED flags across all ED stays mapped to the same hadm
ed_flags_or AS (
    SELECT
        subject_id, hadm_id,
        MAX(ICD10_J9602) AS ICD10_J9602,
        MAX(ICD10_J9612) AS ICD10_J9612,
        MAX(ICD10_J9622) AS ICD10_J9622,
        MAX(ICD10_J9692) AS ICD10_J9692,
        MAX(ICD10_E662 ) AS ICD10_E662,
        MAX(ICD9_27803) AS ICD9_27803
    FROM ed_flags_by_stay
    GROUP BY subject_id, hadm_id
),

-- Earliest ED stay_id per hadm (NO UNNEST of aggregates; use [OFFSET(0)])
ed_earliest AS (
    SELECT
        subject_id,
        hadm_id,
        (ARRAY_AGG(STRUCT(stay_id, ed_intime) ORDER BY ed_intime LIMIT
        ← 1)) [OFFSET(0)].stay_id AS stay_id
    FROM ed_flags_by_stay
    GROUP BY subject_id, hadm_id
),

-- Bring flags and earliest stay_id together
ed_by_hadm AS (
    SELECT
        f.subject_id,
        f.hadm_id,
        e.stay_id,
        f.ICD10_J9602,
        f.ICD10_J9612,
        f.ICD10_J9622,
        f.ICD10_J9692,
        f.ICD10_E662,
        f.ICD9_27803
    FROM ed_flags_or f
    LEFT JOIN ed_earliest e
        USING (subject_id, hadm_id)
),

```

```

-- Combine ED and hospital flags at the admission level
combined AS (
    SELECT
        COALESCE(h.subject_id, e.subject_id) AS subject_id,
        COALESCE(h.hadm_id, e.hadm_id) AS hadm_id,
        GREATEST(IFNULL(h.ICD10_J9602,0), IFNULL(e.ICD10_J9602,0)) AS
        ↵ ICD10_J9602,
        GREATEST(IFNULL(h.ICD10_J9612,0), IFNULL(e.ICD10_J9612,0)) AS
        ↵ ICD10_J9612,
        GREATEST(IFNULL(h.ICD10_J9622,0), IFNULL(e.ICD10_J9622,0)) AS
        ↵ ICD10_J9622,
        GREATEST(IFNULL(h.ICD10_J9692,0), IFNULL(e.ICD10_J9692,0)) AS
        ↵ ICD10_J9692,
        GREATEST(IFNULL(h.ICD10_E662 ,0), IFNULL(e.ICD10_E662 ,0)) AS ICD10_E662,
        GREATEST(IFNULL(h.ICD9_27803,0), IFNULL(e.ICD9_27803,0)) AS ICD9_27803,
        ↵ IF((IFNULL(h.ICD10_J9602,0)+IFNULL(h.ICD10_J9612,0)+IFNULL(h.ICD10_J9622,0)+IFNULL(h.ICD
        ↵ > 0, 1, 0) AS any_hypercap_icd_hosp,
        ↵ IF((IFNULL(e.ICD10_J9602,0)+IFNULL(e.ICD10_J9612,0)+IFNULL(e.ICD10_J9622,0)+IFNULL(e.ICD
        ↵ > 0, 1, 0) AS any_hypercap_icd_ed
    FROM hosp_flags h
    FULL OUTER JOIN ed_by_hadm e
        ON h.hadm_id = e.hadm_id
)
SELECT
    subject_id, hadm_id,
    ICD10_J9602, ICD10_J9612, ICD10_J9622, ICD10_J9692, ICD10_E662, ICD9_27803,
    IF((ICD10_J9602+ICD10_J9612+ICD10_J9622+ICD10_J9692+ICD10_E662+ICD9_27803)
    ↵ > 0, 1, 0) AS any_hypercap_icd,
    any_hypercap_icd_hosp,
    any_hypercap_icd_ed,
    CASE
        WHEN any_hypercap_icd_hosp=1 AND any_hypercap_icd_ed=1 THEN 'ED+HOSP'
        WHEN any_hypercap_icd_ed=1 THEN 'ED'
        WHEN any_hypercap_icd_hosp=1 THEN 'HOSP'
        ELSE 'NONE'
    END AS icd_source
FROM combined
"""

cohort_icd = run_sql_bq(sql("cohort_icd_sql"))

```

```
print("ICD cohort admissions:", len(cohort_icd))
cohort_icd.head(3)
```

ICD cohort admissions: 4237

	subject_id	hadm_id	ICD10_J9602	ICD10_J9612	ICD10_J9622	ICD10_J9692	ICD10_E662	I
0	10485425	21207827	0	0	1	0	1	0
1	16826447	21830147	0	1	0	0	1	0
2	11482582	22869017	0	0	0	0	0	1

2.3.2 2) Blood gas inclusion thresholds — qualifying hypercapnic gases in 0–24h (LAB + POC)

Rationale: Identify physiologic hypercapnia using inclusive ABG/VBG/unknown thresholds tied to first ED arrival timing, retaining both per-type qualifying flags and earliest qualifying event metadata.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪ consistent across notebook stages.

SQL["co2_thresholds_sql"] = f"""
WITH ed_anchor AS (
    SELECT
        hadm_id,
        MIN(intime) AS ed_intime_first
    FROM `{{PHYS}}.{{ED}}.edstays`
    WHERE hadm_id IS NOT NULL
    GROUP BY hadm_id
),
hosp_raw AS (
    SELECT
        le.hadm_id,
        le.charttime,
        le.specimen_id,
        le.itemid,
        COALESCE(
            CAST(le.valuenum AS FLOAT64),
            SAFE_CAST(REGEXP_EXTRACT(LOWER(le.value), r'(-?\d+(?:\.\d+)?)') AS
↪ FLOAT64)
        ) AS val,
        LOWER(REPLACE(COALESCE(le.valueuom, ''), ' ', '')) AS uom_nospace
```

```

FROM `'{PHYS}`.{HOSP}.labevents` le
JOIN `'{PHYS}`.{HOSP}.d_labitems` di
  ON di.itemid = le.itemid
WHERE le.itemid IN UNNEST({LAB_PCO2_SQL_ARRAY})
  AND (le.valuenum IS NOT NULL OR le.value IS NOT NULL)
  AND LOWER(COALESCE(di.fluid, '')) IN ({LAB_BLOOD_FLUID_SQL_LIST})
),
hosp_specimen AS (
  SELECT
    le.specimen_id,
    LOWER(COALESCE(le.value, '')) AS specimen_type_text
  FROM `'{PHYS}`.{HOSP}.labevents` le
  WHERE le.itemid IN UNNEST({LAB_SPECIMEN_SQL_ARRAY})
    AND le.specimen_id IS NOT NULL
    AND le.value IS NOT NULL
  QUALIFY ROW_NUMBER() OVER (
    PARTITION BY le.specimen_id
    ORDER BY le.charttime DESC
  ) = 1
),
hosp_pco2_clean AS (
  SELECT
    r.hadm_id,
    r.charttime,
    r.specimen_id,
    r.itemid,
    CAST(NULL AS INT64) AS stay_id,
    CASE
      WHEN REGEXP_CONTAINS(COALESCE(s.specimen_type_text, ''), '
        ↵ r'arter|\\bart\\b') THEN 'arterial'
      WHEN REGEXP_CONTAINS(COALESCE(s.specimen_type_text, ''), '
        ↵ r'ven|central|mixed') THEN 'venous'
      ELSE 'unknown'
    END AS site,
    CASE WHEN r.uom_nospace = 'kpa' THEN r.val * 7.50062 ELSE r.val END AS
    ↵ pco2_mmhg,
    'LAB' AS source_branch
  FROM hosp_raw r
  LEFT JOIN hosp_specimen s USING (specimen_id)
  WHERE r.val IS NOT NULL
    AND (CASE WHEN r.uom_nospace = 'kpa' THEN r.val * 7.50062 ELSE r.val END)
    ↵ BETWEEN 5 AND 200
),

```

```

icu_pco2_itemids AS (
    SELECT itemid, LOWER(label) AS lbl
    FROM `{{PHYS}}.{{ICU}}.d_items`
    WHERE itemid IN UNNEST({{ICU_PCO2_SQL_ARRAY}})
),
icu_pco2_raw AS (
    SELECT
        ie.hadm_id,
        ce.stay_id,
        ce.itemid,
        ce.charttime,
        LOWER(REPLACE(COALESCE(ce.valueuom, ''), ' ', '')) AS uom_nospace,
        CASE
            WHEN ce.itemid IN UNNEST({{ICU_PCO2_ABG_SQL_ARRAY}}) THEN 'arterial'
            WHEN ce.itemid IN UNNEST({{ICU_PCO2_VBG_SQL_ARRAY}}) THEN 'venous'
            ELSE 'unknown'
        END AS site_from_itemid,
        COALESCE(
            CAST(ce.valuenum AS FLOAT64),
            SAFE_CAST(REGEXP_EXTRACT(LOWER(ce.value), r'(-?\d+(?:\.\d+)?)') AS
            ← FLOAT64)
        ) AS val
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    JOIN icu_pco2_itemids di
        ON di.itemid = ce.itemid
    JOIN `{{PHYS}}.{{ICU}}.icustays` ie
        ON ie.stay_id = ce.stay_id
    WHERE ce.valuenum IS NOT NULL OR ce.value IS NOT NULL
),
icu_specimen AS (
    SELECT
        ce.stay_id,
        ce.charttime,
        LOWER(COALESCE(ce.value, '')) AS specimen_type_text
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    WHERE ce.itemid IN UNNEST({{ICU_SPECIMEN_SQL_ARRAY}})
        AND ce.value IS NOT NULL
),
icu_pco2_ranked AS (
    SELECT
        r.*,
        s.specimen_type_text,
        ROW_NUMBER() OVER (

```

```

        PARTITION BY r.hadm_id, r.stay_id, r.charttime, r.itemid
        ORDER BY ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, SECOND))
    ) AS specimen_rank
FROM icu_pco2_raw r
LEFT JOIN icu_specimen s
    ON s.stay_id = r.stay_id
    AND ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, MINUTE)) <= 120
),
icu_pco2_clean AS (
    SELECT
        hadm_id,
        charttime,
        CAST(NULL AS INT64) AS specimen_id,
        itemid,
        stay_id,
        CASE
            WHEN REGEXP_CONTAINS(COALESCE(specimen_type_text, ''), 
                r'arter|\\bart\\b') THEN 'arterial'
            WHEN REGEXP_CONTAINS(COALESCE(specimen_type_text, ''), 
                r'ven|central|mixed') THEN 'venous'
            WHEN site_from_itemid IN ('arterial', 'venous') THEN site_from_itemid
            ELSE 'unknown'
        END AS site,
        CASE WHEN uom_nospace = 'kpa' THEN val * 7.50062 ELSE val END AS
        pco2_mmhg,
        'POC' AS source_branch
    FROM icu_pco2_ranked
    WHERE val IS NOT NULL
        AND (specimen_rank = 1 OR specimen_type_text IS NULL)
        AND (CASE WHEN uom_nospace = 'kpa' THEN val * 7.50062 ELSE val END)
        BETWEEN 5 AND 200
),
all_pco2 AS (
    SELECT * FROM hosp_pco2_clean
    UNION ALL
    SELECT * FROM icu_pco2_clean
),
windowed_pco2 AS (
    SELECT
        p.hadm_id,
        p.charttime,
        p.specimen_id,
        p.itemid,

```

```

p.stay_id,
p.site,
p.pco2_mmhg,
p.source_branch,
a.ed_intime_first,
TIMESTAMP_DIFF(p.charttime, a.ed_intime_first, SECOND) / 3600.0 AS
↳ dt_hours
FROM all_pco2 p
JOIN ed_anchor a
  ON a.hadm_id = p.hadm_id
WHERE a.ed_intime_first IS NOT NULL
  AND p.charttime BETWEEN a.ed_intime_first AND
↳ TIMESTAMP_ADD(a.ed_intime_first, INTERVAL 24 HOUR)
),
qualifying_candidates AS (
SELECT
  hadm_id,
  charttime,
  pco2_mmhg,
  CASE
    WHEN site = 'arterial' THEN 'arterial'
    WHEN site = 'venous' THEN 'venous'
    ELSE 'unknown'
  END AS site,
  source_branch,
  itemid,
  dt_hours,
  CASE WHEN site = 'arterial' THEN 45.0 ELSE 50.0 END AS threshold_mmhg,
  CONCAT(
    source_branch,
    ':',
    CAST(COALESCE(specimen_id, -1) AS STRING),
    ':',
    CAST(COALESCE(stay_id, -1) AS STRING),
    ':',
    CAST(COALESCE(itemid, -1) AS STRING),
    ':',
    FORMAT_TIMESTAMP('%F %T.%6E', charttime)
  ) AS stable_key
FROM windowed_pco2
WHERE (
  (site = 'arterial' AND pco2_mmhg >= 45.0)
  OR (site IN ('venous', 'unknown') AND pco2_mmhg >= 50.0)
)

```

```

)
),
windowed_extrema AS (
    SELECT
        hadm_id,
        MAX(IF(dt_hours BETWEEN 0 AND 6, pco2_mmhg, NULL)) AS max_pco2_0_6h,
        MAX(IF(dt_hours BETWEEN 0 AND 24, pco2_mmhg, NULL)) AS max_pco2_0_24h
    FROM windowed_pco2
    GROUP BY hadm_id
),
ranked_window_6h AS (
    SELECT
        hadm_id,
        source_branch,
        ROW_NUMBER() OVER (
            PARTITION BY hadm_id
            ORDER BY
                pco2_mmhg DESC,
                charttime,
                CASE source_branch WHEN 'LAB' THEN 0 WHEN 'POC' THEN 1 ELSE 2 END,
                CAST(COALESCE(itemid, -1) AS INT64)
        ) AS rn
    FROM windowed_pco2
    WHERE dt_hours BETWEEN 0 AND 6
),
max_source_6h AS (
    SELECT
        hadm_id,
        source_branch AS max_pco2_0_6h_source_branch
    FROM ranked_window_6h
    WHERE rn = 1
),
ranked_window_24h AS (
    SELECT
        hadm_id,
        source_branch,
        ROW_NUMBER() OVER (
            PARTITION BY hadm_id
            ORDER BY
                pco2_mmhg DESC,
                charttime,
                CASE source_branch WHEN 'LAB' THEN 0 WHEN 'POC' THEN 1 ELSE 2 END,
                CAST(COALESCE(itemid, -1) AS INT64)
)

```

```

) AS rn
FROM windowed_pco2
WHERE dt_hours BETWEEN 0 AND 24
),
max_source_24h AS (
SELECT
    hadm_id,
    source_branch AS max_pco2_0_24h_source_branch
FROM ranked_window_24h
WHERE rn = 1
),
qualifying_flags AS (
SELECT
    hadm_id,
    MAX(IF(site = 'arterial', 1, 0)) AS abg_hypercap_threshold,
    MAX(IF(site = 'venous', 1, 0)) AS vbg_hypercap_threshold,
    MAX(IF(site = 'unknown', 1, 0)) AS unknown_hypercap_threshold
FROM qualifying_candidates
GROUP BY hadm_id
),
ranked_qualifying_overall AS (
SELECT
    q.*,
    ROW_NUMBER() OVER (
        PARTITION BY q.hadm_id
        ORDER BY
            q.charttime,
            CASE q.site WHEN 'arterial' THEN 0 WHEN 'venous' THEN 1 ELSE 2 END,
            CASE q.source_branch WHEN 'LAB' THEN 0 WHEN 'POC' THEN 1 ELSE 2 END,
            q.pco2_mmhg DESC,
            q.stable_key
    ) AS qualifying_rank
FROM qualifying_candidates q
),
qualifying_first AS (
SELECT
    hadm_id,
    charttime AS qualifying_pco2_time,
    pco2_mmhg AS qualifying_pco2_mmhg,
    site AS qualifying_site,
    source_branch AS qualifying_source_branch,
    threshold_mmhg AS qualifying_threshold_mmhg,
    dt_hours AS dt_qualifying_hypercapnia_hours,

```

```

    itemid AS qualifying_pco2_itemid
  FROM ranked_qualifying_overall
 WHERE qualifying_rank = 1
),
ranked_abg AS (
  SELECT
    q.*,
    ROW_NUMBER() OVER (
      PARTITION BY q.hadm_id
      ORDER BY
        q.charttime,
        CASE q.source_branch WHEN 'LAB' THEN 0 WHEN 'POC' THEN 1 ELSE 2 END,
        q.pco2_mmhg DESC,
        q.stable_key
    ) AS rn
  FROM qualifying_candidates q
 WHERE q.site = 'arterial'
),
first_abg AS (
  SELECT
    hadm_id,
    charttime AS first_abg_hypercap_time_0_24h,
    pco2_mmhg AS first_abg_hypercap_pco2_mmhg,
    source_branch AS first_abg_hypercap_source_branch,
    itemid AS first_abg_hypercap_itemid
  FROM ranked_abg
 WHERE rn = 1
),
ranked_vbg AS (
  SELECT
    q.*,
    ROW_NUMBER() OVER (
      PARTITION BY q.hadm_id
      ORDER BY
        q.charttime,
        CASE q.source_branch WHEN 'LAB' THEN 0 WHEN 'POC' THEN 1 ELSE 2 END,
        q.pco2_mmhg DESC,
        q.stable_key
    ) AS rn
  FROM qualifying_candidates q
 WHERE q.site = 'venous'
),
first_vbg AS (

```

```

SELECT
    hadm_id,
    charttime AS first_vbg_hypercap_time_0_24h,
    pco2_mmhg AS first_vbg_hypercap_pco2_mmhg,
    source_branch AS first_vbg_hypercap_source_branch,
    itemid AS first_vbg_hypercap_itemid
FROM ranked_vbg
WHERE rn = 1
),
ranked_other AS (
SELECT
    q.*,
    ROW_NUMBER() OVER (
        PARTITION BY q.hadm_id
        ORDER BY
            q.charttime,
            CASE q.source_branch WHEN 'LAB' THEN 0 WHEN 'POC' THEN 1 ELSE 2 END,
            q.pco2_mmhg DESC,
            q.stable_key
    ) AS rn
FROM qualifying_candidates q
WHERE q.site = 'unknown'
),
first_other AS (
SELECT
    hadm_id,
    charttime AS first_other_hypercap_time_0_24h,
    pco2_mmhg AS first_other_hypercap_pco2_mmhg,
    source_branch AS first_other_hypercap_source_branch,
    itemid AS first_other_hypercap_itemid
FROM ranked_other
WHERE rn = 1
)
SELECT
    q.hadm_id,
    1 AS pco2_threshold_0_24h,
    1 AS pco2_threshold_any,
    f.abg_hypercap_threshold,
    f.vbg_hypercap_threshold,
    f.unknown_hypercap_threshold,
    q.qualifying_pco2_time,
    q.qualifying_pco2_mmhg,
    q.qualifying_site,

```

```

q.qualifying_source_branch,
q.qualifying_threshold_mmhg,
q.dt_qualifying_hypercapnia_hours,
q.qualifying_pco2_itemid,
e.max_pco2_0_6h,
e.max_pco2_0_24h,
s6.max_pco2_0_6h_source_branch,
s24.max_pco2_0_24h_source_branch,
a.first_abg_hypercap_time_0_24h,
a.first_abg_hypercap_pco2_mmhg,
a.first_abg_hypercap_source_branch,
a.first_abg_hypercap_itemid,
v.first_vbg_hypercap_time_0_24h,
v.first_vbg_hypercap_pco2_mmhg,
v.first_vbg_hypercap_source_branch,
v.first_vbg_hypercap_itemid,
o.first_other_hypercap_time_0_24h,
o.first_other_hypercap_pco2_mmhg,
o.first_other_hypercap_source_branch,
o.first_other_hypercap_itemid
FROM qualifying_first q
JOIN qualifying_flags f
  ON f.hadm_id = q.hadm_id
LEFT JOIN windowed_extrema e
  ON e.hadm_id = q.hadm_id
LEFT JOIN max_source_6h s6
  ON s6.hadm_id = q.hadm_id
LEFT JOIN max_source_24h s24
  ON s24.hadm_id = q.hadm_id
LEFT JOIN first_abg a
  ON a.hadm_id = q.hadm_id
LEFT JOIN first_vbg v
  ON v.hadm_id = q.hadm_id
LEFT JOIN first_other o
  ON o.hadm_id = q.hadm_id
"""

co2_thresh = run_sql_bq(sql("co2_thresholds_sql"))
print("Admissions meeting thresholds:", len(co2_thresh))
co2_thresh.head(3)

```

Admissions meeting thresholds: 6246

	hadm_id	pco2_threshold_0_24h	pco2_threshold_any	abg_hypercap_threshold	vbg_hypercap_threshold
0	20868411	1	1	1	1
1	23292518	1	1	0	1
2	23358783	1	1	1	1

2.3.3 3) Cohort enrollment and hadm_list for downstream queries

Rationale: Use ICD OR gas-qualified 0–24h encounters as the primary cohort, with explicit route decomposition.

```
# Purpose: Build gas-qualified 0-24h cohort flags and hadm inclusion list.

# Outer-join to retain ICD metadata and gas qualification metadata.
cohort_any = cohort_icd.merge(co2_thresh, how="outer", on="hadm_id")

# Fill missing flags with 0 where appropriate
icd_cols = ["ICD10_J9602", "ICD10_J9612", "ICD10_J9622", "ICD10_J9692", |
    ↵ "ICD10_E662", "ICD9_27803", "any_hypercap_icd", "any_hypercap_icd_hosp", |
    ↵ "any_hypercap_icd_ed"]
for c in icd_cols:
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].fillna(0).astype(int)

for c in [
    "abg_hypercap_threshold",
    "vbg_hypercap_threshold",
    "pco2_threshold_0_24h",
    "pco2_threshold_any",
    "unknown_hypercap_threshold",
]:
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].fillna(0).astype(int)

if "qualifying_pco2_time" in cohort_any.columns:
    cohort_any["qualifying_pco2_time"] = pd.to_datetime(
        cohort_any["qualifying_pco2_time"], errors="coerce"
    )
for c in [
    "first_abg_hypercap_time_0_24h",
    "first_vbg_hypercap_time_0_24h",
    "first_other_hypercap_time_0_24h",
]:
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].fillna(0).astype(int)
```

```

if c in cohort_any.columns:
    cohort_any[c] = pd.to_datetime(cohort_any[c], errors="coerce")
for c in [
    "qualifying_pco2_mmhg",
    "qualifying_threshold_mmhg",
    "dt_qualifying_hypercapnia_hours",
    "max_pco2_0_6h",
    "max_pco2_0_24h",
    "first_abg_hypercap_pco2_mmhg",
    "first_vbg_hypercap_pco2_mmhg",
    "first_other_hypercap_pco2_mmhg",
]:
    if c in cohort_any.columns:
        cohort_any[c] = pd.to_numeric(cohort_any[c], errors="coerce")
for c in [
    "qualifying_site",
    "qualifying_source_branch",
    "max_pco2_0_6h_source_branch",
    "max_pco2_0_24h_source_branch",
    "first_abg_hypercap_source_branch",
    "first_vbg_hypercap_source_branch",
    "first_other_hypercap_source_branch",
]:
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].astype("string")

# Final enrollment flags (primary definition: ICD OR first qualifying
# hypercapnic gas in 0-24h)
if "pco2_threshold_0_24h" not in cohort_any.columns:
    cohort_any["pco2_threshold_0_24h"] = (
        (cohort_any["abg_hypercap_threshold"] == 1)
        | (cohort_any["vbg_hypercap_threshold"] == 1)
        | (cohort_any["unknown_hypercap_threshold"] == 1)
    ).astype(int)
cohort_any["pco2_threshold_any"] =
    cohort_any["pco2_threshold_0_24h"].astype(int)
if "any_hypercap_icd" not in cohort_any.columns:
    cohort_any["any_hypercap_icd"] = 0
cohort_any["any_hypercap_icd"] = (
    pd.to_numeric(cohort_any["any_hypercap_icd"],
    errors="coerce").fillna(0).astype(int)
)
cohort_any["enrolled_any"] = (

```

```

(cohort_any["any_hypercap_icd"] == 1)
| (cohort_any["pco2_threshold_0_24h"] == 1)
).astype(int)
cohort_any["enrolled_any_icd_union_secondary"] = (
    cohort_any["any_hypercap_icd"] == 1)
| (cohort_any["pco2_threshold_0_24h"] == 1)
).astype(int)
cohort_any["enrollment_route"] = np.select(
    [
        (cohort_any["any_hypercap_icd"] == 0) &
        (cohort_any["pco2_threshold_0_24h"] == 1),
        (cohort_any["any_hypcap_icd"] == 1) &
        (cohort_any["pco2_threshold_0_24h"] == 0),
        (cohort_any["any_hypcap_icd"] == 1) &
        (cohort_any["pco2_threshold_0_24h"] == 1),
        (cohort_any["pco2_threshold_0_24h"] == 1),
    ],
    ["GAS_ONLY", "ICD_ONLY", "ICD+GAS"],
    default="NONE",
)
print("ICD-positive admissions           :",
      int((cohort_any["any_hypcap_icd"] == 1).sum()))
print("Gas-qualified within 24h admissions       :",
      int((cohort_any["pco2_threshold_0_24h"] == 1).sum()))
print(
    "ICD gas union admissions           :",
    int((cohort_any["enrolled_any_icd_union_secondary"] == 1).sum()),
)
print("Final enrolled admissions (ICD OR gas 0-24h):",
      int((cohort_any["enrolled_any"] == 1).sum()))
print("Enrollment routes:")
print(cohort_any["enrollment_route"].value_counts(dropna=False).to_string())
if {"qualifying_source_branch",
     "pco2_threshold_0_24h"}.issubset(cohort_any.columns):
    qualifying_branch = (
        cohort_any.loc[cohort_any["pco2_threshold_0_24h"] == 1,
                      "qualifying_source_branch"]
            .fillna("UNKNOWN")
            .astype("string"))
    qualifying_branch_counts = qualifying_branch.value_counts(dropna=False)
    qualifying_branch_rates = (
        qualifying_branch_counts

```

```

        / max(int(cohort_any["pco2_threshold_0_24h"].sum()), 1)
    ).rename("rate")
print("Qualifying source branch distribution (0-24h):")
print(
    pd.concat(
        [qualifying_branch_counts.rename("count"),
    ↳ qualifying_branch_rates],
        axis=1,
    ).to_string()
)

# New hadm list used for the rest of the notebook
hadm_list = cohort_any.loc[cohort_any["enrolled_any"]==1,
    "hadm_id"].dropna().astype("int64").tolist()
len(hadm_list)

ICD-positive admissions : 4237
Gas-qualified within 24h admissions : 6246
ICD gas union admissions : 9406
Final enrolled admissions (ICD OR gas 0-24h): 9406
Enrollment routes:
enrollment_route
GAS_ONLY      5169
ICD_ONLY      3160
ICD+GAS       1077
Qualifying source branch distribution (0-24h):
            count      rate
qualifying_source_branch
LAB          5383  0.861832
POC          863   0.138168

```

9406

2.3.4 4) First ABG and First VBG (LAB + POC, standardized to mmHg)

Rationale: Extract earliest ABG/VBG measurements to characterize baseline physiology with standardized units.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
    ↳ consistent across notebook stages.
```

```

params = {"hadms": hadm_list}

bg_pairs_sql = rf"""
WITH hadms AS (SELECT hadm_id FROM UNNEST(@hadms) AS hadm_id),

/* ----- LAB (HOSP) ----- */
hosp_cand AS (
SELECT
    le.subject_id, le.hadm_id, le.charttime, le.specimen_id, le.itemid,
    CAST(le.valuenum AS FLOAT64) AS val,
    LOWER(COALESCE(le.valueuom,'')) AS uom,
    LOWER(di.label) AS lbl,
    LOWER(COALESCE(di.fluid,'')) AS fl,
    LOWER(COALESCE(di.category,'')) AS cat
FROM `{{PHYS}}.{{HOSP}}.labevents` le
JOIN `{{PHYS}}.{{HOSP}}.d_labitems` di ON di.itemid = le.itemid
JOIN hadms h ON h.hadm_id = le.hadm_id
WHERE le.valuenum IS NOT NULL
    AND le.itemid IN UNNEST(
        ARRAY_CONCAT({{LAB_PCO2_SQL_ARRAY}}, {{LAB_PH_SQL_ARRAY}},
        {{LAB_P02_SQL_ARRAY}})
    )
    AND LOWER(COALESCE(di.fluid,'')) IN {{LAB_BLOOD_FLUID_SQL_LIST}}
),
hosp_spec AS (
SELECT le.specimen_id, LOWER(COALESCE(le.value,'')) AS spec_val
FROM `{{PHYS}}.{{HOSP}}.labevents` le
WHERE le.specimen_id IS NOT NULL
    AND le.itemid IN UNNEST({{LAB_SPECIMEN_SQL_ARRAY}})
QUALIFY ROW_NUMBER() OVER (
    PARTITION BY le.specimen_id
    ORDER BY le.charttime DESC
) = 1
),
hosp_class AS (
SELECT
    c.hadm_id, c.charttime, c.specimen_id, c.itemid, c.val, c.uom, c.lbl,
    c.fl,
    CASE
        WHEN c.itemid IN UNNEST({{LAB_PH_SQL_ARRAY}}) THEN 'ph'
        WHEN c.itemid IN UNNEST({{LAB_P02_SQL_ARRAY}}) THEN 'po2'
        WHEN c.itemid IN UNNEST({{LAB_PCO2_SQL_ARRAY}}) THEN 'pco2'
        ELSE NULL
)

```

```

        END AS analyte,
        CASE
            WHEN REGEXP_CONTAINS(s.spec_val, r'arter') OR
        ↵  REGEXP_CONTAINS(s.spec_val, r'\bart\b') THEN 'arterial'
            WHEN REGEXP_CONTAINS(s.spec_val, r'ven|mixed|central') THEN 'venous'
            ELSE 'unknown'
        END AS site
    FROM hosp_cand c
    LEFT JOIN hosp_spec s USING (specimen_id)
),
hosp_pairs AS (
    SELECT
        hadm_id, specimen_id,
        MIN(charttime) AS sample_time,
        MAX(IF(analyte='ph', val, NULL)) AS ph,
        MAX(IF(analyte='po2', val, NULL)) AS po2_raw,
        MAX(IF(analyte='pco2', val, NULL)) AS pco2_raw,
        (ARRAY_AGG(IF(analyte='po2', uom, NULL) IGNORE NULLS LIMIT
        ↵  1)) [OFFSET(0)] AS po2_uom,
        (ARRAY_AGG(IF(analyte='pco2', uom, NULL) IGNORE NULLS LIMIT
        ↵  1)) [OFFSET(0)] AS pco2_uom,
        (ARRAY_AGG(site IGNORE NULLS LIMIT 1)) [OFFSET(0)] AS site
    FROM hosp_class
    GROUP BY hadm_id, specimen_id
    HAVING pco2_raw IS NOT NULL AND site IN ('arterial','venous','unknown')
),
hosp_pairs_std AS (
    SELECT
        hadm_id, specimen_id, sample_time, site,
        ph,
        CASE WHEN ph IS NOT NULL THEN 'unitless' ELSE NULL END AS ph_uom,
        CASE WHEN po2_uom = 'kpa' THEN po2_raw * 7.50062 ELSE po2_raw END AS
        ↵  po2_mmHg,
        'mmhg' AS po2_uom_norm,
        CASE WHEN pco2_uom = 'kpa' THEN pco2_raw * 7.50062 ELSE pco2_raw END AS
        ↵  pco2_mmHg,
        'mmhg' AS pco2_uom_norm
    FROM hosp_pairs
    WHERE (ph IS NULL OR (ph BETWEEN 6.3 AND 7.8))
        AND (po2_raw IS NULL OR (CASE WHEN po2_uom='kpa' THEN po2_raw*7.50062
        ↵  ELSE po2_raw END) BETWEEN 5 AND 600)
        AND (pco2_raw IS NULL OR (CASE WHEN pco2_uom='kpa' THEN pco2_raw*7.50062
        ↵  ELSE pco2_raw END) BETWEEN 5 AND 200)

```

```

),
lab_abg AS (
    SELECT hadm_id,
        ph          AS lab_abg_ph,
        ph_uom      AS lab_abg_ph_uom,
        po2_mmHg    AS lab_abg_po2,
        pco2_mmHg   AS lab_abg_paco2,
        'mmhg'       AS lab_abg_paco2_uom,
        sample_time  AS lab_abg_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
        ↵ sample_time) rn
        FROM hosp_pairs_std WHERE site='arterial') WHERE rn=1
),
lab_vbg AS (
    SELECT hadm_id,
        ph          AS lab_vbg_ph,
        ph_uom      AS lab_vbg_ph_uom,
        po2_mmHg    AS lab_vbg_po2,
        pco2_mmHg   AS lab_vbg_paco2,
        'mmhg'       AS lab_vbg_paco2_uom,
        sample_time  AS lab_vbg_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
        ↵ sample_time) rn
        FROM hosp_pairs_std WHERE site='venous') WHERE rn=1
),
lab_other AS (
    SELECT hadm_id,
        ph          AS lab_other_ph,
        ph_uom      AS lab_other_ph_uom,
        po2_mmHg    AS lab_other_po2,
        pco2_mmHg   AS lab_other_paco2,
        'mmhg'       AS lab_other_paco2_uom,
        sample_time  AS lab_other_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
        ↵ sample_time) rn
        FROM hosp_pairs_std WHERE site='unknown') WHERE rn=1
),
/* ----- POC (ICU) ----- */
icu_item_map AS (
    SELECT
        itemid,
        LOWER(label) AS lbl,

```

```

CASE
    WHEN itemid IN UNNEST({ICU_PH_SQL_ARRAY}) THEN 'ph'
    WHEN itemid IN UNNEST({ICU_P02_SQL_ARRAY}) THEN 'po2'
    WHEN itemid IN UNNEST({ICU_PCO2_SQL_ARRAY}) THEN 'pco2'
    ELSE NULL
END AS analyte
FROM `{{PHYS}}.{{ICU}}.d_items`
WHERE itemid IN UNNEST(
    ARRAY_CONCAT({ICU_PH_SQL_ARRAY}, {ICU_P02_SQL_ARRAY},
    {ICU_PCO2_SQL_ARRAY})
)
),
icu_spec AS (
SELECT
    ce.stay_id,
    ce.charttime,
    LOWER(COALESCE(ce.value, '')) AS spec_val
FROM `{{PHYS}}.{{ICU}}.chartevents` ce
WHERE ce.itemid IN UNNEST({ICU_SPECIMEN_SQL_ARRAY})
    AND ce.value IS NOT NULL
),
icu_raw AS (
SELECT
    ie.hadm_id, ce.stay_id, ce.itemid, ce.charttime,
    di.lbl,
    di.analyte,
    CASE
        WHEN ce.itemid IN UNNEST({ICU_PCO2_ABG_SQL_ARRAY}) THEN 'arterial'
        WHEN ce.itemid IN UNNEST({ICU_PCO2_VBG_SQL_ARRAY}) THEN 'venous'
        WHEN ce.itemid IN UNNEST({ICU_P02_ABG_SQL_ARRAY}) THEN 'arterial'
        WHEN ce.itemid IN UNNEST({ICU_P02_VBG_SQL_ARRAY}) THEN 'venous'
        ELSE 'unknown'
    END AS site_itemid,
    LOWER(REPLACE(COALESCE(ce.value uom, ''), ' ', '')) AS uom_nospace,
    LOWER(COALESCE(ce.value, '')) AS valstr,
    COALESCE(
        CAST(ce.valuenum AS FLOAT64),
        SAFE_CAST(REGEXP_EXTRACT(LOWER(ce.value), r'(-?\d+(?:\.\d+)?))') AS
        FLOAT64)
    ) AS val
FROM `{{PHYS}}.{{ICU}}.chartevents` ce
JOIN icu_item_map di ON di.itemid = ce.itemid
JOIN `{{PHYS}}.{{ICU}}.icustays` ie ON ie.stay_id = ce.stay_id

```

```

    JOIN hadms h ON h.hadm_id = ie.hadm_id
),
icu_cand_ranked AS (
    SELECT
        r.*,
        s.spec_val,
        ROW_NUMBER() OVER (
            PARTITION BY r.hadm_id, r.stay_id, r.itemid, r.charttime
            ORDER BY ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, SECOND))
        ) AS spec_rank
    FROM icu_raw r
    LEFT JOIN icu_spec s
        ON s.stay_id = r.stay_id
        AND ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, MINUTE)) <= 120
),
icu_cand AS (
    SELECT
        hadm_id, stay_id, itemid, charttime, lbl, analyte, uom_nospace, valstr,
        val,
        CASE
            WHEN REGEXP_CONTAINS(COALESCE(spec_val, ''), r'arter|\\bart\\b') THEN
            'arterial'
            WHEN REGEXP_CONTAINS(COALESCE(spec_val, ''), r'ven|central|mixed') THEN
            'venous'
            WHEN siteitemid IN ('arterial', 'venous') THEN siteitemid
            ELSE 'unknown'
        END AS site
    FROM icu_cand_ranked
    WHERE val IS NOT NULL
        AND analyte IN ('ph', 'pco2', 'po2')
        AND (spec_rank = 1 OR spec_val IS NULL)
),
icu_ph AS (
    SELECT hadm_id, stay_id, charttime, val AS ph, site AS site_ph
    FROM icu_cand WHERE analyte='ph'
),
icu_po2 AS (
    SELECT hadm_id, stay_id, charttime, val AS po2_raw, uom_nospace, site AS
        site_po2
    FROM icu_cand WHERE analyte='po2'
),
icu_co2 AS (

```

```

SELECT hadm_id, stay_id, charttime, val AS pco2_raw, uom_nospace, valstr,
    site AS site_co2
    FROM icu_cand WHERE analyte='pco2'
),
icu_pair_win AS (
    SELECT
        c.hadm_id, c.stay_id,
        c.site_co2 AS site,
        p.charttime AS ph_time,
        c.charttime AS co2_time,
        p.ph,
        CASE
            WHEN c.uom_nospace='kpa' THEN 'kpa'
            WHEN c.uom_nospace='mmhg' THEN 'mmhg'
            ELSE c.uom_nospace
        END AS pco2_uom_norm_raw,
        c.pco2_raw,
        ABS(TIMESTAMP_DIFF(c.charttime, p.charttime, SECOND)) AS dt_sec_ph
    FROM icu_co2 c
    LEFT JOIN icu_ph p
        ON p.hadm_id = c.hadm_id
        AND p.stay_id = c.stay_id
        AND (
            p.site_ph = c.site_co2
            OR p.site_ph IS NULL
            OR p.site_ph = 'unknown'
        )
        AND ABS(TIMESTAMP_DIFF(c.charttime, p.charttime, MINUTE)) <= 10
    QUALIFY ROW_NUMBER() OVER (
        PARTITION BY c.hadm_id, c.stay_id, c.charttime
        ORDER BY
            CASE WHEN p.charttime = c.charttime THEN 0 ELSE 1 END,
            dt_sec_ph
        ) = 1
),
icu_triplet_win AS (
    SELECT
        cp.hadm_id,
        cp.stay_id,
        cp.site,
        cp.ph_time,
        cp.co2_time,
        cp.ph,

```

```

cp.pco2_uom_norm_raw,
cp.pco2_raw,
o.charttime AS po2_time,
o.po2_raw,
CASE
    WHEN o.uom_nospace='kpa' THEN 'kpa'
    WHEN o.uom_nospace='mmhg' THEN 'mmhg'
    ELSE o.uom_nospace
END AS po2_uom_norm_raw,
ABS(TIMESTAMP_DIFF(cp.co2_time, o.charttime, SECOND)) AS dt_sec_po2
FROM icu_pair_win cp
LEFT JOIN icu_po2 o
    ON o.hadm_id = cp.hadm_id
    AND o.stay_id = cp.stay_id
    AND (
        o.site_po2 = cp.site
        OR o.site_po2 IS NULL
        OR o.site_po2 = 'unknown'
    )
    AND ABS(TIMESTAMP_DIFF(cp.co2_time, o.charttime, MINUTE)) <= 10
QUALIFY ROW_NUMBER() OVER (
    PARTITION BY cp.hadm_id, cp.stay_id, cp.co2_time
    ORDER BY
        CASE WHEN o.charttime = cp.co2_time THEN 0 ELSE 1 END,
        dt_sec_po2
) = 1
),
icu_triplet_std AS (
SELECT
    hadm_id,
    stay_id,
    site,
    COALESCE(ph_time, co2_time, po2_time) AS sample_time,
    ph,
    CASE WHEN ph IS NOT NULL THEN 'unitless' ELSE NULL END AS ph_uom,
    CASE WHEN po2_uom_norm_raw='kpa' THEN po2_raw*7.50062 ELSE po2_raw END AS
    ← po2_mmHg,
    'mmhg' AS po2_uom_norm,
    CASE WHEN pco2_uom_norm_raw='kpa' THEN pco2_raw*7.50062 ELSE pco2_raw END
    ← AS pco2_mmHg,
    'mmhg' AS pco2_uom_norm
FROM icu_triplet_win
WHERE (ph BETWEEN 6.3 AND 7.8 OR ph IS NULL)

```

```

        AND (
            po2_raw IS NULL
            OR (CASE WHEN po2_uom_norm_raw='kpa' THEN po2_raw*7.50062 ELSE po2_raw
        ← END) BETWEEN 5 AND 600
        )
        AND (CASE WHEN pco2_uom_norm_raw='kpa' THEN pco2_raw*7.50062 ELSE
        ← pco2_raw END) BETWEEN 5 AND 200
    ),
    icu_all AS (
        SELECT * FROM icu_triplet_std
    ),

    poc_abg AS (
        SELECT hadm_id,
            ph          AS poc_abg_ph,
            ph_uom      AS poc_abg_ph_uom,
            po2_mmHg    AS poc_abg_po2,
            pco2_mmHg   AS poc_abg_paco2,
            'mmhg'       AS poc_abg_paco2_uom,
            sample_time AS poc_abg_time
        FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
        ← sample_time) rn
            FROM icu_all WHERE site='arterial') WHERE rn=1
    ),
    poc_vbg AS (
        SELECT hadm_id,
            ph          AS poc_vbg_ph,
            ph_uom      AS poc_vbg_ph_uom,
            po2_mmHg    AS poc_vbg_po2,
            pco2_mmHg   AS poc_vbg_paco2,
            'mmhg'       AS poc_vbg_paco2_uom,
            sample_time AS poc_vbg_time
        FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
        ← sample_time) rn
            FROM icu_all WHERE site='venous') WHERE rn=1
    ),
    poc_other AS (
        SELECT hadm_id,
            ph          AS poc_other_ph,
            ph_uom      AS poc_other_ph_uom,
            po2_mmHg    AS poc_other_po2,
            pco2_mmHg   AS poc_other_paco2,
            'mmhg'       AS poc_other_paco2_uom,

```

```

        sample_time AS poc_other_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
    ↵ sample_time) rn
        FROM icu_all WHERE site='unknown') WHERE rn=1
    )

/* ----- Final one row per hadm ----- */
SELECT
    h.hadm_id,
    -- LAB-ABG / LAB-VBG / LAB-OTHER
    la.lab_abg_ph, la.lab_abg_ph_uom, la.lab_abg_po2, la.lab_abg_paco2,
    ↵ la.lab_abg_paco2_uom, la.lab_abg_time,
    lv.lab_vbg_ph, lv.lab_vbg_ph_uom, lv.lab_vbg_po2, lv.lab_vbg_paco2,
    ↵ lv.lab_vbg_paco2_uom, lv.lab_vbg_time,
    lo.lab_other_ph, lo.lab_other_ph_uom, lo.lab_other_po2, lo.lab_other_paco2,
    ↵ lo.lab_other_paco2_uom, lo.lab_other_time,
    -- POC-ABG / POC-VBG / POC-OTHER
    pa.poc_abg_ph, pa.poc_abg_ph_uom, pa.poc_abg_po2, pa.poc_abg_paco2,
    ↵ pa.poc_abg_paco2_uom, pa.poc_abg_time,
    pv.poc_vbg_ph, pv.poc_vbg_ph_uom, pv.poc_vbg_po2, pv.poc_vbg_paco2,
    ↵ pv.poc_vbg_paco2_uom, pv.poc_vbg_time,
    po.poc_other_ph, po.poc_other_ph_uom, po.poc_other_po2, po.poc_other_paco2,
    ↵ po.poc_other_paco2_uom, po.poc_other_time,
    -- First ABG across LAB+POC
    (SELECT AS STRUCT src, t, ph, pco2, po2
        FROM (SELECT 'LAB' AS src, la.lab_abg_time AS t, la.lab_abg_ph AS ph,
    ↵ la.lab_abg_paco2 AS pco2, la.lab_abg_po2 AS po2
            UNION ALL
            SELECT 'POC', pa.poc_abg_time, pa.poc_abg_ph, pa.poc_abg_paco2,
    ↵ pa.poc_abg_po2)
        WHERE t IS NOT NULL
        ORDER BY t LIMIT 1) AS first_abg,
    -- First VBG across LAB+POC
    (SELECT AS STRUCT src, t, ph, pco2, po2
        FROM (SELECT 'LAB' AS src, lv.lab_vbg_time AS t, lv.lab_vbg_ph AS ph,
    ↵ lv.lab_vbg_paco2 AS pco2, lv.lab_vbg_po2 AS po2
            UNION ALL
            SELECT 'POC', pv.poc_vbg_time, pv.poc_vbg_ph, pv.poc_vbg_paco2,
    ↵ pv.poc_vbg_po2)
        WHERE t IS NOT NULL
        ORDER BY t LIMIT 1) AS first_vbg,
    -- First UNKNOWN-source pCO2 across LAB+POC (definitive pCO2, unresolved
    ↵ sample type).

```

```

(SELECT AS STRUCT src, t, ph, pco2, po2
FROM (
    SELECT 'LAB_BG_UNKNOWN' AS src, lo.lab_other_time AS t,
    ↵ lo.lab_other_ph AS ph, lo.lab_other_paco2 AS pco2, lo.lab_other_po2 AS
    ↵ po2
        UNION ALL
        SELECT 'POC_BG_UNKNOWN', po.poc_other_time, po.poc_other_ph,
    ↵ po.poc_other_paco2, po.poc_other_po2
)
WHERE t IS NOT NULL
ORDER BY t LIMIT 1) AS first_other
FROM hadms h
LEFT JOIN lab_abg la USING (hadm_id)
LEFT JOIN lab_vbg lv USING (hadm_id)
LEFT JOIN lab_other lo USING (hadm_id)
LEFT JOIN poc_abg pa USING (hadm_id)
LEFT JOIN poc_vbg pv USING (hadm_id)
LEFT JOIN poc_other po USING (hadm_id)
"""

bg_pairs = run_sql_bq(bg_pairs_sql, params)

# Flatten STRUCTs for first_abg, first_vbg, and first_other
for col in ["first_abg", "first_vbg", "first_other"]:
    if col in bg_pairs.columns:
        bg_pairs[f"{col}_src"] = bg_pairs[col].apply(lambda x: x.get("src"))
    ↵ if isinstance(x, dict) else None)
        bg_pairs[f"{col}_time"] = bg_pairs[col].apply(lambda x: x.get("t"))
    ↵ if isinstance(x, dict) else None)
        bg_pairs[f"{col}_ph"] = bg_pairs[col].apply(lambda x: x.get("ph"))
    ↵ if isinstance(x, dict) else None)
        bg_pairs[f"{col}_pco2"] = bg_pairs[col].apply(lambda x:
    ↵ x.get("pco2")if isinstance(x, dict) else None)
            bg_pairs[f"{col}_po2"] = bg_pairs[col].apply(lambda x: x.get("po2"))
    ↵ if isinstance(x, dict) else None)
        bg_pairs = bg_pairs.drop(columns=[col])

if {"first_other_src", "poc_other_paco2"}.issubset(bg_pairs.columns):
    first_other_src_upper = (
        bg_pairs["first_other_src"].astype("string").str.upper().fillna(""))
    )
    bg_pairs["first_other_src_detail"] = np.select(
        [

```

```

        first_other_src_upper.eq("LAB_BG_UNKNOWN"),
        first_other_src_upper.eq("POC_BG_UNKNOWN"),
        pd.to_numeric(bg_pairs["poc_other_paco2"]),
    ↵  errors="coerce").notna(),
    ],
    [
        "lab_bg_unknown",
        "poc_bg_unknown",
        "poc_bg_unknown_available",
    ],
    default="missing",
)
else:
    bg_pairs["first_other_src_detail"] = "missing"

bg_pairs.head(3)

```

	hadm_id	lab_abg_ph	lab_abg_ph_uom	lab_abg_po2	lab_abg_paco2	lab_abg_paco2_uom	l
0	20001305	NaN	None	NaN	NaN	None	I
1	20002497	7.49	unitless	61.0	44.0	mmhg	2
2	20002950	7.33	unitless	95.0	42.0	mmhg	2

2.3.5 5) Demographics & outcomes

Rationale: Build baseline covariates used for descriptive statistics and potential confounding adjustment.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪  consistent across notebook stages.
```

```

SQL["demo_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
    a.hadm_id,
    a.subject_id,
    a.admittime,
    a.dischtime,
    a.deathtime,
    a.admission_type AS administrative_class,
    a.admission_location AS from_location,

```

```

a.discharge_location,
a.insurance,
-- LOS (days)
TIMESTAMP_DIFF(a.dischtime, a.admittime, HOUR) / 24.0 AS hosp_los_days,
-- in-hospital death
IF(a.deathtime IS NOT NULL, 1, 0) AS death_in_hosp,
-- demographics
p.gender,
SAFE_CAST(ROUND(p.anchor_age + (EXTRACT(YEAR FROM a.admittime) -
→ p.anchor_year), 1) AS FLOAT64) AS age_at_admit,
-- 30-day all-cause mortality from admission
IF(p.dod IS NOT NULL AND DATE_DIFF(DATE(p.dod), DATE(a.admittime), DAY)
→ BETWEEN 0 AND 30, 1, 0) AS death_30d
FROM `{{PHYS}}.{{HOSP}}.admissions` a
JOIN hadms h USING (hadm_id)
JOIN `{{PHYS}}.{{HOSP}}.patients` p USING (subject_id)
"""
demo = run_sql_bq(sql("demo_sql"), {"hadms": hadm_list})
print("Demo rows:", len(demo))
demo.head(3)

```

Demo rows: 9406

	hadm_id	subject_id	admittime	dischtime	deathtime	administrative_class	fr
0	22447711	13352386	2188-03-25 01:17:00	2188-03-26 16:00:00	NaT	EW EMER.	EM
1	20930337	18079481	2152-05-13 20:36:00	2152-05-17 11:00:00	NaT	EW EMER.	EM
2	22834654	11482582	2125-10-30 01:03:00	2125-11-06 18:35:00	NaT	EW EMER.	EM

```

# Purpose: Create reusable data-quality and merge guardrail helpers to
→ prevent silent join errors.

# ===== Drop-in: safe merge utilities (one cell, run once) =====
import pandas as pd
from typing import Iterable, Optional, Literal

def _ensure_Int64(s: pd.Series) -> pd.Series:
    """Coerce to pandas nullable Int64 (preserves NA)."""
    return pd.to_numeric(s, errors="coerce").astype("Int64")

def strip_subject_cols(fr: pd.DataFrame) -> pd.DataFrame:

```

```

"""Remove any subject_id-like columns from a frame (e.g., 'subject_id',
↳   'Subject_ID')."""
return fr.drop(columns=[c for c in fr.columns if
↳   c.lower().startswith("subject_id")],
               errors="ignore")

def safe_merge_on_hadm(
    left: pd.DataFrame,
    right: pd.DataFrame,
    *,
    right_name: str,
    take: Optional[Iterable[str]] = None,
    order_by: Optional[Iterable[str]] = None,
    check_subject: Literal[False, "warn", "raise"] = False,
) -> pd.DataFrame:
    """
    Left-merge 'right' into 'left' on hadm_id, returning a copy of left with
    ↳ right's columns.
    - Dedupes right on hadm_id (optionally using order_by to pick the first
    ↳ row).
    - Optionally restricts right columns via `take`.
    - Optionally audits subject_id agreement before dropping subject_id from
    ↳ right.
    - Always strips subject_id-like columns from the right to prevent *_x/_y
    ↳ suffixes.
    - Raises if any *_x/_y suffixes still appear (indicates overlapping names
    ↳ besides hadm_id).
    """
    if "hadm_id" not in left.columns:
        raise KeyError(f"left frame lacks hadm_id before merging
            ↳ {right_name}")
    if "hadm_id" not in right.columns:
        raise KeyError(f"{right_name} lacks hadm_id")

    L = left.copy()
    R = right.copy()

    # Standardize dtypes of keys
    L["hadm_id"] = _ensure_Int64(L["hadm_id"])
    R["hadm_id"] = _ensure_Int64(R["hadm_id"])
    if "subject_id" in L.columns:
        L["subject_id"] = _ensure_Int64(L["subject_id"])
    if "subject_id" in R.columns:

```

```

R["subject_id"] = _ensure_Int64(R["subject_id"])

# Dedupe RIGHT by hadm_id (optionally order_by first)
if order_by:
    R = (R.sort_values(list(order_by))
          .drop_duplicates(subset=["hadm_id"], keep="first"))
else:
    R = R.drop_duplicates(subset=["hadm_id"], keep="first")

# Optional subject_id consistency audit (before stripping)
if check_subject and ("subject_id" in L.columns) and ("subject_id" in
    R.columns):
    # Join only on hadm_id where both sides have subject_id
    tmp = (L[["hadm_id", "subject_id"]]
            .merge(R[["hadm_id", "subject_id"]],
                   on="hadm_id", how="inner", suffixes=("_L", "_R")))
    mism = (tmp["subject_id_L"].notna() & tmp["subject_id_R"].notna() &
            (tmp["subject_id_L"] != tmp["subject_id_R"]))
    n_mism = int(mism.sum())
    if n_mism > 0:
        sample_ids = tmp.loc[mism, "hadm_id"].head(10).tolist()
        msg = (f"[{right_name}] subject_id mismatch on {n_mism}"
    ↵ hadm_id(s). "
                f"Examples: {sample_ids}")
        if check_subject == "raise":
            raise ValueError(msg)
        else:
            print("WARNING:", msg)

# Limit right columns (avoid accidental overlaps)
if take is not None:
    keep = ["hadm_id"] + [c for c in take if c != "hadm_id"]
    R = R[[c for c in keep if c in R.columns]]

# Always strip subject_id-like columns from right to prevent *_x/_y
R = strip_subject_cols(R)

# Final merge
out = L.merge(R, on="hadm_id", how="left", suffixes=("L", "R"))

# Guard: no suffixes should be present
bad = [c for c in out.columns if c.endswith("_x") or c.endswith("_y")]
if bad:

```

```

        raise RuntimeError(
            f"Merge with {right_name} produced suffixed columns {bad}. "
            "You likely have overlapping column names other than hadm_id."
        )
    return out

print("Safe merge helpers loaded.")

```

Safe merge helpers loaded.

2.3.6 6) NIH/OMB race & ethnicity (ED + Hospital)

Rationale: Harmonize race/ethnicity across sources using NIH/OMB categories for consistent reporting.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪ consistent across notebook stages.

race_eth_sql = rf"""
WITH hadms AS (
    SELECT x AS hadm_id
    FROM UNNEST(@hadms) AS x
),
-- Hospital admission "race" text
hosp AS (
    SELECT a.hadm_id, LOWER(TRIM(a.race)) AS race_hosp_raw
    FROM `{{PHYS}}.{{HOSP}}.admissions` a
    JOIN hadms hm USING (hadm_id)
),
-- Earliest ED stay leading to the admission; take its "race" text if present
ed_first AS (
    SELECT
        e.hadm_id,
        (ARRAY_AGG(STRUCT(e.intime AS intime, LOWER(TRIM(e.race)) AS race_ed_raw)
            ORDER BY e.intime ASC LIMIT 1)) [OFFSET(0)] AS pick
    FROM `{{PHYS}}.{{ED}}.edstays` e
    JOIN hadms hm USING (hadm_id)
    GROUP BY e.hadm_id
),
ed AS (
    SELECT hadm_id, pick.race_ed_raw

```

```

    FROM ed_first
),

-- Combine ED + Hospital for maximum coverage
comb AS (
    SELECT
        hm.hadm_id,
        ho.race_hosp_raw,
        ed.race_ed_raw,
        TRIM(REGEXP_REPLACE(CONCAT(COALESCE(ho.race_hosp_raw,''), ' ',
        COALESCE(ed.race_ed_raw,''))), r'\s+', ' ')) AS race_text_any
    FROM hadms hm
    LEFT JOIN hosp ho USING (hadm_id)
    LEFT JOIN ed ed USING (hadm_id)
),

-- Tokenization to OMB families + Hispanic ethnicity
tok AS (
    SELECT
        hadm_id, race_hosp_raw, race_ed_raw, race_text_any,

        -- Ethnicity (Hispanic)
        REGEXP_CONTAINS(race_text_any, r'\b(hispanic|latinx|latino|latina)\b') AS
        is_hisp,

        -- Race families (use boundaries to reduce false positives)
        REGEXP_CONTAINS(race_text_any, r'americans+indian|\balaska\b') AS
        is_aian,
        REGEXP_CONTAINS(race_text_any, r'\basian\b') AS is_asian,
        REGEXP_CONTAINS(race_text_any, r'\b(black|african\s+american)\b') AS
        is_black,
        REGEXP_CONTAINS(race_text_any, r'hawaiian|pacific\s+islander') AS
        is_nhopi,
        REGEXP_CONTAINS(race_text_any, r'\bwhite\b|caucasian') AS is_white,

        -- Unknown/other indicators
        REGEXP_CONTAINS(race_text_any,
        r'unknown|other|declined|unable|not\s+reported|missing|null') AS
        is_unknown_any,

        -- Multi-race hints
        REGEXP_CONTAINS(race_text_any,
        r'(two|2)\s+or\s+more|multi|biracial|multiracial') AS is_multi_hint
)

```

```

    FROM comb
),

-- Decide ethnicity per NIH
ethn AS (
    SELECT
        hadm_id, race_hosp_raw, race_ed_raw, race_text_any,
        CASE
            WHEN is_hisp THEN 'Hispanic or Latino'
            WHEN (race_text_any IS NULL OR race_text_any = '' OR is_unknown_any)
        ↵ THEN 'Unknown or Not Reported'
            ELSE 'Not Hispanic or Latino'
        END AS nih_ethnicity,
        (CAST(is_aian AS INT64) + CAST(is_asian AS INT64) + CAST(is_black AS
        ↵ INT64)
            + CAST(is_nhopi AS INT64) + CAST(is_white AS INT64)) AS race_hits,
        is_aian, is_asian, is_black, is_nhopi, is_white, is_multi_hint,
        ↵ is_unknown_any
        FROM tok
),
-- Decide race per NIH/OMB (1997)
race_assign AS (
    SELECT
        hadm_id, race_hosp_raw, race_ed_raw, race_text_any, nih_ethnicity,
        CASE
            WHEN race_hits >= 2 OR is_multi_hint THEN 'More than one race'
            WHEN is_aian THEN 'American Indian or Alaska Native'
            WHEN is_asian THEN 'Asian'
            WHEN is_black THEN 'Black or African American'
            WHEN is_nhopi THEN 'Native Hawaiian or Other Pacific Islander'
            WHEN is_white THEN 'White'
            WHEN is_unknown_any OR race_text_any IS NULL OR race_text_any = '' THEN
        ↵ 'Unknown or Not Reported'
            ELSE 'Unknown or Not Reported'
        END AS nih_race
        FROM ethn
)
SELECT hadm_id, race_hosp_raw, race_ed_raw, nih_race, nih_ethnicity
FROM race_assign
"""
race_eth = run_sql_bq(race_eth_sql, {"hadms": hadm_list})

```

```

print("Race/Eth rows:", len(race_eth))
race_eth.head(3)

```

Race/Eth rows: 9406

	hadm_id	race_hosp_raw	race_ed_raw	nih_race	nih_ethnicity
0	20001305	white	white	White	Not Hispanic or L
1	20002497	hispanic/latino - puerto rican	None	Unknown or Not Reported	Hispanic or Latino
2	20002950	white	white	White	Not Hispanic or L

2.3.7 7) ED triage (linked to hadm) and first ED vitals

Rationale: Capture ED presentation features (vitals and chief complaint) for symptom and severity analyses.

```

# Purpose: Pull ED triage/first-vitals by loading ED tables once and
#           filtering to cohort hadm_ids in pandas.

hadms_for_ed = set(
    pd.Series(hadm_list)
    .dropna()
    .astype("int64")
    .tolist()
)

# 1) ED stay map (full table), then restrict to cohort hadm_ids locally.
SQL["edmap_all_sql"] = f"""
SELECT stay_id, hadm_id, intime
FROM `{{PHYS}}.{ED}.edstays`
WHERE hadm_id IS NOT NULL
"""

edmap_all = run_sql_bq(sql("edmap_all_sql"))
edmap_all["hadm_id"] = pd.to_numeric(edmap_all["hadm_id"],
                                      errors="coerce").astype("Int64")
edmap_all["stay_id"] = pd.to_numeric(edmap_all["stay_id"],
                                      errors="coerce").astype("Int64")
edmap_all["intime"] = pd.to_datetime(edmap_all["intime"], errors="coerce")

edmap = edmap_all[edmap_all["hadm_id"].isin(hadms_for_ed)].copy()
edmap = edmap.drop_duplicates(subset=["stay_id", "hadm_id"])
print("ED map rows (cohort):", len(edmap))

```

```

if edmap.empty:
    ed_triage = pd.DataFrame(columns=[
        "hadm_id", "ed_triage_temp", "ed_triage_hr", "ed_triage_rr",
        "ed_triage_o2sat", "ed_triage_sbp", "ed_triage_dbp",
        ↵ "ed_triage_pain",
        "ed_triage_acuity", "ed_triage_cc",
    ])
    ed_first = pd.DataFrame(columns=[
        "hadm_id", "ed_first_vitals_time", "ed_first_temp", "ed_first_hr",
        "ed_first_rr", "ed_first_o2sat", "ed_first_sbp", "ed_first_dbp",
        "ed_first_rhythm", "ed_first_pain",
    ])
else:
    # 2) ED triage (full table), then reduce to earliest ED stay per hadm.
    SQL["ed_triage_all_sql"] = f"""
    SELECT
        stay_id,
        temperature      AS ed_triage_temp,
        heartrate        AS ed_triage_hr,
        resprate         AS ed_triage_rr,
        o2sat            AS ed_triage_o2sat,
        sbp              AS ed_triage_sbp,
        dbp              AS ed_triage_dbp,
        pain             AS ed_triage_pain,
        acuity           AS ed_triage_acuity,
        chiefcomplaint  AS ed_triage_cc
    FROM `'{PHYS}`.{ED}.triage`
    """
    tri_all = run_sql_bq(sql("ed_triage_all_sql"))
    tri_all["stay_id"] = pd.to_numeric(tri_all["stay_id"],
                                      ↵ errors="coerce").astype("Int64")

    tri_merged = (
        edmap[["stay_id", "hadm_id", "intime"]]
        .merge(tri_all, on="stay_id", how="left")
    )
    ed_triage = (
        tri_merged
        .sort_values(["hadm_id", "intime"], na_position="last")
        .drop_duplicates(subset=["hadm_id"], keep="first")
        .drop(columns=["stay_id", "intime"])
        .reset_index(drop=True)
    )

```

```

)
print("ED triage rows:", len(ed_triage))

# 3) First vitals per stay across full ED vitals table, then reduce to
#      ↵ earliest vitals-time per hadm.
SQL["ed_first_vitals_all_sql"] = f"""
WITH vs_ranked AS (
    SELECT
        v.stay_id,
        v.charttime,
        v.temperature,
        v.heartrate,
        v.resprate,
        v.o2sat,
        v.sbp,
        v.dbp,
        v.rhythm,
        v.pain,
        ROW_NUMBER() OVER (PARTITION BY v.stay_id ORDER BY v.charttime) AS rn
    FROM `{{PHYS}}.{{ED}}.vitalsign` v
)
SELECT
    stay_id,
    charttime AS ed_first_vitals_time,
    temperature AS ed_first_temp,
    heartrate AS ed_first_hr,
    resprate AS ed_first_rr,
    o2sat AS ed_first_o2sat,
    sbp AS ed_first_sbp,
    dbp AS ed_first_dbp,
    rhythm AS ed_first_rhythm,
    pain AS ed_first_pain
FROM vs_ranked
WHERE rn = 1
"""

first_stay_all = run_sql_bq(sql("ed_first_vitals_all_sql"))
first_stay_all["stay_id"] = pd.to_numeric(first_stay_all["stay_id"],
                                          errors="coerce").astype("Int64")
first_stay_all["ed_first_vitals_time"] =
pd.to_datetime(first_stay_all["ed_first_vitals_time"], errors="coerce")

first_merged = (
    edmap[["stay_id", "hadm_id"]]
)

```

```

        .merge(first_stay_all, on="stay_id", how="left")
    )
ed_first = (
    first_merged
    .sort_values(["hadm_id", "ed_first_vitals_time"], na_position="last")
    .drop_duplicates(subset=["hadm_id"], keep="first")
    .drop(columns=["stay_id"])
    .reset_index(drop=True)
)
print("ED first vitals rows:", len(ed_first))

ED map rows (cohort): 7166
ED triage rows: 7152
ED first vitals rows: 7152

```

2.3.8 8) ICU meta (first ICU stay, LOS days)

Rationale: Summarize ICU exposure and length of stay to contextualize disease severity.

```

# Purpose: Build first-ICU metadata via full-table pull + local hadm filter
↪ to avoid large ARRAY parameter stalls.

SQL["icu_all_sql"] = f"""
SELECT
    hadm_id,
    stay_id AS first_icu_stay_id,
    intime AS icu_intime,
    outtime AS icu_outtime
FROM `{{PHYS}}.{{ICU}}.icustays`
WHERE hadm_id IS NOT NULL
"""

icu_all = run_sql_bq(sql("icu_all_sql"))
icu_all["hadm_id"] = pd.to_numeric(icu_all["hadm_id"],
↪ errors="coerce").astype("Int64")
icu_all["first_icu_stay_id"] = pd.to_numeric(icu_all["first_icu_stay_id"],
↪ errors="coerce").astype("Int64")
icu_all["icu_intime"] = pd.to_datetime(icu_all["icu_intime"],
↪ errors="coerce")
icu_all["icu_outtime"] = pd.to_datetime(icu_all["icu_outtime"],
↪ errors="coerce")

hadms_for_icu = set(pd.Series(hadm_list).dropna().astype("int64").tolist())

```

```

icu_all = icu_all[icu_all["hadm_id"].isin(hadms_for_icu)].copy()

icu_meta = (
    icu_all
    .sort_values(["hadm_id", "icu_intime", "first_icu_stay_id"],
    ↵ na_position="last")
    .drop_duplicates(subset=["hadm_id"], keep="first")
    .reset_index(drop=True)
)
icu_meta["icu_los_days"] = (
    (icu_meta["icu_outtime"] - icu_meta["icu_intime"]).dt.total_seconds() /
    ↵ 86400.0
)

print("ICU meta rows:", len(icu_meta))

```

ICU meta rows: 7309

2.3.9 9) Ventilation flags (ICD procedures)

Rationale: Identify IMV/NIV exposure as clinically relevant respiratory support indicators.

```
# Purpose: Build IMV/NIV ICD flags from targeted procedure codes, then filter
↪ to cohort hadm_ids locally.
```

```

SQL["vent_proc_sql"] = f"""
SELECT
    hadm_id,
    icd_version,
    REPLACE(icd_code, '.', '') AS code_norm
FROM `{{PHYS}}.{{HOSP}}.procedures_icd`
WHERE
    (icd_version = 10 AND icd_code IN
    ↵ ('5A1935Z','5A1945Z','5A1955Z','0BH17EZ','0BH18EZ','5A09357','5A09457','5A09557'))
    OR
    (icd_version = 9 AND REPLACE(icd_code, '.', '') IN
    ↵ ('9670','9671','9672','9604','9390','9391','9399'))
"""

vent_proc = run_sql_bq(sql("vent_proc_sql"))
vent_proc["hadm_id"] = pd.to_numeric(vent_proc["hadm_id"],
↪ errors="coerce").astype("Int64")

```

```

vent_proc["icd_version"] = pd.to_numeric(vent_proc["icd_version"],
    ↪ errors="coerce").astype("Int64")
vent_proc["code_norm"] = vent_proc["code_norm"].astype(str)

hadms_for_vent = set(pd.Series(hadm_list).dropna().astype("int64").tolist())
vent_proc = vent_proc[vent_proc["hadm_id"].isin(hadms_for_vent)].copy()

imv_codes_10 = {"5A1935Z", "5A1945Z", "5A1955Z", "0BH17EZ", "0BH18EZ"}
imv_codes_9 = {"9670", "9671", "9672", "9604"}
niv_codes_10 = {"5A09357", "5A09457", "5A09557"}
niv_codes_9 = {"9390", "9391", "9399"}

vent_proc["imv_hit"] = (
    ((vent_proc["icd_version"] == 10) &
    ↪ vent_proc["code_norm"].isin(imv_codes_10)) |
    ((vent_proc["icd_version"] == 9) &
    ↪ vent_proc["code_norm"].isin(imv_codes_9))
)
vent_proc["niv_hit"] = (
    ((vent_proc["icd_version"] == 10) &
    ↪ vent_proc["code_norm"].isin(niv_codes_10)) |
    ((vent_proc["icd_version"] == 9) &
    ↪ vent_proc["code_norm"].isin(niv_codes_9))
)

vent = (
    vent_proc
    .groupby("hadm_id", as_index=False)
    .agg(
        imv_flag=("imv_hit", "max"),
        niv_flag=("niv_hit", "max"),
    )
)
vent["imv_flag"] = vent["imv_flag"].astype(int)
vent["niv_flag"] = vent["niv_flag"].astype(int)
vent["any_vent_flag"] = ((vent["imv_flag"] == 1) | (vent["niv_flag"] ==
    ↪ 1)).astype(int)

# Preserve one row per hadm in hadm_list even if no vent procedure codes were
    ↪ found.
vent = pd.DataFrame({"hadm_id": pd.Series(hadm_list,
    ↪ dtype="Int64")}).drop_duplicates().merge(vent, on="hadm_id", how="left")

```

```

vent[["imv_flag", "niv_flag", "any_vent_flag"]] = vent[["imv_flag",
    ↪ "niv_flag", "any_vent_flag"]].fillna(0).astype(int)

print("Vent rows:", len(vent))

Vent rows: 9406

and from charts

# Purpose: Extract earliest NIV/IMV charttimes using prefiltered ventilation
    ↪ itemids for better performance.

SQL["vent_chart_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),

stays AS (
    SELECT DISTINCT hadm_id, stay_id
    FROM `{{PHYS}}.{{ICU}}.icustays`
    WHERE hadm_id IN (SELECT hadm_id FROM hadms)
), 

vent_itemids AS (
    SELECT itemid, LOWER(label) AS lbl
    FROM `{{PHYS}}.{{ICU}}.d_items`
    WHERE REGEXP_CONTAINS(LOWER(label), r'(vent|ventilator|mode|bipap|bi[-
        ]?pap|cpap|nippy|niv|mask|ett|endotracheal)')
), 

cand AS (
    SELECT
        s.hadm_id,
        ce.charttime,
        vi.lbl,
        LOWER(COALESCE(ce.value, '')) AS valstr
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    JOIN stays s ON s.stay_id = ce.stay_id
    JOIN vent_itemids vi ON vi.itemid = ce.itemid
), 

flags AS (
    SELECT
        hadm_id,
        MIN(IF(

```

```

        REGEXP_CONTAINS(lbl, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
        ]?pap|cpap)')
            OR REGEXP_CONTAINS(valstr, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
        ]?pap|cpap)'),
                charttime, NULL)) AS first_niv_time,
        MIN(IF(
            REGEXP_CONTAINS(lbl, r'(invasive
            ventilation|endotracheal|ett|mech(|anical)? vent|ventilator
            mode|ac/|simv|prvc|aprvc|pcv|vcv|assist\s*control)')
                OR REGEXP_CONTAINS(valstr, r'(invasive
            ventilation|endotracheal|ett|ac/|simv|prvc|aprvc|pcv|vcv|assist\s*control'),
                charttime, NULL)) AS first_imv_time,
        MAX(CASE
            WHEN REGEXP_CONTAINS(lbl, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
            ]?pap|cpap)')
                OR REGEXP_CONTAINS(valstr, r'(non[-
            ]?invasive|niv|nippv|bipap|bi[- ]?pap|cpap)')
                    THEN 1 ELSE 0 END) AS niv_chart_flag,
        MAX(CASE
            WHEN REGEXP_CONTAINS(lbl, r'(invasive
            ventilation|endotracheal|ett|mech(|anical)? vent|ventilator
            mode|ac/|simv|prvc|aprvc|pcv|vcv|assist\s*control)')
                OR REGEXP_CONTAINS(valstr, r'(invasive
            ventilation|endotracheal|ett|ac/|simv|prvc|aprvc|pcv|vcv|assist\s*control)')
                    THEN 1 ELSE 0 END) AS imv_chart_flag
    FROM cand
    GROUP BY hadm_id
)
SELECT hadm_id, niv_chart_flag, imv_chart_flag, first_niv_time,
    first_imv_time
FROM flags
"""

try:
    vent_chart = run_sql_bq(sql("vent_chart_sql"), {"hadms": hadm_list})
except Exception as e:
    print("WARNING: vent chart extraction failed; falling back to ICD-only
        vent timing.", e)
    vent_chart = pd.DataFrame(columns=["hadm_id", "niv_chart_flag",
        "imv_chart_flag", "first_niv_time", "first_imv_time"])

```

```

# Purpose: Derive respiratory support flags and timing fields for IMV/NIV
↪ exposure.

# If your existing ICD-only result is called `vent`, rename for clarity:
vent_proc = vent.copy()

# Outer merge so we keep hadm_ids that appear in only one source
vent_combined = vent_proc.merge(vent_chart, on="hadm_id", how="outer")

# Fill missing with 0 before taking maxima
for c in ["imv_flag","niv_flag","any_vent_flag","imv_chart_flag",]
↪ "niv_chart_flag"]:
    if c in vent_combined.columns:
        vent_combined[c] = vent_combined[c].fillna(0).astype("Int64")

# Final "any-source" flags
vent_combined["imv_flag"] =
↪ vent_combined[["imv_flag","imv_chart_flag"]].max(axis=1).astype("Int64")
vent_combined["niv_flag"] =
↪ vent_combined[["niv_flag","niv_chart_flag"]].max(axis=1).astype("Int64")
vent_combined["any_vent_flag"] =
↪ vent_combined[["imv_flag","niv_flag"]].max(axis=1).astype("Int64")

vent_combined = vent_combined[["hadm_id","imv_flag","niv_flag",]
↪ "any_vent_flag","first_imv_time","first_niv_time"]]
print("After combining ICD + chart signals:",
      "\nIMV=1:", int((vent_combined["imv_flag"]==1).sum()),
      "\nNIV=1:", int((vent_combined["niv_flag"]==1).sum()))

```

After combining ICD + chart signals:
IMV=1: 5150
NIV=1: 4161

2.3.10 10) Assemble final DataFrame

Rationale: Merge all derived features into a single analytic table keyed by hadm_id.

```

# Purpose: Extract and standardize first ABG/VBG physiology fields for
↪ baseline characterization.

# Canonical base (carries authoritative subject_id)
df = demo.copy()

```

```

# Cohort flags / thresholds / labs / etc.
df = safe_merge_on_hadm(df, cohort_any, right_name="cohort_any",
    ↵ check_subject="warn")
df = safe_merge_on_hadm(df, bg_pairs,    right_name="bg_pairs")
df = safe_merge_on_hadm(df, race_eth,    right_name="race_eth")
df = safe_merge_on_hadm(df, ed_triage,   right_name="ed_triage")
df = safe_merge_on_hadm(df, ed_first,    right_name="ed_first")
df = safe_merge_on_hadm(df, icu_meta,    right_name="icu_meta")
df = safe_merge_on_hadm(df, vent_combined,    right_name="vent_combined")

# Anchor to first ED presentation (per admission)
if "ed_intime_first" in globals():
    df = safe_merge_on_hadm(df, ed_intime_first,
    ↵ right_name="ed_intime_first")

# Derived timing: first NIV/IMV relative to ED presentation
if "ed_intime_first" in df.columns and "first_imv_time" in df.columns:
    df["dt_first_imv_hours"] = (df["first_imv_time"] -
    ↵ df["ed_intime_first"]).dt.total_seconds() / 3600.0
if "ed_intime_first" in df.columns and "first_niv_time" in df.columns:
    df["dt_first_niv_hours"] = (df["first_niv_time"] -
    ↵ df["ed_intime_first"]).dt.total_seconds() / 3600.0

# ABG/VBG before IMV (hadm-level)
if {"first_abg_time", "first_imv_time"}.issubset(df.columns):
    df["abg_before_imv"] = (
        df["first_abg_time"].notna() & df["first_imv_time"].notna() &
        (df["first_abg_time"] < df["first_imv_time"]))
    ).astype("Int64")
if {"first_vbg_time", "first_imv_time"}.issubset(df.columns):
    df["vbg_before_imv"] = (
        df["first_vbg_time"].notna() & df["first_imv_time"].notna() &
        (df["first_vbg_time"] < df["first_imv_time"]))
    ).astype("Int64")

# Canonical hadm-level threshold normalization.
# Primary cohort contract: earliest hypercapnic blood gas (LAB/POC) within
    ↵ 24h of ED arrival.
def _num_series(frame: pd.DataFrame, col: str) -> pd.Series:
    if col in frame.columns:
        return pd.to_numeric(frame[col], errors="coerce")
    return pd.Series(np.nan, index=frame.index)

```

```

def _any_ge(frame: pd.DataFrame, cols: list[str], threshold: float) ->
    pd.Series:
    avail = [c for c in cols if c in frame.columns]
    if not avail:
        return pd.Series(0, index=frame.index, dtype="Int64")
    mat = pd.concat([_num_series(frame, c) for c in avail], axis=1)
    return mat.ge(threshold).any(axis=1).astype("Int64")

if "pco2_threshold_0_24h" in df.columns:
    df["pco2_threshold_0_24h"] = (
        pd.to_numeric(df["pco2_threshold_0_24h"], errors="coerce")
        .fillna(0)
        .astype("Int64")
    )
    df["pco2_threshold_any"] = df["pco2_threshold_0_24h"].astype("Int64")
    df["abg_hypcap_threshold"] = (
        pd.to_numeric(df.get("abg_hypcap_threshold"),
        errors="coerce").fillna(0).astype("Int64")
    )
    df["vbg_hypcap_threshold"] = (
        pd.to_numeric(df.get("vbg_hypcap_threshold"),
        errors="coerce").fillna(0).astype("Int64")
    )
    df["unknown_hypcap_threshold"] = (
        pd.to_numeric(df.get("unknown_hypcap_threshold"), errors="coerce")
        .fillna(0)
        .astype("Int64")
    )
else:
    existing_abg = _num_series(df,
        "abg_hypcap_threshold").fillna(0).astype(int)
    existing_vbg = _num_series(df,
        "vbg_hypcap_threshold").fillna(0).astype(int)
    existing_unknown = _num_series(df,
        "unknown_hypcap_threshold").fillna(0).astype(int)

    abg_from_values = _any_ge(df, ["lab_abg_paco2", "poc_abg_paco2",
        "first_abg_pco2"], 45.0)
    vbg_from_values = _any_ge(df, ["lab_vbg_paco2", "poc_vbg_paco2",
        "first_vbg_pco2"], 50.0)

```

```

unknown_from_values = _any_ge(df, ["first_other_pco2"], 50.0)

df["abg_hypercap_threshold"] = ((existing_abg == 1) | (abg_from_values ==
↪ 1)).astype("Int64")
df["vbg_hypercap_threshold"] = ((existing_vbg == 1) | (vbg_from_values ==
↪ 1)).astype("Int64")
df["unknown_hypercap_threshold"] = (
    (existing_unknown == 1) | (unknown_from_values == 1)
).astype("Int64")
df["pco2_threshold_0_24h"] = (
    (df["abg_hypercap_threshold"] == 1)
    | (df["vbg_hypercap_threshold"] == 1)
    | (df["unknown_hypercap_threshold"] == 1)
).astype("Int64")
df["pco2_threshold_any"] = df["pco2_threshold_0_24h"].astype("Int64")

for time_col in [
    "first_abg_hypercap_time_0_24h",
    "first_vbg_hypercap_time_0_24h",
    "first_other_hypercap_time_0_24h",
]:
    if time_col in df.columns:
        df[time_col] = pd.to_datetime(df[time_col], errors="coerce")
for value_col in [
    "max_pco2_0_6h",
    "max_pco2_0_24h",
    "first_abg_hypercap_pco2_mmhg",
    "first_vbg_hypercap_pco2_mmhg",
    "first_other_hypercap_pco2_mmhg",
]:
    if value_col in df.columns:
        df[value_col] = pd.to_numeric(df[value_col], errors="coerce")
for source_col in [
    "max_pco2_0_6h_source_branch",
    "max_pco2_0_24h_source_branch",
    "first_abg_hypercap_source_branch",
    "first_vbg_hypercap_source_branch",
    "first_other_hypercap_source_branch",
]:
    if source_col in df.columns:
        df[source_col] = df[source_col].astype("string").str.upper()

if "any_hypercap_icd" not in df.columns:

```

```

df["any_hypercap_icd"] = 0
df["any_hypercap_icd"] = (
    pd.to_numeric(df["any_hypercap_icd"]),
    errors="coerce").fillna(0).astype("Int64")
)
df["enrolled_any"] = (
    (df["any_hypercap_icd"] == 1)
    | (df["pco2_threshold_0_24h"] == 1)
).astype("Int64")
df["enrolled_any_icd_union_secondary"] = (
    (df["any_hypercap_icd"] == 1)
    | (df["pco2_threshold_0_24h"] == 1)
).astype("Int64")
df["enrollment_route"] = np.select(
    [
        (df["any_hypercap_icd"] == 0) & (df["pco2_threshold_0_24h"] == 1),
        (df["any_hypercap_icd"] == 1) & (df["pco2_threshold_0_24h"] == 0),
        (df["any_hypercap_icd"] == 1) & (df["pco2_threshold_0_24h"] == 1),
    ],
    ["GAS_ONLY", "ICD_ONLY", "ICD+GAS"],
    default="NONE",
)
print(
    "Threshold sanity (hadm-level):",
    "ABG=", int(df["abg_hypercap_threshold"].fillna(0).sum()),
    "VBG=", int(df["vbg_hypercap_threshold"].fillna(0).sum()),
    "UNKNOWN=", int(df["unknown_hypercap_threshold"].fillna(0).sum()),
    "0_24H=", int(df["pco2_threshold_0_24h"].fillna(0).sum()),
    "ICD=", int(df["any_hypercap_icd"].fillna(0).sum()),
    "ENROLLED=", int(df["enrolled_any"].fillna(0).sum()),
)
print("Enrollment route distribution (hadm-level):")
print(df["enrollment_route"].value_counts(dropna=False).to_string())
poc_detail_series_for_audit = (
    df["first_other_src_detail"].astype("string")
    if "first_other_src_detail" in df.columns
    else pd.Series("missing", index=df.index, dtype="string")
)
poc_unknown_hadm_n_for_audit = int(
    (
        pd.to_numeric(df["unknown_hypercap_threshold"], errors="coerce")
        .fillna(0)

```

```

        .astype(int)
        .eq(1)
        & poc_detail_series_for_audit
        .str.lower()
        .isin({"poc_bg_unknown", "poc_bg_unknown_available"})
    ).sum()
)
print("POC UNKNOWN qualifying hadm count (audit-only):",
      poc_unknown_hadm_n_for_audit)

# Harmonize NIH race/ethnicity into a collapsed reporting label at
# cohort-build time.
RACE_NIH_MAP = {
    "WHITE": "WHITE",
    "BLACK OR AFRICAN AMERICAN": "BLACK",
    "ASIAN": "ASIAN",
    "AMERICAN INDIAN OR ALASKA NATIVE": "AI/AN",
    "NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER": "NH/PI",
    "MORE THAN ONE RACE": "MULTIRACIAL",
    "UNKNOWN OR NOT REPORTED": "UNKNOWN",
    "OTHER": "OTHER",
    "UNKNOWN": "UNKNOWN",
}
ETH_NIH_MAP = {
    "HISPANIC OR LATINO": "HISPANIC",
    "NOT HISPANIC OR LATINO": "NOT HISPANIC",
    "UNKNOWN OR NOT REPORTED": "UNKNOWN",
    "UNKNOWN": "UNKNOWN",
}
LABEL_MAP = {
    "WHITE": "Non-Hispanic White",
    "BLACK": "Non-Hispanic Black",
    "ASIAN": "Non-Hispanic Asian",
    "AI/AN": "Non-Hispanic American Indian/Alaska Native",
    "NH/PI": "Non-Hispanic Native Hawaiian/Pacific Islander",
    "MULTIRACIAL": "Non-Hispanic Multiracial/Other",
    "OTHER": "Unknown/Other",
    "UNKNOWN": "Unknown/Other",
}

```

```

def addCollapsedRaceEth(
    frame: pd.DataFrame,
    race_col: str = "nih_race",
    eth_col: str = "nih_ethnicity",
    out_col: str = "race_eth_collapsed",
) -> pd.DataFrame:
    """Collapse NIH race/ethnicity fields into publication-friendly
    categories."""
    if race_col not in frame.columns or eth_col not in frame.columns:
        return frame

    def _std(value):
        if pd.isna(value):
            return None
        return re.sub(r"\s+", " ", str(value).strip()).upper()

    race_std = frame[race_col].map(lambda value:
        RACE_NIH_MAP.get(_std(value), "OTHER"))
    eth_std = frame[eth_col].map(lambda value: ETH_NIH_MAP.get(_std(value),
        "UNKNOWN"))

    collapsed = []
    for race_value, eth_value in zip(race_std, eth_std):
        if eth_value == "HISPANIC":
            collapsed.append("Hispanic/Latino")
        else:
            collapsed.append(LABEL_MAP.get(race_value, "Unknown/Other"))

    category_order = [
        "Hispanic/Latino",
        "Non-Hispanic White",
        "Non-Hispanic Black",
        "Non-Hispanic Asian",
        "Non-Hispanic American Indian/Alaska Native",
        "Non-Hispanic Native Hawaiian/Pacific Islander",
        "Non-Hispanic Multiracial/Other",
        "Unknown/Other",
    ]

    frame[out_col] = pd.Categorical(collapsed, categories=category_order,
        ordered=True)
    frame["nih_race_std"] = race_std
    frame["nih_ethnicity_std"] = eth_std

```

```

    return frame

df = add_collapsed_race_eth(df)

print("Final df rows:", len(df), "cols:", len(df.columns))

# Safety checks
assert "subject_id" in df.columns, "subject_id missing from final df"
assert not any(c.endswith("_x") or c.endswith("_y") for c in df.columns),
    "Found suffixed columns"
print("Final df rows:", len(df), "cols:", len(df.columns))
df.head(3)

```

```

Threshold sanity (hadm-level): ABG= 3512 VBG= 3507 UNKNOWN= 396 O_24H= 6246 ICD= 4237 ENROLL
Enrollment route distribution (hadm-level):
enrollment_route
GAS_ONLY      5169
ICD_ONLY      3160
ICD+GAS       1077
POC UNKNOWN qualifying hadm count (audit-only): 0
Final df rows: 9406 cols: 143
Final df rows: 9406 cols: 143

```

	hadm_id	subject_id	admittime	dischtime	deathtime	administrative_class	fr
0	22447711	13352386	2188-03-25 01:17:00	2188-03-26 16:00:00	NaT	EW EMER.	EM
1	20930337	18079481	2152-05-13 20:36:00	2152-05-17 11:00:00	NaT	EW EMER.	EM
2	22834654	11482582	2125-10-30 01:03:00	2125-11-06 18:35:00	NaT	EW EMER.	EM

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
#           consistent across notebook stages.

# --- Cohort flow counts (ED / ICU / blood gas / hypercapnia / CC) ---

# 1) Dataset-level ED counts
SQL["ed_counts_sql"] = f"""
SELECT
    COUNT(*) AS total_ed_encounters,
    COUNTIF(hadm_id IS NOT NULL) AS ed_encounters_with_hadm
FROM `'{PHYS}`.{ED}.edstays`
"""

```

```

SQL["ed_to_icu_sql"] = f"""
SELECT
    COUNT(DISTINCT e.hadm_id) AS ed_to_icu_hadm,
    COUNT(DISTINCT e.stay_id) AS ed_to_icu_edstays
FROM `{{PHYS}}.{{ED}}.edstays` e
JOIN `{{PHYS}}.{{ICU}}.icustays` i USING (hadm_id)
WHERE e.hadm_id IS NOT NULL
"""

try:
    ed_counts = run_sql_bq(sql("ed_counts_sql"))
    ed_to_icu = run_sql_bq(sql("ed_to_icu_sql"))
except Exception as e:
    print("Warning: ED/ICU counts query failed:", e)
    ed_counts = None
    ed_to_icu = None

# 2) Cohort-level counts (admission-level)
cohort_union = int((cohort_any["enrolled_any"] == 1).sum()) if "cohort_any"
    ↵ in globals() and "enrolled_any" in cohort_any.columns else len(hadm_list)
cohort_df_n = len(df)

# Any blood gas present (ABG/VBG, LAB/POC)
co2_cols = [c for c in [
    "lab_abg_paco2", "lab_vbg_paco2", "poc_abg_paco2", "poc_vbg_paco2"
] if c in df.columns]
any_bg = int(df[co2_cols].notna().any(axis=1).sum()) if co2_cols else None

# Hypercapnia thresholds and ICD
icd_count = int((df["any_hypercap_icd"] == 1).sum()) if "any_hypercap_icd"
    ↵ in df.columns else None
threshold_count = int((df["pco2_threshold_any"] == 1).sum()) if
    ↵ "pco2_threshold_any" in df.columns else None

# ED chief complaint missing / present (within cohort)
if "ed_triage_cc" in df.columns:
    mask_cc_present = df["ed_triage_cc"].notna() &
        ↵ (df["ed_triage_cc"].astype(str).str.strip() != "")
    cc_present = int(mask_cc_present.sum())
    cc_missing = int(~mask_cc_present.sum())
else:
    cc_present = None

```

```

cc_missing = None

# ED→ICU within cohort (admissions with ED triage data and ICU stay)
if "first_icu_stay_id" in df.columns and "ed_triage_cc" in df.columns:
    cohort_ed_to_icu = int((df["first_icu_stay_id"].notna() &
    ↵ mask_cc_present).sum())
else:
    cohort_ed_to_icu = None

rows = []
if ed_counts is not None:
    rows.append({"step": "Total ED encounters (edstays)", "count":
    ↵ int(ed_counts.loc[0, "total_ed_encounters"])), "scope": "All ED dataset"})
    rows.append({"step": "ED encounters with hadm_id", "count":
    ↵ int(ed_counts.loc[0, "ed_encounters_with_hadm"])), "scope": "All ED
    ↵ dataset"})
if ed_to_icu is not None:
    rows.append({"step": "ED→ICU admissions (distinct hadm_id)", "count":
    ↵ int(ed_to_icu.loc[0, "ed_to_icu_hadm"])), "scope": "All ED+ICU"})
    rows.append({"step": "ED→ICU ED-stays (distinct stay_id)", "count":
    ↵ int(ed_to_icu.loc[0, "ed_to_icu_edstays"])), "scope": "All ED+ICU"})

rows.append({"step": "Cohort admissions (union ICD thresholds)", "count":
    ↵ cohort_union, "scope": "Cohort"})
rows.append({"step": "Cohort admissions after merges (df rows)", "count":
    ↵ cohort_df_n, "scope": "Cohort"})
if any_bg is not None:
    rows.append({"step": "Cohort with any ABG/VBG (LAB or POC)", "count":
    ↵ any_bg, "scope": "Cohort"})
if threshold_count is not None:
    rows.append({"step": "Cohort meeting hypercapnia thresholds", "count":
    ↵ threshold_count, "scope": "Cohort"})
if icd_count is not None:
    rows.append({"step": "Cohort meeting ICD code criteria", "count":
    ↵ icd_count, "scope": "Cohort"})
if cc_present is not None:
    rows.append({"step": "Cohort with ED chief complaint present", "count":
    ↵ cc_present, "scope": "Cohort"})
if cc_missing is not None:
    rows.append({"step": "Cohort excluded for missing ED chief complaint",
    ↵ "count": cc_missing, "scope": "Cohort"})
if cohort_ed_to_icu is not None:

```

```

    rows.append({"step": "Cohort ED→ICU (ED CC present + ICU stay)", "count": 
    ↵ cohort_ed_to_icu, "scope": "Cohort"})

flow_counts = pd.DataFrame(rows)
flow_counts

```

step	count	scope
0 Total ED encounters (edstays)	425087	All ED dataset
1 ED encounters with hadm_id	203016	All ED dataset
2 ED→ICU admissions (distinct hadm_id)	31862	All ED+ICU
3 ED→ICU ED-stays (distinct stay_id)	31916	All ED+ICU
4 Cohort admissions (union ICD thresholds)	9406	Cohort
5 Cohort admissions after merges (df rows)	9406	Cohort
6 Cohort with any ABG/VBG (LAB or POC)	8681	Cohort
7 Cohort meeting hypercapnia thresholds	6246	Cohort
8 Cohort meeting ICD code criteria	4237	Cohort
9 Cohort with ED chief complaint present	7152	Cohort
10 Cohort excluded for missing ED chief complaint	2254	Cohort
11 Cohort ED→ICU (ED CC present + ICU stay)	5674	Cohort

```

# Purpose: Quantify overlap between ascertainment routes so non-exclusive
↪ cohorts are explicit.

# --- Ascertainment overlap counts (ABG/VBG/ICD) ---

required = ["abg_hypercap_threshold", "vbg_hypercap_threshold",
↪ "unknown_hypercap_threshold", "any_hypercap_icd"]
missing = [c for c in required if c not in df.columns]
if missing:
    raise KeyError(f"Missing required columns for overlap counts: {missing}")

abg = pd.to_numeric(df["abg_hypercap_threshold"],
↪ errors="coerce").fillna(0).astype(int)
vbg = pd.to_numeric(df["vbg_hypercap_threshold"],
↪ errors="coerce").fillna(0).astype(int)
unknown = pd.to_numeric(df["unknown_hypercap_threshold"],
↪ errors="coerce").fillna(0).astype(int)

gas_any = (
    pd.to_numeric(df.get("pco2_threshold_any", None), errors="coerce")
    if "pco2_threshold_any" in df.columns else (abg | vbg | unknown)
)

```

```

)
if hasattr(gas_any, "fillna"):
    gas_any = gas_any.fillna(0).astype(int)
else:
    gas_any = gas_any.astype(int)

icd = pd.to_numeric(df["any_hypercap_icd"],
    ↪ errors="coerce").fillna(0).astype(int)

total_n = len(df)
ngas = int((gas_any == 1).sum())

abg_vbg_overlap = pd.DataFrame([
    {"group": "ABG-only", "count": int(((abg == 1) & (vbg == 0) & (unknown == 0)).sum())},
    {"group": "V ро-only", "count": int(((vbg == 1) & (abg == 0) & (unknown == 0)).sum())},
    {"group": "UNKNOWN-only", "count": int(((unknown == 1) & (abg == 0) & (vbg == 0)).sum())},
    {"group": "Mixed (>=2 routes)", "count": int(((abg + vbg + unknown) >= 2).sum())},
])
if ngas > 0:
    abg_vbg_overlap["pct_of_gas"] = (abg_vbg_overlap["count"] / ngas *
    ↪ * 100).round(1)
else:
    abg_vbg_overlap["pct_of_gas"] = 0.0
abg_vbg_overlap["pct_of_cohort"] = (abg_vbg_overlap["count"] / max(total_n, 1) *
    ↪ * 100).round(1)

icd_gas_overlap = pd.DataFrame([
    {"group": "ICD+Gas", "count": int(((icd==1) & (gas_any==1)).sum())},
    {"group": "ICD-only", "count": int(((icd==1) & (gas_any==0)).sum())},
    {"group": "Gas-only", "count": int(((icd==0) & (gas_any==1)).sum())},
    {"group": "Neither", "count": int(((icd==0) & (gas_any==0)).sum())},
])
icd_gas_overlap["pct_of_cohort"] = (icd_gas_overlap["count"] / max(total_n, 1) *
    ↪ * 100).round(1)

print("ABG/VBG/Other overlap (among gas-positive):")
print(abg_vbg_overlap.to_string(index=False))
print("ICD vs Gas overlap (cohort-level):")
print(icd_gas_overlap.to_string(index=False))

```

```

ABG/VBG/Other overlap (among gas-positive):
    group  count  pct_of_gas  pct_of_cohort
    ABG-only    2463        39.4        26.2
    VBG-only    2512        40.2        26.7
    UNKNOWN-only   158        2.5        1.7
    Mixed (>=2 routes) 1113        17.8        11.8

ICD vs Gas overlap (cohort-level):
    group  count  pct_of_cohort
    ICD+Gas    1077        11.5
    ICD-only    3160        33.6
    Gas-only    5169        55.0
    Neither      0         0.0

# Purpose: Capture ED presentation context (chief complaint, acuity, and
#           ↵ early vitals) at the ED-stay level.

# --- Missingness summary (chief complaint, race/ethnicity, ED triage/vitals)
#           ↵ ---

import numpy as np

summary_rows = []

# Chief complaint missingness (ED triage CC)
if "ed_triage_cc" in df.columns:
    cc_present = df["ed_triage_cc"].notna() &
    ↵ (df["ed_triage_cc"].astype(str).str.strip() != "")
    summary_rows.append({
        "variable": "ed_triage_cc_present",
        "missing_n": int(~cc_present).sum(),
        "missing_pct": float(~cc_present).mean()
    })

# Race/Ethnicity missingness (NIH categories + raw sources)
unknown_tokens = {
    "unknown or not reported",
    "unknown",
    "not reported",
    "missing",
    "declined",
    "unable"
}

```

```

def _missing_rate(series):
    if series is None:
        return None, None
    s = series.astype(str).str.strip()
    is_missing = series.isna() | (s == "") |
    ↵ s.str.lower().isin(unknown_tokens)
    return int(is_missing.sum()), float(is_missing.mean())

for col in ["nih_race", "nih_ethnicity", "race_hosp_raw", "race_ed_raw"]:
    if col in df.columns:
        m_n, m_p = _missing_rate(df[col])
        summary_rows.append({
            "variable": col,
            "missing_n": m_n,
            "missing_pct": m_p
        })

missing_summary = pd.DataFrame(summary_rows)
print("Missingness summary (key variables):")
missing_summary

# ED triage + first ED vitals missingness
triage_cols = [c for c in df.columns if c.startswith("ed_triage_")]
first_cols = [c for c in df.columns if c.startswith("ed_first_")]

vital_cols = triage_cols + first_cols
if vital_cols:
    miss_tbl = (
        pd.DataFrame({"variable": vital_cols})
        .assign(
            missing_n=lambda d: [int(df[c].isna().sum()) for c in
        ↵ d["variable"]],
            missing_pct=lambda d: [float(df[c].isna().mean()) for c in
        ↵ d["variable"]]
        )
        .sort_values("missing_pct", ascending=False)
    )
    print("Missingness summary (ED triage + first ED vitals):")
    miss_tbl
else:
    miss_tbl = pd.DataFrame(columns=["variable", "missing_n", "missing_pct"])

Missingness summary (key variables):

```

Missingness summary (ED triage + first ED vitals):

2.3.11 11) Sanity checks

Rationale: Run QC checks to validate units, flags, and basic data integrity before export.

```
# Purpose: Derive respiratory support flags and timing fields for IMV/NIV
↪ exposure.

# ICU LOS negative?
if {"icu_los_days","first_icu_stay_id"}.issubset(df.columns):
    neg_los = int((df["icu_los_days"] < 0).fillna(False).sum())
    print("Negative ICU LOS rows:", neg_los)

# Vent flags consistency
vent_cols = {"imv_flag","niv_flag","any_vent_flag"}
if vent_cols.issubset(df.columns):
    any_calc = ((df["imv_flag"]==1) |
↪ (df["niv_flag"]==1)).fillna(False).astype(int)
    any_flag = pd.to_numeric(df["any_vent_flag"],
↪ errors="coerce").fillna(0).astype(int)
    mism = int((any_calc != any_flag).sum())
    print("any_vent_flag mismatches vs (imv|niv):", mism)

# UOMs: expect mmhg only
uom_cols = [c for c in df.columns if c.endswith("_paco2_uom")]
for c in uom_cols:
    vals = sorted(pd.Series(df[c]).dropna().astype(str).str.lower().str_
↪ .strip().unique().tolist())
    print(c, vals)

# ABG/VBG coverage QC
def qc_pair(df, ph_col, co2_col, label, ph_lo=6.3, ph_hi=7.8, co2_lo=5,
↪ co2_hi=200):
    ph = pd.to_numeric(df.get(ph_col), errors="coerce")
    co2 = pd.to_numeric(df.get(co2_col), errors="coerce")
    return {
        "pair": label,
        "present_any": int(((ph.notna()) | (co2.notna())).sum()),
        "present_both": int(((ph.notna()) & (co2.notna())).sum()),
        "only_ph": int(((ph.notna()) & (~co2.notna())).sum()),
        "only_pco2": int(((co2.notna()) & (~ph.notna())).sum()),
```

```

    "ph_oob": int(((ph < ph_lo) | (ph > ph_hi)) &
    ↵ ph.notna().sum()),
    "pco2_oob": int(((co2 < co2_lo) | (co2 > co2_hi)) &
    ↵ co2.notna().sum()),
}

qc = pd.DataFrame([
    qc_pair(df, "lab_abg_ph", "lab_abg_paco2", "LAB ABG"),
    qc_pair(df, "lab_vbg_ph", "lab_vbg_paco2", "LAB VBG"),
    qc_pair(df, "poc_abg_ph", "poc_abg_paco2", "POC ABG"),
    qc_pair(df, "poc_vbg_ph", "poc_vbg_paco2", "POC VBG"),
])
qc

Negative ICU LOS rows: 0
any_vent_flag mismatches vs (imv|niv): 0
lab_abg_paco2_uom ['mmhg']
lab_vbg_paco2_uom ['mmhg']
lab_other_paco2_uom ['mmhg']
poc_abg_paco2_uom ['mmhg']
poc_vbg_paco2_uom ['mmhg']
poc_other_paco2_uom []

```

	pair	present_any	present_both	only_ph	only_pco2	ph_oob	pco2_oob
0	LAB ABG	5758	5756	0	2	0	0
1	LAB VBG	6366	6365	0	1	0	0
2	POC ABG	5152	5151	0	1	0	0
3	POC VBG	4836	4836	0	0	0	0

2.3.12 PDF-ready long tables

```

print(
    render_latex_longtable(
        flow_counts,
        caption=f"Cohort flow and key denominator checkpoints
        ↵ (N={len(df)}:,} .",
        label="tab:cohort_flow_counts",
        index=False,
    )
)

```

```

print(
    render_latex_longtable(
        abg_vbg_overlap,
        caption="ABG/VBG/OTHER overlap among gas-positive encounters.",
        label="tab:gas_overlap_abg_vbg_other",
        index=False,
    )
)
print(
    render_latex_longtable(
        icd_gas_overlap,
        caption="ICD versus gas ascertainment overlap in the full cohort.",
        label="tab:icd_gas_overlap",
        index=False,
    )
)
print(
    render_latex_longtable(
        missing_summary,
        caption="Missingness summary for key cohort variables.",
        label="tab:key_missingness",
        index=False,
    )
)
if "miss_tbl" in globals() and not miss_tbl.empty:
    print(
        render_latex_longtable(
            miss_tbl,
            caption="ED triage and first ED vital-sign missingness summary.",
            label="tab:ed_vital_missingness",
            landscape=True,
            index=False,
        )
    )
print(
    render_latex_longtable(
        qc,
        caption="ABG/VBG pH and pCO2 plausibility quality checks.",
        label="tab:abg_vbg_qc",
        landscape=True,
        index=False,
    )
)

```

```

\begin{longtable}{lrl}
\caption{Cohort flow and key denominator checkpoints (N=9,406).} \label{tab:cohort_flow_count}
\toprule
step & count & scope \\
\midrule
\endfirsthead
\caption[] {Cohort flow and key denominator checkpoints (N=9,406).} \\
\toprule
step & count & scope \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
Total ED encounters (edstays) & 425087 & All ED dataset \\
ED encounters with hadm_id & 203016 & All ED dataset \\
ED→ICU admissions (distinct hadm_id) & 31862 & All ED+ICU \\
ED→ICU ED-stays (distinct stay_id) & 31916 & All ED+ICU \\
Cohort admissions (union ICD thresholds) & 9406 & Cohort \\
Cohort admissions after merges (df rows) & 9406 & Cohort \\
Cohort with any ABG/VBG (LAB or POC) & 8681 & Cohort \\
Cohort meeting hypercapnia thresholds & 6246 & Cohort \\
Cohort meeting ICD code criteria & 4237 & Cohort \\
Cohort with ED chief complaint present & 7152 & Cohort \\
Cohort excluded for missing ED chief complaint & 2254 & Cohort \\
Cohort ED→ICU (ED CC present + ICU stay) & 5674 & Cohort \\
\end{longtable}

\begin{longtable}{lrrrr}
\caption{ABG/VBG/OTHER overlap among gas-positive encounters.} \label{tab:gas_overlap_abg_vbg}
\toprule
group & count & pct_of_gas & pct_of_cohort \\
\midrule
\endfirsthead
\caption[] {ABG/VBG/OTHER overlap among gas-positive encounters.} \\
\toprule
group & count & pct_of_gas & pct_of_cohort \\
\midrule
\endhead
\midrule

```

```

\multicolumn{4}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ABG-only & 2463 & 39.400000 & 26.200000 \\
VBG-only & 2512 & 40.200000 & 26.700000 \\
UNKNOWN-only & 158 & 2.500000 & 1.700000 \\
Mixed (>=2 routes) & 1113 & 17.800000 & 11.800000 \\
\end{longtable}

\begin{longtable}{lrr}
\caption{ICD versus gas ascertainment overlap in the full cohort.} \label{tab:icd_gas_overlap}
\toprule
group & count & pct_of_cohort \\
\midrule
\endfirsthead
\caption[] {ICD versus gas ascertainment overlap in the full cohort.} \\
\toprule
group & count & pct_of_cohort \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ICD+Gas & 1077 & 11.500000 \\
ICD-only & 3160 & 33.600000 \\
Gas-only & 5169 & 55.000000 \\
Neither & 0 & 0.000000 \\
\end{longtable}

\begin{longtable}{lrr}
\caption{Missingness summary for key cohort variables.} \label{tab:key_missingness} \\
\toprule
variable & missing_n & missing_pct \\
\midrule
\endfirsthead
\caption[] {Missingness summary for key cohort variables.} \\
\toprule
variable & missing_n & missing_pct \\
\endhead

```

```

\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ed_triage_cc_present & 2254 & 0.239634 \\
nih_race & 1951 & 0.207421 \\
nih_ethnicity & 1818 & 0.193281 \\
race_hosp_raw & 1033 & 0.109824 \\
race_ed_raw & 3192 & 0.339358 \\
\end{longtable}

\begin{landscape}\begin{longtable}{lrr}
\caption{ED triage and first ED vital-sign missingness summary.} \label{tab:ed_vital_missing}
\toprule
variable & missing_n & missing_pct \\
\midrule
\endfirsthead
\caption[]{ED triage and first ED vital-sign missingness summary.} \\
\toprule
variable & missing_n & missing_pct \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ed_first_rhythm & 9028 & 0.959813 \\
ed_first_temp & 5643 & 0.599936 \\
ed_triage_temp & 4951 & 0.526366 \\
ed_triage_rr & 4723 & 0.502126 \\
ed_triage_o2sat & 4704 & 0.500106 \\
ed_triage_dbp & 4699 & 0.499575 \\
ed_triage_sbp & 4666 & 0.496066 \\
ed_triage_hr & 4635 & 0.492771 \\
ed_first_o2sat & 4387 & 0.466404 \\
ed_first_dbp & 4285 & 0.455560 \\
ed_first_sbp & 4285 & 0.455560 \\

```

```

ed_first_rr & 4218 & 0.448437 \\
ed_first_hr & 4146 & 0.440782 \\
ed_first_pain & 3800 & 0.403997 \\
ed_triage_pain & 3653 & 0.388369 \\
ed_triage_acuity & 3311 & 0.352009 \\
ed_first_vitals_time & 3059 & 0.325218 \\
ed_triage_cc & 2254 & 0.239634 \\
\end{longtable}
\n\end{landscape}\n
\begin{landscape}\n\begin{longtable}{lrrrrrrr}
\caption{ABG/VBG pH and pCO2 plausibility quality checks.} \label{tab:abg_vbg_qc} \\
\toprule
pair & present_any & present_both & only_ph & only_pco2 & ph_oob & pco2_oob \\
\midrule
\endfirsthead
\caption[] {ABG/VBG pH and pCO2 plausibility quality checks.} \\
\toprule
pair & present_any & present_both & only_ph & only_pco2 & ph_oob & pco2_oob \\
\midrule
\endhead
\midrule
\multicolumn{7}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
LAB ABG & 5758 & 5756 & 0 & 2 & 0 & 0 \\
LAB VBG & 6366 & 6365 & 0 & 1 & 0 & 0 \\
POC ABG & 5152 & 5151 & 0 & 1 & 0 & 0 \\
POC VBG & 4836 & 4836 & 0 & 0 & 0 & 0 \\
\end{longtable}
\n\end{landscape}\n

```

2.3.13 12) Save to Excel

Rationale: Persist cohort outputs for downstream annotation and NLP analyses.

Purpose: Optional archive export of the full hadm-level cohort workbook.

```
from datetime import datetime
```

```

WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
    "1"

if WRITE_ARCHIVE_XLSX_EXPORTS:
    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    out_path = prior_runs_dir / f"mimic_hypercap_EXT_bq_abg_vbg_{datetime.]
    ↵ now().strftime('%Y%m%d_%H%M%S')}.xlsx"
    with pd.ExcelWriter(out_path, engine="openpyxl") as xw:
        df.to_excel(xw, sheet_name="cohort", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    out_path
else:
    print("Skipping archive hadm-level workbook export (set
    ↵ WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")

```

Skipping archive hadm-level workbook export (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

2.3.14 Create Annotation Dataset

Rationale: Create ED chief-complaint subsets and a reproducible sample for manual annotation.

```

# Purpose: Capture ED presentation context (chief complaint, acuity, and
    ↵ early vitals) at the ED-stay level.

# ---- Extra exports: (1) ED chief-complaint only; (2) random sample of 160
    ↵ patients ----
from datetime import datetime

# Dated artifacts go to archive folder (optional)
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')
WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
    "1"

# 1) Filter to rows with a non-empty ED chief complaint
if "ed_triage_cc" not in df.columns:

```

```

    raise KeyError(
        "Column 'ed_triage_cc' not found in df. "
        "Ensure the ED triage merge cell ran earlier."
    )

mask_cc = df["ed_triage_cc"].notna() &
    ~ (df["ed_triage_cc"].astype(str).str.strip() != "")
df_cc = df.loc[mask_cc].copy()

print(f"ED-CC present rows: {len(df_cc)} of {len(df)} "
      f"({(len(df_cc) / max(len(df), 1)):.1%} of cohort.)")

# Save ED-CC-only cohort only when archive exports are enabled
if WRITE_ARCHIVE_XLSX_EXPORTS:
    out_path_cc = prior_runs_dir /
    f"mimic_hypcap_EXT_EDcc_only_bq_abg_vbg_{timestamp}.xlsx"
    with pd.ExcelWriter(out_path_cc, engine="openpyxl") as xw:
        df_cc.to_excel(xw, sheet_name="cohort_cc_only", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    print("Saved:", out_path_cc)

# 2) Random sample of n = 160 patients (distinct subject_id), one row per
#    patient
if "subject_id" not in df_cc.columns:
    raise KeyError("Column 'subject_id' missing; cannot sample by patient.")

# Make a one-row-per-patient frame by earliest admission
if "admittime" in df_cc.columns:
    df_cc_one = (
        df_cc.sort_values(["subject_id", "admittime"])
            .groupby("subject_id", as_index=False)
            .head(1)
    )
else:
    # Fallback if admittime not present: choose the smallest hadm_id per
    #    patient
    df_cc_one = (
        df_cc.sort_values(["subject_id", "hadm_id"])
            .groupby("subject_id", as_index=False)
            .head(1)

```

```

)
N = 160
n_avail = len(df_cc_one)
n_take = min(N, n_avail)
if n_avail < N:
    print(f"Warning: only {n_avail} unique patients with ED chief complaint;
        sampling all of them.")

RANDOM_SEED = 42
df_cc_sample = df_cc_one.sample(n=n_take, random_state=RANDOM_SEED)

# Save the sample only when archive exports are enabled
if WRITE_ARCHIVE_XLSX_EXPORTS:
    out_path_cc_sample = prior_runs_dir /
    f"mimic_hypcap_EXT_EDcc_sample{n_take}_bq_abg_vbg_{timestamp}.xlsx"
    with pd.ExcelWriter(out_path_cc_sample, engine="openpyxl") as xw:
        df_cc_sample.to_excel(xw, sheet_name="cohort_cc_sample", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    print("Saved:", out_path_cc_sample)
else:
    print("Skipping ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1
        to enable).")

```

ED-CC present rows: 7152 of 9406 (76.0% of cohort).
 Skipping ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

3 ED-stay cohort expansion (timing, severity, comorbidity, outcomes)

Rationale: Build a one-row-per-ED-stay analytic extract with time-anchored gas phenotypes, key comorbidities, and outcomes.

3.0.1 Phase 0 — Inventory & missing-field registry

Rationale: Detect which fields already exist and only add missing fields to avoid redundant extraction or join explosions.

```

# Purpose: Inventory available fields and identify missing target variables
    ↵ before enrichment steps.

from pathlib import Path
import json
import pandas as pd

# Identify current cohort dataframe (admission-level)
if 'df' not in globals():
    raise NameError("Expected admission-level df to exist before inventory
        ↵ step.")

ED_KEY = "ed_stay_id" # target key for ED-level cohort

print("Current admission-level df:", df.shape)
print("Current columns count:", len(df.columns))
if ED_KEY in df.columns:
    print("ED stay unique count:", int(df[ED_KEY].nunique()))
else:
    print("ED stay unique count: ED_KEY not in columns")

# Persist columns snapshot
cols_out = WORK_DIR / "current_columns.json"
cols_out.write_text(json.dumps(sorted(df.columns), indent=2))
print("Wrote:", cols_out)

```

```

# Persist ED-stay columns snapshot (if ed_df exists)
if "ed_df" in globals():
    ed_cols_out = WORK_DIR / "ed_columns.json"
    ed_cols_out.write_text(json.dumps(sorted(ed_df.columns), indent=2))
    print("Wrote:", ed_cols_out)

Current admission-level df: (9406, 143)
Current columns count: 143
ED stay unique count: ED_KEY not in columns
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/current_columns.json

```

```

# Purpose: Capture ED presentation context (chief complaint, acuity, and
    ↵ early vitals) at the ED-stay level.

```

```

# Target field registry
TARGET_RAW_FIELDS = {
    "ed_edstays": [
        "ed_stay_id", "subject_id", "hadm_id", "ed_intime", "ed_outtime", ,
        ↵ "ed_intime_first",
        "arrival_transport", "disposition", "ed_gender", "ed_race",
    ],
    "ed_triage": [
        "ed_triage_temp", "ed_triage_hr", "ed_triage_rr", "ed_triage_o2sat",
        "ed_triage_sbp", "ed_triage_dbp", "ed_triage_pain", "ed_triage_acuity", ,
        ↵ "ed_triage_cc",
    ],
    "ed_vitals_first": [
        "ed_first_vitals_time", "ed_first_temp", "ed_first_hr", "ed_first_rr",
        "ed_first_o2sat", "ed_first_sbp", "ed_first_dbp", "ed_first_rhythm", ,
        ↵ "ed_first_pain",
    ],
    "admissions": [
        "admittime", "dischtime", "deathtime", "hospital_expire_flag",
        "administrative_class", "from_location", "discharge_location",
        "insurance", "language", "marital_status", "hosp_race",
    ],
    "icu": [
        "icu_stay_id", "icu_intime_first", "icu_outtime_last", "icu_los_total", ,
        ↵ "n_icu_stays",
        "first_careunit", "last_careunit",
    ],
    "labs_gas": [
        "first_gas_time", "first_pco2", "first_ph", "first_hco3", ,
        ↵ "first_lactate",
        "qualifying_pco2_time", "qualifying_pco2_mmhg", "qualifying_site", ,
        ↵ "qualifying_source_branch", "qualifying_threshold_mmhg",
        "first_abg_po2", "first_vbg_po2", "first_other_po2",
        "first_abg_hypercap_time_0_24h", "first_abg_hypercap_pco2_mmhg", ,
        ↵ "first_abg_hypercap_source_branch", "first_abg_hypercap_itemid",
        "first_vbg_hypercap_time_0_24h", "first_vbg_hypercap_pco2_mmhg", ,
        ↵ "first_vbg_hypercap_source_branch", "first_vbg_hypercap_itemid",
        "first_other_hypercap_time_0_24h", "first_other_hypercap_pco2_mmhg", ,
        ↵ "first_other_hypercap_source_branch",
        ↵ "first_other_hypercap_itemid",
        "pco2_threshold_0_24h",
        "max_pco2_0_6h", "min_ph_0_6h", "max_pco2_0_24h", "min_ph_0_24h",
        "gas_source_unknown_rate", "gas_source_other_rate",
    ]
}

```

```

    "first_gas_specimen_type", "first_gas_specimen_present", ],
    ↵ "first_gas_pco2_itemid", "first_gas_pco2_fluid",
    "co2_other_is_blood_asserted", "first_gas_anchor_has_pco2", ],
    ↵ "first_gas_anchor_source_validated",
    "dt_first_imv_hours", "dt_first_niv_hours", "first_other_time",
],
"omr": [
    "bmi_closest_pre_ed", "height_closest_pre_ed", "weight_closest_pre_ed",
    "bmi_closest_pre_ed_uom", "height_closest_pre_ed_uom", ],
    ↵ "weight_closest_pre_ed_uom",
    "bmi_closest_pre_ed_time", "height_closest_pre_ed_time", ],
    ↵ "weight_closest_pre_ed_time",
    "bmi_closest_pre_ed_model", "height_closest_pre_ed_model", ],
    ↵ "weight_closest_pre_ed_model",
    "bmi_outlier_flag", "height_outlier_flag", "weight_outlier_flag",
    "anthro_timing_tier", "anthro_days_offset", "anthro_chartdate", ],
    ↵ "anthro_timing_uncertain",
    "anthro_source", "anthro_obstime", "anthro_hours_offset", ],
    ↵ "anthro_timing_basis",
],
"dx_flags": [
    "flag_copd", "flag_osa_ohs", "flag_chf", "flag_neuromuscular",
    "flag_opioid_substance", "flag_pneumonia",
],
"timing": [
    "dt_qualifying_hypercapnia_hours", "dt_first_qualifying_gas_hours", ],
    ↵ "presenting_hypercapnia", "late_hypercapnia",
    "hypercap_timing_class",
    "enrollment_route",
    "dt_first_imv_hours", "dt_first_niv_hours", "abg_before_imv", ],
    ↵ "vbg_before_imv",
    "dt_first_imv_hours_model", "dt_first_niv_hours_model",
    "imv_time_outside_window_flag", "niv_time_outside_window_flag",
    "hospital_los_negative_flag", "admittime_before_ed_intime_flag", ],
    ↵ "dischtime_before_admittime_flag", "time_integrity_any", ],
    ↵ "timing_usable_for_model",
    "ph_band", "hco3_band", "lactate_band",
],
},
TARGET_DERIVED_FIELDS = [
    "hospital_los_hours", "hospital_los_hours_model", "in_hospital_death",
]

```

```

TARGET_FIELDS = sorted({c for v in TARGET_RAW_FIELDS.values() for c in v} |
    set(TARGET_DERIVED_FIELDS))

# Grouped missing report
missing_by_group = {}
for group, cols in TARGET_RAW_FIELDS.items():
    missing_by_group[group] = [c for c in cols if c not in df.columns]

missing_derived = [c for c in TARGET_DERIVED_FIELDS if c not in df.columns]
missing_by_group["derived"] = missing_derived

missing_flat = [c for cols in missing_by_group.values() for c in cols]
print("Missing fields total:", len(missing_flat))
for group, cols in missing_by_group.items():
    if cols:
        print(f"- {group}: {cols}")

Missing fields total: 87
- ed_edstays: ['ed_stay_id', 'ed_intime', 'ed_outtime', 'ed_intime_first', 'arrival_transport'
- admissions: ['hospital_expire_flag', 'language', 'marital_status', 'hosp_race']
- icu: ['icu_stay_id', 'icu_intime_first', 'icu_outtime_last', 'icu_los_total', 'n_icu_stays'
- labs_gas: ['first_gas_time', 'first_pco2', 'first_ph', 'first_hco3', 'first_lactate', 'min_
- omr: ['bmi_closest_pre_ed', 'height_closest_pre_ed', 'weight_closest_pre_ed', 'bmi_closest_
- dx_flags: ['flag_copd', 'flag.osa.ohs', 'flag_chf', 'flag_neuromuscular', 'flag_opioid_sub_
- timing: ['dt_first_qualifying_gas_hours', 'presenting_hypercapnia', 'late_hypercapnia', 'hy_
- derived: ['hospital_los_hours', 'hospital_los_hours_model', 'in_hospital_death']

```

3.0.2 Phase 1 — ED encounter spine and ED enrichment (one row per ED stay)

Rationale: Create a dedicated ED-stay-level cohort with ED-specific attributes to avoid mixing admission- and ED-level grains.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
    ↵ consistent across notebook stages.

# ED stay spine (rename stay_id to ed_stay_id)

SQL["ed_spine_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
    s.stay_id AS ed_stay_id,

```

```

s.subject_id,
s.hadm_id,
s.intime AS ed_intime,
s.outtime AS ed_outtime,
s.arrival_transport,
s.disposition,
s.gender AS ed_gender,
s.race AS ed_race
FROM `'{PHYS}`.{ED}.edstays` s
JOIN hadms h ON h.hadm_id = s.hadm_id
"""

ed_spine = run_sql_bq(sql("ed_spine_sql"), {"hadms": hadm_list})
print("ED spine rows:", len(ed_spine), "unique ed_stay_id:",
      ed_spine["ed_stay_id"].nunique())

# ensure uniqueness
if ed_spine["ed_stay_id"].nunique() != len(ed_spine):
    raise ValueError("ed_stay_id not unique in ED spine")

# First ED presentation time per admission
ed_intime_first = (
    ed_spine.groupby("hadm_id", as_index=False)[["ed_intime"]]
    .min()
    .rename(columns={"ed_intime": "ed_intime_first"})
)

# Start ED-level df
ed_df = ed_spine.copy()
ed_df = ed_df.merge(ed_intime_first, on="hadm_id", how="left")

ED spine rows: 7166 unique ed_stay_id: 7166

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
#           consistent across notebook stages.

# ED triage and first ED vitals (reuse existing logic if present; otherwise
#           join)

def _needs_cols(df, cols):
    return (df is None) or any(c not in df.columns for c in cols)

```

```

# Use existing ed_triage / ed_first if already in memory from earlier cells
try:
    _ = ed_triage
except NameError:
    ed_triage = None

try:
    _ = ed_first
except NameError:
    ed_first = None

# Force re-query if required keys are missing
if _needs_cols(locals().get('ed_triage', None), ['ed_stay_id', 'hadm_id']):
    ed_triage = None
if _needs_cols(locals().get('ed_first', None), ['ed_stay_id']):
    ed_first = None

# If missing, prefer in-memory tables loaded earlier in the notebook.
# This avoids re-running expensive ED queries during nbconvert.
if ed_triage is None:
    if 'edmap' in globals() and 'tri_all' in globals():
        ed_triage = (
            edmap[['stay_id', 'hadm_id']]
            .merge(tri_all, on='stay_id', how='left')
            .rename(columns={'stay_id': 'ed_stay_id'})
        )
        print('ED triage rows (from in-memory tables):', len(ed_triage))
    else:
        ed_triage_sql = f"""
        WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
        SELECT
            s.stay_id AS ed_stay_id,
            s.hadm_id,
            t.temperature AS ed_triage_temp,
            t.heartrate AS ed_triage_hr,
            t.resprate AS ed_triage_rr,
            t.o2sat AS ed_triage_o2sat,
            t.sbp AS ed_triage_sbp,
            t.dbp AS ed_triage_dbp,
            t.pain AS ed_triage_pain,
            t.acuity AS ed_triage_acuity,
            t.chiefcomplaint AS ed_triage_cc
        FROM `{{PHYS}}.{{ED}}.edstays` s

```

```

JOIN hadms h ON h.hadm_id = s.hadm_id
LEFT JOIN `{{PHYS}}.{{ED}}.triage` t ON t.stay_id = s.stay_id
"""
ed_triage = run_sql_bq(ed_triage_sql, {'hadms': hadm_list})
print('ED triage rows (fallback query):', len(ed_triage))

if ed_first is None:
    if 'edmap' in globals() and 'first_stay_all' in globals():
        ed_first = (
            edmap[['stay_id', 'hadm_id']]
            .merge(first_stay_all, on='stay_id', how='left')
            .rename(columns={'stay_id': 'ed_stay_id'})
        )
        print('ED first vitals rows (from in-memory tables):', len(ed_first))
    else:
        ed_first_vitals_sql = f"""
        WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
        edmap AS (
            SELECT stay_id, hadm_id
            FROM `{{PHYS}}.{{ED}}.edstays`
            WHERE hadm_id IN (SELECT hadm_id FROM hadms)
        ),
        vs AS (
            SELECT stay_id, charttime, temperature, heartrate, resprate, o2sat,
        ← sbp, dbp, rhythm, pain
            FROM `{{PHYS}}.{{ED}}.vitalsign`
        ),
        first_vs AS (
            SELECT
                v.stay_id,
                (ARRAY_AGG(STRUCT(v.charttime, v.temperature, v.heartrate,
        ← v.resprate, v.o2sat, v.sbp, v.dbp, v.rhythm, v.pain)
                    ORDER BY v.charttime LIMIT 1)) [OFFSET(0)] AS pick
            FROM vs v
            JOIN edmap m USING (stay_id)
            GROUP BY v.stay_id
        )
        SELECT
            f.stay_id AS ed_stay_id,
            m.hadm_id,
            pick.charttime AS ed_first_vitals_time,
            pick.temperature AS ed_first_temp,
            pick.heartrate AS ed_first_hr,

```

```

pick.resprate AS ed_first_rr,
pick.o2sat AS ed_first_o2sat,
pick.sbp AS ed_first_sbp,
pick.dbp AS ed_first_dbp,
pick.rhythm AS ed_first_rhythm,
pick.pain AS ed_first_pain
FROM first_vs f
JOIN edmap m ON m.stay_id = f.stay_id
"""
ed_first = run_sql_bq(ed_first_vitals_sql, {'hadms': hadm_list})
print('ED first vitals rows (fallback query):', len(ed_first))

# Debug columns before merge
print('ed_triage cols:', list(ed_triage.columns))
print('ed_first cols:', list(ed_first.columns))
print('ed_df cols:', list(ed_df.columns))

if 'ed_stay_id' not in ed_df.columns:
    raise KeyError('ed_df missing ed_stay_id; ensure ED spine cell ran.')

# Merge ED triage + vitals onto ed_df with available keys
merge_keys_triage = [k for k in ["ed_stay_id", "hadm_id"] if k in
    ↪ ed_df.columns and k in ed_triage.columns]
if not merge_keys_triage:
    raise KeyError("No common keys between ed_df and ed_triage")
ed_df = ed_df.merge(ed_triage, on=merge_keys_triage, how="left")
merge_keys_first = [k for k in ["ed_stay_id", "hadm_id"] if k in
    ↪ ed_df.columns and k in ed_first.columns]
if not merge_keys_first:
    raise KeyError("No common keys between ed_df and ed_first")
ed_df = ed_df.merge(ed_first, on=merge_keys_first, how="left")

ED triage rows (from in-memory tables): 7166
ED first vitals rows (from in-memory tables): 7166
ed_triage cols: ['ed_stay_id', 'hadm_id', 'ed_triage_temp', 'ed_triage_hr', 'ed_triage_rr',
ed_first cols: ['ed_stay_id', 'hadm_id', 'ed_first_vitals_time', 'ed_first_temp', 'ed_first_rr',
ed_df cols: ['ed_stay_id', 'subject_id', 'hadm_id', 'ed_intime', 'ed_outtime', 'arrival_transit_time']

```

3.0.3 Phase 2 — Hospital admission context and outcomes

Rationale: Add admission-level outcomes and demographics for downstream stratification and analysis.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↳ consistent across notebook stages.

# Admissions fields
SQL["admit_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
    a.hadm_id,
    a.admittime,
    a.dischtime,
    a.deathtime,
    a.hospital_expire_flag,
    a.admission_type AS administrative_class,
    a.admission_location AS from_location,
    a.discharge_location,
    a.insurance,
    a.language,
    a.marital_status,
    a.race AS hosp_race
FROM `{{PHYS}}.{{HOSP}}.admissions` a
JOIN hadms h USING (hadm_id)
"""

admit = run_sql_bq(sql("admit_sql"), {"hadms": hadm_list})
print("Admissions rows:", len(admit))

ed_df = ed_df.merge(admit, on="hadm_id", how="left")

# Merge ventilation flags/times (hadm-level)
if "vent_combined" in globals():
    ed_df = ed_df.merge(vent_combined, on="hadm_id", how="left")

# Derived outcomes
ed_df["hospital_los_hours"] = (ed_df["dischtime"] -
                                ed_df["admittime"]).dt.total_seconds() / 3600.0
ed_df["in_hospital_death"] = ((ed_df["hospital_expire_flag"] == 1) |
                                ed_df["deathtime"].notna()).astype("int64")

# Concordance check
discord = (
    ((ed_df["hospital_expire_flag"] == 1) & ed_df["deathtime"].isna()) |
    ((ed_df["hospital_expire_flag"] == 0) & ed_df["deathtime"].notna())
)

```

```
print("Admissions discordance (expire_flag vs deathtime):",
    int(discord.sum()))
```

```
Admissions rows: 9406
Admissions discordance (expire_flag vs deathtime): 1
```

3.0.4 Phase 3 — ICU timing and LOS

Rationale: Capture ICU exposure, timing, and total LOS for severity stratification.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
    consistent across notebook stages.
```

```
# ICU stays (aggregate per hadm)
SQL["icu_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
    i.hadm_id,
    i.stay_id AS icu_stay_id,
    i.intime,
    i.outtime,
    i.los,
    i.first_careunit,
    i.last_careunit
FROM `{{PHYS}}.{{ICU}}.icustays` i
JOIN hadms h USING (hadm_id)
"""

icu = run_sql_bq(sql("icu_sql"), {"hadms": hadm_list})
print("ICU stay rows:", len(icu))
```

```
if len(icu) > 0:
    icu_agg = (
        icu.sort_values(["hadm_id", "intime"]).groupby("hadm_id",
    as_index=False)
        .agg(
            icu_intime_first=("intime", "min"),
            icu_outtime_last=("outtime", "max"),
            icu_los_total=("los", "sum"),
            n_icu_stays=("icu_stay_id", "nunique"),
            first_careunit=("first_careunit", "first"),
            last_careunit=("last_careunit", "last"),
        )
)
```

```

)
ed_df = ed_df.merge(icu_agg, on="hadm_id", how="left")
else:
    print("No ICU stays found for cohort.")

```

ICU stay rows: 8469

3.0.5 Phase 4 — ED longitudinal vitals (0–6h)

Rationale: Summarize early ED vitals for severity phenotyping.

```

# Purpose: Build ED vitals features robustly without a single large BigQuery
#           ↵ join that can stall.

# ED vitals long + aggregates (0–6h)

# Build the ED stay map from in-memory cohort tables to keep query scope
#           ↵ tight and deterministic.
if "ed_df" not in globals():
    raise NameError("ed_df is required before ED vitals extraction")

required_cols = ["ed_stay_id", "hadm_id", "ed_intime"]
missing_required = [c for c in required_cols if c not in ed_df.columns]
if missing_required:
    raise KeyError(f"ed_df missing required columns for ED vitals extraction:
                   ↵ {missing_required}")

edmap_local = (
    ed_df[required_cols]
    .dropna(subset=["ed_stay_id"])
    .drop_duplicates(subset=["ed_stay_id"])
    .copy()
)
edmap_local["ed_stay_id"] = pd.to_numeric(edmap_local["ed_stay_id"],
                                         ↵ errors="coerce")
edmap_local = edmap_local.dropna(subset=["ed_stay_id"])
edmap_local["ed_stay_id"] = edmap_local["ed_stay_id"].astype(int)

stay_ids = sorted(edmap_local["ed_stay_id"].unique().tolist())
print("ED stays for vitals pull:", len(stay_ids))

SQL["ed_vitals_sql"] = f"""

```

```

SELECT
    stay_id AS ed_stay_id,
    charttime,
    temperature,
    heartrate,
    resprate,
    o2sat,
    sbp,
    dbp,
    rhythm,
    pain
FROM `'{PHYS}.ED.vitalsign`
WHERE stay_id IN UNNEST(@stay_ids)
"""

# Query in chunks so we avoid one very large parameter payload/job.
chunk_size = 5000
vitals_chunks = []
for i in range(0, len(stay_ids), chunk_size):
    chunk = stay_ids[i:i + chunk_size]
    part = run_sql_bq(sql("ed_vitals_sql"), {"stay_ids": chunk})
    if len(part) > 0:
        part["ed_stay_id"] = pd.to_numeric(part["ed_stay_id"],
                                         errors="coerce")
        part = part.dropna(subset=["ed_stay_id"])
        part["ed_stay_id"] = part["ed_stay_id"].astype(int)
        vitals_chunks.append(part)
    print(f"ED vitals chunk {i // chunk_size + 1}: rows={len(part)})"

if vitals_chunks:
    ed_vitals_long = pd.concat(vitals_chunks, ignore_index=True)
else:
    ed_vitals_long = pd.DataFrame(
        columns=[
            "ed_stay_id", "charttime", "temperature", "heartrate",
            "resprate",
            "o2sat", "sbp", "dbp", "rhythm", "pain"
        ]
    )

ed_vitals_long = ed_vitals_long.merge(edmap_local, on="ed_stay_id",
                                      how="inner")
print("ED vitals long rows:", len(ed_vitals_long))

```

```

# Window filter: 0-6h from ED intime
ed_vitals_long["dt_hours"] = (ed_vitals_long["charttime"] -
    ed_vitals_long["ed_intime"]).dt.total_seconds() / 3600.0
in_6h = ed_vitals_long["dt_hours"].between(0, 6, inclusive="both")

agg = (
    ed_vitals_long.loc[in_6h]
    .groupby("ed_stay_id", as_index=False)
    .agg(
        max_heartrate_0_6h=("heartrate", "max"),
        max_resprate_0_6h=("resprate", "max"),
        min_o2sat_0_6h=("o2sat", "min"),
        min_sbp_0_6h=("sbp", "min"),
        n_vitals_0_6h=("charttime", "count"),
    )
)

ed_df = ed_df.merge(agg, on="ed_stay_id", how="left")

# Range warnings (report only, do not drop)
range_checks = {
    "heartrate": (0, 300),
    "resprate": (0, 80),
    "o2sat": (0, 100),
    "sbp": (0, 300),
}
for col, (lo, hi) in range_checks.items():
    bad = ed_vitals_long[col].notna() & (~ed_vitals_long[col].between(lo,
    hi))
    if bad.any():
        print(f"Warning: {col} out of range count:", int(bad.sum()))

```

```

ED stays for vitals pull: 7166
ED vitals chunk 1: rows=32345
ED vitals chunk 2: rows=13897
ED vitals long rows: 46242
Warning: heartrate out of range count: 1
Warning: resprate out of range count: 1

```

3.0.6 Phase 5 — Robust lab discovery + gas panels

Rationale: Capture blood gas and key chemistry labs with label-robust item discovery and unit normalization.

```
# Purpose: Discover lab item IDs robustly and record cohort-aware item
#           frequency for auditability.

import re
import json

# Discover itemids from d_labitems
SQL["labitems_sql"] = f"""
SELECT itemid, label, fluid, category
FROM `{{PHYS}}.{{HOSP}}.d_labitems`
"""

labitems = run_sql_bq(sql("labitems_sql"))

patterns = {
    "gas_pco2": re.compile(r"\bp\s*co2\b|pco2|pco ", re.I),
    "gas_po2": re.compile(r"\bp\s*o2\b|po2|pao2|po ", re.I),
    "gas_ph": re.compile(r"\bph\b", re.I),
    "gas_hco3": re.compile(r"hco3|bicarbonate", re.I),
    "gas_lactate": re.compile(r"lactate", re.I),
    "gas_specimen": re.compile(r"specimen|source|type", re.I),
    "chem_creatinine": re.compile(r"creatinine", re.I),
    "chem_sodium": re.compile(r"\bsodium\b", re.I),
    "chem_chloride": re.compile(r"\bchloride\b", re.I),
    "chem_total_co2": re.compile(r"carbon dioxide|total co2|\bco2\b", re.I),
    "cbc_hemoglobin": re.compile(r"hemoglobin", re.I),
}

# Category filters limit false matches (for example chemistry CO2 vs
#           blood-gas pCO2)
cat_gas = re.compile(r"blood\s*gas|blood gas|arterial|venous", re.I)
cat_chem = re.compile(r"chemistry|chem|blood", re.I)
cat_cbc = re.compile(r"hematology|cbc", re.I)

manifest_lab_itemid_sets: dict[str, set[int]] = {
    "gas_pco2": set(LAB_PCO2_ITEMIDS),
    "gas_po2": set(LAB_PO2_ITEMIDS),
    "gas_ph": set(LAB_PH_ITEMIDS),
```

```

"gas_hco3": set(LAB_HC03_ITEMIDS),
"gas_specimen": set(LAB_SPECIMEN_TYPE_ITEMIDS),
}
manifest_missing_rows: list[dict[str, Any]] = []
for group_name, itemids in manifest_lab_itemid_sets.items():
    if not itemids:
        continue
    present_itemids = set(
        pd.to_numeric(
            labitems.loc[labitems["itemid"].isin(itemids), "itemid"],
            errors="coerce",
        )
        .dropna()
        .astype(int)
        .tolist()
    )
    missing_itemids = sorted(itemids.difference(present_itemids))
    if missing_itemids:
        manifest_missing_rows.append(
            {
                "group_name": group_name,
                "missing_itemids": ";" .join(str(value) for value in
                                              missing_itemids),
                "status": "missing_from_d_labitems",
            }
        )
if manifest_missing_rows:
    raise ValueError(
        "Manifest lab itemids were not found in d_labitems: "
        f"{manifest_missing_rows}. Update specs/blood_gas_itemids.json."
    )
matches = []
for name, pat in patterns.items():
    if name in manifest_lab_itemid_sets:
        itemids = manifest_lab_itemid_sets[name]
        dfm = labitems.loc[labitems["itemid"].isin(itemids)].copy()
    else:
        dfm = labitems.copy()
        dfm = dfm[dfm["label"].str.contains(pat, na=False)]
        if name.startswith("gas_"):
            dfm = dfm[dfm["category"].str.contains(cat_gas, na=False)]

```

```

    elif name.startswith("chem_"):
        dfm = dfm[dfm["category"].str.contains(cat_chem, na=False)]
    elif name.startswith("cbc_"):
        dfm = dfm[dfm["category"].str.contains(cat_cbc, na=False)]
    matches[name] = dfm[["itemid", "label", "category"]]

# Build lab_item_map with cohort counts. Run hadm counts in chunks to avoid
# one oversized query.
itemids_all = sorted({int(i) for dfm in matches.values() for i in
                     dfm["itemid"].tolist()})

SQL["counts_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT itemid, COUNT(*) AS n
FROM `{{PHYS}}.{{HOSP}}.labevents`
WHERE hadm_id IN (SELECT hadm_id FROM hadms)
    AND itemid IN UNNEST(@itemids)
GROUP BY itemid
"""

counts_cols = ["itemid", "n"]
if itemids_all and len(hadm_list) > 0:
    hadm_chunk_size = 10000
    count_parts = []
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
        part = run_sql_bq(sql("counts_sql"), {"hadms": hadm_chunk, "itemids": itemids_all})
        if len(part) > 0:
            count_parts.append(part)
    print(f"Lab item count chunk {i // hadm_chunk_size + 1}:
          rows={len(part)}")

if count_parts:
    counts = (
        pd.concat(count_parts, ignore_index=True)
        .groupby("itemid", as_index=False)[["n"]]
        .sum()
    )
else:
    counts = pd.DataFrame(columns=counts_cols)
else:
    counts = pd.DataFrame(columns=counts_cols)

```

```

lab_item_map = {}
for name, dfm in matches.items():
    tmp = dfm.merge(counts, on="itemid", how="left").fillna({"n": 0})
    lab_item_map[name] = {
        "pattern": patterns[name].pattern,
        "items": tmp.sort_values("n",
            ↪ ascending=False).to_dict(orient="records"),
    }

lab_item_map_path = WORK_DIR / "lab_item_map.json"
lab_item_map_path.write_text(json.dumps(lab_item_map, indent=2))
print("Wrote:", lab_item_map_path)

manifest_lab_audit_rows: list[dict[str, Any]] = []
for analyte_name, itemids in manifest_lab_itemid_sets.items():
    if not itemids:
        continue
    frame = labitems.loc[labitems["itemid"].isin(itemids), ["itemid",
    ↪ "label", "category", "fluid"]].copy()
    present_ids = set(frame["itemid"].astype(int).tolist()) if not
    ↪ frame.empty else set()
    frame = frame.merge(counts, on="itemid", how="left").fillna({"n": 0})
    record_map = {
        int(record["itemid"]): record for record in
        ↪ frame.to_dict(orient="records")
    }
    for itemid_value in sorted(itemids):
        record = record_map.get(int(itemid_value), {})
        manifest_lab_audit_rows.append(
            {
                "source_branch": "LAB",
                "analyte": analyte_name,
                "itemid": int(itemid_value),
                "label": record.get("label"),
                "category": record.get("category"),
                "fluid": record.get("fluid"),
                "cohort_row_count": int(record.get("n", 0)),
                "manifest_version": BLOOD_GAS_MANIFEST["version"],
                "present_in_dictionary": int(itemid_value in present_ids),
            }
        )
    )
)

```

```

blood_gas_itemid_manifest_audit = pd.DataFrame(manifest_lab_audit_rows)
missing_manifest_rows = blood_gas_itemid_manifest_audit.loc[
    pd.to_numeric(
        blood_gas_itemid_manifest_audit.get("present_in_dictionary"),
        errors="coerce"
    ).fillna(0).astype(int).eq(0)
].copy()
if not missing_manifest_rows.empty:
    missing_repr = ", ".join(
        f"{row.source_branch}:{row.analyte}:{int(row.itemid)}"
        for row in missing_manifest_rows.itertuples(index=False)
    )
    raise ValueError(
        "Blood-gas manifest itemids missing from dictionary tables: "
        f"{missing_repr}"
    )

```

Lab item count chunk 1: rows=35

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/lab_item_map.json

```

# Purpose: Extract cohort labs and apply ED-time windows in pandas for
# predictable runtime.

# Extract labevents and apply ED 0-24h windows locally

# Assemble itemid lists
itemid_sets = {k: [int(x["itemid"]) for x in v["items"]] for k, v in
    lab_item_map.items()}
all_itemids = sorted({i for v in itemid_sets.values() for i in v})

if "ed_df" not in globals():
    raise NameError("ed_df is required before lab window extraction")

ed_windows = (
    ed_df[["ed_stay_id", "hadm_id", "ed_intime"]]
    .dropna(subset=["ed_stay_id", "hadm_id", "ed_intime"])
    .drop_duplicates()
    .copy()
)
ed_windows["hadm_id"] = pd.to_numeric(ed_windows["hadm_id"],
    errors="coerce").astype("Int64")

```

```

ed_windows = ed_windows.dropna(subset=["hadm_id"])
ed_windows["hadm_id"] = ed_windows["hadm_id"].astype(int)

SQL["labs_sql"] = f"""
SELECT
    subject_id,
    hadm_id,
    itemid,
    charttime,
    specimen_id,
    valuenum,
    valueuom,
    value AS value_text
FROM `{{PHYS}}.{HOSP}.labevents`
WHERE hadm_id IN UNNEST(@hadms)
    AND itemid IN UNNEST(@itemids)
"""

# Run in hadm chunks; merge to ED windows in pandas to keep each query
#   lightweight.
labs_parts = []
if all_itemids and len(hadm_list) > 0:
    hadm_chunk_size = 5000
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
        part = run_sql_bq(sql("labs_sql"), {"hadms": hadm_chunk, "itemids": all_itemids})
        raw_n = len(part)

        if raw_n > 0:
            part["hadm_id"] = pd.to_numeric(part["hadm_id"], errors="coerce")
            part = part.dropna(subset=["hadm_id"])
            part["hadm_id"] = part["hadm_id"].astype(int)

            part = part.merge(ed_windows, on="hadm_id", how="inner")
            in_window = (
                (part["charttime"] >= part["ed_intime"]) &
                (part["charttime"] <= (part["ed_intime"] +
                pd.Timedelta(hours=24))))
            )
            part = part.loc[in_window, [
                "ed_stay_id", "subject_id", "hadm_id", "itemid", "charttime",
                "specimen_id", "valuenum", "valueuom", "value_text"
            ]

```

```

        ]]
    if len(part) > 0:
        labs_parts.append(part)

    print(f"Labs window chunk {i // hadm_chunk_size + 1}: raw={raw_n},
        ↵ kept={len(part) if raw_n > 0 else 0}")

if labs_parts:
    labs_long = pd.concat(labs_parts, ignore_index=True)
else:
    labs_long = pd.DataFrame(
        columns=["ed_stay_id", "subject_id", "hadm_id", "itemid",
        ↵ "charttime", "specimen_id", "valuenum", "valueuom", "value_text"]
    )

print("Labs long rows:", len(labs_long))

# Unit audit for pCO2
pco2_ids = itemid_sets.get("gas_pco2", [])
unit_audit = (
    labs_long.loc[labs_long["itemid"].isin(pco2_ids)]
    .groupby("valueuom", dropna=False)
    .size().reset_index(name="n")
)
unit_audit_path = WORK_DIR / "lab_unit_audit.csv"
unit_audit.to_csv(unit_audit_path, index=False)
print("Wrote:", unit_audit_path)

Labs window chunk 1: raw=613144, kept=96044
Labs window chunk 2: raw=560003, kept=84504
Labs long rows: 180548
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/
CC-NLP/lab_unit_audit.csv

# Purpose: Reconstruct gas panels and normalize units before ED-stay level
    ↵ aggregation.

# Reconstruct gas panels by specimen_id within ED stay

pco2_ids = set(itemid_sets.get("gas_pco2", []))
po2_ids = set(itemid_sets.get("gas_po2", []))
ph_ids = set(itemid_sets.get("gas_ph", []))

```

```

hco3_ids = set(itemid_sets.get("gas_hco3", []))
lact_ids = set(itemid_sets.get("gas_lactate", []))

labs = labs_long.copy()

if "fluid" in labitems.columns:
    item_fluid = (
        labitems[["itemid", "fluid"]]
        .drop_duplicates(subset=["itemid"])
        .copy()
    )
    labs = labs.merge(item_fluid, on="itemid", how="left")
else:
    labs["fluid"] = pd.NA

# Convert pCO2 kPa to mmHg when needed
is_kpa = labs["valueuom"].astype(str).str.lower().str.contains("kpa",
    ↵ na=False)
mask_pco2 = labs["itemid"].isin(pco2_ids)
if mask_pco2.any() and is_kpa.any():
    labs.loc[mask_pco2 & is_kpa, "valuenum"] = labs.loc[mask_pco2 & is_kpa,
    ↵ "valuenum"] * 7.50062
    labs.loc[mask_pco2 & is_kpa, "valueuom"] = "mmHg"

# panel by specimen_id
panel = (
    labs.groupby(["ed_stay_id", "specimen_id"], as_index=False)
    .agg(panel_time=("charttime", "min"))
)
# attach analytes

def pick_analyte(
    df: pd.DataFrame,
    ids: set[int],
    name: str,
    *,
    include_item_meta: bool = False,
) -> pd.DataFrame:
    cols = ["ed_stay_id", "specimen_id", "charttime", "valuenum", "itemid",
    ↵ "fluid"]
    tmp = df.loc[df["itemid"].isin(ids), cols].copy()
    tmp["valuenum"] = pd.to_numeric(tmp["valuenum"], errors="coerce")

```

```

tmp = tmp.loc[tmp["valuenum"].notna()].copy()
if tmp.empty:
    base_cols = ["ed_stay_id", "specimen_id", name]
    if include_item_meta:
        base_cols += [f"{name}_itemid", f"{name}_fluid"]
    return pd.DataFrame(columns=base_cols)
tmp = tmp.sort_values(
    ["ed_stay_id", "specimen_id", "charttime", "itemid"], kind="stable"
)
first = tmp.groupby(["ed_stay_id", "specimen_id"]),
↪ as_index=False).first()
first = first.rename(columns={"valuenum": name})
if include_item_meta:
    first = first.rename(
        columns={"itemid": f"{name}_itemid", "fluid": f"{name}_fluid"}
    )
    keep_cols = ["ed_stay_id", "specimen_id", name, f"{name}_itemid",
↪ f"{name}_fluid"]
else:
    keep_cols = ["ed_stay_id", "specimen_id", name]
return first[keep_cols]

if pco2_ids:
    panel = panel.merge(
        pick_analyte(
            labs,
            pco2_ids,
            "pco2",
            include_item_meta=True,
        ),
        on=["ed_stay_id", "specimen_id"],
        how="left",
    )
if po2_ids:
    panel = panel.merge(
        pick_analyte(labs, po2_ids, "po2"),
        on=["ed_stay_id", "specimen_id"],
        how="left",
    )
if ph_ids:
    panel = panel.merge(pick_analyte(labs, ph_ids, "ph"),
↪ on=["ed_stay_id", "specimen_id"], how="left")
if hco3_ids:

```

```

    panel = panel.merge(pick_analyte(labs, hco3_ids, "hco3"),
    ↵ on=["ed_stay_id", "specimen_id"], how="left")
if lact_ids:
    panel = panel.merge(pick_analyte(labs, lact_ids, "lactate"),
    ↵ on=["ed_stay_id", "specimen_id"], how="left")

# First pCO2-qualified panel per ED stay (gas anchor must include pCO2)
panel_anchor_candidates = panel.loc[
    pd.to_numeric(panel.get("pco2"), errors="coerce").notna()
].copy()
if panel_anchor_candidates.empty:
    first_panel = panel.head(0).copy()
else:
    first_panel = (
        panel_anchor_candidates.sort_values(["ed_stay_id", "panel_time"])
        .groupby("ed_stay_id", as_index=False)
        .first()
    )

# 0-6h and 0-24h extrema
panel = panel.merge(ed_df[["ed_stay_id", "ed_intime"]], on="ed_stay_id",
    ↵ how="left")
panel["dt_hours"] = (panel["panel_time"] -
    ↵ panel["ed_intime"]).dt.total_seconds() / 3600.0

# Ensure expected panel columns exist even if analyte is absent
for col in ["pco2", "po2", "ph", "hco3", "lactate"]:
    if col not in panel.columns:
        panel[col] = pd.NA
for col in ["pco2_itemid", "pco2_fluid"]:
    if col not in panel.columns:
        panel[col] = pd.NA

p06 = panel.loc[panel["dt_hours"].between(0, 6, inclusive="both")]
p24 = panel.loc[panel["dt_hours"].between(0, 24, inclusive="both")]

agg06 = p06.groupby("ed_stay_id",
    ↵ as_index=False).agg(max_pco2_0_6h=("pco2", "max"),
    ↵ min_ph_0_6h=("ph", "min"))
agg24 = p24.groupby("ed_stay_id",
    ↵ as_index=False).agg(max_pco2_0_24h=("pco2", "max"),
    ↵ min_ph_0_24h=("ph", "min"))

```

```

first_panel_export_cols = [
    col
    for col in ["ed_stay_id", "panel_time", "pco2", "ph", "hco3", "lactate"]
    if col in first_panel.columns
]
if first_panel_export_cols:
    ed_df = ed_df.merge(
        first_panel[first_panel_export_cols].rename(
            columns={
                "panel_time": "first_gas_time",
                "pco2": "first_pco2",
                "ph": "first_ph",
                "hco3": "first_hco3",
                "lactate": "first_lactate",
            }
        ),
        on="ed_stay_id",
        how="left",
    )
else:
    ed_df["first_gas_time"] = pd.NaT
    ed_df["first_pco2"] = pd.NA
    ed_df["first_ph"] = pd.NA
    ed_df["first_hco3"] = pd.NA
    ed_df["first_lactate"] = pd.NA

# Prioritize LAB blood-gas bicarbonate aligned to pCO2 anchor panel, then
# fallback
# to nearest LAB bicarbonate around first_gas_time when same-panel value is
# missing.
if "first_hco3_source" not in ed_df.columns:
    ed_df["first_hco3_source"] = "missing"

panel_hco3_present = pd.to_numeric(ed_df.get("first_hco3"),
    errors="coerce").notna()
ed_df.loc[panel_hco3_present, "first_hco3_source"] =
    "lab_panel_same_specimen"

lab_hco3_long = pd.DataFrame(columns=["ed_stay_id", "charttime", "valuenum"])
if hco3_ids and "first_gas_time" in ed_df.columns:
    lab_hco3_long = labs.loc[
        labs["itemid"].isin(hco3_ids),
        ["ed_stay_id", "charttime", "valuenum"],
    ]

```

```

].copy()
if not lab_hco3_long.empty:
    lab_hco3_long["valuenum"] = pd.to_numeric(lab_hco3_long["valuenum"], 
    ↵ errors="coerce")
    lab_hco3_long =
    ↵ lab_hco3_long.loc[lab_hco3_long["valuenum"].notna()].copy()
    anchor_missing = ed_df.loc[
        pd.to_numeric(ed_df["first_hco3"], errors="coerce").isna()
        & pd.to_datetime(ed_df["first_gas_time"]),
        ↵ errors="coerce").notna(),
        ["ed_stay_id", "first_gas_time"],
    ].copy()
if not anchor_missing.empty and not lab_hco3_long.empty:
    fallback_candidates = anchor_missing.merge(
        lab_hco3_long,
        on="ed_stay_id",
        how="left",
    )
    fallback_candidates["first_gas_time"] = pd.to_datetime(
        fallback_candidates["first_gas_time"], errors="coerce"
    )
    fallback_candidates["charttime"] = pd.to_datetime(
        fallback_candidates["charttime"], errors="coerce"
    )
    fallback_candidates["dt_abs_hours"] = (
        fallback_candidates["charttime"] -
    ↵ fallback_candidates["first_gas_time"])
        .dt.total_seconds()
        .abs()
        / 3600.0
    )
    fallback_candidates = fallback_candidates.loc[
        fallback_candidates["dt_abs_hours"].notna()
    ].copy()
if not fallback_candidates.empty:
    nearest_hco3 = (
        fallback_candidates.sort_values(
            ["ed_stay_id", "dt_abs_hours", "charttime"],
            kind="stable",
        )
        .groupby("ed_stay_id", as_index=False)
        .first()
        .rename(columns={"valuenum": "first_hco3_fallback_lab"})

```

```

        )
    ed_df = ed_df.merge(
        nearest_hco3[["ed_stay_id", "first_hco3_fallback_lab"]],
        on="ed_stay_id",
        how="left",
    )
    fill_mask = (
        pd.to_numeric(ed_df["first_hco3"]),
        & pd.to_numeric(ed_df["first_hco3_fallback_lab"]),
        & errors="coerce").isna()
    )
    ed_df.loc[fill_mask, "first_hco3"] = ed_df.loc[
        fill_mask, "first_hco3_fallback_lab"
    ]
    ed_df.loc[fill_mask, "first_hco3_source"] =
    "lab_nearest_to_anchor"
    ed_df = ed_df.drop(columns=["first_hco3_fallback_lab"])

if hco3_ids and not lab_hco3_long.empty:
    remaining_missing = pd.to_numeric(ed_df["first_hco3"]),
    & errors="coerce").isna()
    if remaining_missing.any():
        first_hco3_in_window = (
            lab_hco3_long.sort_values(["ed_stay_id", "charttime"],
            kind="stable")
            .groupby("ed_stay_id", as_index=False)
            .first()
            .rename(columns={"valuenum": "first_hco3_fallback_window"})
        )
        ed_df = ed_df.merge(
            first_hco3_in_window[["ed_stay_id",
            "first_hco3_fallback_window"]],
            on="ed_stay_id",
            how="left",
        )
        window_fill_mask = (
            pd.to_numeric(ed_df["first_hco3"], errors="coerce").isna()
            & pd.to_numeric(ed_df["first_hco3_fallback_window"],
            & errors="coerce").notna()
        )
        ed_df.loc[window_fill_mask, "first_hco3"] = ed_df.loc[
            window_fill_mask, "first_hco3_fallback_window"

```

```

        ]
    ed_df.loc[window_fill_mask, "first_hco3_source"] = "lab_first_0_24h"
    ed_df = ed_df.drop(columns=["first_hco3_fallback_window"])

ed_df["first_hco3_source"] = np.where(
    pd.to_numeric(ed_df["first_hco3"], errors="coerce").notna(),
    ed_df["first_hco3_source"].astype("string").fillna("missing"),
    "missing",
)
ed_df["first_gas_anchor_has_pco2"] = ed_df["first_gas_time"].notna()
ed_df["first_gas_anchor_source_validated"] = False

ed_df = ed_df.merge(agg06, on="ed_stay_id", how="left")
ed_df = ed_df.merge(agg24, on="ed_stay_id", how="left")

# Purpose: Reconstruct gas panels and normalize units before ED-stay level
#           aggregation.

# ABG vs VBG classification (specimen-type rows only for core cohort
#           behavior)
panel, gas_source_audit = assign_panel_source_from_specimen(
    panel,
    labs,
    specimen_source_itemids=itemid_sets.get("gas_specimen", []),
)
# Keep heuristic route diagnostics for QA transparency only (non-core
#           behavior).
heuristic_gas_source_audit: dict[str, Any] | None = None
heuristic_mode = os.getenv("COHORT_GAS_SOURCE_INFERENCE_MODE",
                           "metadata_only").strip().lower()
try:
    _, heuristic_gas_source_audit = infer_panel_gas_source_metadata(
        panel_df=panel.copy(),
        labs_df=labs,
        labitems_df=labitems,
        specimen_source_itemids=itemid_sets.get("gas_specimen", []),
        pco2_itemids=itemid_sets.get("gas_pco2", []),
        mode=heuristic_mode,
    )
except Exception as exc:
    heuristic_gas_source_audit = {

```

```

        "status": "error",
        "message": f"heuristic_diagnostics_failed: {exc}",
    }
gas_source_audit["heuristic_reference"] = heuristic_gas_source_audit

fail_on_all_other_source = os.getenv("COHORT_FAIL_ON_ALL_OTHER_SOURCE",
    ↪ "1").strip() == "1"
if gas_source_audit["all_other_or_unknown"]:
    if fail_on_all_other_source:
        assert_gas_source_coverage(
            gas_source_audit,
            fail_on_all_other_source=fail_on_all_other_source,
        )
    else:
        print(
            "WARNING: Gas source attribution collapsed to other/unknown; "
            "continuing because COHORT_FAIL_ON_ALL_OTHER_SOURCE=0."
        )

warn_other_rate_threshold = float(os.getenv("COHORT_WARN_OTHER_RATE",
    ↪ "0.50"))
fail_other_rate_raw = os.getenv("COHORT_FAIL_OTHER_RATE", "").strip()
fail_other_rate_threshold = float(fail_other_rate_raw) if fail_other_rate_raw
    ↪ else None
other_source_rate = float(gas_source_audit.get("source_rates",
    ↪ {}).get("unknown", 0.0))

if fail_other_rate_threshold is not None and other_source_rate >=
    ↪ fail_other_rate_threshold:
    raise ValueError(
        "gas_source_other_rate exceeded configured fail threshold "
        f"{fail_other_rate_threshold:.3f}: {other_source_rate:.4f}."
    )
if other_source_rate >= warn_other_rate_threshold:
    print(
        "WARNING: gas_source_other_rate exceeded warning threshold "
        f"{warn_other_rate_threshold:.3f}: {other_source_rate:.4f}."
    )

print("Gas source rates:", gas_source_audit.get("source_rates", {}))
print("Gas source inference tiers:", gas_source_audit.get("tier_rates", {}))
print("Unresolved specimen IDs:",
    ↪ gas_source_audit.get("unresolved_specimen_id_count", 0))

```

```

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")
gas_source_audit_path = prior_runs_dir / f"{out_date} gas_source_audit.json"

reduction_target = float(os.getenv("COHORT_OTHER_RELATIVE_REDUCTION_MIN",
    ↪ "0.10"))
baseline_other_rate = None
baseline_audit_path = None
baseline_mode_mismatch_path = None
current_source_mode = str(gas_source_audit.get("mode", "")).strip().lower()
for candidate in sorted(prior_runs_dir.glob("* gas_source_audit.json"),
    ↪ reverse=True):
    if candidate.resolve() == gas_source_audit_path.resolve():
        continue
    try:
        baseline_payload = json.loads(candidate.read_text())
        baseline_mode = str(baseline_payload.get("mode", "")).strip().lower()
        if current_source_mode and baseline_mode and baseline_mode !=
            ↪ current_source_mode:
            baseline_mode_mismatch_path = candidate
            continue
        baseline_rates = baseline_payload.get("source_rates", {})
        baseline_other_rate = float(
            baseline_rates.get(
                "unknown",
                baseline_rates.get("other", float("nan"))),
            )
    )
    if math.isnan(baseline_other_rate):
        baseline_other_rate = None
        continue
    baseline_audit_path = candidate
    break
except Exception:
    continue

if baseline_other_rate is None or baseline_other_rate <= 0:
    gas_source_audit["baseline_other_rate"] = None
    gas_source_audit["baseline_path"] = str(baseline_audit_path) if
    ↪ baseline_audit_path else None
    gas_source_audit["baseline_mode_mismatch_path"] = (

```

```

        str(baseline_mode_mismatch_path) if baseline_mode_mismatch_path else
        ↵  None
    )
gas_source_audit["relative_reduction_vs_baseline"] = None
print(
    "WARNING: Unable to evaluate gas_source_other_rate relative reduction
    ↵  "
    "because no valid baseline gas_source_audit was found for the current
    ↵  source mode."
)
else:
    relative_reduction = (baseline_other_rate - other_source_rate) /
    ↵ baseline_other_rate
    gas_source_audit["baseline_other_rate"] = baseline_other_rate
    gas_source_audit["baseline_path"] = str(baseline_audit_path)
    gas_source_audit["baseline_mode_mismatch_path"] = (
        str(baseline_mode_mismatch_path) if baseline_mode_mismatch_path else
        ↵  None
    )
    gas_source_audit["relative_reduction_vs_baseline"] = relative_reduction
    gas_source_audit["relative_reduction_target"] = reduction_target
    print(
        "Gas source OTHER relative reduction vs baseline: "
        f"{relative_reduction:.4f} (target={reduction_target:.4f})"
    )
    reduction_tolerance = 1e-9
    if reduction_target > 0 and (relative_reduction + reduction_tolerance) <
        ↵  reduction_target:
        message = (
            "gas_source_other_rate relative reduction target not met: "
            f"baseline={baseline_other_rate:.4f},
            ↵  current={other_source_rate:.4f}, "
            f"reduction={relative_reduction:.4f},
            ↵  target={reduction_target:.4f}."
        )
        contract_mode = os.getenv("PIPELINE_CONTRACT_MODE",
    ↵  "fail").strip().lower()
        if contract_mode == "fail":
            raise ValueError(message)
        print(f"WARNING: {message}")

gas_source_audit_path.write_text(json.dumps(gas_source_audit, indent=2))
print("Wrote:", gas_source_audit_path)

```

```

panel["pco2_source_branch"] = "LAB"
panel_anchor = panel.loc[pd.to_numeric(panel["pco2"],
    ↪ errors="coerce").notna()].copy()
if panel_anchor.empty:
    first_anchor = ed_df[["ed_stay_id"]].drop_duplicates().copy()
    first_anchor["first_gas_time"] = pd.NaT
    first_anchor["first_pco2"] = pd.NA
    first_anchor["first_ph"] = pd.NA
    first_anchor["first_hco3_anchor"] = pd.NA
    first_anchor["first_lactate"] = pd.NA
    first_anchor["first_gas_specimen_type"] = pd.NA
    first_anchor["first_gas_specimen_present"] = False
    first_anchor["first_gas_pco2_itemid"] = pd.NA
    first_anchor["first_gas_pco2_fluid"] = pd.NA
    first_anchor["co2_other_is_blood asserted"] = False
    first_anchor["first_gas_anchor_has_pco2"] = False
    first_anchor["first_gas_anchor_source_validated"] = False
else:
    first_anchor = (
        panel_anchor.sort_values(["ed_stay_id", "panel_time"])
            .groupby("ed_stay_id", as_index=False)
            .first()
    )
    first_anchor = first_anchor.rename(
        columns={
            "panel_time": "first_gas_time",
            "pco2": "first_pco2",
            "ph": "first_ph",
            "hco3": "first_hco3_anchor",
            "lactate": "first_lactate",
            "source_hint_specimen_text": "first_gas_specimen_type",
            "pco2_itemid": "first_gas_pco2_itemid",
            "pco2_fluid": "first_gas_pco2_fluid",
        }
    )
    first_anchor["first_gas_specimen_present"] = (
        first_anchor["first_gas_specimen_type"]
            .astype("string")
            .fillna("")
            .str.strip()
            .ne("")
    )
)

```

```

first_anchor["co2_other_is_blood_asserted"] = (
    first_anchor["source"].astype("string").str.lower().eq("unknown")
    & first_anchor["first_gas_pco2_fluid"]
    .astype("string")
    .str.strip()
    .str.lower()
    .isin(LAB_BLOOD_FLUID_TERMS)
)
first_anchor["first_gas_anchor_has_pco2"] = True
first_anchor["first_gas_anchor_source_validated"] =
    first_anchor["source"].isin(
        {"arterial", "venous", "unknown"}
    )

anchor_merge_cols = [
    "ed_stay_id",
    "first_gas_time",
    "first_pco2",
    "first_ph",
    "first_hco3_anchor",
    "first_lactate",
    "first_gas_specimen_type",
    "first_gas_specimen_present",
    "first_gas_pco2_itemid",
    "first_gas_pco2_fluid",
    "co2_other_is_blood_asserted",
    "first_gas_anchor_has_pco2",
    "first_gas_anchor_source_validated",
]
drop_existing_anchor_cols = [
    col
    for col in anchor_merge_cols
    if col in ed_df.columns and col not in {"ed_stay_id",
        "first_hco3_anchor"}
]
if drop_existing_anchor_cols:
    ed_df = ed_df.drop(columns=drop_existing_anchor_cols)
ed_df = ed_df.merge(first_anchor[anchor_merge_cols], on="ed_stay_id",
    how="left")
if "first_hco3_source" not in ed_df.columns:
    ed_df["first_hco3_source"] = "missing"
if "first_hco3" not in ed_df.columns:
    ed_df["first_hco3"] = pd.NA

```

```

hco3_anchor_numeric = pd.to_numeric(ed_df.get("first_hco3_anchor"),
    ↵   errors="coerce")
hco3_existing_numeric = pd.to_numeric(ed_df.get("first_hco3"),
    ↵   errors="coerce")
hco3_fill_from_anchor = hco3_existing_numeric.isna() &
    ↵   hco3_anchor_numeric.notna()
ed_df.loc[hco3_fill_from_anchor, "first_hco3"] =
    ↵   hco3_anchor_numeric.loc[hco3_fill_from_anchor]
ed_df.loc[hco3_fill_from_anchor, "first_hco3_source"] =
    ↵   "lab_panel_same_specimen"
ed_df["first_hco3_source"] = np.where(
    pd.to_numeric(ed_df["first_hco3"], errors="coerce").notna(),
    ed_df["first_hco3_source"].astype("string").fillna("missing"),
    "missing",
)
if "first_hco3_anchor" in ed_df.columns:
    ed_df = ed_df.drop(columns=["first_hco3_anchor"])
ed_df["first_gas_anchor_has_pco2"] = (
    ed_df["first_gas_anchor_has_pco2"].fillna(False).astype(bool)
)
ed_df["first_gas_anchor_source_validated"] = (
    ed_df["first_gas_anchor_source_validated"].fillna(False).astype(bool)
)
ed_df["first_gas_specimen_present"] = (
    ed_df["first_gas_specimen_present"].fillna(False).astype(bool)
)
ed_df["co2_other_is_blood_asserted"] = (
    ed_df["co2_other_is_blood_asserted"].fillna(False).astype(bool)
)

# per-stay unknown source rate
if len(panel) > 0:
    unk_rate = (
        panel.assign(_unk=(panel["source"].fillna("unknown") == "unknown"))
            .groupby("ed_stay_id", as_index=False)[["_unk"]].mean()
            .rename(columns={"_unk": "gas_source_unknown_rate"})
    )
else:
    unk_rate = panel[["ed_stay_id"]].drop_duplicates()
    unk_rate["gas_source_unknown_rate"] = 1.0

if len(panel) > 0:
    panel_for_diag = panel.copy()

```

```

panel_for_diag["source_inference_tier"] = (
    panel_for_diag["source_inference_tier"]
        .astype("string")
        .fillna("fallback_unknown")
)
panel_for_diag["source_hint_conflict"] = (
    panel_for_diag["source_hint_conflict"]
        .fillna(False)
        .astype(bool)
)
panel_for_diag["source"] = (
    panel_for_diag["source"].astype("string").fillna("unknown").str.lower()
)
panel_for_diag["source_resolved"] = panel_for_diag["source"].isin(
    {"arterial", "venous"}
)
panel_for_diag["tier_specimen_text"] =
    panel_for_diag["source_inference_tier"].eq(
        "specimen_text"
)
panel_for_diag["tier_label_fluid"] =
    panel_for_diag["source_inference_tier"].eq(
        "label_fluid"
)
panel_for_diag["tier_cluster_inheritance"] = panel_for_diag[
    "source_inference_tier"
].eq("cluster_inheritance")
panel_for_diag["tier_panel_cooccurrence"] = panel_for_diag[
    "source_inference_tier"
].eq("panel_cooccurrence")
panel_for_diag["tier_fallback_unknown"] =
    panel_for_diag["source_inference_tier"].eq(
        "fallback_unknown"
)

def _tier_mode(series: pd.Series) -> str:
    mode_values = series.mode(dropna=True)
    if mode_values.empty:
        return "fallback_unknown"
    return str(mode_values.iloc[0])

source_diag = (

```

```

    panel_for_diag.groupby("ed_stay_id", as_index=False)
    .agg(
        gas_source_inference_primary_tier=("source_inference_tier",
        _tier_mode),
        gas_source_hint_conflict_rate=("source_hint_conflict", "mean"),
        gas_source_resolved_rate=("source_resolved", "mean"),
        gas_source_tier_specimen_text_rate=("tier_specimen_text",
        "mean"),
        gas_source_tier_label_fluid_rate=("tier_label_fluid", "mean"),

    ↵  gas_source_tier_cluster_inheritance_rate=("tier_cluster_inheritance",
    ↵  "mean"),
    ↵  gas_source_tier_panel_cooccurrence_rate=("tier_panel_cooccurrence",
    ↵  "mean"),
        gas_source_tierFallback_unknown_rate=("tierFallback_unknown",
    ↵  "mean"),
    ↵ )
    .copy()
)
else:
    source_diag = ed_df[["ed_stay_id"]].drop_duplicates().copy()
    source_diag["gas_source_inference_primary_tier"] = "fallback_unknown"
    source_diag["gas_source_hint_conflict_rate"] = 0.0
    source_diag["gas_source_resolved_rate"] = 0.0
    source_diag["gas_source_tier_specimen_text_rate"] = 0.0
    source_diag["gas_source_tier_label_fluid_rate"] = 0.0
    source_diag["gas_source_tier_cluster_inheritance_rate"] = 0.0
    source_diag["gas_source_tier_panel_cooccurrence_rate"] = 0.0
    source_diag["gas_source_tierFallback_unknown_rate"] = 1.0

ed_df = ed_df.merge(unk_rate, on="ed_stay_id", how="left")
ed_df = ed_df.merge(source_diag, on="ed_stay_id", how="left")
ed_df["gas_source_unknown_rate"] =
    ↵ ed_df["gas_source_unknown_rate"].fillna(1.0)
ed_df["gas_source_other_rate"] = ed_df["gas_source_unknown_rate"]

```

WARNING: gas_source_other_rate exceeded warning threshold 0.500: 0.6051.

Gas source rates: {'unknown': 0.6051488712390939, 'arterial': 0.24304550256269808, 'venous':

Gas source inference tiers: {'fallback_unknown': 0.6051488712390939, 'specimen_text': 0.39488}

Unresolved specimen IDs: 30903

Gas source OTHER relative reduction vs baseline: 0.0000 (target=0.0000)

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 gas_source_audit.json

3.0.7 Phase 5C — ICU POC blood gases (chartevents, optional)

Rationale: Capture ICU point-of-care gases if central lab labevents miss early measurements.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
#           ↵ consistent across notebook stages.

# Discover ICU POC itemids from d_items

SQL["ditems_sql"] = f"""
SELECT itemid, label, category
FROM `{{PHYS}}.{ICU}.d_items`
"""

ditems = run_sql_bq(sql("ditems_sql"))

icu_patterns = {
    "pco2": re.compile(r"(?:^|[^a-z])p(?:a)?\s*co(?:2|)(?:[a-z]|$)", re.I),
    "po2": re.compile(r"(?:^|[^a-z])p(?:a)?\s*o(?:2|)(?:[a-z]|$)", re.I),
    "ph": re.compile(r"\bph\b", re.I),
    "hco3": re.compile(r"hc03|bicarbonate", re.I),
    "lactate": re.compile(r"lactate", re.I),
    "specimen": re.compile(r"specimen|source|type", re.I),
}

icu_cat = re.compile(r"blood\s*gas|blood gas|resp|arterial|venous", re.I)
icu_pco2_exclusion = re.compile(ICU_LABEL_DENY_REGEX, re.I)

icu_matches = []
manifest_icu_itemids = {
    "pco2": set(ICU_PC02_ITEMIDS),
    "po2": set(ICU_P02_ITEMIDS),
    "ph": set(ICU_PH_ITEMIDS),
    "hco3": set(ICU_HC03_ITEMIDS),
    "specimen": set(ICU_SPECIMEN_TYPE_ITEMIDS),
}
for name, pat in icu_patterns.items():
    manifest_ids = manifest_icu_itemids.get(name, set())
    if manifest_ids:
        dfm = ditems.loc[ditems["itemid"].isin(manifest_ids)].copy()
```

```

else:
    dfm = pd.DataFrame(columns=["itemid", "label", "category"])

    allowFallbackForAnalyte = ICU_ALLOW_PATTERN_FALLBACK and name != "hco3"
    if dfm.empty and allowFallbackForAnalyte:
        dfm = ditems[ditems["label"].str.contains(pat, na=False)].copy()
        if name in {"pco2", "ph", "specimen"}:
            dfm = dfm[dfm["category"].str.contains(icu_cat, na=False)]


    if name == "pco2" and len(dfm) > 0:
        dfm = dfm[~dfm["label"].str.contains(icu_pco2_exclusion, na=False)]
        icu_matches[name] = dfm[["itemid", "label", "category"]]

icu_itemid_map_frames = [
    frame.assign(analyte=name)
    for name, frame in icu_matches.items()
    if not frame.empty
]
if not icu_itemid_map_frames:
    icu_itemid_map = pd.DataFrame(columns=["analyte", "itemid", "label",
                                           "category"])
else:
    icu_itemid_map = pd.concat(icu_itemid_map_frames, ignore_index=True)
    icu_itemid_map = (
        icu_itemid_map[["analyte", "itemid", "label", "category"]]
        .drop_duplicates()
        .sort_values(["analyte", "itemid"], kind="stable")
        .reset_index(drop=True)
    )

icu_itemids = sorted({int(i) for dfm in icu_matches.values() for i in
                     dfm["itemid"].tolist()})
print("ICU POC candidate itemids:", len(icu_itemids))
if not icu_itemid_map.empty:
    print("ICU POC itemid map preview:")
    print(icu_itemid_map.head(20).to_string(index=False))

audit_date = pd.Timestamp.now().strftime("%Y-%m-%d")
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
icu_itemid_map_path = prior_runs_dir / f"{audit_date} icu_poc_itemid_map.csv"
icu_itemid_map.to_csv(icu_itemid_map_path, index=False)

```

```

print("Wrote:", icu_itemid_map_path)

manifest_icu_audit_rows: list[dict[str, Any]] = []
for analyte_name, itemids in manifest_icu_itemids.items():
    if not itemids:
        continue
    frame = ditems.loc[ditems["itemid"].isin(itemids), ["itemid", "label",
    "category"]].copy()
    present_ids = set(frame["itemid"].astype(int).tolist()) if not
    frame.empty else set()
    for itemid in sorted(itemids):
        row = {
            "source_branch": "ICU",
            "analyte": analyte_name,
            "itemid": int(itemid),
            "label": None,
            "category": None,
            "fluid": None,
            "cohort_row_count": 0,
            "manifest_version": BLOOD_GAS_MANIFEST["version"],
            "present_in_dictionary": int(itemid in present_ids),
        }
        if itemid in present_ids:
            record = frame.loc[frame["itemid"].eq(itemid)].iloc[0]
            row["label"] = record.get("label")
            row["category"] = record.get("category")
        manifest_icu_audit_rows.append(row)

if "blood_gas_itemid_manifest_audit" in globals() and not
    blood_gas_itemid_manifest_audit.empty:
    manifest_audit_combined = pd.concat(
        [blood_gas_itemid_manifest_audit,
    pd.DataFrame(manifest_icu_audit_rows)],
        ignore_index=True,
        sort=False,
    )
else:
    manifest_audit_combined = pd.DataFrame(manifest_icu_audit_rows)

blood_gas_itemid_manifest_audit = manifest_audit_combined
manifest_missing = blood_gas_itemid_manifest_audit.loc[
    pd.to_numeric(

```

```

blood_gas_itemid_manifest_audit.get("present_in_dictionary"),
    errors="coerce"
).fillna(1).astype(int).eq(0)
].copy()
if not manifest_missing.empty:
    missing_repr = ", ".join(
        f"{row.source_branch}:{row.analyte}:{int(row.itemid)}" for row in manifest_missing.itertuples(index=False))
)
raise ValueError(
    "Blood-gas manifest references itemids not present in runtime
     dictionaries: "
    f"{missing_repr}"
)
blood_gas_itemid_manifest_audit_path = prior_runs_dir / f"{audit_date}"
    blood_gas_itemid_manifest_audit.csv"
blood_gas_itemid_manifest_audit.to_csv(blood_gas_itemid_manifest_audit_path,
    index=False)
print("Wrote:", blood_gas_itemid_manifest_audit_path)

```

ICU POC candidate itemids: 6

ICU POC itemid map preview:

analyte	itemid	label	category
pco2	220235	Arterial CO2 Pressure	Labs
pco2	226062	Venous CO2 Pressure	Labs
ph	220274	PH (Venous)	Labs
ph	223830	PH (Arterial)	Labs
po2	220224	Arterial O2 pressure	Labs
po2	226063	Venous O2 Pressure	Labs

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 icu_poc_itemid_map.csv

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 blood_gas_itemid_manifest_audit.csv

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
    consistent across notebook stages.
```

```
# Extract ICU chartevents within ED 0-24h window for cohort ICU stays
```

```
if len(icu_itemids) == 0:
    icu_poc_long = pd.DataFrame()
    print("No ICU POC itemids found.")
```

```

else:
    # ensure icu stay ids available
    if 'icu' not in globals():
        raise NameError("ICU stays table 'icu' not found; run ICU phase
                        first.")

icu_stays = icu[["icu_stay_id", "hadm_id", "intime"]].copy()
icu_stays = icu_stays.dropna(subset=["icu_stay_id"])

icu_poc_sql = f"""
WITH icu_stays AS (
    SELECT stay_id AS icu_stay_id, hadm_id
    FROM `{{PHYS}}.{{ICU}}.icustays`
    WHERE stay_id IN UNNEST(@icu_stay_ids)
),
eds AS (
    SELECT stay_id AS ed_stay_id, hadm_id, intime AS ed_intime
    FROM `{{PHYS}}.{{ED}}.edstays`
    WHERE hadm_id IN (SELECT hadm_id FROM icu_stays)
)
SELECT
    s.icu_stay_id,
    e.ed_stay_id,
    e.hadm_id,
    ce.charttime,
    ce.itemid,
    ce.valuenum,
    ce.value uom
FROM `{{PHYS}}.{{ICU}}.chartevents` ce
JOIN icu_stays s ON s.icu_stay_id = ce.stay_id
JOIN eds e ON e.hadm_id = s.hadm_id
WHERE ce.itemid IN UNNEST(@itemids)
    AND ce.charttime BETWEEN e.ed_intime AND TIMESTAMP_ADD(e.ed_intime,
    INTERVAL 24 HOUR)
"""

icu_poc_long = run_sql_bq(icu_poc_sql, {"icu_stay_ids"::
    icu_stays["icu_stay_id"].astype(int).tolist(), "itemids": icu_itemids})
print("ICU POC long rows:", len(icu_poc_long))

# Build panels by 5-minute bins per ICU stay
if len(icu_poc_long) > 0:
    icu_poc_long["time_bin"] = icu_poc_long["charttime"].dt.floor("5min")

```

```

def pick_from_ids(df, ids, name):
    tmp = df.loc[df["itemid"].isin(ids),
    ["icu_stay_id", "time_bin", "valuenum"]]
    tmp = tmp.rename(columns={"valuenum": name})
    return tmp.groupby(["icu_stay_id", "time_bin"],
    as_index=False).first()

pco2_ids = set(icu_matches.get("pco2", pd.DataFrame()).get("itemid",
[]).tolist())
po2_ids = set(icu_matches.get("po2", pd.DataFrame()).get("itemid",
[]).tolist())
ph_ids = set(icu_matches.get("ph", pd.DataFrame()).get("itemid",
[]).tolist())
hco3_ids = set(icu_matches.get("hco3", pd.DataFrame()).get("itemid",
[]).tolist())
lact_ids = set(icu_matches.get("lactate", pd.DataFrame()).get("itemid",
[]).tolist())

panel_poc = (
    icu_poc_long.groupby(["icu_stay_id", "time_bin"], as_index=False)
    .agg(panel_time=("charttime", "min"))
)
if pco2_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, pco2_ids,
"pco2"), on=["icu_stay_id", "time_bin"], how="left")
if po2_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, po2_ids,
"po2"), on=["icu_stay_id", "time_bin"], how="left")
if ph_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, ph_ids,
"ph"), on=["icu_stay_id", "time_bin"], how="left")
if hco3_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, hco3_ids,
"hco3"), on=["icu_stay_id", "time_bin"], how="left")
if lact_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, lact_ids,
"lactate"), on=["icu_stay_id", "time_bin"], how="left")

# map to ED stay via hadm_id
panel_poc = panel_poc.merge(icu[["icu_stay_id", "hadm_id"]],
on="icu_stay_id", how="left")

```

```

panel_poc = panel_poc.merge(ed_df[["ed_stay_id", "hadm_id", "ed_intime"]], 
↪ on="hadm_id", how="left")
panel_poc["dt_hours"] = (panel_poc["panel_time"] - 
↪ panel_poc["ed_intime"]).dt.total_seconds() / 3600.0

icu_itemid_usage = (
    icu_poc_long.groupby("itemid", as_index=False)
    .size()
    .rename(columns={"size": "row_count"})
)
icu_itemid_usage["itemid"] = pd.to_numeric(
    icu_itemid_usage["itemid"], errors="coerce"
).astype("Int64")
if not icu_itemid_map.empty:
    icu_itemid_usage = icu_itemid_usage.merge(
        icu_itemid_map[["analyte", "itemid", "label",
↪ "category"]].drop_duplicates(),
        on="itemid",
        how="left",
    )
    icu_itemid_usage = icu_itemid_usage.sort_values(
        ["row_count", "itemid"], ascending=[False, True], kind="stable"
    ).reset_index(drop=True)
    icu_itemid_usage_path = prior_runs_dir / f"{audit_date}"
    icu_poc_itemid_usage.csv"
    icu_itemid_usage.to_csv(icu_itemid_usage_path, index=False)
    print("Wrote:", icu_itemid_usage_path)
if "blood_gas_itemid_manifest_audit" in globals() and not
    ↪ blood_gas_itemid_manifest_audit.empty:
        usage_merge = icu_itemid_usage[["itemid", "row_count"]].copy()
        usage_merge["itemid"] = pd.to_numeric(usage_merge["itemid"],
↪ errors="coerce").astype("Int64")
        manifest_audit = blood_gas_itemid_manifest_audit.copy()
        manifest_audit["itemid"] = pd.to_numeric(manifest_audit["itemid"],
↪ errors="coerce").astype("Int64")
        manifest_audit = manifest_audit.merge(
            usage_merge.rename(columns={"row_count": "icu_usage_row_count"}),
            on="itemid",
            how="left",
        )
        if "cohort_row_count" in manifest_audit.columns:
            manifest_audit["cohort_row_count"] =
↪ manifest_audit["cohort_row_count"].fillna(0).astype(int)

```

```

    manifest_audit["cohort_row_count"] = (
        manifest_audit["cohort_row_count"] +
    ↵ manifest_audit["icu_usage_row_count"].fillna(0).astype(int)
    )
    manifest_audit = manifest_audit.drop(columns=["icu_usage_row_count"])
    blood_gas_itemid_manifest_audit = manifest_audit

    ↵ blood_gas_itemid_manifest_audit.to_csv(blood_gas_itemid_manifest_audit_path,
    ↵ index=False)
    print("Updated:", blood_gas_itemid_manifest_audit_path)
else:
    panel_poc = pd.DataFrame()

# Ensure expected panel_poc columns exist even if analyte is absent
for col in ["pco2", "ph", "hco3", "lactate"]:
    if col not in panel_poc.columns:
        panel_poc[col] = pd.NA

poc_inclusion_enabled = False
poc_inclusion_reason = "disabled_manifest_or_qc_fail"
pco2_itemid_qc_rows: list[dict[str, Any]] = []

if len(icu_poc_long) > 0 and pco2_ids:
    pco2_long =
    ↵ icu_poc_long.loc[icu_poc_long["itemid"].isin(pco2_ids)].copy()
    if not pco2_long.empty:
        pco2_long["label"] = pco2_long["itemid"].map(
            icu_itemid_map.set_index("itemid")["label"] if not
    ↵ icu_itemid_map.empty else {}
        )
        pco2_long["category"] = pco2_long["itemid"].map(
            icu_itemid_map.set_index("itemid")["category"] if not
    ↵ icu_itemid_map.empty else {}
        )
        pco2_long["valuenum"] = pd.to_numeric(pco2_long["valuenum"],
    ↵ errors="coerce")
        pco2_long["valueuom_norm"] = (

    ↵ pco2_long["valueuom"].astype("string").fillna("").str.strip().str.lower()
        )
        is_kpa = pco2_long["valueuom_norm"].eq("kpa")
        pco2_long.loc[is_kpa, "valuenum"] = pco2_long.loc[is_kpa, "valuenum"]
    ↵ * 7.50062

```

```

pco2_long["pco2_mmhg"] = pco2_long["valuenum"]

    for itemid_value, group in pco2_long.groupby("itemid", dropna=False):
        label_text =
        ↵ str(group["label"].dropna().astype(str).head(1).tolist()[0]) if
        ↵ group["label"].notna().any() else ""
            category_text =
        ↵ str(group["category"].dropna().astype(str).head(1).tolist()[0]) if
        ↵ group["category"].notna().any() else ""
            label_text_l = label_text.lower()
            if int(itemid_value) in set(ICU_PCO2_ABG_ITEMIDS):
                sample_type_hint = "ABG"
            elif int(itemid_value) in set(ICU_PCO2_VBG_ITEMIDS):
                sample_type_hint = "VBG"
            else:
                sample_type_hint = "UNKNOWN"
            value_units =
        ↵ sorted(set(group["valueuom_norm"].dropna().astype(str).tolist()))
            unit_ok = set(value_units).issubset(ICU_PCO2_QC_ALLOWED_UOMS)
            include_ok = any(token in label_text_l for token in
        ↵ ICU_PCO2_QC_INCLUDE_TOKENS)
            exclude_hit = any(token in label_text_l for token in
        ↵ ICU_PCO2_QC_EXCLUDE_TOKENS)
            numeric = pd.to_numeric(group["pco2_mmhg"],
        ↵ errors="coerce").dropna()
            median_value = float(numeric.median()) if len(numeric) else
        ↵ np.nan
            median_ok = bool(
                len(numeric)
                and ICU_PCO2_QC_MEDIAN_MIN <= median_value <=
                ↵ ICU_PCO2_QC_MEDIAN_MAX
            )
            contamination_flag = bool((not include_ok) or exclude_hit or
        ↵ (not unit_ok))

        pco2_itemid_qc_rows.append(
            {
                "itemid": int(itemid_value),
                "label": label_text,
                "category": category_text,
                "source_branch": "POC",
                "sample_type_hint": sample_type_hint,
                "row_count": int(len(group)),

```

```

        "non_null_count": int(len(numeric)),
        "units_observed": ";" .join(value_units),
        "p50_mmhg": median_value if len(numeric) else np.nan,
        "p95_mmhg": float(numeric.quantile(0.95)) if len(numeric)
        ↪ else np.nan,
        "max_mmhg": float(numeric.max()) if len(numeric) else
        ↪ np.nan,
        "label_include_ok": bool(include_ok),
        "label_exclude_hit": bool(exclude_hit),
        "unit_ok": bool(unit_ok),
        "median_in_expected_range": bool(median_ok),
        "contamination_flag": contamination_flag,
        "manifest_whitelisted": int(itemid_value in
        ↪ set(ICU_PCO2_ITEMIDS)),
    }
)
)

pco2_itemid_qc_columns = [
    "itemid",
    "label",
    "category",
    "source_branch",
    "sample_type_hint",
    "row_count",
    "non_null_count",
    "units_observed",
    "p50_mmhg",
    "p95_mmhg",
    "max_mmhg",
    "label_include_ok",
    "label_exclude_hit",
    "unit_ok",
    "median_in_expected_range",
    "contamination_flag",
    "manifest_whitelisted",
]
pco2_itemid_qc_audit = pd.DataFrame(pco2_itemid_qc_rows)
for column_name in pco2_itemid_qc_columns:
    if column_name not in pco2_itemid_qc_audit.columns:
        pco2_itemid_qc_audit[column_name] = pd.NA
pco2_itemid_qc_audit = pco2_itemid_qc_audit[pco2_itemid_qc_columns]
pco2_itemid_qc_audit_path = prior_runs_dir / f"{audit_date}"
↪ pco2_itemid_qc_audit.csv"

```

```

pco2_itemid_qc_audit.to_csv(pco2_itemid_qc_audit_path, index=False)
print("Wrote:", pco2_itemid_qc_audit_path)

if ICU_INCLUSION_MODE == "validated_only":
    if ICU_PC02_ITEMIDS and not pco2_itemid_qc_audit.empty:
        whitelist_mask =
        ↵ pco2_itemid_qc_audit["itemid"].isin(set(ICU_PC02_ITEMIDS))
        whitelist_rows = pco2_itemid_qc_audit.loc[whitelist_mask].copy()
        if (
            len(whitelist_rows) == len(set(ICU_PC02_ITEMIDS))
            and whitelist_rows["contamination_flag"].eq(False).all()
            and whitelist_rows["median_in_expected_range"].eq(True).all()
        ):
            poc_inclusion_enabled = True
            poc_inclusion_reason = "validated_itemids_passed_qc"
        else:
            poc_inclusion_reason = "validated_itemids_failed_qc"
    elif not ICU_PC02_ITEMIDS:
        poc_inclusion_reason = "manifest_has_no_icu_pco2_itemids"
    else:
        poc_inclusion_reason = "missing_itemid_qc_rows"
else:
    poc_inclusion_reason = "icu_inclusion_mode_disabled"

print(
    "POC itemid QC status:",
    {"enabled": int(poc_inclusion_enabled), "reason": poc_inclusion_reason},
)

```

poc_flags = pd.DataFrame(columns=["ed_stay_id", "poc_hypercap_0_24h_flag"])

if len(panel_poc) > 0 and poc_inclusion_enabled:

- p24_poc = panel_poc.loc[panel_poc["dt_hours"].between(0,24,
- ↳ inclusive="both")].copy()
- p24_poc.loc[:, "poc_hypercap_0_24h_flag"] = (p24_poc["pco2"] >=
- ↳ 50).astype("Int64")

poc_flags = p24_poc.groupby("ed_stay_id", as_index=False).agg(
 poc_hypercap_0_24h_flag=("poc_hypercap_0_24h_flag", "max")
)

poc_hypercap_0_24h_edstay_n = (
 int(
 pd.to_numeric(

```

        poc_flags.get("poc_hypercap_0_24h_flag",
    ↵ pd.Series(dtype="float64")),
        errors="coerce",
    ↵ )
    .fillna(0)
    .astype(int)
    .sum()
)
if len(poc_flags) > 0
else 0
)
poc_hypercap_0_24h_edstay_rate = float(
    poc_hypercap_0_24h_edstay_n / max(int(ed_df["ed_stay_id"].nunique()), 1)
)

# Optional fallback for first_hco3 from ICU POC only when explicit manifest
↪ itemids are provided.
if ICU_HC03_ITEMIDS and len(panel_poc) > 0:
    poc_hco3_candidates = panel_poc.loc[
        panel_poc["dt_hours"].between(0, 24, inclusive="both")
    ].copy()
    poc_hco3_candidates["hco3"] = pd.to_numeric(poc_hco3_candidates["hco3"],
↪ errors="coerce")
    poc_hco3_candidates =
↪ poc_hco3_candidates.loc[poc_hco3_candidates["hco3"].notna()].copy()
    if not poc_hco3_candidates.empty:
        first_poc_hco3 = (
            poc_hco3_candidates.sort_values(
                ["ed_stay_id", "dt_hours", "panel_time"],
                kind="stable",
            )
            .groupby("ed_stay_id", as_index=False)
            .first()[["ed_stay_id", "hco3"]]
            .rename(columns={"hco3": "first_hco3_poc_fallback"})
        )
        ed_df = ed_df.merge(first_poc_hco3, on="ed_stay_id", how="left")
        hco3_missing_mask = pd.to_numeric(ed_df.get("first_hco3"),
↪ errors="coerce").isna()
        hco3_fill_mask = hco3_missing_mask & pd.to_numeric(
            ed_df["first_hco3_poc_fallback"], errors="coerce"
        ).notna()
        ed_df.loc[hco3_fill_mask, "first_hco3"] = ed_df.loc[
            hco3_fill_mask, "first_hco3_poc_fallback"

```

```

        ]
    if "first_hco3_source" not in ed_df.columns:
        ed_df["first_hco3_source"] = "missing"
    ed_df.loc[hco3_fill_mask, "first_hco3_source"] =
    "poc_explicit_itemid_fallback"
    ed_df = ed_df.drop(columns=["first_hco3_poc_fallback"])

# optional export
if len(panel_poc) > 0:
    from datetime import datetime

    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    poc_path = prior_runs_dir /
    f"gas_panels_poc_{datetime.now().strftime('%Y%m%d_%H%M%S')}.parquet"
    panel_poc.to_parquet(poc_path, index=False)
    print("Wrote:", poc_path)

```

ICU POC long rows: 52791
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects CC-NLP/MIMIC tabular data/prior runs/2026-02-23 icu_poc_itemid_usage.csv
Updated: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects CC-NLP/MIMIC tabular data/prior runs/2026-02-23 blood_gas_itemid_manifest_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects CC-NLP/MIMIC tabular data/prior runs/2026-02-23 pco2_itemid_qc_audit.csv
POC itemid QC status: {'enabled': 1, 'reason': 'validated_itemids_passed_qc'}
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects CC-NLP/MIMIC tabular data/prior runs/gas_panels_poc_20260223_103325.parquet

3.0.8 Phase 6 — BMI/anthropometrics (OMR)

Rationale: Add BMI/height/weight closest to ED presentation when available.

```

# Purpose: Attach closest anthropometrics with pre-ED preference and bounded
    post-ED fallback.

from datetime import datetime

omr_sql = f"""
SELECT subject_id, chartdate, result_name, result_value
FROM `{{PHYS}}.{{HOSP}}.omr`
WHERE REGEXP_CONTAINS(

```

```

    LOWER(result_name),
    r'(bmi|body\\s*mass\\s*index|height|weight)'
)
"""

allow_omr_query_failure = os.getenv("COHORT_ALLOW_OMR_QUERY_FAILURE",
    ↪ "0").strip() == "1"
try:
    omr_raw = run_sql_bq(omr_sql)
except Exception as exc:
    if allow_omr_query_failure:
        print("OMR query failed; continuing with empty OMR frame because
            ↪ COHORT_ALLOW_OMR_QUERY_FAILURE=1:", exc)
        omr_raw = pd.DataFrame(columns=["subject_id", "chartdate",
    ↪ "result_name", "result_value"])
    else:
        raise

print("OMR rows:", len(omr_raw))
omr_prepared = prepare_omr_records(omr_raw)
ed_df, omr_diagnostics = attach_closest_pre_ed_omr(ed_df, omr_prepared,
    ↪ window_days=365)
omr_diagnostics["source_rows_raw"] = int(len(omr_raw))
omr_diagnostics["prepared_rows"] = int(len(omr_prepared))

diag_table = (
    pd.DataFrame({"metric": list(omr_diagnostics.keys()), "value":
    ↪ list(omr_diagnostics.values())})
    .assign(value=lambda d: d["value"].astype(str))
)
print(diag_table)

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")
omr_diag_path = prior_runs_dir / f"{out_date} omr_diagnostics.json"
omr_diag_path.write_text(json.dumps(omr_diagnostics, indent=2))
print("Wrote:", omr_diag_path)

# Optional charted anthropometric fallback (nearest-anytime selection).
use_chARTED_fallback = os.getenv("COHORT_ANTHRO_CHARTED_FALLBACK",
    ↪ "1").strip() == "1"

```

```

use_nearest_anytime = os.getenv("COHORT_ANTHRO_NEAREST_ANYTIME", "1").strip()
↳ == "1"
charted_anthro_diagnostics = {
    "enabled": bool(use_chartedFallback),
    "nearest_anytime": bool(use_nearestAnytime),
    "charted_rows_input": 0,
    "charted_rows_usable": 0,
    "selected_counts": {},
    "rows_with_any_selected_metric": 0,
}

def _extract_ed_chARTED_anthro(hadms: list[int]) -> tuple[pd.DataFrame,
↳ dict[str, Any]]:
    """Extract ED-documented anthropometrics when present in ED schema."""
    empty = pd.DataFrame(
        columns=[
            "subject_id",
            "obs_time",
            "result_name",
            "result_value_num",
            "result_unit_raw",
            "source",
        ]
    )
    diagnostics: dict[str, Any] = {
        "metadata_query_ok": False,
        "metadata_error": None,
        "triage_candidate_columns": [],
        "vitalsign_candidate_columns": [],
        "triage_rows_raw": 0,
        "vitalsign_rows_raw": 0,
        "rows_long": 0,
        "result_counts": {},
    }
    if not hadms:
        return empty, diagnostics

    try:
        ed_columns = run_sql_bq(
            f"""
            SELECT
                LOWER(table_name) AS table_name,
                LOWER(column_name) AS column_name

```

```

        FROM `.{PHYS}.{ED}.INFORMATION_SCHEMA.COLUMNS`
        WHERE LOWER(table_name) IN ('triage', 'vitalsign')
        """
    )
    diagnostics["metadata_query_ok"] = True
except Exception as exc:
    diagnostics["metadata_error"] = str(exc)
return empty, diagnostics

table_columns: dict[str, set[str]] = {}
for row in ed_columns.itertuples(index=False):
    table_name = str(getattr(row, "table_name", "")).strip().lower()
    column_name = str(getattr(row, "column_name", "")).strip().lower()
    if not table_name or not column_name:
        continue
    table_columns.setdefault(table_name, set()).add(column_name)

    diagnostics["triage_candidate_columns"] =
    ↪ sorted(table_columns.get("triage", set()))
    diagnostics["vitalsign_candidate_columns"] = sorted(
        table_columns.get("vitalsign", set()))
)

def _choose_numeric_expr(
    alias: str, available: set[str], candidates: tuple[str, ...]
) -> str:
    for column_name in candidates:
        if column_name in available:
            return f"SAFE_CAST({alias}.{column_name} AS FLOAT64)"
    return "CAST(NULL AS FLOAT64)"

def _choose_obs_expr(alias: str, available: set[str]) -> str:
    for column_name in ("charttime", "storetime"):
        if column_name in available:
            return f"{alias}.{column_name}"
    return "s.intime"

def _choose_unit_expr(alias: str, available: set[str], candidates:
    ↪ tuple[str, ...]) -> str:
    for column_name in candidates:
        if column_name in available:
            return f"CAST({alias}.{column_name} AS STRING)"
    return "CAST(NULL AS STRING)"

```

```

def _query_table(table_name: str, source_name: str) -> pd.DataFrame:
    available = table_columns.get(table_name, set())
    join_key = next(
        (key_name for key_name in ("stay_id", "ed_stay_id") if key_name
    ↵ in available),
        None,
    )
    value_cols_present = any(
        column_name in available
        for column_name in ("weight", "wt", "height", "ht", "bmi",
        ↵ "body_mass_index")
    )
    if join_key is None or not value_cols_present:
        return pd.DataFrame(
            columns=[
                "subject_id",
                "obs_time",
                "weight_value",
                "weight_unit_raw",
                "height_value",
                "height_unit_raw",
                "bmi_value",
                "bmi_unit_raw",
                "source",
            ]
        )
    alias = "x"
    obs_expr = _choose_obs_expr(alias, available)
    weight_expr = _choose_numeric_expr(alias, available, ("weight",
    ↵ "wt"))
    height_expr = _choose_numeric_expr(alias, available, ("height",
    ↵ "ht"))
    bmi_expr = _choose_numeric_expr(alias, available, ("bmi",
    ↵ "body_mass_index"))
    weight_unit_expr = _choose_unit_expr(
        alias,
        available,
        ("weight_uom", "weight_unit", "wt_uom", "wt_unit"),
    )
    height_unit_expr = _choose_unit_expr(
        alias,

```

```

        available,
        ("height_uom", "height_unit", "ht_uom", "ht_unit"),
    )
bmi_unit_expr = _choose_unit_expr(
    alias,
    available,
    ("bmi_uom", "bmi_unit", "body_mass_index_uom",
    "body_mass_index_unit"),
)
sql_text = f"""
WITH hadms AS (
    SELECT x AS hadm_id
    FROM UNNEST(@hadms) AS x
)
SELECT
    s.subject_id,
    {obs_expr} AS obs_time,
    {weight_expr} AS weight_value,
    {weight_unit_expr} AS weight_unit_raw,
    {height_expr} AS height_value,
    {height_unit_expr} AS height_unit_raw,
    {bmi_expr} AS bmi_value,
    {bmi_unit_expr} AS bmi_unit_raw,
    '{source_name}' AS source
FROM `{{PHYS}}.{{ED}}.{{table_name}}` {{alias}}
JOIN `{{PHYS}}.{{ED}}.edstays` s
    ON s.stay_id = {{alias}}.{{join_key}}
JOIN hadms h
    ON h.hadm_id = s.hadm_id
"""
return run_sql_bq(sql_text, {"hadms": hadms})

triage_wide = _query_table("triage", "ed_charted_triage")
vitalsign_wide = _query_table("vitalsign", "ed_charted_vitalsign")
diagnostics["triage_rows_raw"] = int(len(triage_wide))
diagnostics["vitalsign_rows_raw"] = int(len(vitalsign_wide))

def _to_long(frame: pd.DataFrame) -> pd.DataFrame:
    if frame.empty:
        return empty.copy()
    pieces: list[pd.DataFrame] = []
    for result_name, value_col, unit_col in (
        ("weight", "weight_value", "weight_unit_raw"),

```

```

        ("height", "height_value", "height_unit_raw"),
        ("bmi", "bmi_value", "bmi_unit_raw"),
    ):
        if value_col not in frame.columns:
            continue
        numeric = pd.to_numeric(frame[value_col], errors="coerce")
        mask = numeric.notna()
        if not bool(mask.any()):
            continue
        part = frame.loc[mask, ["subject_id", "obs_time",
        ↵ "source"]].copy()
        part["result_name"] = result_name
        part["result_value_num"] = numeric.loc[mask]
        part["result_unit_raw"] = (
            frame.loc[mask, unit_col].astype("string").replace({"":
        ↵ pd.NA})
        if unit_col in frame.columns
        else pd.NA
    )
        pieces.append(part)
    if not pieces:
        return empty.copy()
    combined = pd.concat(pieces, ignore_index=True, sort=False)
    combined["subject_id"] = pd.to_numeric(
        combined["subject_id"], errors="coerce"
    ).astype("Int64")
    combined["obs_time"] = pd.to_datetime(combined["obs_time"],
    ↵ errors="coerce")
    combined = combined.loc[
        combined["subject_id"].notna()
        & combined["obs_time"].notna()
        & combined["result_value_num"].notna()
    ].copy()
    return combined[
        [
            "subject_id",
            "obs_time",
            "result_name",
            "result_value_num",
            "result_unit_raw",
            "source",
        ]
    ]

```

```

ed_long = pd.concat(
    [_to_long(triage_wide), _to_long(vitalsign_wide)],
    ignore_index=True,
    sort=False,
)
diagnostics["rows_long"] = int(len(ed_long))
diagnostics["result_counts"] = {
    str(key): int(value)
    for key, value in
        ed_long["result_name"].value_counts(dropna=False).items()
}
return ed_long, diagnostics

if use_chARTed_fALLBACK and len(hadm_list) > 0:
    charted_anthro_sql = f"""
    WITH hadms AS (
        SELECT x AS hadm_id
        FROM UNNEST(@hadms) AS x
    ),
    candidate_items AS (
        SELECT itemid, label, category
        FROM `{{PHYS}}.{{ICU}}.d_items`
        WHERE REGEXP_CONTAINS(LOWER(label),
        r'(height|weight|admission\\s+weight)')
    ),
    ce AS (
        SELECT
            ce.subject_id,
            ce.hadm_id,
            ce.charttime AS obs_time,
            ci.label,
            ce.valuenum,
            CAST(ce.valueuom AS STRING) AS valueuom
        FROM `{{PHYS}}.{{ICU}}.chartevents` ce
        JOIN hadms h
        ON ce.hadm_id = h.hadm_id
        JOIN candidate_items ci
        ON ce.itemid = ci.itemid
        WHERE ce.valuenum IS NOT NULL
    )
    SELECT
        subject_id,

```

```

obs_time,
CASE
    WHEN REGEXP_CONTAINS(LOWER(label), r'height') THEN 'height'
    WHEN REGEXP_CONTAINS(LOWER(label), r'weight') THEN 'weight'
    ELSE NULL
END AS result_name,
valuenum AS result_value_num,
value uom AS result_unit_raw,
'ICU' AS source
FROM ce
WHERE REGEXP_CONTAINS(LOWER(label), r'(height|weight)')
"""

icu_charted_raw = run_sql_bq(charted_anthro_sql, {"hadms": hadm_list})
ed_charted_raw, ed_charted_diagnostics =
    _extract_ed_chARTED_anthro(hadm_list)
omr_candidates = omr_prepared[
    [
        "subject_id",
        "obs_time",
        "result_name",
        "result_value_num",
        "result_unit_raw",
        "source",
    ]
].copy()
charted_raw = pd.concat(
    [omr_candidates, icu_charted_raw, ed_charted_raw],
    ignore_index=True,
    sort=False,
)
ed_df, charted_anthro_diagnostics = attach_chARTED_anthro_fallback(
    ed_df,
    charted_raw,
    nearest_anytime=use_nearest_anytime,
    source_preference=("ED", "ICU", "HOSPITAL"),
)
weight_col = ANTHRO_VALUE_COLUMNS["weight"]
height_col = ANTHRO_VALUE_COLUMNS["height"]
bmi_col = ANTHRO_VALUE_COLUMNS["bmi"]
weight_time_col = ANTHRO_TIME_COLUMNS["weight"]
height_time_col = ANTHRO_TIME_COLUMNS["height"]
bmi_time_col = ANTHRO_TIME_COLUMNS["bmi"]

```

```

weight_uom_col = ANTHRO_UNIT_COLUMNS["weight"]
height_uom_col = ANTHRO_UNIT_COLUMNS["height"]
bmi_uom_col = ANTHRO_UNIT_COLUMNS["bmi"]
bmi_source_col = "__anthro_bmi_source"
weight_source_col = "__anthro_weight_source"
height_source_col = "__anthro_height_source"

height_cm = pd.to_numeric(ed_df.get(height_col), errors="coerce")
weight_kg = pd.to_numeric(ed_df.get(weight_col), errors="coerce")
bmi_existing = pd.to_numeric(ed_df.get(bmi_col), errors="coerce")
charted_anthro_diagnostics["bmi_missing_before_pair_fill"] =
    int(bmi_existing.isna().sum())
height_time = pd.to_datetime(ed_df.get(height_time_col), errors="coerce")
weight_time = pd.to_datetime(ed_df.get(weight_time_col), errors="coerce")
pair_gap_hours = ((weight_time - height_time).abs().dt.total_seconds() /
    3600.0).astype("float64")
pair_has_time = height_time.notna() & weight_time.notna()
pair_within_window = pair_has_time &
    pair_gap_hours.le(ANTHRO_BMI_PAIR_WINDOW_HOURS)
pair_outside_window = pair_has_time & ~pair_within_window
pair_missing_time = (~pair_has_time) & (height_cm.notna() &
    weight_kg.notna())
bmi_calc = weight_kg / (height_cm / 100.0) ** 2
bmi_fill_mask = (
    bmi_existing.isna()
    & bmi_calc.notna()
    & height_cm.gt(0)
    & weight_kg.gt(0)
    & pair_within_window
    & ed_df.get(weight_uom_col).astype("string") [
        ].eq(ANTHRO_CANONICAL_UNITS["weight"])
    & ed_df.get(height_uom_col).astype("string") [
        ].eq(ANTHRO_CANONICAL_UNITS["height"]))
ed_df.loc[bmi_fill_mask, bmi_col] = bmi_calc[bmi_fill_mask]
ed_df.loc[bmi_fill_mask, bmi_uom_col] = ANTHRO_CANONICAL_UNITS["bmi"]
ed_df.loc[bmi_fill_mask, bmi_time_col] = pd.to_datetime(
    ed_df.loc[bmi_fill_mask, weight_time_col], errors="coerce")
)
if bmi_source_col in ed_df.columns:
    ed_df.loc[bmi_fill_mask, bmi_source_col] = (
        ed_df.loc[bmi_fill_mask, weight_source_col]
        .astype("string"))

```

```

        .fillna(ed_df.loc[bmi_fill_mask,
↪ height_source_col].astype("string"))
            .map(normalize_anthro_source)
        )

bmi_direct_source = (
    ed_df.get(bmi_source_col, pd.Series("missing", index=ed_df.index))
        .astype("string")
        .map(normalize_anthro_source)
)
weight_source = (
    ed_df.get(weight_source_col, pd.Series("missing", index=ed_df.index))
        .astype("string")
        .map(normalize_anthro_source)
)
height_source = (
    ed_df.get(height_source_col, pd.Series("missing", index=ed_df.index))
        .astype("string")
        .map(normalize_anthro_source)
)
anthro_source = pd.Series("missing", index=ed_df.index, dtype="string")
direct_bmi_mask = pd.to_numeric(ed_df[bmi_col], errors="coerce").notna()
↪ & bmi_direct_source.ne("missing")
anthro_source.loc[direct_bmi_mask] =
↪ bmi_direct_source.loc[direct_bmi_mask]
computed_bmi_mask = bmi_fill_mask & anthro_source.eq("missing")
anthro_source.loc[computed_bmi_mask] =
    weight_source.loc[computed_bmi_mask]
        .where(weight_source.loc[computed_bmi_mask].ne("missing"),
↪ height_source.loc[computed_bmi_mask])
            .fillna("missing")
)
fallback_weight_mask = anthro_source.eq("missing") &
↪ weight_source.ne("missing")
anthro_source.loc[fallback_weight_mask] =
↪ weight_source.loc[fallback_weight_mask]
fallback_height_mask = anthro_source.eq("missing") &
↪ height_source.ne("missing")
anthro_source.loc[fallback_height_mask] =
↪ height_source.loc[fallback_height_mask]
ed_df["anthro_source"] = anthro_source.map(normalize_anthro_source)

```

```

anthro_obstime = pd.Series(pd.NaT, index=ed_df.index,
                           dtype="datetime64[ns]")
bmi_time = pd.to_datetime(ed_df.get(bmi_time_col), errors="coerce")
anthro_obstime.loc[direct_bmi_mask] = bmi_time.loc[direct_bmi_mask]
anthro_obstime.loc[computed_bmi_mask] =
    weight_time.loc[computed_bmi_mask]
anthro_obstime.loc[fallback_weight_mask] =
    weight_time.loc[fallback_weight_mask]
anthro_obstime.loc[fallback_height_mask] =
    height_time.loc[fallback_height_mask]
ed_df["anthro_obstime"] = pd.to_datetime(anthro_obstime, errors="coerce")
ed_df["anthro_hours_offset"] = (
    (ed_df["anthro_obstime"] - pd.to_datetime(ed_df["ed_intime"]),
     errors="coerce"))
    .dt.total_seconds()
    / 3600.0
)
ed_df["anthro_timing_basis"] = "missing"
ed_df.loc[ed_df["anthro_source"].isin({"ED", "ICU"}) &
         ed_df["anthro_obstime"].notna(), "anthro_timing_basis"] =
    "nearest_anytime"
ed_df.loc[
    ed_df["anthro_source"].eq("HOSPITAL") & ed_df["anthro_obstime"].notna() & pd.to_numeric(ed_df["anthro_hours_offset"], errors="coerce").le(0),
    "anthro_timing_basis",
] = "pre"
ed_df.loc[
    ed_df["anthro_source"].eq("HOSPITAL") & ed_df["anthro_obstime"].notna() & pd.to_numeric(ed_df["anthro_hours_offset"], errors="coerce").gt(0),
    "anthro_timing_basis",
] = "post"
ed_df["anthro_timing_uncertain"] = pd.Series(pd.NA, index=ed_df.index,
                                             dtype="boolean")
ed_df.loc[ed_df["anthro_timing_basis"].eq("pre"),
         "anthro_timing_uncertain"] = False
ed_df.loc[
    ed_df["anthro_timing_basis"].isin({"post", "nearest_anytime"}),
    "anthro_timing_uncertain",
] = True
ed_df["anthro_timing_tier"] = "missing"

```

```

within_365d = pd.to_numeric(ed_df["anthro_hours_offset"], 
↪ errors="coerce").abs().le(24.0 * 365.0)
ed_df.loc[
    ed_df["anthro_source"].eq("HOSPITAL")
    & within_365d
    & pd.to_numeric(ed_df["anthro_hours_offset"], errors="coerce").le(0),
    "anthro_timing_tier",
] = "pre_ed_365"
ed_df.loc[
    ed_df["anthro_source"].eq("HOSPITAL")
    & within_365d
    & pd.to_numeric(ed_df["anthro_hours_offset"], errors="coerce").gt(0),
    "anthro_timing_tier",
] = "post_ed_365"
ed_df["anthro_chardate"] = pd.to_datetime(ed_df["anthro_obstime"],
↪ errors="coerce").dt.floor("D")
ed_df["anthro_days_offset"] = (
    pd.to_datetime(ed_df["ed_intime"], errors="coerce").dt.floor("D") -
↪ ed_df["anthro_chardate"]
).dt.days.astype("Int64")

charted_anthro_diagnostics["bmi_filled_from_height_weight"] =
↪ int(bmi_fill_mask.sum())
charted_anthro_diagnostics["bmi_pair_within_7d_n"] =
↪ int(pair_within_window.sum())
charted_anthro_diagnostics["bmi_pair_outside_7d_n"] =
↪ int(pair_outside_window.sum())
charted_anthro_diagnostics["bmi_pair_missing_time_n"] =
↪ int(pair_missing_time.sum())
valid_pair_gaps = pair_gap_hours.loc[pair_has_time]
if not valid_pair_gaps.empty:
    charted_anthro_diagnostics["bmi_pair_gap_hours_quantiles"] = {
        "p50": float(valid_pair_gaps.quantile(0.50)),
        "p95": float(valid_pair_gaps.quantile(0.95)),
        "max": float(valid_pair_gaps.max()),
    }
charted_anthro_diagnostics["source_input_counts"] = {
    str(key): int(value)
    for key, value in charted_raw.get("source",
        ↪ pd.Series(dtype="string"))
    .astype("string")
    .fillna("missing")
    .value_counts(dropna=False)
}

```

```

        .items()
    }
charted_anthro_diagnostics["icu_charted_rows_input"] =
    int(len(icu_charted_raw))
charted_anthro_diagnostics["ed_charted_rows_input"] =
    int(len(ed_charted_raw))
charted_anthro_diagnostics["ed_charted"] = ed_charted_diagnostics
charted_anthro_diagnostics["enabled"] = True
charted_anthro_diagnostics["nearest_anytime"] = bool(use_nearest_anytime)

anthro_model_specs = {
    "bmi_closest_pre_ed": ("bmi_closest_pre_ed_model", "bmi_outlier_flag",
                           ← 0.0, 100.0),
    "height_closest_pre_ed": ("height_closest_pre_ed_model",
                               ← "height_outlier_flag", 100.0, 250.0),
    "weight_closest_pre_ed": ("weight_closest_pre_ed_model",
                               ← "weight_outlier_flag", 0.0, 400.0),
}
anthro_audit_rows: list[dict[str, Any]] = []
for raw_col, (model_col, flag_col, lower_bound, upper_bound) in
    ↵ anthro_model_specs.items():
    raw_numeric = pd.to_numeric(ed_df.get(raw_col), errors="coerce")
    outlier_mask = raw_numeric.notna() & (
        raw_numeric.le(lower_bound) | raw_numeric.gt(upper_bound)
    )
    ed_df[flag_col] = outlier_mask
    ed_df[model_col] = raw_numeric.mask(outlier_mask)
    anthro_audit_rows.append(
        {
            "raw_column": raw_col,
            "model_column": model_col,
            "lower_bound_exclusive": lower_bound,
            "upper_bound_inclusive": upper_bound,
            "raw_non_missing_n": int(raw_numeric.notna().sum()),
            "outlier_n": int(outlier_mask.sum()),
            "model_non_missing_n": int(pd.to_numeric(ed_df[model_col],
                           ← errors="coerce").notna().sum()),
        }
    )
}

anthropometric_cleaning_audit = pd.DataFrame(anthro_audit_rows)
anthropometric_cleaning_audit_path = (
    prior_runs_dir / f"{out_date} anthropometric_cleaning_audit.csv"
)

```

```

)
anthropometric_cleaning_audit.to_csv(anthropometric_cleaning_audit_path,
    ↪ index=False)
print("Wrote:", anthropometric_cleaning_audit_path)

anthro_coverage_audit = build_anthro_coverage_audit(ed_df)
anthro_coverage_audit["omr"] = omr_diagnostics
anthro_coverage_audit["charted_fallback"] = charted_anthro_diagnostics
anthro_coverage_audit["cleaning_audit_path"] =
    ↪ str(anthropometric_cleaning_audit_path)
anthro_coverage_audit_path = prior_runs_dir / f"{out_date}"
    ↪ anthropometrics_coverage_audit.json"
anthro_coverage_audit_path.write_text(json.dumps(anthro_coverage_audit,
    ↪ indent=2))
print("Wrote:", anthro_coverage_audit_path)

attached_non_null = omr_diagnostics.get("attached_non_null_counts", {})
attached_total = int(sum(int(v) for v in attached_non_null.values())) if
    ↪ attached_non_null else 0
subject_overlap = int(omr_diagnostics.get("subject_overlap_count", 0))
pre_window_rows = int(omr_diagnostics.get("pre_window_candidate_rows", 0))
post_window_rows = int(omr_diagnostics.get("post_window_candidate_rows", 0))
within_window_rows = int(omr_diagnostics.get("within_window_candidate_rows",
    ↪ 0))
selected_tier_counts = omr_diagnostics.get("selected_tier_counts", {})
selected_tier_rates = omr_diagnostics.get("selected_tier_rates", {})
allow_empty_omr = os.getenv("COHORT_ALLOW_EMPTY_OMR", "0").strip() == "1"
fail_on_omr_attach_inconsistency = (
    os.getenv("COHORT_FAIL_ON_OMR_ATTACH_INCONSISTENCY", "1").strip() == "1"
)

omr_window_diagnostics = {
    "subject_overlap_count": subject_overlap,
    "pre_window_candidate_rows": pre_window_rows,
    "post_window_candidate_rows": post_window_rows,
    "closest_absolute_candidate_rows":
        ↪ int(omr_diagnostics.get("closest_absolute_candidate_rows", 0)),
    "within_window_candidate_rows": within_window_rows,
    "attached_total_non_null_values": attached_total,
    "selected_tier_counts": selected_tier_counts,
    "selected_tier_rates": selected_tier_rates,
}

```

```

if len(omr_raw) > 0 and subject_overlap > 0:
    if pre_window_rows == 0 and post_window_rows > 0:
        print(
            "INFO: OMR subject overlap exists with no pre-ED candidates but
            ↵ post-ED candidates within the 365-day window. "
            "Anthropometric fallback will use post-ED observations and mark
            ↵ them as timing-uncertain."
        )
    elif pre_window_rows == 0 and post_window_rows == 0:
        print(
            "INFO: OMR subject overlap exists but no candidates are within
            ↵ +/-365 days of ED arrival.
            "Anthropometrics will remain missing for this run."
        )
if attached_total == 0 and (pre_window_rows > 0 or post_window_rows > 0):
    msg = (
        "OMR attachment produced all-null outputs despite available
        ↵ +/-365-day candidates.
        "Set COHORT_ALLOW_EMPTY_OMR=1 or
        ↵ COHORT_FAIL_ON_OMR_ATTACH_INCONSISTENCY=0 to bypass this
        ↵ guard."
    )
    if allow_empty.omr or not fail_on.omr_attach_inconsistency:
        print("WARNING:", msg)
    else:
        raise ValueError(msg)

```

OMR rows: 4908307

	metric \
0	window_days
1	source_rows_prepared
2	parsed_value_rows
3	ed_rows
4	ed_rows_eligible_for_join
5	subject_overlap_count
6	candidate_rows_after_subject_join
7	days_before_min
8	days_before_max
9	nonnegative_candidate_rows
10	pre_window_candidate_rows
11	post_window_candidate_rows
12	closest_absolute_candidate_rows
13	within_window_candidate_rows

```

14 eligible_ed_stays_with_candidates
15     attached_non_null_counts
16     attached_any_non_null_rows
17     selected_tier_counts
18     selected_tier_rates
19     timing_uncertain_count
20     anthro_source_counts
21     source_rows_raw
22     prepared_rows

                                value
0                           365
1                          4908307
2                          4908301
3                           7166
4                           7166
5                           3919
6                          134240
7                          -4034
8                           4721
9                          91549
10                         23777
11                         19595
12                         43372
13                         23777
14                         4869
15 {'bmi_closest_pre_ed': 4525, 'height_closest_p...
16                               4869
17 {'pre_ed_365': 3651, 'post_ed_365': 1218, 'mis...
18 {'pre_ed_365': 0.5094892548144013, 'post_ed_36...
19                               1218
20 {'HOSPITAL': 4869, 'missing': 2297}
21                         4908307
22                         4908307

```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-23 omr_diagnostics.json

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-23 anthropometric_cleaning_audit.csv

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-23 anthropometrics_coverage_audit.json

3.0.9 Phase 7 — ICD comorbidity flags

Rationale: Derive comorbidity indicators from index admission ICD codes for stratified analyses.

```
# Purpose: Compute comorbidity flags from hospital + ED ICD codes with
#           ↴ chunked hadm queries.

# ICD code pulls for comorbidity flags (hospital + ED; combined OR)
# NOTE: Use prefix filters to reduce CPU and avoid regex-heavy scans.

FLAGS = [
    "flag_copd", "flag_osa_ohs", "flag_chf", "flag_neuromuscular",
    "flag_opioid_substance", "flag_pneumonia",
]

SQL["icd_flags_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
icd AS (
    SELECT
        hadm_id,
        UPPER(REPLACE(icd_code, ".", "")) AS code_norm
    FROM `{{PHYS}}.{{HOSP}}.diagnoses_icd`
    WHERE hadm_id IN (SELECT hadm_id FROM hadms)
    AND (
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J43") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J44") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G473") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "E662") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "I50") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G12") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G70") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G71") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "F11") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "T40") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "F13") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J12") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J13") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J14") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J15") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J16") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J17") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J18")
    )
)
```

```

)
SELECT
    hadm_id,
    MAX(IF(STARTS_WITH(code_norm, "J43") OR STARTS_WITH(code_norm, "J44"), 1,
    ↵  0)) AS flag_copd,
    MAX(IF(STARTS_WITH(code_norm, "G473") OR STARTS_WITH(code_norm, "E662"), 1,
    ↵  0)) AS flag_osa_ohs,
    MAX(IF(STARTS_WITH(code_norm, "I50"), 1, 0)) AS flag_chf,
    MAX(IF(STARTS_WITH(code_norm, "G12") OR STARTS_WITH(code_norm, "G70") OR
    ↵  STARTS_WITH(code_norm, "G71"), 1, 0)) AS flag_neuromuscular,
    MAX(IF(STARTS_WITH(code_norm, "F11") OR STARTS_WITH(code_norm, "T40") OR
    ↵  STARTS_WITH(code_norm, "F13"), 1, 0)) AS flag_opioid_substance,
    MAX(IF(STARTS_WITH(code_norm, "J12") OR STARTS_WITH(code_norm, "J13") OR
    ↵  STARTS_WITH(code_norm, "J14") OR STARTS_WITH(code_norm, "J15") OR
    ↵  STARTS_WITH(code_norm, "J16") OR STARTS_WITH(code_norm, "J17") OR
    ↵  STARTS_WITH(code_norm, "J18"), 1, 0)) AS flag_pneumonia
FROM icd
GROUP BY hadm_id
"""

SQL["ed_icd_flags_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
ed_dx AS (
    SELECT
        s.hadm_id,
        UPPER(REPLACE(d.icd_code, ".", "")) AS code_norm
    FROM `{{PHYS}}.{{ED}}.diagnosis` d
    JOIN `{{PHYS}}.{{ED}}.edstays` s
        ON s.stay_id = d.stay_id
    WHERE s.hadm_id IN (SELECT hadm_id FROM hadms)
    AND (
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J43") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J44") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G473") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "E662") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "I50") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G12") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G70") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G71") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "F11") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "T40") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "F13") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J12") OR

```

```

        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J13") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J14") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J15") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J16") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J17") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J18")
    )
)
SELECT
    hadm_id,
    MAX(IF(STARTS_WITH(code_norm, "J43") OR STARTS_WITH(code_norm, "J44"), 1,
    0)) AS flag_copd,
    MAX(IF(STARTS_WITH(code_norm, "G473") OR STARTS_WITH(code_norm, "E662"), 1,
    0)) AS flag_osa_ohs,
    MAX(IF(STARTS_WITH(code_norm, "I50"), 1, 0)) AS flag_chf,
    MAX(IF(STARTS_WITH(code_norm, "G12") OR STARTS_WITH(code_norm, "G70") OR
    STARTS_WITH(code_norm, "G71"), 1, 0)) AS flag_neuromuscular,
    MAX(IF(STARTS_WITH(code_norm, "F11") OR STARTS_WITH(code_norm, "T40") OR
    STARTS_WITH(code_norm, "F13"), 1, 0)) AS flag_opioid_substance,
    MAX(IF(STARTS_WITH(code_norm, "J12") OR STARTS_WITH(code_norm, "J13") OR
    STARTS_WITH(code_norm, "J14") OR STARTS_WITH(code_norm, "J15") OR
    STARTS_WITH(code_norm, "J16") OR STARTS_WITH(code_norm, "J17") OR
    STARTS_WITH(code_norm, "J18"), 1, 0)) AS flag_pneumonia
FROM ed_dx
GROUP BY hadm_id
"""

def _run_flag_query_chunked(sql_name: str, label: str) -> pd.DataFrame:
    parts = []
    hadm_chunk_size = 10000
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
        part = run_sql_bq(sql(sql_name), {"hadms": hadm_chunk})
        if len(part) > 0:
            parts.append(part)
        print(f"{label} chunk {i // hadm_chunk_size + 1}: rows={len(part)}")

    if parts:
        out = pd.concat(parts, ignore_index=True)
        out = out.groupby("hadm_id", as_index=False)[FLAGS].max()
        return out

    return pd.DataFrame(columns=["hadm_id"] + FLAGS)

```

```

flag_hosp = _run_flag_query_chunked("icd_flags_sql", "Hosp ICD flags")
flag_ed = _run_flag_query_chunked("ed_icd_flags_sql", "ED ICD flags")

flag_hosp = flag_hosp.rename(columns={k: f"{k}_hosp" for k in FLAGS})
flag_ed = flag_ed.rename(columns={k: f"{k}_ed" for k in FLAGS})

flag_df = flag_hosp.merge(flag_ed, on="hadm_id", how="outer")
for k in FLAGS:
    hosp_col = f"{k}_hosp"
    ed_col = f"{k}_ed"
    if hosp_col not in flag_df.columns:
        flag_df[hosp_col] = 0
    if ed_col not in flag_df.columns:
        flag_df[ed_col] = 0
    flag_df[k] = ((flag_df[hosp_col].fillna(0).astype(int) == 1) |
    (flag_df[ed_col].fillna(0).astype(int) == 1)).astype(int)

# Merge into admission-level and ED-stay-level frames (override if present)
for _df_name in ["df", "ed_df"]:
    if _df_name in globals():
        _df = globals()[_df_name]
        drop_cols = [c for c in flag_df.columns if c != "hadm_id" and c in
        _df.columns]
        if drop_cols:
            _df = _df.drop(columns=drop_cols)
        _df = _df.merge(flag_df, on="hadm_id", how="left")
        globals()[_df_name] = _df

# Prevalence (combined flags)
for k in FLAGS:
    if k in ed_df.columns:
        print(k, int(ed_df[k].fillna(0).sum()))

```

Hosp ICD flags chunk 1: rows=5070
ED ICD flags chunk 1: rows=1049
flag_copd 1508
flag_osa_ohs 992
flag_chf 1947
flag_neuromuscular 56
flag_opioid_substance 351
flag_pneumonia 1275

3.0.10 Phase 8 — Timing phenotypes and derived bands

Rationale: Compute time-anchored hypercapnia/acidemia phenotypes for ED presentation vs later course.

```
# Purpose: Create time-anchored phenotypes and severity bands relative to
#           ↵ first ED presentation.

# Timing phenotypes

anchor = "ed_intime_first" if "ed_intime_first" in ed_df.columns else
        ↵ "ed_intime"

qualifying_hadm_cols = [
    "hadm_id",
    "any_hypercap_icd",
    "pco2_threshold_0_24h",
    "enrolled_any",
    "enrollment_route",
    "qualifying_pco2_time",
    "qualifying_pco2_mmhg",
    "qualifying_site",
    "qualifying_source_branch",
    "qualifying_threshold_mmhg",
    "dt_qualifying_hypercapnia_hours",
    "max_pco2_0_6h",
    "max_pco2_0_24h",
    "max_pco2_0_6h_source_branch",
    "max_pco2_0_24h_source_branch",
    "first_abg_hypercap_time_0_24h",
    "first_abg_hypercap_pco2_mmhg",
    "first_abg_hypercap_source_branch",
    "first_abg_hypercap_itemid",
    "first_vbg_hypercap_time_0_24h",
    "first_vbg_hypercap_pco2_mmhg",
    "first_vbg_hypercap_source_branch",
    "first_vbg_hypercap_itemid",
    "first_other_hypercap_time_0_24h",
    "first_other_hypercap_pco2_mmhg",
    "first_other_hypercap_source_branch",
    "first_other_hypercap_itemid",
]

if "df" in globals() and "hadm_id" in ed_df.columns and "hadm_id" in
    ↵ df.columns:
```

```

available_qual_cols = [
    col for col in qualifying_hadm_cols if col == "hadm_id" or col in
    df.columns
]
if len(available_qual_cols) > 1:
    hadm_qual = (
        df[available_qual_cols]
        .drop_duplicates(subset=["hadm_id"], keep="first")
        .copy()
    )
    drop_existing_qual_cols = [
        col for col in available_qual_cols if col != "hadm_id" and col
    in ed_df.columns
    ]
    if drop_existing_qual_cols:
        ed_df = ed_df.drop(columns=drop_existing_qual_cols)
    ed_df = ed_df.merge(hadm_qual, on="hadm_id", how="left")

ed_df["qualifying_pco2_time"] = pd.to_datetime(
    ed_df.get("qualifying_pco2_time"), errors="coerce"
)
for time_col in [
    "first_abg_hypercap_time_0_24h",
    "first_vbg_hypercap_time_0_24h",
    "first_other_hypercap_time_0_24h",
]:
    if time_col in ed_df.columns:
        ed_df[time_col] = pd.to_datetime(ed_df[time_col], errors="coerce")
ed_df["qualifying_pco2_mmhg"] = pd.to_numeric(
    ed_df.get("qualifying_pco2_mmhg"), errors="coerce"
)
ed_df["qualifying_threshold_mmhg"] = pd.to_numeric(
    ed_df.get("qualifying_threshold_mmhg"), errors="coerce"
)
for value_col in [
    "first_abg_hypercap_pco2_mmhg",
    "first_vbg_hypercap_pco2_mmhg",
    "first_other_hypercap_pco2_mmhg",
]:
    if value_col in ed_df.columns:
        ed_df[value_col] = pd.to_numeric(ed_df[value_col], errors="coerce")
ed_df["qualifying_site"] = (
    ed_df.get("qualifying_site", pd.Series(pd.NA, index=ed_df.index))
)

```

```

    .astype("string")
    .str.lower()
)
ed_df["qualifying_source_branch"] = (
    ed_df.get("qualifying_source_branch", pd.Series(pd.NA,
    ↵ index=ed_df.index))
    .astype("string")
    .str.upper()
)
for source_col in [
    "first_abg_hypercap_source_branch",
    "first_vbg_hypercap_source_branch",
    "first_other_hypercap_source_branch",
]:
    if source_col in ed_df.columns:
        ed_df[source_col] = ed_df[source_col].astype("string").str.upper()

if "pco2_threshold_0_24h" not in ed_df.columns:
    ed_df["pco2_threshold_0_24h"] =
    ↵ ed_df["qualifying_pco2_time"].notna().astype("Int64")
else:
    ed_df["pco2_threshold_0_24h"] = (
        pd.to_numeric(ed_df["pco2_threshold_0_24h"], errors="coerce")
        .fillna(ed_df["qualifying_pco2_time"].notna().astype(int))
        .astype("Int64")
    )

ed_df["dt_qualifying_hypercapnia_hours"] = (
    ed_df["qualifying_pco2_time"] - pd.to_datetime(ed_df[anchor],
    ↵ errors="coerce")
).dt.total_seconds() / 3600.0
ed_df["dt_first_qualifying_gas_hours"] =
    ↵ ed_df["dt_qualifying_hypercapnia_hours"]
ed_df["first_gas_time"] = ed_df["qualifying_pco2_time"]
ed_df["first_pco2"] = ed_df["qualifying_pco2_mmhg"]
ed_df["first_gas_anchor_has_pco2"] = ed_df["qualifying_pco2_time"].notna()
ed_df["first_gas_anchor_source_validated"] = (
    ed_df["qualifying_pco2_time"].notna()
    & ed_df["qualifying_site"].isin({"arterial", "venous", "unknown"})
)
ed_df["presenting_hypercapnia"] = (

```

```

pd.to_numeric(ed_df["dt_qualifying_hypercapnia_hours"],
    ↪ errors="coerce").le(6)
).astype("Int64")
ed_df["late_hypercapnia"] = (
    pd.to_numeric(ed_df["dt_qualifying_hypercapnia_hours"],
    ↪ errors="coerce").gt(6)
    & pd.to_numeric(ed_df["dt_qualifying_hypercapnia_hours"],
    ↪ errors="coerce").le(24)
).astype("Int64")

qualifying_hours = pd.to_numeric(ed_df["dt_qualifying_hypercapnia_hours"],
    ↪ errors="coerce")
qualifying_observed = qualifying_hours.notna()

threshold_any = (
    pd.to_numeric(ed_df.get("pco2_threshold_0_24h", pd.Series(0,
    ↪ index=ed_df.index)), errors="coerce")
    .fillna(0)
    .astype(int)
)
any_icd = (
    pd.to_numeric(ed_df.get("any_hypercap_icd", pd.Series(0,
    ↪ index=ed_df.index)), errors="coerce")
    .fillna(0)
    .astype(int)
)
ed_df["hypercap_timing_class"] = np.select(
    [
        qualifying_observed & qualifying_hours.le(6),
        qualifying_observed & qualifying_hours.gt(6) &
    ↪ qualifying_hours.le(24),
        (qualifying_observed & qualifying_hours.gt(24))
        | (~qualifying_observed & threshold_any.eq(1)),
    ],
    [
        "presenting_0_6h",
        "late_6_24h",
        "none_within_24h",
    ],
    default="icd_only_or_no_qualifying_gas",
)
# NIV/IMV timing relative to first ED presentation

```

```

if "first_imv_time" in ed_df.columns:
    ed_df["dt_first_imv_hours"] = (ed_df["first_imv_time"] -
    ↵ ed_df[anchor]).dt.total_seconds() / 3600.0
if "first_niv_time" in ed_df.columns:
    ed_df["dt_first_niv_hours"] = (ed_df["first_niv_time"] -
    ↵ ed_df[anchor]).dt.total_seconds() / 3600.0

ed_df["hospital_los_negative_flag"] = (
    pd.to_numeric(ed_df.get("hospital_los_hours"), errors="coerce")
    .lt(0)
    .fillna(False)
)
admittime_series = pd.to_datetime(ed_df.get("admittime"), errors="coerce")
ed_intime_series = pd.to_datetime(ed_df.get("ed_intime"), errors="coerce")
ed_df["admittime_before_ed_intime_flag"] = (
    admittime_series.notna()
    & ed_intime_series.notna()
    & admittime_series.lt(ed_intime_series)
)
dischtime_series = pd.to_datetime(ed_df.get("dischtime"), errors="coerce")
ed_df["dischtime_before_admittime_flag"] = (
    dischtime_series.notna()
    & admittime_series.notna()
    & dischtime_series.lt(admittime_series)
)
ed_df["time_integrity_any"] = (
    ed_df["hospital_los_negative_flag"].fillna(False).astype(bool)
    | ed_df["admittime_before_ed_intime_flag"].fillna(False).astype(bool)
    | ed_df["dischtime_before_admittime_flag"].fillna(False).astype(bool)
)
ed_df["timing_usable_for_model"] =
    (~ed_df["time_integrity_any"].fillna(False).astype(bool)).astype("Int64")
hospital_los_numeric = pd.to_numeric(ed_df.get("hospital_los_hours"),
    ↵ errors="coerce")
ed_df["hospital_los_hours_model"] =
    ↵ hospital_los_numeric.mask(hospital_los_numeric.lt(0))
anchor_series = pd.to_datetime(ed_df.get(anchor), errors="coerce")
for event_col, delta_col, model_col, outside_col in [
    ("first_imv_time", "dt_first_imv_hours", "dt_first_imv_hours_model",
    ↵ "imv_time_outside_window_flag"),
    ("first_niv_time", "dt_first_niv_hours", "dt_first_niv_hours_model",
    ↵ "niv_time_outside_window_flag"),
]:

```

```

if event_col not in ed_df.columns:
    ed_df[model_col] = pd.Series(pd.NA, index=ed_df.index,
↪ dtype="Float64")
    ed_df[outside_col] = False
    continue
event_time = pd.to_datetime(ed_df[event_col], errors="coerce")
delta_hours = pd.to_numeric(ed_df.get(delta_col), errors="coerce")
outside_lower = anchor_series.notna() & event_time.notna() &
↪ event_time.lt(anchor_series)
outside_upper = dischtime_series.notna() & event_time.notna() &
↪ event_time.gt(dischtime_series)
negative_delta = delta_hours.lt(0).fillna(False)
outside_window = outside_lower | outside_upper | negative_delta
ed_df[outside_col] = outside_window.fillna(False)
ed_df[model_col] = delta_hours.mask(outside_window)

# Bands
bins_ph = [-1, 7.20, 7.30, 7.35, 99]
labels_ph = ["<7.20", "7.20-7.29", "7.30-7.34", "7.35"]
ed_df["ph_band"] = pd.cut(ed_df["first_ph"], bins=bins_ph, labels=labels_ph)

bins_hco3 = [-1, 24, 30, 999]
labels_hco3 = ["<24", "24-29", "30"]
ed_df["hco3_band"] = pd.cut(ed_df["first_hco3"], bins=bins_hco3,
↪ labels=labels_hco3)

bins_lac = [-1, 2, 4, 999]
labels_lac = ["<2", "2-4", ">4"]
ed_df["lactate_band"] = pd.cut(ed_df["first_lactate"], bins=bins_lac,
↪ labels=labels_lac)

# Align ED-stay gas source flags to hadm-level thresholds and ICU POC signal.
# This keeps source-specific flags internally consistent in the final ED-stay
↪ export.
if "df" in globals() and "hadm_id" in ed_df.columns and "hadm_id" in
↪ df.columns:
    hadm_thr_cols = [
        c
    for c in [
        "abg_hypercap_threshold",
        "vbg_hypercap_threshold",

```

```

        "unknown_hypercap_threshold",
    ]
    if c in df.columns
]
if hadm_thr_cols:
    hadm_thr = df[["hadm_id"] + hadm_thr_cols].drop_duplicates("hadm_id")
    hadm_thr_map = hadm_thr.set_index("hadm_id")

    for col in hadm_thr_cols:
        mapped = pd.to_numeric(ed_df["hadm_id"].map(hadm_thr_map[col]),
        errors="coerce")
        existing = pd.to_numeric(ed_df[col], errors="coerce") if col in
        ed_df.columns else pd.Series(np.nan, index=ed_df.index)
        ed_df[col] = existing.fillna(mapped).fillna(0).astype("Int64")

ed_df["pco2_threshold_any"] = ed_df["pco2_threshold_0_24h"].astype("Int64")
if "any_hypercap_icd" not in ed_df.columns:
    ed_df["any_hypercap_icd"] = 0
ed_df["any_hypercap_icd"] = (
    pd.to_numeric(ed_df["any_hypercap_icd"],
    errors="coerce").fillna(0).astype("Int64")
)
ed_df["enrolled_any"] = (
    (ed_df["any_hypercap_icd"] == 1)
    | (ed_df["pco2_threshold_0_24h"] == 1)
).astype("Int64")
ed_df["enrolled_any_icd_union_secondary"] =
    ed_df["enrolled_any"].astype("Int64")
ed_df["enrollment_route"] = np.select(
    [
        (ed_df["any_hypercap_icd"] == 0) & (ed_df["pco2_threshold_0_24h"] ==
        1),
        (ed_df["any_hypercap_icd"] == 1) & (ed_df["pco2_threshold_0_24h"] ==
        0),
        (ed_df["any_hypercap_icd"] == 1) & (ed_df["pco2_threshold_0_24h"] ==
        1),
    ],
    ["GAS_ONLY", "ICD_ONLY", "ICD+GAS"],
    default="NONE",
)
# QC invariants for enrolled rows.
max_pco2_0_24h_lt_qualifying_n = 0

```

```

max_pco2_0_6h_lt_qualifying_n = 0
enrolled_mask = pd.to_numeric(ed_df["pco2_threshold_0_24h"],
    ↪ errors="coerce").fillna(0).astype(int).eq(1)
if enrolled_mask.any():
    qualifying_time_missing_n = int(
        pd.to_datetime(ed_df.loc[enrolled_mask, "qualifying_pco2_time"]),
    ↪ errors="coerce").isna().sum()
)
if qualifying_time_missing_n:
    raise ValueError(
        f"Found {qualifying_time_missing_n} enrolled rows without
            ↪ qualifying_pco2_time."
)

qualifying_hours_enrolled = pd.to_numeric(
    ed_df.loc[enrolled_mask, "dt_qualifying_hypercapnia_hours"],
    errors="coerce",
)
out_of_window_n = int(
    (~qualifying_hours_enrolled.between(0, 24,
    ↪ inclusive="both")).fillna(True).sum()
)
if out_of_window_n:
    raise ValueError(
        f"Found {out_of_window_n} enrolled rows with
            ↪ dt_qualifying_hypercapnia_hours outside [0, 24]."
    )

if {"max_pco2_0_24h", "qualifying_pco2_mmhg"}.issubset(ed_df.columns):
    qualifying_values_24h = pd.to_numeric(
        ed_df.loc[enrolled_mask, "qualifying_pco2_mmhg"], errors="coerce"
    )
    max_values_24h = pd.to_numeric(
        ed_df.loc[enrolled_mask, "max_pco2_0_24h"], errors="coerce"
    )
    max_pco2_0_24h_lt_qualifying_n = int(
        (max_values_24h.lt(qualifying_values_24h) |
    ↪ max_values_24h.isna()).sum()
    )
    print(
        "QA: enrolled rows with max_pco2_0_24h < qualifying_pco2_mmhg (or
            ↪ missing max):",
        max_pco2_0_24h_lt_qualifying_n,
    )

```

```

        )
    if max_pco2_0_24h_lt_qualifying_n > 0:
        raise ValueError(
            "max_pco2_0_24h invariant failed: enrolled rows must satisfy
             "
            "max_pco2_0_24h >= qualifying_pco2_mmhg."
        )

    if {"max_pco2_0_6h", "qualifying_pco2_mmhg",
        "dt_qualifying_hypercapnia_hours"}.issubset(ed_df.columns):
        qualifying_in_6h_mask = enrolled_mask & pd.to_numeric(
            ed_df["dt_qualifying_hypercapnia_hours"], errors="coerce"
        ).le(6.0)
        if qualifying_in_6h_mask.any():
            qualifying_values_6h = pd.to_numeric(
                ed_df.loc[qualifying_in_6h_mask, "qualifying_pco2_mmhg"],
                errors="coerce",
            )
            max_values_6h = pd.to_numeric(
                ed_df.loc[qualifying_in_6h_mask, "max_pco2_0_6h"],
                errors="coerce"
            )
            max_pco2_0_6h_lt_qualifying_n = int(
                (max_values_6h.lt(qualifying_values_6h) |
        max_values_6h.isna()).sum()
            )
            print(
                "QA: enrolled dt<=6h rows with max_pco2_0_6h <
                    qualifying_pco2_mmhg (or missing max):",
                max_pco2_0_6h_lt_qualifying_n,
            )
            if max_pco2_0_6h_lt_qualifying_n > 0:
                raise ValueError(
                    "max_pco2_0_6h invariant failed: enrolled rows with
                     dt_qualifying_hypercapnia_hours<=6 "
                    "must satisfy max_pco2_0_6h >= qualifying_pco2_mmhg."
                )
        route_sum = (
            pd.to_numeric(ed_df.loc[enrolled_mask, "abg_hypercap_threshold"],
        errors="coerce").fillna(0).astype(int)
            + pd.to_numeric(ed_df.loc[enrolled_mask, "vbg_hypercap_threshold"],
        errors="coerce").fillna(0).astype(int)
    
```

```

+ pd.to_numeric(ed_df.loc[enrolled_mask,
    ↪ "unknown_hypercap_threshold"], 
    ↪ errors="coerce").fillna(0).astype(int)
)
route_invalid_n = int((route_sum < 1).sum())
if route_invalid_n:
    raise ValueError(
        f"Found {route_invalid_n} gas-qualified rows without any
            ↪ ABG/VBG/UNKNOWN threshold flag."
    )

abg_mask = enrolled_mask & (
    pd.to_numeric(ed_df["abg_hypercap_threshold"],
    ↪ errors="coerce").fillna(0).astype(int).eq(1)
)
if abg_mask.any():
    abg_missing_n = int(
        pd.to_datetime(ed_df.loc[abg_mask,
    ↪ "first_abg_hypercap_time_0_24h"], errors="coerce").isna().sum()
    )
    abg_low_n = int(
        pd.to_numeric(
            ed_df.loc[abg_mask, "first_abg_hypercap_pco2_mmhg"],
            errors="coerce",
        ).lt(45.0).fillna(True).sum()
    )
    if abg_missing_n or abg_low_n:
        raise ValueError(
            "ABG qualifying invariants failed: "
            f"missing_time={abg_missing_n}, below_threshold={abg_low_n}."
        )

vbg_mask = enrolled_mask & (
    pd.to_numeric(ed_df["vbg_hypercap_threshold"],
    ↪ errors="coerce").fillna(0).astype(int).eq(1)
)
if vbg_mask.any():
    vbg_missing_n = int(
        pd.to_datetime(ed_df.loc[vbg_mask,
    ↪ "first_vbg_hypercap_time_0_24h"], errors="coerce").isna().sum()
    )
    vbg_low_n = int(
        pd.to_numeric(

```

```

        ed_df.loc[vbg_mask, "first_vbg_hypercap_pco2_mmhg"],
        errors="coerce",
    ).lt(50.0).fillna(True).sum()
)
if vbg_missing_n or vbg_low_n:
    raise ValueError(
        "VBG qualifying invariants failed: "
        f"missing_time={vbg_missing_n}, below_threshold={vbg_low_n}."
    )

other_mask = enrolled_mask & (
    pd.to_numeric(ed_df["unknown_hypercap_threshold"],
    errors="coerce").fillna(0).astype(int).eq(1)
)
if other_mask.any():
    other_missing_n = int(
        pd.to_datetime(ed_df.loc[other_mask,
    "first_other_hypercap_time_0_24h"], errors="coerce").isna().sum()
    )
    other_low_n = int(
        pd.to_numeric(
            ed_df.loc[other_mask, "first_other_hypercap_pco2_mmhg"],
            errors="coerce",
        ).lt(50.0).fillna(True).sum()
    )
    if other_missing_n or other_low_n:
        raise ValueError(
            "UNKNOWN/OTHER qualifying invariants failed: "
            f"missing_time={other_missing_n},
            below_threshold={other_low_n}."
        )

```

QA: enrolled rows with max_pco2_0_24h < qualifying_pco2_mmhg (or missing max): 0
 QA: enrolled dt<=6h rows with max_pco2_0_6h < qualifying_pco2_mmhg (or missing max): 0

3.1 QA & Data Fidelity

Rationale: Validate joins, missingness, lab completeness, and produce reproducibility artifacts.

```
# Purpose: Create reusable data-quality and merge guardrail helpers to
    prevent silent join errors.
```

```

# --- QA checks (lightweight, deterministic)
STRICT_QA = False

# 1) Key uniqueness
if 'ed_df' in globals() and 'ed_stay_id' in ed_df.columns:
    try:
        assert_unique(ed_df, 'ed_stay_id', 'ed_df')
        print('QA: ed_stay_id unique OK')
    except Exception as e:
        print('QA FAIL:', e)
        if STRICT_QA:
            raise

# 2) Inclusion sanity: enrolled rows must match ICD OR gas-qualified in
#    0-24h.
if 'ed_df' in globals():
    enrolled = pd.to_numeric(ed_df.get('enrolled_any'),
    ↵ errors='coerce').fillna(0).astype(int)
    threshold = pd.to_numeric(ed_df.get('pco2_threshold_0_24h'),
    ↵ errors='coerce').fillna(0).astype(int)
    icd_any = pd.to_numeric(ed_df.get('any_hypcap_icd'),
    ↵ errors='coerce').fillna(0).astype(int)
    enrolled_expected = ((threshold == 1) | (icd_any == 1)).astype(int)
    mismatch_n = int((enrolled != enrolled_expected).sum())
    print('QA: enrolled_any != (any_hypcap_icd OR pco2_threshold_0_24h'
    ↵ count:', mismatch_n)
    if STRICT_QA and mismatch_n > 0:
        raise AssertionError('Found rows where enrolled_any disagrees with
        ↵ ICD OR 0-24h gas rule.')

# 3) Temporal sanity: first_gas_time >= ed_intime (allow small negative
#    drift)
if 'ed_df' in globals() and 'first_gas_time' in ed_df.columns and 'ed_intime'
    ↵ in ed_df.columns:
    dt = (pd.to_datetime(ed_df['first_gas_time']) -
    ↵ pd.to_datetime(ed_df['ed_intime'])).dt.total_seconds() / 3600
    n_neg = int((dt < -1).sum()) # allow 1h clock drift
    print('QA: first_gas_time < ed_intime by >1h:', n_neg)
    if STRICT_QA and n_neg > 0:
        raise AssertionError('first_gas_time before ed_intime by >1h')

# 4) Range checks (report only)

```

```

if 'ed_df' in globals():
    ranges = {
        'first_ph': (6.8, 7.8),
        'first_pco2': (10, 200),
        'first_lactate': (0, 30),
        'creatinine': (0, 20),
    }
    rc = check_ranges(ed_df, ranges)
    if not rc.empty:
        print('QA range check (n out-of-range):')
        print(rc.to_string(index=False))

QA: ed_stay_id unique OK
QA: enrolled_any != (any_hypercap_icd OR pco2_threshold_0_24h) count: 0
QA: first_gas_time < ed_intime by >1h: 2
QA range check (n out-of-range):
      col  n_bad
      first_ph      9
      first_pco2     0
      first_lactate   0

# Purpose: Report missingness for key variables to support transparent
#           data-quality reporting.

import json
from datetime import datetime

# T1 - uniqueness and join explosion guard
if ed_df["ed_stay_id"].nunique() != len(ed_df):
    raise ValueError("ed_stay_id not unique after merges")

# T1b - create cleaned vitals model fields while preserving raw values.
ed_df, vitals_outlier_audit = add_vitals_model_fields(ed_df)
ed_df, gas_outlier_audit = add_gas_model_fields(ed_df)
outliers_audit = pd.concat([vitals_outlier_audit, gas_outlier_audit],
                           ignore_index=True, sort=False)

# T1c - fail-fast cleaned-vitals bounds checks.
clean_bounds = {
    "ed_triage_o2sat_clean": (0.0, 100.0),
    "ed_first_o2sat_clean": (0.0, 100.0),
    "ed_triage_sbp_clean": (20.0, 300.0),
    "ed_first_sbp_clean": (20.0, 300.0),
}

```

```

    "ed_triage_dbp_clean": (10.0, 200.0),
    "ed_first_dbp_clean": (10.0, 200.0),
}
for column_name, (lower_bound, upper_bound) in clean_bounds.items():
    if column_name not in ed_df.columns:
        continue
    numeric = pd.to_numeric(ed_df[column_name], errors="coerce")
    invalid_n = int((numeric.notna() & ((numeric < lower_bound) | (numeric >
    ↵ upper_bound))).sum())
    if invalid_n:
        raise ValueError(
            f"Cleaned vital bounds failure for {column_name}: {invalid_n}"
            ↵     values outside [{lower_bound}, {upper_bound}]"
        )

for temp_clean_column in ("ed_triage_temp_f_clean", "ed_first_temp_f_clean"):
    if temp_clean_column not in ed_df.columns:
        continue
    numeric = pd.to_numeric(ed_df[temp_clean_column], errors="coerce")
    residual_celsius_like_n = int(numeric.between(20.0, 50.0,
    ↵ inclusive="both").sum())
    if residual_celsius_like_n:
        raise ValueError(
            f"Found {residual_celsius_like_n} cleaned values still in 20-50°F"
            ↵     band in {temp_clean_column}."
        )

ed_vitals_distribution_summary, ed_vitals_extreme_examples,
    ↵ ed_vitals_model_delta = build_ed_vitals_audit_artifacts(ed_df)

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")

vitals_outlier_audit_path = prior_runs_dir / f"{out_date}"
    ↵ vitals_outlier_audit.csv"
vitals_outlier_audit.to_csv(vitals_outlier_audit_path, index=False)
print("Wrote:", vitals_outlier_audit_path)
outliers_audit_path = prior_runs_dir / f"{out_date} outliers_audit.csv"
outliers_audit.to_csv(outliers_audit_path, index=False)
print("Wrote:", outliers_audit_path)
ed_vitals_distribution_summary_path = prior_runs_dir / f"{out_date}"
    ↵ ed_vitals_distribution_summary.csv"

```

```

ed_vitals_distribution_summary.to_csv(ed_vitals_distribution_summary_path,
    ↪ index=False)
print("Wrote:", ed_vitals_distribution_summary_path)
ed_vitals_extreme_examples_path = prior_runs_dir / f"{out_date}"
↪ ed_vitals_extreme_examples.csv"
ed_vitals_extreme_examples.to_csv(ed_vitals_extreme_examples_path,
    ↪ index=False)
print("Wrote:", ed_vitals_extreme_examples_path)
ed_vitals_model_delta_path = prior_runs_dir / f"{out_date}"
↪ ed_vitals_model_delta.csv"
ed_vitals_model_delta.to_csv(ed_vitals_model_delta_path, index=False)
print("Wrote:", ed_vitals_model_delta_path)

timing_integrity_audit = pd.DataFrame(
    [
        {
            "rows_total": int(len(ed_df)),
            "time_integrity_any_n": int(ed_df.get("time_integrity_any",
                ↪ pd.Series(False,
                ↪ index=ed_df.index)).fillna(False).astype(bool).sum()),
            "timing_usable_for_model_n":
                ↪ int(pd.to_numeric(ed_df.get("timing_usable_for_model"),
                ↪ errors="coerce").fillna(0).astype(int).sum()),
            "hospital_los_negative_n":
                ↪ int(ed_df.get("hospital_los_negative_flag", pd.Series(False,
                ↪ index=ed_df.index)).fillna(False).astype(bool).sum()),
            "admittime_before_ed_intime_n":
                ↪ int(ed_df.get("admittime_before_ed_intime_flag",
                ↪ pd.Series(False,
                ↪ index=ed_df.index)).fillna(False).astype(bool).sum()),
            "dischtime_before_admittime_n":
                ↪ int(ed_df.get("dischtime_before_admittime_flag",
                ↪ pd.Series(False,
                ↪ index=ed_df.index)).fillna(False).astype(bool).sum()),
            "qualifying_gas_observed_n": int(pd.to_numeric(ed_df.get(
                ↪ "dt_qualifying_hypcapnia_hours"),
                ↪ errors="coerce").notna().sum()),
            "qualifying_gas_missing_n": int(pd.to_numeric(ed_df.get(
                ↪ "dt_qualifying_hypcapnia_hours"),
                ↪ errors="coerce").isna().sum()),
        }
    ]
)

```

```

timing_integrity_audit_path = prior_runs_dir / f"{out_date}"
    ↵ timing_integrity_audit.csv"
timing_integrity_audit.to_csv(timing_integrity_audit_path, index=False)
print("Wrote:", timing_integrity_audit_path)

ventilation_timing_audit = pd.DataFrame(
    [
        {
            "rows_total": int(len(ed_df)),
            "dt_first_imv_hours_negative_n":
                ↵ int(pd.to_numeric(ed_df.get("dt_first_imv_hours")),
                ↵ errors="coerce").lt(0).fillna(False).sum(),
            "dt_first_niv_hours_negative_n":
                ↵ int(pd.to_numeric(ed_df.get("dt_first_niv_hours")),
                ↵ errors="coerce").lt(0).fillna(False).sum(),
            "imv_time_outside_window_n":
                ↵ int(ed_df.get("imv_time_outside_window_flag",
                ↵ pd.Series(False,
                ↵ index=ed_df.index)).fillna(False).astype(bool).sum()),
            "niv_time_outside_window_n":
                ↵ int(ed_df.get("niv_time_outside_window_flag",
                ↵ pd.Series(False,
                ↵ index=ed_df.index)).fillna(False).astype(bool).sum()),
            "dt_first_imv_hours_model_nonnull_n":
                ↵ int(pd.to_numeric(ed_df.get("dt_first_imv_hours_model")),
                ↵ errors="coerce").notna().sum(),
            "dt_first_niv_hours_model_nonnull_n":
                ↵ int(pd.to_numeric(ed_df.get("dt_first_niv_hours_model")),
                ↵ errors="coerce").notna().sum(),
        }
    ]
)
ventilation_timing_audit_path = prior_runs_dir / f"{out_date}"
    ↵ ventilation_timing_audit.csv"
ventilation_timing_audit.to_csv(ventilation_timing_audit_path, index=False)
print("Wrote:", ventilation_timing_audit_path)

hco3_numeric = pd.to_numeric(ed_df.get("first_hco3"), errors="coerce")
hco3_source = ed_df.get("first_hco3_source", pd.Series("missing",
    ↵ index=ed_df.index)).astype("string").fillna("missing")
hco3_source_value_mismatch_n = int(((hco3_source != "missing") &
    ↵ hco3_numeric.isna()).sum())
hco3_source_counts = (

```

```

hco3_source.value_counts(dropna=False)
    .rename_axis("first_hco3_source")
    .reset_index(name="row_count")
)
hco3_source_counts["null_value_count"] =
    hco3_source_counts["first_hco3_source"].map(
        lambda source_name: int(
            ((hco3_source == source_name) & hco3_numeric.isna()).sum()
        )
)
hco3_source_counts["nonnull_value_count"] =
    hco3_source_counts["first_hco3_source"].map(
        lambda source_name: int(
            ((hco3_source == source_name) & hco3_numeric.notna()).sum()
        )
)
hco3_source_counts["null_value_rate"] = np.where(
    hco3_source_counts["row_count"] > 0,
    hco3_source_counts["null_value_count"] / hco3_source_counts["row_count"],
    0.0,
)
hco3_source_counts["source_value_mismatch_n"] = hco3_source_value_mismatch_n
hco3_integrity_audit_path = prior_runs_dir / f"{out_date}"
    hco3_integrity_audit.csv"
hco3_source_counts.to_csv(hco3_integrity_audit_path, index=False)
print("Wrote:", hco3_integrity_audit_path)

print("ED vitals cleaning summary (raw->clean removals):")
display(
    ed_vitals_model_delta[
        [
            "vital_name",
            "raw_n",
            "clean_n",
            "raw_removed_n",
            "raw_removed_pct",
            "converted_celsius_like_n",
        ]
    ].sort_values("vital_name")
)

# T2 - missingness summary with expected-vs-unexpected classifications.
expected_sparse_fields = set()

```

```

omr_pre_window_rows = int(omr_diagnostics.get("pre_window_candidate_rows",
    ↵ 0)) if "omr_diagnostics" in globals() else 0
omr_post_window_rows = int(omr_diagnostics.get("post_window_candidate_rows",
    ↵ 0)) if "omr_diagnostics" in globals() else 0
if omr_pre_window_rows == 0 and omr_post_window_rows == 0:
    expected_sparse_fields.update(
        {
            "bmi_closest_pre_ed",
            "height_closest_pre_ed",
            "weight_closest_pre_ed",
            "bmi_closest_pre_ed_uom",
            "height_closest_pre_ed_uom",
            "weight_closest_pre_ed_uom",
            "bmi_closest_pre_ed_time",
            "height_closest_pre_ed_time",
            "weight_closest_pre_ed_time",
        }
    )
)

miss = classify_missingness_expectations(
    ed_df,
    TARGET_FIELDS,
    expected_sparse_fields=expected_sparse_fields,
)
print(miss.sort_values(["missing_pct", "field"], ascending=[False,
    ↵ True]).head(30))

unexpected_full_null_fields = miss.loc[
    miss["expectation"].eq("unexpected_full_null"), "field"
].tolist()
expected_structural_null_fields = sorted(
    set(EXPECTED_STRUCTURAL_NULL_FIELDS).intersection(ed_df.columns)
)

# Compute UOM checks against canonical export frame contract (or
# deterministic preview).
if "final_cc" in globals():
    uom_scope_df = final_cc
    uom_scope_frame = "final_cc"
elif "df" in globals() and "ed_triage_cc" in df.columns:
    _mask_df_cc = df["ed_triage_cc"].notna() &
    ↵ (df["ed_triage_cc"].astype(str).str.strip() != "")
    _df_cc = df.loc[_mask_df_cc].copy()

```

```

    _mask_cc = ed_df["ed_triage_cc"].notna() &
    ↪ (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
    _ed_cc_only = ed_df.loc[_mask_cc].copy()
    _ed_tmp = _ed_cc_only.copy()
    if "ed_intime" in _ed_tmp.columns:
        _ed_tmp = _ed_tmp.sort_values(["hadm_id", "ed_intime", "ed_stay_id"],
    ↪ na_position="last")
    else:
        _ed_tmp = _ed_tmp.sort_values(["hadm_id", "ed_stay_id"],
    ↪ na_position="last")
    _ed_first = _ed_tmp.drop_duplicates("hadm_id", keep="first")
    _add_ed_cols = [c for c in _ed_first.columns if c not in _df_cc.columns]
    uom_scope_df = _df_cc.merge(_ed_first[["hadm_id"] + _add_ed_cols],
    ↪ on="hadm_id", how="left")
    uom_scope_frame = "final_cc_preview"
else:
    uom_scope_df = ed_df
    uom_scope_frame = "ed_df"

uom_expectation_checks = evaluate_uom_expectations(uom_scope_df)
uom_expectation_scope = {"frame": uom_scope_frame, "row_count":
    ↪ int(len(uom_scope_df))}

# T2b - route-stratified first_other_pco2 audit using canonical QA scope.
first_other_scope_df = (
    uom_scope_df
    if {"first_other_src", "first_other_pco2"}.issubset(uom_scope_df.columns)
    else ed_df
)
first_other_scope_frame = "uom_scope_df" if first_other_scope_df is
    ↪ uom_scope_df else "ed_df"
first_other_pco2_audit = build_first_other_pco2_audit(first_other_scope_df)
first_other_pco2_audit_path = prior_runs_dir / f"{out_date}"
    ↪ first_other_pco2_audit.csv"
first_other_pco2_audit.to_csv(first_other_pco2_audit_path, index=False)
print("Wrote:", first_other_pco2_audit_path)
print("first_other_pco2 audit scope:", first_other_scope_frame)

print("Expected structural null fields:", expected_structural_null_fields)
print("Unexpected full-null fields (top 20):",
    ↪ unexpected_full_null_fields[:20])
print("UOM expectation scope:", uom_expectation_scope)

```

```

print("UOM expectation checks:", json.dumps(uom_expectation_checks,
    indent=2))

# T3 - lab capture completeness.
pct_any_6h = float(p06["ed_stay_id"].nunique() /
    ↵ max(ed_df["ed_stay_id"].nunique(),1))
pct_any_24h = float(p24["ed_stay_id"].nunique() /
    ↵ max(ed_df["ed_stay_id"].nunique(),1))
print("% any gas panel 0-6h:", round(pct_any_6h*100,1))
print("% any gas panel 0-24h:", round(pct_any_24h*100,1))

if "gas_source_audit" not in globals():
    gas_source_audit = summarize_gas_source(panel if "panel" in globals()
    ↵ else pd.DataFrame())

gas_overlap_summary = build_gas_source_overlap_summary(ed_df)
gas_overlap_summary_path = prior_runs_dir / f"{out_date}"
    ↵ gas_source_overlap_summary.csv"
gas_overlap_summary.to_csv(gas_overlap_summary_path, index=False)
print("Wrote:", gas_overlap_summary_path)

gas_source_unknown_rate = float(gas_source_audit.get("source_rates",
    ↵ {}).get("unknown", 0.0))
gas_source_other_rate = gas_source_unknown_rate
print("% source unknown (panel-level):", round(gas_source_unknown_rate * 100,
    ↵ 1))

# Blood-gas source distribution audit (LAB + POC definitive pCO2 by source
    ↵ class).
pco2_source_distribution_rows: list[dict[str, Any]] = []
PCC_DIST_COLUMNS = [
    "scope",
    "source_branch",
    "source",
    "n",
    "median",
    "p95",
    "max",
    "pct_ge_50",
]
    ↵

def _append_distribution_rows(source_branch: str, source: str, values:
    ↵ pd.Series) -> None:

```

```

        numeric = pd.to_numeric(values, errors="coerce").dropna()
        if numeric.empty:
            return
        pco2_source_distribution_rows.append(
            {
                "scope": "ALL_CANDIDATE_PCO2_LAB_POC",
                "source_branch": source_branch,
                "source": source,
                "n": int(numeric.shape[0]),
                "median": float(numeric.median()),
                "p95": float(numeric.quantile(0.95)),
                "max": float(numeric.max()),
                "pct_ge_50": float((numeric >= 50).mean()),
            }
        )

    pco2_distribution_scope = ed_df
    if "uom_scope_df" in globals() and isinstance(uom_scope_df, pd.DataFrame):
        poc_scope_columns = ["poc_abg_paco2", "poc_vbg_paco2", "poc_other_paco2"]
        has_poc_values = any(
            col in uom_scope_df.columns
            and pd.to_numeric(uom_scope_df[col], errors="coerce").notna().any()
            for col in poc_scope_columns
        )
        if has_poc_values:
            pco2_distribution_scope = uom_scope_df

    if "panel" in globals() and len(panel) > 0 and "pco2" in panel.columns:
        panel_pco2 = panel.copy()
        panel_pco2["pco2"] = pd.to_numeric(panel_pco2["pco2"], errors="coerce")
        panel_pco2 = panel_pco2.loc[panel_pco2["pco2"].notna()].copy()
        if len(panel_pco2) > 0:
            for source_name, subset in panel_pco2.groupby("source",
                ↳ dropna=False):
                _append_distribution_rows("LAB", str(source_name),
                ↳ subset["pco2"])

    if "poc_abg_paco2" in pco2_distribution_scope.columns:
        _append_distribution_rows("POC", "arterial",
        ↳ pco2_distribution_scope["poc_abg_paco2"])
    if "poc_vbg_paco2" in pco2_distribution_scope.columns:
        _append_distribution_rows("POC", "venous",
        ↳ pco2_distribution_scope["poc_vbg_paco2"])

```

```

if "poc_other_paco2" in pco2_distribution_scope.columns:
    _append_distribution_rows("POC", "unknown",
    ↵ pco2_distribution_scope["poc_other_paco2"])

pco2_source_distribution_audit = pd.DataFrame(pco2_source_distribution_rows)
for column_name in PCC_DIST_COLUMNS:
    if column_name not in pco2_source_distribution_audit.columns:
        pco2_source_distribution_audit[column_name] = pd.NA
pco2_source_distribution_audit =
    ↵ pco2_source_distribution_audit[PCC_DIST_COLUMNS]
pco2_source_distribution_audit_path = prior_runs_dir / f"{out_date}"
    ↵ pco2_source_distribution_audit.csv"
pco2_source_distribution_audit.to_csv(pco2_source_distribution_audit_path,
    ↵ index=False)
print("Wrote:", pco2_source_distribution_audit_path)

# Max-window pCO2 contributor source audit (branch responsible for hadm-level
    ↵ max values).
max_contributor_rows: list[dict[str, Any]] = []
for window_name, max_col, source_col in [
    ("0_6h", "max_pco2_0_6h", "max_pco2_0_6h_source_branch"),
    ("0_24h", "max_pco2_0_24h", "max_pco2_0_24h_source_branch"),
]:
    max_values = pd.to_numeric(ed_df.get(max_col), errors="coerce")
    source_values = (
        ed_df.get(source_col, pd.Series(pd.NA, index=ed_df.index))
        .astype("string")
        .fillna("UNKNOWN")
        .str.upper()
    )
    present_mask = max_values.notna()
    present_n = int(present_mask.sum())
    if present_n == 0:
        max_contributor_rows.append(
            {
                "window": window_name,
                "source_branch": "NONE",
                "max_value_rows": 0,
                "source_branch_rows": 0,
                "source_branch_rate": 0.0,
            }
        )
    continue

```

```

branch_counts =
    ↵ source_values.loc[present_mask].value_counts(dropna=False)
    for branch, count in branch_counts.items():
        max_contributor_rows.append(
            {
                "window": window_name,
                "source_branch": str(branch),
                "max_value_rows": present_n,
                "source_branch_rows": int(count),
                "source_branch_rate": float(count / present_n),
            }
        )
    )

pco2_window_max_contributor_audit = pd.DataFrame(max_contributor_rows)
pco2_window_max_contributor_audit_path =
    prior_runs_dir / f"{out_date} pco2_window_max_contributor_audit.csv"
)
pco2_window_max_contributor_audit.to_csv(
    pco2_window_max_contributor_audit_path, index=False
)
print("Wrote:", pco2_window_max_contributor_audit_path)

# Blood-gas triplet completeness audit (pCO2+pH+pO2) by branch/site.
triplet_rows: list[dict[str, Any]] = []
triplet_scope_df = (
    uom_scope_df
    if "uom_scope_df" in globals() and isinstance(uom_scope_df, pd.DataFrame)
    else ed_df
)
triplet_scope_name = (
    uom_scope_frame
    if "uom_scope_frame" in globals()
    else "ed_df"
)
triplet_specs = [
    ("LAB", "arterial", "lab_abg_paco2", "lab_abg_ph", "lab_abg_po2"),
    ("LAB", "venous", "lab_vbg_paco2", "lab_vbg_ph", "lab_vbg_po2"),
    ("LAB", "unknown", "lab_other_paco2", "lab_other_ph", "lab_other_po2"),
    ("POC", "arterial", "poc_abg_paco2", "poc_abg_ph", "poc_abg_po2"),
    ("POC", "venous", "poc_vbg_paco2", "poc_vbg_ph", "poc_vbg_po2"),
    ("POC", "unknown", "poc_other_paco2", "poc_other_ph", "poc_other_po2"),
]

```

```

for source_branch, site_name, pco2_column, ph_column, po2_column in
    ↵ triplet_specs:
        pco2_values = pd.to_numeric(
            triplet_scope_df.get(
                pco2_column,
                pd.Series(index=triplet_scope_df.index, dtype="float64"),
            ),
            errors="coerce",
        )
        ph_values = pd.to_numeric(
            triplet_scope_df.get(
                ph_column,
                pd.Series(index=triplet_scope_df.index, dtype="float64"),
            ),
            errors="coerce",
        )
        po2_values = pd.to_numeric(
            triplet_scope_df.get(
                po2_column,
                pd.Series(index=triplet_scope_df.index, dtype="float64"),
            ),
            errors="coerce",
        )
        pco2_present = pco2_values.notna()
        pco2_n = int(pco2_present.sum())
        ph_present_n = int((pco2_present & ph_values.notna()).sum())
        po2_present_n = int((pco2_present & po2_values.notna()).sum())
        ph_po2_present_n = int((pco2_present & ph_values.notna() &
    ↵ po2_values.notna()).sum())
        triplet_rows.append(
            {
                "scope_frame": triplet_scope_name,
                "source_branch": source_branch,
                "site": site_name,
                "pco2_n": pco2_n,
                "ph_present_n": ph_present_n,
                "po2_present_n": po2_present_n,
                "ph_po2_present_n": ph_po2_present_n,
                "ph_present_rate_given_pco2": float(ph_present_n / pco2_n) if
                    ↵ pco2_n else 0.0,
                "po2_present_rate_given_pco2": float(po2_present_n / pco2_n) if
                    ↵ pco2_n else 0.0,
            }
        )
    
```

```

        "triplet_present_rate_given_pco2": float(ph_po2_present_n /
            ↪ pco2_n) if pco2_n else 0.0,
    }
)

blood_gas_triplet_completeness_audit = pd.DataFrame(triplet_rows)
blood_gas_triplet_completeness_audit_path = (
    prior_runs_dir / f"{out_date} blood_gas_triplet_completeness_audit.csv"
)
blood_gas_triplet_completeness_audit.to_csv(
    blood_gas_triplet_completeness_audit_path, index=False
)
print("Wrote:", blood_gas_triplet_completeness_audit_path)

# Qualifying-event distribution audit (first qualifying by type + overall
# earliest qualifying).
qualifying_pco2_distribution_rows: list[dict[str, Any]] = []
QUALIFYING_DIST_COLUMNS = ["metric", "n", "median", "p5", "p95", "min",
    ↪ "max"]

def _append_qualifying_distribution(metric: str, values: pd.Series) -> None:
    numeric = pd.to_numeric(values, errors="coerce").dropna()
    qualifying_pco2_distribution_rows.append(
        {
            "metric": metric,
            "n": int(numeric.shape[0]),
            "median": float(numeric.median()) if not numeric.empty else
                ↪ pd.NA,
            "p5": float(numeric.quantile(0.05)) if not numeric.empty else
                ↪ pd.NA,
            "p95": float(numeric.quantile(0.95)) if not numeric.empty else
                ↪ pd.NA,
            "min": float(numeric.min()) if not numeric.empty else pd.NA,
            "max": float(numeric.max()) if not numeric.empty else pd.NA,
        }
    )

    _append_qualifying_distribution(
        "first_abg_hypercap_pco2_mmhg",
        ed_df.get("first_abg_hypercap_pco2_mmhg", pd.Series(dtype="float64")),
    )
    _append_qualifying_distribution(
        "first_vbg_hypercap_pco2_mmhg",

```

```

        ed_df.get("first_vbg_hypercap_pco2_mmhg", pd.Series(dtype="float64")),
    )
    _append_qualifying_distribution(
        "first_other_hypercap_pco2_mmhg",
        ed_df.get("first_other_hypercap_pco2_mmhg", pd.Series(dtype="float64")),
    )
    _append_qualifying_distribution(
        "qualifying_pco2_mmhg",
        ed_df.get("qualifying_pco2_mmhg", pd.Series(dtype="float64")),
    )
qualifying_pco2_distribution_audit =
    pd.DataFrame(qualifying_pco2_distribution_rows)
for column_name in QUALIFYING_DIST_COLUMNS:
    if column_name not in qualifying_pco2_distribution_audit.columns:
        qualifying_pco2_distribution_audit[column_name] = pd.NA
qualifying_pco2_distribution_audit = qualifying_pco2_distribution_audit[
    QUALIFYING_DIST_COLUMNS
]
qualifying_pco2_distribution_audit_path = (
    prior_runs_dir / f"{out_date}"
    ← qualifying_pco2_distribution_by_type_audit.csv"
)
qualifying_pco2_distribution_audit.to_csv(
    qualifying_pco2_distribution_audit_path, index=False
)
print("Wrote:", qualifying_pco2_distribution_audit_path)

poc_other_values_all = (
    pd.to_numeric(ed_df["poc_other_paco2"], errors="coerce")
    if "poc_other_paco2" in ed_df.columns
    else pd.Series(dtype="float64")
)
unknown_threshold_series = (
    pd.to_numeric(ed_df["unknown_hypercap_threshold"], errors="coerce")
    if "unknown_hypercap_threshold" in ed_df.columns
    else pd.Series(0, index=ed_df.index, dtype="float64")
)
first_other_detail_series = (
    ed_df["first_other_src_detail"].astype("string")
    if "first_other_src_detail" in ed_df.columns
    else pd.Series("missing", index=ed_df.index, dtype="string")
)
poc_other_candidate_n = int(poc_other_values_all.notna().sum())

```

```

poc_other_hypercap_n = int((poc_other_values_all >= 50).sum())
unknown_from_poc_detail_n = int(
    (
        unknown_threshold_series.fillna(0).astype(int).eq(1)
        & first_other_detail_series.str.lower().isin({"poc_bg_unknown",
            ↪ "poc_bg_unknown_available"})
    ).sum()
)
other_route_quarantine_audit = pd.DataFrame(
    [
        {
            "poc_other_candidate_n": poc_other_candidate_n,
            "poc_other_ge50_n": poc_other_hypercap_n,
            "poc_other_ge50_rate": float(poc_other_hypercap_n /
                ↪ poc_other_candidate_n)
            if poc_other_candidate_n
            else 0.0,
            "poc_unknown_hadm_n":
                ↪ int(globals().get("poc_unknown_hadm_n_for_audit",
                    ↪ unknown_from_poc_detail_n)),
            "unknown_hypercap_threshold_hadm_n":
                ↪ int(unknown_threshold_series.fillna(0).astype(int).sum()),
            "unknown_from_poc_detail_hadm_n": unknown_from_poc_detail_n,
            "poc_hypercap_0_24h_edstay_n":
                ↪ int(globals().get("poc_hypercap_0_24h_edstay_n", 0)),
            "poc_hypercap_0_24h_edstay_rate":
                ↪ float(globals().get("poc_hypercap_0_24h_edstay_rate", 0.0)),
        }
    ]
)
other_route_quarantine_audit_path = prior_runs_dir / f"{out_date}"
    ↪ other_route_quarantine_audit.csv"
other_route_quarantine_audit.to_csv(other_route_quarantine_audit_path,
    ↪ index=False)
print("Wrote:", other_route_quarantine_audit_path)

first_gas_anchor_audit = pd.DataFrame(
    [
        {
            "row_count": int(len(ed_df)),
            "first_gas_time_nonnull_n":
                ↪ int(pd.to_datetime(ed_df.get("first_gas_time"),
                    ↪ errors="coerce").notna().sum())

```

```

        if "first_gas_time" in ed_df.columns
        else 0,
        "first_gas_anchor_has_pco2_n":
            ↵ int(ed_df.get("first_gas_anchor_has_pco2", pd.Series(False,
            ↵ index=ed_df.index)).fillna(False).astype(bool).sum()),
        "first_gas_anchor_source_validated_n":
            ↵ int(ed_df.get("first_gas_anchor_source_validated",
            ↵ pd.Series(False,
            ↵ index=ed_df.index)).fillna(False).astype(bool).sum(),
        }
    ]
)
first_gas_anchor_audit["first_gas_without_pco2_anchor_n"] = (
    first_gas_anchor_audit["first_gas_time_nonnull_n"] -
    ↵ first_gas_anchor_audit["first_gas_anchor_has_pco2_n"]
)
first_gas_anchor_audit_path = prior_runs_dir / f"{out_date}"
    ↵ first_gas_anchor_audit.csv"
first_gas_anchor_audit.to_csv(first_gas_anchor_audit_path, index=False)
print("Wrote:", first_gas_anchor_audit_path)

# Drop diagnostic-only POC unknown threshold fields from exported cohort
    ↵ sheets.
for audit_only_column in [
    "poc_unknown_hypercap_threshold",
    "poc_other_quarantined_hypercap_threshold",
    "flag_any_gas_hypercapnia_poc",
    "flag_any_gas_hypercapnia",
    "__anthro_bmi_source",
    "__anthro_height_source",
    "__anthro_weight_source",
    "__anthro_bmi_hours_offset",
    "__anthro_height_hours_offset",
    "__anthro_weight_hours_offset",
]:
    if audit_only_column in ed_df.columns:
        ed_df = ed_df.drop(columns=[audit_only_column])
    if "df" in globals() and audit_only_column in df.columns:
        df = df.drop(columns=[audit_only_column])

hadm_rows = int(len(df)) if "df" in globals() else None
hadm_cc_rows = None
if "df" in globals() and "ed_triage_cc" in df.columns:

```

```

    _hadm_cc_mask = df["ed_triage_cc"].notna() &
    ↵ (df["ed_triage_cc"].astype(str).str.strip() != "")
    hadm_cc_rows = int(_hadm_cc_mask.sum())

uom_failed_checks = []
for uom_column, details in uom_expectation_checks.get("paco2_uom_checks",
    ↵ {}).items():
    if details.get("present") and not details.get("passes", True):
        uom_failed_checks.append(uom_column)

qa_warnings = []
qa_infos = []
if unexpected_full_null_fields:
    qa_warnings.append("unexpected_full_null_fields_detected")
if gas_source_audit.get("all_other_or_unknown", False):
    qa_warnings.append("gas_source_all_other_or_unknown")
if omr_pre_window_rows == 0 and omr_post_window_rows > 0:
    qa_infos.append("omr_pre_window_empty_post_windowFallbackUsed")
if omr_pre_window_rows == 0 and omr_post_window_rows == 0:
    qa_infos.append("omr_pre_post_windowEmptyExpectedSparse")
if uom_failed_checks:
    qa_warnings.append("uom_checks_failed")
if "status" in first_other_pco2_audit.columns and
    ↵ (first_other_pco2_audit["status"] == "missing_columns").any():
    qa_warnings.append("first_other_pco2_audit_missing_columns")
if "unknown_from_poc_detail_hadm_n" in other_route_quarantine_audit.columns:
    if int(other_route_quarantine_audit["unknown_from_poc_detail_hadm_n"]
        ↵ ].iloc[0]) > 0:
        qa_infos.append("unknown_route_includes_poc_detail")
if "first_gas_without_pco2_anchor_n" in first_gas_anchor_audit.columns:
    if int(first_gas_anchor_audit["first_gas_without_pco2_anchor_n"].iloc[0])
        ↵ > 0:
            qa_warnings.append("first_gas_anchor_without_pco2_detected")
if "source_value_mismatch_n" in hco3_source_counts.columns:
    if int(hco3_source_counts["source_value_mismatch_n"].max()) > 0:
        qa_warnings.append("first_hco3_source_value_mismatch_detected")
qa_other_warn_threshold = float(os.getenv("COHORT_WARN_OTHER_RATE", "0.50"))
qa_other_fail_raw = os.getenv("COHORT_FAIL_OTHER_RATE", "").strip()
qa_other_fail_threshold = float(qa_other_fail_raw) if qa_other_fail_raw else
    ↵ None
if qa_other_fail_threshold is not None and gas_source_other_rate >=
    ↵ qa_other_fail_threshold:
    qa_warnings.append("gas_source_other_rate_fail_threshold_exceeded")

```

```

elif gas_source_other_rate >= qa_other_warn_threshold:
    qa_infos.append("gas_source_other_rate_high")

first_other_pco2_audit_records =
    json.loads(first_other_pco2_audit.to_json(orient="records"))
vitals_outlier_audit_records =
    json.loads(vitals_outlier_audit.to_json(orient="records"))
outliers_audit_records = json.loads(outliers_audit.to_json(orient="records"))
ed_vitals_distribution_summary_records =
    json.loads(ed_vitals_distribution_summary.to_json(orient="records"))
ed_vitals_model_delta_records =
    json.loads(ed_vitals_model_delta.to_json(orient="records"))
gas_overlap_summary_records =
    json.loads(gas_overlap_summary.to_json(orient="records")) if
    "gas_overlap_summary" in globals() else []
pco2_source_distribution_audit_records = json.loads(
    pco2_source_distribution_audit.to_json(orient="records")
)
pco2_window_max_contributor_audit_records = json.loads(
    pco2_window_max_contributor_audit.to_json(orient="records")
)
blood_gas_triplet_completeness_audit_records = json.loads(
    blood_gas_triplet_completeness_audit.to_json(orient="records")
)
qualifying_pco2_distribution_audit_records = json.loads(
    qualifying_pco2_distribution_audit.to_json(orient="records")
)
other_route_quarantine_audit_records = json.loads(
    other_route_quarantine_audit.to_json(orient="records")
)
first_gas_anchor_audit_records =
    json.loads(first_gas_anchor_audit.to_json(orient="records"))
pco2_itemid_qc_audit_records = (
    json.loads(pco2_itemid_qc_audit.to_json(orient="records"))
    if "pco2_itemid_qc_audit" in globals() and not pco2_itemid_qc_audit.empty
    else []
)
timing_integrity_audit_records = (
    json.loads(timing_integrity_audit.to_json(orient="records"))
    if "timing_integrity_audit" in globals()
    else []
)
ventilation_timing_audit_records = (

```

```

        json.loads(ventilation_timing_audit.to_json(orient="records"))
        if "ventilation_timing_audit" in globals()
        else []
    )
anthropometric_cleaning_audit_records = (
    json.loads(anthropometric_cleaning_audit.to_json(orient="records"))
    if "anthropometric_cleaning_audit" in globals()
    else []
)
hco3_integrity_audit_records = (
    json.loads(hco3_source_counts.to_json(orient="records"))
    if "hco3_source_counts" in globals()
    else []
)
blood_gas_itemid_manifest_audit_records = (
    json.loads(blood_gas_itemid_manifest_audit.to_json(orient="records"))
    if "blood_gas_itemid_manifest_audit" in globals() and not
    ↵ blood_gas_itemid_manifest_audit.empty
    else []
)

qa_status = "warning" if qa_warnings else "pass"

# T4 - QA summary artifact.
qa_summary = {
    "ed_spine_rows": int(len(ed_df)),
    "ed_rows": int(len(ed_df)),
    "ed_unique": int(ed_df["ed_stay_id"].nunique()),
    "hadm_rows": hadm_rows,
    "hadm_cc_rows": hadm_cc_rows,
    "icu_link_rate": float(ed_df["icu_intime_first"].notna().mean()) if
    ↵ "icu_intime_first" in ed_df.columns else None,
    "pct_any_gas_0_6h": pct_any_6h,
    "pct_any_gas_0_24h": pct_any_24h,
    "pco2_threshold_0_24h_n":
    ↵ int(pd.to_numeric(ed_df.get("pco2_threshold_0_24h"),
    ↵ errors="coerce").fillna(0).astype(int).sum()),
    "pco2_threshold_0_24h_rate":
    ↵ float(pd.to_numeric(ed_df.get("pco2_threshold_0_24h"),
    ↵ errors="coerce").fillna(0).astype(int).mean()),
    "max_pco2_0_24h_lt_qualifying_n": int(max_pco2_0_24h_lt_qualifying_n),
    "max_pco2_0_6h_lt_qualifying_n": int(max_pco2_0_6h_lt_qualifying_n),
}

```

```

"enrolled_any_n": int(pd.to_numeric(ed_df.get("enrolled_any")),
    ↵   errors="coerce").fillna(0).astype(int).sum()),
"enrolled_any_rate": float(pd.to_numeric(ed_df.get("enrolled_any")),
    ↵   errors="coerce").fillna(0).astype(int).mean()),
"enrollment_route_counts": (
    ed_df.get("enrollment_route", pd.Series(dtype="string"))
        .fillna("NONE")
        .astype("string")
        .value_counts(dropna=False)
        .to_dict()
),
"gas_source_other_rate": gas_source_other_rate,
"gas_source_unknown_rate": gas_source_unknown_rate,
"gas_source_other_rate_warn_threshold": qa_other_warn_threshold,
"gas_source_other_rate_fail_threshold": qa_other_fail_threshold,
"source_unknown_rate": gas_source_unknown_rate,
"gas_source_audit": gas_source_audit,
"gas_source_overlap_summary": gas_overlap_summary_records,
"blood_gas_manifest": {
    "version": BLOOD_GAS_MANIFEST.get("version"),
    "path": BLOOD_GAS_MANIFEST.get("path"),
},
"blood_gas_itemid_manifest_audit":
    ↵ blood_gas_itemid_manifest_audit_records,
"pco2_source_distribution_audit": pco2_source_distribution_audit_records,
"pco2_source_distribution_scope": "ALL_CANDIDATE_PC02_LAB_POC",
"pco2_window_max_contributor_audit":
    ↵ pco2_window_max_contributor_audit_records,
"blood_gas_triplet_completeness_audit":
    ↵ blood_gas_triplet_completeness_audit_records,
"qualifying_pco2_distribution_by_type_audit":
    ↵ qualifying_pco2_distribution_audit_records,
"other_route_quarantine_audit": other_route_quarantine_audit_records,
"first_gas_anchor_audit": first_gas_anchor_audit_records,
"pco2_itemid_qc_audit": pco2_itemid_qc_audit_records,
"poc_itemid_qc_passed": bool(poc_inclusion_enabled),
"poc_itemid_qc_reason": str(poc_inclusion_reason),
"poc_hypcap_0_24h_edstay_n":
    ↵ int(globals().get("poc_hypcap_0_24h_edstay_n", 0)),
"poc_hypcap_0_24h_edstay_rate":
    ↵ float(globals().get("poc_hypcap_0_24h_edstay_rate", 0.0)),
"qualifying_source_branch_counts": (
    ed_df.loc[

```

```

        pd.to_numeric(ed_df.get("pco2_threshold_0_24h"),
←   errors="coerce").fillna(0).astype(int).eq(1),
        "qualifying_source_branch",
    ]
    .fillna("UNKNOWN")
    .astype("string")
    .value_counts(dropna=False)
    .to_dict()
    if "qualifying_source_branch" in ed_df.columns
    else {}
),
"omr_diagnostics": omr_diagnostics if "omr_diagnostics" in globals() else
←   None,
"omr_window_diagnostics": omr_window_diagnostics if
←   "omr_window_diagnostics" in globals() else None,
"anthro_timing_tier_counts": omr_diagnostics.get("selected_tier_counts",
←   {}) if "omr_diagnostics" in globals() else {},
"anthro_timing_tier_rates": omr_diagnostics.get("selected_tier_rates",
←   {}) if "omr_diagnostics" in globals() else {},
"anthro_coverage_audit": anthro_coverage_audit if "anthro_coverage_audit"
←   in globals() else None,
"anthropometric_cleaning_audit": anthropometric_cleaning_audit_records,
"hco3_integrity_audit": hco3_integrity_audit_records,
"charted_anthro_diagnostics": charted_anthro_diagnostics if
←   "charted_anthro_diagnostics" in globals() else None,
"timing_integrity_audit": timing_integrity_audit_records,
"ventilation_timing_audit": ventilation_timing_audit_records,
"unexpected_full_null_fields": unexpected_full_null_fields,
"expected_structural_null_fields": expected_structural_null_fields,
"uom_expectation_scope": uom_expectation_scope,
"uom_expectation_checks": uom_expectation_checks,
"first_other_pco2_audit_scope": first_other_scope_frame,
"first_other_pco2_audit": first_other_pco2_audit_records,
"vitals_outlier_audit": vitals_outlier_audit_records,
"ed_vitals_distribution_summary": ed_vitals_distribution_summary_records,
"ed_vitals_model_delta": ed_vitals_model_delta_records,
"outliers_audit": outliers_audit_records,
"ed_vitals_audit_paths": {
    "distribution_summary": str(ed_vitals_distribution_summary_path),
    "extreme_examples": str(ed_vitals_extreme_examples_path),
    "model_delta": str(ed_vitals_model_delta_path),
},
"blood_gas_audit_paths": {

```

```

"manifest_itemids": str(blood_gas_itemid_manifest_audit_path)
if "blood_gas_itemid_manifest_audit_path" in globals()
else None,
"pco2_itemid_qc": str(pco2_itemid_qc_audit_path)
if "pco2_itemid_qc_audit_path" in globals()
else None,
"pco2_source_distribution": str(pco2_source_distribution_audit_path),
"pco2_window_max_contributor":
    ↳ str(pco2_window_max_contributor_audit_path),
"blood_gas_triplet_completeness": str(
    blood_gas_triplet_completeness_audit_path
),
"qualifying_pco2_distribution_by_type": str(
    qualifying_pco2_distribution_audit_path
),
"other_route_quarantine": str(other_route_quarantine_audit_path),
"first_gas_anchor": str(first_gas_anchor_audit_path),
"hco3_integrity": str(hco3_integrity_audit_path)
if "hco3_integrity_audit_path" in globals()
else None,
},
"timing_integrity_audit_path": str(timing_integrity_audit_path)
if "timing_integrity_audit_path" in globals()
else None,
"ventilation_timing_audit_path": str(ventilation_timing_audit_path)
if "ventilation_timing_audit_path" in globals()
else None,
"anthropometric_cleaning_audit_path":
    ↳ str(anthropometric_cleaning_audit_path)
if "anthropometric_cleaning_audit_path" in globals()
else None,
"qa_status": qa_status,
"qa_warnings": qa_warnings,
"qa_infos": qa_infos,
"run_metadata": RUN_METADATA if "RUN_METADATA" in globals() else None,
}
qa_path = WORK_DIR / "qa_summary.json"
qa_path.write_text(json.dumps(qa_summary, indent=2))
print("Wrote:", qa_path)

```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-23 vitals_outlier_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects

CC-NLP/MIMIC tabular data/prior runs/2026-02-23 outliers_audit.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 ed_vitals_distribution_summary.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 ed_vitals_extreme_examples.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 ed_vitals_model_delta.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 timing_integrity_audit.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 ventilation_timing_audit.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 hco3_integrity_audit.csv
 ED vitals cleaning summary (raw->clean removals):

	vital_name	raw_n	clean_n	raw_removed_n	raw_removed_pct	converted_celsius_like_n
7	first_dbp	5129	5121	8	0.001560	0
9	first_o2sat	5027	5027	0	0.000000	0
3	first_pain	3805	3805	0	0.000000	0
6	first_sbp	5129	5123	6	0.001170	0
1	first_temp	3770	3770	0	0.000000	47
5	triage_dbp	4716	4705	11	0.002332	0
8	triage_o2sat	4711	4711	0	0.000000	0
2	triage_pain	4354	3994	360	0.082683	0
4	triage_sbp	4749	4745	4	0.000842	0
0	triage_temp	4464	4464	0	0.000000	17

	field	missing_n	missing_pct	\
0	abg_before_imv	7166	1.000000	
59	first_abg_po2	7166	1.000000	
74	first_other_po2	7166	1.000000	
75	first_other_time	7166	1.000000	
82	first_vbg_po2	7166	1.000000	
108	icu_stay_id	7166	1.000000	
134	vbg_before_imv	7166	1.000000	
68	first_hco3	6964	0.971811	
93	hco3_band	6964	0.971811	
34	ed_first_rhythm	6788	0.947251	
70	first_other_hypercap_itemid	6768	0.944460	
71	first_other_hypercap_pco2_mmhg	6768	0.944460	
72	first_other_hypercap_source_branch	6768	0.944460	

73	first_other_hypercap_time_0_24h	6768	0.944460
19	deathtime	5921	0.826263
120	min_ph_0_6h	4422	0.617081
118	max_pco2_0_6h	4004	0.558750
27	dt_first_niv_hours_model	3999	0.558052
26	dt_first_niv_hours	3988	0.556517
83	flag_chf	3874	0.540608
84	flag_copd	3874	0.540608
85	flag_neuromuscular	3874	0.540608
86	flag opioid substance	3874	0.540608
87	flag_osa_ohs	3874	0.540608
88	flag_pneumonia	3874	0.540608
78	first_vbg_hypercap_itemid	3655	0.510047
79	first_vbg_hypercap_pco2_mmhg	3655	0.510047
80	first_vbg_hypercap_source_branch	3655	0.510047
81	first_vbg_hypercap_time_0_24h	3655	0.510047
55	first_abg_hypercap_itemid	3649	0.509210
	expectation		
0	missing_column		
59	missing_column		
74	missing_column		
75	missing_column		
82	missing_column		
108	missing_column		
134	missing_column		
68	conditional_sparse		
93	conditional_sparse		
34	conditional_sparse		
70	conditional_sparse		
71	conditional_sparse		
72	conditional_sparse		
73	conditional_sparse		
19	conditional_sparse		
120	conditional_sparse		
118	conditional_sparse		
27	conditional_sparse		
26	conditional_sparse		
83	conditional_sparse		
84	conditional_sparse		
85	conditional_sparse		
86	conditional_sparse		
87	conditional_sparse		

```

88 conditional_sparse
78 conditional_sparse
79 conditional_sparse
80 conditional_sparse
81 conditional_sparse
55 conditional_sparse
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 first_other_pco2_audit.csv
first_other_pco2 audit scope: uom_scope_df
Expected structural null fields: []
Unexpected full-null fields (top 20): []
UOM expectation scope: {'frame': 'final_cc_preview', 'row_count': 7152}
UOM expectation checks: {
    "expected_structural_null_fields": [],
    "structural_null_checks": {},
    "paco2_uom_checks": {
        "poc_abg_paco2_uom": {
            "present": true,
            "paired_value_column": "poc_abg_paco2",
            "value_rows": 4097,
            "missing_uom_when_value_present": 0,
            "non_mmhg_uom_when_value_present": 0,
            "passes": true
        },
        "poc_vbg_paco2_uom": {
            "present": true,
            "paired_value_column": "poc_vbg_paco2",
            "value_rows": 3561,
            "missing_uom_when_value_present": 0,
            "non_mmhg_uom_when_value_present": 0,
            "passes": true
        },
        "poc_other_paco2_uom": {
            "present": true,
            "paired_value_column": "poc_other_paco2",
            "value_rows": 0,
            "missing_uom_when_value_present": 0,
            "non_mmhg_uom_when_value_present": 0,
            "passes": true
        }
    }
}
% any gas panel 0-6h: 50.5

```

```
% any gas panel 0-24h: 98.3
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 gas_source_overlap_summary.csv
% source unknown (panel-level): 60.5
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 pco2_source_distribution_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 pco2_window_max_contributor_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 blood_gas_triplet_completeness_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 qualifying_pco2_distribution_by_type_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 other_route_quarantine_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 first_gas_anchor_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/qa_summary.json
```

3.2 Outputs (ED-stay cohort + long tables)

Rationale: Persist ED-stay analytic datasets and supporting long tables.

```
# Purpose: Write finalized cohort outputs and derivative tables for
↳ downstream analysis workflows.

# Save outputs (dated parquet artifacts archived under prior runs)
from datetime import datetime

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")

cohort_path = prior_runs_dir / f"cohort_ed_stay_{timestamp}.parquet"
ed_vitals_path = prior_runs_dir / f"ed_vitals_long_{timestamp}.parquet"
labs_path = prior_runs_dir / f"labs_long_{timestamp}.parquet"
gas_panels_path = prior_runs_dir / f"gas_panels_{timestamp}.parquet"

ed_df.to_parquet(cohort_path, index=False)
ed_vitals_long.to_parquet(ed_vitals_path, index=False)
labs_long.to_parquet(labs_path, index=False)
panel.to_parquet(gas_panels_path, index=False)
```

```

print("Wrote:", cohort_path)
print("Wrote:", ed_vitals_path)
print("Wrote:", labs_path)
print("Wrote:", gas_panels_path)

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/cohort_ed_stay_20260223_103527.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/ed_vitals_long_20260223_103527.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/labs_long_20260223_103527.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/gas_panels_20260223_103527.parquet

# Purpose: Report missingness for key variables to support transparent
#           data-quality reporting.

# Also export to Excel (tabular dataset) - optional archive output
from datetime import datetime

WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
                             "1"
if WRITE_ARCHIVE_XLSX_EXPORTS:
    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')

    xlsx_path = prior_runs_dir /
    f"mimic_hypcap_EXT_EDstay_bq_gas_{timestamp}.xlsx"

    with pd.ExcelWriter(xlsx_path, engine="openpyxl") as xw:
        ed_df.to_excel(xw, sheet_name="cohort_ed_stay", index=False)
        # Optional: include QA tables if present
        try:
            miss.to_excel(xw, sheet_name="missingness", index=False)
        except Exception:
            pass

    print("Saved:", xlsx_path)
else:
    print("Skipping ED-stay archive workbook export (set
          WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")

```

```

Skipping ED-stay archive workbook export (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

# Purpose: Capture ED presentation context (chief complaint, acuity, and
#           ↵ early vitals) at the ED-stay level.

# Final Excel outputs requested: all encounters + ED chief-complaint only
from datetime import datetime

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')

# 1) All encounters meeting inclusion criteria (admission-level), with ED
#     linkage flags
ed_link = (
    ed_df.groupby("hadm_id", as_index=False)
        .agg(ed_stay_id_first=("ed_stay_id", "first"),
             ↵ n_ed_stays=("ed_stay_id", "nunique"))
)
all_encounters = df.merge(ed_link, on="hadm_id", how="left")
all_encounters["has_ed_encounter"] =
    ↵ all_encounters["n_ed_stays"].fillna(0).astype(int).gt(0).astype(int)
all_encounters["encounter_source"] =
    ↵ all_encounters["has_ed_encounter"].map({1: "ED-linked", 0:
        ↵ "Inpatient-only"})

if WRITE_ARCHIVE_XLSX_EXPORTS:
    all_path = prior_runs_dir /
    ↵ f"mimic_hypercap_EXT_all_encounters_bq_{timestamp}.xlsx"
    with pd.ExcelWriter(all_path, engine="openpyxl") as xw:
        all_encounters.to_excel(xw, sheet_name="all_encounters", index=False)
    print("Saved:", all_path)

# 2) ED chief-complaint-only (ED-stay level)
if "ed_triage_cc" not in ed_df.columns:
    raise KeyError("Column 'ed_triage_cc' not found in ed_df. Ensure ED
        ↵ triage merge ran.")

mask_cc = ed_df["ed_triage_cc"].notna() &
    ↵ (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
ed_cc_only = ed_df.loc[mask_cc].copy()

if WRITE_ARCHIVE_XLSX_EXPORTS:

```

```

cc_path = prior_runs_dir /
↳ f"mimic_hypercap_EXT_EDcc_only_edstay_bq_{timestamp}.xlsx"
with pd.ExcelWriter(cc_path, engine="openpyxl") as xw:
    ed_cc_only.to_excel(xw, sheet_name="ed_cc_only", index=False)
    print("Saved:", cc_path)
else:
    print("Skipping all-encounters/ED-CC archive exports (set
↳ WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")

Skipping all-encounters/ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

# Purpose: Capture ED presentation context (chief complaint, acuity, and
↳ early vitals) at the ED-stay level.

# Final CC output (cohort-only; excludes NLP-derived columns)
from datetime import datetime

# Ensure cc-only frames exist (build from df if upstream optional export
↳ cells were skipped)
if "df_cc" not in globals():
    if "df" not in globals():
        raise NameError("df not found. Run cohort assembly cells first.")
    if "ed_triage_cc" not in df.columns:
        raise KeyError("ed_triage_cc missing in df; cannot build df_cc.")
    _mask_df_cc = df["ed_triage_cc"].notna() &
↳ (df["ed_triage_cc"].astype(str).str.strip() != "")
    df_cc = df.loc[_mask_df_cc].copy()

if "ed_cc_only" not in globals():
    if "ed_df" not in globals():
        raise NameError("ed_df not found. Run the ED-stay build cells
↳ first.")
    if "ed_triage_cc" not in ed_df.columns:
        raise KeyError("ed_triage_cc missing in ed_df; cannot build
↳ ed_cc_only.")
    _mask_cc = ed_df["ed_triage_cc"].notna() &
↳ (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
    ed_cc_only = ed_df.loc[_mask_cc].copy()

# Add ED-stay columns (dedup to earliest ED stay per hadm_id)
ed_tmp = ed_cc_only.copy()
if "ed_intime" in ed_tmp.columns:

```

```

    ed_tmp = ed_tmp.sort_values(["hadm_id", "ed_intime", "ed_stay_id"],
                                na_position="last")
else:
    ed_tmp = ed_tmp.sort_values(["hadm_id", "ed_stay_id"],
                                na_position="last")
ed_first = ed_tmp.drop_duplicates("hadm_id", keep="first")

add_ed_cols = [c for c in ed_first.columns if c not in df_cc.columns]

final_cc = df_cc.merge(ed_first[["hadm_id"] + add_ed_cols], on="hadm_id",
                       how="left")
pre_filter_n = int(len(final_cc))
if "pco2_threshold_0_24h" not in final_cc.columns:
    raise KeyError("Expected pco2_threshold_0_24h in final_cc before
                    export.")
final_cc["pco2_threshold_0_24h"] = (
    pd.to_numeric(final_cc["pco2_threshold_0_24h"], errors="coerce")
    .fillna(0)
    .astype("Int64")
)
if "any_hypercap_icd" not in final_cc.columns:
    final_cc["any_hypercap_icd"] = 0
final_cc["any_hypercap_icd"] = (
    pd.to_numeric(final_cc["any_hypercap_icd"], errors="coerce")
    .fillna(0)
    .astype("Int64")
)
final_cc["enrolled_any"] = (
    (final_cc["pco2_threshold_0_24h"] == 1)
    | (final_cc["any_hypercap_icd"] == 1)
).astype("Int64")
final_cc["enrollment_route"] = np.select(
    [
        (final_cc["any_hypercap_icd"] == 0) &
        (final_cc["pco2_threshold_0_24h"] == 1),
        (final_cc["any_hypercap_icd"] == 1) &
        (final_cc["pco2_threshold_0_24h"] == 0),
        (final_cc["any_hypercap_icd"] == 1) &
        (final_cc["pco2_threshold_0_24h"] == 1),
    ],
    ["GAS_ONLY", "ICD_ONLY", "ICD+GAS"],
    default="NONE",
)

```

```

final_cc = final_cc.loc[final_cc["enrolled_any"] == 1].copy()
post_filter_n = int(len(final_cc))
print(
    "Cohort size comparison (pre vs final ICD OR 0-24h qualifying):",
    f"pre={pre_filter_n:,}",
    f"post={post_filter_n:,}",
    f"removed={pre_filter_n - post_filter_n:,}",
)
print("Final enrollment routes:")
print(final_cc["enrollment_route"].value_counts(dropna=False).to_string())
if "qualifying_source_branch" in final_cc.columns:
    qualifying_source_counts = (
        final_cc["qualifying_source_branch"]
        .fillna("UNKNOWN")
        .astype("string")
        .value_counts(dropna=False)
    )
    qualifying_source_props = (
        qualifying_source_counts / max(len(final_cc), 1)
    ).rename("proportion")
    print("Final qualifying source mix (LAB vs POC):")
    print(
        pd.concat(
            [qualifying_source_counts.rename("count"),
            ← qualifying_source_props],
            axis=1,
            ).to_string()
    )
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")

cohort_contract = validate_cohort_contract(final_cc)
cohort_contract_report = {
    "generated_utc": datetime.utcnow().isoformat() + "Z",
    "status": cohort_contract["status"],
    "contracts": {"cohort": cohort_contract},
    "findings": cohort_contract["findings"],
}
cohort_contract_paths = write_contract_report(cohort_contract_report,
← work_dir=WORK_DIR)
print("Wrote:", cohort_contract_paths["contract_report_path"])

```

```

if cohort_contract_paths["failed_contract_path"].exists():
    print("Wrote:", cohort_contract_paths["failed_contract_path"])

contract_warning_codes = [
    finding.get("code")
    for finding in cohort_contract.get("findings", [])
    if str(finding.get("severity", "")).lower() == "warning"
]
contract_error_codes = [
    finding.get("code")
    for finding in cohort_contract.get("findings", [])
    if str(finding.get("severity", "")).lower() == "error"
]
if "qa_path" not in globals():
    qa_path = WORK_DIR / "qa_summary.json"
if qa_path.exists():
    qa_payload = json.loads(qa_path.read_text())
else:
    qa_payload = {}
local_status = str(qa_payload.get("qa_status", "pass")).strip().lower()
local_warning_present = bool(qa_payload.get("qa_warnings", [])) or
    local_status == "warning"
local_error_present = local_status == "fail"
contract_status = str(cohort_contract.get("status", "pass")).strip().lower()
if local_error_present or contract_status == "fail":
    qa_status_final = "fail"
elif local_warning_present or contract_status == "warning":
    qa_status_final = "warning"
else:
    qa_status_final = "pass"
qa_payload["contract_status"] = contract_status
qa_payload["contract_warning_codes"] = contract_warning_codes
qa_payload["contract_error_codes"] = contract_error_codes
qa_payload["qa_status_final"] = qa_status_final
qa_payload["cohort_contract_report_path"] =
    str(cohort_contract_paths["contract_report_path"])
qa_path.write_text(json.dumps(qa_payload, indent=2))
print("Updated:", qa_path, f"(qa_status_final={qa_status_final})")

contract_mode = os.getenv("PIPELINE_CONTRACT_MODE", "fail").strip().lower()
if contract_mode not in {"fail", "warn"}:
    raise ValueError("PIPELINE_CONTRACT_MODE must be one of {'fail', 'warn'}")

```

```

if cohort_contract["status"] == "fail" and contract_mode == "fail":
    raise ValueError(
        "Cohort contract checks failed. "
        f"See {cohort_contract_paths['failed_contract_path']}"))
)

# Keep cohort export focused on analysis-facing composition fields.
# Row-level gas-source QA diagnostics and redundant enrollment aliases stay
# in QA artifacts.
cohort_export_drop_columns = [
    "enrolled_any",
    "enrolled_any_icd_union_secondary",
    "poc_inclusion_enabled",
    "poc_inclusion_reason",
    "race_hosp_raw",
    "race_ed_raw",
    "gas_source_unknown_rate",
    "gas_source_other_rate",
    "gas_source_inference_primary_tier",
    "gas_source_hint_conflict_rate",
    "gas_source_resolved_rate",
    "gas_source_tier_specimen_text_rate",
    "gas_source_tier_label_fluid_rate",
    "gas_source_tier_cluster_inheritance_rate",
    "gas_source_tier_panel_cooccurrence_rate",
    "gas_source_tier_fallback_unknown_rate",
    "lab_abg_po2",
    "lab_vbg_po2",
    "lab_other_po2",
    "poc_abg_po2",
    "poc_vbg_po2",
    "poc_other_po2",
    "max_pco2_0_6h_source_branch",
    "max_pco2_0_24h_source_branch",
]
cohort_export_drop_existing = [
    column for column in cohort_export_drop_columns if column in
    final_cc.columns
]
if cohort_export_drop_existing:
    final_cc = final_cc.drop(columns=cohort_export_drop_existing)
    print(
        "Dropped export-only QA/redundant columns:",

```

```

        ", ".join(cohort_export_drop_existing),
    )

dated_out_path = prior_runs_dir / f"{out_date} MIMICIV all with CC.xlsx"
canonical_out_path = DATA_DIR / "MIMICIV all with CC.xlsx"

# Write via temporary files then atomically replace targets to avoid
# zero-byte outputs on interruption.
def _atomic_write_excel(df_obj, target_path: Path) -> None:
    tmp_path = target_path.with_suffix(target_path.suffix + ".tmp")
    df_obj.to_excel(tmp_path, index=False)
    os.replace(tmp_path, target_path)

def _sha256_file(path: Path) -> str:
    hasher = hashlib.sha256()
    with path.open("rb") as file_handle:
        for chunk in iter(lambda: file_handle.read(1024 * 1024), b ""):
            hasher.update(chunk)
    return hasher.hexdigest()

_atomic_write_excel(final_cc, dated_out_path)
_atomic_write_excel(final_cc, canonical_out_path)
canonical_sha256 = _sha256_file(canonical_out_path)
dated_sha256 = _sha256_file(dated_out_path)
print(
    "EXPORT_SHAPE cohort",
    f"rows={len(final_cc)}:,}",
    f"cols={final_cc.shape[1]}:,}",
    f"canonical_sha256={canonical_sha256}",
)
cohort_manifest = collect_run_manifest(
    WORK_DIR,
    run_id=f"cohort_{datetime.now().strftime('%Y%m%d_%H%M%S')}",
)
cohort_manifest["stage"] = "cohort"
cohort_manifest["outputs"] = {
    "canonical_output_path": str(canonical_out_path),
    "canonical_output_sha256": canonical_sha256,
    "dated_output_path": str(dated_out_path),
    "dated_output_sha256": dated_sha256,
}

```

```

"gas_source_audit_path": str(gas_source_audit_path) if
    "gas_source_audit_path" in globals() else None,
"gas_source_overlap_summary_path": str(gas_overlap_summary_path) if
    "gas_overlap_summary_path" in globals() else None,
"pco2_source_distribution_audit_path":
    str(pco2_source_distribution_audit_path)
if "pco2_source_distribution_audit_path" in globals()
else None,
"pco2_window_max_contributor_audit_path": str(
    pco2_window_max_contributor_audit_path
)
if "pco2_window_max_contributor_audit_path" in globals()
else None,
"qualifying_pco2_distribution_by_type_audit_path": str(
    qualifying_pco2_distribution_audit_path
)
if "qualifying_pco2_distribution_audit_path" in globals()
else None,
"anthropometrics_coverage_audit_path": str(anthro_coverage_audit_path) if
    "anthro_coverage_audit_path" in globals() else None,
"ed_vitals_distribution_summary_path":
    str(ed_vitals_distribution_summary_path) if
    "ed_vitals_distribution_summary_path" in globals() else None,
"ed_vitals_extreme_examples_path": str(ed_vitals_extreme_examples_path)
    if "ed_vitals_extreme_examples_path" in globals() else None,
"ed_vitals_model_delta_path": str(ed_vitals_model_delta_path) if
    "ed_vitals_model_delta_path" in globals() else None,
"outliers_audit_path": str(outliers_audit_path) if "outliers_audit_path"
    in globals() else None,
}
cohort_manifest["gas_source"] = {
    "mode": gas_source_mode if "gas_source_mode" in globals() else None,
    "source_rates": gas_source_audit.get("source_rates", {}) if
        "gas_source_audit" in globals() else {},
    "tier_rates": gas_source_audit.get("tier_rates", {}) if
        "gas_source_audit" in globals() else {},
}
cohort_manifest["anthropometrics"] = (
    anthro_coverage_audit if "anthro_coverage_audit" in globals() else None
)
cohort_manifest_path = prior_runs_dir / f"{out_date}"
    cohort_run_manifest.json"
cohort_manifest_path.write_text(json.dumps(cohort_manifest, indent=2))

```

```

print("Saved:", dated_out_path)
print("Saved:", canonical_out_path)
print("Wrote:", cohort_manifest_path)
print("Base rows:", len(df_cc), "| Added from ED-stay:", len(add_ed_cols))
print("Total columns:", len(final_cc.columns))

Cohort size comparison (pre vs final ICD OR 0-24h qualifying): pre=7,152 post=7,152 removed=0
Final enrollment routes:
enrollment_route
GAS_ONLY      5169
ICD+GAS       1077
ICD_ONLY      906
Final qualifying source mix (LAB vs POC):
              count    proportion
qualifying_source_branch
LAB             5383     0.752657
UNKNOWN         906      0.126678
POC             863      0.120666
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project/CC-NLP/debug/contracts/20260223_173528/contract_report.json
Updated: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project/CC-NLP/qa_summary.json (qa_status_final=warning)
Dropped export-only QA/redundant columns: enrolled_any, enrolled_any_icd_union_secondary, raw_EXPORT_SHAPE cohort rows=7,152 cols=297 canonical_sha256=9c0820da8a1ba4d5b9831685ab34befd357
Saved: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 MIMICIV all with CC.xlsx
Saved: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project/CC-NLP/MIMIC tabular data/MIMICIV all with CC.xlsx
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project/CC-NLP/MIMIC tabular data/prior runs/2026-02-23 cohort_run_manifest.json
Base rows: 7152 | Added from ED-stay: 176
Total columns: 297

```

```

# Purpose: Extract and standardize first ABG/VBG physiology fields for
#           → baseline characterization.

# Data dictionary (OSF-style) for final cohort output
from datetime import datetime
import numpy as np

# Choose target dataset for dictionary

```

```

if "final_cc" in globals():
    dd_target = final_cc
elif "df_cc" in globals():
    dd_target = df_cc
elif "df" in globals():
    dd_target = df
else:
    raise NameError("No cohort dataframe found for data dictionary (expected
        final_cc/df_cc/df).")

# Helper functions

def _readable(name: str) -> str:
    return name.replace("_", " ").strip().title()

def _infer_units(col: str) -> str:
    if col.endswith("_hours"):
        return "hours"
    if col.endswith("_time") or col.endswith("_intime") or
       col.endswith("_outtime") or col.endswith("_date"):
        return "datetime"
    if col.startswith("flag_") or col.endswith("_flag") or
       col.endswith("_before_imv"):
        return "binary (0/1)"
    if col.endswith("_paco2") or col.endswith("_pco2"):
        return "mmHg"
    if col.endswith("_ph"):
        return "pH"
    if col.endswith("_lactate"):
        return "mmol/L"
    return ""

def _allowed_values(series):
    s = series.dropna()
    if s.empty:
        return ""
    if s.dtype.kind in "biu":
        vals = sorted(map(int, s.unique()))
        if len(vals) <= 15:
            return ", ".join(map(str, vals))
        return f"{min(vals)}-{max(vals)}"

```

```

if s.dtype.kind in "f":
    return f"{{np.nanmin(s):.3g}-{np.nanmax(s):.3g}}"
# object/categorical: use bounded sampling to keep dictionary generation
# fast.
sample_n = 5000
s_str = s.astype(str)
if len(s_str) > sample_n:
    s_sample = s_str.head(sample_n)
    sample_note = f"sampled {sample_n}/{len(s_str)} rows"
else:
    s_sample = s_str
    sample_note = None
unique_n = int(s_sample.nunique(dropna=True))
top_vals = s_sample.value_counts(dropna=False).head(15).index.tolist()
if unique_n > len(top_vals):
    tail = f" ... ({unique_n} unique"
    if sample_note:
        tail += f"; {sample_note}"
    tail += ")"
    return ", ".join(top_vals) + tail
if sample_note:
    return ", ".join(top_vals) + f" ({sample_note})"
return ", ".join(top_vals)

META = {
    # IDs
    "subject_id": {
        "label": "Subject identifier",
        "definition": "Identifier for a unique patient in MIMIC-IV.",
        "source": "mimiciv_hosp.patients",
        "derivation": "as recorded",
    },
    "hadm_id": {
        "label": "Hospital admission identifier",
        "definition": "Identifier for a hospital admission.",
        "source": "mimiciv_hosp.admissions",
        "derivation": "as recorded",
    },
    "ed_stay_id": {
        "label": "ED stay identifier",
        "definition": "Identifier for an ED stay.",
        "source": "mimiciv_ed.edstays",
    }
}

```

```

    "derivation": "as recorded",
},
"ed_intime": {
    "label": "ED arrival time",
    "definition": "ED arrival time for this ED stay.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
    "units": "datetime",
},
"ed_outtime": {
    "label": "ED departure time",
    "definition": "ED departure time for this ED stay.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
    "units": "datetime",
},
"ed_intime_first": {
    "label": "First ED arrival time (admission)",
    "definition": "Earliest ED arrival time among ED stays linked to the
    ↵ admission.",
    "source": "mimiciv_ed.edstays",
    "derivation": "min(edstays.intime) per hadm_id",
    "units": "datetime",
},
# Demographics / triage
"ed_gender": {
    "label": "ED gender",
    "definition": "Gender as recorded in ED stay table.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
},
"ed_race": {
    "label": "ED race",
    "definition": "Race as recorded in ED stay table.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
},
"hosp_race": {
    "label": "Hospital race",
    "definition": "Race as recorded in admissions table.",
    "source": "mimiciv_hosp.admissions",
    "derivation": "as recorded",
},
}

```

```

"ed_triage_cc": {
    "label": "ED chief complaint",
    "definition": "Chief complaint recorded at ED triage.",
    "source": "mimiciv_ed.triage",
    "derivation": "as recorded",
},
"ed_triage_acuity": {
    "label": "ED triage acuity",
    "definition": "Triage acuity level recorded at ED presentation.",
    "source": "mimiciv_ed.triage",
    "derivation": "as recorded",
},
# Outcomes
"hospital_expire_flag": {
    "label": "Hospital expire flag",
    "definition": "Hospital mortality indicator from admissions table.",
    "source": "mimiciv_hosp.admissions",
    "derivation": "as recorded",
    "units": "binary (0/1)",
},
"in_hospital_death": {
    "label": "In-hospital death",
    "definition": "Indicator for in-hospital death during admission.",
    "source": "mimiciv_hosp.admissions",
    "derivation": "(hospital_expire_flag == 1) OR (deathtime not null)",
    "units": "binary (0/1)",
},
# Ventilation
"imv_flag": {
    "label": "IMV flag",
    "definition": "Indicator of invasive mechanical ventilation during admission.",
    "source": "ICD procedures + ICU chartevents",
    "derivation": "max(IMV ICD flag, IMV chart flag)",
    "units": "binary (0/1)",
},
"niv_flag": {
    "label": "NIV flag",
    "definition": "Indicator of noninvasive ventilation during admission.",
    "source": "ICD procedures + ICU chartevents",
    "derivation": "max(NIV ICD flag, NIV chart flag)",
    "units": "binary (0/1)",
}

```

```

},
"first_imv_time": {
    "label": "First IMV time",
    "definition": "Earliest charted time of IMV in ICU chartevents (if
    ↵ present).",
    "source": "mimiciv_icu.chartevents",
    "derivation": "min charttime among IMV charted events",
    "units": "datetime",
},
"first_niv_time": {
    "label": "First NIV time",
    "definition": "Earliest charted time of NIV in ICU chartevents (if
    ↵ present).",
    "source": "mimiciv_icu.chartevents",
    "derivation": "min charttime among NIV charted events",
    "units": "datetime",
},
"dt_first_imv_hours": {
    "label": "Hours to first IMV",
    "definition": "Hours from first ED arrival to first IMV time.",
    "source": "Derived",
    "derivation": "(first_imv_time - ed_intime_first) in hours",
    "units": "hours",
},
"dt_first_niv_hours": {
    "label": "Hours to first NIV",
    "definition": "Hours from first ED arrival to first NIV time.",
    "source": "Derived",
    "derivation": "(first_niv_time - ed_intime_first) in hours",
    "units": "hours",
},
"abg_before_imv": {
    "label": "ABG before IMV",
    "definition": "Indicator that first ABG time precedes first IMV
    ↵ time.",
    "source": "Derived",
    "derivation": "first_abg_time < first_imv_time",
    "units": "binary (0/1)",
},
"vbg_before_imv": {
    "label": "VBG before IMV",
    "definition": "Indicator that first VBG time precedes first IMV
    ↵ time.",
}

```

```

    "source": "Derived",
    "derivation": "first_vbg_time < first_imv_time",
    "units": "binary (0/1)",
},
# Gas/ABG/VBG
"abg_hypercap_threshold": {
    "label": "ABG hypercapnia threshold met",
    "definition": "Indicator that any arterial pCO2  45 mmHg in hosp/ICU
    ↳ pCO2 extraction.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "max(arterial pCO2  45) per hadm_id",
    "units": "binary (0/1)",
},
"vbg_hypercap_threshold": {
    "label": "VBG hypercapnia threshold met",
    "definition": "Indicator that any venous pCO2  50 mmHg in hosp/ICU
    ↳ pCO2 extraction.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "max(venous pCO2  50) per hadm_id",
    "units": "binary (0/1)",
},
"unknown_hypercap_threshold": {
    "label": "Unknown-source hypercapnia threshold met",
    "definition": "Indicator that any unknown/indeterminate-source pCO2
    ↳ 50 mmHg in hosp/ICU pCO2 extraction.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "max(unknown-source pCO2  50) per hadm_id",
    "units": "binary (0/1)",
},
"first_gas_time": {
    "label": "First gas panel time",
    "definition": "Earliest gas panel time in the ED 0-24h window.",
    "source": "mimiciv_hosp.labevents (gas panels)",
    "derivation": "min panel_time within ED 0-24h window",
    "units": "datetime",
},
"dt_first_qualifying_gas_hours": {
    "label": "Hours to first gas panel",
    "definition": "Hours from first ED arrival to first gas panel time.",
    "source": "Derived",
    "derivation": "(first_gas_time - ed_intime_first) in hours",
    "units": "hours",
}
},

```

```

"first_other_time": {
    "label": "First other-source gas time",
    "definition": "Earliest time of a pCO2 measurement with
        ↳ other/unspecified source classification.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "min(other-source pCO2 charttime) across LAB and POC",
    "units": "datetime",
},
"presenting_hypercapnia": {
    "label": "Presenting (6h) gas timing",
    "definition": "Indicator that first gas panel occurred within 6h of
        ↳ first ED arrival.",
    "source": "Derived",
    "derivation": "dt_first_qualifying_gas_hours <= 6",
    "units": "binary (0/1)",
},
"late_hypercapnia": {
    "label": "Late (>6h) gas timing",
    "definition": "Indicator that first gas panel occurred >6h after
        ↳ first ED arrival.",
    "source": "Derived",
    "derivation": "dt_first_qualifying_gas_hours > 6",
    "units": "binary (0/1)",
},
"gas_source_other_rate": {
    "label": "Gas source other/unknown rate",
    "definition": "Proportion of gas panels classified as
        ↳ other/unspecified source per ED stay.",
    "source": "Derived from gas panel source inference",
    "derivation": "mean(source == 'other') per ed_stay_id",
    "units": "proportion",
},
# Comorbidities
"flag_copd": {
    "label": "COPD/emphysema flag",
    "definition": "Indicator for COPD/emphysema ICD codes in ED or
        ↳ hospital diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes J43/J44 (ED OR hospital)",
    "units": "binary (0/1)",
},
"flag_osa_ohs": {
    "label": "OSA/OHS flag",

```

```

    "definition": "Indicator for OSA/OHS ICD codes in ED or hospital
      ↳ diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes G473/E662 (ED OR hospital)",
    "units": "binary (0/1)",
  },
  "flag_chf": {
    "label": "CHF flag",
    "definition": "Indicator for CHF ICD codes in ED or hospital
      ↳ diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefix I50 (ED OR hospital)",
    "units": "binary (0/1)",
  },
  "flag_neuromuscular": {
    "label": "Neuromuscular flag",
    "definition": "Indicator for neuromuscular ICD codes in ED or
      ↳ hospital diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes G12/G70/G71 (ED OR hospital)",
    "units": "binary (0/1)",
  },
  "flag_opioid_substance": {
    "label": "Opioid/substance flag",
    "definition": "Indicator for opioid/substance ICD codes in ED or
      ↳ hospital diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes F11/T40/F13 (ED OR hospital)",
    "units": "binary (0/1)",
  },
  "flag_pneumonia": {
    "label": "Pneumonia flag",
    "definition": "Indicator for pneumonia ICD codes in ED or hospital
      ↳ diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes J12-J18 (ED OR hospital)",
    "units": "binary (0/1)",
  },
}

rows = []
for col in dd_target.columns:
  info = META.get(col, {})

```

```

series = dd_target[col]
rows.append({
    "variable": col,
    "label": info.get("label", _readable(col)),
    "units": info.get("units", _infer_units(col)),
    "allowed_values": info.get("allowed_values",
        ↪ _allowed_values(series)),
    "definition": info.get("definition", "UNCONFIRMED"),
    "synonyms": info.get("synonyms", ""),
    "description": info.get("description", ""),
    "source": info.get("source", "UNCONFIRMED"),
    "derivation": info.get("derivation", "UNCONFIRMED"),
    "dtype": str(series.dtype),
    "example_value": series.dropna().iloc[0] if series.dropna().shape[0]
        ↪ else None,
})
data_dict = pd.DataFrame(rows)

out_date = datetime.now().strftime("%Y-%m-%d")
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_xlsx = prior_runs_dir / f"{out_date} MIMICIV all with
    ↪ CC_data_dictionary.xlsx"
out_csv = prior_runs_dir / f"{out_date} MIMICIV all with
    ↪ CC_data_dictionary.csv"

data_dict.to_excel(out_xlsx, index=False)
data_dict.to_csv(out_csv, index=False)

print("Saved data dictionary:", out_xlsx)
print("Saved data dictionary:", out_csv)
print("Rows:", len(data_dict))

```