

MIMICIV_hypercap_EXT_cohort

Table of contents

1 MIMIC-IV Hypercapnia Cohort — ICD Physiologic Thresholds (BigQuery)	2
1.1 Pipeline stages (readable map)	3
1.2 SQL registry (centralized query templates)	3
2 MIMIC-IV on BigQuery	4
2.1 Environment Bootstrap & Smoke Test	4
2.2 0. Prerequisites (one-time)	4
2.3 Data Generation	66
2.3.1 1) ICD cohort flags (hypercapnic respiratory failure)	66
2.3.2 2) Blood gas inclusion thresholds — ANY PaCO meeting cutoffs (LAB + POC)	70
2.3.3 3) Cohort union (ICD thresholds) and hadm_list for downstream queries	74
2.3.4 4) First ABG and First VBG (LAB + POC, standardized to mmHg) . .	76
2.3.5 5) Demographics & outcomes	84
2.3.6 6) NIH/OMB race & ethnicity (ED + Hospital)	88
2.3.7 7) ED triage (linked to hadm) and first ED vitals	91
2.3.8 8) ICU meta (first ICU stay, LOS days)	94
2.3.9 9) Ventilation flags (ICD procedures)	95
2.3.10 10) Assemble final DataFrame	100
2.3.11 11) Sanity checks	112
2.3.12 PDF-ready long tables	113
2.3.13 12) Save to Excel	119
2.3.14 Create Annotation Dataset	119
3 ED-stay cohort expansion (timing, severity, comorbidity, outcomes)	121
3.0.1 Phase 0 — Inventory & missing-field registry	122
3.0.2 Phase 1 — ED encounter spine and ED enrichment (one row per ED stay)	125
3.0.3 Phase 2 — Hospital admission context and outcomes	129
3.0.4 Phase 3 — ICU timing and LOS	131
3.0.5 Phase 4 — ED longitudinal vitals (0–6h)	132

3.0.6	Phase 5 — Robust lab discovery + gas panels	134
3.0.7	Phase 5C — ICU POC blood gases (chartevents, optional)	154
3.0.8	Phase 6 — BMI/anthropometrics (OMR)	166
3.0.9	Phase 7 — ICD comorbidity flags	173
3.0.10	Phase 8 — Timing phenotypes and derived bands	177
3.1	QA & Data Fidelity	181
3.2	Outputs (ED-stay cohort + long tables)	199

1 MIMIC-IV Hypercapnia Cohort — ICD Physiologic Thresholds (BigQuery)

TODO:

[] want to add ED-rendered diagnoses - a flag to split off whether the hypercapnic respiratory failure ICD was rendered in the ED or during the hospital stay. [] are the installation instructions at the beginning of this notebook consistent with the way we intend the notebook to be used? is it also consistent with the README.md instructions?

Goal: Build an admissions-level tabular dataset that **enrolls** any hospital admission (`hadm_id`) meeting **any** of:

1. **ICD** codes for hypercapnic respiratory failure (legacy cohort).
2. **Any arterial blood gas** (LAB or POC) with **PaCO 45.0 mmHg** anywhere during the episode.
3. **Any venous blood gas** (LAB or POC) with **PaCO 50.0 mmHg** anywhere during the episode.

Then, keep all downstream columns/logic from the current workflow: - Per-code ICD indicators and an `any_hypercap_icd` flag. - Robust extraction of **first ABG** and **first VBG** (across LAB + POC) with standardized units (mmHg) and pairing logic. - Demographics/outcomes, NIH/OMB race & ethnicity, ED triage + first ED vitals, ICU meta (first stay + LOS), ventilation flags. - Sanity checks.

Assumptions - You already configured BigQuery auth (`gcloud auth application-default login`) and `.env` variables as in the previous notebook. - The PhysioNet hosting project is `physionet-data`. - Datasets exist (e.g., `mimiciv_3_1_hosp`, `mimiciv_3_1_icu`, and an ED dataset such as `mimiciv_ed`). This notebook auto-detects the ED dataset.

Execution timing note: To capture per-cell runtimes, enable cell execution timing in VS Code (Notebook: Show Cell Execution Time) or enable ExecuteTime in Jupyter. Then re-run and save to allow runtime summaries.

1.1 Pipeline stages (readable map)

1. Stage A — Config & helpers: dataset IDs, thresholds, shared helper functions.
2. Stage B — Cohort spine: ICD + gas thresholds → hadm_list / ed_stay_list.
3. Stage C — Bulk pulls: one query per table where possible.
4. Stage D — Transforms: panels, flags, timing, comorbidities, vitals.
5. Stage E — QA checks: deterministic assertions and range checks.
6. Stage F — Outputs: parquet + Excel exports.

1.2 SQL registry (centralized query templates)

All BigQuery SQL is defined in one place below, then referenced by name in subsequent cells.

```
# Purpose: Build ABG/VBG hypercapnia threshold flags from lab and ICU POC
#           ↳ pCO2 measurements.

import sys
print(sys.executable)

# Central SQL registry (define all query templates here)
SQL = {}

def sql(name: str) -> str:
    if name not in SQL:
        raise KeyError(f"SQL template not found: {name}")
    return SQL[name]

# SQL templates (populated below in-place to keep notebook linear)
# Names: admit_sql, co2_thresholds_sql, cohort_icd_sql, counts_sql, demo_sql,
#         ↳ ditems_sql, ed_counts_sql, ed_first_vitals_sql, ed_spine_sql,
#         ↳ ed_to_icu_sql, ed_triage_sql, ed_vitals_sql, icd_sql, icu_meta_sql,
#         ↳ icu_sql, labitems_sql, labs_sql, vent_chart_sql, vent_sql

/Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/HyperCC-NLP/.venv/bin/python
```

Rationale: Define a reproducible, admission-level cohort that captures hypercapnia using complementary diagnostic (ICD) and physiologic (blood gas) criteria.

2 MIMIC-IV on BigQuery

2.1 Environment Bootstrap & Smoke Test

Purpose: make a clean, reproducible start on a new machine.

Outcome: verify auth, project config, and dataset access; provide a reusable BigQuery runner for the build notebook.

Rationale: Establish the BigQuery environment and dataset configuration so queries are consistent and reproducible across runs.

2.2 0. Prerequisites (one-time)

Accounts & access - PhysioNet access to MIMIC-IV on BigQuery; in BigQuery Console star project `physionet-data`. - A Google Cloud **Project ID** with BigQuery API enabled (this is your **billing** project).

CLI & environment - Google Cloud SDK (`gcloud`) installed and on PATH. - Python environment created with `uv` (see README) and Jupyter kernel selected. - A project-local `.env` with the variables below.

.env variables

```
MIMIC_BACKEND=bigquery
WORK_PROJECT=<your-billing-project-id>
BQ_PHYSIONET_PROJECT=physionet-data
BQ_DATASET_HOSP=mimiciv_3_1_hosp
BQ_DATASET_ICU=mimiciv_3_1_icu
BQ_DATASET_ED=mimiciv_ed
WORK_DIR=/path/to/Hypercap-CC-NLP
# GOOGLE_APPLICATION_CREDENTIALS=/path/to/service-account.json
```

Command line quickstart

```
brew install --cask google-cloud-sdk
gcloud init
gcloud auth application-default login
gcloud services enable bq.googleapis.com --project
  ↳ <your-billing-project-id>
ls -l ~/.config/gcloud/application_default_credentials.json
```

Rationale: Verify access and credentials up front to prevent silent failures later in the pipeline.

```

# Purpose: Set up project paths, environment variables, and BigQuery client
    ↵ connections for reproducible execution.

# --- Imports & environment
import os, re, json, math, textwrap, sys, hashlib
from pathlib import Path
from typing import Any, Mapping, Sequence

import numpy as np
import pandas as pd

from google.cloud import bigquery
from google.oauth2 import service_account
from dotenv import load_dotenv

load_dotenv()

WORK_DIR = Path(os.getenv("WORK_DIR", Path.cwd())).expanduser().resolve()
DATA_DIR = WORK_DIR / "MIMIC tabular data"
DATA_DIR.mkdir(parents=True, exist_ok=True)
SRC_DIR = WORK_DIR / "src"
if SRC_DIR.exists() and str(SRC_DIR) not in sys.path:
    sys.path.insert(0, str(SRC_DIR))

# Local helper functions (core cohort logic must remain embedded in this
    ↵ notebook)
OMR_RESULT_NAMES = ("bmi", "height", "weight")
OMR_OUTPUT_COLUMNS = (
    "bmi_closest_pre_ed",
    "height_closest_pre_ed",
    "weight_closest_pre_ed",
)
OMR_PROVENANCE_COLUMNS = (
    "anthro_timing_tier",
    "anthro_days_offset",
    "anthro_chartdate",
    "anthro_timing_uncertain",
    "anthro_source",
    "anthro_obstime",
    "anthro_hours_offset",
    "anthro_timing_basis",
)

```

```

OMR_TIMING_TIERS = ("pre_ed_365", "post_ed_365", "missing")

EXPECTED_STRUCTURAL_NULL_FIELDS = (
    "poc_abg_ph_uom",
    "poc_vbg_ph_uom",
    "poc_other_ph_uom",
)
PACO2_VALUE_UOM_PAIRS = (
    ("poc_abg_paco2", "poc_abg_paco2_uom"),
    ("poc_vbg_paco2", "poc_vbg_paco2_uom"),
    ("poc_other_paco2", "poc_other_paco2_uom"),
)
DEFAULT_VITALS_MODEL_RANGES: dict[str, tuple[float, float]] = {
    "ed_triage_hr": (20.0, 250.0),
    "ed_first_hr": (20.0, 250.0),
    "ed_triage_rr": (4.0, 80.0),
    "ed_first_rr": (4.0, 80.0),
}
DEFAULT_GAS_MODEL_RANGES: dict[str, tuple[float, float]] = {
    "first_ph": (6.8, 7.8),
    "first_pco2": (10.0, 200.0),
    "first_other_pco2": (10.0, 200.0),
    "first_lactate": (0.0, 30.0),
}
_NUMERIC_TOKEN_PATTERN = re.compile(r"(-?\d+(?:\.\d+)?)")
_ARTERIAL_HINT_PATTERN = re.compile(
    r"(arterial|abg|a[- ]?line|art line|\bart\b|\bartery\b)",
    re.I,
)
_VENOUS_HINT_PATTERN = re.compile(
    r"(venous|vbg|cvbg|mixed venous|central venous|\bven\b)",
    re.I,
)
_P02_LABEL_DENY_PATTERN_DEFAULT = re.compile(
    r"^(?:(^|[a-z])p(?:(a)?\s*\s*2(?:(^|[a-z]|$)|\bpo2\b|\bpao2\b|oxygen|o2\s*)"
    "\s*(?:(sat|saturation)|sao2|spo2|fio2|et|\s*co2|end|"
    "\s*tidal|t|\s*co2|tco2|total|\s*co2|hco3|bicar|v|\s*co2|vco2|co2|\s*(?:""
    "\s*prod|production|elimin|elimination))",
)

```

```

        re.I,
    )

BLOOD_GAS_SPEC_PATH = WORK_DIR / "specs" / "blood_gas_itemids.json"

def _coerce_itemid_list(values: Sequence[Any], *, key_name: str) ->
    list[int]:
    cleaned: list[int] = []
    for raw in values:
        try:
            cleaned.append(int(raw))
        except (TypeError, ValueError) as exc:
            raise ValueError(f"Invalid itemid in {key_name}: {raw!r}") from
            exc
    return sorted(set(cleaned))

def _sql_int_array(values: Sequence[int], *, empty_sentinel: int =
    -999999999) -> str:
    """Return a BigQuery INT64 array literal from Python integers."""
    resolved = [int(v) for v in values]
    if not resolved:
        resolved = [int(empty_sentinel)]
    return "[" + ", ".join(str(v) for v in resolved) + "]"

def load_blood_gas_itemid_manifest(path: Path) -> dict[str, Any]:
    """Load and validate versioned blood-gas itemid manifest."""
    if not path.exists():
        raise FileNotFoundError(
            f"Blood-gas manifest missing at {path}. Add"
            "specs/blood_gas_itemids.json."
        )
    payload = json.loads(path.read_text())
    if not isinstance(payload, dict):
        raise ValueError("Blood-gas manifest must be a JSON object.")

    for section in ("lab", "icu", "denylist"):
        if section not in payload or not isinstance(payload[section], dict):
            raise ValueError(f"Blood-gas manifest missing required section"
                            '{section}')."

```

```

lab = payload["lab"]
icu = payload["icu"]
denylist = payload["denylist"]

lab_pco2 = _coerce_itemid_list(lab.get("pco2_itemids", []),
↪ key_name="lab.pco2_itemids")
lab_pco2_excluded = _coerce_itemid_list(
    lab.get("excluded_pco2_itemids", []),
    key_name="lab.excluded_pco2_itemids",
)
lab_ph = _coerce_itemid_list(lab.get("ph_itemids", []),
↪ key_name="lab.ph_itemids")
lab_hco3 = _coerce_itemid_list(lab.get("hco3_itemids", []),
↪ key_name="lab.hco3_itemids")
lab_spec = _coerce_itemid_list(
    lab.get("specimen_type_itemids", []),
    key_name="lab.specimen_type_itemids",
)
if not lab_pco2:
    raise ValueError("Blood-gas manifest requires at least one
↪ lab.pco2_itemids value.")
if not lab_ph:
    raise ValueError("Blood-gas manifest requires at least one
↪ lab.ph_itemids value.")
if not lab_hco3:
    raise ValueError("Blood-gas manifest requires at least one
↪ lab.hco3_itemids value.")
if not lab_spec:
    raise ValueError(
        "Blood-gas manifest requires at least one
        ↪ lab.specimen_type_itemids value."
    )

blood_fluid_terms = [str(term).strip().lower() for term in
↪ lab.get("blood_fluid_terms", ["blood"])]
blood_fluid_terms = sorted(set(term for term in blood_fluid_terms if
↪ term))
if not blood_fluid_terms:
    blood_fluid_terms = ["blood"]

icu_pco2 = _coerce_itemid_list(icu.get("pco2_itemids", []),
↪ key_name="icu.pco2_itemids")

```

```

icu_ph = _coerce_itemid_list(icu.get("ph_itemids", []),
↪ key_name="icu.ph_itemids")
icu_hco3 = _coerce_itemid_list(icu.get("hco3_itemids", []),
↪ key_name="icu.hco3_itemids")
icu_spec = _coerce_itemid_list(
    icu.get("specimen_type_itemids", []),
    key_name="icu.specimen_type_itemids",
)
deny_tokens = [
    str(token).strip().lower()
    for token in denylist.get("label_tokens", [])
    if str(token).strip()
]
allow_patternFallback = bool(icu.get("allow_patternFallback", True))
inclusion_mode = str(icu.get("inclusion_mode",
↪ "validated_only")).strip().lower()
if inclusion_mode not in {"validated_only", "disabled"}:
    raise ValueError(
        "Blood-gas manifest icu.inclusion_mode must be one of "
        "{'validated_only', 'disabled'}."
)
pco2_qc_payload = icu.get("pco2_qc", {})
if pco2_qc_payload is None:
    pco2_qc_payload = {}
if not isinstance(pco2_qc_payload, dict):
    raise ValueError("Blood-gas manifest icu.pco2_qc must be a JSON
↪ object.")
qc_allowed_uoms = [
    str(value).strip().lower()
    for value in pco2_qc_payload.get("allowed_uoms", ["", "mmhg", "kpa"])
]
qc_allowed_uoms = sorted(set(qc_allowed_uoms))
if not qc_allowed_uoms:
    qc_allowed_uoms = ["", "mmhg", "kpa"]
qc_label_include = [
    str(value).strip().lower()
    for value in pco2_qc_payload.get("label_must_include_tokens",
↪ ["pco2"])
    if str(value).strip()
]
qc_label_include = sorted(set(qc_label_include))

```

```

qc_label_exclude = [
    str(value).strip().lower()
    for value in pco2_qc_payload.get(
        "label_must_exclude_tokens",
        ["po2", "pao2", "oxygen", "sao2", "spo2", "fio2"],
    )
    if str(value).strip()
]

qc_label_exclude = sorted(set(qc_label_exclude))
qc_median_min = float(pco2_qc_payload.get("median_min", 45.0))
qc_median_max = float(pco2_qc_payload.get("median_max", 80.0))
if qc_median_min > qc_median_max:
    raise ValueError(
        "Blood-gas manifest icu.pco2_qc median_min must be <=
        ↳ median_max."
)

manifest = {
    "version": str(payload.get("version", "UNSPECIFIED")),
    "lab": {
        "pco2_itemids": lab_pco2,
        "excluded_pco2_itemids": lab_pco2_excluded,
        "ph_itemids": lab_ph,
        "hco3_itemids": lab_hco3,
        "specimen_type_itemids": lab_spec,
        "blood_fluid_terms": blood_fluid_terms,
    },
    "icu": {
        "pco2_itemids": icu_pco2,
        "ph_itemids": icu_ph,
        "hco3_itemids": icu_hco3,
        "specimen_type_itemids": icu_spec,
        "allow_patternFallback": allow_patternFallback,
        "inclusion_mode": inclusion_mode,
        "pco2_qc": {
            "allowed_uoms": qc_allowed_uoms,
            "median_min": qc_median_min,
            "median_max": qc_median_max,
            "label_must_include_tokens": qc_label_include,
            "label_must_exclude_tokens": qc_label_exclude,
        },
    },
    "denylist": {"label_tokens": deny_tokens},
}

```

```

        "path": str(path),
    }
    return manifest

BLOOD_GAS_MANIFEST = load_blood_gas_itemid_manifest(BLOOD_GAS_SPEC_PATH)
LAB_PCO2_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["pco2_itemids"]
LAB_EXCLUDED_PCO2_ITEMIDS =
    ↪ BLOOD_GAS_MANIFEST["lab"]["excluded_pco2_itemids"]
if LAB_EXCLUDED_PCO2_ITEMIDS:
    LAB_PCO2_ITEMIDS = sorted(
        set(LAB_PCO2_ITEMIDS).difference(set(LAB_EXCLUDED_PCO2_ITEMIDS)))
)
if not LAB_PCO2_ITEMIDS:
    raise ValueError(
        "All lab.pco2_itemids were excluded by lab.excluded_pco2_itemids. "
        "Update specs/blood_gas_itemids.json."
    )
LAB_PH_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["ph_itemids"]
LAB_HCO3_ITEMIDS = BLOOD_GAS_MANIFEST["lab"]["hco3_itemids"]
LAB_SPECIMEN_TYPE_ITEMIDS =
    ↪ BLOOD_GAS_MANIFEST["lab"]["specimen_type_itemids"]
LAB_BLOOD_FLUID_TERMS = BLOOD_GAS_MANIFEST["lab"]["blood_fluid_terms"]
ICU_PCO2_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["pco2_itemids"]
ICU_PH_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["ph_itemids"]
ICU_HCO3_ITEMIDS = BLOOD_GAS_MANIFEST["icu"]["hco3_itemids"]
ICU_SPECIMEN_TYPE_ITEMIDS =
    ↪ BLOOD_GAS_MANIFEST["icu"]["specimen_type_itemids"]
ICU_ALLOW_PATTERN_FALLBACK =
    ↪ bool(BLOOD_GAS_MANIFEST["icu"]["allow_patternFallback"])
ICU_INCLUSION_MODE =
    ↪ str(BLOOD_GAS_MANIFEST["icu"]["inclusion_mode"]).strip().lower()
ICU_PCO2_QC = BLOOD_GAS_MANIFEST["icu"]["pco2_qc"]
ICU_LABEL_DENY_TOKENS = BLOOD_GAS_MANIFEST["denylist"]["label_tokens"]
ICU_PCO2_QC_ALLOWED_UOMS = [
    str(value).strip().lower() for value in ICU_PCO2_QC["allowed_uoms"]
]
ICU_PCO2_QC_MEDIAN_MIN = float(ICU_PCO2_QC["median_min"])
ICU_PCO2_QC_MEDIAN_MAX = float(ICU_PCO2_QC["median_max"])
ICU_PCO2_QC_INCLUDE_TOKENS = [
    str(value).strip().lower()
    for value in ICU_PCO2_QC["label_must_include_tokens"]
]

```

```

ICU_PC02_QC_EXCLUDE_TOKENS = [
    str(value).strip().lower()
    for value in ICU_PC02_QC["label_must_exclude_tokens"]
]

LAB_PC02_SQL_ARRAY = _sql_int_array(LAB_PC02_ITEMIDS)
LAB_PH_SQL_ARRAY = _sql_int_array(LAB_PH_ITEMIDS)
LAB_HCO3_SQL_ARRAY = _sql_int_array(LAB_HCO3_ITEMIDS)
LAB_SPECIMEN_SQL_ARRAY = _sql_int_array(LAB_SPECIMEN_TYPE_ITEMIDS)
ICU_PC02_SQL_ARRAY = _sql_int_array(ICU_PC02_ITEMIDS)
ICU_PH_SQL_ARRAY = _sql_int_array(ICU_PH_ITEMIDS)
ICU_HCO3_SQL_ARRAY = _sql_int_array(ICU_HCO3_ITEMIDS)
ICU_SPECIMEN_SQL_ARRAY = _sql_int_array(ICU_SPECIMEN_TYPE_ITEMIDS)
LAB_BLOOD_FLUID_SQL_LIST = ", ".join(f"'{term}'" for term in
    LAB_BLOOD_FLUID_TERMS)
ICU_LABEL_DENY_REGEX = (
    "|".join(re.escape(token) for token in ICU_LABEL_DENY_TOKENS)
    if ICU_LABEL_DENY_TOKENS
    else _PO2_LABEL_DENY_PATTERN_DEFAULT.pattern
)

def _to_int64(series: pd.Series) -> pd.Series:
    """Return pandas nullable integer series for key columns."""
    return pd.to_numeric(series, errors="coerce").astype("Int64")

def prepare_omr_records(omr_df: pd.DataFrame) -> pd.DataFrame:
    """Normalize OMR records into a deterministic schema.

    Required source columns: ``subject_id``, ``chartdate``, ``result_name``,
    ``result_value``.

    Returns:
        DataFrame with columns:
        ``subject_id`` (Int64), ``chartdate_dt`` (datetime64),
        ``result_name`` (lowercase), and ``result_value_num`` (float).
    """
    required = {"subject_id", "chartdate", "result_name", "result_value"}
    missing = sorted(required.difference(omr_df.columns))
    if missing:
        raise KeyError(f"prepare_omr_records missing required columns:
            {missing}")

```

```

prepared = omr_df.copy()
prepared["subject_id"] = _to_int64(prepared["subject_id"])
prepared["chartdate_dt"] = pd.to_datetime(prepared["chartdate"],
                                         errors="coerce")
prepared["result_name"] = (
    prepared["result_name"].astype(str).str.strip().str.lower())
prepared["result_value_num"] = pd.to_numeric(
    prepared["result_value"]
    .astype(str)
    .str.extract(_NUMERIC_TOKEN_PATTERN, expand=False),
    errors="coerce",
)
prepared =
prepared.loc[prepared["result_name"].isin(OMR_RESULT_NAMES)].copy()
prepared = prepared.loc[
    prepared["subject_id"].notna() & prepared["chartdate_dt"].notna()]
].copy()

return prepared[["subject_id", "chartdate_dt", "result_name",
                 "result_value_num"]]

def attach_closest_pre_ed_omr(
    ed_df: pd.DataFrame,
    omr_df: pd.DataFrame,
    window_days: int = 365,
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Attach closest anthropometrics with tiered timing fallback.

    Args:
        ed_df: ED-stay grain dataframe with ``ed_stay_id``, ``subject_id```,
               and ``ed_intime``.
        omr_df: Prepared OMR records from ``prepare_omr_records``.
        window_days: Inclusion window in days before/after ED arrival.
    """
    def _with_default_anthro_columns(frame: pd.DataFrame) -> pd.DataFrame:
        updated = frame.copy()
        for column_name in OMR_OUTPUT_COLUMNS:
            if column_name not in updated.columns:

```

```

        updated[column_name] = pd.NA
    if "anthro_timing_tier" not in updated.columns:
        updated["anthro_timing_tier"] = "missing"
    else:
        updated["anthro_timing_tier"] =
    ↵ updated["anthro_timing_tier"].fillna("missing")
        if "anthro_days_offset" not in updated.columns:
            updated["anthro_days_offset"] = pd.Series([pd.NA] * len(updated),
    ↵ dtype="Int64")
    else:
        updated["anthro_days_offset"] = pd.to_numeric(
            updated["anthro_days_offset"], errors="coerce"
        ).astype("Int64")
    if "anthro_chartdate" not in updated.columns:
        updated["anthro_chartdate"] = pd.NaT
    else:
        updated["anthro_chartdate"] = pd.to_datetime(
            updated["anthro_chartdate"], errors="coerce"
        )
    if "anthro_timing_uncertain" not in updated.columns:
        updated["anthro_timing_uncertain"] = pd.Series(
            [pd.NA] * len(updated), dtype="boolean"
        )
    else:
        updated["anthro_timing_uncertain"] = updated[
            "anthro_timing_uncertain"
        ].astype("boolean")
    if "anthro_source" not in updated.columns:
        updated["anthro_source"] = "missing"
    else:
        updated["anthro_source"] = (
            updated["anthro_source"].astype(str).replace({"":
    ↵ "missing"}) .fillna("missing")
        )
    if "anthro_obstime" not in updated.columns:
        updated["anthro_obstime"] = pd.NaT
    else:
        updated["anthro_obstime"] =
    ↵ pd.to_datetime(updated["anthro_obstime"], errors="coerce")
        if "anthro_hours_offset" not in updated.columns:
            updated["anthro_hours_offset"] = pd.NA
        updated["anthro_hours_offset"] = pd.to_numeric(
            updated["anthro_hours_offset"], errors="coerce"

```

```

        )
    if "anthro_timing_basis" not in updated.columns:
        updated["anthro_timing_basis"] = "missing"
    else:
        updated["anthro_timing_basis"] = (
            updated["anthro_timing_basis"].astype(str).replace({"":
    ↪ "missing"}).fillna("missing")
        )

    missing_mask = updated["anthro_timing_tier"].eq("missing")
    pre_mask = updated["anthro_timing_tier"].eq("pre_ed_365")
    post_mask = updated["anthro_timing_tier"].eq("post_ed_365")
    updated.loc[pre_mask, "anthro_timing_uncertain"] = False
    updated.loc[post_mask, "anthro_timing_uncertain"] = True
    updated.loc[missing_mask, "anthro_timing_uncertain"] = pd.NA
    updated.loc[pre_mask | post_mask, "anthro_source"] = "omr"
    updated.loc[missing_mask, "anthro_source"] = "missing"
    updated.loc[pre_mask, "anthro_timing_basis"] = "pre"
    updated.loc[post_mask, "anthro_timing_basis"] = "post"
    updated.loc[missing_mask, "anthro_timing_basis"] = "missing"
    updated.loc[missing_mask, "anthro_obstime"] = pd.NaT
    updated.loc[missing_mask, "anthro_hours_offset"] = pd.NA

    return updated

def _tier_counts(frame: pd.DataFrame) -> dict[str, int]:
    tier_series =
    ↪ frame["anthro_timing_tier"].astype(str).fillna("missing")
    return {
        "pre_ed_365": int((tier_series == "pre_ed_365").sum()),
        "post_ed_365": int((tier_series == "post_ed_365").sum()),
        "missing": int((tier_series == "missing").sum()),
    }

required_ed = {"ed_stay_id", "subject_id", "ed_intime"}
missing_ed = sorted(required_ed.difference(ed_df.columns))
if missing_ed:
    raise KeyError(f"attach_closest_pre_ed.omr missing ED columns:
    ↪ {missing_ed}"))

required_omr = {"subject_id", "chartdate_dt", "result_name",
    ↪ "result_value_num"}
missing_omr = sorted(required_omr.difference(omr_df.columns))

```

```

if missing_omr:
    raise KeyError(f"attach_closest_pre_ed_omr missing OMR columns:
        {missing_omr}")

diagnostics: dict[str, Any] = {
    "window_days": int(window_days),
    "source_rows_prepared": int(len(omr_df)),
    "parsed_value_rows": int(omr_df["result_value_num"].notna().sum()),
    "ed_rows": int(len(ed_df)),
}

empty_window_diagnostics = {
    "nonnegative_candidate_rows": 0,
    "pre_window_candidate_rows": 0,
    "post_window_candidate_rows": 0,
    "closest_absolute_candidate_rows": 0,
    "within_window_candidate_rows": 0,
}

ed_norm = ed_df[["ed_stay_id", "subject_id", "ed_intime"]].copy()
ed_norm["subject_id"] = _to_int64(ed_norm["subject_id"])
ed_norm["ed_intime_dt"] = pd.to_datetime(ed_norm["ed_intime"],
    errors="coerce")
ed_norm["ed_date_dt"] = ed_norm["ed_intime_dt"].dt.floor("D")
ed_norm = ed_norm.loc[ed_norm["subject_id"].notna() &
    ed_norm["ed_date_dt"].notna()]

if omr_df.empty or ed_norm.empty:
    updated = _with_default_anthro_columns(ed_df)
    tier_counts = _tier_counts(updated)
    total_rows = max(int(len(updated)), 1)
    diagnostics.update(
        {
            "ed_rows_eligible_for_join": int(len(ed_norm)),
            "subject_overlap_count": 0,
            "candidate_rows_after_subject_join": 0,
            **empty_window_diagnostics,
            "within_window_candidate_rows": 0,
            "eligible_ed_stays_with_candidates": 0,
            "attached_non_null_counts": {
                column_name: int(updated[column_name].notna().sum())
                for column_name in OMR_OUTPUT_COLUMNS
            },
        }
    )

```

```

        "attached_any_non_null_rows": int(
    ↵     updated[list(OMR_OUTPUT_COLUMNS)].notna().any(axis=1).sum()
        ),
        "selected_tier_counts": tier_counts,
        "selected_tier_rates": {
            key: float(value / total_rows) for key, value in
    ↵     tier_counts.items()
        },
        "timing_uncertain_count": int(
            updated["anthro_timing_uncertain"].fillna(False).sum()
        ),
    ↵ )
)
return updated, diagnostics

omr_pivot = (
    omr_df.pivot_table(
        index=["subject_id", "chartdate_dt"],
        columns="result_name",
        values="result_value_num",
        aggfunc="first",
    )
    .reset_index()
    .copy()
)

shared_subjects =
    set(ed_norm["subject_id"].dropna().astype(int)).intersection(
        set(omr_pivot["subject_id"].dropna().astype(int))
    )
diagnostics["ed_rows_eligible_for_join"] = int(len(ed_norm))
diagnostics["subject_overlap_count"] = int(len(shared_subjects))

merged = ed_norm.merge(omr_pivot, on="subject_id", how="left")
diagnostics["candidate_rows_after_subject_join"] = int(
    merged["chartdate_dt"].notna().sum()
)

merged["days_before"] = (
    merged["ed_date_dt"] - pd.to_datetime(merged["chartdate_dt"],
    ↵ errors="coerce")
).dt.days

```

```

valid_days = merged["days_before"].dropna()
diagnostics["days_before_min"] = (
    int(valid_days.min()) if not valid_days.empty else None
)
diagnostics["days_before_max"] = (
    int(valid_days.max()) if not valid_days.empty else None
)
pre_window_mask = (
    merged["days_before"].notna()
    & merged["days_before"].ge(0)
    & merged["days_before"].le(window_days)
)
post_window_mask = (
    merged["days_before"].notna()
    & merged["days_before"].lt(0)
    & merged["days_before"].ge(-window_days)
)
abs_window_mask = (
    merged["days_before"].notna()
    & merged["days_before"].abs().le(window_days)
)

diagnostics["nonnegative_candidate_rows"] = int((merged["days_before"] >
↪ 0).sum())
diagnostics["pre_window_candidate_rows"] = int(pre_window_mask.sum())
diagnostics["post_window_candidate_rows"] = int(post_window_mask.sum())
diagnostics["closest_absolute_candidate_rows"] =
↪ int(abs_window_mask.sum())

pre_candidates = merged.loc[pre_window_mask].copy()
post_candidates = merged.loc[post_window_mask].copy()
diagnostics["within_window_candidate_rows"] = int(len(pre_candidates))

selected_parts: list[pd.DataFrame] = []
selected_stays: set[Any] = set()

if not pre_candidates.empty:
    pre_selected = (
        pre_candidates.sort_values(
            ["ed_stay_id", "days_before", "chartdate_dt"],
            ascending=[True, True, False],
        )
        .groupby("ed_stay_id", as_index=False)
    )

```

```

        .first()
    )
    pre_selected["anthro_timing_tier"] = "pre_ed_365"
    selected_parts.append(pre_selected)
    selected_stays.update(pre_selected["ed_stay_id"].tolist())

    if not post_candidates.empty:
        post_candidates["days_after_abs"] =
        post_candidates["days_before"].abs()
        post_selected = (
            post_candidates.sort_values(
                ["ed_stay_id", "days_after_abs", "chartdate_dt"],
                ascending=[True, True, True],
            )
            .groupby("ed_stay_id", as_index=False)
            .first()
        )
        post_selected = post_selected.loc[
            ~post_selected["ed_stay_id"].isin(selected_stays)
        ].copy()
        if not post_selected.empty:
            post_selected["anthro_timing_tier"] = "post_ed_365"
            selected_parts.append(post_selected)

    if selected_parts:
        selected = pd.concat(selected_parts, ignore_index=True)
        selected = selected.rename(
            columns={
                "bmi": "bmi_closest_pre_ed",
                "height": "height_closest_pre_ed",
                "weight": "weight_closest_pre_ed",
                "chartdate_dt": "anthro_chartdate",
                "days_before": "anthro_days_offset",
            }
        )
        selected["anthro_chartdate"] = pd.to_datetime(
            selected["anthro_chartdate"], errors="coerce"
        )
        selected["anthro_days_offset"] = pd.to_numeric(
            selected["anthro_days_offset"], errors="coerce"
        ).astype("Int64")
        selected["anthro_timing_uncertain"] =
        selected["anthro_timing_tier"].eq(

```

```

        "post_ed_365"
    )
selected["anthro_source"] = "omr"
selected["anthro_obstime"] = pd.to_datetime(
    selected["anthro_chartdate"], errors="coerce"
)
selected["anthro_hours_offset"] = pd.to_numeric(
    selected["anthro_days_offset"], errors="coerce"
) * 24.0
selected["anthro_timing_basis"] = selected["anthro_timing_tier"].map(
    {"pre_ed_365": "pre", "post_ed_365": "post"}
).fillna("missing")

updates = selected[
    [
        "ed_stay_id",
        *OMR_OUTPUT_COLUMNS,
        *OMR_PROVENANCE_COLUMNS,
    ]
].copy()
updated = ed_df.merge(updates, on="ed_stay_id", how="left")
else:
    updated = ed_df.copy()

updated = _with_default_anthro_columns(updated)

diagnostics["eligible_ed_stays_with_candidates"] = int(
    updated["anthro_timing_tier"].isin({"pre_ed_365",
    "post_ed_365"}).sum()
)
diagnostics["attached_non_null_counts"] = {
    column_name: int(updated[column_name].notna().sum())
    for column_name in OMR_OUTPUT_COLUMNS
}
diagnostics["attached_any_non_null_rows"] = int(
    updated[list(OMR_OUTPUT_COLUMNS)].notna().any(axis=1).sum()
)
tier_counts = _tier_counts(updated)
diagnostics["selected_tier_counts"] = tier_counts
total_rows = max(int(len(updated)), 1)
diagnostics["selected_tier_rates"] = {
    key: float(value / total_rows) for key, value in tier_counts.items()
}

```

```

diagnostics["timing_uncertain_count"] = int(
    updated["anthro_timing_uncertain"].fillna(False).sum()
)
diagnostics["anthro_source_counts"] = {
    str(key): int(value)
    for key, value in updated["anthro_source"].fillna("missing")
        .astype(str)
        .value_counts(dropna=False)
        .items()
}
return updated, diagnostics

def _infer_route_hint_text(value: object) -> str | None:
    text = "" if pd.isna(value) else str(value).strip().lower()
    if not text or text in {"nan", "none"}:
        return None
    if _ARTERIAL_HINT_PATTERN.search(text):
        return "arterial"
    if _VENOUS_HINT_PATTERN.search(text):
        return "venous"
    return None

def _resolve_route_hints(values: Sequence[str]) -> tuple[str | None, bool,
    int]:
    hints = [str(value).strip().lower() for value in values if
    str(value).strip()]
    arterial_n = sum(1 for value in hints if value == "arterial")
    venous_n = sum(1 for value in hints if value == "venous")
    conflict = arterial_n > 0 and venous_n > 0
    if conflict:
        return None, True, int(len(hints))
    if arterial_n > 0:
        return "arterial", False, int(len(hints))
    if venous_n > 0:
        return "venous", False, int(len(hints))
    return None, False, int(len(hints))

def _classify_specimen_type_text(value: object) -> str | None:
    text = "" if pd.isna(value) else str(value).strip().lower()

```

```

if not text:
    return None
if re.search(r"arter|(?:^|\b)art(?:\b|$)", text):
    return "arterial"
if re.search(r"ven|mixed|central", text):
    return "venous"
return None

def assign_panel_source_from_specimen(
    panel_df: pd.DataFrame,
    labs_df: pd.DataFrame,
    *,
    specimen_source_itemids: Sequence[int] | None = None,
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Assign gas source using specimen-type rows only (deterministic core
    ↵ policy)."""
    required_panel = {"ed_stay_id", "specimen_id"}
    missing_panel = sorted(required_panel.difference(panel_df.columns))
    if missing_panel:
        raise KeyError(
            "assign_panel_source_from_specimen missing panel columns: "
            f"{missing_panel}"
        )
    required_labs = {"ed_stay_id", "specimen_id", "itemid"}
    missing_labs = sorted(required_labs.difference(labs_df.columns))
    if missing_labs:
        raise KeyError(
            "assign_panel_source_from_specimen missing labs columns: "
            f"{missing_labs}"
        )

    key_cols = ["ed_stay_id", "specimen_id"]
    panel = panel_df.copy()
    panel["source"] = "other"
    panel["source_inference_tier"] = "fallback_other"
    panel["source_hint_conflict"] = False
    panel["source_hint_count"] = pd.Series(0, index=panel.index,
    ↵ dtype="Int64")
    panel["source_hint_specimen_text"] = pd.Series(pd.NA, index=panel.index,
    ↵ dtype="string")

    specimen_ids = {

```

```

        int(value) for value in (specimen_source_itemids or []) if
        ↵ pd.notna(value)
    }
    if not specimen_ids or "value_text" not in labs_df.columns:
        diagnostics = summarize_gas_source(panel)
        diagnostics.update(
            {
                "mode": "specimen_only",
                "resolved_rows": int(panel["source"].isin({"arterial",
                ↵ "venous"}).sum()),
                "resolved_rate": float(panel["source"].isin({"arterial",
                ↵ "venous"}).mean())
                if len(panel)
                else 0.0,
                "unresolved_specimen_id_count": int(
                    pd.to_numeric(panel["specimen_id"]),
        ↵ errors="coerce").nunique(dropna=True)
                ),
                "unresolved_specimen_id_examples": [] ,
            }
        )
    return panel, diagnostics

labs = labs_df.copy()
labs["itemid"] = pd.to_numeric(labs["itemid"],
    ↵ errors="coerce").astype("Int64")
labs["specimen_id"] = pd.to_numeric(labs["specimen_id"],
    ↵ errors="coerce").astype("Int64")
labs = labs.loc[labs["itemid"].isin(specimen_ids) &
    ↵ labs["specimen_id"].notna()].copy()
if labs.empty:
    diagnostics = summarize_gas_source(panel)
    diagnostics.update(
        {
            "mode": "specimen_only",
            "resolved_rows": int(panel["source"].isin({"arterial",
            ↵ "venous"}).sum()),
            "resolved_rate": float(panel["source"].isin({"arterial",
            ↵ "venous"}).mean())
            if len(panel)
            else 0.0,
            "unresolved_specimen_id_count": int(

```

```

        pd.to_numeric(panel["specimen_id"] ,
←  errors="coerce").nunique(dropna=True)
        ),
        "unresolved_specimen_id_examples": [] ,
    }
)
return panel, diagnostics

labs["specimen_text_raw"] = (
    labs["value_text"].astype("string").fillna("").str.strip()
)
labs["specimen_text_norm"] = labs["specimen_text_raw"].str.lower()
labs["specimen_route_hint"] =
← labs["specimen_text_norm"].map(_classify_specimen_type_text)

grouped = (
    labs.groupby(key_cols, dropna=False)
    .agg(
        route_hints= ("specimen_route_hint", list),
        specimen_texts= ("specimen_text_raw", list),
    )
    .reset_index()
)

mapping_rows: list[dict[str, Any]] = []
for _, grouped_row in grouped.iterrows():
    source, conflict, hint_count =
← _resolve_route_hints(grouped_row["route_hints"])
    specimen_texts = [
        str(value).strip()
        for value in grouped_row["specimen_texts"]
        if str(value).strip()
    ]
    specimen_text = specimen_texts[0] if specimen_texts else None
    mapping_rows.append(
    {
        "ed_stay_id": grouped_row["ed_stay_id"],
        "specimen_id": grouped_row["specimen_id"],
        "source": source if source is not None else "other",
        "source_inference_tier": "specimen_text" if source is not
        ← None else "fallback_other",
        "source_hint_conflict": bool(conflict),
        "source_hint_count": int(hint_count),
    }
)

```

```

        "source_hint_specimen_text": specimen_text,
    }
)

mapping = pd.DataFrame(mapping_rows)
if not mapping.empty:
    panel = panel.merge(mapping, on=key_cols, how="left", suffixes=("",
"specimen"))
    panel["source"] = (
        panel.get("source_specimen")
        .fillna(panel["source"])
        .astype("string")
        .fillna("other")
    )
    panel["source_inference_tier"] = (
        panel.get("source_inference_tier_specimen")
        .fillna(panel["source_inference_tier"])
        .astype("string")
        .fillna("fallback_other")
    )
    panel["source_hint_conflict"] = (
        panel.get("source_hint_conflict_specimen")
        .fillna(panel["source_hint_conflict"])
        .astype(bool)
    )
    panel["source_hint_count"] = pd.to_numeric(
        panel.get("source_hint_count_specimen").fillna(panel[
"source_hint_count"]),
        errors="coerce",
    ).fillna(0).astype("Int64")
    panel["source_hint_specimen_text"] = (
        panel.get("source_hint_specimen_text_specimen")
        .fillna(panel["source_hint_specimen_text"])
        .astype("string")
    )
    panel = panel.drop(
        columns=[
            column_name
            for column_name in (
                "source_specimen",
                "source_inference_tier_specimen",
                "source_hint_conflict_specimen",
                "source_hint_count_specimen",
            )
        ]
    )

```

```

        "source_hint_specimen_text_specimen",
    )
    if column_name in panel.columns
]
)

diagnostics = summarize_gas_source(panel)
unresolved = panel.loc[panel["source"].eq("other"),
↪ key_cols].drop_duplicates()
diagnostics.update(
{
    "mode": "specimen_only",
    "resolved_rows": int(panel["source"].isin({"arterial",
↪ "venous"}).sum()),
    "resolved_rate": float(panel["source"].isin({"arterial",
↪ "venous"}).mean())
    if len(panel)
    else 0.0,
    "unresolved_specimen_id_count": int(
        pd.to_numeric(unresolved["specimen_id"],
↪ errors="coerce").nunique(dropna=True)
    ),
    "unresolved_specimen_id_examples": unresolved["specimen_id"]
        .dropna()
        .astype(int)
        .head(20)
        .tolist(),
}
)
return panel, diagnostics

def infer_panel_gas_source_metadata(
    panel_df: pd.DataFrame,
    labs_df: pd.DataFrame,
    labitems_df: pd.DataFrame,
    *,
    specimen_source_itemids: Sequence[int] | None = None,
    pco2_itemids: Sequence[int] | None = None,
    mode: str = "metadata_only",
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Infer gas panel source from metadata/text hints with tier diagnostics.

```

```

Tier precedence:
    1. Specimen/source row text (``value_text`` from specimen/source
   ↵ itemids),
    2. ``d_labitems`` label/fluid hints,
    3. Panel co-occurrence hints (pCO2-labeled rows + free-text),
    4. Fallback ``other``.

"""
if mode != "metadata_only":
    raise ValueError(
        "infer_panel_gas_source_metadata supports only
         ↵ mode='metadata_only'."
    )

required_panel = {"ed_stay_id", "specimen_id"}
missing_panel = sorted(required_panel.difference(panel_df.columns))
if missing_panel:
    raise KeyError(
        "infer_panel_gas_source_metadata missing panel columns: "
        f"{missing_panel}"
    )
required_labs = {"ed_stay_id", "specimen_id", "itemid"}
missing_labs = sorted(required_labs.difference(labs_df.columns))
if missing_labs:
    raise KeyError(
        "infer_panel_gas_source_metadata missing labs columns: "
        f"{missing_labs}"
    )
required_labitems = {"itemid"}
missing_labitems =
   ↵ sorted(required_labitems.difference(labitems_df.columns))
if missing_labitems:
    raise KeyError(
        "infer_panel_gas_source_metadata missing labitems columns: "
        f"{missing_labitems}"
    )

key_cols = ["ed_stay_id", "specimen_id"]
panel = panel_df.copy()
panel["source"] = pd.Series(pd.NA, index=panel.index, dtype="string")
panel["source_inference_tier"] = "fallback_other"
panel["source_hint_conflict"] = pd.Series(False, index=panel.index,
   ↵ dtype="boolean")

```

```

panel["source_hint_count"] = pd.Series(pd.NA, index=panel.index,
                                         dtype="Int64")

    ↵ labs = labs_df.copy()
    ↵ labs["itemid"] = pd.to_numeric(labs["itemid"],
                                     errors="coerce").astype("Int64")
    ↵ labs["specimen_id"] = pd.to_numeric(labs["specimen_id"],
                                         errors="coerce").astype(
                                             "Int64"
                                         )
    ↵ labs = labs.loc[labs["itemid"].notna() &
    ↵ labs["specimen_id"].notna()].copy()
    ↵ if labs.empty:
        panel["source"] = "other"
        diagnostics = summarize_gas_source(panel)
        diagnostics.update(
            {
                "mode": mode,
                "resolved_rows": 0,
                "resolved_rate": 0.0,
                "unresolved_specimen_id_count": 0,
                "unresolved_specimen_id_examples": [],
                "unresolved_value_text_top": {},
                "unresolved_label_top": {}
            }
        )
    ↵ return panel, diagnostics

if "value_text" in labs.columns:
    labs["value_text_norm"] = (
        labs["value_text"]
        .astype("string")
        .fillna("")
        .str.strip()
        .str.lower()
    )
else:
    labs["value_text_norm"] = ""

labitems = labitems_df.copy()
labitems["itemid"] = pd.to_numeric(labitems["itemid"],
                                     errors="coerce").astype(
                                         "Int64"

```

```

)
for text_col in ("label", "fluid"):
    if text_col not in labitems.columns:
        labitems[text_col] = ""
labitems["label_norm"] =
    labitems["label"].astype("string").fillna("").str.strip()
labitems["fluid_norm"] =
    labitems["fluid"].astype("string").fillna("").str.strip()
labitems["item_route_hint"] = (
    (labitems["label_norm"] + " " + labitems["fluid_norm"])
    .str.strip()
    .map(_infer_route_hint_text)
)
item_meta = (
    labitems.dropna(subset=["itemid"])
    .drop_duplicates(subset=["itemid"])
    .set_index("itemid")[["label_norm", "fluid_norm", "item_route_hint"]]
)
labs = labs.merge(item_meta, left_on="itemid", right_index=True,
    how="left")
labs["item_route_hint"] = labs["item_route_hint"].astype("string")
labs["text_route_hint"] =
    labs["value_text_norm"].map(_infer_route_hint_text).astype(
        "string"
)
labs["label_fluid_text"] = (
    labs["label_norm"].fillna("").astype(str).str.strip()
    + " | "
    + labs["fluid_norm"].fillna("").astype(str).str.strip()
).str.strip(" |")

def _build_tier_mapping(
    frame: pd.DataFrame,
    *,
    hint_column: str,
    tier_name: str,
) -> pd.DataFrame:
    subset = frame[key_cols + [hint_column]].copy()
    subset = subset.loc[subset[hint_column].notna()].copy()
    if subset.empty:
        return pd.DataFrame(
            columns=[
                *key_cols,

```

```

        "tier_source",
        "tier_conflict",
        "tier_hint_count",
        "tier_name",
    ]
)
grouped = (
    subset.groupby(key_cols, dropna=False)[hint_column]
    .agg(list)
    .reset_index(name="hints")
)
rows: list[dict[str, Any]] = []
for _, grouped_row in grouped.iterrows():
    source, conflict, hint_count =
    ↪ _resolve_route_hints(grouped_row["hints"])
    rows.append(
        {
            "ed_stay_id": grouped_row["ed_stay_id"],
            "specimen_id": grouped_row["specimen_id"],
            "tier_source": source,
            "tier_conflict": bool(conflict),
            "tier_hint_count": int(hint_count),
            "tier_name": tier_name,
        }
    )
resolved = pd.DataFrame(rows)
return resolved.loc[resolved["tier_source"] |
    ↪ ].notna()].reset_index(drop=True)

specimen_source_itemid_set = {
    int(item) for item in (specimen_source_itemids or [])
    ↪ pd.notna(item)
}
pco2_itemid_set = {int(item) for item in (pco2_itemids or [])
    ↪ pd.notna(item)}

tier1 = _build_tier_mapping(
    labs.loc[labs["itemid"].astype("Int64") |
    ↪ ).isin(specimen_source_itemid_set)],
    hint_column="text_route_hint",
    tier_name="specimen_text",
)

```

```

tier2 = _build_tier_mapping(
    labs,
    hint_column="item_route_hint",
    tier_name="label_fluid",
)
tier_cluster = pd.DataFrame(
    columns=[
        *key_cols,
        "tier_source",
        "tier_conflict",
        "tier_hint_count",
        "tier_name",
    ]
)
if "panel_time" in panel_df.columns:
    panel_times = panel_df[key_cols + ["panel_time"]].copy()
    panel_times["panel_time"] = pd.to_datetime(panel_times["panel_time"],
    errors="coerce")

    tier_seed = pd.concat([tier1, tier2], ignore_index=True)
    if not tier_seed.empty:
        tier_seed = (
            tier_seed[key_cols + ["tier_source"]]
            .dropna(subset=["tier_source"])
            .drop_duplicates(subset=key_cols)
        )
        panel_seed = panel_times.merge(tier_seed, on=key_cols,
    how="left")
        known = panel_seed.loc[panel_seed["tier_source"].notna()].copy()
        unknown = panel_seed.loc[panel_seed["tier_source"].isna()].copy()
        if not known.empty and not unknown.empty:
            pairs = unknown.merge(
                known[["ed_stay_id", "specimen_id", "panel_time",
    "tier_source"]],
                on="ed_stay_id",
                how="left",
                suffixes=("_unknown", "_known"),
            )
            pairs["delta_minutes"] = (
                (pairs["panel_time_known"] -
    pairs["panel_time_unknown"]).abs().dt.total_seconds()
                / 60.0
            )

```

```

        pairs = pairs.loc[pairs["delta_minutes"].notna() &
↪  pairs["delta_minutes"].le(60.0)].copy()
        if not pairs.empty:
            keys = ["ed_stay_id", "specimen_id_unknown"]
            min_delta =
↪  pairs.groupby(keys)["delta_minutes"].transform("min")
            nearest =
↪  pairs.loc[pairs["delta_minutes"].eq(min_delta)].copy()
            conflict_stats = (
                nearest.groupby(keys)[["tier_source_known"]]
                .nunique(dropna=True)
                .reset_index(name="source_nunique")
            )
            nearest = nearest.merge(conflict_stats, on=keys,
↪  how="left")
            nearest =
↪  nearest.loc[nearest["source_nunique"].eq(1)].copy()
            if not nearest.empty:
                nearest = nearest.sort_values(keys +
↪  ["delta_minutes"])
                nearest = nearest.groupby(keys,
↪  as_index=False).first()
                tier_cluster = nearest.rename(
                    columns={
                        "specimen_id_unknown": "specimen_id",
                        "tier_source_known": "tier_source",
                    }
                )[["ed_stay_id", "specimen_id", "tier_source"]]
                .copy()
                tier_cluster["tier_conflict"] = False
                tier_cluster["tier_hint_count"] = 1
                tier_cluster["tier_name"] = "cluster_inheritance"

tier3_candidates = labs.copy()
if pco2_itemid_set:
    tier3_candidates = tier3_candidates.loc[
        tier3_candidates["itemid"].astype("Int64").isin(pco2_itemid_set)
        | tier3_candidates["text_route_hint"].notna()
    ].copy()
    tier3_candidates["panel_route_hint"] =
↪  tier3_candidates["item_route_hint"].fillna(
        tier3_candidates["text_route_hint"]

```

```

)
tier3 = _build_tier_mapping(
    tier3_candidates,
    hint_column="panel_route_hint",
    tier_name="panel_cooccurrence",
)
for tier_frame in (tier1, tier2, tier_cluster, tier3):
    if tier_frame.empty:
        continue
    panel = panel.merge(
        tier_frame[
            key_cols
            + ["tier_source", "tier_conflict", "tier_hint_count",
               "tier_name"]
        ],
        on=key_cols,
        how="left",
    )
    assign_mask = panel["source"].isna() & panel["tier_source"].notna()
    panel.loc[assign_mask, "source"] = panel.loc[assign_mask,
        "tier_source"]
    panel.loc[assign_mask, "source_inference_tier"] = panel.loc[
        assign_mask, "tier_name"
    ]
    selected_conflict = panel.loc[assign_mask,
        "tier_conflict"].astype("boolean")
    panel.loc[assign_mask, "source_hint_conflict"] =
    selected_conflict.fillna(False)
    panel.loc[assign_mask, "source_hint_count"] = pd.to_numeric(
        panel.loc[assign_mask, "tier_hint_count"], errors="coerce"
    ).astype("Int64")
    panel = panel.drop(
        columns=["tier_source", "tier_conflict", "tier_hint_count",
        "tier_name"]
    )
    panel["source"] = panel["source"].fillna("other").astype("string")
    panel["source"] = panel["source"].replace({"unknown": "other"})
    panel.loc[panel["source"].eq("other"), "source_inference_tier"] =
    panel.loc[
        panel["source"].eq("other"), "source_inference_tier"
    ].replace({"": "fallback_other"}).fillna("fallback_other")

```

```

panel["source_inference_tier"] =
↪ panel["source_inference_tier"].astype("string")
panel["source_hint_conflict"] = (
    panel["source_hint_conflict"].fillna(False).astype("boolean")
)

diagnostics = summarize_gas_source(panel)
total = max(int(len(panel)), 1)
diagnostics["mode"] = mode
diagnostics["resolved_rows"] = int(
    panel["source"].isin({"arterial", "venous"}).sum()
)
diagnostics["resolved_rate"] = float(diagnostics["resolved_rows"] /
total)

unresolved = panel.loc[panel["source"].eq("other"),
↪ key_cols].drop_duplicates()
diagnostics["unresolved_specimen_id_count"] =
↪ int(unresolved["specimen_id"].nunique())
diagnostics["unresolved_specimen_id_examples"] = (
    unresolved["specimen_id"].dropna().astype(int).head(20).tolist()
)
if unresolved.empty:
    diagnostics["unresolved_value_text_top"] = {}
    diagnostics["unresolved_label_top"] = {}
return panel, diagnostics

unresolved_labs = labs.merge(unresolved, on=key_cols, how="inner")
unresolved_value_text = (
    unresolved_labs["value_text_norm"]
    .replace({"": pd.NA})
    .dropna()
    .astype(str)
    .value_counts()
    .head(15)
)
unresolved_label_text = (
    unresolved_labs["label_fluid_text"]
    .replace({"": pd.NA})
    .dropna()
    .astype(str)
    .value_counts()
    .head(15)
)

```

```

)
diagnostics["unresolved_value_text_top"] = {
    str(key): int(value) for key, value in unresolved_value_text.items()
}
diagnostics["unresolved_label_top"] = {
    str(key): int(value) for key, value in unresolved_label_text.items()
}
return panel, diagnostics

def summarize_gas_source(panel_df: pd.DataFrame) -> dict[str, Any]:
    """Summarize gas source composition from panel-level records."""
    total_rows = int(len(panel_df))
    if total_rows == 0:
        return {
            "panel_rows": 0,
            "source_present": bool("source" in panel_df.columns),
            "source_counts": {},
            "source_rates": {},
            "all_other_or_unknown": False,
            "tier_counts": {},
            "tier_rates": {},
        }
    if "source" not in panel_df.columns:
        return {
            "panel_rows": total_rows,
            "source_present": False,
            "source_counts": {},
            "source_rates": {},
            "all_other_or_unknown": True,
            "tier_counts": {},
            "tier_rates": {},
        }
    source = (
        panel_df["source"]
        .astype(str)
        .str.strip()
        .str.lower()
        .replace({"": "other", "nan": "other", "none": "other"})
        .fillna("other")
    )

```

```

source_counts = {str(key): int(value) for key, value in
← source.value_counts(dropna=False).items()}
source_rates = {key: float(value / total_rows) for key, value in
← source_counts.items()}
all_other_or_unknown = bool(
    total_rows > 0 and set(source_counts.keys()).issubset({"other",
← "unknown"})
)
tier_counts: dict[str, int] = {}
tier_rates: dict[str, float] = {}
if "source_inference_tier" in panel_df.columns:
    tier = (
        panel_df["source_inference_tier"]
        .astype(str)
        .str.strip()
        .str.lower()
        .replace({"": "unknown", "nan": "unknown", "none": "unknown"})
        .fillna("unknown")
    )
    tier_counts = {
        str(key): int(value) for key, value in
        ← tier.value_counts(dropna=False).items()
    }
    tier_rates = {key: float(value / total_rows) for key, value in
    ← tier_counts.items()}

return {
    "panel_rows": total_rows,
    "source_present": True,
    "source_counts": source_counts,
    "source_rates": source_rates,
    "all_other_or_unknown": all_other_or_unknown,
    "tier_counts": tier_counts,
    "tier_rates": tier_rates,
}
}

def assert_gas_source_coverage(
    gas_source_audit: Mapping[str, Any],
    *,
    fail_on_all_other_source: bool = True,
) -> None:
    """Fail fast when source attribution collapses to all other/unknown."""

```

```

if not fail_on_all_other_source:
    return
if int(gas_source_audit.get("panel_rows", 0)) <= 0:
    return
if bool(gas_source_audit.get("all_other_or_unknown", False)):
    raise ValueError(
        "Gas source attribution classified all panel rows as
        ↵ other/unknown."
        "Set COHORT_FAIL_ON_ALL_OTHER_SOURCE=0 to bypass this guard."
    )

def build_gas_source_overlap_summary(ed_df: pd.DataFrame) -> pd.DataFrame:
    """Build ABG/VBG/OTHER overlap counts and percentages."""
    frame = ed_df.copy()
    abg = pd.to_numeric(frame.get("flag_abg_hypcapnia", 0),
    ↵ errors="coerce").fillna(0).astype(int)
    vbg = pd.to_numeric(frame.get("flag_vbg_hypcapnia", 0),
    ↵ errors="coerce").fillna(0).astype(int)
    other = pd.to_numeric(frame.get("flag_other_hypcapnia", 0),
    ↵ errors="coerce").fillna(0).astype(int)

    labels = []
    for a, v, o in zip(abg.tolist(), vbg.tolist(), other.tolist(),
    ↵ strict=False):
        parts: list[str] = []
        if a == 1:
            parts.append("ABG")
        if v == 1:
            parts.append("V ро")
        if o == 1:
            parts.append("OTHER")
        labels.append("+".join(parts) if parts else "NO_GAS")

    counts = pd.Series(labels, dtype="string").value_counts(dropna=False,
    ↵ ).rename_axis("gas_overlap").reset_index(name="count")
    total = max(int(counts["count"].sum()), 1)
    counts["percent"] = counts["count"].astype(float) / total * 100.0
    counts["percent"] = counts["percent"].round(2)
    return counts.sort_values(["count", "gas_overlap"], ascending=[False,
    ↵ True]).reset_index(drop=True)

```

```

def add_gas_model_fields(
    ed_df: pd.DataFrame,
    *,
    ranges: Mapping[str, tuple[float, float]] | None = None,
) -> tuple[pd.DataFrame, pd.DataFrame]:
    """Create cleaned gas model fields and a field-level outlier audit."""
    resolved_ranges = dict(ranges or DEFAULT_GAS_MODEL_RANGES)
    updated = ed_df.copy()
    audit_rows: list[dict[str, Any]] = []

    for raw_column, (lower_bound, upper_bound) in resolved_ranges.items():
        if raw_column not in updated.columns:
            continue
        numeric = pd.to_numeric(updated[raw_column], errors="coerce")
        out_of_range = numeric.notna() & ((numeric < lower_bound) | (numeric
        ↵ > upper_bound))
        model_column = f"{raw_column}_model"
        outlier_flag_column = f"{raw_column}_outlier_flag"
        updated[model_column] = numeric.where(~out_of_range)
        updated[outlier_flag_column] = out_of_range.astype("boolean")
        nonnull_n = int(numeric.notna().sum())
        outlier_n = int(out_of_range.sum())
        examples = (
            numeric.loc[out_of_range]
            .round(4)
            .value_counts()
            .head(5)
            .index.astype(str)
            .tolist()
        )
        audit_rows.append(
            {
                "domain": "gas",
                "raw_column": raw_column,
                "model_column": model_column,
                "outlier_flag_column": outlier_flag_column,
                "lower_bound": float(lower_bound),
                "upper_bound": float(upper_bound),
                "nonnull_n": nonnull_n,
                "out_of_range_n": outlier_n,
                "out_of_range_pct": float(outlier_n / nonnull_n) if nonnull_n
                ↵ else 0.0,
                "example_outlier_values": "; ".join(examples),
            }
        )
    return updated, pd.DataFrame(audit_rows)

```

```

        }
    )

audit =
↳ pd.DataFrame(audit_rows).sort_values("raw_column").reset_index(drop=True)
return updated, audit

def attach_chARTed_anthro_fallback(
    ed_df: pd.DataFrame,
    charted_df: pd.DataFrame,
    *,
    nearest_anytime: bool = True,
) -> tuple[pd.DataFrame, dict[str, Any]]:
    """Fill missing anthropometrics from charted records with nearest-time
    selection."""
    required_ed = {"ed_stay_id", "subject_id", "ed_intime"}
    missing_ed = sorted(required_ed.difference(ed_df.columns))
    if missing_ed:
        raise KeyError(f"attach_chARTed_anthro_fallback missing ED columns:
            {missing_ed}")

    required_chARTed = {"subject_id", "obs_time", "result_name",
    ↳ "result_value_num"}
    missing_chARTed = sorted(required_chARTed.difference(chARTed_df.columns))
    if missing_chARTed:
        raise KeyError(
            f"attach_chARTed_anthro_fallback missing charted columns: "
            f"{missing_chARTed}"
        )

    updated = ed_df.copy()
    for output_column in OMR_OUTPUT_COLUMNS:
        if output_column not in updated.columns:
            updated[output_column] = pd.NA
    if "anthro_source" not in updated.columns:
        updated["anthro_source"] = "missing"
    else:
        updated["anthro_source"] = (
            updated["anthro_source"].astype(str).replace({"":
    ↳ "missing"}).fillna("missing")
        )
    if "anthro_obstime" not in updated.columns:

```

```

        updated["anthro_obstime"] = pd.NaT
    else:
        updated["anthro_obstime"] = pd.to_datetime(updated["anthro_obstime"],
        ↵ errors="coerce")
    if "anthro_hours_offset" not in updated.columns:
        updated["anthro_hours_offset"] = pd.NA
    updated["anthro_hours_offset"] =
    ↵ pd.to_numeric(updated["anthro_hours_offset"], errors="coerce")
    if "anthro_timing_basis" not in updated.columns:
        updated["anthro_timing_basis"] = "missing"
    if "anthro_timing_uncertain" not in updated.columns:
        updated["anthro_timing_uncertain"] = pd.Series([pd.NA] *
    ↵ len(updated), dtype="boolean")
    else:
        updated["anthro_timing_uncertain"] =
    ↵ updated["anthro_timing_uncertain"].astype("boolean")

charted = charted_df.copy()
charted["subject_id"] = _to_int64(charted["subject_id"])
charted["obs_time"] = pd.to_datetime(charted["obs_time"],
    ↵ errors="coerce")
charted["result_name"] =
    ↵ charted["result_name"].astype(str).str.strip().str.lower()
charted["result_value_num"] = pd.to_numeric(charted["result_value_num"],
    ↵ errors="coerce")
if "source" not in charted.columns:
    charted["source"] = "icu_charted"
charted["source"] = charted["source"].astype(str).replace({"":
    ↵ "icu_charted"}).fillna("icu_charted")
charted = charted.loc[
    charted["subject_id"].notna()
    & charted["obs_time"].notna()
    & charted["result_name"].isin(OMR_RESULT_NAMES)
    & charted["result_value_num"].notna()]
].copy()

diagnostics: dict[str, Any] = {
    "charted_rows_input": int(len(charted_df)),
    "charted_rows_usable": int(len(charted)),
    "nearest_anytime": bool(nearest_anytime),
    "filled_counts": {column: 0 for column in OMR_OUTPUT_COLUMNS},
    "rows_with_any_charted_fill": 0,
}

```

```

if charted.empty:
    return updated, diagnostics

ed_norm = updated[["ed_stay_id", "subject_id", "ed_intime"]].copy()
ed_norm["subject_id"] = _to_int64(ed_norm["subject_id"])
ed_norm["ed_intime_dt"] = pd.to_datetime(ed_norm["ed_intime"],
                                         errors="coerce")
ed_norm = ed_norm.loc[ed_norm["subject_id"].notna() &
                     ed_norm["ed_intime_dt"].notna()].copy()
if ed_norm.empty:
    return updated, diagnostics

merged = ed_norm.merge(charted, on="subject_id", how="inner")
if merged.empty:
    diagnostics["subject_overlap_count"] = 0
    return updated, diagnostics
diagnostics["subject_overlap_count"] =
    int(merged["subject_id"].nunique())
merged["hours_offset"] = (
    (merged["obs_time"] - merged["ed_intime_dt"]).dt.total_seconds() /
    3600.0
)
merged["abs_hours_offset"] = merged["hours_offset"].abs()
if not nearest_anytime:
    merged = merged.loc[merged["abs_hours_offset"] <= 24.0].copy()

output_to_name = {
    "bmi_closest_pre_ed": "bmi",
    "height_closest_pre_ed": "height",
    "weight_closest_pre_ed": "weight",
}
any_fill_mask = pd.Series(False, index=updated.index)
for output_column, result_name in output_to_name.items():
    subset = merged.loc[merged["result_name"] == result_name].copy()
    if subset.empty:
        continue
    selected = (
        subset.sort_values(
            ["ed_stay_id", "abs_hours_offset", "obs_time"],
            ascending=[True, True, True],
        )
        .groupby("ed_stay_id", as_index=False)
        .first()
    )
    any_fill_mask |= subset["abs_hours_offset"] <= 24.0

```

```

        )
selected = selected.rename(
    columns={
        "result_value_num": f"{output_column}__candidate",
        "obs_time": f"{output_column}__obs_time",
        "hours_offset": f"{output_column}__hours_offset",
        "source": f"{output_column}__source",
    }
)
keep_columns = [
    "ed_stay_id",
    f"{output_column}__candidate",
    f"{output_column}__obs_time",
    f"{output_column}__hours_offset",
    f"{output_column}__source",
]
updated = updated.merge(selected[keep_columns], on="ed_stay_id",
↪ how="left")

fill_mask = updated[output_column].isna() &
↪ updated[f"{output_column}__candidate"].notna()
    updated.loc[fill_mask, output_column] = updated.loc[fill_mask,
↪ f"{output_column}__candidate"]

provenance_mask = fill_mask &
↪ updated["anthro_source"].isin({"missing", "nan"})
    updated.loc[provenance_mask, "anthro_source"] = updated.loc[
        provenance_mask, f"{output_column}__source"
    ]
    updated.loc[provenance_mask, "anthro_obs_time"] = pd.to_datetime(
        updated.loc[provenance_mask, f"{output_column}__obs_time"],
↪ errors="coerce"
    )
    updated.loc[provenance_mask, "anthro_hours_offset"] = pd.to_numeric(
        updated.loc[provenance_mask, f"{output_column}__hours_offset"],
↪ errors="coerce"
    )
    updated.loc[provenance_mask, "anthro_timing_basis"] =
↪ "nearest_anytime"
    updated.loc[provenance_mask, "anthro_timing_uncertain"] = True

any_fill_mask = any_fill_mask | fill_mask
diagnostics["filled_counts"][output_column] = int(fill_mask.sum())

```

```

drop_columns = [
    f"{{output_column}}_candidate",
    f"{{output_column}}_obs_time",
    f"{{output_column}}_hours_offset",
    f"{{output_column}}_source",
]
updated = updated.drop(columns=drop_columns)

diagnostics["rows_with_any_chARTed_fill"] = int(any_fill_mask.sum())
diagnostics["fallback_source_counts"] = {
    str(key): int(value)
    for key, value in updated.loc[any_fill_mask, "anthro_source"]
        .fillna("missing")
        .astype(str)
        .value_counts(dropna=False)
        .items()
}
return updated, diagnostics

def build_anthro_coverage_audit(ed_df: pd.DataFrame) -> dict[str, Any]:
    """Summarize anthropometric coverage and provenance rates."""
    total_rows = max(int(len(ed_df)), 1)
    field_counts = {
        column: int(pd.to_numeric(ed_df[column],
        ↵ errors="coerce").notna().sum())
        for column in OMR_OUTPUT_COLUMNS
        if column in ed_df.columns
    }
    field_rates = {column: float(count / total_rows) for column, count in
    ↵ field_counts.items()}

    source_counts: dict[str, int] = {}
    source_rates: dict[str, float] = {}
    if "anthro_source" in ed_df.columns:
        source_counts = {
            str(key): int(value)
            for key, value in ed_df["anthro_source"]
                .fillna("missing")
                .astype(str)
                .value_counts(dropna=False)
                .items()
        }

```

```

        }
    source_rates = {key: float(value / total_rows) for key, value in
↪  source_counts.items()}

    timing_basis_counts: dict[str, int] = {}
    timing_basis_rates: dict[str, float] = {}
    if "anthro_timing_basis" in ed_df.columns:
        timing_basis_counts = {
            str(key): int(value)
            for key, value in ed_df["anthro_timing_basis"]
            .fillna("missing")
            .astype(str)
            .value_counts(dropna=False)
            .items()
        }
        timing_basis_rates = {
            key: float(value / total_rows) for key, value in
↪  timing_basis_counts.items()
        }

    return {
        "row_count": int(len(ed_df)),
        "field_nonnull_counts": field_counts,
        "field_nonnull_rates": field_rates,
        "source_counts": source_counts,
        "source_rates": source_rates,
        "timing_basis_counts": timing_basis_counts,
        "timing_basis_rates": timing_basis_rates,
    }

def build_first_other_pco2_audit(ed_df: pd.DataFrame) -> pd.DataFrame:
    """Build route-stratified audit summary for first_other_pco2 values.

    This helper is intentionally tolerant for notebook QA execution: if the
    required fields are not present, it returns a sentinel audit row instead
    of
    raising.
    """
    columns = [
        "source",
        "count_nonnull",
        "mean",

```

```

        "median",
        "q25",
        "q75",
        "p95",
        "max",
        "pct_ge_80",
        "pct_ge_100",
        "pct_ge_150",
        "pct_eq_160",
        "top_values",
        "status",
        "missing_columns",
    ]
required = {"first_other_pco2", "first_other_src"}
missing = sorted(required.difference(ed_df.columns))
if missing:
    return pd.DataFrame(
        [
            {
                "source": "UNAVAILABLE",
                "count_nonnull": 0,
                "mean": None,
                "median": None,
                "q25": None,
                "q75": None,
                "p95": None,
                "max": None,
                "pct_ge_80": 0.0,
                "pct_ge_100": 0.0,
                "pct_ge_150": 0.0,
                "pct_eq_160": 0.0,
                "top_values": {},
                "status": "missing_columns",
                "missing_columns": ",".join(missing),
            }
        ],
        columns=columns,
    )

frame = ed_df[["first_other_src", "first_other_pco2"]].copy()
frame["first_other_src"] = (

```

```

    ↵ frame["first_other_src"].astype(str).str.strip().str.upper().replace({"": "UNKNOWN", "NAN": "UNKNOWN"})
    ↵ )
    ↵ frame["first_other_pco2"] = pd.to_numeric(frame["first_other_pco2"],
    ↵ errors="coerce")
    ↵ frame = frame.loc[frame["first_other_pco2"].notna()].copy()

if frame.empty:
    return pd.DataFrame(
        [
            {
                "source": "UNAVAILABLE",
                "count_nonnull": 0,
                "mean": None,
                "median": None,
                "q25": None,
                "q75": None,
                "p95": None,
                "max": None,
                "pct_ge_80": 0.0,
                "pct_ge_100": 0.0,
                "pct_ge_150": 0.0,
                "pct_eq_160": 0.0,
                "top_values": {},
                "status": "no_nonnull_values",
                "missing_columns": ""
            }
        ],
        columns=columns,
    )

rows: list[dict[str, Any]] = []
for source_name, group in frame.groupby("first_other_src"):
    values = group["first_other_pco2"]
    rows.append(
        {
            "source": source_name,
            "count_nonnull": int(values.shape[0]),
            "mean": float(values.mean()),
            "median": float(values.median()),
            "q25": float(values.quantile(0.25)),
            "q75": float(values.quantile(0.75)),
        }
    )

```

```

        "p95": float(values.quantile(0.95)),
        "max": float(values.max()),
        "pct_ge_80": float((values >= 80).mean()),
        "pct_ge_100": float((values >= 100).mean()),
        "pct_ge_150": float((values >= 150).mean()),
        "pct_eq_160": float((values == 160).mean()),
        "top_values": [
            str(key): int(value)
            for key, value in values.value_counts().head(10).items()
        ],
        "status": "ok",
        "missing_columns": ""
    }
)
return pd.DataFrame(rows,
    columns=columns).sort_values("source").reset_index(drop=True)

def normalize_temperature_to_f(temp: pd.Series) -> pd.DataFrame:
    """Normalize mixed-unit temperature values to cleaned Fahrenheit output.

    Rules:
    - values in ``[20, 50]`` are treated as Celsius-like and converted to
    ← Fahrenheit,
    - all other numeric values are treated as Fahrenheit already,
    - cleaned output is valid only for Fahrenheit ``[50, 120]``.
    """
    numeric = pd.to_numeric(temp, errors="coerce")
    celsius_like = numeric.between(20.0, 50.0, inclusive="both")
    temp_f = numeric.where(~celsius_like, numeric * 9.0 / 5.0 + 32.0)
    out_of_range = ~temp_f.between(50.0, 120.0, inclusive="both")
    temp_f_clean = temp_f.where(~out_of_range)
    return pd.DataFrame(
        {
            "temp_f_clean": temp_f_clean,
            "temp_was_celsius_like": celsius_like.astype("boolean"),
            "temp_out_of_range": out_of_range.astype("boolean"),
        }
    )

def clean_pain_score(pain: pd.Series) -> pd.DataFrame:
    """Clean pain score with explicit sentinel handling."""

```

```

        numeric = pd.to_numeric(pain, errors="coerce")
        sentinel = numeric.eq(13)
        out_of_range = numeric.notna() & (~sentinel) & ((numeric < 0) | (numeric
        > 10))
        clean = numeric.where(~(sentinel | out_of_range))
        return pd.DataFrame(
            {
                "pain_clean": clean,
                "pain_is_sentinel_13": sentinel.astype("boolean"),
                "pain_out_of_range": out_of_range.astype("boolean"),
            }
        )

def clean_bp(sbp: pd.Series, dbp: pd.Series) -> pd.DataFrame:
    """Range-clean SBP/DBP values with deterministic bounds."""
    sbp_num = pd.to_numeric(sbp, errors="coerce")
    dbp_num = pd.to_numeric(dbp, errors="coerce")
    sbp_out_of_range = sbp_num.notna() & ((sbp_num < 20) | (sbp_num > 300))
    dbp_out_of_range = dbp_num.notna() & ((dbp_num < 10) | (dbp_num > 200))
    return pd.DataFrame(
        {
            "sbp_clean": sbp_num.where(~sbp_out_of_range),
            "dbp_clean": dbp_num.where(~dbp_out_of_range),
            "sbp_out_of_range": sbp_out_of_range.astype("boolean"),
            "dbp_out_of_range": dbp_out_of_range.astype("boolean"),
        }
    )

def clean_o2sat(o2sat: pd.Series) -> pd.DataFrame:
    """Clean oxygen saturation and keep explicit >100 and zero flags."""
    numeric = pd.to_numeric(o2sat, errors="coerce")
    gt_100 = numeric > 100
    out_of_range = numeric.notna() & (numeric < 0)
    clean = numeric.where(~(gt_100 | out_of_range))
    return pd.DataFrame(
        {
            "o2sat_clean": clean,
            "o2sat_gt_100": gt_100.astype("boolean"),
            "o2sat_out_of_range": out_of_range.astype("boolean"),
            "o2sat_zero": numeric.eq(0).astype("boolean"),
        }
    )

```

```

    )

def _series_distribution_stats(values: pd.Series) -> dict[str, Any]:
    """Return compact distribution metrics for audit CSV artifacts."""
    numeric = pd.to_numeric(values, errors="coerce")
    total_rows = int(len(numeric))
    non_missing = int(numeric.notna().sum())
    if non_missing == 0:
        return {
            "n_non_missing": 0,
            "missing_pct": 100.0 if total_rows else 0.0,
            "min": np.nan,
            "p1": np.nan,
            "median": np.nan,
            "mean": np.nan,
            "p99": np.nan,
            "max": np.nan,
        }
    return {
        "n_non_missing": non_missing,
        "missing_pct": round(float((1.0 - non_missing / max(total_rows, 1)) *
                                     100.0), 2),
        "min": float(numeric.min()),
        "p1": float(numeric.quantile(0.01)),
        "median": float(numeric.median()),
        "mean": float(numeric.mean()),
        "p99": float(numeric.quantile(0.99)),
        "max": float(numeric.max()),
    }

def add_vitals_model_fields(
    ed_df: pd.DataFrame,
    *,
    ranges: Mapping[str, tuple[float, float]] | None = None,
) -> tuple[pd.DataFrame, pd.DataFrame]:
    """Create cleaned ED vitals fields and backward-compatible ``*_model`` aliases."""
    updated = ed_df.copy()
    resolved_ranges = dict(ranges or DEFAULT_VITALS_MODEL_RANGES)
    audit_rows: list[dict[str, Any]] = []

```

```

def _append_audit_row(
    *,
    raw_column: str,
    model_column: str,
    outlier_flag_column: str,
    lower_bound: float | None,
    upper_bound: float | None,
    numeric: pd.Series,
    outlier_flag: pd.Series,
) -> None:
    outlier_flag_numeric = outlier_flag.fillna(False).astype(bool)
    nonnull_n = int(numeric.notna().sum())
    outlier_n = int(outlier_flag_numeric.sum())
    examples = (
        numeric.loc[outlier_flag_numeric]
        .round(4)
        .value_counts()
        .head(5)
        .index.astype(str)
        .tolist()
    )
    audit_rows.append(
        {
            "domain": "vitals",
            "raw_column": raw_column,
            "model_column": model_column,
            "outlier_flag_column": outlier_flag_column,
            "lower_bound": lower_bound,
            "upper_bound": upper_bound,
            "nonnull_n": nonnull_n,
            "out_of_range_n": outlier_n,
            "out_of_range_pct": float(outlier_n / nonnull_n) if nonnull_n
                ↵ else 0.0,
            "example_outlier_values": "; ".join(examples),
        }
    )

for raw_column, (lower_bound, upper_bound) in resolved_ranges.items():
    if raw_column not in updated.columns:
        continue
    numeric = pd.to_numeric(updated[raw_column], errors="coerce")
    out_of_range = numeric.notna() & ((numeric < lower_bound) | (numeric
    ↵ > upper_bound))

```

```

model_column = f"{raw_column}_model"
outlier_flag_column = f"{raw_column}_outlier_flag"
updated[model_column] = numeric.where(~out_of_range)
updated[outlier_flag_column] = out_of_range.astype("boolean")
_append_audit_row(
    raw_column=raw_column,
    model_column=model_column,
    outlier_flag_column=outlier_flag_column,
    lower_bound=float(lower_bound),
    upper_bound=float(upper_bound),
    numeric=numeric,
    outlier_flag=out_of_range,
)
for prefix in ("ed_triage", "ed_first"):
    temp_raw_col = f"{prefix}_temp"
    pain_raw_col = f"{prefix}_pain"
    sbp_raw_col = f"{prefix}_sbp"
    dbp_raw_col = f"{prefix}_dbp"
    o2sat_raw_col = f"{prefix}_o2sat"

    if temp_raw_col in updated.columns:
        temp_clean = normalize_temperature_to_f(updated[temp_raw_col])
        temp_clean_col = f"{prefix}_temp_f_clean"
        temp_celsius_flag_col = f"{prefix}_temp_was_celsius_like"
        temp_range_flag_col = f"{prefix}_temp_out_of_range"
        updated[temp_clean_col] = temp_clean["temp_f_clean"]
        updated[temp_celsius_flag_col] =
        ↵ temp_clean["temp_was_celsius_like"]
        updated[temp_range_flag_col] = temp_clean["temp_out_of_range"]
        updated[f"{prefix}_temp_f_model"] = updated[temp_clean_col]
        updated[f"{prefix}_temp_c_model"] =
        ↵ (updated[f"{prefix}_temp_f_model"] - 32.0) * (5.0 / 9.0)
        updated[f"{prefix}_temp_model"] =
        ↵ updated[f"{prefix}_temp_f_model"]
        updated[f"{prefix}_temp_outlier_flag"] =
        ↵ updated[temp_range_flag_col].astype("boolean")
        _append_audit_row(
            raw_column=temp_raw_col,
            model_column=f"{prefix}_temp_model",
            outlier_flag_column=f"{prefix}_temp_outlier_flag",
            lower_bound=50.0,
            upper_bound=120.0,

```

```

        numeric=pd.to_numeric(updated[temp_raw_col],
↪  errors="coerce"),
          outlier_flag=updated[temp_range_flag_col],
    )

if pain_raw_col in updated.columns:
    pain_clean = clean_pain_score(updated[pain_raw_col])
    pain_clean_col = f"{prefix}_pain_clean"
    pain_sentinel_col = f"{prefix}_pain_is_sentinel_13"
    pain_range_col = f"{prefix}_pain_out_of_range"
    updated[pain_clean_col] = pain_clean["pain_clean"]
    updated[pain_sentinel_col] = pain_clean["pain_is_sentinel_13"]
    updated[pain_range_col] = pain_clean["pain_out_of_range"]
    updated[f"{prefix}_pain_model"] = updated[pain_clean_col]
    pain_outlier = updated[pain_sentinel_col].fillna(False) |
↪  updated[pain_range_col].fillna(False)
        updated[f"{prefix}_pain_outlier_flag"] =
↪  pain_outlier.astype("boolean")
        _append_audit_row(
            raw_column=pain_raw_col,
            model_column=f"{prefix}_pain_model",
            outlier_flag_column=f"{prefix}_pain_outlier_flag",
            lower_bound=0.0,
            upper_bound=10.0,
            numeric=pd.to_numeric(updated[pain_raw_col],
↪  errors="coerce"),
            outlier_flag=pain_outlier,
        )

if sbp_raw_col in updated.columns and dbp_raw_col in updated.columns:
    bp_clean = clean_bp(updated[sbp_raw_col], updated[dbp_raw_col])
    sbp_clean_col = f"{prefix}_sbp_clean"
    dbp_clean_col = f"{prefix}_dbp_clean"
    sbp_range_col = f"{prefix}_sbp_out_of_range"
    dbp_range_col = f"{prefix}_dbp_out_of_range"
    updated[sbp_clean_col] = bp_clean["sbp_clean"]
    updated[dbp_clean_col] = bp_clean["dbp_clean"]
    updated[sbp_range_col] = bp_clean["sbp_out_of_range"]
    updated[dbp_range_col] = bp_clean["dbp_out_of_range"]
    updated[f"{prefix}_sbp_model"] = updated[sbp_clean_col]
    updated[f"{prefix}_dbp_model"] = updated[dbp_clean_col]
    updated[f"{prefix}_sbp_outlier_flag"] =
↪  updated[sbp_range_col].astype("boolean")

```

```

        updated[f"{prefix}_dbp_outlier_flag"] =
    ↵ updated[dbp_range_col].astype("boolean")
        _append_audit_row(
            raw_column=sbp_raw_col,
            model_column=f"{prefix}_sbp_model",
            outlier_flag_column=f"{prefix}_sbp_outlier_flag",
            lower_bound=20.0,
            upper_bound=300.0,
            numeric=pd.to_numeric(updated[sbp_raw_col], errors="coerce"),
            outlier_flag=updated[dbp_range_col],
        )
        _append_audit_row(
            raw_column=dbp_raw_col,
            model_column=f"{prefix}_dbp_model",
            outlier_flag_column=f"{prefix}_dbp_outlier_flag",
            lower_bound=10.0,
            upper_bound=200.0,
            numeric=pd.to_numeric(updated[dbp_raw_col], errors="coerce"),
            outlier_flag=updated[dbp_range_col],
        )
    )

if o2sat_raw_col in updated.columns:
    o2sat_clean = clean_o2sat(updated[o2sat_raw_col])
    o2sat_clean_col = f"{prefix}_o2sat_clean"
    o2sat_gt_100_col = f"{prefix}_o2sat_gt_100"
    o2sat_range_col = f"{prefix}_o2sat_out_of_range"
    o2sat_zero_col = f"{prefix}_o2sat_zero"
    updated[o2sat_clean_col] = o2sat_clean["o2sat_clean"]
    updated[o2sat_gt_100_col] = o2sat_clean["o2sat_gt_100"]
    updated[o2sat_range_col] = o2sat_clean["o2sat_out_of_range"]
    updated[o2sat_zero_col] = o2sat_clean["o2sat_zero"]
    updated[f"{prefix}_o2sat_model"] = updated[o2sat_clean_col]
    o2sat_outlier = updated[o2sat_gt_100_col].fillna(False) |
    ↵ updated[o2sat_range_col].fillna(False)
        updated[f"{prefix}_o2sat_outlier_flag"] =
    ↵ o2sat_outlier.astype("boolean")
        _append_audit_row(
            raw_column=o2sat_raw_col,
            model_column=f"{prefix}_o2sat_model",
            outlier_flag_column=f"{prefix}_o2sat_outlier_flag",
            lower_bound=0.0,
            upper_bound=100.0,

```

```

        numeric=pd.to_numeric(updated[o2sat_raw_col] ,
    ↵  errors="coerce"),
        outlier_flag=o2sat_outlier,
    )

audit =
pd.DataFrame(audit_rows).sort_values("raw_column").reset_index(drop=True)
return updated, audit

def build_ed_vitals_audit_artifacts(
    ed_df: pd.DataFrame,
) -> tuple[pd.DataFrame, pd.DataFrame, pd.DataFrame]:
    """Build ED-vitals distribution, extreme-example, and raw-to-clean delta
    ↵  artifacts."""
    specs = [
        {
            "vital_name": "triage_temp",
            "raw_col": "ed_triage_temp",
            "clean_col": "ed_triage_temp_f_clean",
            "flags": ["ed_triage_temp_was_celsius_like",
                      ↵  "ed_triage_temp_out_of_range"],
            "extra_cols": [],
        },
        {
            "vital_name": "first_temp",
            "raw_col": "ed_first_temp",
            "clean_col": "ed_first_temp_f_clean",
            "flags": ["ed_first_temp_was_celsius_like",
                      ↵  "ed_first_temp_out_of_range"],
            "extra_cols": ["ed_first_vitals_time"],
        },
        {
            "vital_name": "triage_pain",
            "raw_col": "ed_triage_pain",
            "clean_col": "ed_triage_pain_clean",
            "flags": ["ed_triage_pain_is_sentinel_13",
                      ↵  "ed_triage_pain_out_of_range"],
            "extra_cols": [],
        },
        {
            "vital_name": "first_pain",
            "raw_col": "ed_first_pain",

```

```

"clean_col": "ed_first_pain_clean",
"flags": ["ed_first_pain_is_sentinel_13",
    ↳ "ed_first_pain_out_of_range"],
"extra_cols": ["ed_first_vitals_time"],
},
{
    "vital_name": "triage_sbp",
    "raw_col": "ed_triage_sbp",
    "clean_col": "ed_triage_sbp_clean",
    "flags": ["ed_triage_sbp_out_of_range"],
    "extra_cols": [],
},
{
    "vital_name": "triage_dbp",
    "raw_col": "ed_triage_dbp",
    "clean_col": "ed_triage_dbp_clean",
    "flags": ["ed_triage_dbp_out_of_range"],
    "extra_cols": [],
},
{
    "vital_name": "first_sbp",
    "raw_col": "ed_first_sbp",
    "clean_col": "ed_first_sbp_clean",
    "flags": ["ed_first_sbp_out_of_range"],
    "extra_cols": ["ed_first_vitals_time"],
},
{
    "vital_name": "first_dbp",
    "raw_col": "ed_first_dbp",
    "clean_col": "ed_first_dbp_clean",
    "flags": ["ed_first_dbp_out_of_range"],
    "extra_cols": ["ed_first_vitals_time"],
},
{
    "vital_name": "triage_o2sat",
    "raw_col": "ed_triage_o2sat",
    "clean_col": "ed_triage_o2sat_clean",
    "flags": ["ed_triage_o2sat_gt_100",
        ↳ "ed_triage_o2sat_out_of_range", "ed_triage_o2sat_zero"],
    "extra_cols": [],
},
{
    "vital_name": "first_o2sat",

```

```

    "raw_col": "ed_first_o2sat",
    "clean_col": "ed_first_o2sat_clean",
    "flags": ["ed_first_o2sat_gt_100", "ed_first_o2sat_out_of_range",
              "ed_first_o2sat_zero"],
    "extra_cols": ["ed_first_vitals_time"],
},
]

summary_rows: list[dict[str, Any]] = []
delta_rows: list[dict[str, Any]] = []
extreme_frames: list[pd.DataFrame] = []

for spec in specs:
    raw_col = spec["raw_col"]
    clean_col = spec["clean_col"]
    if raw_col not in ed_df.columns or clean_col not in ed_df.columns:
        continue

    raw = pd.to_numeric(ed_df[raw_col], errors="coerce")
    clean = pd.to_numeric(ed_df[clean_col], errors="coerce")
    raw_stats = _series_distribution_stats(raw)
    clean_stats = _series_distribution_stats(clean)
    row: dict[str, Any] = {
        "vital_name": spec["vital_name"],
        "raw_column": raw_col,
        "clean_column": clean_col,
        "raw_n_non_missing": raw_stats["n_non_missing"],
        "raw_missing_pct": raw_stats["missing_pct"],
        "raw_min": raw_stats["min"],
        "raw_p1": raw_stats["p1"],
        "raw_median": raw_stats["median"],
        "raw_mean": raw_stats["mean"],
        "raw_p99": raw_stats["p99"],
        "raw_max": raw_stats["max"],
        "clean_n_non_missing": clean_stats["n_non_missing"],
        "clean_missing_pct": clean_stats["missing_pct"],
        "clean_min": clean_stats["min"],
        "clean_p1": clean_stats["p1"],
        "clean_median": clean_stats["median"],
        "clean_mean": clean_stats["mean"],
        "clean_p99": clean_stats["p99"],
        "clean_max": clean_stats["max"],
    }
}

```

```

delta_rows.append(
    {
        "vital_name": spec["vital_name"],
        "raw_column": raw_col,
        "clean_column": clean_col,
        "raw_n": int(raw.notna().sum()),
        "clean_n": int(clean.notna().sum()),
        "raw_removed_n": int(raw.notna().sum() -
            ↪ clean.notna().sum()),
        "raw_removed_pct": (
            float((raw.notna().sum() - clean.notna().sum()) /
            ↪ max(int(raw.notna().sum()), 1))
            if int(raw.notna().sum()) > 0
            else 0.0
        ),
        "converted_celsius_like_n": int(
            ed_df.get(spec["flags"][0], pd.Series(False,
            ↪ index=ed_df.index))
                .fillna(False)
                .astype(bool)
                .sum()
        )
        if "temp_was_celsius_like" in spec["flags"][0]
        else 0,
    }
)

available_flags = [flag for flag in spec["flags"] if flag in
↪ ed_df.columns]
for flag_col in available_flags:
    flag_series = ed_df[flag_col].fillna(False).astype(bool)
    row[f"{flag_col}_count"] = int(flag_series.sum())
    row[f"{flag_col}_pct"] = float(flag_series.mean()) if
↪ len(flag_series) else 0.0
    if not flag_series.any():
        continue
    include_cols = ["hadm_id", "ed_stay_id", raw_col, clean_col,
↪ flag_col]
    include_cols += [col for col in spec["extra_cols"] if col in
↪ ed_df.columns]
    include_cols = [col for col in include_cols if col in
↪ ed_df.columns]

```

```

        flagged = ed_df.loc[flag_series, include_cols].copy()
        flagged = flagged.rename(columns={raw_col: "raw_value",
        ← clean_col: "clean_value", flag_col: "flag_value"})
        low = flagged.sort_values("raw_value", ascending=True,
        ← na_position="last").head(5).copy()
        low["tail"] = "low"
        high = flagged.sort_values("raw_value", ascending=False,
        ← na_position="last").head(5).copy()
        high["tail"] = "high"
        extremes = pd.concat([low, high],
        ← ignore_index=True).drop_duplicates()
        extremes.insert(0, "vital_name", spec["vital_name"])
        extremes.insert(1, "raw_column", raw_col)
        extremes.insert(2, "clean_column", clean_col)
        extremes.insert(3, "flag_type", flag_col)
        extreme_frames.append(extremes)

    summary_rows.append(row)

distribution = pd.DataFrame(summary_rows)
delta = pd.DataFrame(delta_rows)
if extreme_frames:
    extremes = pd.concat(extreme_frames, ignore_index=True, sort=False)
else:
    extremes = pd.DataFrame(
        columns=[
            "vital_name",
            "raw_column",
            "clean_column",
            "flag_type",
            "hadm_id",
            "ed_stay_id",
            "raw_value",
            "clean_value",
            "flag_value",
            "tail",
        ]
    )
return distribution, extremes, delta

def evaluate_uom_expectations(ed_df: pd.DataFrame) -> dict[str, Any]:
    """Evaluate expected-null and value/uom consistency rules."""

```

```

structural_nulls: dict[str, dict[str, Any]] = {}
for field_name in EXPECTED_STRUCTURAL_NULL_FIELDS:
    if field_name not in ed_df.columns:
        continue
    structural_nulls[field_name] = {
        "present": True,
        "missing_n": int(ed_df[field_name].isna().sum()),
        "missing_pct": float(ed_df[field_name].isna().mean()),
        "all_null": bool(ed_df[field_name].isna().all()),
    }
}

paco2_checks: dict[str, dict[str, Any]] = {}
for value_column, uom_column in PACO2_VALUE_UOM_PAIRS:
    if value_column not in ed_df.columns or uom_column not in
    ↪ ed_df.columns:
        paco2_checks[uom_column] = {
            "present": False,
            "reason": "missing_value_or_uom_column",
        }
        continue

    value_present = ed_df[value_column].notna()
    uom_lower = ed_df[uom_column].astype(str).str.strip().str.lower()
    missing_uom_with_value = int((value_present &
    ↪ ed_df[uom_column].isna()).sum())
    non_mmhg_uom_with_value = int(
        (value_present & ed_df[uom_column].notna() &
    ↪ uom_lower.ne("mmhg")).sum()
    )

    paco2_checks[uom_column] = {
        "present": True,
        "paired_value_column": value_column,
        "value_rows": int(value_present.sum()),
        "missing_uom_when_value_present": missing_uom_with_value,
        "non_mmhg_uom_when_value_present": non_mmhg_uom_with_value,
        "passes": bool(
            missing_uom_with_value == 0 and non_mmhg_uom_with_value == 0
        ),
    }

return {

```

```

    "expected_structural_null_fields":
        ↵  list(EXPECTED_STRUCTURAL_NULL_FIELDS),
    "structural_null_checks": structural_nulls,
    "paco2_uom_checks": paco2_checks,
}

def classify_missingness_expectations(
    ed_df: pd.DataFrame,
    target_fields: list[str],
    *,
    expected_sparse_fields: set[str] | None = None,
) -> pd.DataFrame:
    """Classify field-level missingness into expected and unexpected
    categories."""
    rows: list[dict[str, Any]] = []
    total_rows = max(int(len(ed_df)), 1)
    expected_structural = set(EXPECTED_STRUCTURAL_NULL_FIELDS)
    expected_sparse = set(expected_sparse_fields or set())

    for field_name in target_fields:
        if field_name not in ed_df.columns:
            rows.append(
                {
                    "field": field_name,
                    "missing_n": int(len(ed_df)),
                    "missing_pct": 1.0,
                    "expectation": "missing_column",
                }
            )
        continue

        missing_n = int(ed_df[field_name].isna().sum())
        missing_pct = float(missing_n / total_rows)

        if field_name in expected_structural:
            expectation = "expected_structural_null"
        elif field_name in expected_sparse and missing_pct >= 1.0:
            expectation = "expected_sparse"
        elif missing_pct >= 1.0:
            expectation = "unexpected_full_null"
        elif missing_pct > 0.0:
            expectation = "conditional_sparse"

```

```

        else:
            expectation = "complete"

        rows.append(
            {
                "field": field_name,
                "missing_n": missing_n,
                "missing_pct": missing_pct,
                "expectation": expectation,
            }
        )

    return pd.DataFrame(rows)

def render_latex_longtable(
    table_df: pd.DataFrame,
    *,
    caption: str,
    label: str,
    landscape: bool = False,
    index: bool = True,
) -> str:
    latex_text = table_df.to_latex(
        index=index,
        escape=False,
        longtable=True,
        caption=caption,
        label=label,
    )
    if landscape:
        latex_text = "\\begin{landscape}\\n" + latex_text +
        "\\n\\end{landscape}\\n"
    return latex_text

from hypercap_cc_nlp.pipeline_audit import collect_run_manifest
from hypercap_cc_nlp.pipeline_contracts import validate_cohort_contract,
    write_contract_report

# ---- Backend selection (we use BigQuery)
BACKEND = os.getenv("MIMIC_BACKEND", "bigquery").strip().lower()
assert BACKEND == "bigquery", "This notebook is BigQuery-specific."

```

```

WORK_PROJECT = os.getenv("WORK_PROJECT", "").strip() # your billing project
PHYS = os.getenv("BQ_PHYSIONET_PROJECT", "physionet-data").strip() # hosting
    ↵ project (read-only)

# Dataset preferences: resolved to accessible datasets in the next setup
    ↵ cell.
HOSP = os.getenv("BQ_DATASET_HOSP", "mimiciv_3_1_hosp").strip()
ICU  = os.getenv("BQ_DATASET_ICU",  "mimiciv_3_1_icu").strip()
ED   = os.getenv("BQ_DATASET_ED",   "").strip()

# BigQuery client
client = bigquery.Client(project=WORK_PROJECT)

print("Project:", WORK_PROJECT)
print("PhysioNet host:", PHYS)
print("HOSP (pref):", HOSP, "ICU (pref):", ICU, "ED (pref):", ED)
print("WORK_DIR:", WORK_DIR)

Project: mimic-hypercapnia
PhysioNet host: physionet-data
HOSP (pref): mimiciv_3_1_hosp ICU (pref): mimiciv_3_1_icu ED (pref): mimiciv_ed
WORK_DIR: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Proj
CC-NLP

# Purpose: Define reusable BigQuery helpers and resolve dataset names across
    ↵ known naming variants.

from datetime import datetime, timezone

# Fail fast on long-running queries instead of hanging indefinitely in
    ↵ nbconvert.
BQ_QUERY_TIMEOUT_SECS = int(os.getenv("BQ_QUERY_TIMEOUT_SECS", "1800"))

# --- Helper: run SQL with optional named parameters
def run_sql_bq(sql: str, params: dict | None = None) -> pd.DataFrame:
    job_config = bigquery.QueryJobConfig()
    job_config.labels = {"pipeline": "hypercapcc", "notebook": "cohort"}
    if params:
        bq_params = []
        for k, v in params.items():
            if isinstance(v, (list, tuple, np.ndarray, pd.Series)):
                # BigQuery ARRAY<INT64> if all ints; else ARRAY<STRING>

```

```

        v_list = list(v)
        if all(isinstance(x, (int, np.integer)) for x in v_list):
            bq_params.append(bigquery.ArrayQueryParameter(k, "INT64",
        ↵  list(map(int, v_list))))
        else:
            bq_params.append(bigquery.ArrayQueryParameter(k,
        ↵  "STRING", list(map(str, v_list))))
        else:
            # scalar
            if isinstance(v, (int, np.integer)):
                bq_params.append(bigquery.ScalarQueryParameter(k,
        ↵  "INT64", int(v)))
            elif isinstance(v, float):
                bq_params.append(bigquery.ScalarQueryParameter(k,
        ↵  "FLOAT64", float(v)))
            else:
                bq_params.append(bigquery.ScalarQueryParameter(k,
        ↵  "STRING", str(v)))
        job_config.query_parameters = bq_params

    job = client.query(sql, job_config=job_config)
    try:
        result = job.result(timeout=BQ_QUERY_TIMEOUT_SECS)
    except Exception as exc:
        try:
            job.cancel()
        except Exception:
            pass
        raise RuntimeError(
            f"BigQuery query failed or timed out after
            ↵  {BQ_QUERY_TIMEOUT_SECS}s (job_id={job.job_id})."
        ) from exc

    try:
        return result.to_dataframe(create_bqstorage_client=True)
    except TypeError:
        return result.to_dataframe()

# --- Helper: test if a fully-qualified table exists and is accessible
def table_exists(fqtn: str) -> bool:
    try:
        _ = run_sql_bq(f"SELECT 1 FROM `fqtn` LIMIT 1")
    return True

```

```

except Exception:
    return False

def resolve_dataset(preferred: str, candidates: list[str], probe_table: str,
                   label: str) -> tuple[str, str]:
    # Keep preferred first, then try known aliases.
    ordered = []
    if preferred:
        ordered.append(preferred)
    for cand in candidates:
        if cand not in ordered:
            ordered.append(cand)

    for dataset in ordered:
        fqtn = f"{PHYS}.{dataset}.{probe_table}"
        if table_exists(fqtn):
            return dataset, fqtn

    raise RuntimeError(
        f"No accessible {label} dataset found for probe table "
        f'{probe_table}'. Tried: {ordered}"
    )

# Resolve HOSP/ICU/ED names so notebook runs across v3.1 naming variants.
HOSP, HOSP_PROBE = resolve_dataset(
    preferred=HOSP,
    candidates=["mimiciv_3_1_hosp", "mimiciv_v3_1_hosp", "mimiciv_hosp"],
    probe_table="admissions",
    label="HOSP",
)
ICU, ICU_PROBE = resolve_dataset(
    preferred=ICU,
    candidates=["mimiciv_3_1_icu", "mimiciv_v3_1_icu", "mimiciv_icu"],
    probe_table="icustays",
    label="ICU",
)
if not ED:
    ED = "mimiciv_ed"
ED, ED_PROBE = resolve_dataset(
    preferred=ED,
    candidates=["mimiciv_ed", "mimiciv_3_1_ed", "mimiciv_v3_1_ed"],
    probe_table="edstays",
    label="ED",
)

```

```
)
```

```
RUN_METADATA = {
    "run_utc": datetime.now(timezone.utc).isoformat(),
    "work_project": WORK_PROJECT,
    "physionet_project": PHYS,
    "datasets": {"hosp": HOSP, "icu": ICU, "ed": ED},
    "dataset_probes": {"hosp": HOSP_PROBE, "icu": ICU_PROBE, "ed": ED_PROBE},
    "notebook": "MIMICIV_hypercap_EXT_cohort.ipynb",
    "blood_gas_manifest_path": BLOOD_GAS_MANIFEST["path"],
    "blood_gas_manifest_version": BLOOD_GAS_MANIFEST["version"],
}

print("Resolved datasets:", RUN_METADATA["datasets"])
print("Dataset probes:", RUN_METADATA["dataset_probes"])

Resolved datasets: {'hosp': 'mimiciv_3_1_hosp', 'icu': 'mimiciv_3_1_icu', 'ed': 'mimiciv_ed'}
Dataset probes: {'hosp': 'physionet-data.mimiciv_3_1_hosp.admissions', 'icu': 'physionet-data.mimiciv_3_1_icu.icustays', 'ed': 'physionet-data.mimiciv_ed.edstays'}
```

```
# Purpose: Create reusable data-quality and merge guardrail helpers to
# prevent silent join errors.

# --- Helper utilities for reproducibility and safe joins
def require_cols(df: pd.DataFrame, cols: list[str], name: str) -> None:
    missing = [c for c in cols if c not in df.columns]
    if missing:
        raise KeyError(f"{name} missing columns: {missing}")

def assert_unique(df: pd.DataFrame, key: str, name: str) -> None:
    if df[key].duplicated().any():
        n = int(df[key].duplicated().sum())
        raise ValueError(f"{name} has {n} duplicate {key} values")

def safe_merge(left: pd.DataFrame, right: pd.DataFrame, on: list[str] | str,
# how: str, name: str) -> pd.DataFrame:
    # guard against accidental duplicate columns
    overlap = set(left.columns) & set(right.columns)
    if isinstance(on, str):
        on_cols = {on}
    else:
        on_cols = set(on)
```

```

overlap = overlap - on_cols
if overlap:
    raise ValueError(f"{name} merge would duplicate columns:
        ↵ {sorted(overlap)}")
return left.merge(right, on=on, how=how)

def check_ranges(df: pd.DataFrame, ranges: dict[str, tuple[float, float]])-
    ↵ -> pd.DataFrame:
    rows = []
    for col, (lo, hi) in ranges.items():
        if col not in df.columns:
            continue
        bad = df[col].notna() & ((df[col] < lo) | (df[col] > hi))
        rows.append({"col": col, "n_bad": int(bad.sum())})
    return pd.DataFrame(rows)

```

2.3 Data Generation

2.3.1 1) ICD cohort flags (hypercapnic respiratory failure)

Rationale: Capture diagnosis-based hypercapnia from ED and hospital discharge codes to define a broad, clinically recognized cohort.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
    ↵ consistent across notebook stages.

# Target ICD codes (dotless, uppercase)
ICD10_CODES = ['J9602', 'J9612', 'J9622', 'J9692', 'E662']
ICD9_CODES = ['27803']

SQL["cohort_icd_sql"] = f"""
-- ICD-based cohort flags per admission
WITH target_codes AS (
    SELECT 'J9602' AS code, 10 AS ver UNION ALL
    SELECT 'J9612', 10 UNION ALL
    SELECT 'J9622', 10 UNION ALL
    SELECT 'J9692', 10 UNION ALL
    SELECT 'E662', 10 UNION ALL
    SELECT '27803', 9
),
-- Hospital ICDs restricted to target codes
hosp_dx AS (

```

```

SELECT
    d.subject_id,
    d.hadm_id,
    UPPER(REPLACE(d.icd_code, '.', '')) AS code_norm,
    d.icd_version
FROM `'{PHYS}`.{HOSP}.diagnoses_icd` d
JOIN target_codes t
    ON t.ver = d.icd_version AND t.code = UPPER(REPLACE(d.icd_code, '.', ''))
WHERE d.hadm_id IS NOT NULL
),

-- Hospital flags per admission
hosp_flags AS (
    SELECT
        subject_id, hadm_id,
        MAX(IF(icd_version=10 AND code_norm='J9602',1,0)) AS ICD10_J9602,
        MAX(IF(icd_version=10 AND code_norm='J9612',1,0)) AS ICD10_J9612,
        MAX(IF(icd_version=10 AND code_norm='J9622',1,0)) AS ICD10_J9622,
        MAX(IF(icd_version=10 AND code_norm='J9692',1,0)) AS ICD10_J9692,
        MAX(IF(icd_version=10 AND code_norm='E662', 1,0)) AS ICD10_E662,
        MAX(IF(icd_version=9  AND code_norm='27803',1,0)) AS ICD9_27803
    FROM hosp_dx
    GROUP BY subject_id, hadm_id
),

-- ED ICDs restricted to target codes (map to hadm via edstays)
ed_dx AS (
    SELECT
        s.subject_id,
        s.hadm_id,
        s.stay_id,
        s.intime AS ed_intime,
        UPPER(REPLACE(d.icd_code, '.', '')) AS code_norm,
        d.icd_version
    FROM `'{PHYS}`.{ED}.diagnosis` d
    JOIN `'{PHYS}`.{ED}.edstays` s
        ON s.subject_id = d.subject_id AND s.stay_id = d.stay_id
    JOIN target_codes t
        ON t.ver = d.icd_version AND t.code = UPPER(REPLACE(d.icd_code, '.', ''))
    WHERE s.hadm_id IS NOT NULL
),

```

```

-- ED flags per ED stay (so we can both: OR flags across stays and also pick
-- earliest stay_id)
ed_flags_by_stay AS (
    SELECT
        subject_id, hadm_id, stay_id, MIN(ed_intime) AS ed_intime,
        MAX(IF(icd_version=10 AND code_norm='J9602',1,0)) AS ICD10_J9602,
        MAX(IF(icd_version=10 AND code_norm='J9612',1,0)) AS ICD10_J9612,
        MAX(IF(icd_version=10 AND code_norm='J9622',1,0)) AS ICD10_J9622,
        MAX(IF(icd_version=10 AND code_norm='J9692',1,0)) AS ICD10_J9692,
        MAX(IF(icd_version=10 AND code_norm='E662', 1,0)) AS ICD10_E662,
        MAX(IF(icd_version=9  AND code_norm='27803',1,0)) AS ICD9_27803
    FROM ed_dx
    GROUP BY subject_id, hadm_id, stay_id
),

-- OR the ED flags across all ED stays mapped to the same hadm
ed_flags_or AS (
    SELECT
        subject_id, hadm_id,
        MAX(ICD10_J9602) AS ICD10_J9602,
        MAX(ICD10_J9612) AS ICD10_J9612,
        MAX(ICD10_J9622) AS ICD10_J9622,
        MAX(ICD10_J9692) AS ICD10_J9692,
        MAX(ICD10_E662 ) AS ICD10_E662,
        MAX(ICD9_27803) AS ICD9_27803
    FROM ed_flags_by_stay
    GROUP BY subject_id, hadm_id
),

-- Earliest ED stay_id per hadm (NO UNNEST of aggregates; use [OFFSET(0)])
ed_earliest AS (
    SELECT
        subject_id,
        hadm_id,
        (ARRAY_AGG(STRUCT(stay_id, ed_intime) ORDER BY ed_intime LIMIT
        ← 1))[OFFSET(0)].stay_id AS stay_id
    FROM ed_flags_by_stay
    GROUP BY subject_id, hadm_id
),

-- Bring flags and earliest stay_id together
ed_by_hadm AS (
    SELECT

```

```

f.subject_id,
f.hadm_id,
e.stay_id,
f.ICD10_J9602,
f.ICD10_J9612,
f.ICD10_J9622,
f.ICD10_J9692,
f.ICD10_E662,
f.ICD9_27803
FROM ed_flags_or f
LEFT JOIN ed_earliest e
    USING (subject_id, hadm_id)
),
-- Combine ED and hospital flags at the admission level
combined AS (
    SELECT
        COALESCE(h.subject_id, e.subject_id) AS subject_id,
        COALESCE(h.hadm_id, e.hadm_id) AS hadm_id,
        GREATEST(IFNULL(h.ICD10_J9602,0), IFNULL(e.ICD10_J9602,0)) AS
        ↵ ICD10_J9602,
        GREATEST(IFNULL(h.ICD10_J9612,0), IFNULL(e.ICD10_J9612,0)) AS
        ↵ ICD10_J9612,
        GREATEST(IFNULL(h.ICD10_J9622,0), IFNULL(e.ICD10_J9622,0)) AS
        ↵ ICD10_J9622,
        GREATEST(IFNULL(h.ICD10_J9692,0), IFNULL(e.ICD10_J9692,0)) AS
        ↵ ICD10_J9692,
        GREATEST(IFNULL(h.ICD10_E662 ,0), IFNULL(e.ICD10_E662 ,0)) AS ICD10_E662,
        GREATEST(IFNULL(h.ICD9_27803,0), IFNULL(e.ICD9_27803,0)) AS ICD9_27803,
        ↵ IF((IFNULL(h.ICD10_J9602,0)+IFNULL(h.ICD10_J9612,0)+IFNULL(h.ICD10_J9622,0)+IFNULL(h.ICD10_J9692,0)+IFNULL(h.ICD10_E662 ,0)+IFNULL(h.ICD9_27803,0)) > 0, 1, 0) AS any_hypercap_icd_hosp,
        ↵ IF((IFNULL(e.ICD10_J9602,0)+IFNULL(e.ICD10_J9612,0)+IFNULL(e.ICD10_J9622,0)+IFNULL(e.ICD10_J9692,0)+IFNULL(e.ICD10_E662 ,0)+IFNULL(e.ICD9_27803,0)) > 0, 1, 0) AS any_hypercap_icd_ed
    FROM hosp_flags h
    FULL OUTER JOIN ed_by_hadm e
        ON h.hadm_id = e.hadm_id
)
SELECT
    subject_id, hadm_id,
    ICD10_J9602, ICD10_J9612, ICD10_J9622, ICD10_J9692, ICD10_E662, ICD9_27803,

```

```

IF((ICD10_J9602+ICD10_J9612+ICD10_J9622+ICD10_J9692+ICD10_E662+ICD9_27803)
↳ > 0, 1, 0) AS any_hypercap_icd,
any_hypercap_icd_hosp,
any_hypercap_icd_ed,
CASE
    WHEN any_hypercap_icd_hosp=1 AND any_hypercap_icd_ed=1 THEN 'ED+HOSP'
    WHEN any_hypercap_icd_ed=1 THEN 'ED'
    WHEN any_hypercap_icd_hosp=1 THEN 'HOSP'
    ELSE 'NONE'
END AS icd_source
FROM combined
"""

cohort_icd = run_sql_bq(sql("cohort_icd_sql"))
print("ICD cohort admissions:", len(cohort_icd))
cohort_icd.head(3)

ICD cohort admissions: 4237

```

	subject_id	hadm_id	ICD10_J9602	ICD10_J9612	ICD10_J9622	ICD10_J9692	ICD10_E662	I
0	10485425	21207827	0	0	1	0	1	0
1	16826447	21830147	0	1	0	0	1	0
2	11482582	22869017	0	0	0	0	0	1

2.3.2 2) Blood gas inclusion thresholds — ANY PaCO meeting cutoffs (LAB + POC)

Rationale: Identify physiologic hypercapnia using ABG/VBG thresholds, independent of diagnostic coding.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↳ consistent across notebook stages.
```

```
SQL["co2_thresholds_sql"] = f"""
WITH hosp_raw AS (
    SELECT
        le.hadm_id,
        le.charttime,
        le.specimen_id,
        COALESCE(
            CAST(le.valuenum AS FLOAT64),
```

```

        SAFE_CAST(REGEXP_EXTRACT(LOWER(le.value), r'(-?\d+(?:\.\d+)?)) AS
↳ FLOAT64)
) AS val,
LOWER(REPLACE(COALESCE(le.valueuom, ''), ' ', '')) AS uom_nospace,
LOWER(COALESCE(di.label, '')) AS lbl,
LOWER(COALESCE(di.fluid, '')) AS fluid_text
FROM `{{PHYS}}.{{HOSP}}.labevents` le
JOIN `{{PHYS}}.{{HOSP}}.d_labitems` di
    ON di.itemid = le.itemid
WHERE le.itemid IN UNNEST({{LAB_PCO2_SQL_ARRAY}})
    AND (le.valuenum IS NOT NULL OR le.value IS NOT NULL)
    AND LOWER(COALESCE(di.fluid, '')) IN {{LAB_BLOOD_FLUID_SQL_LIST}}
),
hosp_specimen AS (
    SELECT
        le.specimen_id,
        LOWER(COALESCE(le.value, '')) AS specimen_type_text
    FROM `{{PHYS}}.{{HOSP}}.labevents` le
    WHERE le.itemid IN UNNEST({{LAB_SPECIMEN_SQL_ARRAY}})
        AND le.specimen_id IS NOT NULL
        AND le.value IS NOT NULL
    QUALIFY ROW_NUMBER() OVER (
        PARTITION BY le.specimen_id
        ORDER BY le.charttime DESC
    ) = 1
),
hosp_pco2_clean AS (
    SELECT
        r.hadm_id,
        CASE
            WHEN REGEXP_CONTAINS(COALESCE(s.specimen_type_text, ''), '
↳ r'arter|\\bart\\b') THEN 'arterial'
            WHEN REGEXP_CONTAINS(COALESCE(s.specimen_type_text, ''), '
↳ r'ven|central|mixed') THEN 'venous'
            ELSE 'other'
        END AS site,
        r.charttime,
        CASE WHEN r.uom_nospace = 'kpa' THEN r.val * 7.50062 ELSE r.val END AS
↳ pco2_mmhg,
        'LAB' AS source_branch
    FROM hosp_raw r
    LEFT JOIN hosp_specimen s USING (specimen_id)
    WHERE r.val IS NOT NULL

```

```

        AND (CASE WHEN r.uom_nospace = 'kpa' THEN r.val * 7.50062 ELSE r.val END)
        ↵ BETWEEN 5 AND 200
),
icu_pco2_itemids AS (
    SELECT itemid, LOWER(label) AS lbl
    FROM `{{PHYS}}.{{ICU}}.d_items`
    WHERE itemid IN UNNEST({{ICU_PCO2_SQL_ARRAY}})
),
icu_pco2_raw AS (
    SELECT
        ie.hadm_id,
        ce.stay_id,
        ce.charttime,
        di.lbl,
        LOWER(REPLACE(COALESCE(ce.value uom, ''), ' ', '')) AS uom_nospace,
        COALESCE(
            CAST(ce.valuenum AS FLOAT64),
            SAFE_CAST(REGEXP_EXTRACT(LOWER(ce.value), r'(-?\d+(?:\.\d+)?|)') AS
        ↵ FLOAT64)
        ) AS val
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    JOIN icu_pco2_itemids di
        ON di.itemid = ce.itemid
    JOIN `{{PHYS}}.{{ICU}}.icustays` ie
        ON ie.stay_id = ce.stay_id
    WHERE ce.valuenum IS NOT NULL OR ce.value IS NOT NULL
),
icu_specimen AS (
    SELECT
        ce.stay_id,
        ce.charttime,
        LOWER(COALESCE(ce.value, '')) AS specimen_type_text
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    WHERE ce.itemid IN UNNEST({{ICU_SPECIMEN_SQL_ARRAY}})
        AND ce.value IS NOT NULL
),
icu_pco2_ranked AS (
    SELECT
        r.*,
        s.specimen_type_text,
        ROW_NUMBER() OVER (
            PARTITION BY r.hadm_id, r.stay_id, r.charttime
            ORDER BY ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, SECOND))
)

```

```

) AS specimen_rank
FROM icu_pco2_raw r
LEFT JOIN icu_specimen s
    ON s.stay_id = r.stay_id
    AND ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, MINUTE)) <= 120
),
icu_pco2_clean AS (
SELECT
    hadm_id,
    CASE
        WHEN REGEXP_CONTAINS(COALESCE(specimen_type_text, ''), 
        ↵ r'arter|\\bart\\b') THEN 'arterial'
        WHEN REGEXP_CONTAINS(COALESCE(specimen_type_text, ''), 
        ↵ r'ven|central|mixed') THEN 'venous'
        ELSE 'other'
    END AS site,
    charttime,
    CASE WHEN uom_nospace = 'kpa' THEN val * 7.50062 ELSE val END AS
    ↵ pco2_mmhg,
    'POC' AS source_branch
FROM icu_pco2_ranked
WHERE val IS NOT NULL
    AND (specimen_rank = 1 OR specimen_type_text IS NULL)
    AND (CASE WHEN uom_nospace = 'kpa' THEN val * 7.50062 ELSE val END)
    ↵ BETWEEN 5 AND 200
),
all_pco2 AS (
SELECT * FROM hosp_pco2_clean
UNION ALL
SELECT * FROM icu_pco2_clean
),
thresh AS (
SELECT
    hadm_id,
    MAX(IF(site = 'arterial' AND pco2_mmhg >= 45.0, 1, 0)) AS
    ↵ abg_hypercap_threshold,
    MAX(IF(site = 'venous' AND pco2_mmhg >= 50.0, 1, 0)) AS
    ↵ vbg_hypercap_threshold,
    MAX(IF(site = 'other' AND source_branch = 'LAB' AND pco2_mmhg >= 50.0,
    ↵ 1, 0)) AS other_hypercap_threshold,
    MAX(IF(site = 'other' AND source_branch = 'POC' AND pco2_mmhg >= 50.0,
    ↵ 1, 0)) AS poc_other_quarantined_hypercap_threshold
FROM all_pco2

```

```

        GROUP BY hadm_id
    )
SELECT * FROM thresh
"""

co2_thresh = run_sql_bq(sql("co2_thresholds_sql"))
print("Admissions meeting thresholds:", len(co2_thresh))
co2_thresh.head(3)

```

Admissions meeting thresholds: 82874

	hadm_id	abg_hypercap_threshold	vbg_hypercap_threshold	other_hypercap_threshold	poc_other_quarantined_hypercap_threshold
0	25143777	0	0	0	0
1	20278008	0	0	0	0
2	22459966	0	0	0	0

2.3.3 3) Cohort union (ICD thresholds) and hadm_list for downstream queries

Rationale: Combine ICD and physiologic routes to maximize sensitivity, then define a stable hadm_id list for downstream joins.

```

# Purpose: Combine ICD and gas-threshold ascertainment routes to produce the
#           final hadm inclusion list.

# Outer-join because thresholds can identify hadm_id with no ICD codes and
#           vice versa
cohort_any = cohort_icd.merge(co2_thresh, how="outer", on="hadm_id")

# Fill missing flags with 0 where appropriate
icd_cols = ["ICD10_J9602", "ICD10_J9612", "ICD10_J9622", "ICD10_J9692", ]
#           "ICD10_E662", "ICD9_27803", "any_hypercap_icd", "any_hypercap_icd_hosp", ]
#           "any_hypercap_icd_ed"]
for c in icd_cols:
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].fillna(0).astype(int)

for c in [
    "abg_hypercap_threshold",
    "vbg_hypercap_threshold",
    "other_hypercap_threshold",
    "poc_other_quarantined_hypercap_threshold",

```

```

] :
    if c in cohort_any.columns:
        cohort_any[c] = cohort_any[c].fillna(0).astype(int)

# Final enrollment flag
cohort_any["pco2_threshold_any"] = ((cohort_any["abg_hypercap_threshold"]==1)
    ↵ | (cohort_any["vbg_hypercap_threshold"]==1) |
    ↵ (cohort_any["other_hypercap_threshold"]==1)).astype(int)
cohort_any["enrolled_any"] = ((cohort_any["any_hypercap_icd"]==1) |
    ↵ (cohort_any["pco2_threshold_any"]==1)).astype(int)

print("ICD-only admissions      :",
    ↵ int((cohort_any["any_hypercap_icd"]==1).sum()))
print("Threshold-only admissions :",
    ↵ int(((cohort_any["pco2_threshold_any"]==1) &
    ↵ (cohort_any["any_hypercap_icd"]==0)).sum()))
print("Both ICD and threshold   :",
    ↵ int(((cohort_any["pco2_threshold_any"]==1) &
    ↵ (cohort_any["any_hypercap_icd"]==1)).sum()))
print("Total enrolled (union)   :",
    ↵ int((cohort_any["enrolled_any"]==1).sum()))
if "poc_other_quarantined_hypercap_threshold" in cohort_any.columns:
    print(
        "POC OTHER quarantined hadm count:",
        int(cohort_any["poc_other_quarantined_hypercap_threshold"].sum()),
    )

# New hadm list used for the rest of the notebook
hadm_list = cohort_any.loc[cohort_any["enrolled_any"]==1,
    ↵ "hadm_id"].dropna().astype("int64").tolist()
len(hadm_list)

```

```

ICD-only admissions      : 4237
Threshold-only admissions : 35931
Both ICD and threshold   : 3270
Total enrolled (union)   : 40168
POC OTHER quarantined hadm count: 0

```

40167

2.3.4 4) First ABG and First VBG (LAB + POC, standardized to mmHg)

Rationale: Extract earliest ABG/VBG measurements to characterize baseline physiology with standardized units.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪ consistent across notebook stages.

params = {"hadms": hadm_list}

bg_pairs_sql = rf"""
WITH hadms AS (SELECT hadm_id FROM UNNEST(@hadms) AS hadm_id),

/* ----- LAB (HOSP) ----- */
hosp_cand AS (
    SELECT
        le.subject_id, le.hadm_id, le.charttime, le.specimen_id, le.itemid,
        CAST(le.valuenum AS FLOAT64) AS val,
        LOWER(COALESCE(le.valueuom,'')) AS uom,
        LOWER(di.label) AS lbl,
        LOWER(COALESCE(di.fluid,'')) AS fl,
        LOWER(COALESCE(di.category,'')) AS cat
    FROM `{{PHYS}}.{{HOSP}}.labevents` le
    JOIN `{{PHYS}}.{{HOSP}}.d_labitems` di ON di.itemid = le.itemid
    JOIN hadms h ON h.hadm_id = le.hadm_id
    WHERE le.valuenum IS NOT NULL
        AND le.itemid IN UNNEST(ARRAY_CONCAT({{LAB_PCO2_SQL_ARRAY}},
↪ {{LAB_PH_SQL_ARRAY}}))
        AND LOWER(COALESCE(di.fluid,'')) IN {{LAB_BLOOD_FLUID_SQL_LIST}})
),
hosp_spec AS (
    SELECT le.specimen_id, LOWER(COALESCE(le.value,'')) AS spec_val
    FROM `{{PHYS}}.{{HOSP}}.labevents` le
    WHERE le.specimen_id IS NOT NULL
        AND le.itemid IN UNNEST({{LAB_SPECIMEN_SQL_ARRAY}})
    QUALIFY ROW_NUMBER() OVER (
        PARTITION BY le.specimen_id
        ORDER BY le.charttime DESC
    ) = 1
),
hosp_class AS (
    SELECT
        c.hadm_id, c.charttime, c.specimen_id, c.itemid, c.val, c.uom, c.lbl,
↪ c.fl,
```

```

CASE
    WHEN c.itemid IN UNNEST({LAB_PH_SQL_ARRAY}) THEN 'ph'
    WHEN c.itemid IN UNNEST({LAB_PCO2_SQL_ARRAY}) THEN 'pco2'
    ELSE NULL
END AS analyte,
CASE
    WHEN REGEXP_CONTAINS(s.spec_val, r'arter') OR
        REGEXP_CONTAINS(s.spec_val, r'\bart\b') THEN 'arterial'
    WHEN REGEXP_CONTAINS(s.spec_val, r'ven|mixed|central') THEN 'venous'
    ELSE 'other'
END AS site
FROM hosp_cand c
LEFT JOIN hosp_spec s USING (specimen_id)
),
hosp_pairs AS (
SELECT
    hadm_id, specimen_id,
    MIN(charttime) AS sample_time,
    MAX(IF(analyte='ph', val, NULL)) AS ph,
    MAX(IF(analyte='pco2', val, NULL)) AS pco2_raw,
    (ARRAY_AGG(IF(analyte='pco2', uom, NULL) IGNORE NULLS LIMIT
    1)) [OFFSET(0)] AS pco2_uom,
    (ARRAY_AGG(IF(analyte='ph', uom, NULL) IGNORE NULLS LIMIT
    1)) [OFFSET(0)] AS ph_uom,
    (ARRAY_AGG(site IGNORE NULLS LIMIT 1)) [OFFSET(0)] AS site
FROM hosp_class
GROUP BY hadm_id, specimen_id
HAVING pco2_raw IS NOT NULL AND site IN ('arterial','venous','other')
),
hosp_pairs_std AS (
SELECT
    hadm_id, specimen_id, sample_time, site,
    ph, ph_uom,
    CASE WHEN pco2_uom = 'kpa' THEN pco2_raw * 7.50062 ELSE pco2_raw END AS
    pco2_mmHg,
    'mmhg' AS pco2_uom_norm
FROM hosp_pairs
WHERE (ph IS NULL OR (ph BETWEEN 6.3 AND 7.8))
    AND (pco2_raw IS NULL OR (CASE WHEN pco2_uom='kpa' THEN pco2_raw*7.50062
    ELSE pco2_raw END) BETWEEN 5 AND 200)
),
lab_abg AS (
SELECT hadm_id,

```

```

    ph          AS lab_abg_ph,
    ph_uom     AS lab_abg_ph_uom,
    pco2_mmHg   AS lab_abg_paco2,
    'mmhg'      AS lab_abg_paco2_uom,
    sample_time AS lab_abg_time
  FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
  ↵ sample_time) rn
        FROM hosp_pairs_std WHERE site='arterial') WHERE rn=1
),
lab_vbg AS (
  SELECT hadm_id,
    ph          AS lab_vbg_ph,
    ph_uom     AS lab_vbg_ph_uom,
    pco2_mmHg   AS lab_vbg_paco2,
    'mmhg'      AS lab_vbg_paco2_uom,
    sample_time AS lab_vbg_time
  FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
  ↵ sample_time) rn
        FROM hosp_pairs_std WHERE site='venous') WHERE rn=1
),
lab_other AS (
  SELECT hadm_id,
    ph          AS lab_other_ph,
    ph_uom     AS lab_other_ph_uom,
    pco2_mmHg   AS lab_other_paco2,
    'mmhg'      AS lab_other_paco2_uom,
    sample_time AS lab_other_time
  FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
  ↵ sample_time) rn
        FROM hosp_pairs_std WHERE site='other') WHERE rn=1
),
/* ----- POC (ICU) ----- */
icu_item_map AS (
  SELECT
    itemid,
    LOWER(label) AS lbl,
    CASE
      WHEN itemid IN UNNEST({ICU_PH_SQL_ARRAY}) THEN 'ph'
      WHEN itemid IN UNNEST({ICU_PCO2_SQL_ARRAY}) THEN 'pco2'
      ELSE NULL
    END AS analyte
  FROM `{{PHYS}}.{{ICU}}.d_items`
```

```

    WHERE itemid IN UNNEST(ARRAY_CONCAT({ICU_PH_SQL_ARRAY},
    ↵ {ICU_PCO2_SQL_ARRAY}))
),
icu_spec AS (
    SELECT
        ce.stay_id,
        ce.charttime,
        LOWER(COALESCE(ce.value, '')) AS spec_val
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    WHERE ce.itemid IN UNNEST({ICU_SPECIMEN_SQL_ARRAY})
        AND ce.value IS NOT NULL
),
icu_raw AS (
    SELECT
        ie.hadm_id, ce.stay_id, ce.itemid, ce.charttime,
        di.lbl,
        di.analyte,
        LOWER(REPLACE(COALESCE(ce.valueuom,''), ' ', '')) AS uom_nospace,
        LOWER(COALESCE(ce.value,'')) AS valstr,
        COALESCE(
            CAST(ce.valuenum AS FLOAT64),
            SAFE_CAST(REGEXP_EXTRACT(LOWER(ce.value), r'(-?\d+(?:\.\d+)?))') AS
        ↵ FLOAT64)
        ) AS val
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    JOIN icu_item_map di ON di.itemid = ce.itemid
    JOIN `{{PHYS}}.{{ICU}}.icustays` ie ON ie.stay_id = ce.stay_id
    JOIN hadms h ON h.hadm_id = ie.hadm_id
),
icu_cand_ranked AS (
    SELECT
        r.*,
        s.spec_val,
        ROW_NUMBER() OVER (
            PARTITION BY r.hadm_id, r.stay_id, r.itemid, r.charttime
            ORDER BY ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, SECOND))
        ) AS spec_rank
    FROM icu_raw r
    LEFT JOIN icu_spec s
        ON s.stay_id = r.stay_id
        AND ABS(TIMESTAMP_DIFF(r.charttime, s.charttime, MINUTE)) <= 120
),
icu_cand AS (

```

```

SELECT
    hadm_id, stay_id, itemid, charttime, lbl, analyte, uom_nospace, valstr,
    val,
    CASE
        WHEN REGEXP_CONTAINS(COALESCE(spec_val, ''), r'arter|\\bart\\b') THEN
        'arterial'
        WHEN REGEXP_CONTAINS(COALESCE(spec_val, ''), r'ven|central|mixed') THEN
        'venous'
        ELSE 'other'
    END AS site
FROM icu_cand_ranked
WHERE val IS NOT NULL
    AND analyte IN ('ph', 'pco2')
    AND (spec_rank = 1 OR spec_val IS NULL)
),
icu_ph AS (
    SELECT hadm_id, stay_id, charttime, val AS ph, site AS site_ph
    FROM icu_cand WHERE analyte='ph'
),
icu_co2 AS (
    SELECT hadm_id, stay_id, charttime, val AS pco2_raw, uom_nospace, valstr,
    site AS site_co2
    FROM icu_cand WHERE analyte='pco2'
),
icu_pair_win AS (
    SELECT
        p.hadm_id, p.stay_id,
        COALESCE(p.site_ph, c.site_co2) AS site,
        p.charttime AS ph_time, c.charttime AS co2_time,
        p.ph,
        CASE
            WHEN c.uom_nospace='kpa' THEN 'kpa'
            WHEN c.uom_nospace='mmhg' THEN 'mmhg'
            ELSE c.uom_nospace
        END AS pco2_uom_norm_raw,
        c.pco2_raw,
        ABS(TIMESTAMP_DIFF(c.charttime, p.charttime, SECOND)) AS dt_sec
    FROM icu_ph p
    JOIN icu_co2 c
        ON c.hadm_id = p.hadm_id
        AND c.stay_id = p.stay_id
        AND (COALESCE(p.site_ph, c.site_co2) IN ('arterial','venous','other'))
        AND ABS(TIMESTAMP_DIFF(c.charttime, p.charttime, MINUTE)) <= 10

```

```

QUALIFY ROW_NUMBER() OVER (
    PARTITION BY p.hadm_id, p.stay_id, p.charttime
    ORDER BY dt_sec
) = 1
),
icu_pairs_std AS (
    SELECT
        hadm_id, stay_id, site,
        LEAST(ph_time, co2_time) AS sample_time,
        ph,
        CAST(NULL AS STRING) AS ph_uom, -- POC pH is unitless/null
        CASE WHEN pco2_uom_norm_raw='kpa' THEN pco2_raw*7.50062 ELSE pco2_raw END
        AS pco2_mmHg,
        'mmhg' AS pco2_uom_norm
    FROM icu_pair_win
    WHERE (ph BETWEEN 6.3 AND 7.8 OR ph IS NULL)
        AND (CASE WHEN pco2_uom_norm_raw='kpa' THEN pco2_raw*7.50062 ELSE
        pco2_raw END) BETWEEN 5 AND 200
),
icu_solo_pco2_std AS (
    SELECT
        hadm_id, stay_id, site_co2 AS site,
        charttime AS sample_time,
        CAST(NULL AS FLOAT64) AS ph,
        CAST(NULL AS STRING) AS ph_uom, -- no pH here
        CASE WHEN uom_nospace='kpa' THEN pco2_raw*7.50062 ELSE pco2_raw END AS
        pco2_mmHg,
        'mmhg' AS pco2_uom_norm
    FROM icu_co2
    WHERE site_co2 IN ('arterial','venous','other')
        AND pco2_raw BETWEEN 5 AND 200
),
icu_all AS (
    SELECT * FROM icu_pairs_std
    UNION ALL
    SELECT * FROM icu_solo_pco2_std
),
poc_abg AS (
    SELECT hadm_id,
        ph AS poc_abg_ph,
        ph_uom AS poc_abg_ph_uom,
        pco2_mmHg AS poc_abg_paco2,

```

```

        'mmhg'           AS poc_abg_paco2_uom,
        sample_time      AS poc_abg_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
    ↵  sample_time) rn
          FROM icu_all WHERE site='arterial') WHERE rn=1
),
poc_vbg AS (
    SELECT hadm_id,
        ph           AS poc_vbg_ph,
        ph_uom       AS poc_vbg_ph_uom,
        pco2_mmHg    AS poc_vbg_paco2,
        'mmhg'       AS poc_vbg_paco2_uom,
        sample_time   AS poc_vbg_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
    ↵  sample_time) rn
          FROM icu_all WHERE site='venous') WHERE rn=1
),
poc_other AS (
    SELECT hadm_id,
        ph           AS poc_other_ph,
        ph_uom       AS poc_other_ph_uom,
        pco2_mmHg    AS poc_other_paco2,
        'mmhg'       AS poc_other_paco2_uom,
        sample_time   AS poc_other_time
    FROM (SELECT *, ROW_NUMBER() OVER (PARTITION BY hadm_id ORDER BY
    ↵  sample_time) rn
          FROM icu_all WHERE site='other') WHERE rn=1
)
/* ----- Final one row per hadm ----- */
SELECT
    h.hadm_id,
    -- LAB-ABG / LAB-VBG / LAB-OTHER
    la.lab_abg_ph, la.lab_abg_ph_uom, la.lab_abg_paco2, la.lab_abg_paco2_uom,
    ↵  la.lab_abg_time,
    lv.lab_vbg_ph, lv.lab_vbg_ph_uom, lv.lab_vbg_paco2, lv.lab_vbg_paco2_uom,
    ↵  lv.lab_vbg_time,
    lo.lab_other_ph, lo.lab_other_ph_uom, lo.lab_other_paco2,
    ↵  lo.lab_other_paco2_uom, lo.lab_other_time,
    -- POC-ABG / POC-VBG / POC-OTHER
    pa.poc_abg_ph, pa.poc_abg_ph_uom, pa.poc_abg_paco2, pa.poc_abg_paco2_uom,
    ↵  pa.poc_abg_time,

```

```

pv.poc_vbg_ph, pv.poc_vbg_ph_uom, pv.poc_vbg_paco2, pv.poc_vbg_paco2_uom,
↳ pv.poc_vbg_time,
po.poc_other_ph, po.poc_other_ph_uom, po.poc_other_paco2,
↳ po.poc_other_paco2_uom, po.poc_other_time,
-- First ABG across LAB+POC
(SELECT AS STRUCT src, t, ph, pco2
  FROM (SELECT 'LAB' AS src, la.lab_abg_time AS t, la.lab_abg_ph AS ph,
↳ la.lab_abg_paco2 AS pco2
        UNION ALL
        SELECT 'POC', pa.poc_abg_time, pa.poc_abg_ph, pa.poc_abg_paco2)
 WHERE t IS NOT NULL
 ORDER BY t LIMIT 1) AS first_abg,
-- First VBG across LAB+POC
(SELECT AS STRUCT src, t, ph, pco2
  FROM (SELECT 'LAB' AS src, lv.lab_vbg_time AS t, lv.lab_vbg_ph AS ph,
↳ lv.lab_vbg_paco2 AS pco2
        UNION ALL
        SELECT 'POC', pv.poc_vbg_time, pv.poc_vbg_ph, pv.poc_vbg_paco2)
 WHERE t IS NOT NULL
 ORDER BY t LIMIT 1) AS first_vbg,
-- First OTHER-source pCO2 (LAB blood-gas unknown specimen only; POC OTHER
↳ is quarantined)
(SELECT AS STRUCT src, t, ph, pco2
  FROM (SELECT 'LAB_BG_UNKNOWN' AS src, lo.lab_other_time AS t,
↳ lo.lab_other_ph AS ph, lo.lab_other_paco2 AS pco2)
 WHERE t IS NOT NULL
 ORDER BY t LIMIT 1) AS first_other
FROM hadms h
LEFT JOIN lab_abg la USING (hadm_id)
LEFT JOIN lab_vbg lv USING (hadm_id)
LEFT JOIN lab_other lo USING (hadm_id)
LEFT JOIN poc_abg pa USING (hadm_id)
LEFT JOIN poc_vbg pv USING (hadm_id)
LEFT JOIN poc_other po USING (hadm_id)
"""

bg_pairs = run_sql_bq(bg_pairs_sql, params)

# Flatten STRUCTs for first_abg, first_vbg, and first_other
for col in ["first_abg", "first_vbg", "first_other"]:
    if col in bg_pairs.columns:
        bg_pairs[f"{col}_src"] = bg_pairs[col].apply(lambda x: x.get("src"))
    ↳ if isinstance(x, dict) else None)

```

```

        bg_pairs[f"{col}_time"] = bg_pairs[col].apply(lambda x: x.get("t"))
    ↵ if isinstance(x, dict) else None)
        bg_pairs[f"{col}_ph"] = bg_pairs[col].apply(lambda x: x.get("ph"))
    ↵ if isinstance(x, dict) else None)
        bg_pairs[f"{col}_pco2"] = bg_pairs[col].apply(lambda x:
    ↵ x.get("pco2")if isinstance(x, dict) else None)
        bg_pairs = bg_pairs.drop(columns=[col])

if {"first_other_src", "poc_other_paco2"}.issubset(bg_pairs.columns):
    first_other_src_upper =
        bg_pairs["first_other_src"].astype("string").str.upper().fillna("")
    )
    bg_pairs["first_other_src_detail"] = np.select(
        [
            first_other_src_upper.eq("LAB_BG_UNKNOWN"),
            pd.to_numeric(bg_pairs["poc_other_paco2"]),
    ↵ errors="coerce").notna(),
        ],
        [
            "lab_bg_unknown",
            "poc_quarantined",
        ],
        default="missing",
    )
else:
    bg_pairs["first_other_src_detail"] = "missing"

bg_pairs.head(3)

```

	hadm_id	lab_abg_ph	lab_abg_ph_uom	lab_abg_paco2	lab_abg_paco2_uom	lab_abg_time
0	20000147	7.41	units	35.0	mmhg	2121-08-30 17:38
1	20001305	NaN	None	NaN	None	NaT
2	20001494	7.39	units	45.0	mmhg	2125-10-26 11:27

2.3.5 5) Demographics & outcomes

Rationale: Build baseline covariates used for descriptive statistics and potential confounding adjustment.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪ consistent across notebook stages.
```

```

SQL["demo_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
    a.hadm_id,
    a.subject_id,
    a.admittime,
    a.dischtime,
    a.deathtime,
    a.admission_type,
    a.admission_location,
    a.discharge_location,
    a.insurance,
    -- LOS (days)
    TIMESTAMP_DIFF(a.dischtime, a.admittime, HOUR) / 24.0 AS hosp_los_days,
    -- in-hospital death
    IF(a.deathtime IS NOT NULL, 1, 0) AS death_in_hosp,
    -- demographics
    p.gender,
    SAFE_CAST(ROUND(p.anchor_age + (EXTRACT(YEAR FROM a.admittime) -
    ↪ p.anchor_year), 1) AS FLOAT64) AS age_at_admit,
    -- 30-day all-cause mortality from admission
    IF(p.dod IS NOT NULL AND DATE_DIFF(DATE(p.dod), DATE(a.admittime), DAY)
    ↪ BETWEEN 0 AND 30, 1, 0) AS death_30d
FROM `{{PHYS}}.{{HOSP}}.admissions` a
JOIN hadms h USING (hadm_id)
JOIN `{{PHYS}}.{{HOSP}}.patients` p USING (subject_id)
"""
demo = run_sql_bq(sql("demo_sql"), {"hadms": hadm_list})
print("Demo rows:", len(demo))
demo.head(3)

```

Demo rows: 40167

	hadm_id	subject_id	admittime	dischtime	deathtime	admission_type
0	23485217	10584718	2165-02-12 15:41:00	2165-03-06 08:20:00	2165-03-06 08:20:00	EW EMER.
1	27004173	10297948	2135-08-23 18:46:00	2135-09-08 15:50:00	NaT	URGENT
2	22661627	10032409	2130-01-12 18:42:00	2130-01-21 14:32:00	NaT	EW EMER.

```

# Purpose: Create reusable data-quality and merge guardrail helpers to
↪ prevent silent join errors.

```

```

# ===== Drop-in: safe merge utilities (one cell, run once) =====
import pandas as pd
from typing import Iterable, Optional, Literal

def _ensure_Int64(s: pd.Series) -> pd.Series:
    """Coerce to pandas nullable Int64 (preserves NA)."""
    return pd.to_numeric(s, errors="coerce").astype("Int64")

def strip_subject_cols(fr: pd.DataFrame) -> pd.DataFrame:
    """Remove any subject_id-like columns from a frame (e.g., 'subject_id',
    'Subject_ID')."""
    return fr.drop(columns=[c for c in fr.columns if
                           c.lower().startswith("subject_id")],
                  errors="ignore")

def safe_merge_on_hadm(
    left: pd.DataFrame,
    right: pd.DataFrame,
    *,
    right_name: str,
    take: Optional[Iterable[str]] = None,
    order_by: Optional[Iterable[str]] = None,
    check_subject: Literal[False, "warn", "raise"] = False,
) -> pd.DataFrame:
    """
    Left-merge 'right' into 'left' on hadm_id, returning a copy of left with
    right's columns.
    - Dedupes right on hadm_id (optionally using order_by to pick the first
    row).
    - Optionally restricts right columns via `take`.
    - Optionally audits subject_id agreement before dropping subject_id from
    right.
    - Always strips subject_id-like columns from the right to prevent *_x/_y
    suffixes.
    - Raises if any *_x/_y suffixes still appear (indicates overlapping names
    besides hadm_id).
    """
    if "hadm_id" not in left.columns:
        raise KeyError(f"left frame lacks hadm_id before merging\n"
                      f"  {right_name}")
    if "hadm_id" not in right.columns:
        raise KeyError(f"{right_name} lacks hadm_id")

```

```

L = left.copy()
R = right.copy()

# Standardize dtypes of keys
L["hadm_id"] = _ensure_Int64(L["hadm_id"])
R["hadm_id"] = _ensure_Int64(R["hadm_id"])
if "subject_id" in L.columns:
    L["subject_id"] = _ensure_Int64(L["subject_id"])
if "subject_id" in R.columns:
    R["subject_id"] = _ensure_Int64(R["subject_id"])

# Dedupe RIGHT by hadm_id (optionally order_by first)
if order_by:
    R = (R.sort_values(list(order_by))
        .drop_duplicates(subset=["hadm_id"], keep="first"))
else:
    R = R.drop_duplicates(subset=["hadm_id"], keep="first")

# Optional subject_id consistency audit (before stripping)
if check_subject and ("subject_id" in L.columns) and ("subject_id" in
    R.columns):
    # Join only on hadm_id where both sides have subject_id
    tmp = (L[["hadm_id", "subject_id"]]
        .merge(R[["hadm_id", "subject_id"]],
            on="hadm_id", how="inner", suffixes=("_L", "_R")))
    mism = (tmp["subject_id_L"].notna() & tmp["subject_id_R"].notna() &
        (tmp["subject_id_L"] != tmp["subject_id_R"]))
    n_mism = int(mism.sum())
    if n_mism > 0:
        sample_ids = tmp.loc[mism, "hadm_id"].head(10).tolist()
        msg = (f"[{right_name}] subject_id mismatch on {n_mism}"
    ↵  hadm_id(s). "
                f"Examples: {sample_ids}")
        if check_subject == "raise":
            raise ValueError(msg)
        else:
            print("WARNING:", msg)

# Limit right columns (avoid accidental overlaps)
if take is not None:
    keep = ["hadm_id"] + [c for c in take if c != "hadm_id"]
    R = R[[c for c in keep if c in R.columns]]

```

```

# Always strip subject_id-like columns from right to prevent *_x/_y
R = strip_subject_cols(R)

# Final merge
out = L.merge(R, on="hadm_id", how="left", suffixes=("",""))

# Guard: no suffixes should be present
bad = [c for c in out.columns if c.endswith("_x") or c.endswith("_y")]
if bad:
    raise RuntimeError(
        f"Merge with {right_name} produced suffixed columns {bad}. "
        "You likely have overlapping column names other than hadm_id."
    )
return out

print("Safe merge helpers loaded.")

```

Safe merge helpers loaded.

2.3.6 6) NIH/OMB race & ethnicity (ED + Hospital)

Rationale: Harmonize race/ethnicity across sources using NIH/OMB categories for consistent reporting.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
#           consistent across notebook stages.

race_eth_sql = rf"""
WITH hadms AS (
    SELECT x AS hadm_id
    FROM UNNEST(@hadms) AS x
),
-- Hospital admission "race" text
hosp AS (
    SELECT a.hadm_id, LOWER(TRIM(a.race)) AS race_hosp_raw
    FROM `{{PHYS}}.{{HOSP}}.admissions` a
    JOIN hadms hm USING (hadm_id)
),
-- Earliest ED stay leading to the admission; take its "race" text if present
ed_first AS (

```

```

SELECT
    e.hadm_id,
    (ARRAY_AGG(STRUCT(e.intime AS intime, LOWER(TRIM(e.race)) AS race_ed_raw)
        ORDER BY e.intime ASC LIMIT 1)) [OFFSET(0)] AS pick
FROM `{{PHYS}}.{{ED}}.edstays` e
JOIN hadms hm USING (hadm_id)
GROUP BY e.hadm_id
),
ed AS (
    SELECT hadm_id, pick.race_ed_raw
    FROM ed_first
),
-- Combine ED + Hospital for maximum coverage
comb AS (
    SELECT
        hm.hadm_id,
        ho.race_hosp_raw,
        ed.race_ed_raw,
        TRIM(REGEXP_REPLACE(CONCAT(COALESCE(ho.race_hosp_raw,''), ' ',
        COALESCE(ed.race_ed_raw,'')), r'\s+', ' ')) AS race_text_any
    FROM hadms hm
    LEFT JOIN hosp ho USING (hadm_id)
    LEFT JOIN ed ed USING (hadm_id)
),
-- Tokenization to OMB families + Hispanic ethnicity
tok AS (
    SELECT
        hadm_id, race_hosp_raw, race_ed_raw, race_text_any,
        -- Ethnicity (Hispanic)
        REGEXP_CONTAINS(race_text_any, r'\b(hispanic|latinx|latino|latina)\b') AS is_hisp,
        -- Race families (use boundaries to reduce false positives)
        REGEXP_CONTAINS(race_text_any, r'americans+indian|\balaska\b') AS is_aian,
        REGEXP_CONTAINS(race_text_any, r'\basian\b') AS is_asian,
        REGEXP_CONTAINS(race_text_any, r'\b(black|african\s+american)\b') AS is_black,
        REGEXP_CONTAINS(race_text_any, r'hawaiian|pacific\s+islander') AS is_nhopi,

```

```

    REGEXP_CONTAINS(race_text_any, r'\bwhite\b|caucasian') AS is_white,
    -- Unknown/other indicators
    REGEXP_CONTAINS(race_text_any,
    ↵ r'unknown|other|declined|unable|not\s+reported|missing|null') AS
    ↵ is_unknown_any,
    -- Multi-race hints
    REGEXP_CONTAINS(race_text_any,
    ↵ r'(two|2)\s+or\s+more|multi|biracial|multiracial') AS is_multi_hint
    FROM comb
),
-- Decide ethnicity per NIH
ethn AS (
    SELECT
        hadm_id, race_hosp_raw, race_ed_raw, race_text_any,
        CASE
            WHEN is_hisp THEN 'Hispanic or Latino'
            WHEN (race_text_any IS NULL OR race_text_any = '' OR is_unknown_any)
            ↵ THEN 'Unknown or Not Reported'
            ELSE 'Not Hispanic or Latino'
        END AS nih_ethnicity,
        (CAST(is_aian AS INT64) + CAST(is_asian AS INT64) + CAST(is_black AS
        ↵ INT64)
        + CAST(is_nhopi AS INT64) + CAST(is_white AS INT64)) AS race_hits,
        is_aian, is_asian, is_black, is_nhopi, is_white, is_multi_hint,
        ↵ is_unknown_any
    FROM tok
),
-- Decide race per NIH/OMB (1997)
race_assign AS (
    SELECT
        hadm_id, race_hosp_raw, race_ed_raw, race_text_any, nih_ethnicity,
        CASE
            WHEN race_hits >= 2 OR is_multi_hint THEN 'More than one race'
            WHEN is_aian THEN 'American Indian or Alaska Native'
            WHEN is_asian THEN 'Asian'
            WHEN is_black THEN 'Black or African American'
            WHEN is_nhopi THEN 'Native Hawaiian or Other Pacific Islander'
            WHEN is_white THEN 'White'

```

```

        WHEN is_unknown_any OR race_text_any IS NULL OR race_text_any = '' THEN
    ↵  'Unknown or Not Reported'
        ELSE 'Unknown or Not Reported'
    END AS nih_race
    FROM ethn
)
SELECT hadm_id, race_hosp_raw, race_ed_raw, nih_race, nih_ethnicity
FROM race_assign
"""
race_eth = run_sql_bq(race_eth_sql, {"hadms": hadm_list})
print("Race/Eth rows:", len(race_eth))
race_eth.head(3)

```

Race/Eth rows: 40167

	hadm_id	race_hosp_raw	race_ed_raw	nih_race	nih_ethnicity
0	20000147	white - other european	None	White	Unknown or Not Reported
1	20001305	white	white	White	Not Hispanic or Latino
2	20001494	white	None	White	Not Hispanic or Latino

2.3.7 7) ED triage (linked to hadm) and first ED vitals

Rationale: Capture ED presentation features (vitals and chief complaint) for symptom and severity analyses.

```
# Purpose: Pull ED triage/first-vitals by loading ED tables once and
↪ filtering to cohort hadm_ids in pandas.
```

```

hadms_for_ed = set(
    pd.Series(hadm_list)
    .dropna()
    .astype("int64")
    .tolist()
)

# 1) ED stay map (full table), then restrict to cohort hadm_ids locally.
SQL["edmap_all_sql"] = f"""
SELECT stay_id, hadm_id, intime
FROM `{{PHYS}}.{{ED}}.edstays`
WHERE hadm_id IS NOT NULL

```

```

"""
edmap_all = run_sql_bq(sql("edmap_all_sql"))
edmap_all["hadm_id"] = pd.to_numeric(edmap_all["hadm_id"],
    ↪ errors="coerce").astype("Int64")
edmap_all["stay_id"] = pd.to_numeric(edmap_all["stay_id"],
    ↪ errors="coerce").astype("Int64")
edmap_all["intime"] = pd.to_datetime(edmap_all["intime"], errors="coerce")

edmap = edmap_all[edmap_all["hadm_id"].isin(hadms_for_ed)].copy()
edmap = edmap.drop_duplicates(subset=["stay_id", "hadm_id"])
print("ED map rows (cohort):", len(edmap))

if edmap.empty:
    ed_triage = pd.DataFrame(columns=[
        "hadm_id", "ed_triage_temp", "ed_triage_hr", "ed_triage_rr",
        "ed_triage_o2sat", "ed_triage_sbp", "ed_triage_dbp",
        ↪ "ed_triage_pain",
        "ed_triage_acuity", "ed_triage_cc",
    ])
    ed_first = pd.DataFrame(columns=[
        "hadm_id", "ed_first_vitals_time", "ed_first_temp", "ed_first_hr",
        "ed_first_rr", "ed_first_o2sat", "ed_first_sbp", "ed_first_dbp",
        "ed_first_rhythm", "ed_first_pain",
    ])
else:
    # 2) ED triage (full table), then reduce to earliest ED stay per hadm.
    SQL["ed_triage_all_sql"] = f"""
    SELECT
        stay_id,
        temperature      AS ed_triage_temp,
        heartrate       AS ed_triage_hr,
        resprate        AS ed_triage_rr,
        o2sat           AS ed_triage_o2sat,
        sbp              AS ed_triage_sbp,
        dbp              AS ed_triage_dbp,
        pain             AS ed_triage_pain,
        acuity           AS ed_triage_acuity,
        chiefcomplaint AS ed_triage_cc
    FROM `{{PHYS}}.{{ED}}.triage`
"""

    tri_all = run_sql_bq(sql("ed_triage_all_sql"))
    tri_all["stay_id"] = pd.to_numeric(tri_all["stay_id"],
    ↪ errors="coerce").astype("Int64")

```

```

tri_merged = (
    edmap[["stay_id", "hadm_id", "intime"]]
    .merge(tri_all, on="stay_id", how="left")
)
ed_triage = (
    tri_merged
    .sort_values(["hadm_id", "intime"], na_position="last")
    .drop_duplicates(subset=["hadm_id"], keep="first")
    .drop(columns=["stay_id", "intime"])
    .reset_index(drop=True)
)
print("ED triage rows:", len(ed_triage))

# 3) First vitals per stay across full ED vitals table, then reduce to
#     ↵ earliest vitals-time per hadm.
SQL["ed_first_vitals_all_sql"] = f"""
WITH vs_ranked AS (
    SELECT
        v.stay_id,
        v.charttime,
        v.temperature,
        v.heartrate,
        v.resprate,
        v.o2sat,
        v.sbp,
        v.dbp,
        v.rhythm,
        v.pain,
        ROW_NUMBER() OVER (PARTITION BY v.stay_id ORDER BY v.charttime) AS rn
    FROM `{{PHYS}}.{{ED}}.vitalsign` v
)
SELECT
    stay_id,
    charttime AS ed_first_vitals_time,
    temperature AS ed_first_temp,
    heartrate AS ed_first_hr,
    resprate AS ed_first_rr,
    o2sat AS ed_first_o2sat,
    sbp AS ed_first_sbp,
    dbp AS ed_first_dbp,
    rhythm AS ed_first_rhythm,
    pain AS ed_first_pain

```

```

FROM vs_ranked
WHERE rn = 1
"""
first_stay_all = run_sql_bq(sql("ed_first_vitals_all_sql"))
first_stay_all["stay_id"] = pd.to_numeric(first_stay_all["stay_id"], errors="coerce").astype("Int64")
first_stay_all["ed_first_vitals_time"] =
    pd.to_datetime(first_stay_all["ed_first_vitals_time"], errors="coerce")

first_merged = (
    edmap[["stay_id", "hadm_id"]]
    .merge(first_stay_all, on="stay_id", how="left")
)
ed_first = (
    first_merged
    .sort_values(["hadm_id", "ed_first_vitals_time"], na_position="last")
    .drop_duplicates(subset=["hadm_id"], keep="first")
    .drop(columns=["stay_id"])
    .reset_index(drop=True)
)
print("ED first vitals rows:", len(ed_first))

```

```

ED map rows (cohort): 11786
ED triage rows: 11769
ED first vitals rows: 11769

```

2.3.8 8) ICU meta (first ICU stay, LOS days)

Rationale: Summarize ICU exposure and length of stay to contextualize disease severity.

```
# Purpose: Build first-ICU metadata via full-table pull + local hadm filter
    ↳ to avoid large ARRAY parameter stalls.
```

```

SQL["icu_all_sql"] = f"""
SELECT
    hadm_id,
    stay_id AS first_icu_stay_id,
    intime AS icu_intime,
    outtime AS icu_outtime
FROM `{{PHYS}}.{{ICU}}.icustays`
WHERE hadm_id IS NOT NULL
"""

```

```

icu_all = run_sql_bq(sql("icu_all_sql"))
icu_all["hadm_id"] = pd.to_numeric(icu_all["hadm_id"],
    ↪ errors="coerce").astype("Int64")
icu_all["first_icu_stay_id"] = pd.to_numeric(icu_all["first_icu_stay_id"],
    ↪ errors="coerce").astype("Int64")
icu_all["icu_intime"] = pd.to_datetime(icu_all["icu_intime"],
    ↪ errors="coerce")
icu_all["icu_outtime"] = pd.to_datetime(icu_all["icu_outtime"],
    ↪ errors="coerce")

hadms_for_icu = set(pd.Series(hadm_list).dropna().astype("int64").tolist())
icu_all = icu_all[icu_all["hadm_id"].isin(hadms_for_icu)].copy()

icu_meta = (
    icu_all
        .sort_values(["hadm_id", "icu_intime", "first_icu_stay_id"],
    ↪ na_position="last")
        .drop_duplicates(subset=["hadm_id"], keep="first")
        .reset_index(drop=True)
)
icu_meta["icu_los_days"] = (
    (icu_meta["icu_outtime"] - icu_meta["icu_intime"]).dt.total_seconds() /
    ↪ 86400.0
)

print("ICU meta rows:", len(icu_meta))

```

ICU meta rows: 30571

2.3.9 9) Ventilation flags (ICD procedures)

Rationale: Identify IMV/NIV exposure as clinically relevant respiratory support indicators.

```
# Purpose: Build IMV/NIV ICD flags from targeted procedure codes, then filter
    ↪ to cohort hadm_ids locally.
```

```
SQL["vent_proc_sql"] = f"""
SELECT
    hadm_id,
    icd_version,
    REPLACE(icd_code, '.', '') AS code_norm
FROM `{{PHYS}}.{{HOSP}}.procedures_icd`
```

```

WHERE
  (icd_version = 10 AND icd_code IN
  ↪ ('5A1935Z','5A1945Z','5A1955Z','0BH17EZ','0BH18EZ','5A09357','5A09457','5A09557'))
  OR
  (icd_version = 9 AND REPLACE(icd_code, '.', '') IN
  ↪ ('9670','9671','9672','9604','9390','9391','9399'))
"""

vent_proc = run_sql_bq(sql("vent_proc_sql"))
vent_proc["hadm_id"] = pd.to_numeric(vent_proc["hadm_id"],
  ↪ errors="coerce").astype("Int64")
vent_proc["icd_version"] = pd.to_numeric(vent_proc["icd_version"],
  ↪ errors="coerce").astype("Int64")
vent_proc["code_norm"] = vent_proc["code_norm"].astype(str)

hadms_for_vent = set(pd.Series(hadm_list).dropna().astype("int64").tolist())
vent_proc = vent_proc[vent_proc["hadm_id"].isin(hadms_for_vent)].copy()

imv_codes_10 = {"5A1935Z", "5A1945Z", "5A1955Z", "0BH17EZ", "0BH18EZ"}
imv_codes_9 = {"9670", "9671", "9672", "9604"}
niv_codes_10 = {"5A09357", "5A09457", "5A09557"}
niv_codes_9 = {"9390", "9391", "9399"}

vent_proc["imv_hit"] = (
  ((vent_proc["icd_version"] == 10) &
  ↪ vent_proc["code_norm"].isin(imv_codes_10)) |
  ((vent_proc["icd_version"] == 9) &
  ↪ vent_proc["code_norm"].isin(imv_codes_9))
)
vent_proc["niv_hit"] = (
  ((vent_proc["icd_version"] == 10) &
  ↪ vent_proc["code_norm"].isin(niv_codes_10)) |
  ((vent_proc["icd_version"] == 9) &
  ↪ vent_proc["code_norm"].isin(niv_codes_9))
)

vent = (
  vent_proc
  .groupby("hadm_id", as_index=False)
  .agg(
    imv_flag=("imv_hit", "max"),
    niv_flag=("niv_hit", "max"),
  )
)

```

```

)
vent["imv_flag"] = vent["imv_flag"].astype(int)
vent["niv_flag"] = vent["niv_flag"].astype(int)
vent["any_vent_flag"] = ((vent["imv_flag"] == 1) | (vent["niv_flag"] ==
    ↵ 1)).astype(int)

# Preserve one row per hadm in hadm_list even if no vent procedure codes were
    ↵ found.
vent = pd.DataFrame({"hadm_id": pd.Series(hadm_list,
    ↵ dtype="Int64")}).drop_duplicates().merge(vent, on="hadm_id", how="left")
vent[["imv_flag", "niv_flag", "any_vent_flag"]] = vent[["imv_flag",
    ↵ "niv_flag", "any_vent_flag"]].fillna(0).astype(int)

print("Vent rows:", len(vent))

```

Vent rows: 40167

and from charts

```

# Purpose: Extract earliest NIV/IMV charttimes using prefiltered ventilation
    ↵ itemids for better performance.

SQL["vent_chart_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),

stays AS (
    SELECT DISTINCT hadm_id, stay_id
    FROM `{{PHYS}}.{{ICU}}.icustays`
    WHERE hadm_id IN (SELECT hadm_id FROM hadms)
),

vent_itemids AS (
    SELECT itemid, LOWER(label) AS lbl
    FROM `{{PHYS}}.{{ICU}}.d_items`
    WHERE REGEXP_CONTAINS(LOWER(label), r'(vent|ventilator|mode|bipap|bi[-
        ↵ ]?pap|cpap|nippv|niv|mask|ett|endotracheal)')
),

cand AS (
    SELECT
        s.hadm_id,
        ce.charttime,
        vi.lbl,

```

```

        LOWER(COALESCE(ce.value, '')) AS valstr
    FROM `.{PHYS}.{ICU}.chartevents` ce
    JOIN stays s ON s.stay_id = ce.stay_id
    JOIN vent_itemids vi ON vi.itemid = ce.itemid
),
flags AS (
SELECT
    hadm_id,
    MIN(IF(
        REGEXP_CONTAINS(lbl, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
        ]?pap|cpap)')
        OR REGEXP_CONTAINS(valstr, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
        ]?pap|cpap)'),
        charttime, NULL)) AS first_niv_time,
    MIN(IF(
        REGEXP_CONTAINS(lbl, r'(invasive
        ventilation|endotracheal|ett|mech(|anical)? vent|ventilator
        mode|ac/|simv|prvc|aprvc|pcv|vcv|assist\$\*control)')
        OR REGEXP_CONTAINS(valstr, r'(invasive
        ventilation|endotracheal|ett|ac/|simv|prvc|aprvc|pcv|vcv|assist\$\*
        s\*control)'),
        charttime, NULL)) AS first_imv_time,
    MAX(CASE
        WHEN REGEXP_CONTAINS(lbl, r'(non[- ]?invasive|niv|nippv|bipap|bi[-
        ]?pap|cpap)')
            OR REGEXP_CONTAINS(valstr, r'(non[-
            ]?invasive|niv|nippv|bipap|bi[- ]?pap|cpap)')
            THEN 1 ELSE 0 END) AS niv_chart_flag,
    MAX(CASE
        WHEN REGEXP_CONTAINS(lbl, r'(invasive
        ventilation|endotracheal|ett|mech(|anical)? vent|ventilator
        mode|ac/|simv|prvc|aprvc|pcv|vcv|assist\$\*control)')
            OR REGEXP_CONTAINS(valstr, r'(invasive
            ventilation|endotracheal|ett|ac/|simv|prvc|aprvc|pcv|vcv|assist\$\*
            s\*control)')
            THEN 1 ELSE 0 END) AS imv_chart_flag
    FROM cand
    GROUP BY hadm_id
)
SELECT hadm_id, niv_chart_flag, imv_chart_flag, first_niv_time,
    first_imv_time

```

```

FROM flags
"""

try:
    vent_chart = run_sql_bq(sql("vent_chart_sql"), {"hadms": hadm_list})
except Exception as e:
    print("WARNING: vent chart extraction failed; falling back to ICD-only
          vent timing.", e)
    vent_chart = pd.DataFrame(columns=["hadm_id", "niv_chart_flag",
                                      "imv_chart_flag", "first_niv_time", "first_imv_time"])

# Purpose: Derive respiratory support flags and timing fields for IMV/NIV
# exposure.

# If your existing ICD-only result is called `vent`, rename for clarity:
vent_proc = vent.copy()

# Outer merge so we keep hadm_ids that appear in only one source
vent_combined = vent_proc.merge(vent_chart, on="hadm_id", how="outer")

# Fill missing with 0 before taking maxima
for c in ["imv_flag", "niv_flag", "any_vent_flag", "imv_chart_flag", ]
    ↵ "niv_chart_flag"]:
    if c in vent_combined.columns:
        vent_combined[c] = vent_combined[c].fillna(0).astype("Int64")

# Final "any-source" flags
vent_combined["imv_flag"] =
    ↵ vent_combined[["imv_flag", "imv_chart_flag"]].max(axis=1).astype("Int64")
vent_combined["niv_flag"] =
    ↵ vent_combined[["niv_flag", "niv_chart_flag"]].max(axis=1).astype("Int64")
vent_combined["any_vent_flag"] =
    ↵ vent_combined[["imv_flag", "niv_flag"]].max(axis=1).astype("Int64")

vent_combined = vent_combined[["hadm_id", "imv_flag", "niv_flag", ]
    ↵ "any_vent_flag", "first_imv_time", "first_niv_time"]]
print("After combining ICD + chart signals:",
      "\nIMV=1:", int((vent_combined["imv_flag"]==1).sum()),
      "\nNIV=1:", int((vent_combined["niv_flag"]==1).sum()))

After combining ICD + chart signals:
IMV=1: 24745
NIV=1: 16574

```

2.3.10 10) Assemble final DataFrame

Rationale: Merge all derived features into a single analytic table keyed by hadm_id.

```
# Purpose: Extract and standardize first ABG/VBG physiology fields for
#           baseline characterization.

# Canonical base (carries authoritative subject_id)
df = demo.copy()

# Cohort flags / thresholds / labs / etc.
df = safe_merge_on_hadm(df, cohort_any, right_name="cohort_any",
                        ↵ check_subject="warn")
df = safe_merge_on_hadm(df, bg_pairs,    right_name="bg_pairs")
df = safe_merge_on_hadm(df, race_eth,    right_name="race_eth")
df = safe_merge_on_hadm(df, ed_triage,   right_name="ed_triage")
df = safe_merge_on_hadm(df, ed_first,    right_name="ed_first")
df = safe_merge_on_hadm(df, icu_meta,    right_name="icu_meta")
df = safe_merge_on_hadm(df, vent_combined, right_name="vent_combined")

# Anchor to first ED presentation (per admission)
if "ed_intime_first" in globals():
    df = safe_merge_on_hadm(df, ed_intime_first,
                           ↵ right_name="ed_intime_first")

# Derived timing: first NIV/IMV relative to ED presentation
if "ed_intime_first" in df.columns and "first_imv_time" in df.columns:
    df["dt_first_imv_hours"] = (df["first_imv_time"] -
                                ↵ df["ed_intime_first"]).dt.total_seconds() / 3600.0
if "ed_intime_first" in df.columns and "first_niv_time" in df.columns:
    df["dt_first_niv_hours"] = (df["first_niv_time"] -
                                ↵ df["ed_intime_first"]).dt.total_seconds() / 3600.0

# ABG/VBG before IMV (hadm-level)
if {"first_abg_time", "first_imv_time"}.issubset(df.columns):
    df["abg_before_imv"] = (
        df["first_abg_time"].notna() & df["first_imv_time"].notna() &
        (df["first_abg_time"] < df["first_imv_time"]))
    ).astype("Int64")
if {"first_vbg_time", "first_imv_time"}.issubset(df.columns):
    df["vbg_before_imv"] = (
        df["first_vbg_time"].notna() & df["first_imv_time"].notna() &
        (df["first_vbg_time"] < df["first_imv_time"]))
```

```

).astype("Int64")

# Canonical hadm-level threshold normalization.
# This consolidates explicit threshold columns with parsed ABG/VBG/other pCO2
# fields
# so later merges cannot silently zero-out source-specific thresholds.
def _num_series(frame: pd.DataFrame, col: str) -> pd.Series:
    if col in frame.columns:
        return pd.to_numeric(frame[col], errors="coerce")
    return pd.Series(np.nan, index=frame.index)

def _any_ge(frame: pd.DataFrame, cols: list[str], threshold: float) ->
    pd.Series:
    avail = [c for c in cols if c in frame.columns]
    if not avail:
        return pd.Series(0, index=frame.index, dtype="Int64")
    mat = pd.concat([_num_series(frame, c) for c in avail], axis=1)
    return mat.ge(threshold).any(axis=1).astype("Int64")

existing_abg = _num_series(df,
    "abg_hypercap_threshold").fillna(0).astype(int)
existing_vbg = _num_series(df,
    "vbg_hypercap_threshold").fillna(0).astype(int)
existing_other = _num_series(df,
    "other_hypercap_threshold").fillna(0).astype(int)

abg_from_values = _any_ge(df, ["lab_abg_paco2", "poc_abg_paco2",
    "first_abg_pco2"], 45.0)
vbg_from_values = _any_ge(df, ["lab_vbg_paco2", "poc_vbg_paco2",
    "first_vbg_pco2"], 50.0)
other_from_values = _any_ge(df, ["first_other_pco2"], 50.0)
if "first_other_src_detail" in df.columns:
    other_detail = df["first_other_src_detail"].astype("string").str.lower()
    other_from_values = ((other_from_values == 1) &
    other_detail.eq("lab_bg_unknown")).astype(
        "Int64"
    )

df["abg_hypercap_threshold"] = ((existing_abg == 1) | (abg_from_values ==
    1)).astype("Int64")

```

```

df["vbg_hypercap_threshold"] = ((existing_vbg == 1) | (vbg_from_values ==
↪ 1)).astype("Int64")
df["other_hypercap_threshold"] = ((existing_other == 1) | (other_from_values
↪ == 1)).astype("Int64")

df["pco2_threshold_any"] = (
    (df["abg_hypercap_threshold"] == 1)
    | (df["vbg_hypercap_threshold"] == 1)
    | (df["other_hypercap_threshold"] == 1)
).astype("Int64")

if "any_hypercap_icd" in df.columns:
    df["enrolled_any"] = ((pd.to_numeric(df["any_hypercap_icd"]),
↪ errors="coerce").fillna(0).astype(int) == 1) | (df["pco2_threshold_any"]
↪ == 1)).astype("Int64")

print(
    "Threshold sanity (hadm-level):",
    "ABG=", int(df["abg_hypercap_threshold"].fillna(0).sum()),
    "VBG=", int(df["vbg_hypercap_threshold"].fillna(0).sum()),
    "OTHER=", int(df["other_hypercap_threshold"].fillna(0).sum()),
    "ANY=", int(df["pco2_threshold_any"].fillna(0).sum()),
)
if "poc_other_quarantined_hypercap_threshold" in df.columns:
    poc_other_quarantine_n = int(
        pd.to_numeric(df["poc_other_quarantined_hypercap_threshold"]),
↪ errors="coerce")
    .fillna(0)
    .astype(int)
    .sum()
)
print("POC OTHER quarantined hadm count:", poc_other_quarantine_n)
poc_other_threshold_leak_n = int(
(
    pd.to_numeric(df["poc_other_quarantined_hypercap_threshold"]),
↪ errors="coerce")
    .fillna(0)
    .astype(int)
    .eq(1)
    & pd.to_numeric(df["other_hypercap_threshold"], errors="coerce")
    .fillna(0)
    .astype(int)
    .eq(1)
)

```

```

        ).sum()
    )
    if poc_other_threshold_leak_n > 0:
        raise ValueError(
            "Detected quarantined POC-other leakage into
            ↵ other_hypercap_threshold: "
            f"{poc_other_threshold_leak_n} rows."
        )

if "first_other_src" in df.columns:
    first_other_poc_n =
    ↵ int(df["first_other_src"].astype("string").str.upper().eq("POC").sum())
    if first_other_poc_n > 0:
        raise ValueError(
            "first_other_src includes POC rows after LAB-only OTHER
            ↵ quarantine policy: "
            f"{first_other_poc_n} rows."
        )

# Harmonize NIH race/ethnicity into a collapsed reporting label at
# cohort-build time.
RACE_NIH_MAP = {
    "WHITE": "WHITE",
    "BLACK OR AFRICAN AMERICAN": "BLACK",
    "ASIAN": "ASIAN",
    "AMERICAN INDIAN OR ALASKA NATIVE": "AI/AN",
    "NATIVE HAWAIIAN OR OTHER PACIFIC ISLANDER": "NH/PI",
    "MORE THAN ONE RACE": "MULTIRACIAL",
    "UNKNOWN OR NOT REPORTED": "UNKNOWN",
    "OTHER": "OTHER",
    "UNKNOWN": "UNKNOWN",
}
ETH_NIH_MAP = {
    "HISPANIC OR LATINO": "HISPANIC",
    "NOT HISPANIC OR LATINO": "NOT HISPANIC",
    "UNKNOWN OR NOT REPORTED": "UNKNOWN",
    "UNKNOWN": "UNKNOWN",
}
LABEL_MAP = {
    "WHITE": "Non-Hispanic White",
    "BLACK": "Non-Hispanic Black",
}
```

```

    "ASIAN": "Non-Hispanic Asian",
    "AI/AN": "Non-Hispanic American Indian/Alaska Native",
    "NH/PI": "Non-Hispanic Native Hawaiian/Pacific Islander",
    "MULTIRACIAL": "Non-Hispanic Multiracial/Other",
    "OTHER": "Unknown/Other",
    "UNKNOWN": "Unknown/Other",
}

def addCollapsedRaceEth(
    frame: pd.DataFrame,
    race_col: str = "nih_race",
    eth_col: str = "nih_ethnicity",
    out_col: str = "race_eth_collapsed",
) -> pd.DataFrame:
    """Collapse NIH race/ethnicity fields into publication-friendly
    categories."""
    if race_col not in frame.columns or eth_col not in frame.columns:
        return frame

    def _std(value):
        if pd.isna(value):
            return None
        return re.sub(r"\s+", " ", str(value).strip()).upper()

    race_std = frame[race_col].map(lambda value:
    ↵ RACE_NIH_MAP.get(_std(value), "OTHER"))
    eth_std = frame[eth_col].map(lambda value: ETH_NIH_MAP.get(_std(value),
    ↵ "UNKNOWN"))

    collapsed = []
    for race_value, eth_value in zip(race_std, eth_std):
        if eth_value == "HISPANIC":
            collapsed.append("Hispanic/Latino")
        else:
            collapsed.append(LABEL_MAP.get(race_value, "Unknown/Other"))

    category_order = [
        "Hispanic/Latino",
        "Non-Hispanic White",
        "Non-Hispanic Black",
        "Non-Hispanic Asian",
        "Non-Hispanic American Indian/Alaska Native",
    ]

```

```

    "Non-Hispanic Native Hawaiian/Pacific Islander",
    "Non-Hispanic Multiracial/Other",
    "Unknown/Other",
]

frame[out_col] = pd.Categorical(collapsed, categories=category_order,
↪ ordered=True)
frame["nih_race_std"] = race_std
frame["nih_ethnicity_std"] = eth_std
return frame

```

df = addCollapsedRaceEth(df)

print("Final df rows:", len(df), "cols:", len(df.columns))

Safety checks

```

assert "subject_id" in df.columns, "subject_id missing from final df"
assert not any(c.endswith("_x") or c.endswith("_y") for c in df.columns),
↪ "Found suffixed columns"
print("Final df rows:", len(df), "cols:", len(df.columns))
df.head(3)

```

Threshold sanity (hadm-level): ABG= 30257 VBG= 14057 OTHER= 4148 ANY= 39200
POC OTHER quarantined hadm count: 0
Final df rows: 40167 cols: 109
Final df rows: 40167 cols: 109

	hadm_id	subject_id	admittime	dischtime	deathtime	admission_type
0	23485217	10584718	2165-02-12 15:41:00	2165-03-06 08:20:00	2165-03-06 08:20:00	EW EMER.
1	27004173	10297948	2135-08-23 18:46:00	2135-09-08 15:50:00	NaT	URGENT
2	22661627	10032409	2130-01-12 18:42:00	2130-01-21 14:32:00	NaT	EW EMER.

Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪ consistent across notebook stages.

--- Cohort flow counts (ED / ICU / blood gas / hypercapnia / CC) ---

```

# 1) Dataset-level ED counts
SQL["ed_counts_sql"] = f"""
SELECT

```

```

        COUNT(*) AS total_ed_encounters,
        COUNTIF(hadm_id IS NOT NULL) AS ed_encounters_with_hadm
    FROM `'{PHYS}`.{ED}.edstays`
"""

SQL["ed_to_icu_sql"] = f"""
SELECT
    COUNT(DISTINCT e.hadm_id) AS ed_to_icu_hadm,
    COUNT(DISTINCT e.stay_id) AS ed_to_icu_edstays
FROM `'{PHYS}`.{ED}.edstays` e
JOIN `'{PHYS}`.{ICU}.icustays` i USING (hadm_id)
WHERE e.hadm_id IS NOT NULL
"""

try:
    ed_counts = run_sql_bq(sql("ed_counts_sql"))
    ed_to_icu = run_sql_bq(sql("ed_to_icu_sql"))
except Exception as e:
    print("Warning: ED/ICU counts query failed:", e)
    ed_counts = None
    ed_to_icu = None

# 2) Cohort-level counts (admission-level)
cohort_union = int((cohort_any["enrolled_any"] == 1).sum()) if "cohort_any" \
    in globals() and "enrolled_any" in cohort_any.columns else len(hadm_list)
cohort_df_n = len(df)

# Any blood gas present (ABG/VBG, LAB/POC)
co2_cols = [c for c in [
    "lab_abg_paco2", "lab_vbg_paco2", "poc_abg_paco2", "poc_vbg_paco2"
] if c in df.columns]
any_bg = int(df[co2_cols].notna().any(axis=1).sum()) if co2_cols else None

# Hypercapnia thresholds and ICD
icd_count = int((df["any_hypercap_icd"] == 1).sum()) if "any_hypercap_icd" \
    in df.columns else None
threshold_count = int((df["pco2_threshold_any"] == 1).sum()) if \
    "pco2_threshold_any" in df.columns else None

# ED chief complaint missing / present (within cohort)
if "ed_triage_cc" in df.columns:
    mask_cc_present = df["ed_triage_cc"].notna() &
        (df["ed_triage_cc"].astype(str).str.strip() != "")

```

```

    cc_present = int(mask_cc_present.sum())
    cc_missing = int(~mask_cc_present.sum())
else:
    cc_present = None
    cc_missing = None

# ED→ICU within cohort (admissions with ED triage data and ICU stay)
if "first_icu_stay_id" in df.columns and "ed_triage_cc" in df.columns:
    cohort_ed_to_icu = int((df["first_icu_stay_id"].notna() &
    ~ mask_cc_present).sum())
else:
    cohort_ed_to_icu = None

rows = []
if ed_counts is not None:
    rows.append({"step": "Total ED encounters (edstays)", "count":
    int(ed_counts.loc[0, "total_ed_encounters"]), "scope": "All ED dataset"})
    rows.append({"step": "ED encounters with hadm_id", "count":
    int(ed_counts.loc[0, "ed_encounters_with_hadm"]), "scope": "All ED
    dataset"})
if ed_to_icu is not None:
    rows.append({"step": "ED→ICU admissions (distinct hadm_id)", "count":
    int(ed_to_icu.loc[0, "ed_to_icu_hadm"]), "scope": "All ED+ICU"})
    rows.append({"step": "ED→ICU ED-stays (distinct stay_id)", "count":
    int(ed_to_icu.loc[0, "ed_to_icu_edstays"]), "scope": "All ED+ICU"})

rows.append({"step": "Cohort admissions (union ICD thresholds)", "count":
    cohort_union, "scope": "Cohort"})
rows.append({"step": "Cohort admissions after merges (df rows)", "count":
    cohort_df_n, "scope": "Cohort"})
if any_bg is not None:
    rows.append({"step": "Cohort with any ABG/VBG (LAB or POC)", "count":
    any_bg, "scope": "Cohort"})
if threshold_count is not None:
    rows.append({"step": "Cohort meeting hypercapnia thresholds", "count":
    threshold_count, "scope": "Cohort"})
if icd_count is not None:
    rows.append({"step": "Cohort meeting ICD code criteria", "count":
    icd_count, "scope": "Cohort"})
if cc_present is not None:
    rows.append({"step": "Cohort with ED chief complaint present", "count":
    cc_present, "scope": "Cohort"})
if cc_missing is not None:

```

```

    rows.append({"step": "Cohort excluded for missing ED chief complaint",
    ↵ "count": cc_missing, "scope": "Cohort"})
if cohort_ed_to_icu is not None:
    rows.append({"step": "Cohort ED→ICU (ED CC present + ICU stay)", "count":
    ↵ cohort_ed_to_icu, "scope": "Cohort"})

flow_counts = pd.DataFrame(rows)
flow_counts

```

step	count	scope
0 Total ED encounters (edstays)	425087	All ED dataset
1 ED encounters with hadm_id	203016	All ED dataset
2 ED→ICU admissions (distinct hadm_id)	31862	All ED+ICU
3 ED→ICU ED-stays (distinct stay_id)	31916	All ED+ICU
4 Cohort admissions (union ICD thresholds)	40168	Cohort
5 Cohort admissions after merges (df rows)	40167	Cohort
6 Cohort with any ABG/VBG (LAB or POC)	38991	Cohort
7 Cohort meeting hypercapnia thresholds	39200	Cohort
8 Cohort meeting ICD code criteria	4237	Cohort
9 Cohort with ED chief complaint present	11769	Cohort
10 Cohort excluded for missing ED chief complaint	28398	Cohort
11 Cohort ED→ICU (ED CC present + ICU stay)	9035	Cohort

```

# Purpose: Quantify overlap between ascertainment routes so non-exclusive
↪ cohorts are explicit.

# --- Ascertainment overlap counts (ABG/VBG/ICD) ---

required = ["abg_hypercap_threshold", "vbg_hypercap_threshold",
    ↵ "other_hypercap_threshold", "any_hypercap_icd"]
missing = [c for c in required if c not in df.columns]
if missing:
    raise KeyError(f"Missing required columns for overlap counts: {missing}")

abg = pd.to_numeric(df["abg_hypercap_threshold"],
    ↵ errors="coerce").fillna(0).astype(int)
vbg = pd.to_numeric(df["vbg_hypercap_threshold"],
    ↵ errors="coerce").fillna(0).astype(int)
other = pd.to_numeric(df["other_hypercap_threshold"],
    ↵ errors="coerce").fillna(0).astype(int)

```

```

gas_any = (
    pd.to_numeric(df.get("pco2_threshold_any", None), errors="coerce")
    if "pco2_threshold_any" in df.columns else (abg | vbg | other)
)
if hasattr(gas_any, "fillna"):
    gas_any = gas_any.fillna(0).astype(int)
else:
    gas_any = gas_any.astype(int)

icd = pd.to_numeric(df["any_hypercap_icd"],
    errors="coerce").fillna(0).astype(int)

total_n = len(df)
ngas = int((gas_any == 1).sum())

abg_vbg_overlap = pd.DataFrame([
    {"group": "ABG-only", "count": int(((abg==1) & (vbg==0) &
    (other==0)).sum())},
    {"group": "VBG-only", "count": int(((vbg==1) & (abg==0) &
    (other==0)).sum())},
    {"group": "Other-only", "count": int(((other==1) & (abg==0) &
    (vbg==0)).sum())},
    {"group": "Mixed (>=2 routes)", "count": int(((abg + vbg + other) >=
    2).sum())},
])
if ngas > 0:
    abg_vbg_overlap["pct_of_gas"] = (abg_vbg_overlap["count"] / ngas *
    100).round(1)
else:
    abg_vbg_overlap["pct_of_gas"] = 0.0
abg_vbg_overlap["pct_of_cohort"] = (abg_vbg_overlap["count"] / max(total_n,1)
    * 100).round(1)

icd_gas_overlap = pd.DataFrame([
    {"group": "ICD+Gas", "count": int(((icd==1) & (gas_any==1)).sum())},
    {"group": "ICD-only", "count": int(((icd==1) & (gas_any==0)).sum())},
    {"group": "Gas-only", "count": int(((icd==0) & (gas_any==1)).sum())},
    {"group": "Neither", "count": int(((icd==0) & (gas_any==0)).sum())},
])
icd_gas_overlap["pct_of_cohort"] = (icd_gas_overlap["count"] / max(total_n,1)
    * 100).round(1)

print("ABG/VBG/Other overlap (among gas-positive):")

```

```

print(abg_vbg_overlap.to_string(index=False))
print("ICD vs Gas overlap (cohort-level):")
print(icd_gas_overlap.to_string(index=False))

ABG/VBG/Other overlap (among gas-positive):
    group  count  pct_of_gas  pct_of_cohort
    ABG-only  22895      58.4      57.0
    VBG-only   7517      19.2      18.7
    Other-only   848       2.2       2.1
    Mixed (>=2 routes)  7940      20.3      19.8
ICD vs Gas overlap (cohort-level):
    group  count  pct_of_cohort
    ICD+Gas   3270       8.1
    ICD-only   967       2.4
    Gas-only  35930      89.5
    Neither     0       0.0

# Purpose: Capture ED presentation context (chief complaint, acuity, and
#           ↳ early vitals) at the ED-stay level.

# --- Missingness summary (chief complaint, race/ethnicity, ED triage/vitals)
#           ↳ ---

import numpy as np

summary_rows = []

# Chief complaint missingness (ED triage CC)
if "ed_triage_cc" in df.columns:
    cc_present = df["ed_triage_cc"].notna() &
    ↳ (df["ed_triage_cc"].astype(str).str.strip() != "")
    summary_rows.append({
        "variable": "ed_triage_cc_present",
        "missing_n": int(~cc_present).sum(),
        "missing_pct": float(~cc_present).mean()
    })

# Race/Ethnicity missingness (NIH categories + raw sources)
unknown_tokens = [
    "unknown or not reported",
    "unknown",
    "not reported",

```

```

    "missing",
    "declined",
    "unable"
}

def _missing_rate(series):
    if series is None:
        return None, None
    s = series.astype(str).str.strip()
    is_missing = series.isna() | (s == "") |
    ↵ s.str.lower().isin(unknown_tokens)
    return int(is_missing.sum()), float(is_missing.mean())

for col in ["nih_race", "nih_ethnicity", "race_hosp_raw", "race_ed_raw"]:
    if col in df.columns:
        m_n, m_p = _missing_rate(df[col])
        summary_rows.append({
            "variable": col,
            "missing_n": m_n,
            "missing_pct": m_p
        })

missing_summary = pd.DataFrame(summary_rows)
print("Missingness summary (key variables):")
missing_summary

# ED triage + first ED vitals missingness
triage_cols = [c for c in df.columns if c.startswith("ed_triage_")]
first_cols = [c for c in df.columns if c.startswith("ed_first_")]

vital_cols = triage_cols + first_cols
if vital_cols:
    miss_tbl = (
        pd.DataFrame({"variable": vital_cols})
        .assign(
            missing_n=lambda d: [int(df[c].isna().sum()) for c in
        ↵ d["variable"]],
            missing_pct=lambda d: [float(df[c].isna().mean()) for c in
        ↵ d["variable"]]
        )
        .sort_values("missing_pct", ascending=False)
    )
    print("Missingness summary (ED triage + first ED vitals):")

```

```

    miss_tbl
else:
    miss_tbl = pd.DataFrame(columns=["variable", "missing_n", "missing_pct"])

```

```

Missingness summary (key variables):
Missingness summary (ED triage + first ED vitals):

```

2.3.11 11) Sanity checks

Rationale: Run QC checks to validate units, flags, and basic data integrity before export.

```
# Purpose: Derive respiratory support flags and timing fields for IMV/NIV
→ exposure.
```

```
# ICU LOS negative?
if {"icu_los_days", "first_icu_stay_id"}.issubset(df.columns):
    neg_los = int((df["icu_los_days"] < 0).fillna(False).sum())
    print("Negative ICU LOS rows:", neg_los)
```

```
# Vent flags consistency
vent_cols = {"imv_flag", "niv_flag", "any_vent_flag"}
if vent_cols.issubset(df.columns):
    any_calc = ((df["imv_flag"]==1) |
    → (df["niv_flag"]==1)).fillna(False).astype(int)
    any_flag = pd.to_numeric(df["any_vent_flag"],
    → errors="coerce").fillna(0).astype(int)
    mism = int((any_calc != any_flag).sum())
    print("any_vent_flag mismatches vs (imv|niv):", mism)
```

```
# UOMs: expect mmhg only
uom_cols = [c for c in df.columns if c.endswith("_paco2_uom")]
for c in uom_cols:
    vals = sorted(pd.Series(df[c]).dropna().astype(str).str.lower().str_
    → .strip().unique().tolist())
    print(c, vals)
```

```
# ABG/VBG coverage QC
def qc_pair(df, ph_col, co2_col, label, ph_lo=6.3, ph_hi=7.8, co2_lo=5,
    → co2_hi=200):
    ph = pd.to_numeric(df.get(ph_col), errors="coerce")
    co2 = pd.to_numeric(df.get(co2_col), errors="coerce")
```

```

    return {
        "pair": label,
        "present_any": int(((ph.notna() | (co2.notna())).sum())),
        "present_both": int(((ph.notna()) & (co2.notna())).sum()),
        "only_ph": int(((ph.notna()) & (~co2.notna())).sum()),
        "only_pco2": int(((co2.notna()) & (~ph.notna())).sum()),
        "ph_oob": int(((ph < ph_lo) | (ph > ph_hi)) &
                      ph.notna()).sum(),
        "pco2_oob": int(((co2 < co2_lo) | (co2 > co2_hi)) &
                      co2.notna()).sum(),
    }
}

qc = pd.DataFrame([
    qc_pair(df, "lab_abg_ph", "lab_abg_paco2", "LAB ABG"),
    qc_pair(df, "lab_vbg_ph", "lab_vbg_paco2", "LAB VBG"),
    qc_pair(df, "poc_abg_ph", "poc_abg_paco2", "POC ABG"),
    qc_pair(df, "poc_vbg_ph", "poc_vbg_paco2", "POC VBG"),
])
qc

```

```

Negative ICU LOS rows: 0
any_vent_flag mismatches vs (imv|niv): 0
lab_abg_paco2_uom ['mmhg']
lab_vbg_paco2_uom ['mmhg']
lab_other_paco2_uom ['mmhg']
poc_abg_paco2_uom []
poc_vbg_paco2_uom []
poc_other_paco2_uom []

```

	pair	present_any	present_both	only_ph	only_pco2	ph_oob	pco2_oob
0	LAB ABG	32122	32115	0	7	0	0
1	LAB VBG	17802	17797	0	5	0	0
2	POC ABG	0	0	0	0	0	0
3	POC VBG	0	0	0	0	0	0

2.3.12 PDF-ready long tables

```

print(
    render_latex_longtable(
        flow_counts,

```

```

        caption=f"Cohort flow and key denominator checkpoints
    ↵  (N={len(df):,}).",
        label="tab:cohort_flow_counts",
        index=False,
    )
)
print(
    render_latex_longtable(
        abg_vbg_overlap,
        caption="ABG/VBG/OTHER overlap among gas-positive encounters.",
        label="tab:gas_overlap_abg_vbg_other",
        index=False,
    )
)
print(
    render_latex_longtable(
        icd_gas_overlap,
        caption="ICD versus gas ascertainment overlap in the full cohort.",
        label="tab:icd_gas_overlap",
        index=False,
    )
)
print(
    render_latex_longtable(
        missing_summary,
        caption="Missingness summary for key cohort variables.",
        label="tab:key_missingness",
        index=False,
    )
)
if "miss_tbl" in globals() and not miss_tbl.empty:
    print(
        render_latex_longtable(
            miss_tbl,
            caption="ED triage and first ED vital-sign missingness summary.",
            label="tab:ed_vital_missingness",
            landscape=True,
            index=False,
        )
    )
print(
    render_latex_longtable(
        qc,

```

```

        caption="ABG/VBG pH and pCO2 plausibility quality checks.",  

        label="tab:abg_vbg_qc",  

        landscape=True,  

        index=False,  

    )  

)  
  

\begin{longtable}{lrl}
\caption{Cohort flow and key denominator checkpoints (N=40,167).} \label{tab:cohort_flow_count}
\toprule
step & count & scope \\
\midrule
\endfirsthead
\caption[]{Cohort flow and key denominator checkpoints (N=40,167).} \\
\toprule
step & count & scope \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
Total ED encounters (edstays) & 425087 & All ED dataset \\
ED encounters with hadm_id & 203016 & All ED dataset \\
ED→ICU admissions (distinct hadm_id) & 31862 & All ED+ICU \\
ED→ICU ED-stays (distinct stay_id) & 31916 & All ED+ICU \\
Cohort admissions (union ICD thresholds) & 40168 & Cohort \\
Cohort admissions after merges (df rows) & 40167 & Cohort \\
Cohort with any ABG/VBG (LAB or POC) & 38991 & Cohort \\
Cohort meeting hypercapnia thresholds & 39200 & Cohort \\
Cohort meeting ICD code criteria & 4237 & Cohort \\
Cohort with ED chief complaint present & 11769 & Cohort \\
Cohort excluded for missing ED chief complaint & 28398 & Cohort \\
Cohort ED→ICU (ED CC present + ICU stay) & 9035 & Cohort \\
\end{longtable}

\begin{longtable}{lrrr}
\caption{ABG/VBG/OTHER overlap among gas-positive encounters.} \label{tab:gas_overlap_abg_vbg}
\toprule
group & count & pct_of_gas & pct_of_cohort \\
\midrule

```

```

\endfirsthead
\caption[]{ABG/VBG/OTHER overlap among gas-positive encounters.} \\
\toprule
group & count & pct_of_gas & pct_of_cohort \\
\midrule
\endhead
\midrule
\multicolumn{4}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ABG-only & 22895 & 58.400000 & 57.000000 \\
VBG-only & 7517 & 19.200000 & 18.700000 \\
Other-only & 848 & 2.200000 & 2.100000 \\
Mixed (>=2 routes) & 7940 & 20.300000 & 19.800000 \\
\end{longtable}

\begin{longtable}{lrr}
\caption{ICD versus gas ascertainment overlap in the full cohort.} \label{tab:icd_gas_overlap}
\toprule
group & count & pct_of_cohort \\
\midrule
\endfirsthead
\caption{ICD versus gas ascertainment overlap in the full cohort.} \\
\toprule
group & count & pct_of_cohort \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ICD+Gas & 3270 & 8.100000 \\
ICD-only & 967 & 2.400000 \\
Gas-only & 35930 & 89.500000 \\
Neither & 0 & 0.000000 \\
\end{longtable}

\begin{longtable}{lrr}
\caption{Missingness summary for key cohort variables.} \label{tab:key_missingness} \\

```

```

\toprule
variable & missing_n & missing_pct \\
\midrule
\endfirsthead
\caption[]{Missingness summary for key cohort variables.} \\
\toprule
variable & missing_n & missing_pct \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ed_triage_cc_present & 28398 & 0.706998 \\
nih_race & 7473 & 0.186048 \\
nih_ethnicity & 6984 & 0.173874 \\
race_hosp_raw & 3357 & 0.083576 \\
race_ed_raw & 29569 & 0.736152 \\
\end{longtable}

\begin{landscape}\begin{longtable}{lrr}
\caption{ED triage and first ED vital-sign missingness summary.} \label{tab:ed_vital_missing}
\toprule
variable & missing_n & missing_pct \\
\midrule
\endfirsthead
\caption[]{ED triage and first ED vital-sign missingness summary.} \\
\toprule
variable & missing_n & missing_pct \\
\midrule
\endhead
\midrule
\multicolumn{3}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
ed_first_rhythm & 39511 & 0.983668 \\
ed_first_temp & 33266 & 0.828192 \\
ed_triage_temp & 31836 & 0.792591 \\
ed_triage_rr & 31533 & 0.785047 \\

```

```

ed_triage_o2sat & 31499 & 0.784201 \\
ed_triage_dbp & 31441 & 0.782757 \\
ed_triage_sbp & 31402 & 0.781786 \\
ed_triage_hr & 31372 & 0.781039 \\
ed_first_o2sat & 31259 & 0.778226 \\
ed_first_dbp & 31018 & 0.772226 \\
ed_first_sbp & 31018 & 0.772226 \\
ed_first_rr & 30943 & 0.770359 \\
ed_first_hr & 30829 & 0.767521 \\
ed_first_pain & 30660 & 0.763313 \\
ed_triage_pain & 30146 & 0.750517 \\
ed_triage_acuity & 29683 & 0.738990 \\
ed_first_vitals_time & 29399 & 0.731919 \\
ed_triage_cc & 28398 & 0.706998 \\
\end{longtable}
\n\end{landscape}\n
\begin{landscape}\n\begin{longtable}{lrrrrrrr}
\caption{ABG/VBG pH and pCO2 plausibility quality checks.} \label{tab:abg_vbg_qc} \\
\toprule
pair & present_any & present_both & only_ph & only_pco2 & ph_oob & pco2_oob \\
\midrule
\endfirsthead
\caption[] {ABG/VBG pH and pCO2 plausibility quality checks.} \\
\toprule
pair & present_any & present_both & only_ph & only_pco2 & ph_oob & pco2_oob \\
\midrule
\endhead
\midrule
\multicolumn{7}{r}{Continued on next page} \\
\midrule
\endfoot
\bottomrule
\endlastfoot
LAB ABG & 32122 & 32115 & 0 & 7 & 0 & 0 \\
LAB VBG & 17802 & 17797 & 0 & 5 & 0 & 0 \\
POC ABG & 0 & 0 & 0 & 0 & 0 & 0 \\
POC VBG & 0 & 0 & 0 & 0 & 0 & 0 \\
\end{longtable}
\n\end{landscape}\n

```

2.3.13 12) Save to Excel

Rationale: Persist cohort outputs for downstream annotation and NLP analyses.

```
# Purpose: Optional archive export of the full hadm-level cohort workbook.

from datetime import datetime

WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
    "1"

if WRITE_ARCHIVE_XLSX_EXPORTS:
    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    out_path = prior_runs_dir / f"mimic_hypercap_EXT_bq_abg_vbg_{datetime.}_
    now().strftime('%Y%m%d_%H%M%S')}.xlsx"
    with pd.ExcelWriter(out_path, engine="openpyxl") as xw:
        df.to_excel(xw, sheet_name="cohort", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    out_path
else:
    print("Skipping archive hadm-level workbook export (set
        WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")
```

Skipping archive hadm-level workbook export (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

2.3.14 Create Annotation Dataset

Rationale: Create ED chief-complaint subsets and a reproducible sample for manual annotation.

```
# Purpose: Capture ED presentation context (chief complaint, acuity, and
    early vitals) at the ED-stay level.

# ---- Extra exports: (1) ED chief-complaint only; (2) random sample of 160
    patients ----
from datetime import datetime

# Dated artifacts go to archive folder (optional)
prior_runs_dir = DATA_DIR / "prior runs"
```

```

prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')
WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
    ↪ "1"

# 1) Filter to rows with a non-empty ED chief complaint
if "ed_triage_cc" not in df.columns:
    raise KeyError(
        "Column 'ed_triage_cc' not found in df. "
        "Ensure the ED triage merge cell ran earlier."
    )

mask_cc = df["ed_triage_cc"].notna() &
    ↪ (df["ed_triage_cc"].astype(str).str.strip() != "")
df_cc = df.loc[mask_cc].copy()

print(f"ED-CC present rows: {len(df_cc)} of {len(df)} "
      f"({{(len(df_cc) / max(len(df), 1)):.1%}} of cohort.)")

# Save ED-CC-only cohort only when archive exports are enabled
if WRITE_ARCHIVE_XLSX_EXPORTS:
    out_path_cc = prior_runs_dir /
    ↪ f"mimic_hypcap_EXT_EDcc_only_bq_abg_vbg_{timestamp}.xlsx"
    with pd.ExcelWriter(out_path_cc, engine="openpyxl") as xw:
        df_cc.to_excel(xw, sheet_name="cohort_cc_only", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    print("Saved:", out_path_cc)

# 2) Random sample of n = 160 patients (distinct subject_id), one row per
    ↪ patient
if "subject_id" not in df_cc.columns:
    raise KeyError("Column 'subject_id' missing; cannot sample by patient.")

# Make a one-row-per-patient frame by earliest admission
if "admittime" in df_cc.columns:
    df_cc_one = (
        df_cc.sort_values(["subject_id", "admittime"])
            .groupby("subject_id", as_index=False)
            .head(1)
    )

```

```

else:
    # Fallback if admittime not present: choose the smallest hadm_id per
    # patient
    df_cc_one = (
        df_cc.sort_values(["subject_id", "hadm_id"])
            .groupby("subject_id", as_index=False)
            .head(1)
    )

N = 160
n_avail = len(df_cc_one)
n_take = min(N, n_avail)
if n_avail < N:
    print(f"Warning: only {n_avail} unique patients with ED chief complaint;
          sampling all of them.")

RANDOM_SEED = 42
df_cc_sample = df_cc_one.sample(n=n_take, random_state=RANDOM_SEED)

# Save the sample only when archive exports are enabled
if WRITE_ARCHIVE_XLSX_EXPORTS:
    out_path_cc_sample = prior_runs_dir /
    f"mimic_hypercap_EXT_EDcc_sample{n_take}_bq_abg_vbg_{timestamp}.xlsx"
    with pd.ExcelWriter(out_path_cc_sample, engine="openpyxl") as xw:
        df_cc_sample.to_excel(xw, sheet_name="cohort_cc_sample", index=False)
        try:
            qc.to_excel(xw, sheet_name="qc_abg_vbg", index=False)
        except Exception:
            pass
    print("Saved:", out_path_cc_sample)
else:
    print("Skipping ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1
          to enable).")

ED-CC present rows: 11769 of 40167 (29.3% of cohort).
Skipping ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

```

3 ED-stay cohort expansion (timing, severity, comorbidity, outcomes)

Rationale: Build a one-row-per-ED-stay analytic extract with time-anchored gas phenotypes, key comorbidities, and outcomes.

3.0.1 Phase 0 — Inventory & missing-field registry

Rationale: Detect which fields already exist and only add missing fields to avoid redundant extraction or join explosions.

```
# Purpose: Inventory available fields and identify missing target variables
#           before enrichment steps.

from pathlib import Path
import json
import pandas as pd

# Identify current cohort dataframe (admission-level)
if 'df' not in globals():
    raise NameError("Expected admission-level df to exist before inventory
                    step.")

ED_KEY = "ed_stay_id" # target key for ED-level cohort

print("Current admission-level df:", df.shape)
print("Current columns count:", len(df.columns))
if ED_KEY in df.columns:
    print("ED stay unique count:", int(df[ED_KEY].nunique()))
else:
    print("ED stay unique count: ED_KEY not in columns")

# Persist columns snapshot
cols_out = WORK_DIR / "current_columns.json"
cols_out.write_text(json.dumps(sorted(df.columns), indent=2))
print("Wrote:", cols_out)

# Persist ED-stay columns snapshot (if ed_df exists)
if "ed_df" in globals():
    ed_cols_out = WORK_DIR / "ed_columns.json"
    ed_cols_out.write_text(json.dumps(sorted(ed_df.columns), indent=2))
    print("Wrote:", ed_cols_out)

Current admission-level df: (40167, 109)
Current columns count: 109
ED stay unique count: ED_KEY not in columns
```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/current_columns.json

```
# Purpose: Capture ED presentation context (chief complaint, acuity, and
#           early vitals) at the ED-stay level.

# Target field registry
TARGET_RAW_FIELDS = {

    "ed_edstays": [
        "ed_stay_id", "subject_id", "hadm_id", "ed_intime", "ed_outtime",
        ↵ "ed_intime_first",
        "arrival_transport", "disposition", "ed_gender", "ed_race",
    ],
    "ed_triage": [
        "ed_triage_temp", "ed_triage_hr", "ed_triage_rr", "ed_triage_o2sat",
        "ed_triage_sbp", "ed_triage_dbp", "ed_triage_pain", "ed_triage_acuity",
        ↵ "ed_triage_cc",
    ],
    "ed_vitals_first": [
        "ed_first_vitals_time", "ed_first_temp", "ed_first_hr", "ed_first_rr",
        "ed_first_o2sat", "ed_first_sbp", "ed_first_dbp", "ed_first_rhythm",
        ↵ "ed_first_pain",
    ],
    "admissions": [
        "admittime", "dischtime", "deathtime", "hospital_expire_flag",
        "admission_type", "admission_location", "discharge_location",
        "insurance", "language", "marital_status", "hosp_race",
    ],
    "icu": [
        "icu_stay_id", "icu_intime_first", "icu_outtime_last", "icu_los_total",
        ↵ "n_icu_stays",
        "first_careunit", "last_careunit",
    ],
    "labs_gas": [
        "first_gas_time", "first_pco2", "first_ph", "first_hco3",
        ↵ "first_lactate",
        "max_pco2_0_6h", "min_ph_0_6h", "max_pco2_0_24h", "min_ph_0_24h",
        "flag_abg_hypercapnia", "flag_vbg_hypercapnia",
        ↵ "flag_other_hypercapnia", "flag_any_gas_hypercapnia",
        "gas_source_other_rate", "poc_inclusion_enabled",
        ↵ "poc_inclusion_reason",
        "first_gas_specimen_type", "first_gas_specimen_present",
        ↵ "first_gas_pco2_itemid", "first_gas_pco2_fluid",
    ]
}
```

```

    "co2_other_is_blood_asserted", "first_gas_anchor_has_pco2", ]
    ↵  "first_gas_anchor_source_validated",
    "dt_first_imv_hours", "dt_first_niv_hours", "first_other_time",
],
"omr": [
    "bmi_closest_pre_ed", "height_closest_pre_ed", "weight_closest_pre_ed",
    "bmi_closest_pre_ed_model", "height_closest_pre_ed_model", ]
    ↵  "weight_closest_pre_ed_model",
    "bmi_outlier_flag", "height_outlier_flag", "weight_outlier_flag",
    "anthro_timing_tier", "anthro_days_offset", "anthro_chartdate",
    ↵  "anthro_timing_uncertain",
    "anthro_source", "anthro_obstime", "anthro_hours_offset",
    ↵  "anthro_timing_basis",
],
"dx_flags": [
    "flag_copd", "flag_osa_ohs", "flag_chf", "flag_neuromuscular",
    "flag_opioid_substance", "flag_pneumonia",
],
"timing": [
    "dt_first_qualifying_gas_hours", "presenting_hypercapnia",
    ↵  "late_hypercapnia",
    "qualifying_gas_observed", "presenting_hypercapnia_tri",
    ↵  "late_hypercapnia_tri", "hypercap_timing_class",
    "dt_first_imv_hours", "dt_first_niv_hours", "abg_before_imv",
    ↵  "vbg_before_imv",
    "dt_first_imv_hours_model", "dt_first_niv_hours_model",
    "imv_time_outside_window_flag", "niv_time_outside_window_flag",
    "hospital_los_negative_flag", "admittime_before_ed_intime_flag",
    "ph_band", "hco3_band", "lactate_band",
],
}
}

TARGET_DERIVED_FIELDS = [
    "hospital_los_hours", "hospital_los_hours_model", "in_hospital_death",
]

TARGET_FIELDS = sorted({c for v in TARGET_RAW_FIELDS.values() for c in v} |
    ↵  set(TARGET_DERIVED_FIELDS))

# Grouped missing report
missing_by_group = {}
for group, cols in TARGET_RAW_FIELDS.items():
    missing_by_group[group] = [c for c in cols if c not in df.columns]

```

```

missing_derived = [c for c in TARGET_DERIVED_FIELDS if c not in df.columns]
missing_by_group["derived"] = missing_derived

missing_flat = [c for cols in missing_by_group.values() for c in cols]
print("Missing fields total:", len(missing_flat))
for group, cols in missing_by_group.items():
    if cols:
        print(f"- {group}: {cols}")

```

Missing fields total: 88

- ed_edstays: ['ed_stay_id', 'ed_intime', 'ed_outtime', 'ed_intime_first', 'arrival_transport']
- admissions: ['hospital_expire_flag', 'language', 'marital_status', 'hosp_race']
- icu: ['icu_stay_id', 'icu_intime_first', 'icu_outtime_last', 'icu_los_total', 'n_icu_stays']
- labs_gas: ['first_gas_time', 'first_pco2', 'first_ph', 'first_hco3', 'first_lactate', 'max']
- omr: ['bmi_closest_pre_ed', 'height_closest_pre_ed', 'weight_closest_pre_ed', 'bmi_closest']
- dx_flags: ['flag_copd', 'flag.osa.ohs', 'flag_chf', 'flag_neuromuscular', 'flag_opioid_sub']
- timing: ['dt_first_qualifying_gas_hours', 'presenting_hypercapnia', 'late_hypercapnia', 'qu']
- derived: ['hospital_los_hours', 'hospital_los_hours_model', 'in_hospital_death']

3.0.2 Phase 1 — ED encounter spine and ED enrichment (one row per ED stay)

Rationale: Create a dedicated ED-stay-level cohort with ED-specific attributes to avoid mixing admission- and ED-level grains.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
#           ↴ consistent across notebook stages.

# ED stay spine (rename stay_id to ed_stay_id)

SQL["ed_spine_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
    s.stay_id AS ed_stay_id,
    s.subject_id,
    s.hadm_id,
    s.intime AS ed_intime,
    s.outtime AS ed_outtime,
    s.arrival_transport,
    s.disposition,
    s.gender AS ed_gender,
    s.race   AS ed_race

```

```

FROM `'{PHYS}`.{ED}`.edstays` s
JOIN hadms h ON h.hadm_id = s.hadm_id
"""

ed_spine = run_sql_bq(sql("ed_spine_sql"), {"hadms": hadm_list})
print("ED spine rows:", len(ed_spine), "unique ed_stay_id:",
      ed_spine["ed_stay_id"].nunique())

# ensure uniqueness
if ed_spine["ed_stay_id"].nunique() != len(ed_spine):
    raise ValueError("ed_stay_id not unique in ED spine")

# First ED presentation time per admission
ed_intime_first = (
    ed_spine.groupby("hadm_id", as_index=False)[["ed_intime"]]
    .min()
    .rename(columns={"ed_intime": "ed_intime_first"})
)

# Start ED-level df
ed_df = ed_spine.copy()
ed_df = ed_df.merge(ed_intime_first, on="hadm_id", how="left")

ED spine rows: 11786 unique ed_stay_id: 11786

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
#           consistent across notebook stages.

# ED triage and first ED vitals (reuse existing logic if present; otherwise
#           join)

def _needs_cols(df, cols):
    return (df is None) or any(c not in df.columns for c in cols)

# Use existing ed_triage / ed_first if already in memory from earlier cells
try:
    _ = ed_triage
except NameError:
    ed_triage = None

try:
    _ = ed_first

```

```

except NameError:
    ed_first = None

# Force re-query if required keys are missing
if _needs_cols(locals().get('ed_triage', None), ['ed_stay_id', 'hadm_id']):
    ed_triage = None
if _needs_cols(locals().get('ed_first', None), ['ed_stay_id']):
    ed_first = None

# If missing, prefer in-memory tables loaded earlier in the notebook.
# This avoids re-running expensive ED queries during nbconvert.
if ed_triage is None:
    if 'edmap' in globals() and 'tri_all' in globals():
        ed_triage = (
            edmap[['stay_id', 'hadm_id']]
            .merge(tri_all, on='stay_id', how='left')
            .rename(columns={'stay_id': 'ed_stay_id'})
        )
        print('ED triage rows (from in-memory tables):', len(ed_triage))
    else:
        ed_triage_sql = f"""
        WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
        SELECT
            s.stay_id AS ed_stay_id,
            s.hadm_id,
            t.temperature AS ed_triage_temp,
            t.heartrate AS ed_triage_hr,
            t.resprate AS ed_triage_rr,
            t.o2sat AS ed_triage_o2sat,
            t.sbp AS ed_triage_sbp,
            t.dbp AS ed_triage_dbp,
            t.pain AS ed_triage_pain,
            t.acuity AS ed_triage_acuity,
            t.chiefcomplaint AS ed_triage_cc
        FROM `{{PHYS}}.{{ED}}.edstays` s
        JOIN hadms h ON h.hadm_id = s.hadm_id
        LEFT JOIN `{{PHYS}}.{{ED}}.triage` t ON t.stay_id = s.stay_id
        """
        ed_triage = run_sql_bq(ed_triage_sql, {'hadms': hadm_list})
        print('ED triage rows (fallback query):', len(ed_triage))

if ed_first is None:
    if 'edmap' in globals() and 'first_stay_all' in globals():

```

```

ed_first = (
    edmap[['stay_id', 'hadm_id']]
    .merge(first_stay_all, on='stay_id', how='left')
    .rename(columns={'stay_id': 'ed_stay_id'})
)
print('ED first vitals rows (from in-memory tables):', len(ed_first))
else:
    ed_first_vitals_sql = """
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
edmap AS (
    SELECT stay_id, hadm_id
    FROM `{{PHYS}}.{{ED}}.edstays`
    WHERE hadm_id IN (SELECT hadm_id FROM hadms)
),
vs AS (
    SELECT stay_id, charttime, temperature, heartrate, resprate, o2sat,
    sbp, dbp, rhythm, pain
    FROM `{{PHYS}}.{{ED}}.vitalsign`
),
first_vs AS (
    SELECT
        v.stay_id,
        (ARRAY_AGG(STRUCT(v.charttime, v.temperature, v.heartrate,
    v.resprate, v.o2sat, v.sbp, v.dbp, v.rhythm, v.pain)
    ORDER BY v.charttime LIMIT 1)) [OFFSET(0)] AS pick
    FROM vs v
    JOIN edmap m USING (stay_id)
    GROUP BY v.stay_id
)
SELECT
    f.stay_id AS ed_stay_id,
    m.hadm_id,
    pick.charttime AS ed_first_vitals_time,
    pick.temperature AS ed_first_temp,
    pick.heartrate AS ed_first_hr,
    pick.resprate AS ed_first_rr,
    pick.o2sat AS ed_first_o2sat,
    pick.sbp AS ed_first_sbp,
    pick.dbp AS ed_first_dbp,
    pick.rhythm AS ed_first_rhythm,
    pick.pain AS ed_first_pain
FROM first_vs f
JOIN edmap m ON m.stay_id = f.stay_id
"""

```

```

"""
ed_first = run_sql_bq(ed_first_vitals_sql, {'hadms': hadm_list})
print('ED first vitals rows (fallback query):', len(ed_first))

# Debug columns before merge
print('ed_triage cols:', list(ed_triage.columns))
print('ed_first cols:', list(ed_first.columns))
print('ed_df cols:', list(ed_df.columns))

if 'ed_stay_id' not in ed_df.columns:
    raise KeyError('ed_df missing ed_stay_id; ensure ED spine cell ran.')

# Merge ED triage + vitals onto ed_df with available keys
merge_keys_triage = [k for k in ["ed_stay_id", "hadm_id"] if k in
    ↪ ed_df.columns and k in ed_triage.columns]
if not merge_keys_triage:
    raise KeyError("No common keys between ed_df and ed_triage")
ed_df = ed_df.merge(ed_triage, on=merge_keys_triage, how="left")
merge_keys_first = [k for k in ["ed_stay_id", "hadm_id"] if k in
    ↪ ed_df.columns and k in ed_first.columns]
if not merge_keys_first:
    raise KeyError("No common keys between ed_df and ed_first")
ed_df = ed_df.merge(ed_first, on=merge_keys_first, how="left")

ED triage rows (from in-memory tables): 11786
ED first vitals rows (from in-memory tables): 11786
ed_triage cols: ['ed_stay_id', 'hadm_id', 'ed_triage_temp', 'ed_triage_hr', 'ed_triage_rr',
ed_first cols: ['ed_stay_id', 'hadm_id', 'ed_first_vitals_time', 'ed_first_temp', 'ed_first_rr',
ed_df cols: ['ed_stay_id', 'subject_id', 'hadm_id', 'ed_intime', 'ed_outtime', 'arrival_transit_time', 'ed_triage_rr', 'ed_triage_hr', 'ed_triage_temp', 'ed_first_vitals_time', 'ed_first_temp', 'ed_first_rr']

```

3.0.3 Phase 2 — Hospital admission context and outcomes

Rationale: Add admission-level outcomes and demographics for downstream stratification and analysis.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
    ↪ consistent across notebook stages.

# Admissions fields
SQL["admit_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT

```

```

a.hadm_id,
a.admittime,
a.dischtime,
a.deathtime,
a.hospital_expire_flag,
a.admission_type,
a.admission_location,
a.discharge_location,
a.insurance,
a.language,
a.marital_status,
a.race AS hosp_race
FROM `{{PHYS}}.{{HOSP}}.admissions` a
JOIN hadms h USING (hadm_id)
"""

admit = run_sql_bq(sql("admit_sql"), {"hadms": hadm_list})
print("Admissions rows:", len(admit))

ed_df = ed_df.merge(admit, on="hadm_id", how="left")

# Merge ventilation flags/times (hadm-level)
if "vent_combined" in globals():
    ed_df = ed_df.merge(vent_combined, on="hadm_id", how="left")

# Derived outcomes
ed_df["hospital_los_hours"] = (ed_df["dischtime"] -
                                ed_df["admittime"]).dt.total_seconds() / 3600.0
ed_df["in_hospital_death"] = ((ed_df["hospital_expire_flag"] == 1) |
                                ed_df["deathtime"].notna()).astype("int64")

# Concordance check
discord = (
    ((ed_df["hospital_expire_flag"] == 1) & ed_df["deathtime"].isna()) |
    ((ed_df["hospital_expire_flag"] == 0) & ed_df["deathtime"].notna())
)
print("Admissions discordance (expire_flag vs deathtime):",
      int(discord.sum()))

Admissions rows: 40167
Admissions discordance (expire_flag vs deathtime): 3

```

3.0.4 Phase 3 — ICU timing and LOS

Rationale: Capture ICU exposure, timing, and total LOS for severity stratification.

```
# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↪ consistent across notebook stages.

# ICU stays (aggregate per hadm)
SQL["icu_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT
    i.hadm_id,
    i.stay_id AS icu_stay_id,
    i.intime,
    i.outtime,
    i.los,
    i.first_careunit,
    i.last_careunit
FROM `'{PHYS}`.{ICU}`.icustays` i
JOIN hadms h USING (hadm_id)
"""

icu = run_sql_bq(sql("icu_sql"), {"hadms": hadm_list})
print("ICU stay rows:", len(icu))

if len(icu) > 0:
    icu_agg = (
        icu.sort_values(["hadm_id", "intime"]).groupby("hadm_id",
↪ as_index=False)
        .agg(
            icu_intime_first=("intime", "min"),
            icu_outtime_last=("outtime", "max"),
            icu_los_total=("los", "sum"),
            n_icu_stays=("icu_stay_id", "nunique"),
            first_careunit=("first_careunit", "first"),
            last_careunit=("last_careunit", "last"),
        )
    )
    ed_df = ed_df.merge(icu_agg, on="hadm_id", how="left")
else:
    print("No ICU stays found for cohort.")

ICU stay rows: 36260
```

3.0.5 Phase 4 — ED longitudinal vitals (0–6h)

Rationale: Summarize early ED vitals for severity phenotyping.

```
# Purpose: Build ED vitals features robustly without a single large BigQuery
↪ join that can stall.

# ED vitals long + aggregates (0-6h)

# Build the ED stay map from in-memory cohort tables to keep query scope
↪ tight and deterministic.
if "ed_df" not in globals():
    raise NameError("ed_df is required before ED vitals extraction")

required_cols = ["ed_stay_id", "hadm_id", "ed_intime"]
missing_required = [c for c in required_cols if c not in ed_df.columns]
if missing_required:
    raise KeyError(f"ed_df missing required columns for ED vitals extraction:
↪ {missing_required}")

edmap_local = (
    ed_df[required_cols]
    .dropna(subset=["ed_stay_id"])
    .drop_duplicates(subset=["ed_stay_id"])
    .copy()
)
edmap_local["ed_stay_id"] = pd.to_numeric(edmap_local["ed_stay_id"],
↪ errors="coerce")
edmap_local = edmap_local.dropna(subset=["ed_stay_id"])
edmap_local["ed_stay_id"] = edmap_local["ed_stay_id"].astype(int)

stay_ids = sorted(edmap_local["ed_stay_id"].unique().tolist())
print("ED stays for vitals pull:", len(stay_ids))

SQL["ed_vitals_sql"] = f"""
SELECT
    stay_id AS ed_stay_id,
    charttime,
    temperature,
    heartrate,
    resprate,
    o2sat,
    sbp,
```

```

    dbp,
    rhythm,
    pain
FROM `{{PHYS}}.{{ED}}.vitalsign`
WHERE stay_id IN UNNEST(@stay_ids)
"""

# Query in chunks so we avoid one very large parameter payload/job.
chunk_size = 5000
vitals_chunks = []
for i in range(0, len(stay_ids), chunk_size):
    chunk = stay_ids[i:i + chunk_size]
    part = run_sql_bq(sql("ed_vitals_sql"), {"stay_ids": chunk})
    if len(part) > 0:
        part["ed_stay_id"] = pd.to_numeric(part["ed_stay_id"],
                                         errors="coerce")
        part = part.dropna(subset=["ed_stay_id"])
        part["ed_stay_id"] = part["ed_stay_id"].astype(int)
        vitals_chunks.append(part)
    print(f"ED vitals chunk {i // chunk_size + 1}: rows={len(part)})")

if vitals_chunks:
    ed_vitals_long = pd.concat(vitals_chunks, ignore_index=True)
else:
    ed_vitals_long = pd.DataFrame(
        columns=[
            "ed_stay_id", "charttime", "temperature", "heartrate",
            "resprate",
            "o2sat", "sbp", "dbp", "rhythm", "pain"
        ]
    )

ed_vitals_long = ed_vitals_long.merge(edmap_local, on="ed_stay_id",
                                       how="inner")
print("ED vitals long rows:", len(ed_vitals_long))

# Window filter: 0-6h from ED intime
ed_vitals_long["dt_hours"] = (ed_vitals_long["charttime"] -
                               ed_vitals_long["ed_intime"]).dt.total_seconds() / 3600.0
in_6h = ed_vitals_long["dt_hours"].between(0, 6, inclusive="both")

agg = (
    ed_vitals_long.loc[in_6h]

```

```

    .groupby("ed_stay_id", as_index=False)
    .agg(
        max_heartrate_0_6h=("heartrate", "max"),
        max_resprate_0_6h=("resprate", "max"),
        min_o2sat_0_6h=("o2sat", "min"),
        min_sbp_0_6h=("sbp", "min"),
        n_vitals_0_6h=("charttime", "count"),
    )
)

ed_df = ed_df.merge(agg, on="ed_stay_id", how="left")

# Range warnings (report only, do not drop)
range_checks = {
    "heartrate": (0, 300),
    "resprate": (0, 80),
    "o2sat": (0, 100),
    "sbp": (0, 300),
}
for col, (lo, hi) in range_checks.items():
    bad = ed_vitals_long[col].notna() & (~ed_vitals_long[col].between(lo,
    ↵ hi))
    if bad.any():
        print(f"Warning: {col} out of range count:", int(bad.sum()))

```

```

ED stays for vitals pull: 11786
ED vitals chunk 1: rows=30899
ED vitals chunk 2: rows=31867
ED vitals chunk 3: rows=10815
ED vitals long rows: 73581
Warning: heartrate out of range count: 1
Warning: resprate out of range count: 2
Warning: o2sat out of range count: 1

```

3.0.6 Phase 5 — Robust lab discovery + gas panels

Rationale: Capture blood gas and key chemistry labs with label-robust item discovery and unit normalization.

```
# Purpose: Discover lab item IDs robustly and record cohort-aware item
    ↵ frequency for auditability.
```

```

import re
import json

# Discover itemids from d_labitems
SQL["labitems_sql"] = f"""
SELECT itemid, label, fluid, category
FROM `{{PHYS}}.{{HOSP}}.d_labitems`"""
"""

labitems = run_sql_bq(sql("labitems_sql"))

patterns = {
    "gas_pco2": re.compile(r"\bp\bs*co2\b|pco2|pco ", re.I),
    "gas_ph": re.compile(r"\bph\b", re.I),
    "gas_hco3": re.compile(r"hc03|bicarbonate", re.I),
    "gas_lactate": re.compile(r"lactate", re.I),
    "gas_specimen": re.compile(r"specimen|source|type", re.I),
    "chem_creatinine": re.compile(r"creatinine", re.I),
    "chem_sodium": re.compile(r"\bsodium\b", re.I),
    "chem_chloride": re.compile(r"\bchloride\b", re.I),
    "chem_total_co2": re.compile(r"carbon dioxide|total co2|\bco2\b", re.I),
    "cbc_hemoglobin": re.compile(r"hemoglobin", re.I),
}

# Category filters limit false matches (for example chemistry CO2 vs
#   blood-gas PCO2)
cat_gas = re.compile(r"blood\s*gas|blood gas|arterial|venous", re.I)
cat_chem = re.compile(r"chemistry|chem|blood", re.I)
cat_cbc = re.compile(r"hematology|cbc", re.I)

manifest_lab_itemid_sets: dict[str, set[int]] = {
    "gas_pco2": set(LAB_PCO2_ITEMIDS),
    "gas_ph": set(LAB_PH_ITEMIDS),
    "gas_hco3": set(LAB_HC03_ITEMIDS),
    "gas_specimen": set(LAB_SPECIMEN_TYPE_ITEMIDS),
}
manifest_missing_rows: list[dict[str, Any]] = []
for group_name, itemids in manifest_lab_itemid_sets.items():
    if not itemids:
        continue
    present_itemids = set(
        pd.to_numeric(
            labitems.loc[labitems["itemid"].isin(itemids), "itemid"],

```

```

        errors="coerce",
    )
    .dropna()
    .astype(int)
    .tolist()
)
missing_itemids = sorted(itemids.difference(present_itemids))
if missing_itemids:
    manifest_missing_rows.append(
        {
            "group_name": group_name,
            "missing_itemids": ";".join(str(value) for value in
                ↪ missing_itemids),
            "status": "missing_from_d_labitems",
        }
    )
)

if manifest_missing_rows:
    raise ValueError(
        "Manifest lab itemids were not found in d_labitems: "
        f"{manifest_missing_rows}. Update specs/blood_gas_itemids.json."
    )

matches = {}
for name, pat in patterns.items():
    if name in manifest_lab_itemid_sets:
        itemids = manifest_lab_itemid_sets[name]
        dfm = labitems.loc[labitems["itemid"].isin(itemids)].copy()
    else:
        dfm = labitems.copy()
        dfm = dfm[dfm["label"].str.contains(pat, na=False)]
        if name.startswith("gas_"):
            dfm = dfm[dfm["category"].str.contains(cat_gas, na=False)]
        elif name.startswith("chem_"):
            dfm = dfm[dfm["category"].str.contains(cat_chem, na=False)]
        elif name.startswith("cbc_"):
            dfm = dfm[dfm["category"].str.contains(cat_cbc, na=False)]
        matches[name] = dfm[["itemid", "label", "category"]]

# Build lab_item_map with cohort counts. Run hadm counts in chunks to avoid
# one oversized query.
itemids_all = sorted({int(i) for dfm in matches.values() for i in
    ↪ dfm["itemid"].tolist()})

```

```

SQL["counts_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x)
SELECT itemid, COUNT(*) AS n
FROM `{{PHYS}}.{{HOSP}}.labevents`
WHERE hadm_id IN (SELECT hadm_id FROM hadms)
    AND itemid IN UNNEST(@itemids)
GROUP BY itemid
"""

counts_cols = ["itemid", "n"]
if itemids_all and len(hadm_list) > 0:
    hadm_chunk_size = 10000
    count_parts = []
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
        part = run_sql_bq(sql("counts_sql"), {"hadms": hadm_chunk, "itemids": itemids_all})
        if len(part) > 0:
            count_parts.append(part)
    print(f"Lab item count chunk {i // hadm_chunk_size + 1}:
        rows={len(part)}")

if count_parts:
    counts = (
        pd.concat(count_parts, ignore_index=True)
        .groupby("itemid", as_index=False)[["n"]]
        .sum()
    )
else:
    counts = pd.DataFrame(columns=counts_cols)
else:
    counts = pd.DataFrame(columns=counts_cols)

lab_item_map = {}
for name, dfm in matches.items():
    tmp = dfm.merge(counts, on="itemid", how="left").fillna({"n": 0})
    lab_item_map[name] = {
        "pattern": patterns[name].pattern,
        "items": tmp.sort_values("n",
            ascending=False).to_dict(orient="records"),
    }
}

```

```

lab_item_map_path = WORK_DIR / "lab_item_map.json"
lab_item_map_path.write_text(json.dumps(lab_item_map, indent=2))
print("Wrote:", lab_item_map_path)

manifest_lab_audit_rows: list[dict[str, Any]] = []
for analyte_name, itemids in manifest_lab_itemid_sets.items():
    if not itemids:
        continue
    frame = labitems.loc[labitems["itemid"].isin(itemids), ["itemid",
    "label", "category", "fluid"]].copy()
    frame = frame.merge(counts, on="itemid", how="left").fillna({"n": 0})
    for record in frame.to_dict(orient="records"):
        manifest_lab_audit_rows.append(
            {
                "source_branch": "LAB",
                "analyte": analyte_name,
                "itemid": int(record["itemid"]),
                "label": record.get("label"),
                "category": record.get("category"),
                "fluid": record.get("fluid"),
                "cohort_row_count": int(record.get("n", 0)),
                "manifest_version": BLOOD_GAS_MANIFEST["version"],
            }
        )
    )

blood_gas_itemid_manifest_audit = pd.DataFrame(manifest_lab_audit_rows)

Lab item count chunk 1: rows=38
Lab item count chunk 2: rows=38
Lab item count chunk 3: rows=36
Lab item count chunk 4: rows=37
Lab item count chunk 5: rows=22
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project:
CC-NLP/lab_item_map.json

# Purpose: Extract cohort labs and apply ED-time windows in pandas for
#           predictable runtime.

# Extract labevents and apply ED 0-24h windows locally

# Assemble itemid lists
itemid_sets = {k: [int(x["itemid"]) for x in v["items"]] for k, v in
    lab_item_map.items()}

```

```

all_itemids = sorted({i for v in itemid_sets.values() for i in v})

if "ed_df" not in globals():
    raise NameError("ed_df is required before lab window extraction")

ed_windows = (
    ed_df[["ed_stay_id", "hadm_id", "ed_intime"]]
    .dropna(subset=["ed_stay_id", "hadm_id", "ed_intime"])
    .drop_duplicates()
    .copy()
)
ed_windows["hadm_id"] = pd.to_numeric(ed_windows["hadm_id"],
                                       errors="coerce").astype("Int64")
ed_windows = ed_windows.dropna(subset=["hadm_id"])
ed_windows["hadm_id"] = ed_windows["hadm_id"].astype(int)

SQL["labs_sql"] = f"""
SELECT
    subject_id,
    hadm_id,
    itemid,
    charttime,
    specimen_id,
    valuenum,
    valueuom,
    value AS value_text
FROM `{{PHYS}}.{{HOSP}}.labevents`
WHERE hadm_id IN UNNEST(@hadms)
    AND itemid IN UNNEST(@itemids)
"""

# Run in hadm chunks; merge to ED windows in pandas to keep each query
# lightweight.
labs_parts = []
if all_itemids and len(hadm_list) > 0:
    hadm_chunk_size = 5000
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]
        part = run_sql_bq(sql("labs_sql"), {"hadms": hadm_chunk, "itemids": all_itemids})
        raw_n = len(part)

        if raw_n > 0:

```

```

part["hadm_id"] = pd.to_numeric(part["hadm_id"], errors="coerce")
part = part.dropna(subset=["hadm_id"])
part["hadm_id"] = part["hadm_id"].astype(int)

part = part.merge(ed_windows, on="hadm_id", how="inner")
in_window = (
    (part["charttime"] >= part["ed_intime"]) &
    (part["charttime"] <= (part["ed_intime"] +
    pd.Timedelta(hours=24)))
)
part = part.loc[in_window, [
    "ed_stay_id", "subject_id", "hadm_id", "itemid", "charttime",
    "specimen_id", "valuenum", "valueuom", "value_text"
]]
if len(part) > 0:
    labs_parts.append(part)

print(f"Labs window chunk {i // hadm_chunk_size + 1}: raw={raw_n},
      kept={len(part) if raw_n > 0 else 0}")

if labs_parts:
    labs_long = pd.concat(labs_parts, ignore_index=True)
else:
    labs_long = pd.DataFrame(
        columns=["ed_stay_id", "subject_id", "hadm_id", "itemid",
        "charttime", "specimen_id", "valuenum", "valueuom", "value_text"]
    )

print("Labs long rows:", len(labs_long))

# Unit audit for pCO2
pco2_ids = itemid_sets.get("gas_pco2", [])
unit_audit = (
    labs_long.loc[labs_long["itemid"].isin(pco2_ids)]
    .groupby("valueuom", dropna=False)
    .size().reset_index(name="n")
)
unit_audit_path = WORK_DIR / "lab_unit_audit.csv"
unit_audit.to_csv(unit_audit_path, index=False)
print("Wrote:", unit_audit_path)

```

Labs window chunk 1: raw=646473, kept=25741
Labs window chunk 2: raw=653510, kept=26117

```

Labs window chunk 3: raw=654739, kept=27046
Labs window chunk 4: raw=651910, kept=25179
Labs window chunk 5: raw=664900, kept=25348
Labs window chunk 6: raw=674175, kept=25805
Labs window chunk 7: raw=656655, kept=26628
Labs window chunk 8: raw=640120, kept=24835
Labs window chunk 9: raw=21053, kept=884
Labs long rows: 207583
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/lab_unit_audit.csv

# Purpose: Reconstruct gas panels and normalize units before ED-stay level
# aggregation.

# Reconstruct gas panels by specimen_id within ED stay

pco2_ids = set(itemid_sets.get("gas_pco2", []))
ph_ids = set(itemid_sets.get("gas_ph", []))
hco3_ids = set(itemid_sets.get("gas_hco3", []))
lact_ids = set(itemid_sets.get("gas_lactate", []))

labs = labs_long.copy()

if "fluid" in labitems.columns:
    item_fluid = (
        labitems[["itemid", "fluid"]]
        .drop_duplicates(subset=["itemid"])
        .copy()
    )
    labs = labs.merge(item_fluid, on="itemid", how="left")
else:
    labs["fluid"] = pd.NA

# Convert pCO2 kPa to mmHg when needed
is_kpa = labs["valueuom"].astype(str).str.lower().str.contains("kpa",
    na=False)
mask_pco2 = labs["itemid"].isin(pco2_ids)
if mask_pco2.any() and is_kpa.any():
    labs.loc[mask_pco2 & is_kpa, "valuenum"] = labs.loc[mask_pco2 & is_kpa,
        "valuenum"] * 7.50062
    labs.loc[mask_pco2 & is_kpa, "valueuom"] = "mmHg"

# panel by specimen_id

```

```

panel = (
    labs.groupby(["ed_stay_id", "specimen_id"], as_index=False)
    .agg(panel_time=("charttime", "min"))
)
# attach analytes

def pick_analyte(
    df: pd.DataFrame,
    ids: set[int],
    name: str,
    *,
    include_item_meta: bool = False,
) -> pd.DataFrame:
    cols = ["ed_stay_id", "specimen_id", "charttime", "valuenum", "itemid",
    ← "fluid"]
    tmp = df.loc[df["itemid"].isin(ids), cols].copy()
    tmp["valuenum"] = pd.to_numeric(tmp["valuenum"], errors="coerce")
    tmp = tmp.loc[tmp["valuenum"].notna()].copy()
    if tmp.empty:
        base_cols = ["ed_stay_id", "specimen_id", name]
        if include_item_meta:
            base_cols += [f"{name}_itemid", f"{name}_fluid"]
        return pd.DataFrame(columns=base_cols)
    tmp = tmp.sort_values(
        ["ed_stay_id", "specimen_id", "charttime", "itemid"], kind="stable"
    )
    first = tmp.groupby(["ed_stay_id", "specimen_id"],
    ← as_index=False).first()
    first = first.rename(columns={"valuenum": name})
    if include_item_meta:
        first = first.rename(
            columns={"itemid": f"{name}_itemid", "fluid": f"{name}_fluid"}
        )
        keep_cols = ["ed_stay_id", "specimen_id", name, f"{name}_itemid",
    ← f"{name}_fluid"]
    else:
        keep_cols = ["ed_stay_id", "specimen_id", name]
    return first[keep_cols]

if pco2_ids:
    panel = panel.merge(
        pick_analyte(

```

```

        labs,
        pco2_ids,
        "pco2",
        include_item_meta=True,
    ),
    on=["ed_stay_id", "specimen_id"],
    how="left",
)
if ph_ids:
    panel = panel.merge(pick_analyte(labs, ph_ids, "ph"),
    ↵ on=["ed_stay_id", "specimen_id"], how="left")
if hco3_ids:
    panel = panel.merge(pick_analyte(labs, hco3_ids, "hco3"),
    ↵ on=["ed_stay_id", "specimen_id"], how="left")
if lact_ids:
    panel = panel.merge(pick_analyte(labs, lact_ids, "lactate"),
    ↵ on=["ed_stay_id", "specimen_id"], how="left")

# First pCO2-qualified panel per ED stay (gas anchor must include pCO2)
panel_anchor_candidates = panel.loc[
    pd.to_numeric(panel.get("pco2"), errors="coerce").notna()
].copy()
if panel_anchor_candidates.empty:
    first_panel = panel.head(0).copy()
else:
    first_panel = (
        panel_anchor_candidates.sort_values(["ed_stay_id", "panel_time"])
        .groupby("ed_stay_id", as_index=False)
        .first()
    )

# 0-6h and 0-24h extrema
panel = panel.merge(ed_df[["ed_stay_id", "ed_intime"]], on="ed_stay_id",
    ↵ how="left")
panel["dt_hours"] = (panel["panel_time"] -
    ↵ panel["ed_intime"]).dt.total_seconds() / 3600.0

# Ensure expected panel columns exist even if analyte is absent
for col in ["pco2", "ph", "hco3", "lactate"]:
    if col not in panel.columns:
        panel[col] = pd.NA
for col in ["pco2itemid", "pco2fluid"]:
    if col not in panel.columns:

```

```

    panel[col] = pd.NA

p06 = panel.loc[panel["dt_hours"].between(0, 6, inclusive="both")]
p24 = panel.loc[panel["dt_hours"].between(0, 24, inclusive="both")]

agg06 = p06.groupby("ed_stay_id",
    ↪ as_index=False).agg(max_pco2_0_6h=("pco2", "max"),
    ↪ min_ph_0_6h=("ph", "min"))
agg24 = p24.groupby("ed_stay_id",
    ↪ as_index=False).agg(max_pco2_0_24h=("pco2", "max"),
    ↪ min_ph_0_24h=("ph", "min"))

first_panel_export_cols = [
    col
    for col in ["ed_stay_id", "panel_time", "pco2", "ph", "hco3", "lactate"]
    if col in first_panel.columns
]
if first_panel_export_cols:
    ed_df = ed_df.merge(
        first_panel[first_panel_export_cols].rename(
            columns={
                "panel_time": "first_gas_time",
                "pco2": "first_pco2",
                "ph": "first_ph",
                "hco3": "first_hco3",
                "lactate": "first_lactate",
            }
        ),
        on="ed_stay_id",
        how="left",
    )
else:
    ed_df["first_gas_time"] = pd.NaT
    ed_df["first_pco2"] = pd.NA
    ed_df["first_ph"] = pd.NA
    ed_df["first_hco3"] = pd.NA
    ed_df["first_lactate"] = pd.NA

# Prioritize LAB blood-gas bicarbonate aligned to pCO2 anchor panel, then
    ↪ fallback
# to nearest LAB bicarbonate around first_gas_time when same-panel value is
    ↪ missing.
if "first_hco3_source" not in ed_df.columns:

```

```

ed_df["first_hco3_source"] = "missing"

panel_hco3_present = pd.to_numeric(ed_df.get("first_hco3"),
    ↪ errors="coerce").notna()
ed_df.loc[panel_hco3_present, "first_hco3_source"] =
    ↪ "lab_panel_same_specimen"

if hco3_ids and "first_gas_time" in ed_df.columns:
    lab_hco3_long = labs.loc[
        labs["itemid"].isin(hco3_ids),
        ["ed_stay_id", "charttime", "valuenum"],
    ].copy()
    if not lab_hco3_long.empty:
        lab_hco3_long["valuenum"] = pd.to_numeric(lab_hco3_long["valuenum"], 
    ↪ errors="coerce")
        lab_hco3_long =
            lab_hco3_long.loc[lab_hco3_long["valuenum"].notna()].copy()
        anchor_missing = ed_df.loc[
            pd.to_numeric(ed_df["first_hco3"], errors="coerce").isna()
            & pd.to_datetime(ed_df["first_gas_time"]),
            ↪ errors="coerce").notna(),
            ["ed_stay_id", "first_gas_time"],
        ].copy()
        if not anchor_missing.empty and not lab_hco3_long.empty:
            fallback_candidates = anchor_missing.merge(
                lab_hco3_long,
                on="ed_stay_id",
                how="left",
            )
            fallback_candidates["first_gas_time"] = pd.to_datetime(
                fallback_candidates["first_gas_time"], errors="coerce"
            )
            fallback_candidates["charttime"] = pd.to_datetime(
                fallback_candidates["charttime"], errors="coerce"
            )
            fallback_candidates["dt_abs_hours"] = (
                (fallback_candidates["charttime"] -
            ↪ fallback_candidates["first_gas_time"])
                .dt.total_seconds()
                .abs()
                / 3600.0
            )
            fallback_candidates = fallback_candidates.loc[

```

```

        fallback_candidates["dt_abs_hours"].notna()
    ].copy()
    if not fallback_candidates.empty:
        nearest_hco3 = (
            fallback_candidates.sort_values(
                ["ed_stay_id", "dt_abs_hours", "charttime"],
                kind="stable",
            )
            .groupby("ed_stay_id", as_index=False)
            .first()
            .rename(columns={"valuenum": "first_hco3Fallback_lab"})
        )
        ed_df = ed_df.merge(
            nearest_hco3[["ed_stay_id", "first_hco3Fallback_lab"]],
            on="ed_stay_id",
            how="left",
        )
        fill_mask = (
            pd.to_numeric(ed_df["first_hco3"]),
            & pd.to_numeric(ed_df["first_hco3Fallback_lab"]),
            & errors="coerce").isna()
        )
        ed_df.loc[fill_mask, "first_hco3"] = ed_df.loc[
            fill_mask, "first_hco3Fallback_lab"
        ]
        ed_df.loc[fill_mask, "first_hco3_source"] =
    "lab_nearest_to_anchor"
    ed_df = ed_df.drop(columns=["first_hco3Fallback_lab"])

ed_df["first_gas_anchor_has_pco2"] = ed_df["first_gas_time"].notna()
ed_df["first_gas_anchor_source_validated"] = False

ed_df = ed_df.merge(agg06, on="ed_stay_id", how="left")
ed_df = ed_df.merge(agg24, on="ed_stay_id", how="left")

# Purpose: Reconstruct gas panels and normalize units before ED-stay level
# aggregation.

# ABG vs VBG classification (specimen-type rows only for core cohort
# behavior)
panel, gas_source_audit = assign_panel_source_from_specimen(
    panel,

```

```

    labs,
    specimen_source_itemids=itemid_sets.get("gas_specimen", []),
)

# Keep heuristic route diagnostics for QA transparency only (non-core
# behavior).
heuristic_gas_source_audit: dict[str, Any] | None = None
heuristic_mode = os.getenv("COHORT_GAS_SOURCE_INFERENCE_MODE",
                           "metadata_only").strip().lower()
try:
    _, heuristic_gas_source_audit = infer_panel_gas_source_metadata(
        panel_df=panel.copy(),
        labs_df=labs,
        labitems_df=labitems,
        specimen_source_itemids=itemid_sets.get("gas_specimen", []),
        pco2_itemids=itemid_sets.get("gas_pco2", []),
        mode=heuristic_mode,
    )
except Exception as exc:
    heuristic_gas_source_audit = {
        "status": "error",
        "message": f"heuristic_diagnostics_failed: {exc}",
    }
gas_source_audit["heuristic_reference"] = heuristic_gas_source_audit

fail_on_all_other_source = os.getenv("COHORT_FAIL_ON_ALL_OTHER_SOURCE",
                                     "1").strip() == "1"
if gas_source_audit["all_other_or_unknown"]:
    if fail_on_all_other_source:
        assert_gas_source_coverage(
            gas_source_audit,
            fail_on_all_other_source=fail_on_all_other_source,
        )
    else:
        print(
            "WARNING: Gas source attribution collapsed to other/unknown; "
            "continuing because COHORT_FAIL_ON_ALL_OTHER_SOURCE=0."
        )
warn_other_rate_threshold = float(os.getenv("COHORT_WARN_OTHER_RATE",
                                           "0.50"))
fail_other_rate_raw = os.getenv("COHORT_FAIL_OTHER_RATE", "").strip()

```

```

fail_other_rate_threshold = float(fail_other_rate_raw) if fail_other_rate_raw
↪ else None
other_source_rate = float(gas_source_audit.get("source_rates",
↪ {}).get("other", 0.0))

if fail_other_rate_threshold is not None and other_source_rate >=
↪ fail_other_rate_threshold:
    raise ValueError(
        "gas_source_other_rate exceeded configured fail threshold "
        f"{fail_other_rate_threshold:.3f}: {other_source_rate:.4f}."
    )
if other_source_rate >= warn_other_rate_threshold:
    print(
        "WARNING: gas_source_other_rate exceeded warning threshold "
        f"{warn_other_rate_threshold:.3f}: {other_source_rate:.4f}."
    )

print("Gas source rates:", gas_source_audit.get("source_rates", {}))
print("Gas source inference tiers:", gas_source_audit.get("tier_rates", {}))
print("Unresolved specimen IDs:",
↪ gas_source_audit.get("unresolved_specimen_id_count", 0))

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")
gas_source_audit_path = prior_runs_dir / f"{out_date} gas_source_audit.json"

reduction_target = float(os.getenv("COHORT_OTHER_RELATIVE_REDUCTION_MIN",
↪ "0.10"))
baseline_other_rate = None
baseline_audit_path = None
for candidate in sorted(prior_runs_dir.glob("* gas_source_audit.json"),
↪ reverse=True):
    if candidate.resolve() == gas_source_audit_path.resolve():
        continue
    try:
        baseline_payload = json.loads(candidate.read_text())
        baseline_other_rate = float(
            baseline_payload.get("source_rates", {}).get("other",
↪ float("nan")))
        if math.isnan(baseline_other_rate):
            baseline_other_rate = None

```

```

        continue
    baseline_audit_path = candidate
    break
except Exception:
    continue

if baseline_other_rate is None or baseline_other_rate <= 0:
    gas_source_audit["baseline_other_rate"] = None
    gas_source_audit["baseline_path"] = str(baseline_audit_path) if
    ↪ baseline_audit_path else None
    gas_source_audit["relative_reduction_vs_baseline"] = None
    print(
        "WARNING: Unable to evaluate gas_source_other_rate relative reduction
        ↪ "
        "because no valid baseline gas_source_audit was found."
    )
else:
    relative_reduction = (baseline_other_rate - other_source_rate) /
    ↪ baseline_other_rate
    gas_source_audit["baseline_other_rate"] = baseline_other_rate
    gas_source_audit["baseline_path"] = str(baseline_audit_path)
    gas_source_audit["relative_reduction_vs_baseline"] = relative_reduction
    gas_source_audit["relative_reduction_target"] = reduction_target
    print(
        "Gas source OTHER relative reduction vs baseline: "
        f"{relative_reduction:.4f} (target={reduction_target:.4f})"
    )
    if relative_reduction < reduction_target:
        message = (
            "gas_source_other_rate relative reduction target not met: "
            f"baseline={baseline_other_rate:.4f},
            ↪ current={other_source_rate:.4f}, "
            f"reduction={relative_reduction:.4f},
            ↪ target={reduction_target:.4f}."
        )
        contract_mode = os.getenv("PIPELINE_CONTRACT_MODE",
        ↪ "fail").strip().lower()
        if contract_mode == "fail":
            raise ValueError(message)
        print(f"WARNING: {message}")

gas_source_audit_path.write_text(json.dumps(gas_source_audit, indent=2))
print("Wrote:", gas_source_audit_path)

```

```

panel["pco2_source_branch"] = "LAB"
panel_anchor = panel.loc[pd.to_numeric(panel["pco2"],
    ↪ errors="coerce").notna()].copy()
if panel_anchor.empty:
    first_anchor = ed_df[["ed_stay_id"]].drop_duplicates().copy()
    first_anchor["first_gas_time"] = pd.NaT
    first_anchor["first_pco2"] = pd.NA
    first_anchor["first_ph"] = pd.NA
    first_anchor["first_hco3"] = pd.NA
    first_anchor["first_lactate"] = pd.NA
    first_anchor["first_gas_specimen_type"] = pd.NA
    first_anchor["first_gas_specimen_present"] = False
    first_anchor["first_gas_pco2_itemid"] = pd.NA
    first_anchor["first_gas_pco2_fluid"] = pd.NA
    first_anchor["co2_other_is_blood asserted"] = False
    first_anchor["first_gas_anchor_has_pco2"] = False
    first_anchor["first_gas_anchor_source_validated"] = False
else:
    first_anchor = (
        panel_anchor.sort_values(["ed_stay_id", "panel_time"])
            .groupby("ed_stay_id", as_index=False)
            .first()
    )
    first_anchor = first_anchor.rename(
        columns={
            "panel_time": "first_gas_time",
            "pco2": "first_pco2",
            "ph": "first_ph",
            "hco3": "first_hco3",
            "lactate": "first_lactate",
            "source_hint_specimen_text": "first_gas_specimen_type",
            "pco2_itemid": "first_gas_pco2_itemid",
            "pco2_fluid": "first_gas_pco2_fluid",
        }
    )
    first_anchor["first_gas_specimen_present"] = (
        first_anchor["first_gas_specimen_type"]
            .astype("string")
            .fillna("")
            .str.strip()
            .ne("")
    )
)

```

```

first_anchor["co2_other_is_blood_asserted"] = (
    first_anchor["source"].astype("string").str.lower().eq("other")
    & first_anchor["first_gas_pco2_fluid"]
    .astype("string")
    .str.strip()
    .str.lower()
    .isin(LAB_BLOOD_FLUID_TERMS)
)
first_anchor["first_gas_anchor_has_pco2"] = True
first_anchor["first_gas_anchor_source_validated"] =
    first_anchor["source"].isin(
        {"arterial", "venous", "other"}
    )

anchor_merge_cols = [
    "ed_stay_id",
    "first_gas_time",
    "first_pco2",
    "first_ph",
    "first_hco3",
    "first_lactate",
    "first_gas_specimen_type",
    "first_gas_specimen_present",
    "first_gas_pco2_itemid",
    "first_gas_pco2_fluid",
    "co2_other_is_blood_asserted",
    "first_gas_anchor_has_pco2",
    "first_gas_anchor_source_validated",
]
drop_existing_anchor_cols = [col for col in anchor_merge_cols if col in
    ed_df.columns and col != "ed_stay_id"]
if drop_existing_anchor_cols:
    ed_df = ed_df.drop(columns=drop_existing_anchor_cols)
ed_df = ed_df.merge(first_anchor[anchor_merge_cols], on="ed_stay_id",
    how="left")
ed_df["first_gas_anchor_has_pco2"] = (
    ed_df["first_gas_anchor_has_pco2"].fillna(False).astype(bool)
)
ed_df["first_gas_anchor_source_validated"] = (
    ed_df["first_gas_anchor_source_validated"].fillna(False).astype(bool)
)
ed_df["first_gas_specimen_present"] = (
    ed_df["first_gas_specimen_present"].fillna(False).astype(bool)
)

```

```

)
ed_df["co2_other_is_blood_asserted"] = (
    ed_df["co2_other_is_blood_asserted"].fillna(False).astype(bool)
)

# flags
panel["flag_abg_hypercapnia"] = ((panel["source"]=="arterial") &
    (panel["pcO2"]>=45)).astype(int)
panel["flag_vbg_hypercapnia"] = ((panel["source"]=="venous") &
    (panel["pcO2"]>=50)).astype(int)
panel["flag_other_hypercapnia"] = ((panel["source"]=="other") &
    (panel["pcO2"]>=50)).astype(int)
panel["flag_any_gas_hypercapnia"] = ((panel["flag_abg_hypercapnia"]==1) |
    (panel["flag_vbg_hypercapnia"]==1) |
    (panel["flag_other_hypercapnia"]==1)).astype(int)

# per-stay unknown source rate
if len(panel) > 0:
    unk_rate = (
        panel.assign(_unk=(panel["source"].fillna("other") == "other"))
            .groupby("ed_stay_id", as_index=False)[["_unk"]].mean()
            .rename(columns={"_unk": "gas_source_other_rate"})
    )
else:
    unk_rate = panel[["ed_stay_id"]].drop_duplicates()
    unk_rate["gas_source_other_rate"] = 1.0

if len(panel) > 0:
    panel_for_diag = panel.copy()
    panel_for_diag["source_inference_tier"] = (
        panel_for_diag["source_inference_tier"]
            .astype("string")
            .fillna("fallback_other")
    )
    panel_for_diag["source_hint_conflict"] = (
        panel_for_diag["source_hint_conflict"]
            .fillna(False)
            .astype(bool)
    )
    panel_for_diag["source"] = (
        panel_for_diag["source"].astype("string").fillna("other").str.lower()
    )
    panel_for_diag["source_resolved"] = panel_for_diag["source"].isin(

```

```

        {"arterial", "venous"}
    )
panel_for_diag["tier_specimen_text"] =
↪ panel_for_diag["source_inference_tier"].eq(
    "specimen_text"
)
panel_for_diag["tier_label_fluid"] =
↪ panel_for_diag["source_inference_tier"].eq(
    "label_fluid"
)
panel_for_diag["tier_cluster_inheritance"] = panel_for_diag[
    "source_inference_tier"
].eq("cluster_inheritance")
panel_for_diag["tier_panel_cooccurrence"] = panel_for_diag[
    "source_inference_tier"
].eq("panel_cooccurrence")
panel_for_diag["tierFallback_other"] =
↪ panel_for_diag["source_inference_tier"].eq(
    "fallback_other"
)

def _tier_mode(series: pd.Series) -> str:
    mode_values = series.mode(dropna=True)
    if mode_values.empty:
        return "fallback_other"
    return str(mode_values.iloc[0])

source_diag = (
    panel_for_diag.groupby("ed_stay_id", as_index=False)
    .agg(
        gas_source_inference_primary_tier=("source_inference_tier",
↪ _tier_mode),
        gas_source_hint_conflict_rate=("source_hint_conflict", "mean"),
        gas_source_resolved_rate=("source_resolved", "mean"),
        gas_source_tier_specimen_text_rate=("tier_specimen_text",
↪ "mean"),
        gas_source_tier_label_fluid_rate=("tier_label_fluid", "mean"),
↪
        gas_source_tier_cluster_inheritance_rate=("tier_cluster_inheritance",
↪ "mean"),
↪
        gas_source_tier_panel_cooccurrence_rate=("tier_panel_cooccurrence",
↪ "mean"),
↪
    )
)

```

```

        gas_source_tierFallback_other_rate="tierFallback_other",
        "mean"),
    )
    .copy()
)
else:
    source_diag = ed_df[["ed_stay_id"]].drop_duplicates().copy()
    source_diag["gas_source_inference_primary_tier"] = "fallback_other"
    source_diag["gas_source_hint_conflict_rate"] = 0.0
    source_diag["gas_source_resolved_rate"] = 0.0
    source_diag["gas_source_tier_specimen_text_rate"] = 0.0
    source_diag["gas_source_tier_label_fluid_rate"] = 0.0
    source_diag["gas_source_tier_cluster_inheritance_rate"] = 0.0
    source_diag["gas_source_tier_panel_cooccurrence_rate"] = 0.0
    source_diag["gas_source_tierFallback_other_rate"] = 1.0

# collapse to ED stay flags
flags = panel.groupby("ed_stay_id", as_index=False).agg(
    flag_abg_hypercapnia=("flag_abg_hypercapnia", "max"),
    flag_vbg_hypercapnia=("flag_vbg_hypercapnia", "max"),
    flag_other_hypercapnia=("flag_other_hypercapnia", "max"),
    flag_any_gas_hypercapnia=("flag_any_gas_hypercapnia", "max"),
)
ed_df = ed_df.merge(flags, on="ed_stay_id", how="left")
ed_df = ed_df.merge(unk_rate, on="ed_stay_id", how="left")
ed_df = ed_df.merge(source_diag, on="ed_stay_id", how="left")
ed_df["gas_source_other_rate"] = ed_df["gas_source_other_rate"].fillna(1.0)

WARNING: gas_source_other_rate exceeded warning threshold 0.500: 0.6618.
Gas source rates: {'other': 0.6617519019140463, 'arterial': 0.20549235666062432, 'venous': 0
Gas source inference tiers: {'fallback_other': 0.6617519019140463, 'specimen_text': 0.338248
Unresolved specimen IDs: 46332
Gas source OTHER relative reduction vs baseline: -0.8256 (target=0.0000)
WARNING: gas_source_other_rate relative reduction target not met: baseline=0.3625, current=0
0.8256, target=0.0000.
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 gas_source_audit.json

```

3.0.7 Phase 5C — ICU POC blood gases (chartevents, optional)

Rationale: Capture ICU point-of-care gases if central lab labevents miss early measurements.

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↳ consistent across notebook stages.

# Discover ICU POC itemids from d_items

SQL["ditems_sql"] = f"""
SELECT itemid, label, category
FROM `{{PHYS}}.{{ICU}}.d_items`
"""

ditems = run_sql_bq(sql("ditems_sql"))

icu_patterns = {
    "pco2": re.compile(r"(?:^|[^a-z])p(?:a)?\s*co(?:2|)(?:[a-z]|$)", re.I),
    "ph": re.compile(r"\bph\b", re.I),
    "hco3": re.compile(r"hc03|bicarbonate", re.I),
    "lactate": re.compile(r"lactate", re.I),
    "specimen": re.compile(r"specimen|source|type", re.I),
}

icu_cat = re.compile(r"blood\s*gas|blood gas|resp|arterial|venous", re.I)
icu_pco2_exclusion = re.compile(ICU_LABEL_DENY_REGEX, re.I)

icu_matches = []
manifest_icu_itemids = {
    "pco2": set(ICU_PC02_ITEMIDS),
    "ph": set(ICU_PH_ITEMIDS),
    "hco3": set(ICU_HC03_ITEMIDS),
    "specimen": set(ICU_SPECIMEN_TYPE_ITEMIDS),
}
for name, pat in icu_patterns.items():
    manifest_ids = manifest_icu_itemids.get(name, set())
    if manifest_ids:
        dfm = ditems.loc[ditems["itemid"].isin(manifest_ids)].copy()
    else:
        dfm = pd.DataFrame(columns=["itemid", "label", "category"])

    allowFallbackForAnalyte = ICU_ALLOW_PATTERN_FALLBACK and name !=
    ↳ "hco3"
    if dfm.empty and allowFallbackForAnalyte:
        dfm = ditems[ditems["label"].str.contains(pat, na=False)].copy()
        if name in {"pco2", "ph", "specimen"}:
            dfm = dfm[dfm["category"].str.contains(icu_cat, na=False)]

```

```

if name == "pco2" and len(dfm) > 0:
    dfm = dfm[~dfm["label"].str.contains(icu_pco2_exclusion, na=False)]
icu_matches[name] = dfm[["itemid", "label", "category"]]

icu_itemid_map_frames = [
    frame.assign(analyte=name)
    for name, frame in icu_matches.items()
    if not frame.empty
]
if not icu_itemid_map_frames:
    icu_itemid_map = pd.DataFrame(columns=["analyte", "itemid", "label",
                                           "category"])
else:
    icu_itemid_map = pd.concat(icu_itemid_map_frames, ignore_index=True)
    icu_itemid_map = (
        icu_itemid_map[["analyte", "itemid", "label", "category"]]
        .drop_duplicates()
        .sort_values(["analyte", "itemid"], kind="stable")
        .reset_index(drop=True)
    )

icu_itemids = sorted({int(i) for dfm in icu_matches.values() for i in
                     dfm["itemid"].tolist()})
print("ICU POC candidate itemids:", len(icu_itemids))
if not icu_itemid_map.empty:
    print("ICU POC itemid map preview:")
    print(icu_itemid_map.head(20).to_string(index=False))

audit_date = pd.Timestamp.now().strftime("%Y-%m-%d")
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
icu_itemid_map_path = prior_runs_dir / f"{audit_date}_icu_poc_itemid_map.csv"
icu_itemid_map.to_csv(icu_itemid_map_path, index=False)
print("Wrote:", icu_itemid_map_path)

manifest_icu_audit_rows: list[dict[str, Any]] = []
for analyte_name, itemids in manifest_icu_itemids.items():
    if not itemids:
        continue
    frame = ditems.loc[ditems["itemid"].isin(itemids), ["itemid", "label",
                                                       "category"]].copy()

```

```

present_ids = set(frame["itemid"].astype(int).tolist()) if not
↪ frame.empty else set()
for itemid in sorted(itemids):
    row = {
        "source_branch": "ICU",
        "analyte": analyte_name,
        "itemid": int(itemid),
        "label": None,
        "category": None,
        "fluid": None,
        "cohort_row_count": 0,
        "manifest_version": BLOOD_GAS_MANIFEST["version"],
        "present_in_dictionary": int(itemid in present_ids),
    }
    if itemid in present_ids:
        record = frame.loc[frame["itemid"].eq(itemid)].iloc[0]
        row["label"] = record.get("label")
        row["category"] = record.get("category")
    manifest_icu_audit_rows.append(row)

if "blood_gas_itemid_manifest_audit" in globals() and not
↪ blood_gas_itemid_manifest_audit.empty:
    manifest_audit_combined = pd.concat(
        [blood_gas_itemid_manifest_audit,
    ↪ pd.DataFrame(manifest_icu_audit_rows)],
        ignore_index=True,
        sort=False,
    )
else:
    manifest_audit_combined = pd.DataFrame(manifest_icu_audit_rows)

blood_gas_itemid_manifest_audit = manifest_audit_combined
blood_gas_itemid_manifest_audit_path = prior_runs_dir / f"{audit_date}"
↪ blood_gas_itemid_manifest_audit.csv"
blood_gas_itemid_manifest_audit.to_csv(blood_gas_itemid_manifest_audit_path,
↪ index=False)
print("Wrote:", blood_gas_itemid_manifest_audit_path)

```

ICU POC candidate itemids: 0

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 icu_poc_itemid_map.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 blood_gas_itemid_manifest_audit.csv

```

# Purpose: Define a reusable BigQuery execution helper so SQL calls are
↳ consistent across notebook stages.

# Extract ICU chartevents within ED 0-24h window for cohort ICU stays

if len(icu_itemids) == 0:
    icu_poc_long = pd.DataFrame()
    print("No ICU POC itemids found.")
else:
    # ensure icu stay ids available
    if 'icu' not in globals():
        raise NameError("ICU stays table 'icu' not found; run ICU phase
            ↳ first.")

icu_stays = icu[["icu_stay_id", "hadm_id", "intime"]].copy()
icu_stays = icu_stays.dropna(subset=["icu_stay_id"])

icu_poc_sql = f"""
WITH icu_stays AS (
    SELECT stay_id AS icu_stay_id, hadm_id
    FROM `{{PHYS}}.{{ICU}}.icustays`
    WHERE stay_id IN UNNEST(@icu_stay_ids)
),
eds AS (
    SELECT stay_id AS ed_stay_id, hadm_id, intime AS ed_intime
    FROM `{{PHYS}}.{{ED}}.edstays`
    WHERE hadm_id IN (SELECT hadm_id FROM icu_stays)
)
SELECT
    s.icu_stay_id,
    e.ed_stay_id,
    e.hadm_id,
    ce.charttime,
    ce.itemid,
    ce.valuenum,
    ce.valueuom
FROM `{{PHYS}}.{{ICU}}.chartevents` ce
JOIN icu_stays s ON s.icu_stay_id = ce.stay_id
JOIN eds e ON e.hadm_id = s.hadm_id
WHERE ce.itemid IN UNNEST(@itemids)
    AND ce.charttime BETWEEN e.ed_intime AND TIMESTAMP_ADD(e.ed_intime,
    ↳ INTERVAL 24 HOUR)
"""

```

```

icu_poc_long = run_sql_bq(icu_poc_sql, {"icu_stay_ids":
    ↵ icu_stays["icu_stay_id"].astype(int).tolist(), "itemids": icu_itemids})
print("ICU POC long rows:", len(icu_poc_long))

# Build panels by 5-minute bins per ICU stay
if len(icu_poc_long) > 0:
    icu_poc_long["time_bin"] = icu_poc_long["charttime"].dt.floor("5min")

def pick_from_ids(df, ids, name):
    tmp = df.loc[df["itemid"].isin(ids),
    ↵ ["icu_stay_id", "time_bin", "valuenum"]]
    tmp = tmp.rename(columns={"valuenum": name})
    return tmp.groupby(["icu_stay_id", "time_bin"],
        ↵ as_index=False).first()

pco2_ids = set(icu_matches.get("pco2", pd.DataFrame()).get("itemid",
    ↵ []).tolist())
ph_ids = set(icu_matches.get("ph", pd.DataFrame()).get("itemid",
    ↵ []).tolist())
hco3_ids = set(icu_matches.get("hco3", pd.DataFrame()).get("itemid",
    ↵ []).tolist())
lact_ids = set(icu_matches.get("lactate", pd.DataFrame()).get("itemid",
    ↵ []).tolist())

panel_poc = (
    icu_poc_long.groupby(["icu_stay_id", "time_bin"], as_index=False)
        .agg(panel_time=("charttime", "min"))
)
if pco2_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, pco2_ids,
    ↵ "pco2"), on=["icu_stay_id", "time_bin"], how="left")
if ph_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, ph_ids,
    ↵ "ph"), on=["icu_stay_id", "time_bin"], how="left")
if hco3_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, hco3_ids,
    ↵ "hco3"), on=["icu_stay_id", "time_bin"], how="left")
if lact_ids:
    panel_poc = panel_poc.merge(pick_from_ids(icu_poc_long, lact_ids,
    ↵ "lactate"), on=["icu_stay_id", "time_bin"], how="left")

# map to ED stay via hadm_id

```

```

panel_poc = panel_poc.merge(icu[["icu_stay_id", "hadm_id"]],  

    ↵ on="icu_stay_id", how="left")  

panel_poc = panel_poc.merge(ed_df[["ed_stay_id", "hadm_id", "ed_intime"]],  

    ↵ on="hadm_id", how="left")  

panel_poc["dt_hours"] = (panel_poc["panel_time"] -  

    ↵ panel_poc["ed_intime"]).dt.total_seconds() / 3600.0

icu_itemid_usage = (  

    icu_poc_long.groupby("itemid", as_index=False)  

    .size()  

    .rename(columns={"size": "row_count"})  

)  

icu_itemid_usage["itemid"] = pd.to_numeric(  

    icu_itemid_usage["itemid"], errors="coerce"  

).astype("Int64")  

if not icu_itemid_map.empty:  

    icu_itemid_usage = icu_itemid_usage.merge(  

        icu_itemid_map[["analyte", "itemid", "label",  

    ↵ "category"]].drop_duplicates(),  

        on="itemid",  

        how="left",  

    )  

icu_itemid_usage = icu_itemid_usage.sort_values(  

    ["row_count", "itemid"], ascending=[False, True], kind="stable"  

).reset_index(drop=True)  

icu_itemid_usage_path = prior_runs_dir / f"{audit_date}  

    ↵ icu_poc_itemid_usage.csv"  

icu_itemid_usage.to_csv(icu_itemid_usage_path, index=False)  

print("Wrote:", icu_itemid_usage_path)  

if "blood_gas_itemid_manifest_audit" in globals() and not  

    ↵ blood_gas_itemid_manifest_audit.empty:  

    usage_merge = icu_itemid_usage[["itemid", "row_count"]].copy()  

    usage_merge["itemid"] = pd.to_numeric(usage_merge["itemid"],  

    ↵ errors="coerce").astype("Int64")  

    manifest_audit = blood_gas_itemid_manifest_audit.copy()  

    manifest_audit["itemid"] = pd.to_numeric(manifest_audit["itemid"],  

    ↵ errors="coerce").astype("Int64")  

    manifest_audit = manifest_audit.merge(  

        usage_merge.rename(columns={"row_count": "icu_usage_row_count"}),  

        on="itemid",  

        how="left",  

    )  

    if "cohort_row_count" in manifest_audit.columns:

```

```

        manifest_audit["cohort_row_count"] =
    ↵ manifest_audit["cohort_row_count"].fillna(0).astype(int)
        manifest_audit["cohort_row_count"] = (
            manifest_audit["cohort_row_count"] +
    ↵ manifest_audit["icu_usage_row_count"].fillna(0).astype(int)
        )
        manifest_audit = manifest_audit.drop(columns=["icu_usage_row_count"])
        blood_gas_itemid_manifest_audit = manifest_audit

    ↵ blood_gas_itemid_manifest_audit.to_csv(blood_gas_itemid_manifest_audit_path,
    ↵ index=False)
        print("Updated:", blood_gas_itemid_manifest_audit_path)
else:
    panel_poc = pd.DataFrame()

# Ensure expected panel_poc columns exist even if analyte is absent
for col in ["pco2", "ph", "hco3", "lactate"]:
    if col not in panel_poc.columns:
        panel_poc[col] = pd.NA

poc_inclusion_enabled = False
poc_inclusion_reason = "disabled_manifest_or_qc_fail"
pco2_itemid_qc_rows: list[dict[str, Any]] = []

if len(icu_poc_long) > 0 and pco2_ids:
    pco2_long =
    ↵ icu_poc_long.loc[icu_poc_long["itemid"].isin(pco2_ids)].copy()
    if not pco2_long.empty:
        pco2_long["label"] = pco2_long["itemid"].map(
            icu_itemid_map.set_index("itemid")["label"] if not
    ↵ icu_itemid_map.empty else {}
        )
        pco2_long["category"] = pco2_long["itemid"].map(
            icu_itemid_map.set_index("itemid")["category"] if not
    ↵ icu_itemid_map.empty else {}
        )
        pco2_long["valuenum"] = pd.to_numeric(pco2_long["valuenum"],
    ↵ errors="coerce")
        pco2_long["valueuom_norm"] = (
    ↵ pco2_long["valueuom"].astype("string").fillna("").str.strip().str.lower()
        )
        is_kpa = pco2_long["valueuom_norm"].eq("kpa")

```

```

    pco2_long.loc[is_kpa, "valuenum"] = pco2_long.loc[is_kpa, "valuenum"]

    ↵ * 7.50062
    pco2_long["pco2_mmhg"] = pco2_long["valuenum"]

        for itemid_value, group in pco2_long.groupby("itemid", dropna=False):
            label_text =
    ↵ str(group["label"].dropna().astype(str).head(1).tolist()[0]) if
    ↵ group["label"].notna().any() else ""
            category_text =
    ↵ str(group["category"].dropna().astype(str).head(1).tolist()[0]) if
    ↵ group["category"].notna().any() else ""
            label_text_l = label_text.lower()
            value_units =
    ↵ sorted(set(group["valueuom_norm"].dropna().astype(str).tolist()))
            unit_ok = set(value_units).issubset(ICU_PC02_QC_ALLOWED_UOMS)
            include_ok = any(token in label_text_l for token in
    ↵ ICU_PC02_QC_INCLUDE_TOKENS)
            exclude_hit = any(token in label_text_l for token in
    ↵ ICU_PC02_QC_EXCLUDE_TOKENS)
            numeric = pd.to_numeric(group["pco2_mmhg"],
    ↵ errors="coerce").dropna()
            median_value = float(numeric.median()) if len(numeric) else
    ↵ np.nan
            median_ok = bool(
                len(numeric)
                and ICU_PC02_QC_MEDIAN_MIN <= median_value <=
    ↵ ICU_PC02_QC_MEDIAN_MAX
            )
            contamination_flag = bool((not include_ok) or exclude_hit or
    ↵ (not unit_ok))

    pco2_itemid_qc_rows.append(
    {
        "itemid": int(itemid_value),
        "label": label_text,
        "category": category_text,
        "row_count": int(len(group)),
        "non_null_count": int(len(numeric)),
        "units_observed": ";" . join(value_units),
        "p50_mmhg": median_value if len(numeric) else np.nan,
        "p95_mmhg": float(numeric.quantile(0.95)) if len(numeric)
    ↵ else np.nan,

```

```

        "max_mmhg": float(numeric.max()) if len(numeric) else
        ↪ np.nan,
        "label_include_ok": bool(include_ok),
        "label_exclude_hit": bool(exclude_hit),
        "unit_ok": bool(unit_ok),
        "median_in_expected_range": bool(median_ok),
        "contamination_flag": contamination_flag,
        "manifest_whitelisted": int(itemid_value in
        ↪ set(ICU_PC02_ITEMIDS)),
    }
)

pco2_itemid_qc_audit = pd.DataFrame(pco2_itemid_qc_rows)
pco2_itemid_qc_audit_path = prior_runs_dir / f"{audit_date}"
↪ pco2_itemid_qc_audit.csv"
pco2_itemid_qc_audit.to_csv(pco2_itemid_qc_audit_path, index=False)
print("Wrote:", pco2_itemid_qc_audit_path)

if ICU_INCLUSION_MODE == "validated_only":
    if ICU_PC02_ITEMIDS and not pco2_itemid_qc_audit.empty:
        whitelist_mask =
↪ pco2_itemid_qc_audit["itemid"].isin(set(ICU_PC02_ITEMIDS))
        whitelist_rows = pco2_itemid_qc_audit.loc[whitelist_mask].copy()
        if (
            len(whitelist_rows) == len(set(ICU_PC02_ITEMIDS))
            and whitelist_rows["contamination_flag"].eq(False).all()
            and whitelist_rows["median_in_expected_range"].eq(True).all()
        ):
            poc_inclusion_enabled = True
            poc_inclusion_reason = "validated_itemids_passed_qc"
        else:
            poc_inclusion_reason = "validated_itemids_failed_qc"
    elif not ICU_PC02_ITEMIDS:
        poc_inclusion_reason = "manifest_has_no_icu_pco2_itemids"
    else:
        poc_inclusion_reason = "missing_itemid_qc_rows"
else:
    poc_inclusion_reason = "icu_inclusion_mode_disabled"

print(
    "POC inclusion status:",
    {"enabled": int(poc_inclusion_enabled), "reason": poc_inclusion_reason},
)

```

```

poc_flags = pd.DataFrame(columns=["ed_stay_id",
    ↵ "flag_any_gas_hypercapnia_poc"])
if len(panel_poc) > 0 and poc_inclusion_enabled:
    p24_poc = panel_poc.loc[panel_poc["dt_hours"].between(0,24,
    ↵ inclusive="both")].copy()
    p24_poc.loc[:, "flag_any_gas_hypercapnia_poc"] = (p24_poc["pco2"] >=
    ↵ 50).astype("Int64")

    poc_flags = p24_poc.groupby("ed_stay_id", as_index=False).agg(
        flag_any_gas_hypercapnia_poc=("flag_any_gas_hypercapnia_poc", "max"))
    )

# Avoid duplicate column on re-run
if "flag_any_gas_hypercapnia_poc" in ed_df.columns:
    ed_df = ed_df.drop(columns=["flag_any_gas_hypercapnia_poc"])

if len(poc_flags) > 0:
    ed_df = ed_df.merge(poc_flags, on="ed_stay_id", how="left")
else:
    ed_df["flag_any_gas_hypercapnia_poc"] = 0
ed_df["flag_any_gas_hypercapnia_poc"] = (
    pd.to_numeric(ed_df["flag_any_gas_hypercapnia_poc"], errors="coerce")
    .fillna(0)
    .astype("Int64")
)

ed_df["poc_inclusion_enabled"] = bool(poc_inclusion_enabled)
ed_df["poc_inclusion_reason"] = str(poc_inclusion_reason)

# Optional fallback for first_hco3 from ICU POC only when explicit manifest
    ↵ itemids are provided.
if ICU_HCO3_ITEMIDS and len(panel_poc) > 0:
    poc_hco3_candidates = panel_poc.loc[
        panel_poc["dt_hours"].between(0, 24, inclusive="both")]
    ].copy()
    poc_hco3_candidates["hco3"] = pd.to_numeric(poc_hco3_candidates["hco3"],
    ↵ errors="coerce")
    poc_hco3_candidates =
    ↵ poc_hco3_candidates.loc[poc_hco3_candidates["hco3"].notna()].copy()
    if not poc_hco3_candidates.empty:
        first_poc_hco3 = (
            poc_hco3_candidates.sort_values(

```

```

        ["ed_stay_id", "dt_hours", "panel_time"],
        kind="stable",
    )
    .groupby("ed_stay_id", as_index=False)
    .first()[["ed_stay_id", "hco3"]]
    .rename(columns={"hco3": "first_hco3_poc_fallback"})
)
ed_df = ed_df.merge(first_poc_hco3, on="ed_stay_id", how="left")
hco3_missing_mask = pd.to_numeric(ed_df.get("first_hco3"),
↪ errors="coerce").isna()
hco3_fill_mask = hco3_missing_mask & pd.to_numeric(
    ed_df["first_hco3_poc_fallback"], errors="coerce"
).notna()
ed_df.loc[hco3_fill_mask, "first_hco3"] = ed_df.loc[
    hco3_fill_mask, "first_hco3_poc_fallback"
]
if "first_hco3_source" not in ed_df.columns:
    ed_df["first_hco3_source"] = "missing"
ed_df.loc[hco3_fill_mask, "first_hco3_source"] =
↪ "poc_explicit_itemid_fallback"
ed_df = ed_df.drop(columns=["first_hco3_poc_fallback"])

# incremental yield
base = ed_df.get("flag_any_gas_hypercapnia", pd.Series(0,
↪ index=ed_df.index)).fillna(0).astype(int)
poc = ed_df.get("flag_any_gas_hypercapnia_poc", pd.Series(0,
↪ index=ed_df.index)).fillna(0).astype(int)
inc = ((base == 0) & (poc == 1)).sum()
print("ICU POC incremental hypercapnia cases (ED stays):", int(inc))

# optional export
if len(panel_poc) > 0:
    from datetime import datetime

    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    poc_path = prior_runs_dir /
↪ f"gas_panels_poc_{datetime.now().strftime('%Y%m%d_%H%M%S')}.parquet"
    panel_poc.to_parquet(poc_path, index=False)
    print("Wrote:", poc_path)

```

No ICU POC itemids found.

```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
CC-NLP/MIMIC tabular data/prior runs/2026-02-20 pco2_itemid_qc_audit.csv
POC inclusion status: {'enabled': 0, 'reason': 'manifest_has_no_icu_pco2_itemids'}
ICU POC incremental hypercapnia cases (ED stays): 0

```

3.0.8 Phase 6 — BMI/anthropometrics (OMR)

Rationale: Add BMI/height/weight closest to ED presentation when available.

```

# Purpose: Attach closest anthropometrics with pre-ED preference and bounded
↪ post-ED fallback.

from datetime import datetime

omr_sql = f"""
SELECT subject_id, chartdate, result_name, result_value
FROM `'{PHYS}`.{HOSP}`.omr`
WHERE LOWER(result_name) IN ('bmi','height','weight')
"""

allow_omr_query_failure = os.getenv("COHORT_ALLOW_OMR_QUERY_FAILURE",
↪ "0").strip() == "1"
try:
    omr_raw = run_sql_bq(omr_sql)
except Exception as exc:
    if allow_omr_query_failure:
        print("OMR query failed; continuing with empty OMR frame because
↪ COHORT_ALLOW_OMR_QUERY_FAILURE=1:", exc)
        omr_raw = pd.DataFrame(columns=["subject_id", "chartdate",
↪ "result_name", "result_value"])
    else:
        raise

print("OMR rows:", len(omr_raw))
omr_prepared = prepare_omr_records(omr_raw)
ed_df, omr_diagnostics = attach_closest_pre_ed_omr(ed_df, omr_prepared,
↪ window_days=365)
omr_diagnostics["source_rows_raw"] = int(len(omr_raw))
omr_diagnostics["prepared_rows"] = int(len(omr_prepared))

diag_table = (
    pd.DataFrame({"metric": list(omr_diagnostics.keys()), "value":'
↪ list(omr_diagnostics.values())})

```

```

        .assign(value=lambda d: d["value"].astype(str))
    )
print(diag_table)

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")
omr_diag_path = prior_runs_dir / f"{out_date} omr_diagnostics.json"
omr_diag_path.write_text(json.dumps(omr_diagnostics, indent=2))
print("Wrote:", omr_diag_path)

# Optional charted anthropometric fallback (nearest-anytime selection).
use_chARTed_fallback = os.getenv("COHORT_ANTHRO_CHARTED_FALLBACK",
    ↪ "1").strip() == "1"
use_nearest_anytime = os.getenv("COHORT_ANTHRO_NEAREST_ANYTIME", "1").strip()
    ↪ == "1"
charted_anthro_diagnostics = {
    "enabled": bool(use_chARTed_fallback),
    "nearest_anytime": bool(use_nearest_anytime),
    "charted_rows_input": 0,
    "charted_rows_usable": 0,
    "filled_counts": {},
    "rows_with_any_charted_fill": 0,
}

if use_chARTed_fallback and len(hadm_list) > 0:
    charted_anthro_sql = f"""
    WITH hadms AS (
        SELECT x AS hadm_id
        FROM UNNEST(@hadms) AS x
    ),
    candidate_items AS (
        SELECT itemid, label, category
        FROM `{{PHYS}}.{{ICU}}.d_items`
        WHERE REGEXP_CONTAINS(LOWER(label),
    ↪ r'(height|weight|admission\\s+weight)')
    ),
    ce AS (
        SELECT
            ce.subject_id,
            ce.hadm_id,
            ce.charttime AS obs_time,
            ci.label,

```

```

        ce.valuenum,
        LOWER(REPLACE(COALESCE(ce.valueuom, ''), ' ', '')) AS valueuom
    FROM `{{PHYS}}.{{ICU}}.chartevents` ce
    JOIN hadms h
        ON ce.hadm_id = h.hadm_id
    JOIN candidate_items ci
        ON ce.itemid = ci.itemid
    WHERE ce.valuenum IS NOT NULL
)
SELECT
    subject_id,
    obs_time,
    CASE
        WHEN REGEXP_CONTAINS(LOWER(label), r'height') THEN 'height'
        WHEN REGEXP_CONTAINS(LOWER(label), r'weight') THEN 'weight'
        ELSE NULL
    END AS result_name,
    CASE
        WHEN REGEXP_CONTAINS(LOWER(label), r'height') THEN
            CASE
                WHEN valueuom IN ('cm', 'centimeter', 'centimeters') THEN
                    ← valuenum
                        WHEN valueuom IN ('m', 'meter', 'meters') THEN valuenum * 100.0
                        WHEN valueuom IN ('in', 'inch', 'inches') THEN valuenum * 2.54
                        ELSE valuenum
            END
        WHEN REGEXP_CONTAINS(LOWER(label), r'weight') THEN
            CASE
                WHEN valueuom IN ('kg', 'kilogram', 'kilograms') THEN valuenum
                WHEN valueuom IN ('lb', 'lbs', 'pound', 'pounds') THEN valuenum *
                    ← 0.45359237
                    WHEN valueuom IN ('oz', 'ounce', 'ounces') THEN valuenum *
                    ← 0.0283495231
                    ELSE valuenum
            END
        ELSE NULL
    END AS result_value_num,
    'icu_charted' AS source
FROM ce
WHERE REGEXP_CONTAINS(LOWER(label), r'(height|weight)')
"""

charted_raw = run_sql_bq(charted_anthro_sql, {"hadms": hadm_list})
ed_df, charted_anthro_diagnostics = attach_charted_anthro_fallback(

```

```

        ed_df,
        charted_raw,
        nearest_anytime=use_nearest_anytime,
    )

    height_cm = pd.to_numeric(ed_df.get("height_closest_pre_ed"),
    ↵ errors="coerce")
    weight_kg = pd.to_numeric(ed_df.get("weight_closest_pre_ed"),
    ↵ errors="coerce")
    bmi_existing = pd.to_numeric(ed_df.get("bmi_closest_pre_ed"),
    ↵ errors="coerce")
    bmi_calc = weight_kg / (height_cm / 100.0) ** 2
    bmi_fill_mask = bmi_existing.isna() & bmi_calc.notna() & height_cm.gt(0)
    ed_df.loc[bmi_fill_mask, "bmi_closest_pre_ed"] = bmi_calc[bmi_fill_mask]
    charted_anthro_diagnostics["bmi_filled_from_height_weight"] =
    ↵ int(bmi_fill_mask.sum())
    charted_anthro_diagnostics["enabled"] = True
    charted_anthro_diagnostics["nearest_anytime"] = bool(use_nearest_anytime)

anthro_model_specs = {
    "bmi_closest_pre_ed": ("bmi_closest_pre_ed_model", "bmi_outlier_flag",
    ↵ 0.0, 100.0),
    "height_closest_pre_ed": ("height_closest_pre_ed_model",
    ↵ "height_outlier_flag", 100.0, 250.0),
    "weight_closest_pre_ed": ("weight_closest_pre_ed_model",
    ↵ "weight_outlier_flag", 0.0, 400.0),
}
anthro_audit_rows: list[dict[str, Any]] = []
for raw_col, (model_col, flag_col, lower_bound, upper_bound) in
    ↵ anthro_model_specs.items():
    raw_numeric = pd.to_numeric(ed_df.get(raw_col), errors="coerce")
    outlier_mask = raw_numeric.notna() &
        raw_numeric.le(lower_bound) | raw_numeric.gt(upper_bound)
    )
    ed_df[flag_col] = outlier_mask
    ed_df[model_col] = raw_numeric.mask(outlier_mask)
    anthro_audit_rows.append(
        {
            "raw_column": raw_col,
            "model_column": model_col,
            "lower_bound_exclusive": lower_bound,
            "upper_bound_inclusive": upper_bound,
            "raw_non_missing_n": int(raw_numeric.notna().sum()),
        }
    )
}

```

```

        "outlier_n": int(outlier_mask.sum()),
        "model_non_missing_n": int(pd.to_numeric(ed_df[model_col],
            ↪ errors="coerce").notna().sum()),
    }
)

anthropometric_cleaning_audit = pd.DataFrame(anthro_audit_rows)
anthropometric_cleaning_audit_path = (
    prior_runs_dir / f"{out_date} anthropometric_cleaning_audit.csv"
)
anthropometric_cleaning_audit.to_csv(anthropometric_cleaning_audit_path,
    ↪ index=False)
print("Wrote:", anthropometric_cleaning_audit_path)

anthro_coverage_audit = build_anthro_coverage_audit(ed_df)
anthro_coverage_audit["omr"] = omr_diagnostics
anthro_coverage_audit["chartedFallback"] = charted_anthro_diagnostics
anthro_coverage_audit["cleaning_audit_path"] =
    ↪ str(anthropometric_cleaning_audit_path)
anthro_coverage_audit_path = prior_runs_dir / f"{out_date}"
    ↪ anthropometrics_coverage_audit.json"
anthro_coverage_audit_path.write_text(json.dumps(anthro_coverage_audit,
    ↪ indent=2))
print("Wrote:", anthro_coverage_audit_path)

attached_non_null = omr_diagnostics.get("attached_non_null_counts", {})
attached_total = int(sum(int(v) for v in attached_non_null.values())) if
    ↪ attached_non_null else 0
subject_overlap = int(omr_diagnostics.get("subject_overlap_count", 0))
pre_window_rows = int(omr_diagnostics.get("pre_window_candidate_rows", 0))
post_window_rows = int(omr_diagnostics.get("post_window_candidate_rows", 0))
within_window_rows = int(omr_diagnostics.get("within_window_candidate_rows",
    ↪ 0))
selected_tier_counts = omr_diagnostics.get("selected_tier_counts", {})
selected_tier_rates = omr_diagnostics.get("selected_tier_rates", {})
allow_empty_omr = os.getenv("COHORT_ALLOW_EMPTY_OMR", "0").strip() == "1"
fail_on_omr_attach_inconsistency =
    os.getenv("COHORT_FAIL_ON_OMR_ATTACH_INCONSISTENCY", "1").strip() == "1"
)

omr_window_diagnostics = {
    "subject_overlap_count": subject_overlap,
    "pre_window_candidate_rows": pre_window_rows,

```

```

"post_window_candidate_rows": post_window_rows,
"closest_absolute_candidate_rows":
    ↵ int(omr_diagnostics.get("closest_absolute_candidate_rows", 0)),
"within_window_candidate_rows": within_window_rows,
"attached_total_non_null_values": attached_total,
"selected_tier_counts": selected_tier_counts,
"selected_tier_rates": selected_tier_rates,
}

if len(omr_raw) > 0 and subject_overlap > 0:
    if pre_window_rows == 0 and post_window_rows > 0:
        print(
            "INFO: OMR subject overlap exists with no pre-ED candidates but
            ↵ post-ED candidates within the 365-day window. "
            "Anthropometric fallback will use post-ED observations and mark
            ↵ them as timing-uncertain."
        )
    elif pre_window_rows == 0 and post_window_rows == 0:
        print(
            "INFO: OMR subject overlap exists but no candidates are within
            ↵ +/-365 days of ED arrival. "
            "Anthropometrics will remain missing for this run."
        )
    if attached_total == 0 and (pre_window_rows > 0 or post_window_rows > 0):
        msg = (
            "OMR attachment produced all-null outputs despite available
            ↵ +/-365-day candidates. "
            "Set COHORT_ALLOW_EMPTY_OMR=1 or
            ↵ COHORT_FAIL_ON_OMR_ATTACH_INCONSISTENCY=0 to bypass this
            ↵ guard."
        )
        if allow_empty_omr or not fail_on_omr_attach_inconsistency:
            print("WARNING:", msg)
        else:
            raise ValueError(msg)

OMR rows: 46494
               metric \
0              window_days
1      source_rows_prepared
2          parsed_value_rows
3                  ed_rows
4   ed_rows_eligible_for_join

```

```

5           subject_overlap_count
6 candidate_rows_after_subject_join
7             days_before_min
8             days_before_max
9     nonnegative_candidate_rows
10    pre_window_candidate_rows
11    post_window_candidate_rows
12 closest_absolute_candidate_rows
13 within_window_candidate_rows
14 eligible_ed_stays_with_candidates
15   attached_non_null_counts
16   attached_any_non_null_rows
17     selected_tier_counts
18     selected_tier_rates
19 timing_uncertain_count
20 anthro_source_counts
21   source_rows_raw
22     prepared_rows

                           value
0                      365
1                     46494
2                     46489
3                     11786
4                     11786
5                      626
6                     1940
7                     -4034
8                     -275
9                      0
10                     0
11                     21
12                     21
13                     0
14                     19
15 {'bmi_closest_pre_ed': 8, 'height_closest_pre_... 19
16                                         19
17 {'pre_ed_365': 0, 'post_ed_365': 19, 'missing'... 19
18 {'pre_ed_365': 0.0, 'post_ed_365': 0.001612082... 19
19                                         19
20                                         {'missing': 11767, 'omr': 19}
21                                         46494
22                                         46494

```

```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 omr_diagnostics.json
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 anthropometric_cleaning_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 anthropometrics_coverage_audit.json
INFO: OMR subject overlap exists with no pre-ED candidates but post-ED candidates within the day window. Anthropometric fallback will use post-ED observations and mark them as timing-uncertain.

```

3.0.9 Phase 7 — ICD comorbidity flags

Rationale: Derive comorbidity indicators from index admission ICD codes for stratified analyses.

```

# Purpose: Compute comorbidity flags from hospital + ED ICD codes with
#           ↳ chunked hadm queries.

# ICD code pulls for comorbidity flags (hospital + ED; combined OR)
# NOTE: Use prefix filters to reduce CPU and avoid regex-heavy scans.

FLAGS = [
    "flag_copd", "flag_osa_ohs", "flag_chf", "flag_neuromuscular",
    "flag_opioid_substance", "flag_pneumonia",
]

SQL["icd_flags_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
icd AS (
    SELECT
        hadm_id,
        UPPER(REPLACE(icd_code, ".", "")) AS code_norm
    FROM `{{PHYS}}.{{HOSP}}.diagnoses_icd`
    WHERE hadm_id IN (SELECT hadm_id FROM hadms)
    AND (
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J43") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J44") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G473") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "E662") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "I50") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G12") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G70") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "G71") OR

```

```

        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "F11") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "T40") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "F13") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J12") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J13") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J14") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J15") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J16") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J17") OR
        STARTS_WITH(UPPER(REPLACE(icd_code, ".", "")), "J18")
    )
)
SELECT
    hadm_id,
    MAX(IF(STARTS_WITH(code_norm, "J43") OR STARTS_WITH(code_norm, "J44"), 1,
    0)) AS flag_copd,
    MAX(IF(STARTS_WITH(code_norm, "G473") OR STARTS_WITH(code_norm, "E662"), 1,
    0)) AS flag_osa_ohs,
    MAX(IF(STARTS_WITH(code_norm, "I50"), 1, 0)) AS flag_chf,
    MAX(IF(STARTS_WITH(code_norm, "G12") OR STARTS_WITH(code_norm, "G70") OR
    STARTS_WITH(code_norm, "G71"), 1, 0)) AS flag_neuromuscular,
    MAX(IF(STARTS_WITH(code_norm, "F11") OR STARTS_WITH(code_norm, "T40") OR
    STARTS_WITH(code_norm, "F13"), 1, 0)) AS flag_opioid_substance,
    MAX(IF(STARTS_WITH(code_norm, "J12") OR STARTS_WITH(code_norm, "J13") OR
    STARTS_WITH(code_norm, "J14") OR STARTS_WITH(code_norm, "J15") OR
    STARTS_WITH(code_norm, "J16") OR STARTS_WITH(code_norm, "J17") OR
    STARTS_WITH(code_norm, "J18"), 1, 0)) AS flag_pneumonia
FROM icd
GROUP BY hadm_id
"""

SQL["ed_icd_flags_sql"] = f"""
WITH hadms AS (SELECT x AS hadm_id FROM UNNEST(@hadms) AS x),
ed_dx AS (
    SELECT
        s.hadm_id,
        UPPER(REPLACE(d.icd_code, ".", "")) AS code_norm
    FROM `{{PHYS}}.{{ED}}.diagnosis` d
    JOIN `{{PHYS}}.{{ED}}.edstays` s
    ON s.stay_id = d.stay_id
    WHERE s.hadm_id IN (SELECT hadm_id FROM hadms)
    AND (
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J43") OR

```

```

        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J44") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G473") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "E662") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "I50") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G12") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G70") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "G71") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "F11") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "T40") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "F13") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J12") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J13") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J14") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J15") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J16") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J17") OR
        STARTS_WITH(UPPER(REPLACE(d.icd_code, ".", "")), "J18")
    )
)
SELECT
    hadm_id,
    MAX(IF(STARTS_WITH(code_norm, "J43") OR STARTS_WITH(code_norm, "J44"), 1,
    0)) AS flag_copd,
    MAX(IF(STARTS_WITH(code_norm, "G473") OR STARTS_WITH(code_norm, "E662"), 1,
    0)) AS flag_osa_ohs,
    MAX(IF(STARTS_WITH(code_norm, "I50"), 1, 0)) AS flag_chf,
    MAX(IF(STARTS_WITH(code_norm, "G12") OR STARTS_WITH(code_norm, "G70") OR
    STARTS_WITH(code_norm, "G71"), 1, 0)) AS flag_neuromuscular,
    MAX(IF(STARTS_WITH(code_norm, "F11") OR STARTS_WITH(code_norm, "T40") OR
    STARTS_WITH(code_norm, "F13"), 1, 0)) AS flag_opioid_substance,
    MAX(IF(STARTS_WITH(code_norm, "J12") OR STARTS_WITH(code_norm, "J13") OR
    STARTS_WITH(code_norm, "J14") OR STARTS_WITH(code_norm, "J15") OR
    STARTS_WITH(code_norm, "J16") OR STARTS_WITH(code_norm, "J17") OR
    STARTS_WITH(code_norm, "J18"), 1, 0)) AS flag_pneumonia
FROM ed_dx
GROUP BY hadm_id
"""

def _run_flag_query_chunked(sql_name: str, label: str) -> pd.DataFrame:
    parts = []
    hadm_chunk_size = 10000
    for i in range(0, len(hadm_list), hadm_chunk_size):
        hadm_chunk = hadm_list[i:i + hadm_chunk_size]

```

```

part = run_sql_bq(sql(sql_name), {"hadms": hadm_chunk})
if len(part) > 0:
    parts.append(part)
print(f"{label} chunk {i // hadm_chunk_size + 1}: rows={len(part)}")

if parts:
    out = pd.concat(parts, ignore_index=True)
    out = out.groupby("hadm_id", as_index=False)[FLAGS].max()
    return out

return pd.DataFrame(columns=["hadm_id"] + FLAGS)

flag_hosp = _run_flag_query_chunked("icd_flags_sql", "Hosp ICD flags")
flag_ed = _run_flag_query_chunked("ed_icd_flags_sql", "ED ICD flags")

flag_hosp = flag_hosp.rename(columns={k: f"{k}_hosp" for k in FLAGS})
flag_ed = flag_ed.rename(columns={k: f"{k}_ed" for k in FLAGS})

flag_df = flag_hosp.merge(flag_ed, on="hadm_id", how="outer")
for k in FLAGS:
    hosp_col = f"{k}_hosp"
    ed_col = f"{k}_ed"
    if hosp_col not in flag_df.columns:
        flag_df[hosp_col] = 0
    if ed_col not in flag_df.columns:
        flag_df[ed_col] = 0
    flag_df[k] = ((flag_df[hosp_col].fillna(0).astype(int) == 1) |
    (flag_df[ed_col].fillna(0).astype(int) == 1)).astype(int)

# Merge into admission-level and ED-stay-level frames (override if present)
for _df_name in ["df", "ed_df"]:
    if _df_name in globals():
        _df = globals()[_df_name]
        drop_cols = [c for c in flag_df.columns if c != "hadm_id" and c in
        _df.columns]
        if drop_cols:
            _df = _df.drop(columns=drop_cols)
        _df = _df.merge(flag_df, on="hadm_id", how="left")
        globals()[_df_name] = _df

# Prevalence (combined flags)
for k in FLAGS:
    if k in ed_df.columns:

```

```

print(k, int(ed_df[k].fillna(0).sum()))

Hosp ICD flags chunk 1: rows=3538
Hosp ICD flags chunk 2: rows=3613
Hosp ICD flags chunk 3: rows=3544
Hosp ICD flags chunk 4: rows=3527
Hosp ICD flags chunk 5: rows=56
ED ICD flags chunk 1: rows=329
ED ICD flags chunk 2: rows=407
ED ICD flags chunk 3: rows=378
ED ICD flags chunk 4: rows=407
ED ICD flags chunk 5: rows=5
flag_copd 2168
flag_osa_ohs 1450
flag_chf 3153
flag_neuromuscular 78
flag_opioid_substance 521
flag_pneumonia 1997

```

3.0.10 Phase 8 — Timing phenotypes and derived bands

Rationale: Compute time-anchored hypercapnia/acidemia phenotypes for ED presentation vs later course.

```

# Purpose: Create time-anchored phenotypes and severity bands relative to
#           ↵ first ED presentation.

# Timing phenotypes

anchor = "ed_intime_first" if "ed_intime_first" in ed_df.columns else
        ↵ "ed_intime"

ed_df["dt_first_qualifying_gas_hours"] = (ed_df["first_gas_time"] -
                                         ↵ ed_df[anchor]).dt.total_seconds() / 3600.0
ed_df["presenting_hypercapnia"] = (ed_df["dt_first_qualifying_gas_hours"] <=
        ↵ 6).astype("Int64")
ed_df["late_hypercapnia"] = (ed_df["dt_first_qualifying_gas_hours"] >
        ↵ 6).astype("Int64")

qualifying_hours = pd.to_numeric(ed_df["dt_first_qualifying_gas_hours"],
        ↵ errors="coerce")
qualifying_observed = qualifying_hours.notna()

```

```

ed_df["qualifying_gas_observed"] = qualifying_observed.astype("Int64")
ed_df["presenting_hypercapnia_tri"] = pd.Series(pd.NA, index=ed_df.index,
↪ dtype="Int64")
ed_df["late_hypercapnia_tri"] = pd.Series(pd.NA, index=ed_df.index,
↪ dtype="Int64")
ed_df.loc[qualifying_observed, "presenting_hypercapnia_tri"] = (
    qualifying_hours.loc[qualifying_observed].le(6).astype("Int64")
)
ed_df.loc[qualifying_observed, "late_hypercapnia_tri"] = (
    qualifying_hours.loc[qualifying_observed].gt(6)
    & qualifying_hours.loc[qualifying_observed].le(24)
).astype("Int64")

threshold_any = (
    pd.to_numeric(ed_df.get("pco2_threshold_any", pd.Series(0,
↪ index=ed_df.index)), errors="coerce")
    .fillna(0)
    .astype(int)
)
any_icd = (
    pd.to_numeric(ed_df.get("any_hypercap_icd", pd.Series(0,
↪ index=ed_df.index)), errors="coerce")
    .fillna(0)
    .astype(int)
)
ed_df["hypercap_timing_class"] = np.select(
    [
        qualifying_observed & qualifying_hours.le(6),
        qualifying_observed & qualifying_hours.gt(6) &
↪ qualifying_hours.le(24),
        (qualifying_observed & qualifying_hours.gt(24))
        | (~qualifying_observed & threshold_any.eq(1)),
    ],
    [
        "presenting_0_6h",
        "late_6_24h",
        "none_within_24h",
    ],
    default="icd_only_or_no_qualifying_gas",
)
# NIV/IMV timing relative to first ED presentation
if "first_imv_time" in ed_df.columns:

```

```

ed_df["dt_first_imv_hours"] = (ed_df["first_imv_time"] -
    ed_df[anchor]).dt.total_seconds() / 3600.0
if "first_niv_time" in ed_df.columns:
    ed_df["dt_first_niv_hours"] = (ed_df["first_niv_time"] -
        ed_df[anchor]).dt.total_seconds() / 3600.0

ed_df["hospital_los_negative_flag"] = (
    pd.to_numeric(ed_df.get("hospital_los_hours"), errors="coerce")
    .lt(0)
    .fillna(False)
)
admittime_series = pd.to_datetime(ed_df.get("admittime"), errors="coerce")
ed_intime_series = pd.to_datetime(ed_df.get("ed_intime"), errors="coerce")
ed_df["admittime_before_ed_intime_flag"] = (
    admittime_series.notna()
    & ed_intime_series.notna()
    & admittime_series.lt(ed_intime_series)
)
hospital_los_numeric = pd.to_numeric(ed_df.get("hospital_los_hours"),
    errors="coerce")
ed_df["hospital_los_hours_model"] =
    hospital_los_numeric.mask(hospital_los_numeric.lt(0))

dischtime_series = pd.to_datetime(ed_df.get("dischtime"), errors="coerce")
anchor_series = pd.to_datetime(ed_df.get(anchor), errors="coerce")
for event_col, delta_col, model_col, outside_col in [
    ("first_imv_time", "dt_first_imv_hours", "dt_first_imv_hours_model",
     "imv_time_outside_window_flag"),
    ("first_niv_time", "dt_first_niv_hours", "dt_first_niv_hours_model",
     "niv_time_outside_window_flag"),
]:
    if event_col not in ed_df.columns:
        ed_df[model_col] = pd.Series(pd.NA, index=ed_df.index,
            dtype="Float64")
        ed_df[outside_col] = False
        continue
    event_time = pd.to_datetime(ed_df[event_col], errors="coerce")
    delta_hours = pd.to_numeric(ed_df.get(delta_col), errors="coerce")
    outside_lower = anchor_series.notna() & event_time.notna() &
    event_time.lt(anchor_series)
    outside_upper = discharge_time_series.notna() & event_time.notna() &
    event_time.gt(discharge_time_series)
    negative_delta = delta_hours.lt(0).fillna(False)

```

```

outside_window = outside_lower | outside_upper | negative_delta
ed_df[outside_col] = outside_window.fillna(False)
ed_df[model_col] = delta_hours.mask(outside_window)

# Bands
bins_ph = [-1, 7.20, 7.30, 7.35, 99]
labels_ph = ["<7.20", "7.20-7.29", "7.30-7.34", "7.35"]
ed_df["ph_band"] = pd.cut(ed_df["first_ph"], bins=bins_ph, labels=labels_ph)

bins_hco3 = [-1, 24, 30, 999]
labels_hco3 = ["<24", "24-29", "30"]
ed_df["hco3_band"] = pd.cut(ed_df["first_hco3"], bins=bins_hco3,
                           labels=labels_hco3)

bins_lac = [-1, 2, 4, 999]
labels_lac = ["<2", "2-4", ">4"]
ed_df["lactate_band"] = pd.cut(ed_df["first_lactate"], bins=bins_lac,
                               labels=labels_lac)

# Align ED-stay gas source flags to hadm-level thresholds and ICU POC signal.
# This keeps source-specific flags internally consistent in the final ED-stay
# export.
if "df" in globals() and "hadm_id" in ed_df.columns and "hadm_id" in
    df.columns:
    hadm_thr_cols = [c for c in ["abg_hypercap_threshold",
        "vbg_hypercap_threshold", "other_hypercap_threshold"] if c in df.columns]
    if hadm_thr_cols:
        hadm_thr = df[["hadm_id"] + hadm_thr_cols].drop_duplicates("hadm_id")
        hadm_thr_map = hadm_thr.set_index("hadm_id")

        for col in hadm_thr_cols:
            mapped = pd.to_numeric(ed_df["hadm_id"].map(hadm_thr_map[col]),
                errors="coerce")
            existing = pd.to_numeric(ed_df[col], errors="coerce") if col in
                ed_df.columns else pd.Series(np.nan, index=ed_df.index)
            ed_df[col] = existing.fillna(mapped).fillna(0).astype("Int64")

# Source-specific flags must include any matching hadm-level threshold.
flag_specs = [
    ("flag_abg_hypercapnia", "abg_hypercap_threshold"),
    ("flag_vbg_hypercapnia", "vbg_hypercap_threshold"),

```

```

        ("flag_other_hypercapnia", "other_hypercap_threshold"),
    ]
for flag_col, thr_col in flag_specs:
    base = pd.to_numeric(ed_df.get(flag_col, pd.Series(0,
        index=ed_df.index)), errors="coerce").fillna(0).astype(int)
    thr = pd.to_numeric(ed_df.get(thr_col, pd.Series(0, index=ed_df.index)),
        errors="coerce").fillna(0).astype(int)
    ed_df[flag_col] = ((base == 1) | (thr == 1)).astype("Int64")

any_from_sources = (
    (pd.to_numeric(ed_df.get("flag_abg_hypercapnia", 0),
        errors="coerce").fillna(0).astype(int) == 1)
    | (pd.to_numeric(ed_df.get("flag_vbg_hypercapnia", 0),
        errors="coerce").fillna(0).astype(int) == 1)
    | (pd.to_numeric(ed_df.get("flag_other_hypercapnia", 0),
        errors="coerce").fillna(0).astype(int) == 1)
)
poc_any = pd.to_numeric(ed_df.get("flag_any_gas_hypercapnia_poc",
    pd.Series(0, index=ed_df.index)), errors="coerce").fillna(0).astype(int)
    == 1
ed_df["flag_any_gas_hypercapnia"] = (any_from_sources |
    poc_any).astype("Int64")

```

3.1 QA & Data Fidelity

Rationale: Validate joins, missingness, lab completeness, and produce reproducibility artifacts.

```

# Purpose: Create reusable data-quality and merge guardrail helpers to
# prevent silent join errors.

# --- QA checks (lightweight, deterministic)
STRICT_QA = False

# 1) Key uniqueness
if 'ed_df' in globals() and 'ed_stay_id' in ed_df.columns:
    try:
        assert_unique(ed_df, 'ed_stay_id', 'ed_df')
        print('QA: ed_stay_id unique OK')
    except Exception as e:
        print('QA FAIL:', e)
        if STRICT_QA:

```

```

    raise

# 2) Inclusion sanity: any_hypercap_icd==0 implies gas criteria met
if 'ed_df' in globals() and 'any_hypercap_icd' in ed_df.columns:
    gas_flag = ed_df.get('flag_any_gas_hypercapnia', pd.Series(0,
    ↵ index=ed_df.index)).fillna(0).astype(int)
    icd_flag = ed_df['any_hypercap_icd'].fillna(0).astype(int)
    viol = (icd_flag == 0) & (gas_flag == 0)
    n_viol = int(viol.sum())
    print('QA: ICD==0 & Gas==0 count:', n_viol)
    if STRICT_QA and n_viol > 0:
        raise AssertionError('Found rows without ICD and without gas
        ↵ criteria.')
else:
    print('QA: No rows with ICD==0 & Gas==0')

# 3) Temporal sanity: first_gas_time >= ed_intime (allow small negative
    ↵ drift)
if 'ed_df' in globals() and 'first_gas_time' in ed_df.columns and 'ed_intime'
    ↵ in ed_df.columns:
    dt = (pd.to_datetime(ed_df['first_gas_time']) -
    ↵ pd.to_datetime(ed_df['ed_intime'])).dt.total_seconds() / 3600
    n_neg = int((dt < -1).sum()) # allow 1h clock drift
    print('QA: first_gas_time < ed_intime by >1h:', n_neg)
    if STRICT_QA and n_neg > 0:
        raise AssertionError('first_gas_time before ed_intime by >1h')
else:
    print('QA: No rows with first_gas_time < ed_intime by >1h')

# 4) Range checks (report only)
if 'ed_df' in globals():
    ranges = {
        'first_ph': (6.8, 7.8),
        'first_pco2': (10, 200),
        'first_lactate': (0, 30),
        'creatinine': (0, 20),
    }
    rc = check_ranges(ed_df, ranges)
    if not rc.empty:
        print('QA range check (n out-of-range):')
        print(rc.to_string(index=False))

QA: ed_stay_id unique OK
QA: first_gas_time < ed_intime by >1h: 0
QA range check (n out-of-range):
    col  n_bad
    first_ph      9

```

```

first_pco2      2
first_lactate   0

# Purpose: Report missingness for key variables to support transparent
#           data-quality reporting.

import json
from datetime import datetime

# T1 - uniqueness and join explosion guard
if ed_df["ed_stay_id"].nunique() != len(ed_df):
    raise ValueError("ed_stay_id not unique after merges")

# T1b - create cleaned vitals model fields while preserving raw values.
ed_df, vitals_outlier_audit = add_vitals_model_fields(ed_df)
ed_df, gas_outlier_audit = add_gas_model_fields(ed_df)
outliers_audit = pd.concat([vitals_outlier_audit, gas_outlier_audit],
                           ignore_index=True, sort=False)

# T1c - fail-fast cleaned-vitals bounds checks.
clean_bounds = {
    "ed_triage_o2sat_clean": (0.0, 100.0),
    "ed_first_o2sat_clean": (0.0, 100.0),
    "ed_triage_sbp_clean": (20.0, 300.0),
    "ed_first_sbp_clean": (20.0, 300.0),
    "ed_triage_dbp_clean": (10.0, 200.0),
    "ed_first_dbp_clean": (10.0, 200.0),
}
for column_name, (lower_bound, upper_bound) in clean_bounds.items():
    if column_name not in ed_df.columns:
        continue
    numeric = pd.to_numeric(ed_df[column_name], errors="coerce")
    invalid_n = int((numeric.notna() & ((numeric < lower_bound) | (numeric >
    upper_bound))).sum())
    if invalid_n:
        raise ValueError(
            f"Cleaned vital bounds failure for {column_name}: {invalid_n}"
            " values outside [{lower_bound}, {upper_bound}]"
        )

for temp_clean_column in ("ed_triage_temp_f_clean", "ed_first_temp_f_clean"):
    if temp_clean_column not in ed_df.columns:
        continue
    numeric = pd.to_numeric(ed_df[temp_clean_column], errors="coerce")

```

```

    residual_celsius_like_n = int(numeric.between(20.0, 50.0,
    ↵ inclusive="both").sum())
    if residual_celsius_like_n:
        raise ValueError(
            f"Found {residual_celsius_like_n} cleaned values still in 20-50°F
            ↵ band in {temp_clean_column}."
        )

ed_vitals_distribution_summary, ed_vitals_extreme_examples,
    ↵ ed_vitals_model_delta = build_ed_vitals_audit_artifacts(ed_df)

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")

vitals_outlier_audit_path = prior_runs_dir / f"{out_date}
    ↵ vitals_outlier_audit.csv"
vitals_outlier_audit.to_csv(vitals_outlier_audit_path, index=False)
print("Wrote:", vitals_outlier_audit_path)
outliers_audit_path = prior_runs_dir / f"{out_date} outliers_audit.csv"
outliers_audit.to_csv(outliers_audit_path, index=False)
print("Wrote:", outliers_audit_path)
ed_vitals_distribution_summary_path = prior_runs_dir / f"{out_date}
    ↵ ed_vitals_distribution_summary.csv"
ed_vitals_distribution_summary.to_csv(ed_vitals_distribution_summary_path,
    ↵ index=False)
print("Wrote:", ed_vitals_distribution_summary_path)
ed_vitals_extreme_examples_path = prior_runs_dir / f"{out_date}
    ↵ ed_vitals_extreme_examples.csv"
ed_vitals_extreme_examples.to_csv(ed_vitals_extreme_examples_path,
    ↵ index=False)
print("Wrote:", ed_vitals_extreme_examples_path)
ed_vitals_model_delta_path = prior_runs_dir / f"{out_date}
    ↵ ed_vitals_model_delta.csv"
ed_vitals_model_delta.to_csv(ed_vitals_model_delta_path, index=False)
print("Wrote:", ed_vitals_model_delta_path)

timing_integrity_audit = pd.DataFrame(
    [
        {
            "rows_total": int(len(ed_df)),

```

```

    "hospital_los_negative_n":
    ↵ int(ed_df.get("hospital_los_negative_flag", pd.Series(False,
    ↵ index=ed_df.index)).fillna(False).astype(bool).sum()),
    "admittime_before_ed_intime_n":
    ↵ int(ed_df.get("admittime_before_ed_intime_flag",
    ↵ pd.Series(False,
    ↵ index=ed_df.index)).fillna(False).astype(bool).sum()),
    "qualifying_gas_observed_n":
    ↵ int(pd.to_numeric(ed_df.get("qualifying_gas_observed")),
    ↵ errors="coerce").fillna(0).astype(int).sum()),
    "qualifying_gas_missing_n":
    ↵ int(pd.to_numeric(ed_df.get("qualifying_gas_observed")),
    ↵ errors="coerce").fillna(0).astype(int).eq(0).sum(),
}
]
)
timing_integrity_audit_path = prior_runs_dir / f"{out_date}"
↪ timing_integrity_audit.csv"
timing_integrity_audit.to_csv(timing_integrity_audit_path, index=False)
print("Wrote:", timing_integrity_audit_path)

ventilation_timing_audit = pd.DataFrame(
[
{
    "rows_total": int(len(ed_df)),
    "dt_first_imv_hours_negative_n":
    ↵ int(pd.to_numeric(ed_df.get("dt_first_imv_hours"),
    ↵ errors="coerce").lt(0).fillna(False).sum()),
    "dt_first_niv_hours_negative_n":
    ↵ int(pd.to_numeric(ed_df.get("dt_first_niv_hours"),
    ↵ errors="coerce").lt(0).fillna(False).sum()),
    "imv_time_outside_window_n":
    ↵ int(ed_df.get("imv_time_outside_window_flag",
    ↵ pd.Series(False,
    ↵ index=ed_df.index)).fillna(False).astype(bool).sum()),
    "niv_time_outside_window_n":
    ↵ int(ed_df.get("niv_time_outside_window_flag",
    ↵ pd.Series(False,
    ↵ index=ed_df.index)).fillna(False).astype(bool).sum()),
    "dt_first_imv_hours_model_nonnull_n":
    ↵ int(pd.to_numeric(ed_df.get("dt_first_imv_hours_model"),
    ↵ errors="coerce").notna().sum()),
}
]
)

```

```

        "dt_first_niv_hours_model_nonnull_n":
        ↵  int(pd.to_numeric(ed_df.get("dt_first_niv_hours_model")),
        ↵  errors="coerce").notna().sum(),
    }
]
)
ventilation_timing_audit_path = prior_runs_dir / f"{out_date}"
↪ ventilation_timing_audit.csv"
ventilation_timing_audit.to_csv(ventilation_timing_audit_path, index=False)
print("Wrote:", ventilation_timing_audit_path)

print("ED vitals cleaning summary (raw->clean removals):")
display(
    ed_vitals_model_delta[
        [
            "vital_name",
            "raw_n",
            "clean_n",
            "raw_removed_n",
            "raw_removed_pct",
            "converted_celsius_like_n",
        ]
    ].sort_values("vital_name")
)

# T2 - missingness summary with expected-vs-unexpected classifications.
expected_sparse_fields = set()
omr_pre_window_rows = int(omr_diagnostics.get("pre_window_candidate_rows",
    ↵  0)) if "omr_diagnostics" in globals() else 0
omr_post_window_rows = int(omr_diagnostics.get("post_window_candidate_rows",
    ↵  0)) if "omr_diagnostics" in globals() else 0
if omr_pre_window_rows == 0 and omr_post_window_rows == 0:
    expected_sparse_fields.update({"bmi_closest_pre_ed",
    ↵  "height_closest_pre_ed", "weight_closest_pre_ed"})

miss = classify_missingness_expectations(
    ed_df,
    TARGET_FIELDS,
    expected_sparse_fields=expected_sparse_fields,
)
print(miss.sort_values(["missing_pct", "field"], ascending=[False,
    ↵  True]).head(30))

```

```

unexpected_full_null_fields = miss.loc[
    miss["expectation"].eq("unexpected_full_null"), "field"
].tolist()
expected_structural_null_fields = sorted(
    set(EXPECTED_STRUCTURAL_NULL_FIELDS).intersection(ed_df.columns)
)

# Compute UOM checks against canonical export frame contract (or
# deterministic preview).
if "final_cc" in globals():
    uom_scope_df = final_cc
    uom_scope_frame = "final_cc"
elif "df" in globals() and "ed_triage_cc" in df.columns:
    _mask_df_cc = df["ed_triage_cc"].notna() &
    (df["ed_triage_cc"].astype(str).str.strip() != "")
    _df_cc = df.loc[_mask_df_cc].copy()
    _mask_cc = ed_df["ed_triage_cc"].notna() &
    (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
    _ed_cc_only = ed_df.loc[_mask_cc].copy()
    _ed_tmp = _ed_cc_only.copy()
    if "ed_intime" in _ed_tmp.columns:
        _ed_tmp = _ed_tmp.sort_values(["hadm_id", "ed_intime", "ed_stay_id"],
        na_position="last")
    else:
        _ed_tmp = _ed_tmp.sort_values(["hadm_id", "ed_stay_id"],
        na_position="last")
    _ed_first = _ed_tmp.drop_duplicates("hadm_id", keep="first")
    _add_ed_cols = [c for c in _ed_first.columns if c not in _df_cc.columns]
    uom_scope_df = _df_cc.merge(_ed_first[["hadm_id"] + _add_ed_cols],
    on="hadm_id", how="left")
    uom_scope_frame = "final_cc_preview"
else:
    uom_scope_df = ed_df
    uom_scope_frame = "ed_df"

uom_expectation_checks = evaluate_uom_expectations(uom_scope_df)
uom_expectation_scope = {"frame": uom_scope_frame, "row_count":
    int(len(uom_scope_df))}

# T2b - route-stratified first_other_pco2 audit using canonical QA scope.
first_other_scope_df = (
    uom_scope_df
    if {"first_other_src", "first_other_pco2"}.issubset(uom_scope_df.columns)
)

```

```

        else ed_df
    )
first_other_scope_frame = "uom_scope_df" if first_other_scope_df is
    ↵ uom_scope_df else "ed_df"
first_other_pco2_audit = build_first_other_pco2_audit(first_other_scope_df)
first_other_pco2_audit_path = prior_runs_dir / f"{out_date}"
    ↵ first_other_pco2_audit.csv"
first_other_pco2_audit.to_csv(first_other_pco2_audit_path, index=False)
print("Wrote:", first_other_pco2_audit_path)
print("first_other_pco2 audit scope:", first_other_scope_frame)

print("Expected structural null fields:", expected_structural_null_fields)
print("Unexpected full-null fields (top 20):",
    ↵ unexpected_full_null_fields[:20])
print("UOM expectation scope:", uom_expectation_scope)
print("UOM expectation checks:", json.dumps(uom_expectation_checks,
    ↵ indent=2))

# T3 - lab capture completeness.
pct_any_6h = float(p06["ed_stay_id"].nunique() /
    ↵ max(ed_df["ed_stay_id"].nunique(),1))
pct_any_24h = float(p24["ed_stay_id"].nunique() /
    ↵ max(ed_df["ed_stay_id"].nunique(),1))
print("% any gas panel 0-6h:", round(pct_any_6h*100,1))
print("% any gas panel 0-24h:", round(pct_any_24h*100,1))

if "gas_source_audit" not in globals():
    gas_source_audit = summarize_gas_source(panel if "panel" in globals()
    ↵ else pd.DataFrame())

gas_overlap_summary = build_gas_source_overlap_summary(ed_df)
gas_overlap_summary_path = prior_runs_dir / f"{out_date}"
    ↵ gas_source_overlap_summary.csv"
gas_overlap_summary.to_csv(gas_overlap_summary_path, index=False)
print("Wrote:", gas_overlap_summary_path)

gas_source_other_rate = float(gas_source_audit.get("source_rates",
    ↵ {}).get("other", 0.0))
source_unknown_rate = float(gas_source_audit.get("source_rates",
    ↵ {}).get("unknown", 0.0))
print("% source other (panel-level):", round(gas_source_other_rate * 100, 1))
print("% source unknown (panel-level):", round(source_unknown_rate * 100, 1))

```

```

# Blood-gas source distribution audit (validated LAB + quarantined POC OTHER
    ↵ monitoring).
pco2_source_distribution_rows: list[dict[str, Any]] = []
if "panel" in globals() and len(panel) > 0 and "pco2" in panel.columns:
    panel_pco2 = panel.copy()
    panel_pco2["pco2"] = pd.to_numeric(panel_pco2["pco2"], errors="coerce")
    panel_pco2 = panel_pco2.loc[panel_pco2["pco2"].notna()].copy()
    if len(panel_pco2) > 0:
        for source_name, subset in panel_pco2.groupby("source",
            ↵ dropna=False):
            values = pd.to_numeric(subset["pco2"], errors="coerce").dropna()
            if values.empty:
                continue
            pco2_source_distribution_rows.append(
                {
                    "source_branch": "LAB",
                    "source": str(source_name),
                    "n": int(values.shape[0]),
                    "median": float(values.median()),
                    "p95": float(values.quantile(0.95)),
                    "max": float(values.max()),
                    "pct_ge_50": float((values >= 50).mean()),
                }
            )
    )

if "poc_other_paco2" in ed_df.columns:
    poc_other_values = pd.to_numeric(ed_df["poc_other_paco2"],
        ↵ errors="coerce").dropna()
    if len(poc_other_values) > 0:
        pco2_source_distribution_rows.append(
            {
                "source_branch": "POC_QUARANTINED",
                "source": "other",
                "n": int(poc_other_values.shape[0]),
                "median": float(poc_other_values.median()),
                "p95": float(poc_other_values.quantile(0.95)),
                "max": float(poc_other_values.max()),
                "pct_ge_50": float((poc_other_values >= 50).mean()),
            }
        )
    )

pco2_source_distribution_audit = pd.DataFrame(pco2_source_distribution_rows)

```

```

pco2_source_distribution_audit_path = prior_runs_dir / f"{out_date}"
    ↵ pco2_source_distribution_audit.csv"
pco2_source_distribution_audit.to_csv(pco2_source_distribution_audit_path,
    ↵ index=False)
print("Wrote:", pco2_source_distribution_audit_path)

poc_other_values_all = (
    pd.to_numeric(ed_df["poc_other_paco2"], errors="coerce")
    if "poc_other_paco2" in ed_df.columns
    else pd.Series(dtype="float64")
)
poc_other_quarantine_series = (
    pd.to_numeric(ed_df["poc_other_quarantined_hypercap_threshold"],
    ↵ errors="coerce")
    if "poc_other_quarantined_hypercap_threshold" in ed_df.columns
    else pd.Series(0, index=ed_df.index, dtype="float64")
)
other_threshold_series = (
    pd.to_numeric(ed_df["other_hypercap_threshold"], errors="coerce")
    if "other_hypercap_threshold" in ed_df.columns
    else pd.Series(0, index=ed_df.index, dtype="float64")
)
first_other_detail_series = (
    ed_df["first_other_src_detail"].astype("string")
    if "first_other_src_detail" in ed_df.columns
    else pd.Series("missing", index=ed_df.index, dtype="string")
)
poc_other_candidate_n = int(poc_other_values_all.notna().sum())
poc_other_hypercap_n = int((poc_other_values_all >= 50).sum())
poc_other_quarantined_hadm_n = int(
    poc_other_quarantine_series.fillna(0).astype(int).sum()
)
poc_other_leak_into_other_threshold_n = int(
    (
        other_threshold_series.fillna(0).astype(int).eq(1)
        & first_other_detail_series.str.lower().eq("poc_quarantined")
    ).sum()
)
other_route_quarantine_audit = pd.DataFrame(
    [
        {
            "poc_other_candidate_n": poc_other_candidate_n,
            "poc_other_ge50_n": poc_other_hypercap_n,

```

```

    "poc_other_ge50_rate": float(poc_other_hypercap_n /
        ↪ poc_other_candidate_n)
    if poc_other_candidate_n
        else 0.0,
    "poc_other_quarantined_hadm_n": poc_other_quarantined_hadm_n,
    "other_hypercap_threshold_hadm_n":
        ↪ int(other_threshold_series.fillna(0).astype(int).sum()),
    "poc_other_leak_into_other_threshold_n":
        ↪ poc_other_leak_into_other_threshold_n,
    }
]
)
other_route_quarantine_audit_path = prior_runs_dir / f"{out_date}"
    ↪ other_route_quarantine_audit.csv"
other_route_quarantine_audit.to_csv(other_route_quarantine_audit_path,
    ↪ index=False)
print("Wrote:", other_route_quarantine_audit_path)

first_gas_anchor_audit = pd.DataFrame(
    [
        {
            "row_count": int(len(ed_df)),
            "first_gas_time_nonnull_n":
                ↪ int(pd.to_datetime(ed_df.get("first_gas_time"),
                    ↪ errors="coerce").notna().sum())
            if "first_gas_time" in ed_df.columns
            else 0,
            "first_gas_anchor_has_pco2_n":
                ↪ int(ed_df.get("first_gas_anchor_has_pco2", pd.Series(False,
                    ↪ index=ed_df.index)).fillna(False).astype(bool).sum()),
            "first_gas_anchor_source_validated_n":
                ↪ int(ed_df.get("first_gas_anchor_source_validated",
                    ↪ pd.Series(False,
                    ↪ index=ed_df.index)).fillna(False).astype(bool).sum()),
        }
    ]
)
first_gas_anchor_audit["first_gas_without_pco2_anchor_n"] = (
    first_gas_anchor_audit["first_gas_time_nonnull_n"] -
    ↪ first_gas_anchor_audit["first_gas_anchor_has_pco2_n"]
)
first_gas_anchor_audit_path = prior_runs_dir / f"{out_date}"
    ↪ first_gas_anchor_audit.csv"

```

```

first_gas_anchor_audit.to_csv(first_gas_anchor_audit_path, index=False)
print("Wrote:", first_gas_anchor_audit_path)

hadm_rows = int(len(df)) if "df" in globals() else None
hadm_cc_rows = None
if "df" in globals() and "ed_triage_cc" in df.columns:
    _hadm_cc_mask = df["ed_triage_cc"].notna() &
    ↵ (df["ed_triage_cc"].astype(str).str.strip() != "")
    hadm_cc_rows = int(_hadm_cc_mask.sum())

uom_failed_checks = []
for uom_column, details in uom_expectation_checks.get("paco2_uom_checks",
    ↵ {}).items():
    if details.get("present") and not details.get("passes", True):
        uom_failed_checks.append(uom_column)

qa_warnings = []
qa_infos = []
if unexpected_full_null_fields:
    qa_warnings.append("unexpected_full_null_fields_detected")
if gas_source_audit.get("all_other_or_unknown", False):
    qa_warnings.append("gas_source_all_other_or_unknown")
if omr_pre_window_rows == 0 and omr_post_window_rows > 0:
    qa_infos.append("omr_pre_window_empty_post_window_fallback_used")
if omr_pre_window_rows == 0 and omr_post_window_rows == 0:
    qa_infos.append("omr_pre_post_window_empty_expected_sparse")
if uom_failed_checks:
    qa_warnings.append("uom_checks_failed")
if "status" in first_other_pco2_audit.columns and
    ↵ (first_other_pco2_audit["status"] == "missing_columns").any():
    qa_warnings.append("first_other_pco2_audit_missing_columns")
if "poc_other_leak_into_other_threshold_n" in
    ↵ other_route_quarantine_audit.columns:
    if int(other_route_quarantine_audit[
        ↵ "poc_other_leak_into_other_threshold_n"].iloc[0]) > 0:
        qa_warnings.append("poc_other_quarantine_leakage_detected")
if "first_gas_without_pco2_anchor_n" in first_gas_anchor_audit.columns:
    if int(first_gas_anchor_audit["first_gas_without_pco2_anchor_n"].iloc[0])
        ↵ > 0:
        qa_warnings.append("first_gas_anchor_without_pco2_detected")
qa_other_warn_threshold = float(os.getenv("COHORT_WARN_OTHER_RATE", "0.50"))
qa_other_fail_raw = os.getenv("COHORT_FAIL_OTHER_RATE", "").strip()

```

```

qa_other_fail_threshold = float(qa_other_fail_raw) if qa_other_fail_raw else
    None
if qa_other_fail_threshold is not None and gas_source_other_rate >=
    qa_other_fail_threshold:
    qa_warnings.append("gas_source_other_rate_fail_threshold_exceeded")
elif gas_source_other_rate >= qa_other_warn_threshold:
    qa_infos.append("gas_source_other_rate_high")

first_other_pco2_audit_records =
    json.loads(first_other_pco2_audit.to_json(orient="records"))
vitals_outlier_audit_records =
    json.loads(vitals_outlier_audit.to_json(orient="records"))
outliers_audit_records = json.loads(outliers_audit.to_json(orient="records"))
ed_vitals_distribution_summary_records =
    json.loads(ed_vitals_distribution_summary.to_json(orient="records"))
ed_vitals_model_delta_records =
    json.loads(ed_vitals_model_delta.to_json(orient="records"))
gas_overlap_summary_records =
    json.loads(gas_overlap_summary.to_json(orient="records")) if
    "gas_overlap_summary" in globals() else []
pco2_source_distribution_audit_records = json.loads(
    pco2_source_distribution_audit.to_json(orient="records"))
)
other_route_quarantine_audit_records = json.loads(
    other_route_quarantine_audit.to_json(orient="records"))
)
first_gas_anchor_audit_records =
    json.loads(first_gas_anchor_audit.to_json(orient="records"))
pco2_itemid_qc_audit_records = (
    json.loads(pco2_itemid_qc_audit.to_json(orient="records"))
    if "pco2_itemid_qc_audit" in globals() and not pco2_itemid_qc_audit.empty
    else []
)
timing_integrity_audit_records = (
    json.loads(timing_integrity_audit.to_json(orient="records"))
    if "timing_integrity_audit" in globals()
    else []
)
ventilation_timing_audit_records = (
    json.loads(ventilation_timing_audit.to_json(orient="records"))
    if "ventilation_timing_audit" in globals()
    else []
)

```

```

anthropometric_cleaning_audit_records = (
    json.loads(anthropometric_cleaning_audit.to_json(orient="records"))
    if "anthropometric_cleaning_audit" in globals()
    else []
)
blood_gas_itemid_manifest_audit_records = (
    json.loads(blood_gas_itemid_manifest_audit.to_json(orient="records"))
    if "blood_gas_itemid_manifest_audit" in globals() and not
        blood_gas_itemid_manifest_audit.empty
    else []
)
qa_status = "warning" if qa_warnings else "pass"

# T4 - QA summary artifact.
qa_summary = {
    "ed_spine_rows": int(len(ed_df)),
    "ed_rows": int(len(ed_df)),
    "ed_unique": int(ed_df["ed_stay_id"].nunique()),
    "hadm_rows": hadm_rows,
    "hadm_cc_rows": hadm_cc_rows,
    "icu_link_rate": float(ed_df["icu_intime_first"].notna().mean()) if
        "icu_intime_first" in ed_df.columns else None,
    "pct_any_gas_0_6h": pct_any_6h,
    "pct_any_gas_0_24h": pct_any_24h,
    "gas_source_other_rate": gas_source_other_rate,
    "gas_source_other_rate_warn_threshold": qa_other_warn_threshold,
    "gas_source_other_rate_fail_threshold": qa_other_fail_threshold,
    "source_unknown_rate": source_unknown_rate,
    "gas_source_audit": gas_source_audit,
    "gas_source_overlap_summary": gas_overlap_summary_records,
    "blood_gas_manifest": {
        "version": BLOOD_GAS_MANIFEST.get("version"),
        "path": BLOOD_GAS_MANIFEST.get("path"),
    },
    "blood_gas_itemid_manifest_audit": [
        blood_gas_itemid_manifest_audit_records,
    ],
    "pco2_source_distribution_audit": pco2_source_distribution_audit_records,
    "other_route_quarantine_audit": other_route_quarantine_audit_records,
    "first_gas_anchor_audit": first_gas_anchor_audit_records,
    "pco2_itemid_qc_audit": pco2_itemid_qc_audit_records,
    "poc_inclusion_enabled": bool(poc_inclusion_enabled),
    "poc_inclusion_reason": str(poc_inclusion_reason),
}

```

```

"omr_diagnostics": omr_diagnostics if "omr_diagnostics" in globals() else
    None,
"omr_window_diagnostics": omr_window_diagnostics if
    "omr_window_diagnostics" in globals() else None,
"anthro_timing_tier_counts": omr_diagnostics.get("selected_tier_counts",
    {}) if "omr_diagnostics" in globals() else {},
"anthro_timing_tier_rates": omr_diagnostics.get("selected_tier_rates",
    {}) if "omr_diagnostics" in globals() else {},
"anthro_coverage_audit": anthro_coverage_audit if "anthro_coverage_audit"
    in globals() else None,
"anthropometric_cleaning_audit": anthropometric_cleaning_audit_records,
"charted_anthro_diagnostics": charted_anthro_diagnostics if
    "charted_anthro_diagnostics" in globals() else None,
"timing_integrity_audit": timing_integrity_audit_records,
"ventilation_timing_audit": ventilation_timing_audit_records,
"unexpected_full_null_fields": unexpected_full_null_fields,
"expected_structural_null_fields": expected_structural_null_fields,
"uom_expectation_scope": uom_expectation_scope,
"uom_expectation_checks": uom_expectation_checks,
"first_other_pco2_audit_scope": first_other_scope_frame,
"first_other_pco2_audit": first_other_pco2_audit_records,
"vitals_outlier_audit": vitals_outlier_audit_records,
"ed_vitals_distribution_summary": ed_vitals_distribution_summary_records,
"ed_vitals_model_delta": ed_vitals_model_delta_records,
"outliers_audit": outliers_audit_records,
"ed_vitals_audit_paths": {
    "distribution_summary": str(ed_vitals_distribution_summary_path),
    "extreme_examples": str(ed_vitals_extreme_examples_path),
    "model_delta": str(ed_vitals_model_delta_path),
},
"blood_gas_audit_paths": {
    "manifest_itemids": str(blood_gas_itemid_manifest_audit_path)
        if "blood_gas_itemid_manifest_audit_path" in globals()
        else None,
    "pco2_itemid_qc": str(pco2_itemid_qc_audit_path)
        if "pco2_itemid_qc_audit_path" in globals()
        else None,
    "pco2_source_distribution": str(pco2_source_distribution_audit_path),
    "other_route_quarantine": str(other_route_quarantine_audit_path),
    "first_gas_anchor": str(first_gas_anchor_audit_path),
},
"timing_integrity_audit_path": str(timing_integrity_audit_path)
if "timing_integrity_audit_path" in globals()

```

```

    else None,
    "ventilation_timing_audit_path": str(ventilation_timing_audit_path)
    if "ventilation_timing_audit_path" in globals()
    else None,
    "anthropometric_cleaning_audit_path":
        ↵ str(anthropometric_cleaning_audit_path)
    if "anthropometric_cleaning_audit_path" in globals()
    else None,
    "qa_status": qa_status,
    "qa_warnings": qa_warnings,
    "qa_infos": qa_infos,
    "run_metadata": RUN_METADATA if "RUN_METADATA" in globals() else None,
}
qa_path = WORK_DIR / "qa_summary.json"
qa_path.write_text(json.dumps(qa_summary, indent=2))
print("Wrote:", qa_path)

```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 vitals_outlier_audit.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 outliers_audit.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 ed_vitals_distribution_summary.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 ed_vitals_extreme_examples.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 ed_vitals_model_delta.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 timing_integrity_audit.csv
 Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 ventilation_timing_audit.csv
 ED vitals cleaning summary (raw->clean removals):

	vital_name	raw_n	clean_n	raw_removed_n	raw_removed_pct	converted_celsius_like_n
7	first_dbp	9161	9152	9	0.000982	0
9	first_o2sat	8920	8919	1	0.000112	0
3	first_pain	7077	7077	0	0.000000	0
6	first_sbp	9161	9155	6	0.000655	0
1	first_temp	6911	6911	0	0.000000	55
5	triage_dbp	8738	8727	11	0.001259	0
8	triage_o2sat	8680	8679	1	0.000115	0

vital_name	raw_n	clean_n	raw_removed_n	raw_removed_pct	converted_celsius_like_n
2 triage_pain	8151	7595	556	0.068212	0
4 triage_sbp	8777	8771	6	0.000684	0
0 triage_temp	8343	8342	1	0.000120	20

	field	missing_n	missing_pct	expectation
0	abg_before_imv	11786	1.000000	missing_column
61	first_other_time	11786	1.000000	missing_column
89	icu_stay_id	11786	1.000000	missing_column
112	vbg_before_imv	11786	1.000000	missing_column
5	anthro_chartdate	11767	0.998388	conditional_sparse
6	anthro_days_offset	11767	0.998388	conditional_sparse
59	first_hco3	11651	0.988546	conditional_sparse
76	hco3_band	11651	0.988546	conditional_sparse
31	ed_first_rhythm	11130	0.944341	conditional_sparse
18	deathtime	9675	0.820889	conditional_sparse
100	max_pco2_0_6h	8717	0.739606	conditional_sparse
102	min_ph_0_6h	8676	0.736128	conditional_sparse
25	dt_first_niv_hours_model	6728	0.570847	conditional_sparse
24	dt_first_niv_hours	6715	0.569744	conditional_sparse
66	flag_chf	6522	0.553368	conditional_sparse
67	flag_copd	6522	0.553368	conditional_sparse
68	flag_neuromuscular	6522	0.553368	conditional_sparse
69	flag_opioid_substance	6522	0.553368	conditional_sparse
70	flag.osa_ohs	6522	0.553368	conditional_sparse
72	flag_pneumonia	6522	0.553368	conditional_sparse
60	first_lactate	5561	0.471831	conditional_sparse
93	lactate_band	5561	0.471831	conditional_sparse
23	dt_first_imv_hours_model	5243	0.444850	conditional_sparse
22	dt_first_imv_hours	5215	0.442474	conditional_sparse
34	ed_first_temp	4875	0.413626	conditional_sparse
15	bmi_closest_pre_ed_model	4641	0.393772	conditional_sparse
78	height_closest_pre_ed_model	4558	0.386730	conditional_sparse
14	bmi_closest_pre_ed	4529	0.384269	conditional_sparse
77	height_closest_pre_ed	4525	0.383930	conditional_sparse
57	first_gas_specimen_type	4265	0.361870	conditional_sparse

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects CC-NLP/MIMIC tabular data/prior runs/2026-02-20 first_other_pco2_audit.csv

first_other_pco2 audit scope: uom_scope_df

Expected structural null fields: []

Unexpected full-null fields (top 20): []

```

UOM expectation scope: {'frame': 'final_cc_preview', 'row_count': 11769}
UOM expectation checks: {
    "expected_structural_null_fields": [
        "poc_abg_ph_uom",
        "poc_vbg_ph_uom",
        "poc_other_ph_uom"
    ],
    "structural_null_checks": {
        "poc_abg_ph_uom": {
            "present": true,
            "missing_n": 11769,
            "missing_pct": 1.0,
            "all_null": true
        },
        "poc_vbg_ph_uom": {
            "present": true,
            "missing_n": 11769,
            "missing_pct": 1.0,
            "all_null": true
        },
        "poc_other_ph_uom": {
            "present": true,
            "missing_n": 11769,
            "missing_pct": 1.0,
            "all_null": true
        }
    },
    "paco2_uom_checks": {
        "poc_abg_paco2_uom": {
            "present": true,
            "paired_value_column": "poc_abg_paco2",
            "value_rows": 0,
            "missing_uom_when_value_present": 0,
            "non_mmhg_uom_when_value_present": 0,
            "passes": true
        },
        "poc_vbg_paco2_uom": {
            "present": true,
            "paired_value_column": "poc_vbg_paco2",
            "value_rows": 0,
            "missing_uom_when_value_present": 0,
            "non_mmhg_uom_when_value_present": 0,
            "passes": true
        }
    }
}

```

```

},
"poc_other_paco2_uom": {
    "present": true,
    "paired_value_column": "poc_other_paco2",
    "value_rows": 0,
    "missing_uom_when_value_present": 0,
    "non_mmhg_uom_when_value_present": 0,
    "passes": true
}
}
}

% any gas panel 0-6h: 41.9
% any gas panel 0-24h: 95.9
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 gas_source_overlap_summary.csv
% source other (panel-level): 66.2
% source unknown (panel-level): 0.0
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 pco2_source_distribution_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 other_route_quarantine_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 first_gas_anchor_audit.csv
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/qa_summary.json

```

3.2 Outputs (ED-stay cohort + long tables)

Rationale: Persist ED-stay analytic datasets and supporting long tables.

```

# Purpose: Write finalized cohort outputs and derivative tables for
#           downstream analysis workflows.

# Save outputs (dated parquet artifacts archived under prior runs)
from datetime import datetime

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")

cohort_path = prior_runs_dir / f"cohort_ed_stay_{timestamp}.parquet"
ed_vitals_path = prior_runs_dir / f"ed_vitals_long_{timestamp}.parquet"

```

```

labs_path = prior_runs_dir / f"labs_long_{timestamp}.parquet"
gas_panels_path = prior_runs_dir / f"gas_panels_{timestamp}.parquet"

ed_df.to_parquet(cohort_path, index=False)
ed_vitals_long.to_parquet(ed_vitals_path, index=False)
labs_long.to_parquet(labs_path, index=False)
panel.to_parquet(gas_panels_path, index=False)

print("Wrote:", cohort_path)
print("Wrote:", ed_vitals_path)
print("Wrote:", labs_path)
print("Wrote:", gas_panels_path)

```

```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/cohort_ed_stay_20260220_081357.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/ed_vitals_long_20260220_081357.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/labs_long_20260220_081357.parquet
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/gas_panels_20260220_081357.parquet

```

```

# Purpose: Report missingness for key variables to support transparent
↪ data-quality reporting.

```

```

# Also export to Excel (tabular dataset) - optional archive output
from datetime import datetime

```

```

WRITE_ARCHIVE_XLSX_EXPORTS = os.getenv("WRITE_ARCHIVE_XLSX_EXPORTS", "0") ==
↪ "1"
if WRITE_ARCHIVE_XLSX_EXPORTS:
    prior_runs_dir = DATA_DIR / "prior runs"
    prior_runs_dir.mkdir(parents=True, exist_ok=True)
    timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')

    xlsx_path = prior_runs_dir /
↪ f"mimic_hypcap_EXT_EDstay_bq_gas_{timestamp}.xlsx"

    with pd.ExcelWriter(xlsx_path, engine="openpyxl") as xw:
        ed_df.to_excel(xw, sheet_name="cohort_ed_stay", index=False)
        # Optional: include QA tables if present
        try:

```

```

        miss.to_excel(xw, sheet_name="missingness", index=False)
    except Exception:
        pass

    print("Saved:", xlsx_path)
else:
    print("Skipping ED-stay archive workbook export (set
        WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")

Skipping ED-stay archive workbook export (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

# Purpose: Capture ED presentation context (chief complaint, acuity, and
    ↳ early vitals) at the ED-stay level.

# Final Excel outputs requested: all encounters + ED chief-complaint only
from datetime import datetime

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
timestamp = datetime.now().strftime('%Y%m%d_%H%M%S')

# 1) All encounters meeting inclusion criteria (admission-level), with ED
    ↳ linkage flags
ed_link = (
    ed_df.groupby("hadm_id", as_index=False)
        .agg(ed_stay_id_first=("ed_stay_id", "first"),
            ↳ n_ed_stays=("ed_stay_id", "nunique"))
)
all_encounters = df.merge(ed_link, on="hadm_id", how="left")
all_encounters["has_ed_encounter"] =
    ↳ all_encounters["n_ed_stays"].fillna(0).astype(int).gt(0).astype(int)
all_encounters["encounter_source"] =
    ↳ all_encounters["has_ed_encounter"].map({1: "ED-linked", 0:
        ↳ "Inpatient-only"})

if WRITE_ARCHIVE_XLSX_EXPORTS:
    all_path = prior_runs_dir /
    ↳ f"mimic_hypercap_EXT_all_encounters_bq_{timestamp}.xlsx"
    with pd.ExcelWriter(all_path, engine="openpyxl") as xw:
        all_encounters.to_excel(xw, sheet_name="all_encounters", index=False)
    print("Saved:", all_path)

```

```

# 2) ED chief-complaint-only (ED-stay level)
if "ed_triage_cc" not in ed_df.columns:
    raise KeyError("Column 'ed_triage_cc' not found in ed_df. Ensure ED
    triage merge ran.")

mask_cc = ed_df["ed_triage_cc"].notna() &
    (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
ed_cc_only = ed_df.loc[mask_cc].copy()

if WRITE_ARCHIVE_XLSX_EXPORTS:
    cc_path = prior_runs_dir /
    f"mimic_hypcap_EXT_EDcc_only_edstay_bq_{timestamp}.xlsx"
    with pd.ExcelWriter(cc_path, engine="openpyxl") as xw:
        ed_cc_only.to_excel(xw, sheet_name="ed_cc_only", index=False)
        print("Saved:", cc_path)
else:
    print("Skipping all-encounters/ED-CC archive exports (set
    WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).")

Skipping all-encounters/ED-CC archive exports (set WRITE_ARCHIVE_XLSX_EXPORTS=1 to enable).

# Purpose: Capture ED presentation context (chief complaint, acuity, and
# early vitals) at the ED-stay level.

# Final CC output (cohort-only; excludes NLP-derived columns)
from datetime import datetime

# Ensure cc-only frames exist (build from df if upstream optional export
# cells were skipped)
if "df_cc" not in globals():
    if "df" not in globals():
        raise NameError("df not found. Run cohort assembly cells first.")
    if "ed_triage_cc" not in df.columns:
        raise KeyError("ed_triage_cc missing in df; cannot build df_cc.")
    _mask_df_cc = df["ed_triage_cc"].notna() &
    (df["ed_triage_cc"].astype(str).str.strip() != "")
    df_cc = df.loc[_mask_df_cc].copy()

if "ed_cc_only" not in globals():
    if "ed_df" not in globals():
        raise NameError("ed_df not found. Run the ED-stay build cells
        first.")

```

```

if "ed_triage_cc" not in ed_df.columns:
    raise KeyError("ed_triage_cc missing in ed_df; cannot build
                   ed_cc_only.")
    _mask_cc = ed_df["ed_triage_cc"].notna() &
    (ed_df["ed_triage_cc"].astype(str).str.strip() != "")
    ed_cc_only = ed_df.loc[_mask_cc].copy()

# Add ED-stay columns (dedup to earliest ED stay per hadm_id)
ed_tmp = ed_cc_only.copy()
if "ed_intime" in ed_tmp.columns:
    ed_tmp = ed_tmp.sort_values(["hadm_id", "ed_intime", "ed_stay_id"],
                                na_position="last")
else:
    ed_tmp = ed_tmp.sort_values(["hadm_id", "ed_stay_id"],
                                na_position="last")
ed_first = ed_tmp.drop_duplicates("hadm_id", keep="first")

add_ed_cols = [c for c in ed_first.columns if c not in df_cc.columns]

final_cc = df_cc.merge(ed_first[["hadm_id"] + add_ed_cols], on="hadm_id",
                       how="left")

prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_date = datetime.now().strftime("%Y-%m-%d")

cohort_contract = validate_cohort_contract(final_cc)
cohort_contract_report = {
    "generated_utc": datetime.utcnow().isoformat() + "Z",
    "status": cohort_contract["status"],
    "contracts": {"cohort": cohort_contract},
    "findings": cohort_contract["findings"],
}
cohort_contract_paths = write_contract_report(cohort_contract_report,
                                              work_dir=WORK_DIR)
print("Wrote:", cohort_contract_paths["contract_report_path"])
if cohort_contract_paths["failed_contract_path"].exists():
    print("Wrote:", cohort_contract_paths["failed_contract_path"])

contract_mode = os.getenv("PIPELINE_CONTRACT_MODE", "fail").strip().lower()
if contract_mode not in {"fail", "warn"}:
    raise ValueError("PIPELINE_CONTRACT_MODE must be one of {'fail',
                  'warn'}")

```

```

if cohort_contract["status"] == "fail" and contract_mode == "fail":
    raise ValueError(
        "Cohort contract checks failed. "
        f"See {cohort_contract_paths['failed_contract_path']}"
    )

dated_out_path = prior_runs_dir / f"{out_date} MIMICIV all with CC.xlsx"
canonical_out_path = DATA_DIR / "MIMICIV all with CC.xlsx"

# Write via temporary files then atomically replace targets to avoid
# zero-byte outputs on interruption.
def _atomic_write_excel(df_obj, target_path: Path) -> None:
    tmp_path = target_path.with_suffix(target_path.suffix + ".tmp")
    df_obj.to_excel(tmp_path, index=False)
    os.replace(tmp_path, target_path)

def _sha256_file(path: Path) -> str:
    hasher = hashlib.sha256()
    with path.open("rb") as file_handle:
        for chunk in iter(lambda: file_handle.read(1024 * 1024), b ""):
            hasher.update(chunk)
    return hasher.hexdigest()

_atomic_write_excel(final_cc, dated_out_path)
_atomic_write_excel(final_cc, canonical_out_path)
canonical_sha256 = _sha256_file(canonical_out_path)
dated_sha256 = _sha256_file(dated_out_path)
print(
    "EXPORT_SHAPE cohort",
    f"rows={len(final_cc)}:,}",
    f"cols={final_cc.shape[1]}:,}",
    f"canonical_sha256={canonical_sha256}",
)
cohort_manifest = collect_run_manifest(
    WORK_DIR,
    run_id=f"cohort_{datetime.now().strftime('%Y%m%d_%H%M%S')}",
)
cohort_manifest["stage"] = "cohort"
cohort_manifest["outputs"] = {
    "canonical_output_path": str(canonical_out_path),
    "canonical_output_sha256": canonical_sha256,
}

```

```

    "dated_output_path": str(dated_out_path),
    "dated_output_sha256": dated_sha256,
    "gas_source_audit_path": str(gas_source_audit_path) if
        "gas_source_audit_path" in globals() else None,
    "gas_overlap_summary_path": str(gas_overlap_summary_path) if
        "gas_overlap_summary_path" in globals() else None,
    "anthropometrics_coverage_audit_path": str(anthro_coverage_audit_path) if
        "anthro_coverage_audit_path" in globals() else None,
    "ed_vitals_distribution_summary_path":
        str(ed_vitals_distribution_summary_path) if
        "ed_vitals_distribution_summary_path" in globals() else None,
    "ed_vitals_extreme_examples_path": str(ed_vitals_extreme_examples_path)
        if "ed_vitals_extreme_examples_path" in globals() else None,
    "ed_vitals_model_delta_path": str(ed_vitals_model_delta_path) if
        "ed_vitals_model_delta_path" in globals() else None,
    "outliers_audit_path": str(outliers_audit_path) if "outliers_audit_path"
        in globals() else None,
}
cohort_manifest["gas_source"] = {
    "mode": gas_source_mode if "gas_source_mode" in globals() else None,
    "source_rates": gas_source_audit.get("source_rates", {}) if
        "gas_source_audit" in globals() else {},
    "tier_rates": gas_source_audit.get("tier_rates", {}) if
        "gas_source_audit" in globals() else {},
}
cohort_manifest["anthropometrics"] = (
    anthro_coverage_audit if "anthro_coverage_audit" in globals() else None
)
cohort_manifest_path = prior_runs_dir / f"{out_date}"
    cohort_run_manifest.json"
cohort_manifest_path.write_text(json.dumps(cohort_manifest, indent=2))

print("Saved:", dated_out_path)
print("Saved:", canonical_out_path)
print("Wrote:", cohort_manifest_path)
print("Base rows:", len(df_cc), "| Added from ED-stay:", len(add_ed_cols))
print("Total columns:", len(final_cc.columns))

```

Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project
 CC-NLP/debug/contracts/20260220_151358/contract_report.json
 EXPORT_SHAPE cohort rows=11,769 cols=287 canonical_sha256=7405dff1273105034fdfcd035c6547136f
 Saved: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Project
 CC-NLP/MIMIC tabular data/prior runs/2026-02-20 MIMICIV all with CC.xlsx

```

Saved: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/MIMICIV all with CC.xlsx
Wrote: /Users/blocke/Box Sync/Residency Personal Files/Scholarly Work/Locke Research Projects/CC-NLP/MIMIC tabular data/prior runs/2026-02-20 cohort_run_manifest.json
Base rows: 11769 | Added from ED-stay: 178
Total columns: 287

# Purpose: Extract and standardize first ABG/VBG physiology fields for
#           baseline characterization.

# Data dictionary (OSF-style) for final cohort output
from datetime import datetime
import numpy as np

# Choose target dataset for dictionary
if "final_cc" in globals():
    dd_target = final_cc
elif "df_cc" in globals():
    dd_target = df_cc
elif "df" in globals():
    dd_target = df
else:
    raise NameError("No cohort dataframe found for data dictionary (expected
                     final_cc/df_cc/df).")

# Helper functions

def _readable(name: str) -> str:
    return name.replace("_", " ").strip().title()

def _infer_units(col: str) -> str:
    if col.endswith("_hours"):
        return "hours"
    if col.endswith("_time") or col.endswith("_intime") or
       col.endswith("_outtime") or col.endswith("_date"):
        return "datetime"
    if col.startswith("flag_") or col.endswith("_flag") or
       col.endswith("_before_imv"):
        return "binary (0/1)"
    if col.endswith("_paco2") or col.endswith("_pco2"):
        return "mmHg"
    if col.endswith("_ph"):


```

```

        return "pH"
    if col.endswith("_lactate"):
        return "mmol/L"
    return ""

def _allowed_values(series):
    s = series.dropna()
    if s.empty:
        return ""
    if s.dtype.kind in "biu":
        vals = sorted(map(int, s.unique()))
        if len(vals) <= 15:
            return ", ".join(map(str, vals))
        return f"{min(vals)}-{max(vals)}"
    if s.dtype.kind in "f":
        return f"{np.nanmin(s):.3g}-{np.nanmax(s):.3g}"
    # object/categorical: use bounded sampling to keep dictionary generation
    # ↵ fast.
    sample_n = 5000
    s_str = s.astype(str)
    if len(s_str) > sample_n:
        s_sample = s_str.head(sample_n)
        sample_note = f"sampled {sample_n}/{len(s_str)} rows"
    else:
        s_sample = s_str
        sample_note = None
    unique_n = int(s_sample.nunique(dropna=True))
    top_vals = s_sample.value_counts(dropna=False).head(15).index.tolist()
    if unique_n > len(top_vals):
        tail = f" ... ({unique_n} unique"
        if sample_note:
            tail += f"; {sample_note}"
        tail += ")"
        return ", ".join(top_vals) + tail
    if sample_note:
        return ", ".join(top_vals) + f" ({sample_note})"
    return ", ".join(top_vals)

META = {
    # IDs
    "subject_id": {

```

```

    "label": "Subject identifier",
    "definition": "Identifier for a unique patient in MIMIC-IV.",
    "source": "mimiciv_hosp.patients",
    "derivation": "as recorded",
},
"hadm_id": {
    "label": "Hospital admission identifier",
    "definition": "Identifier for a hospital admission.",
    "source": "mimiciv_hosp.admissions",
    "derivation": "as recorded",
},
"ed_stay_id": {
    "label": "ED stay identifier",
    "definition": "Identifier for an ED stay.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
},
"ed_intime": {
    "label": "ED arrival time",
    "definition": "ED arrival time for this ED stay.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
    "units": "datetime",
},
"ed_outtime": {
    "label": "ED departure time",
    "definition": "ED departure time for this ED stay.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
    "units": "datetime",
},
"ed_intime_first": {
    "label": "First ED arrival time (admission)",
    "definition": "Earliest ED arrival time among ED stays linked to the
    ↳ admission.",
    "source": "mimiciv_ed.edstays",
    "derivation": "min(edstays.intime) per hadm_id",
    "units": "datetime",
},
# Demographics / triage
"ed_gender": {
    "label": "ED gender",
    "definition": "Gender as recorded in ED stay table.",

```

```

    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
},
"ed_race": {
    "label": "ED race",
    "definition": "Race as recorded in ED stay table.",
    "source": "mimiciv_ed.edstays",
    "derivation": "as recorded",
},
"hospt_race": {
    "label": "Hospital race",
    "definition": "Race as recorded in admissions table.",
    "source": "mimiciv_hosp.admissions",
    "derivation": "as recorded",
},
"ed_triage_cc": {
    "label": "ED chief complaint",
    "definition": "Chief complaint recorded at ED triage.",
    "source": "mimiciv_ed.triage",
    "derivation": "as recorded",
},
"ed_triage_acuity": {
    "label": "ED triage acuity",
    "definition": "Triage acuity level recorded at ED presentation.",
    "source": "mimiciv_ed.triage",
    "derivation": "as recorded",
},
# Outcomes
"hospital_expire_flag": {
    "label": "Hospital expire flag",
    "definition": "Hospital mortality indicator from admissions table.",
    "source": "mimiciv_hosp.admissions",
    "derivation": "as recorded",
    "units": "binary (0/1)",
},
"in_hospital_death": {
    "label": "In-hospital death",
    "definition": "Indicator for in-hospital death during admission.",
    "source": "mimiciv_hosp.admissions",
    "derivation": "(hospital_expire_flag == 1) OR (deathtime not null)",
    "units": "binary (0/1)",
},
# Ventilation

```

```

"imv_flag": {
    "label": "IMV flag",
    "definition": "Indicator of invasive mechanical ventilation during admission.",
    "source": "ICD procedures + ICU chartevents",
    "derivation": "max(IMV ICD flag, IMV chart flag)",
    "units": "binary (0/1)",
},
"niv_flag": {
    "label": "NIV flag",
    "definition": "Indicator of noninvasive ventilation during admission.",
    "source": "ICD procedures + ICU chartevents",
    "derivation": "max(NIV ICD flag, NIV chart flag)",
    "units": "binary (0/1)",
},
"first_imv_time": {
    "label": "First IMV time",
    "definition": "Earliest charted time of IMV in ICU chartevents (if present).",
    "source": "mimiciv_icu.chartevents",
    "derivation": "min charttime among IMV charted events",
    "units": "datetime",
},
"first_niv_time": {
    "label": "First NIV time",
    "definition": "Earliest charted time of NIV in ICU chartevents (if present).",
    "source": "mimiciv_icu.chartevents",
    "derivation": "min charttime among NIV charted events",
    "units": "datetime",
},
"dt_first_imv_hours": {
    "label": "Hours to first IMV",
    "definition": "Hours from first ED arrival to first IMV time.",
    "source": "Derived",
    "derivation": "(first_imv_time - ed_intime_first) in hours",
    "units": "hours",
},
"dt_first_niv_hours": {
    "label": "Hours to first NIV",
    "definition": "Hours from first ED arrival to first NIV time.",
    "source": "Derived",
}

```

```

    "derivation": "(first_niv_time - ed_intime_first) in hours",
    "units": "hours",
},
"abg_before_imv": {
    "label": "ABG before IMV",
    "definition": "Indicator that first ABG time precedes first IMV
    ↵ time.",
    "source": "Derived",
    "derivation": "first_abg_time < first_imv_time",
    "units": "binary (0/1)",
},
"vbg_before_imv": {
    "label": "VBG before IMV",
    "definition": "Indicator that first VBG time precedes first IMV
    ↵ time.",
    "source": "Derived",
    "derivation": "first_vbg_time < first_imv_time",
    "units": "binary (0/1)",
},
# Gas/ABG/VBG
"abg_hypercap_threshold": {
    "label": "ABG hypercapnia threshold met",
    "definition": "Indicator that any arterial pCO2  45 mmHg in hosp/ICU
    ↵ pCO2 extraction.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "max(arterial pCO2  45) per hadm_id",
    "units": "binary (0/1)",
},
"vbg_hypercap_threshold": {
    "label": "VBG hypercapnia threshold met",
    "definition": "Indicator that any venous pCO2  50 mmHg in hosp/ICU
    ↵ pCO2 extraction.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "max(venous pCO2  50) per hadm_id",
    "units": "binary (0/1)",
},
"other_hypercap_threshold": {
    "label": "Other-source hypercapnia threshold met",
    "definition": "Indicator that any other/unspecified-source pCO2  50
    ↵ mmHg in hosp/ICU pCO2 extraction.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "max(other-source pCO2  50) per hadm_id",
    "units": "binary (0/1)",
}

```

```

},
"first_gas_time": {
    "label": "First gas panel time",
    "definition": "Earliest gas panel time in the ED 0-24h window.",
    "source": "mimiciv_hosp.labevents (gas panels)",
    "derivation": "min panel_time within ED 0-24h window",
    "units": "datetime",
},
"dt_first_qualifying_gas_hours": {
    "label": "Hours to first gas panel",
    "definition": "Hours from first ED arrival to first gas panel time.",
    "source": "Derived",
    "derivation": "(first_gas_time - ed_intime_first) in hours",
    "units": "hours",
},
"first_other_time": {
    "label": "First other-source gas time",
    "definition": "Earliest time of a pCO2 measurement with
    ↳ other/unspecified source classification.",
    "source": "mimiciv_hosp.labevents + mimiciv_icu.chartevents",
    "derivation": "min(other-source pCO2 charttime) across LAB and POC",
    "units": "datetime",
},
"presenting_hypercapnia": {
    "label": "Presenting (6h) gas timing",
    "definition": "Indicator that first gas panel occurred within 6h of
    ↳ first ED arrival.",
    "source": "Derived",
    "derivation": "dt_first_qualifying_gas_hours <= 6",
    "units": "binary (0/1)",
},
"late_hypercapnia": {
    "label": "Late (>6h) gas timing",
    "definition": "Indicator that first gas panel occurred >6h after
    ↳ first ED arrival.",
    "source": "Derived",
    "derivation": "dt_first_qualifying_gas_hours > 6",
    "units": "binary (0/1)",
},
"gas_source_other_rate": {
    "label": "Gas source other/unknown rate",
    "definition": "Proportion of gas panels classified as
    ↳ other/unspecified source per ED stay.",
}

```

```

    "source": "Derived from gas panel source inference",
    "derivation": "mean(source == 'other') per ed_stay_id",
    "units": "proportion",
},
# Comorbidities
"flag_copd": {
    "label": "COPD/emphysema flag",
    "definition": "Indicator for COPD/emphysema ICD codes in ED or
    ↳ hospital diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes J43/J44 (ED OR hospital)",
    "units": "binary (0/1)",
},
"flag_osa_ohs": {
    "label": "OSA/OHS flag",
    "definition": "Indicator for OSA/OHS ICD codes in ED or hospital
    ↳ diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes G473/E662 (ED OR hospital)",
    "units": "binary (0/1)",
},
"flag_chf": {
    "label": "CHF flag",
    "definition": "Indicator for CHF ICD codes in ED or hospital
    ↳ diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefix I50 (ED OR hospital)",
    "units": "binary (0/1)",
},
"flag_neuromuscular": {
    "label": "Neuromuscular flag",
    "definition": "Indicator for neuromuscular ICD codes in ED or
    ↳ hospital diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes G12/G70/G71 (ED OR hospital)",
    "units": "binary (0/1)",
},
"flag_opioid_substance": {
    "label": "Opioid/substance flag",
    "definition": "Indicator for opioid/substance ICD codes in ED or
    ↳ hospital diagnoses.",
    "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
    "derivation": "ICD prefixes F11/T40/F13 (ED OR hospital)",
}

```

```

        "units": "binary (0/1)",
    },
    "flag_pneumonia": {
        "label": "Pneumonia flag",
        "definition": "Indicator for pneumonia ICD codes in ED or hospital
        ↵ diagnoses.",
        "source": "mimiciv_hosp.diagnoses_icd + mimiciv_ed.diagnosis",
        "derivation": "ICD prefixes J12-J18 (ED OR hospital)",
        "units": "binary (0/1)",
    },
}

rows = []
for col in dd_target.columns:
    info = META.get(col, {})
    series = dd_target[col]
    rows.append({
        "variable": col,
        "label": info.get("label", _readable(col)),
        "units": info.get("units", _infer_units(col)),
        "allowed_values": info.get("allowed_values",
        ↵ _allowed_values(series)),
        "definition": info.get("definition", "UNCONFIRMED"),
        "synonyms": info.get("synonyms", ""),
        "description": info.get("description", ""),
        "source": info.get("source", "UNCONFIRMED"),
        "derivation": info.get("derivation", "UNCONFIRMED"),
        "dtype": str(series.dtype),
        "example_value": series.dropna().iloc[0] if series.dropna().shape[0]
        ↵ else None,
    })
)

data_dict = pd.DataFrame(rows)

out_date = datetime.now().strftime("%Y-%m-%d")
prior_runs_dir = DATA_DIR / "prior runs"
prior_runs_dir.mkdir(parents=True, exist_ok=True)
out_xlsx = prior_runs_dir / f"{out_date} MIMICIV all with
    ↵ CC_data_dictionary.xlsx"
out_csv = prior_runs_dir / f"{out_date} MIMICIV all with
    ↵ CC_data_dictionary.csv"

data_dict.to_excel(out_xlsx, index=False)

```

```
data_dict.to_csv(out_csv, index=False)

print("Saved data dictionary:", out_xlsx)
print("Saved data dictionary:", out_csv)
print("Rows:", len(data_dict))
```